



Lecture 2: 3D Representations

Li Yi

2025.02.27

Homework 1 is coming!

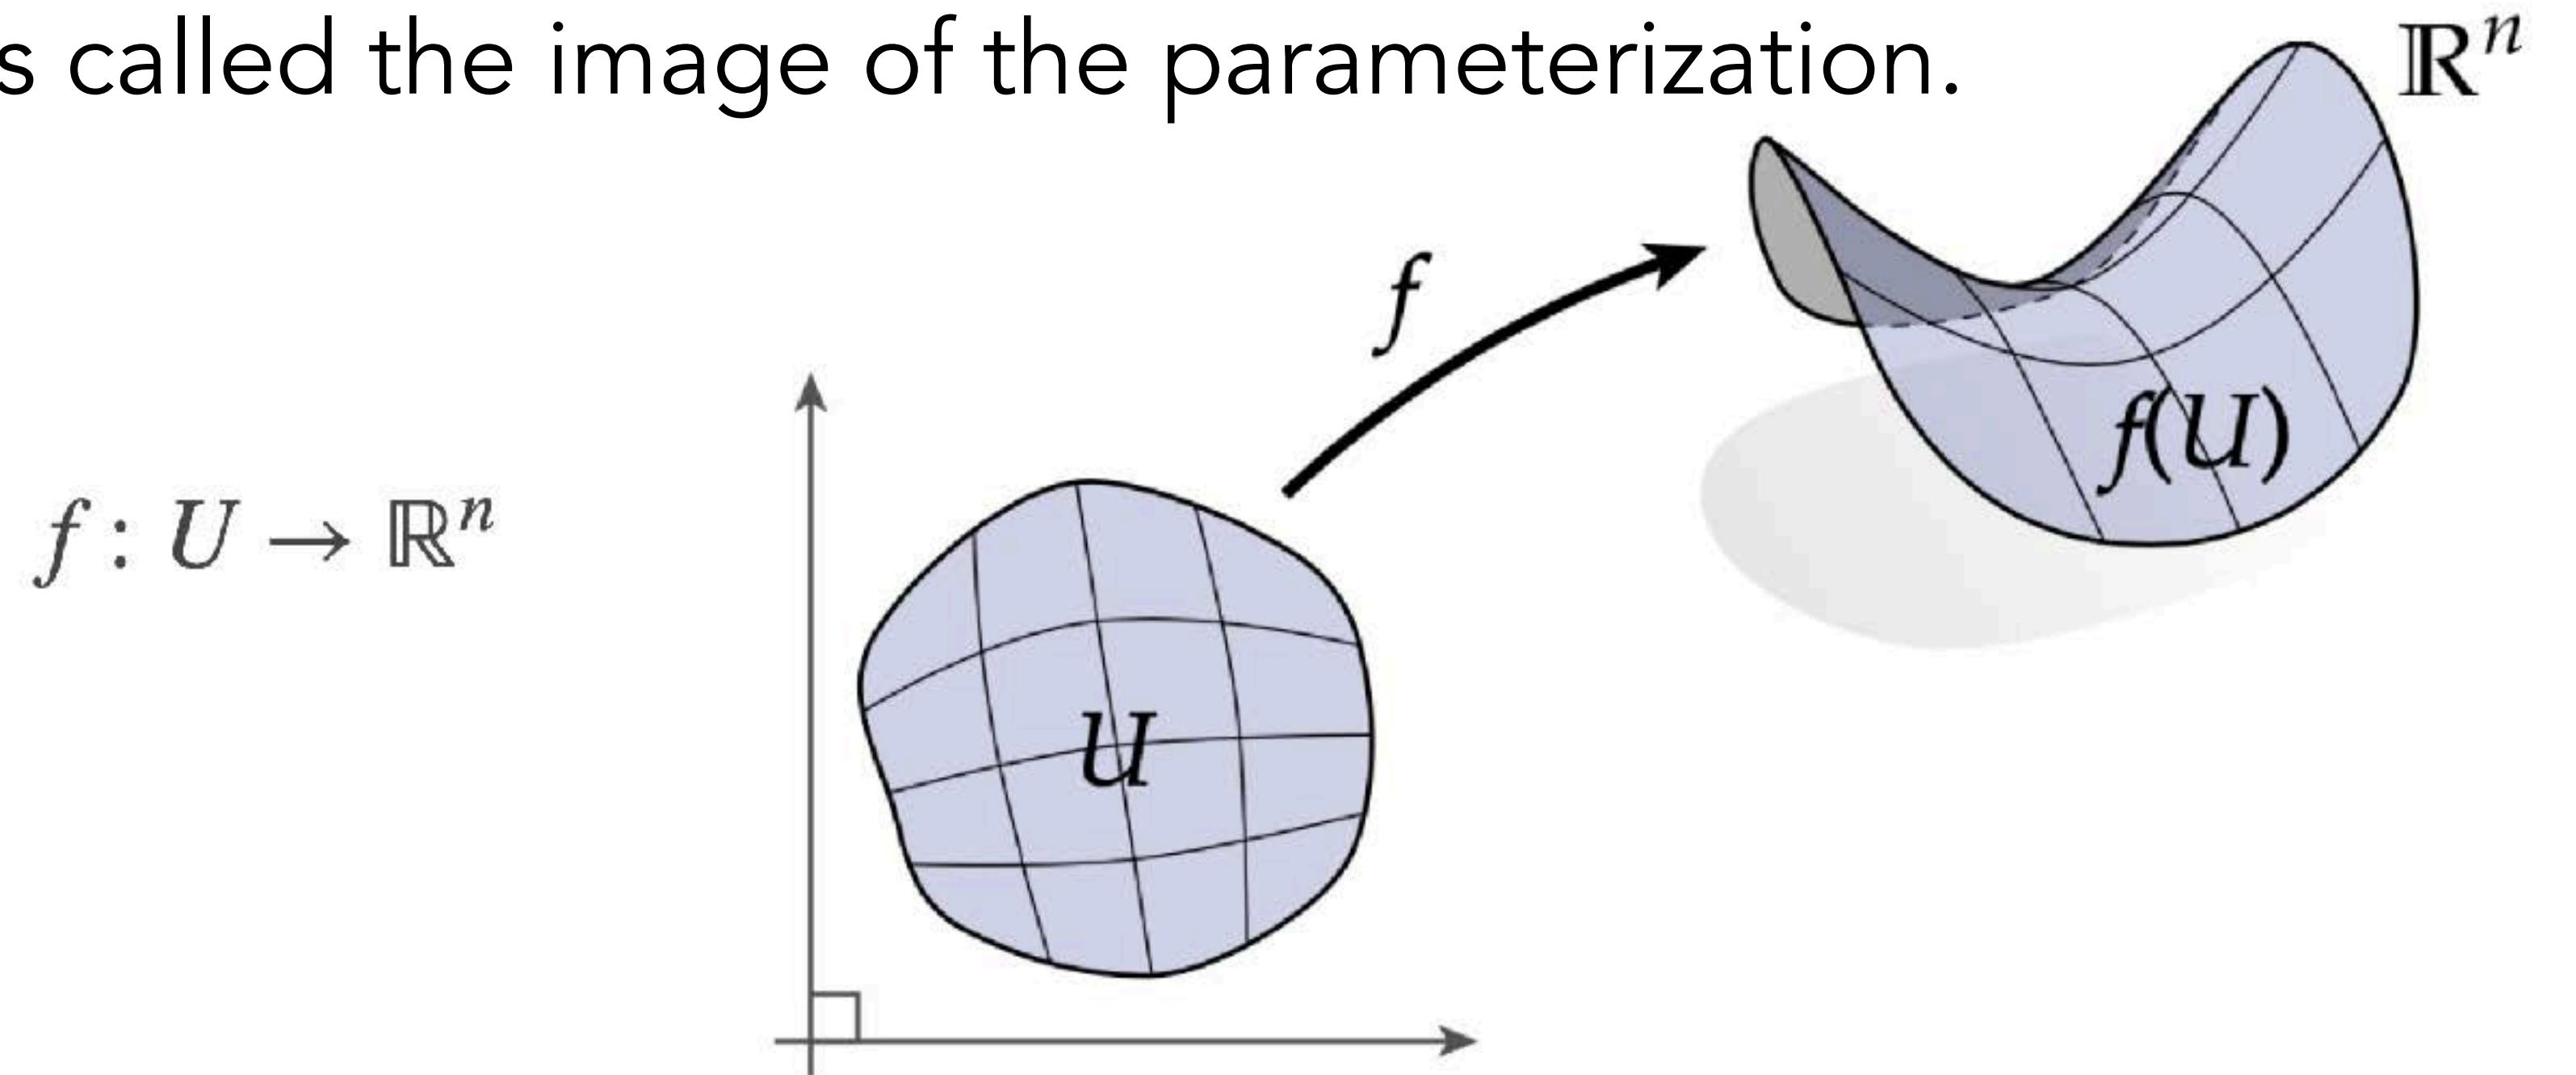
- The first assignment is going to be released this week!

Recap

- We can use parametrized form to represent smooth curves and surfaces
- Such parametrized form allows us to define local properties such as tangent, normal and curvature
- We usually require “smoothness” for easier analysis.

Previously: parametric surfaces

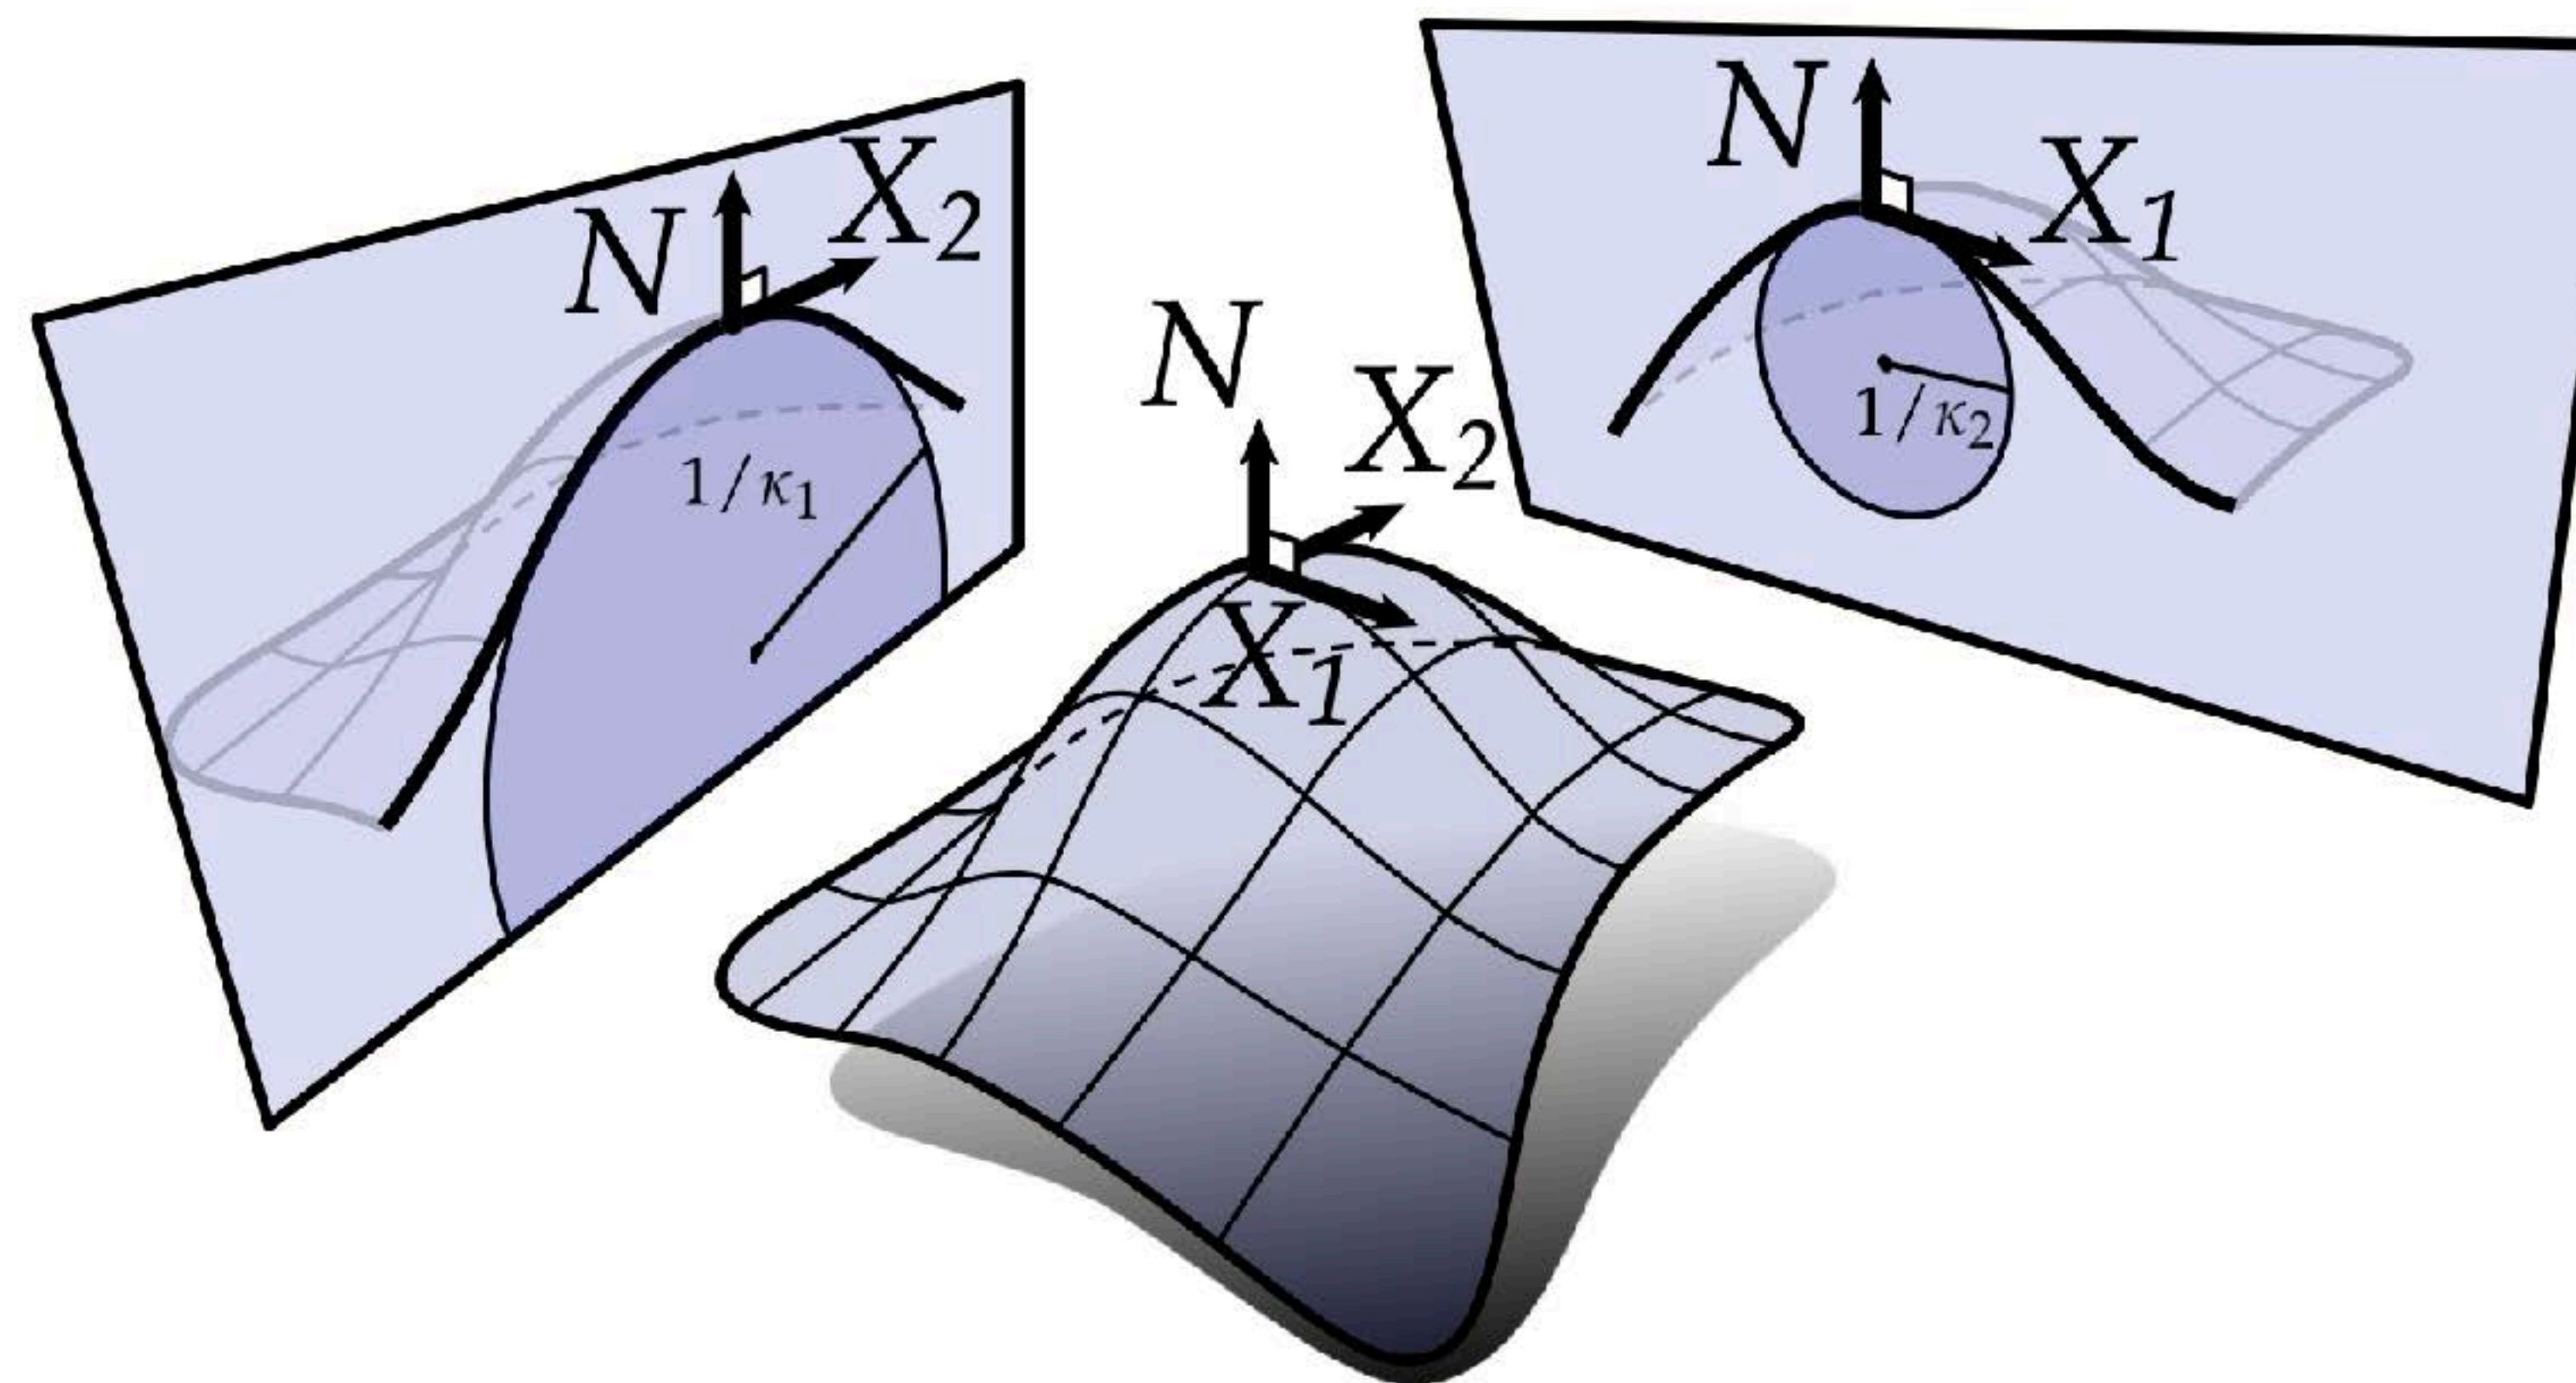
- A parameterized surface is a map f from a two-dimensional region U into \mathbb{R}^n
- $U \subset \mathbb{R}^2$ is the parameter domain
- The set of points $f(U)$ is called the image of the parameterization.



Previously: principal curvatures

Maximal curvature: $\kappa_1 = \kappa_{\max} = \max_{\varphi} \kappa_n(\varphi)$

Minimal curvature: $\kappa_2 = \kappa_{\min} = \min_{\varphi} \kappa_n(\varphi)$

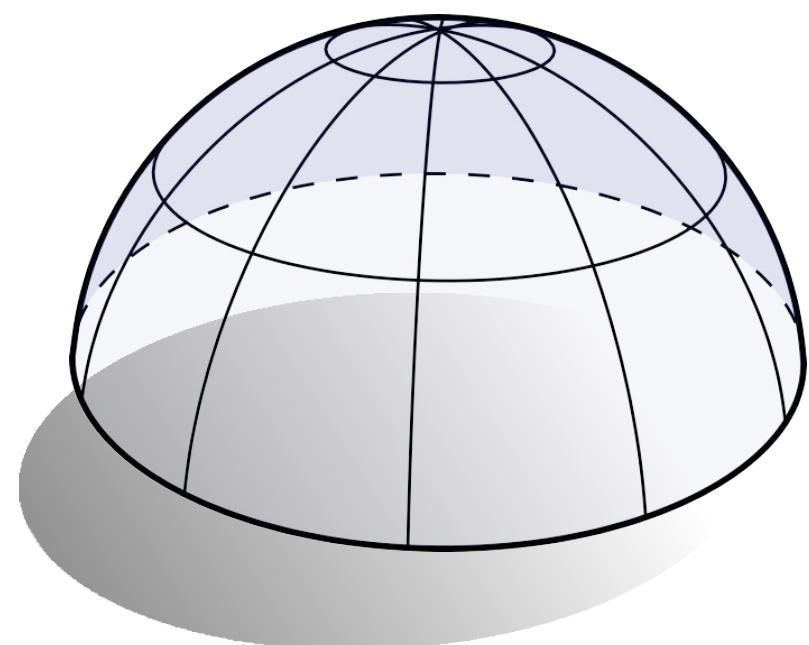


Previously: Gaussian and mean curvature

- Gaussian and mean curvature also fully describe local bending:

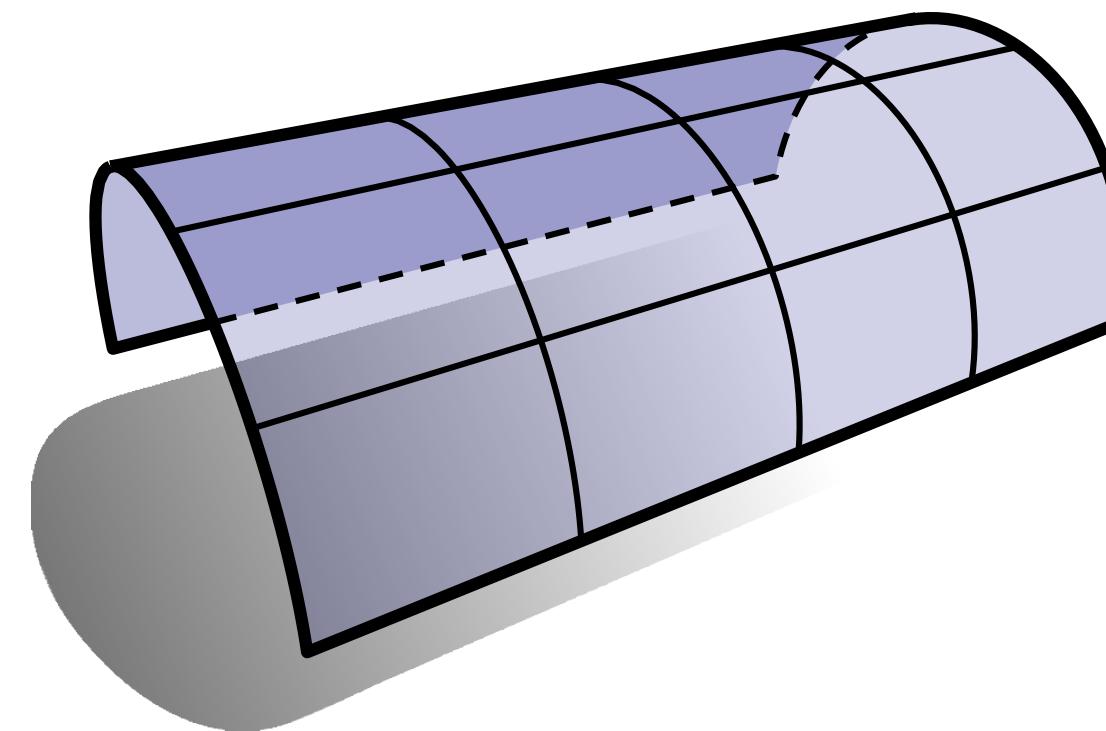
$$\text{Gaussian: } K := \kappa_1 \kappa_2$$

$$\text{mean: } H := \frac{1}{2}(\kappa_1 + \kappa_2)$$



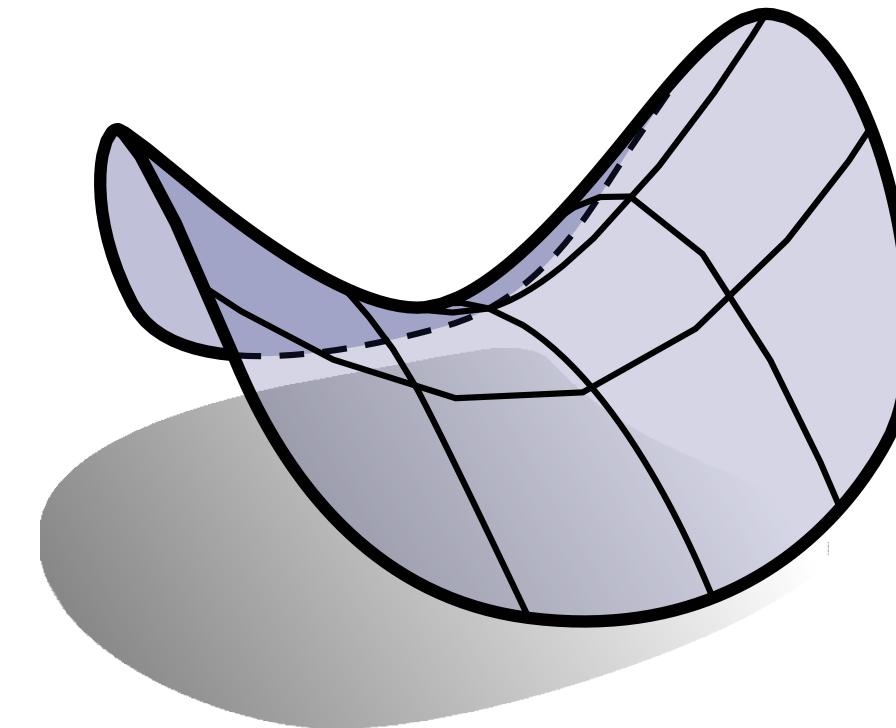
$$K > 0$$

$$H \neq 0$$



$$\text{"developable"} \quad K = 0$$

$$H \neq 0$$



$$K < 0$$

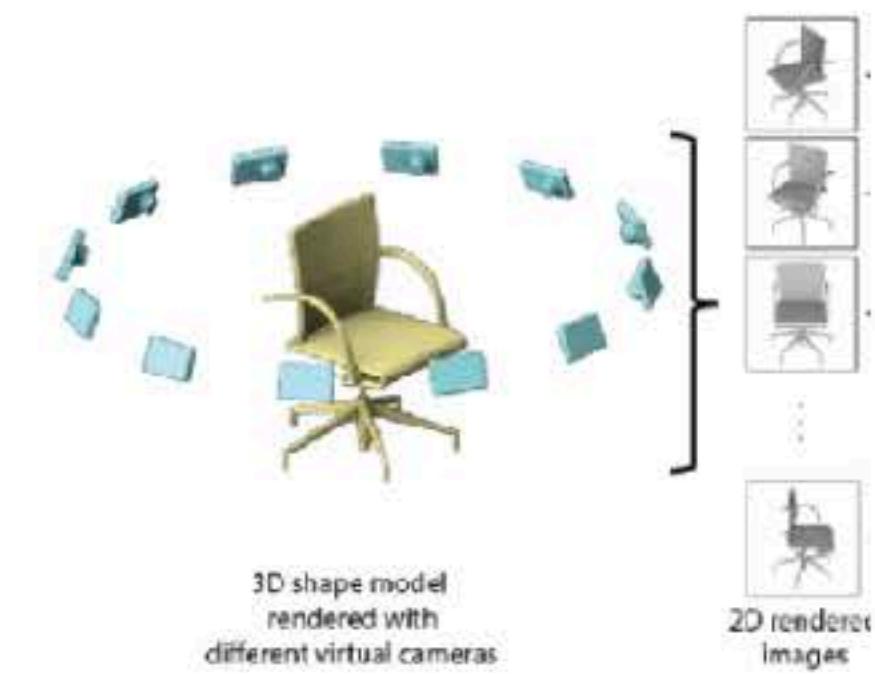
$$\text{"minimal"} \quad H = 0$$





Today's focus: 3D representations

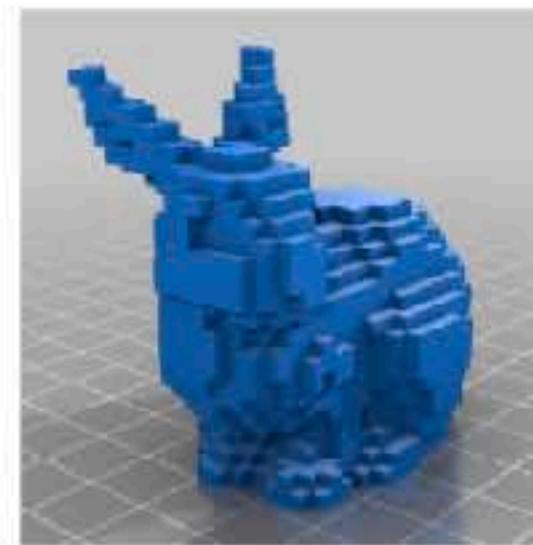
Rasterized form (regular grids)



Multi-view

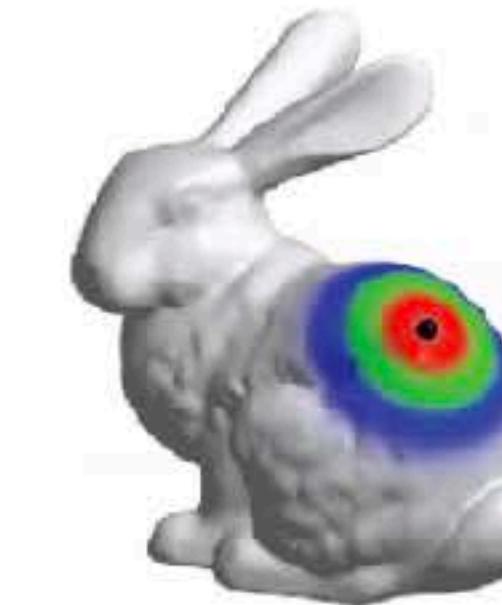


Depth Map

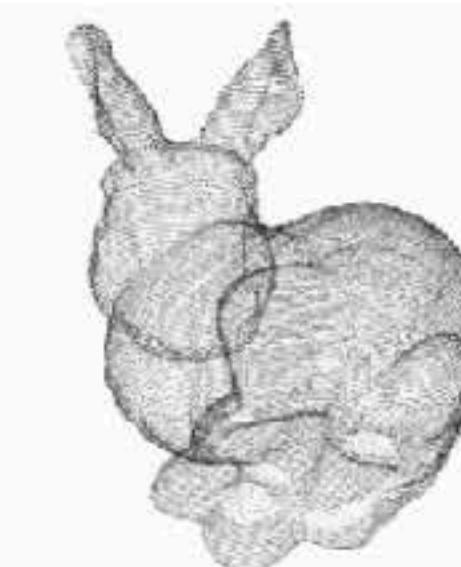


Volumetric

Geometric form (irregular)



Mesh



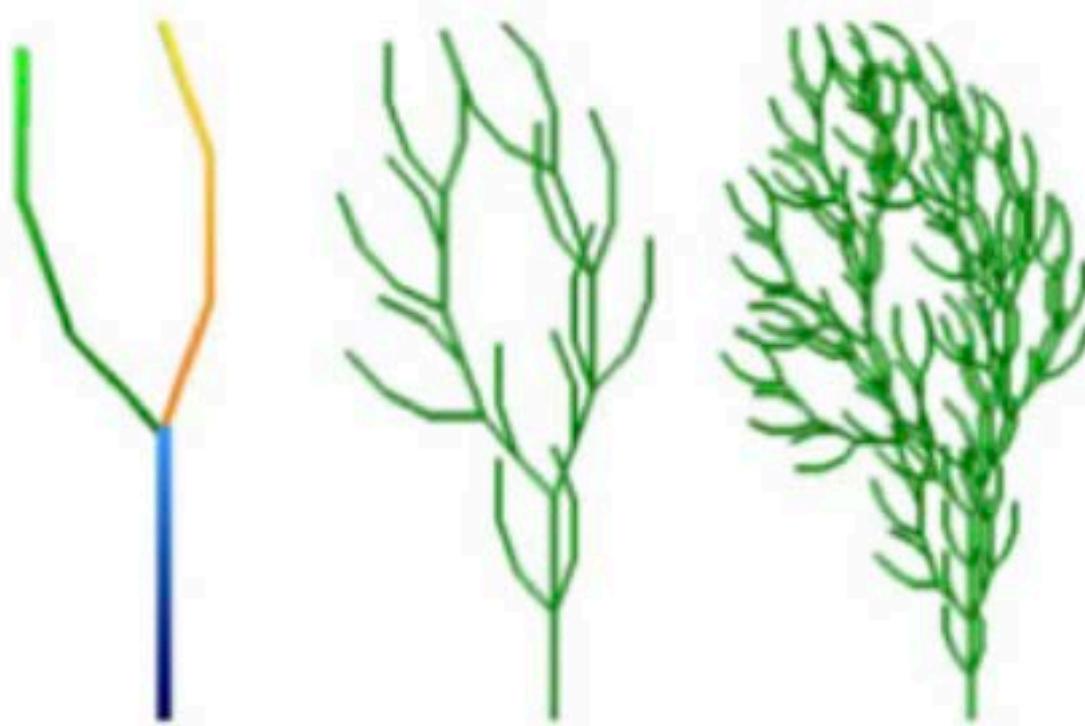
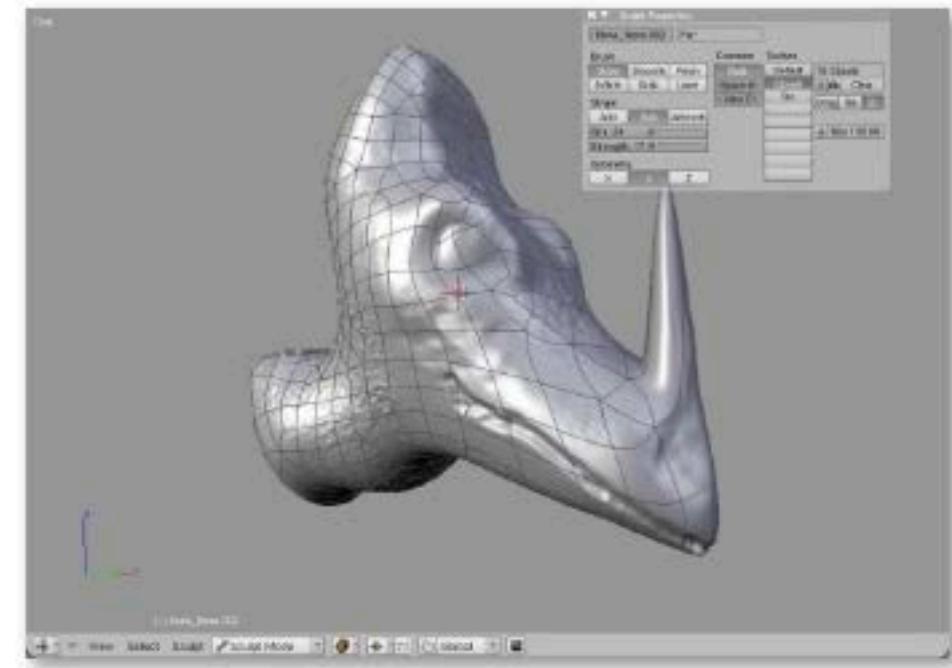
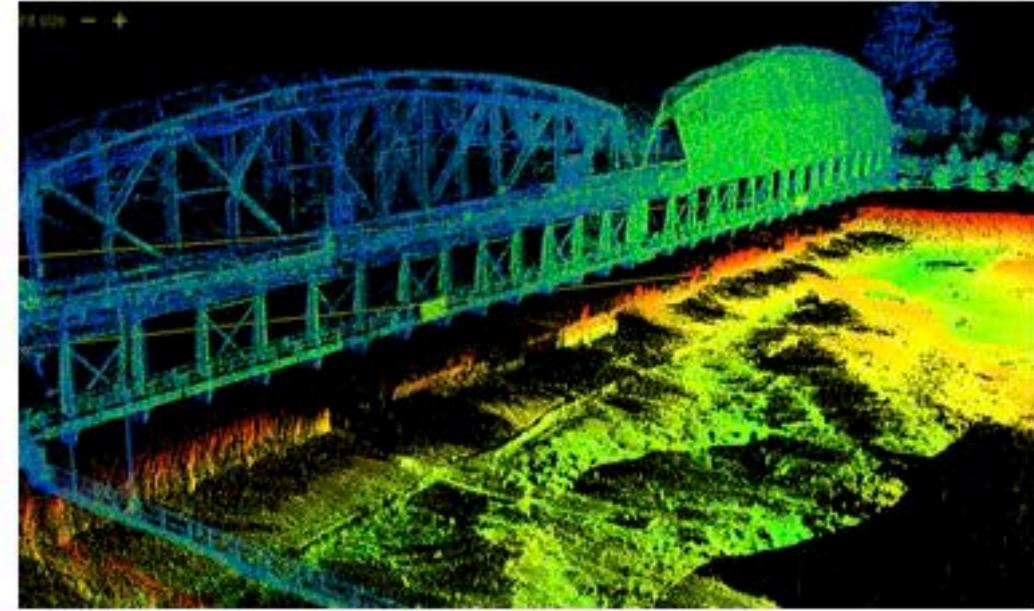
Point Cloud

$$F(x) = 0$$

Implicit Shape

3D representations: origin-dependent

- Acquired real-world object
- Modeling “by hand”
- Procedural modeling
- ...



3D representations: application-dependent

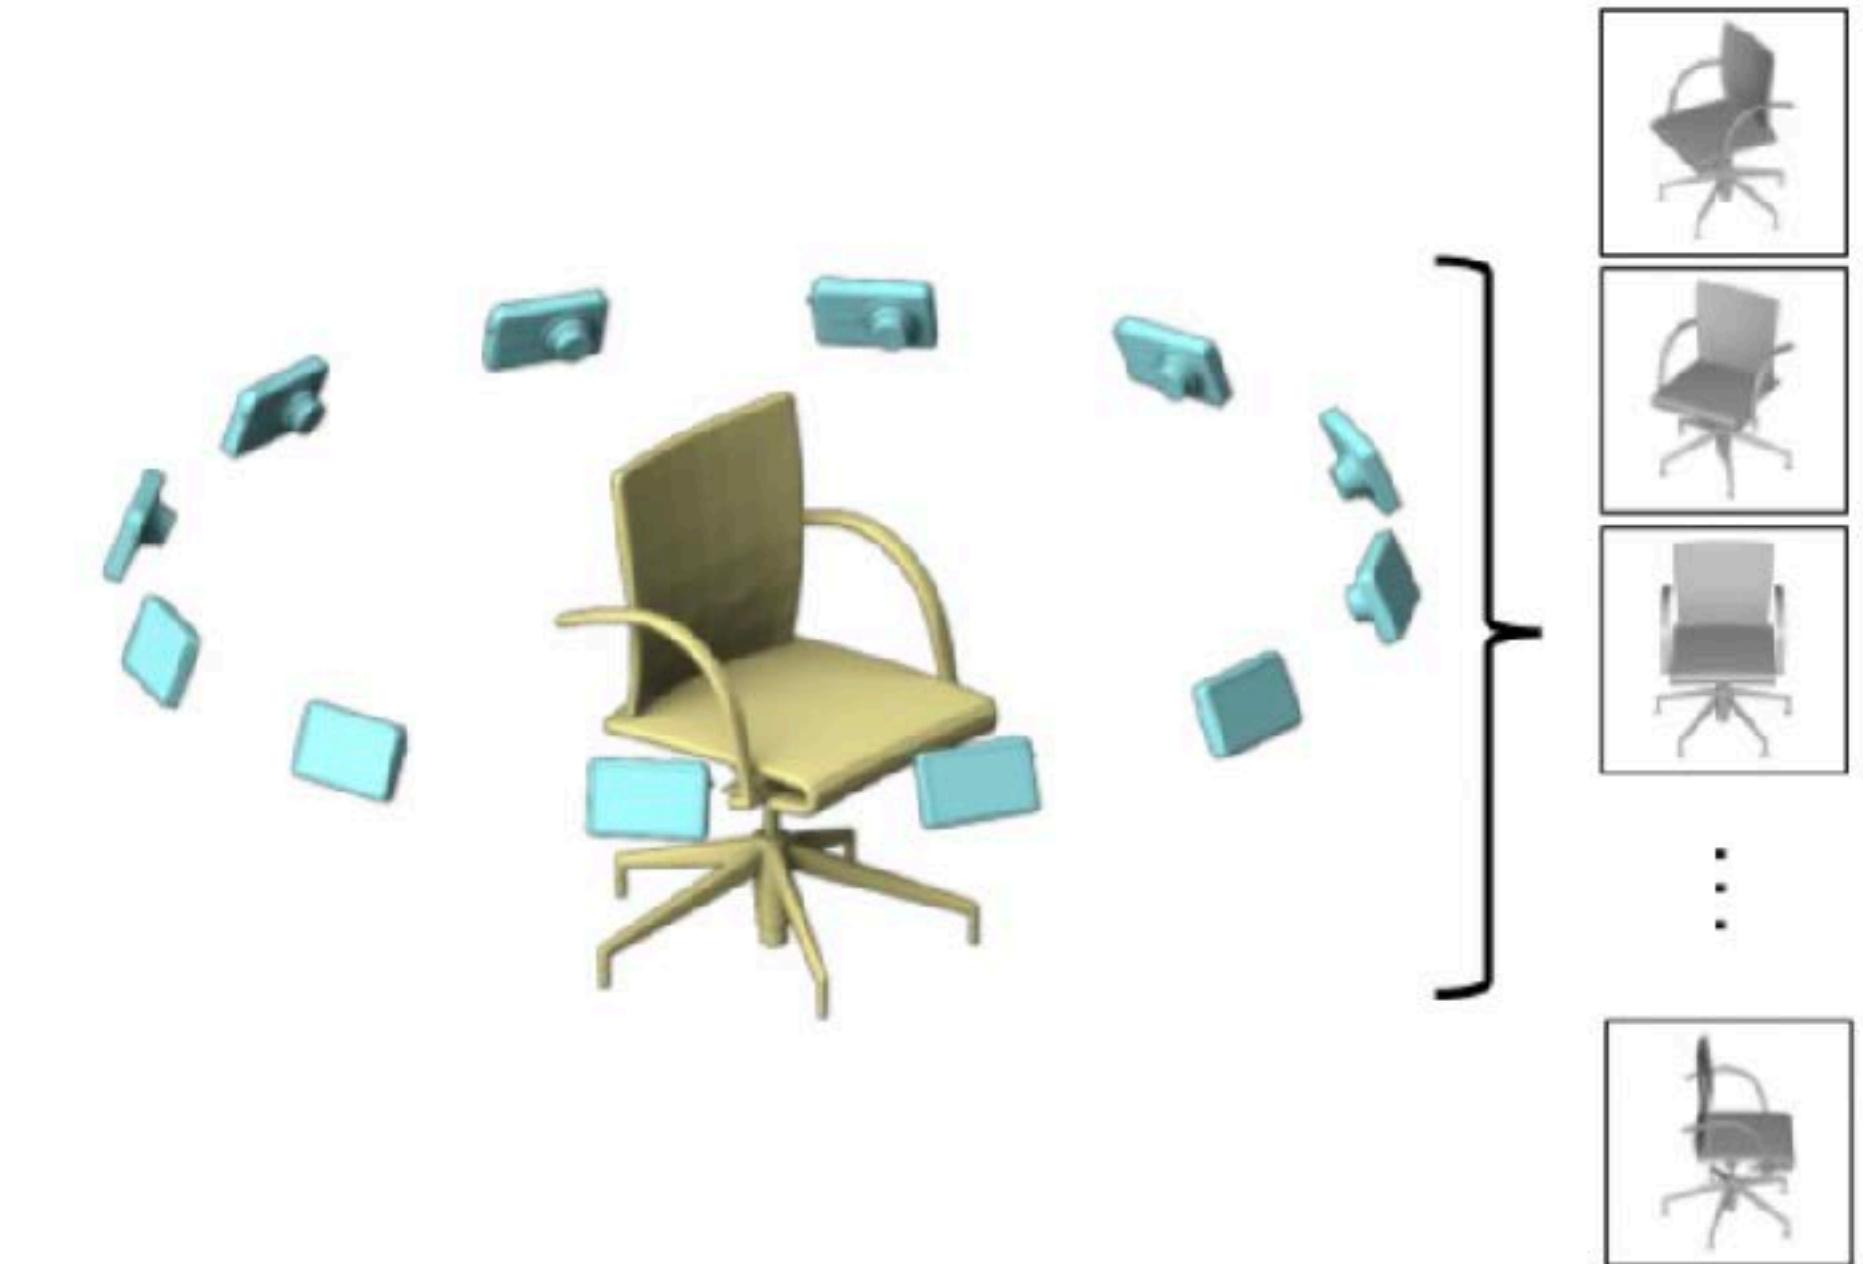
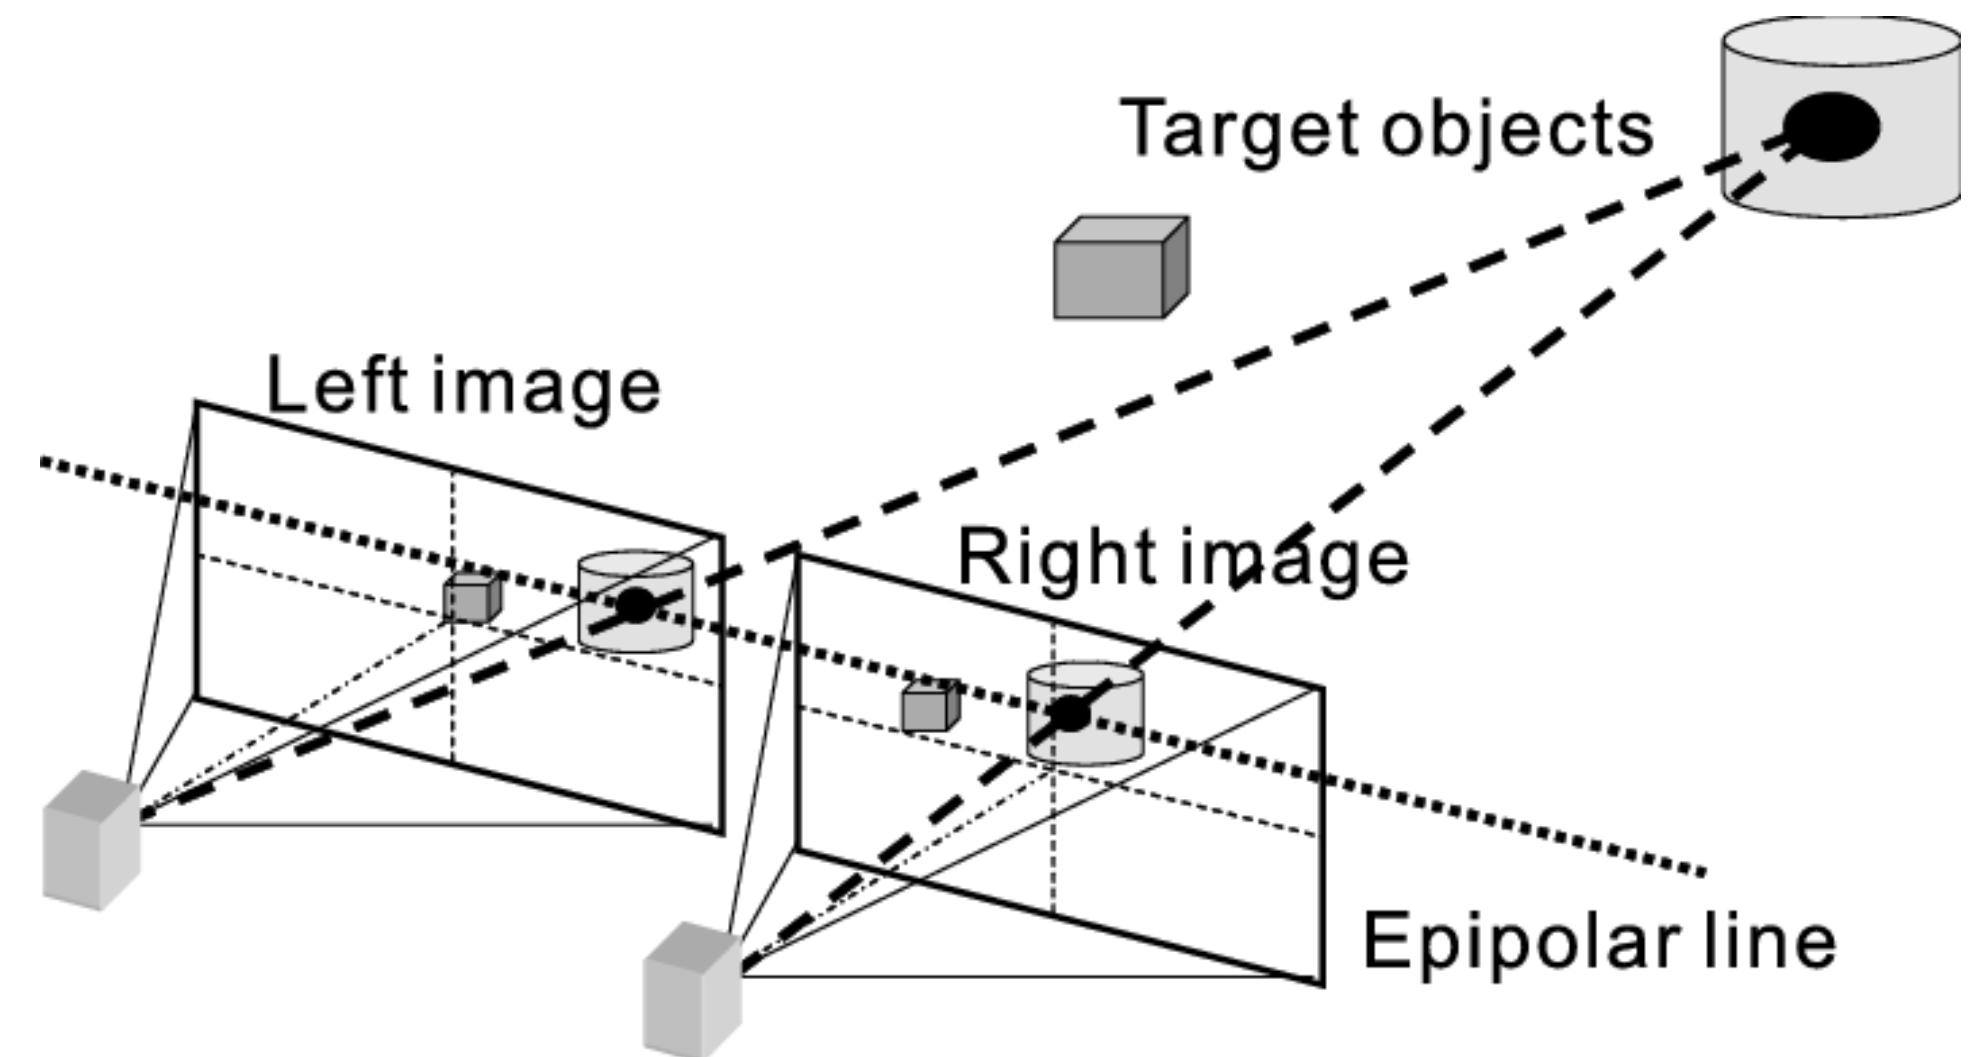
- Needs to be stored in the computer
- Creation of new shapes — input metaphors, interfaces...
- Operations — editing, simplification, smoothing, filtering, repairing...
- Rendering — rasterization, ray tracing...
- Animation

Outline

- *Rasterized 3D representations*
- Mesh
- Point cloud
- Implicit representations

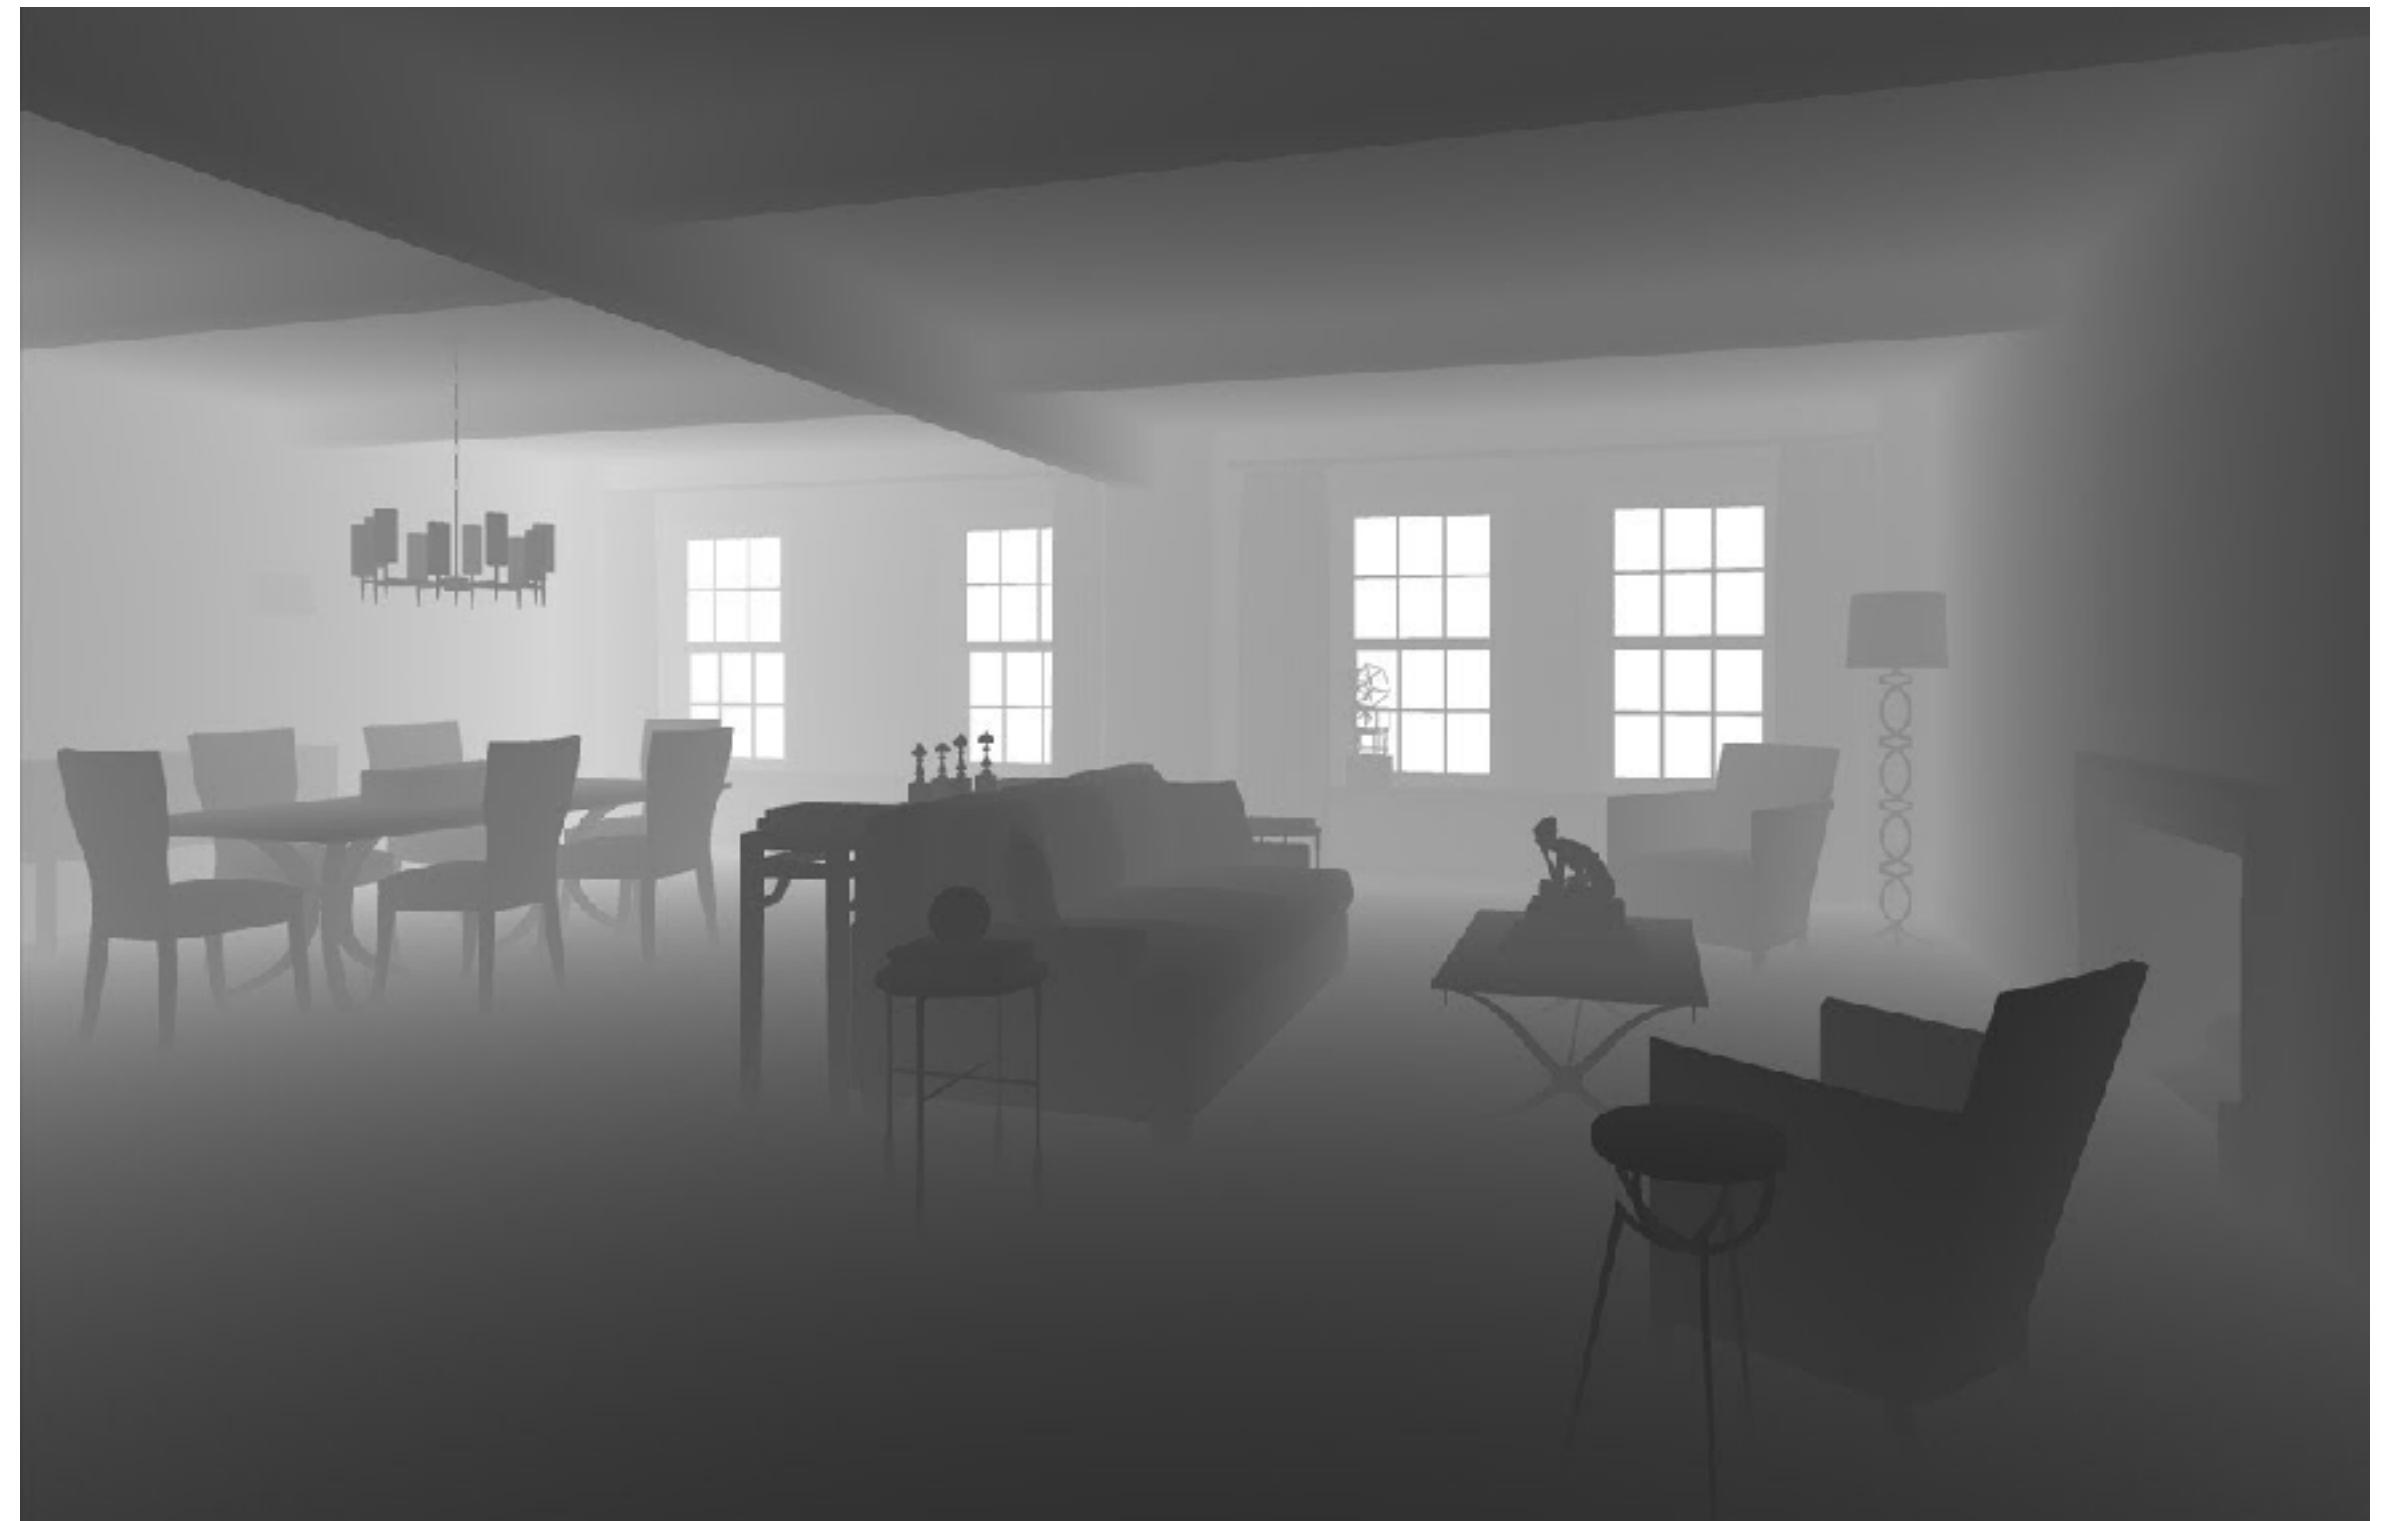
Multiview 3D representation

- No explicit 3D geometry
- Multiview reconstruction possible (non-trivial though)
- Viewpoint matters a lot
- How many views are sufficient?



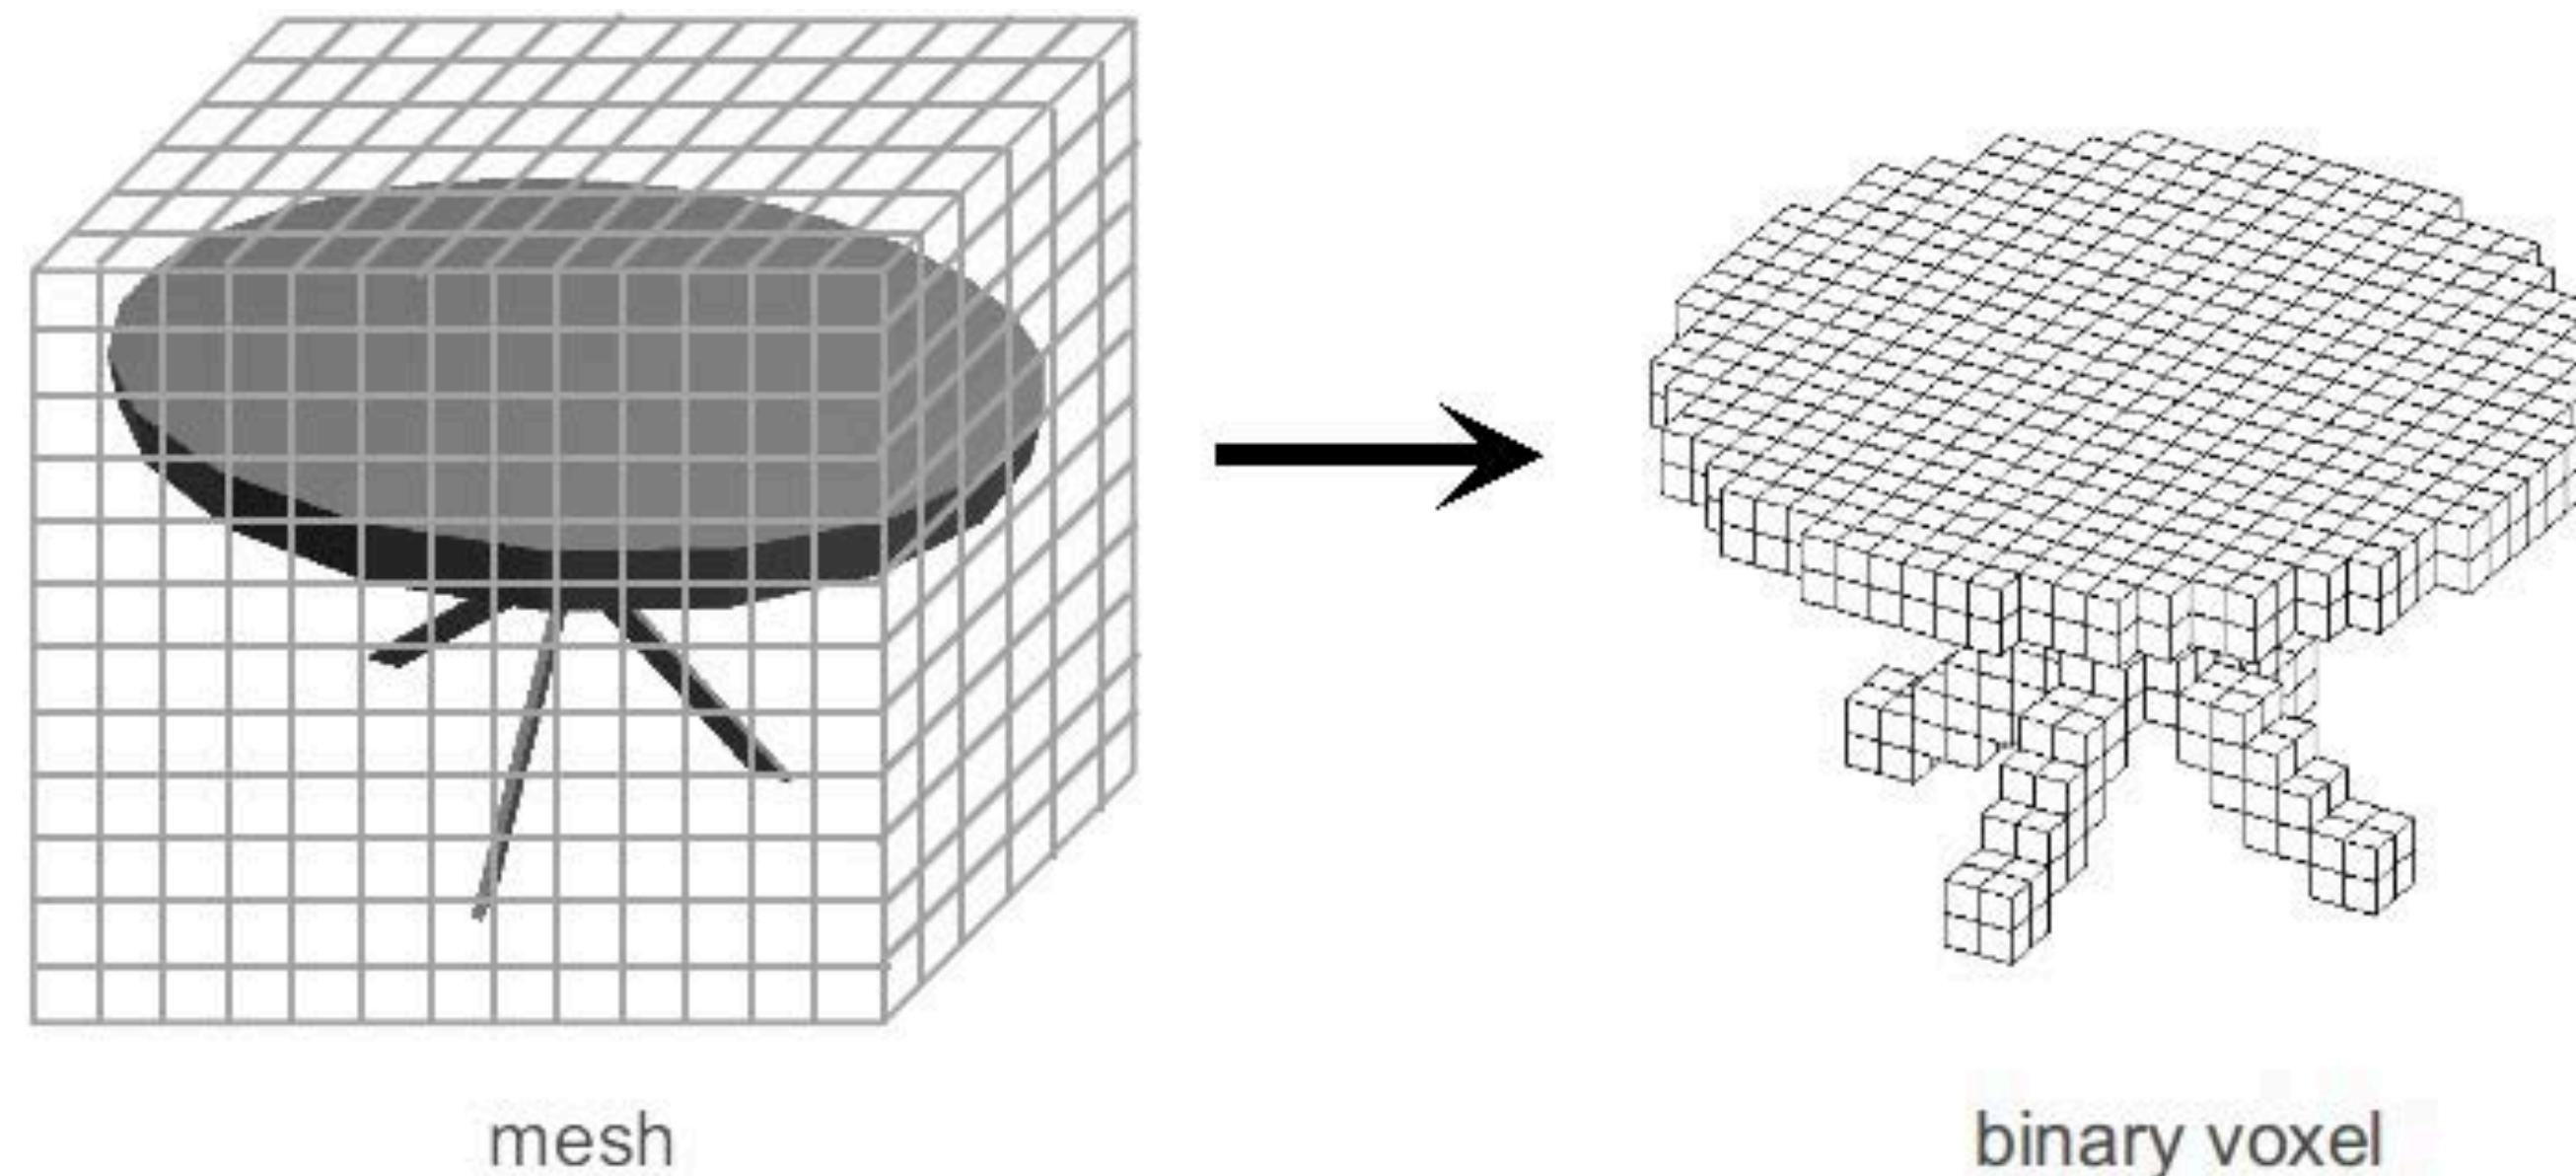
Depth map

- Incomplete 3D
- Need camera information for true 3D
- Resolution issue
- Object properties matter



Volumetric representation

- Storage and computation cost
- Obvious artifacts (w. low resolution)
- Sparse representation could be a remedy

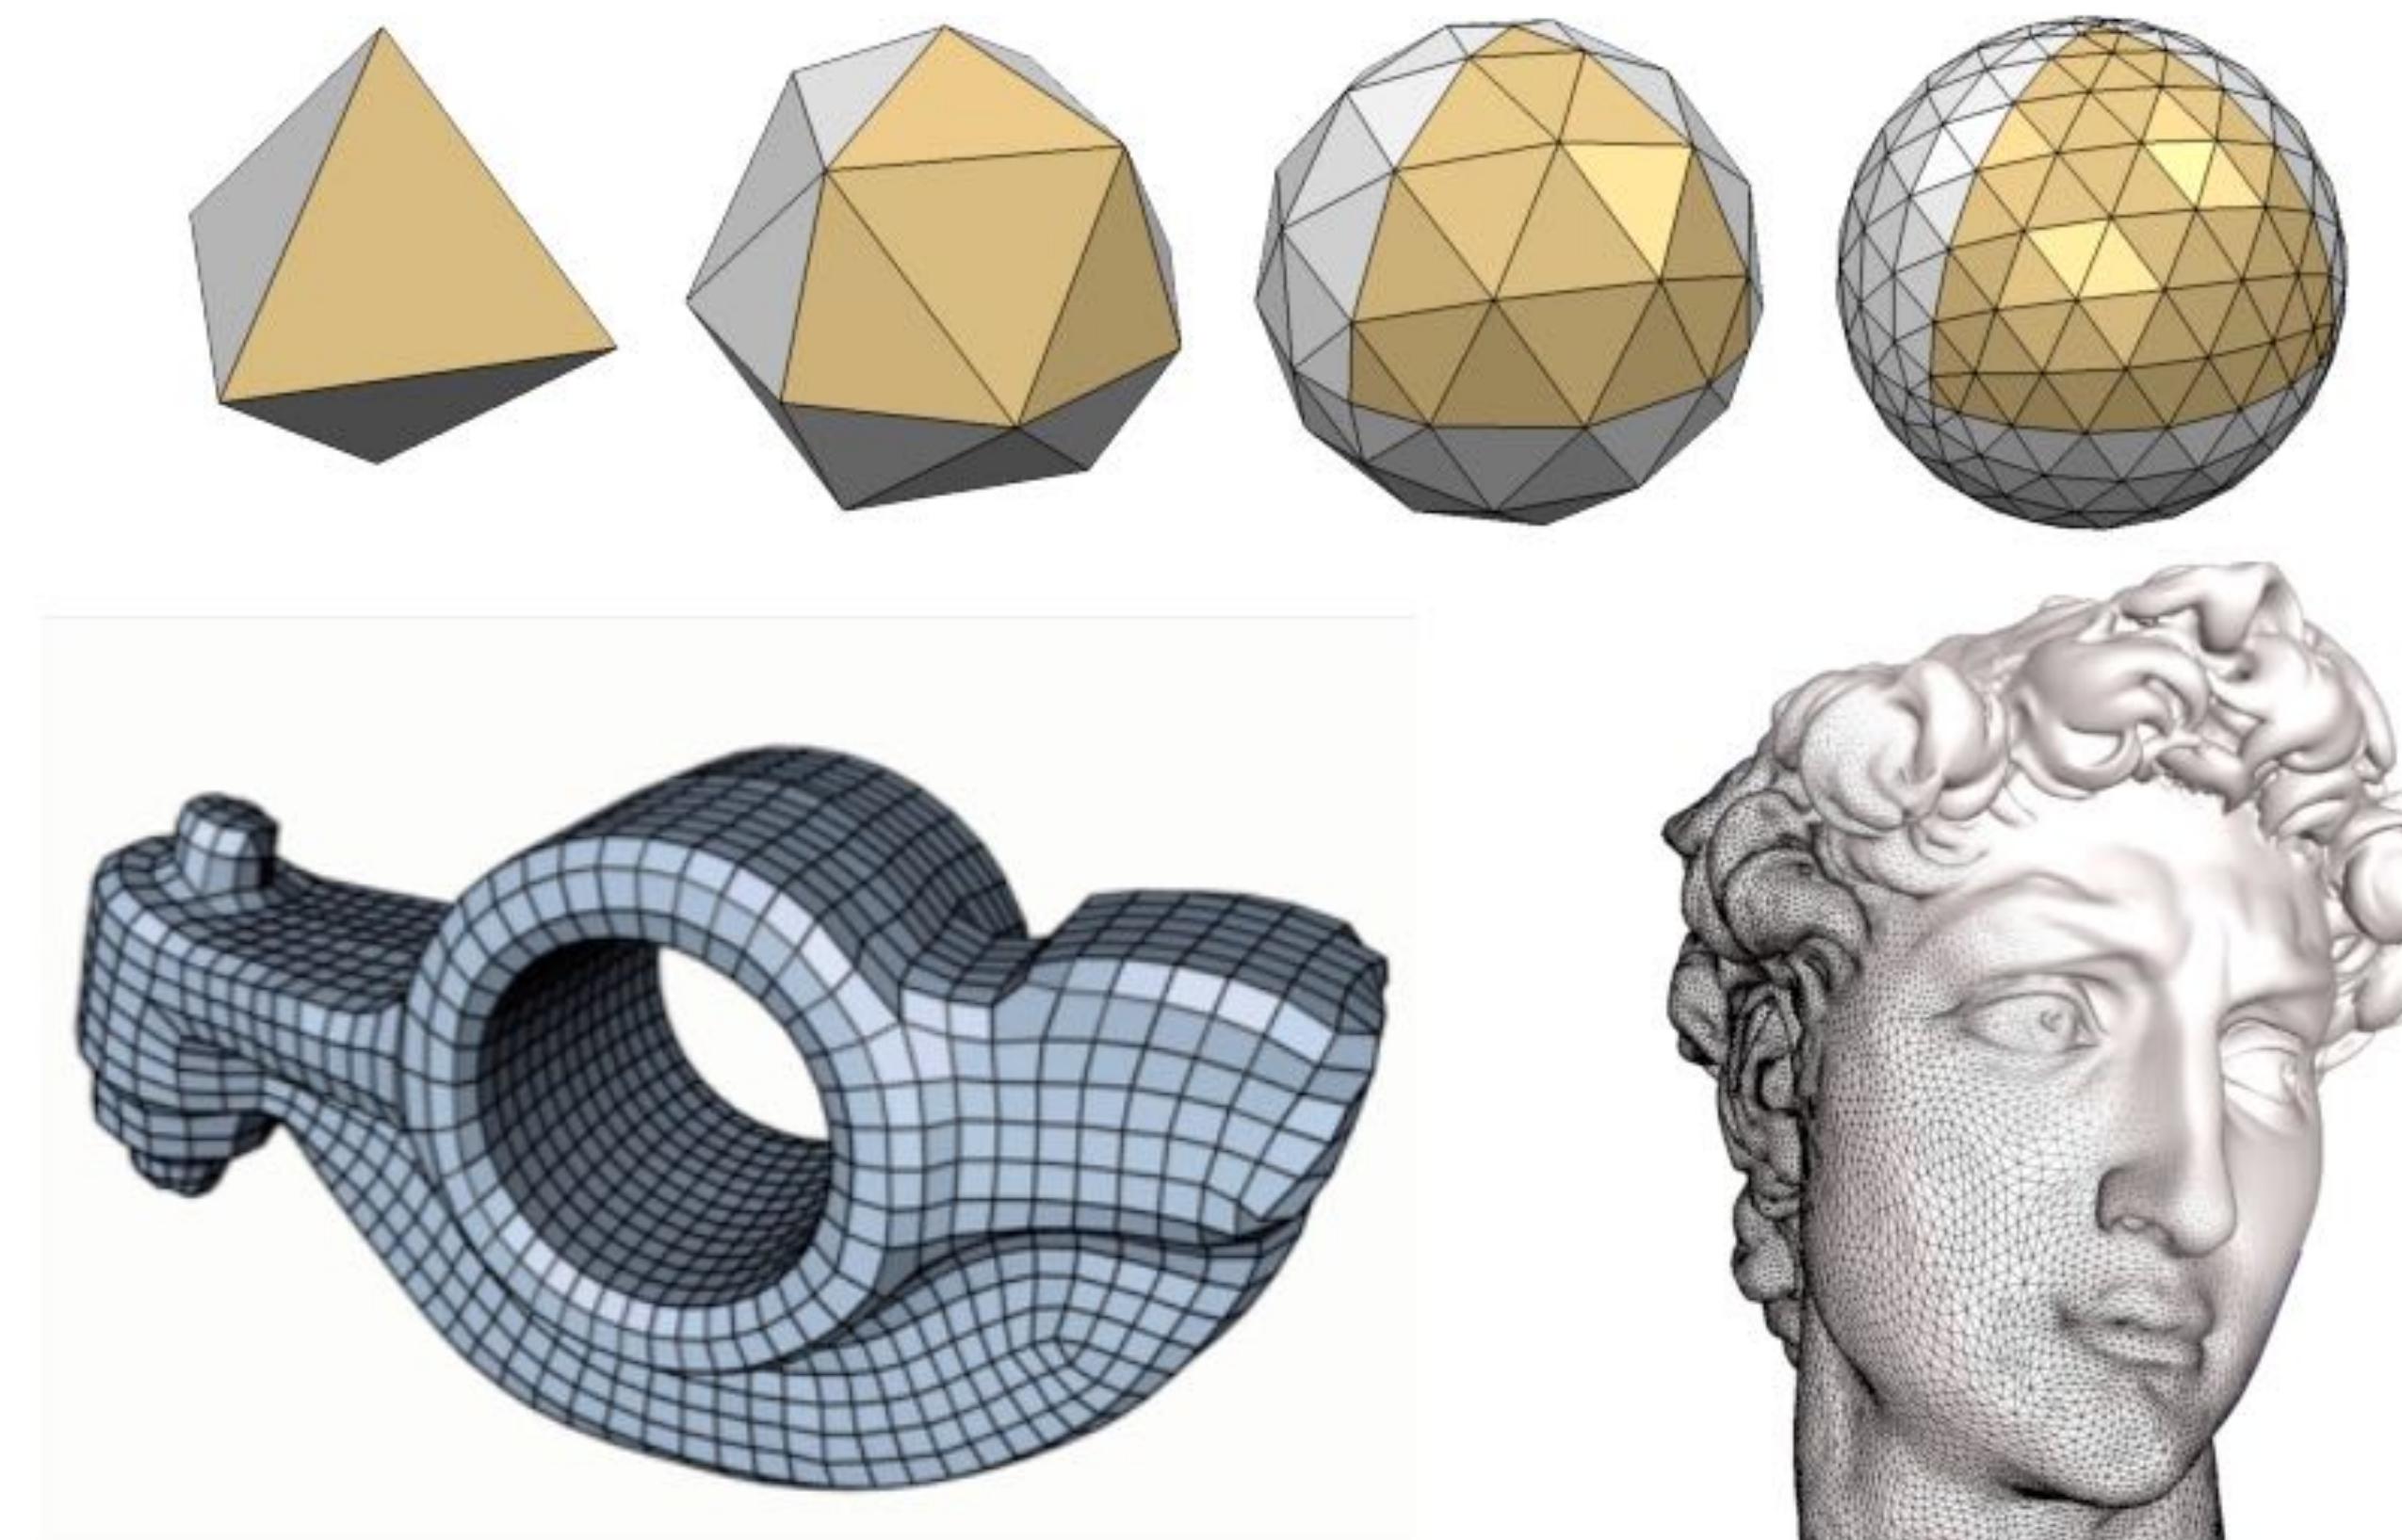


Outline

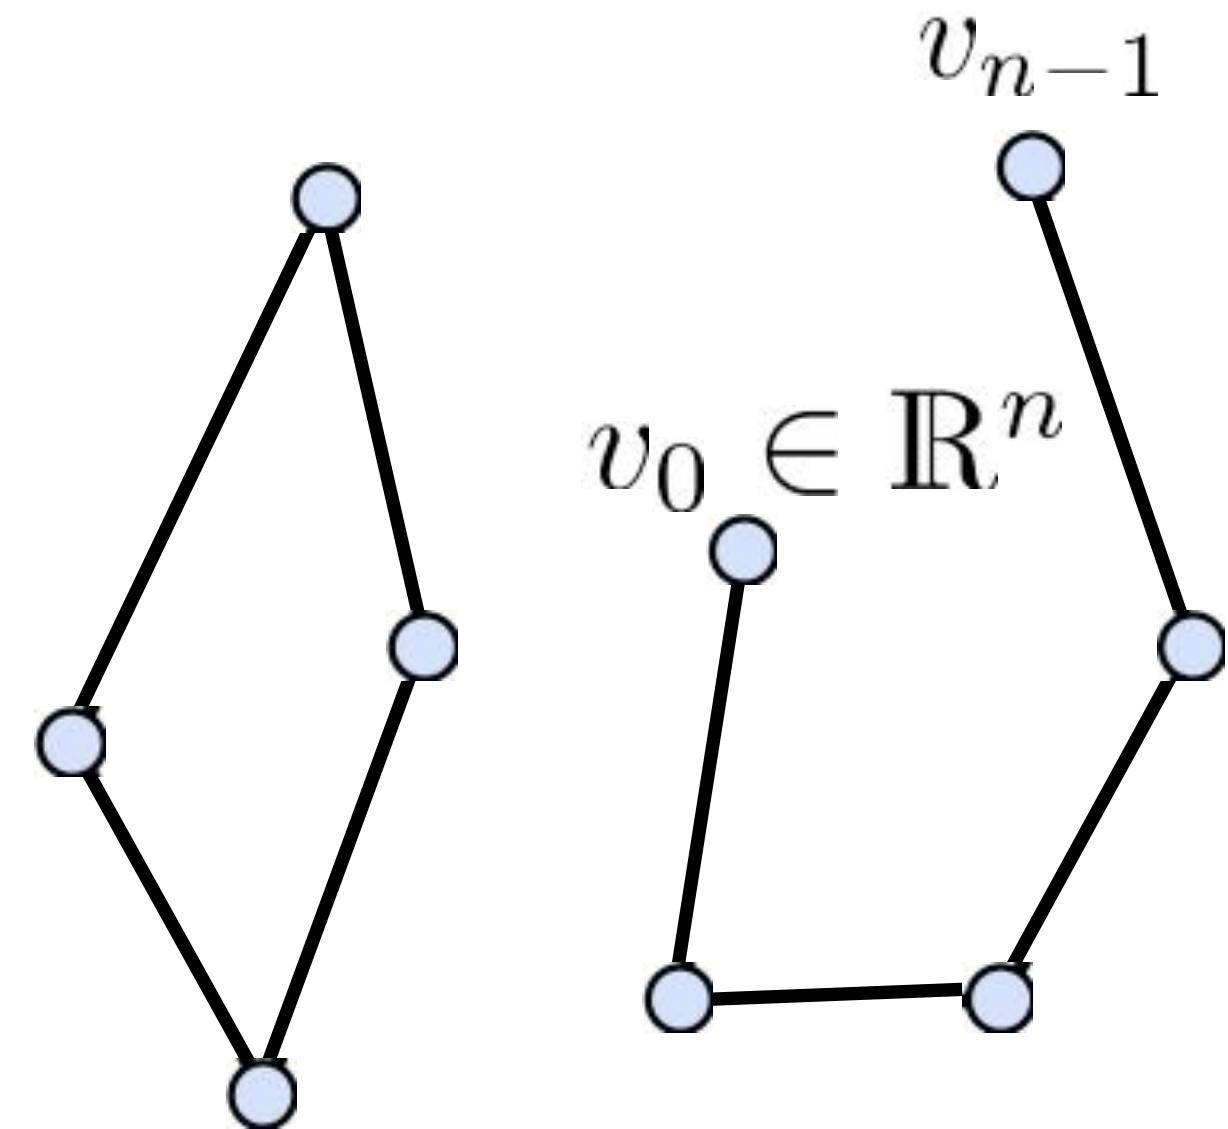
- Rasterized 3D representations
- ***Mesh***
 - ***Representation***
 - Storage
 - Curvature computation
- Point cloud
- Implicit representations

Polygonal Meshes

- Piece-wise linear surface representation



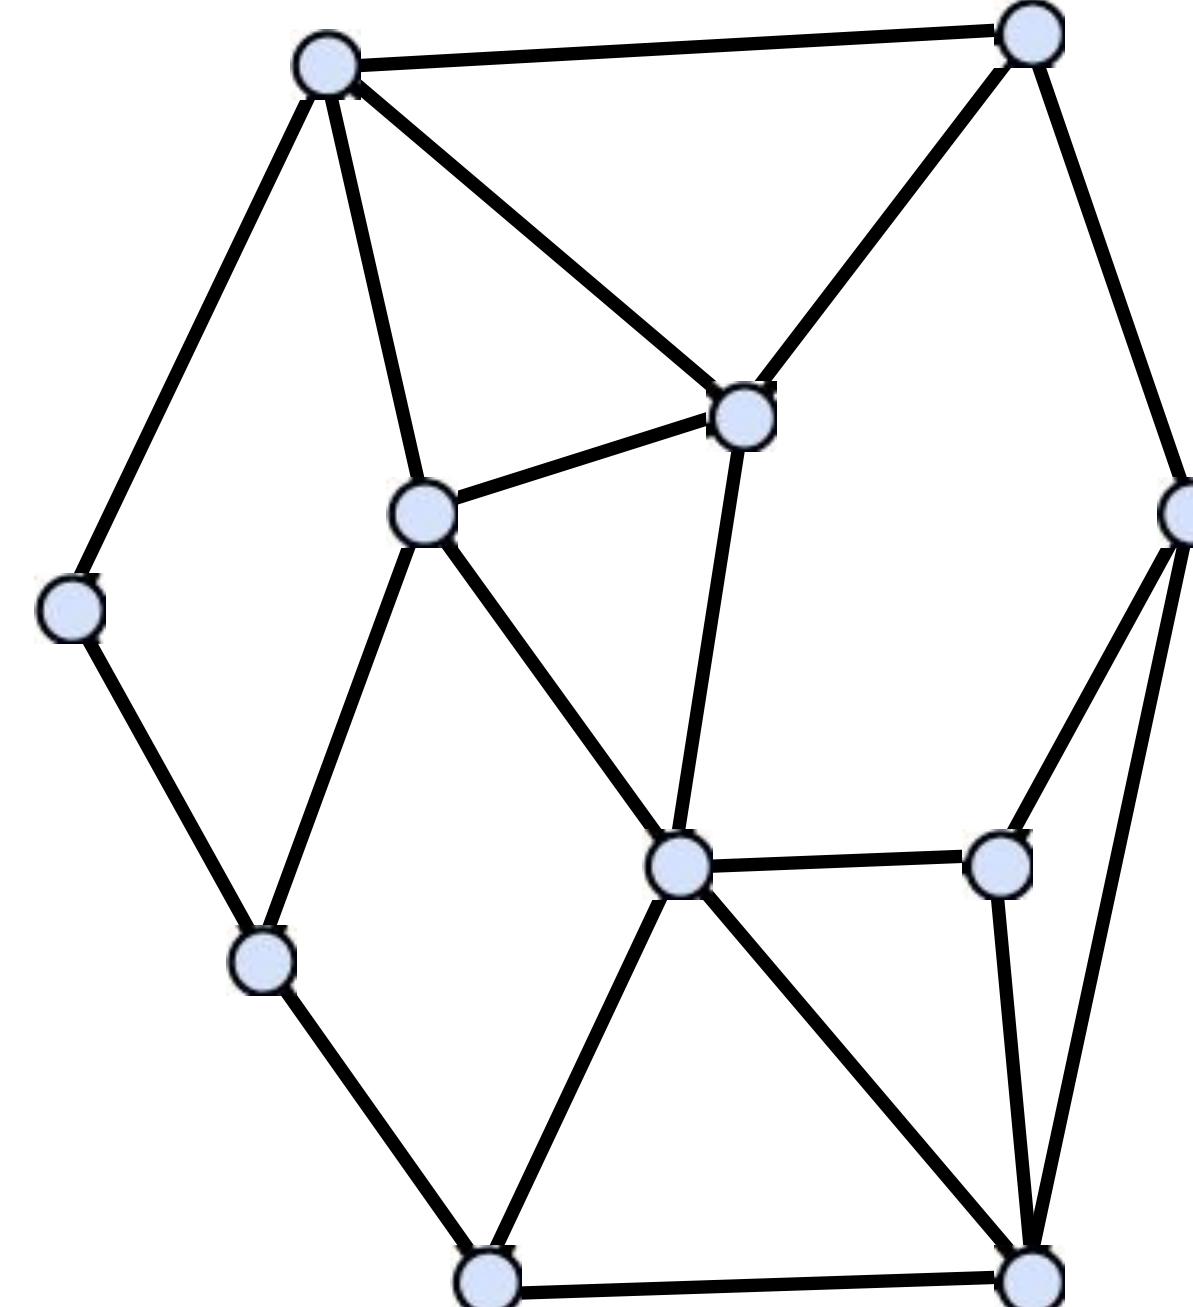
Polygon



- Vertices: v_0, v_1, \dots, v_{n-1}
- Edges: $\{(v_0, v_1), \dots, (v_{n-2}, v_{n-1})\}$
- Closed: $v_0 = v_{n-1}$
- Planar: all vertices on a plane
- Simple: not self-intersecting

Polygonal Mesh

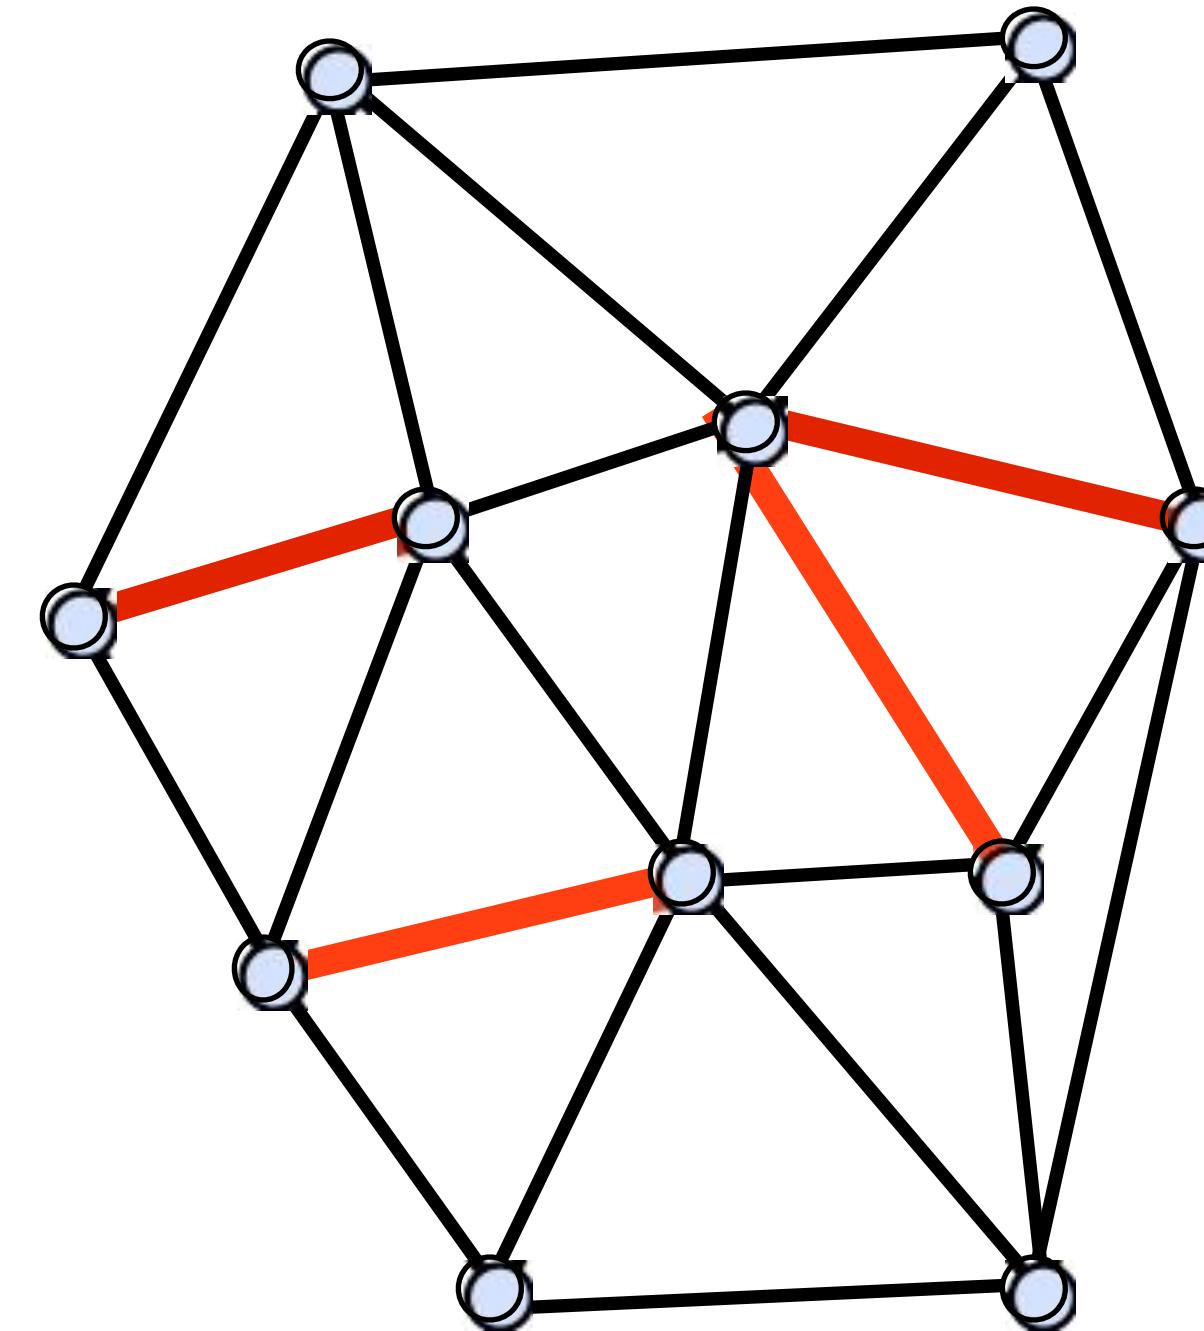
- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge
(for triangle meshes)



$$M = \langle V, E, F \rangle$$

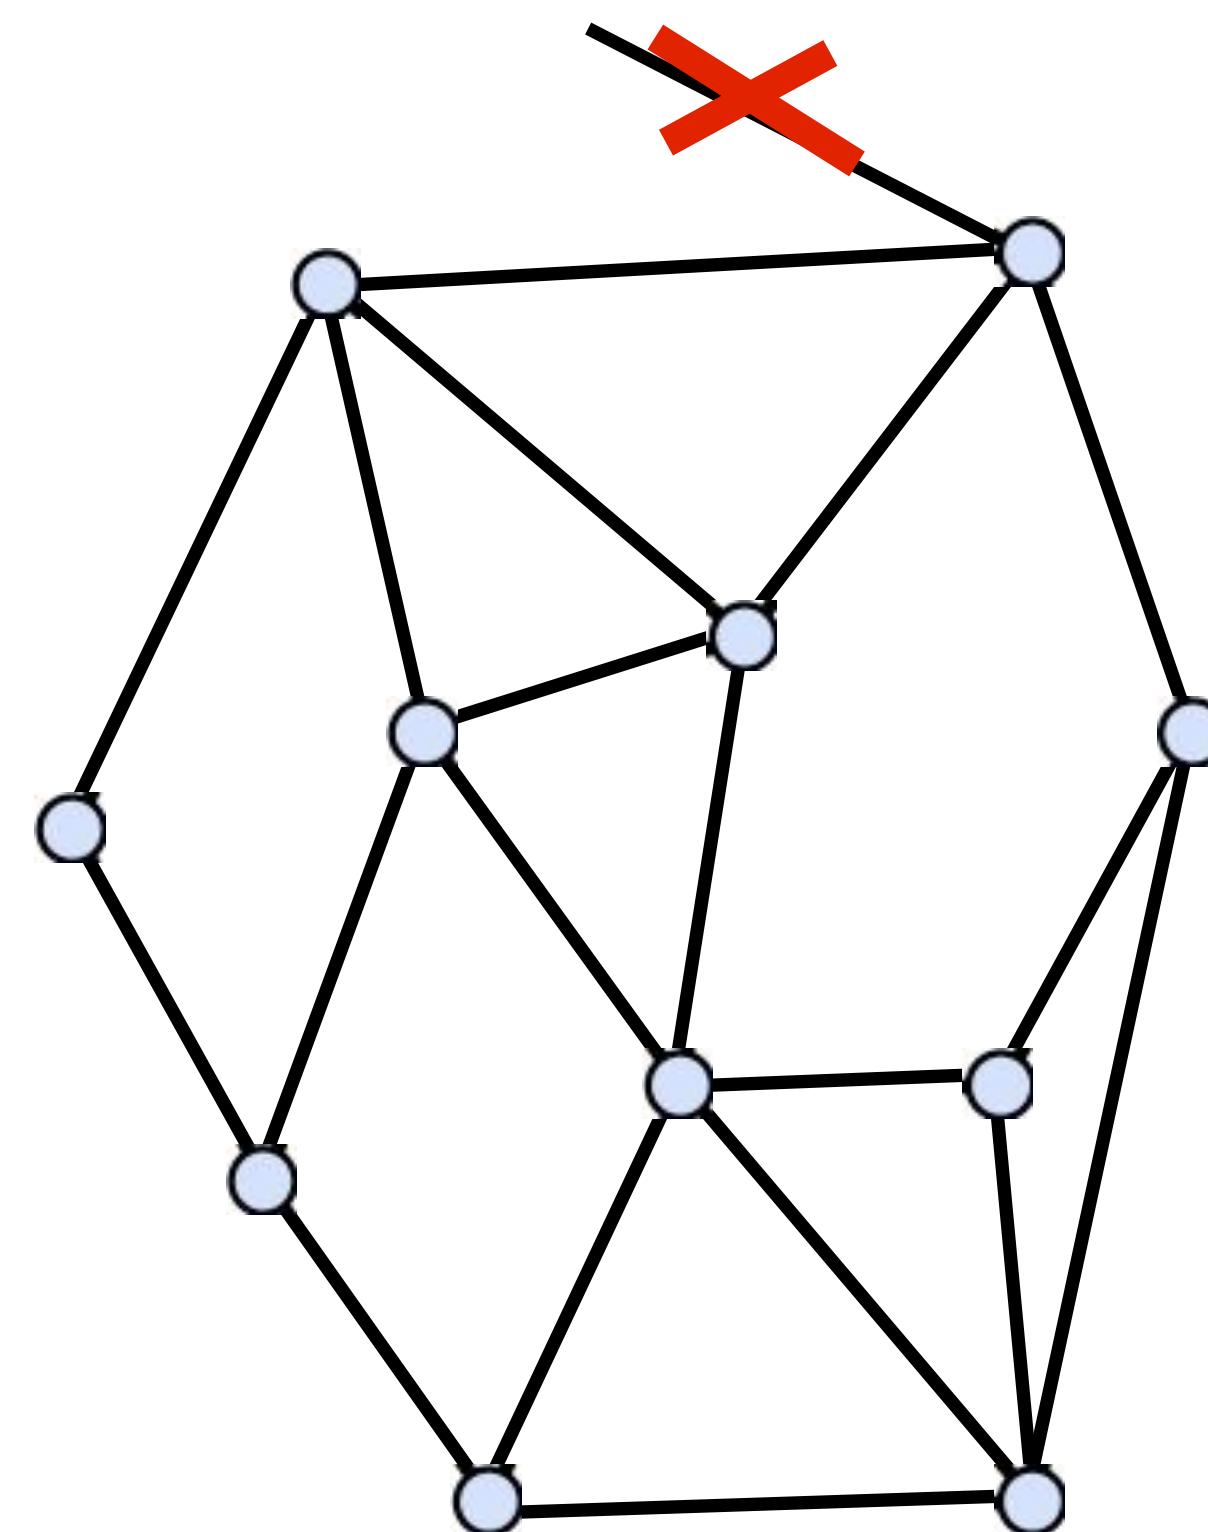
vertices edges faces

Triangulation



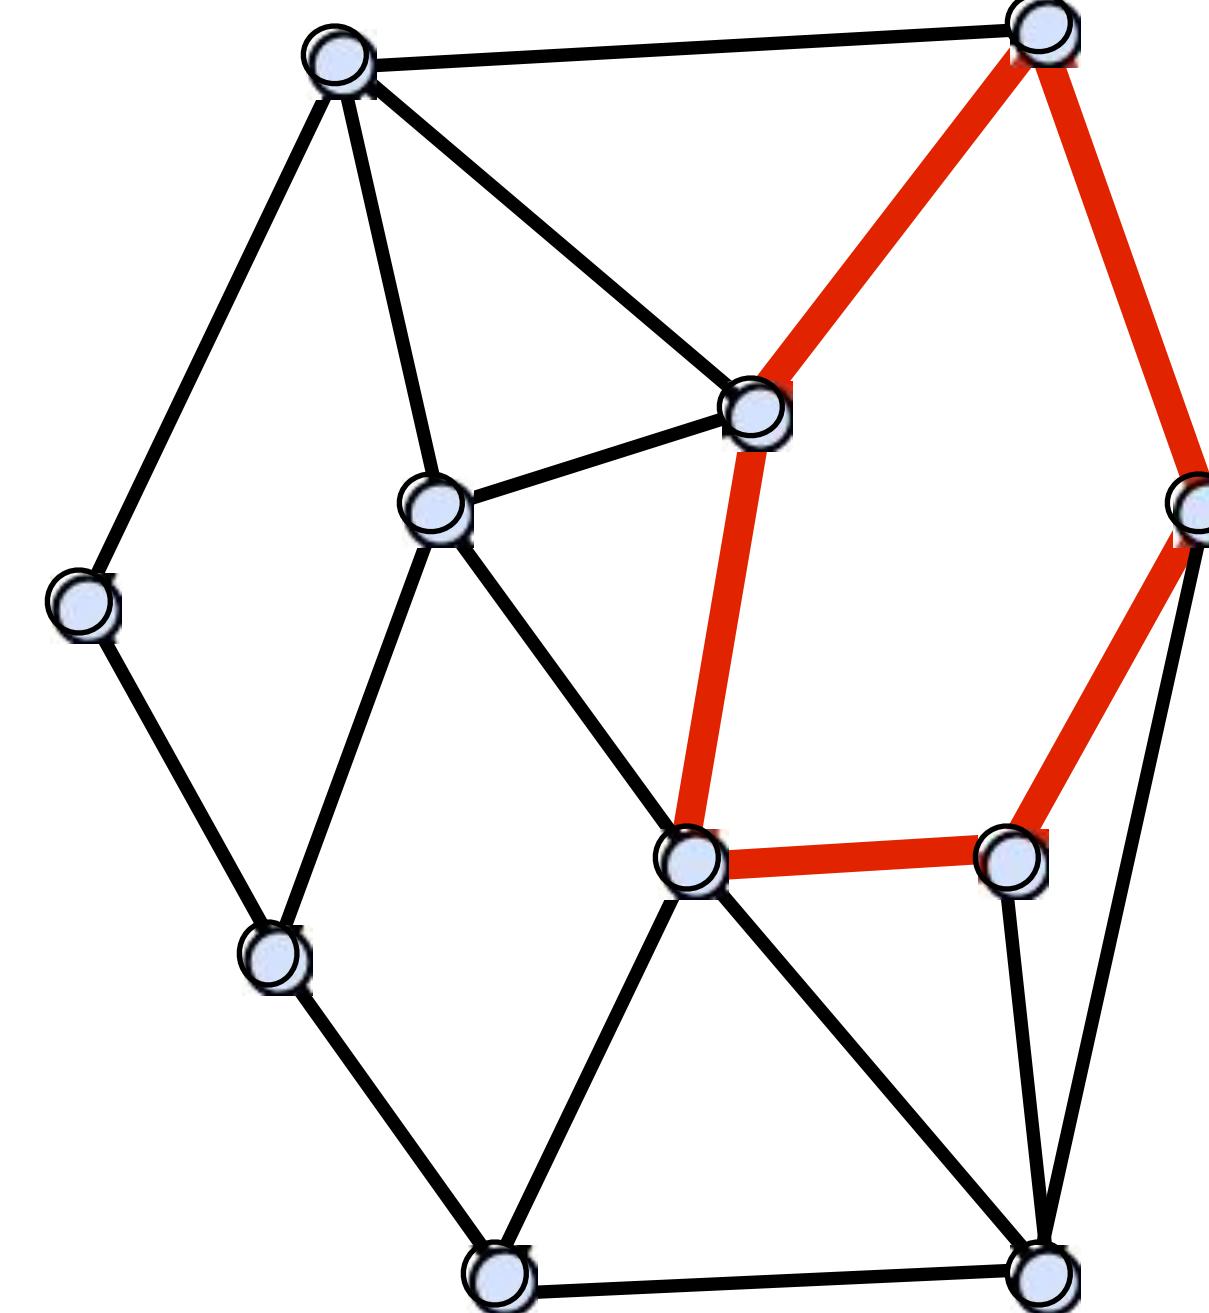
- Polygonal mesh where every face is a triangle
- Simplifies data structures
- Simplifies rendering
- Simplifies algorithms
- Each face planar and convex
- Any polygon can be triangulated

Polygonal Mesh



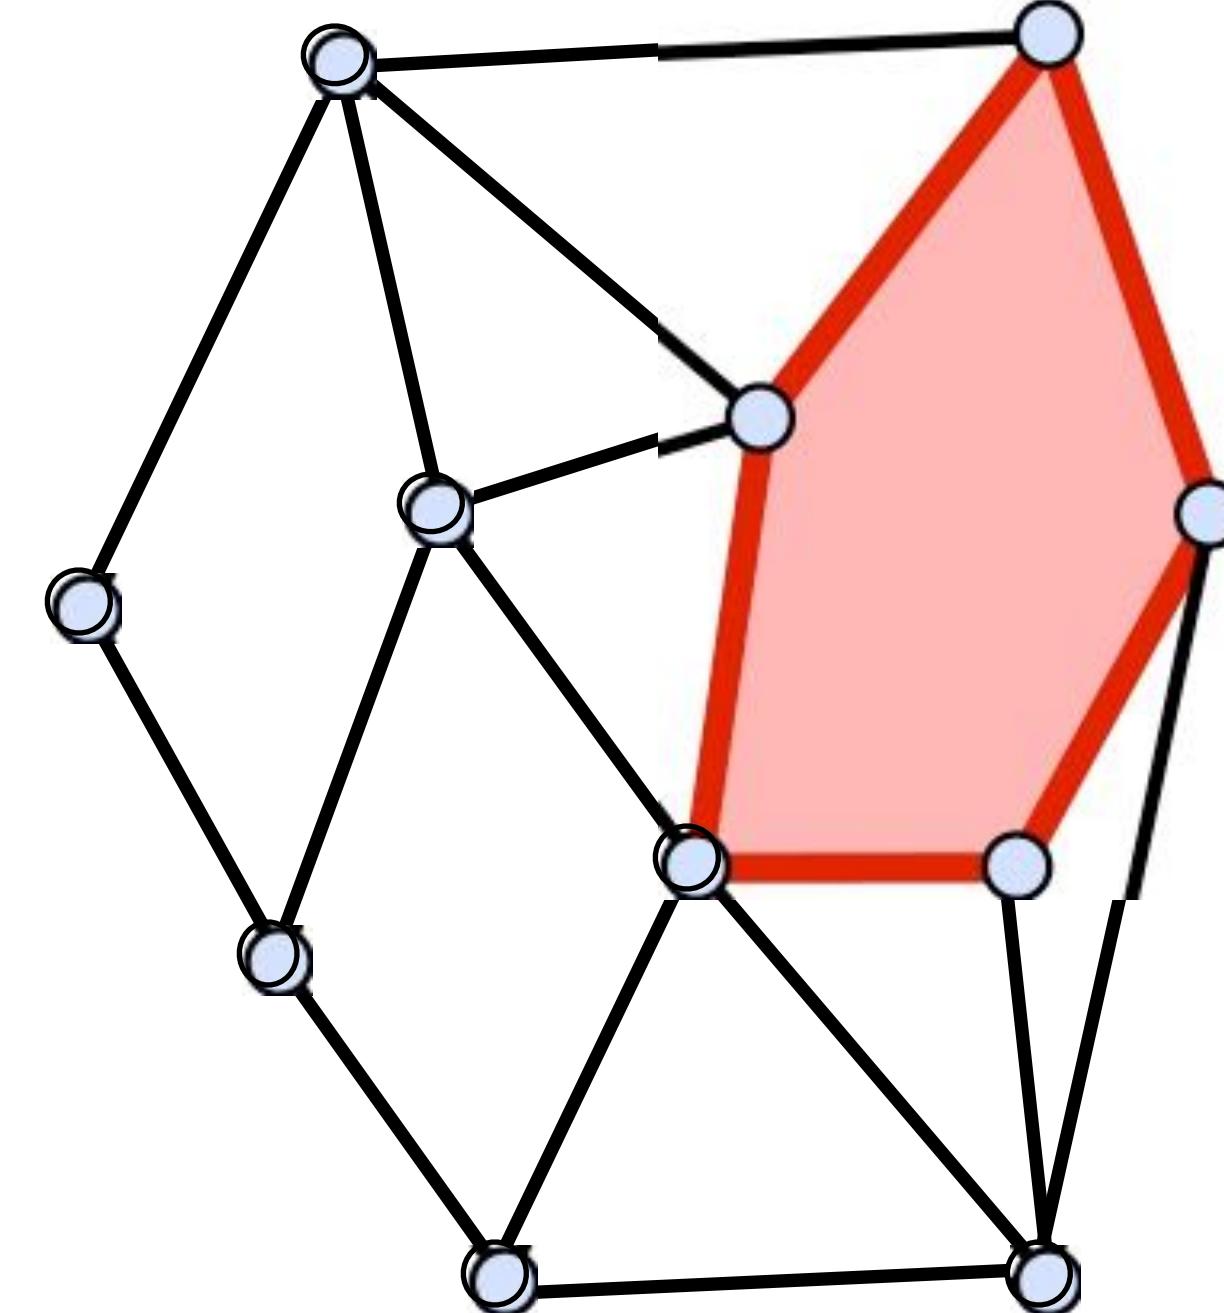
- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon

Polygonal Mesh



- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a **face** of the polygonal mesh

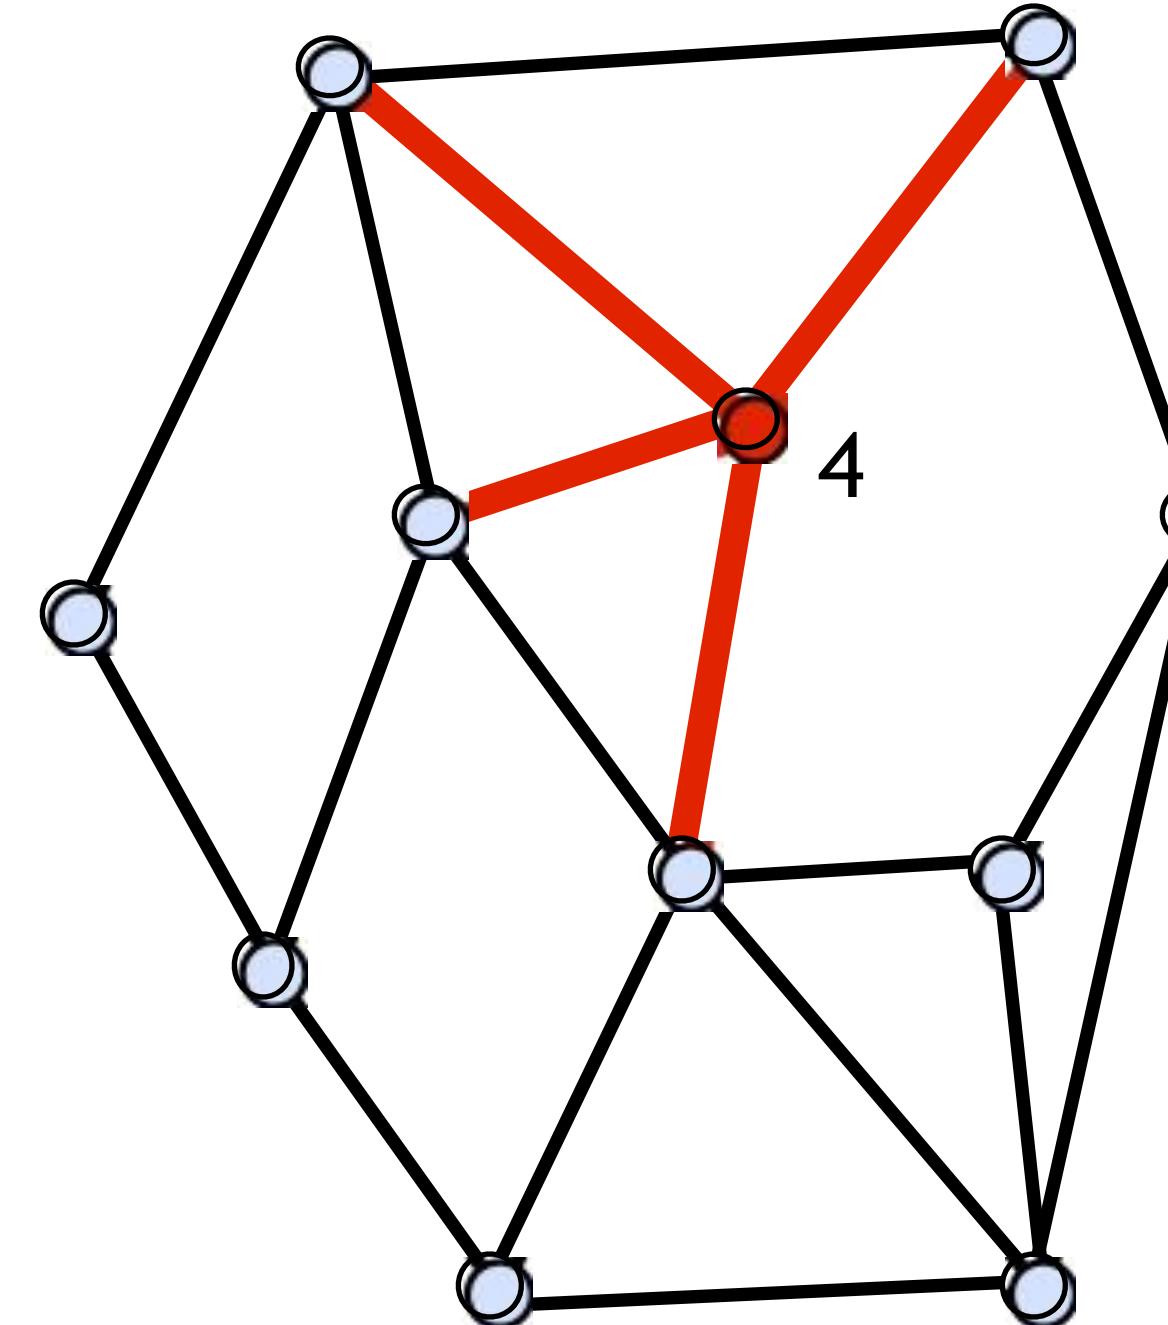
Polygonal Mesh



- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a face of the polygonal mesh

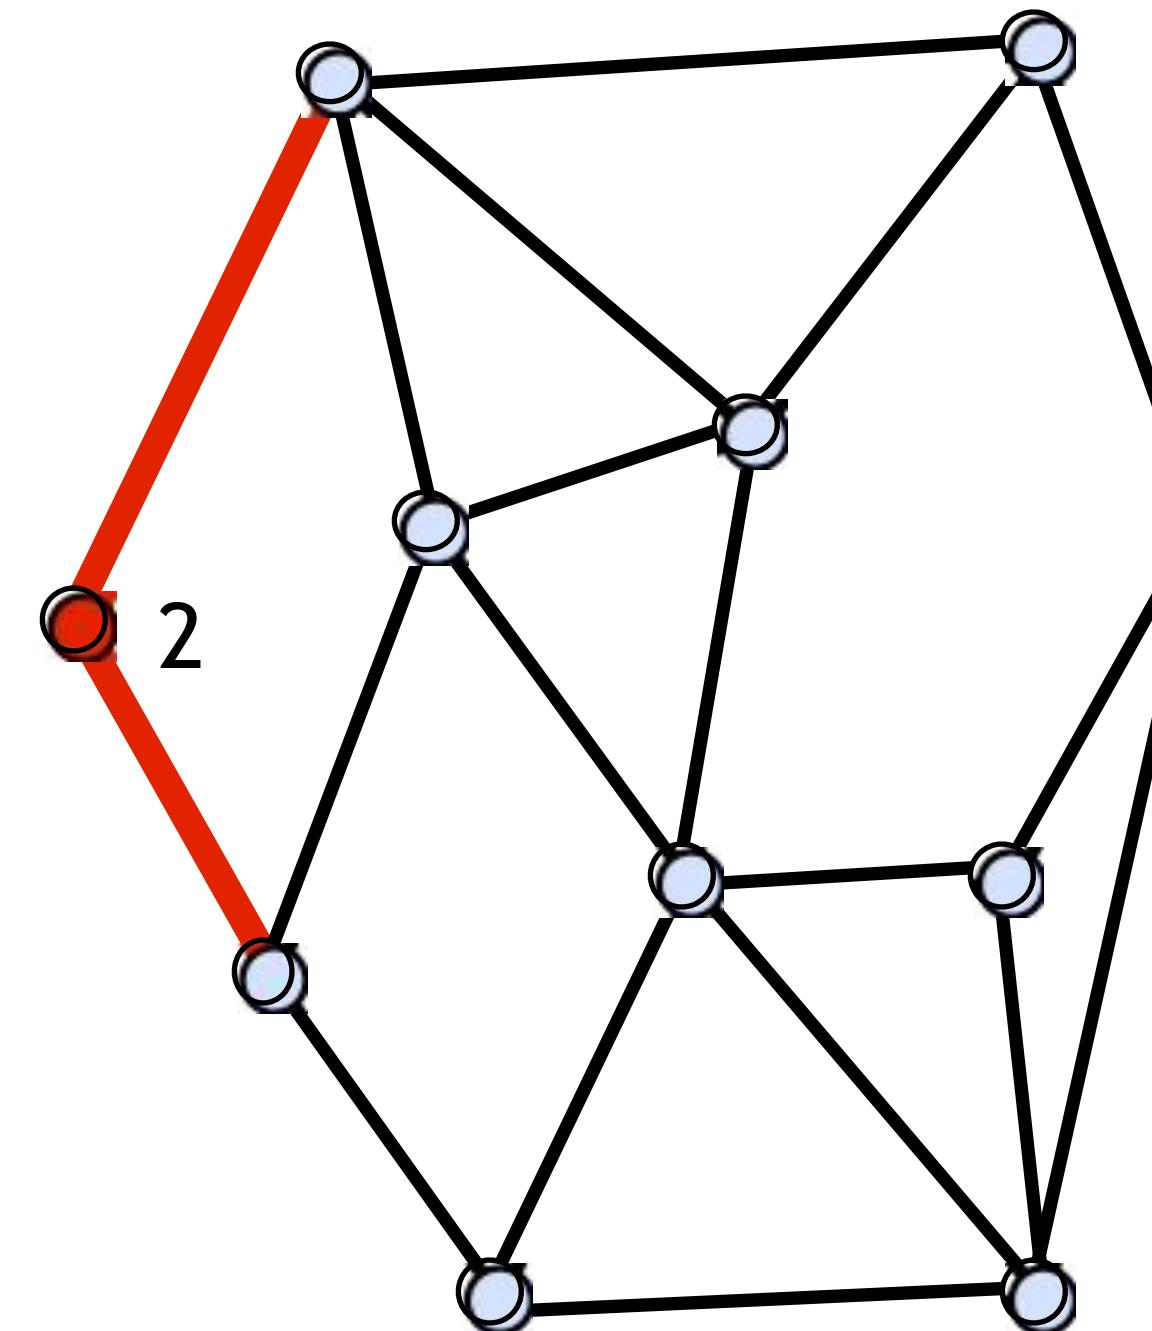
Polygonal Mesh

- Vertex **degree or valence** = number of incident edges

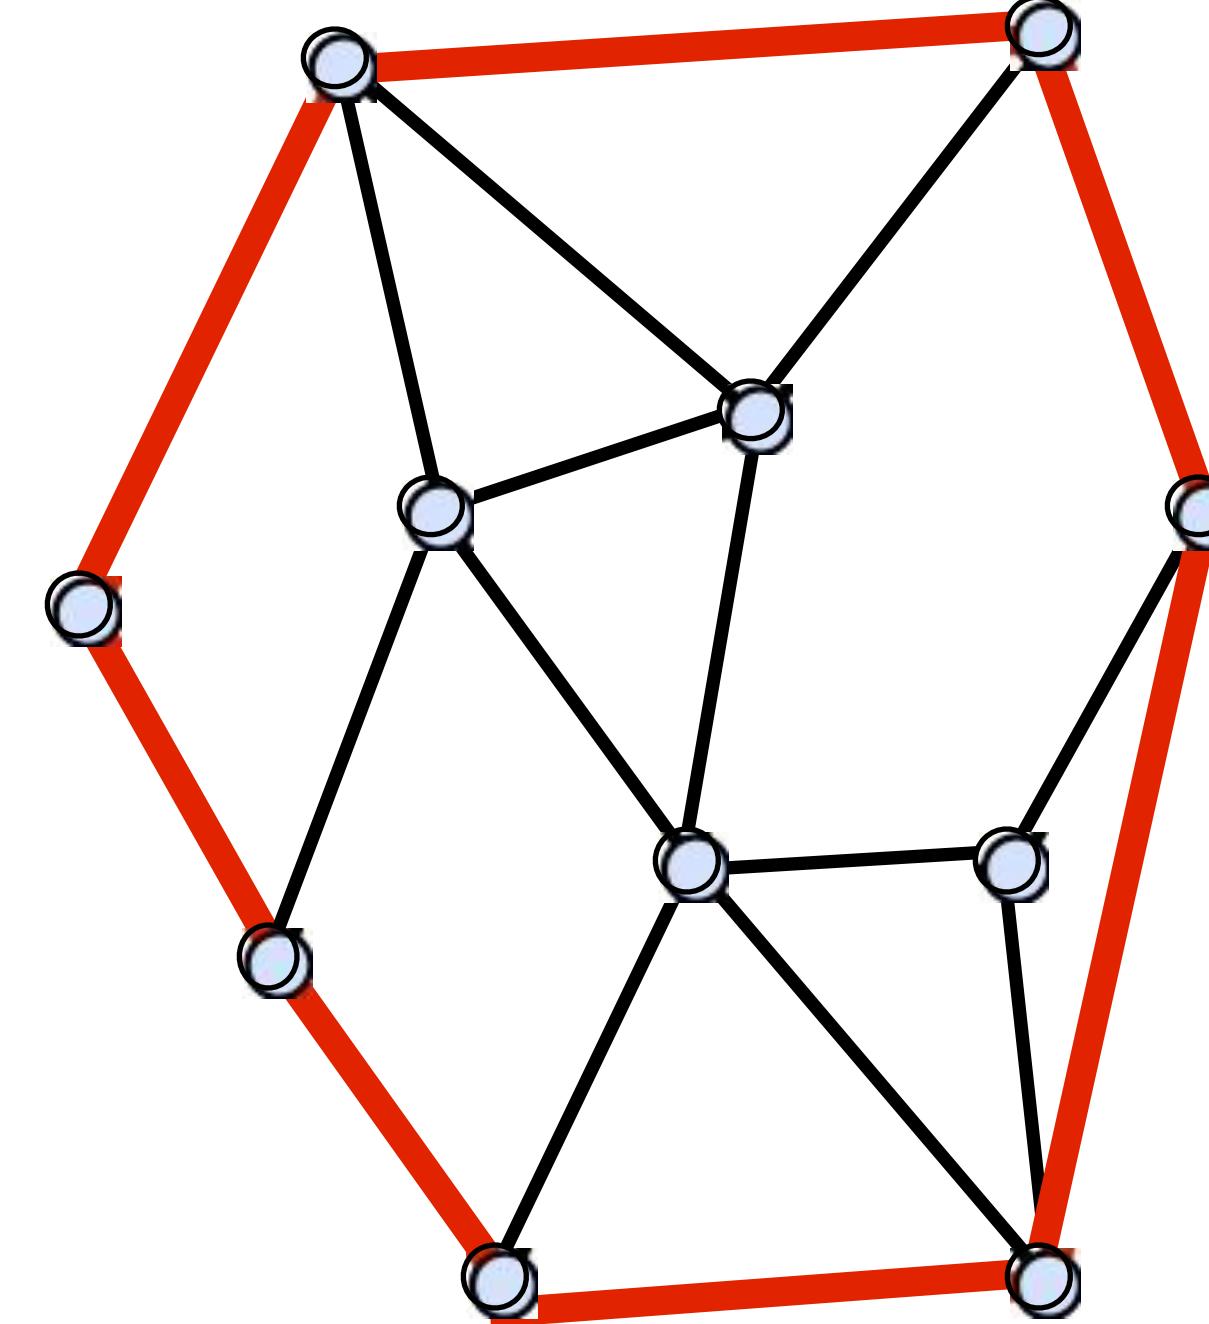


Polygonal Mesh

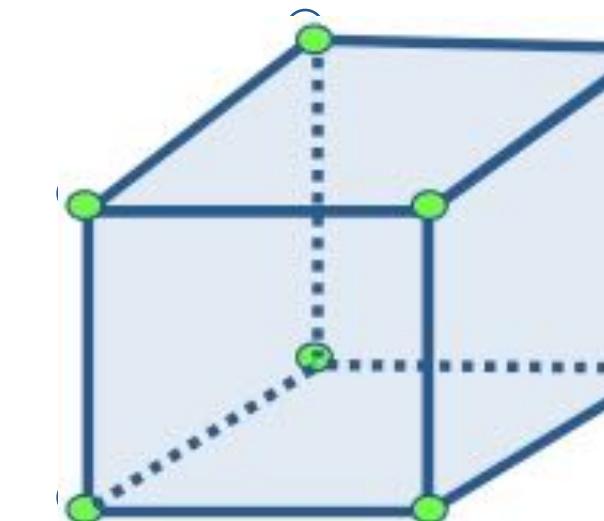
- Vertex degree or valence = number of incident edges



Polygonal Mesh



- Boundary: the set of all edges that belong to only one polygon
 - Either empty or forms closed loops
 - If empty, then the polygonal mesh is closed



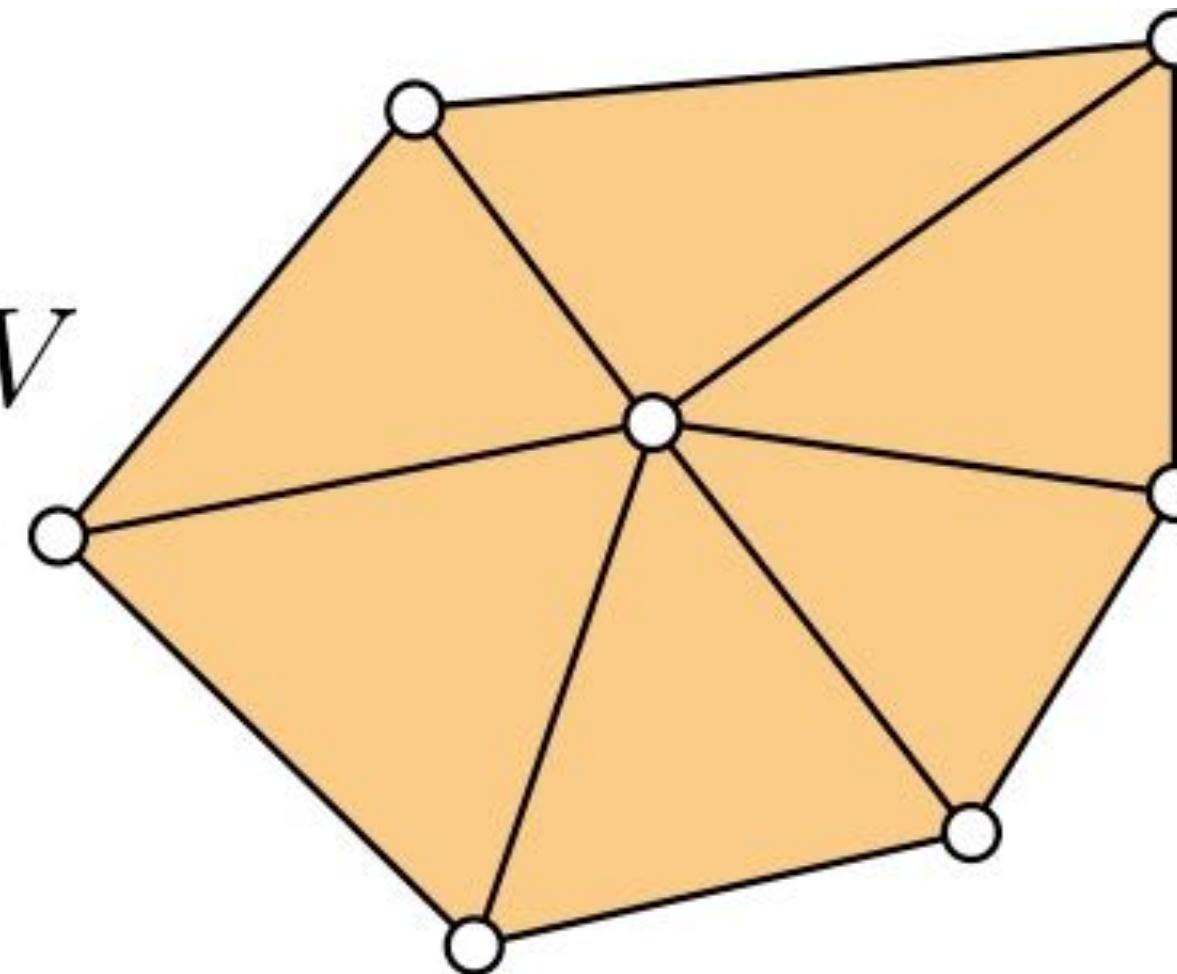
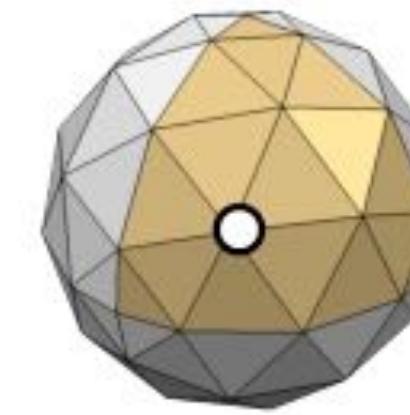
Triangle Mesh

- Connectivity: vertices, edges, triangles
- Geometry: vertex positions

$$V = \{v_1, \dots, v_n\}$$

$$E = \{e_1, \dots, e_k\}, \quad e_i \in V \times V$$

$$F = \{f_1, \dots, f_m\}, \quad f_i \in V \times V \times V$$



Triangle Mesh

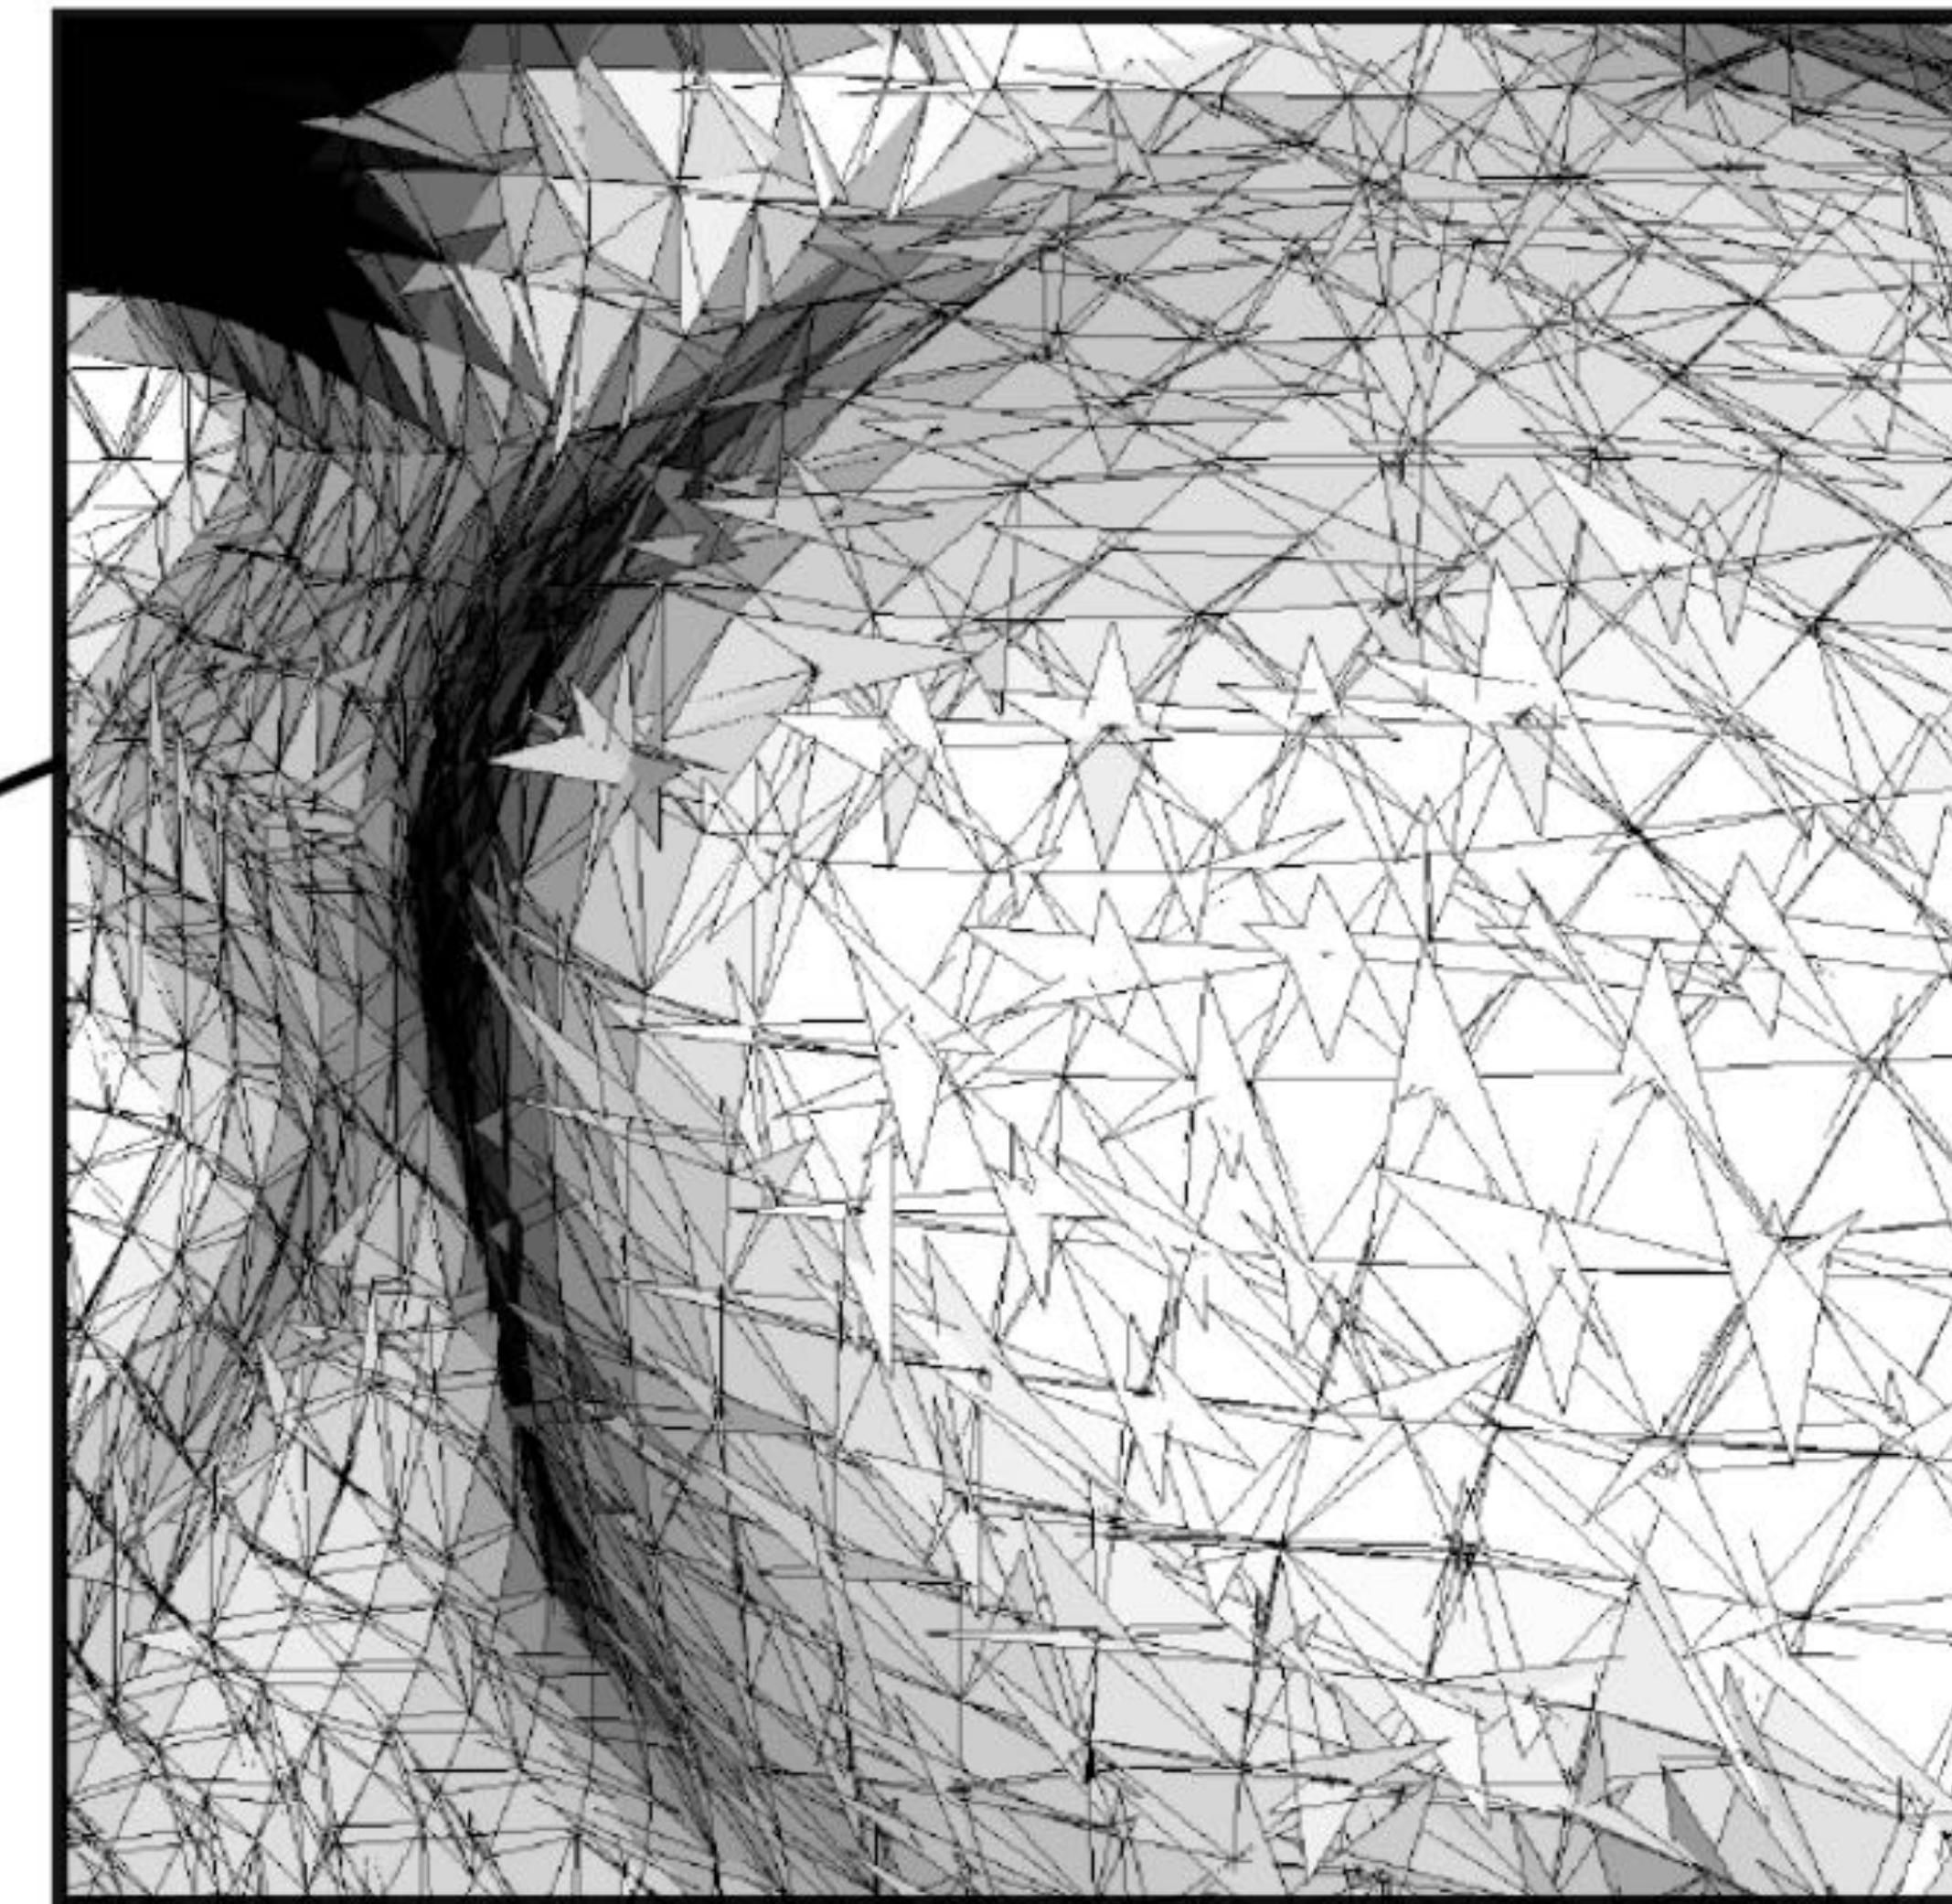
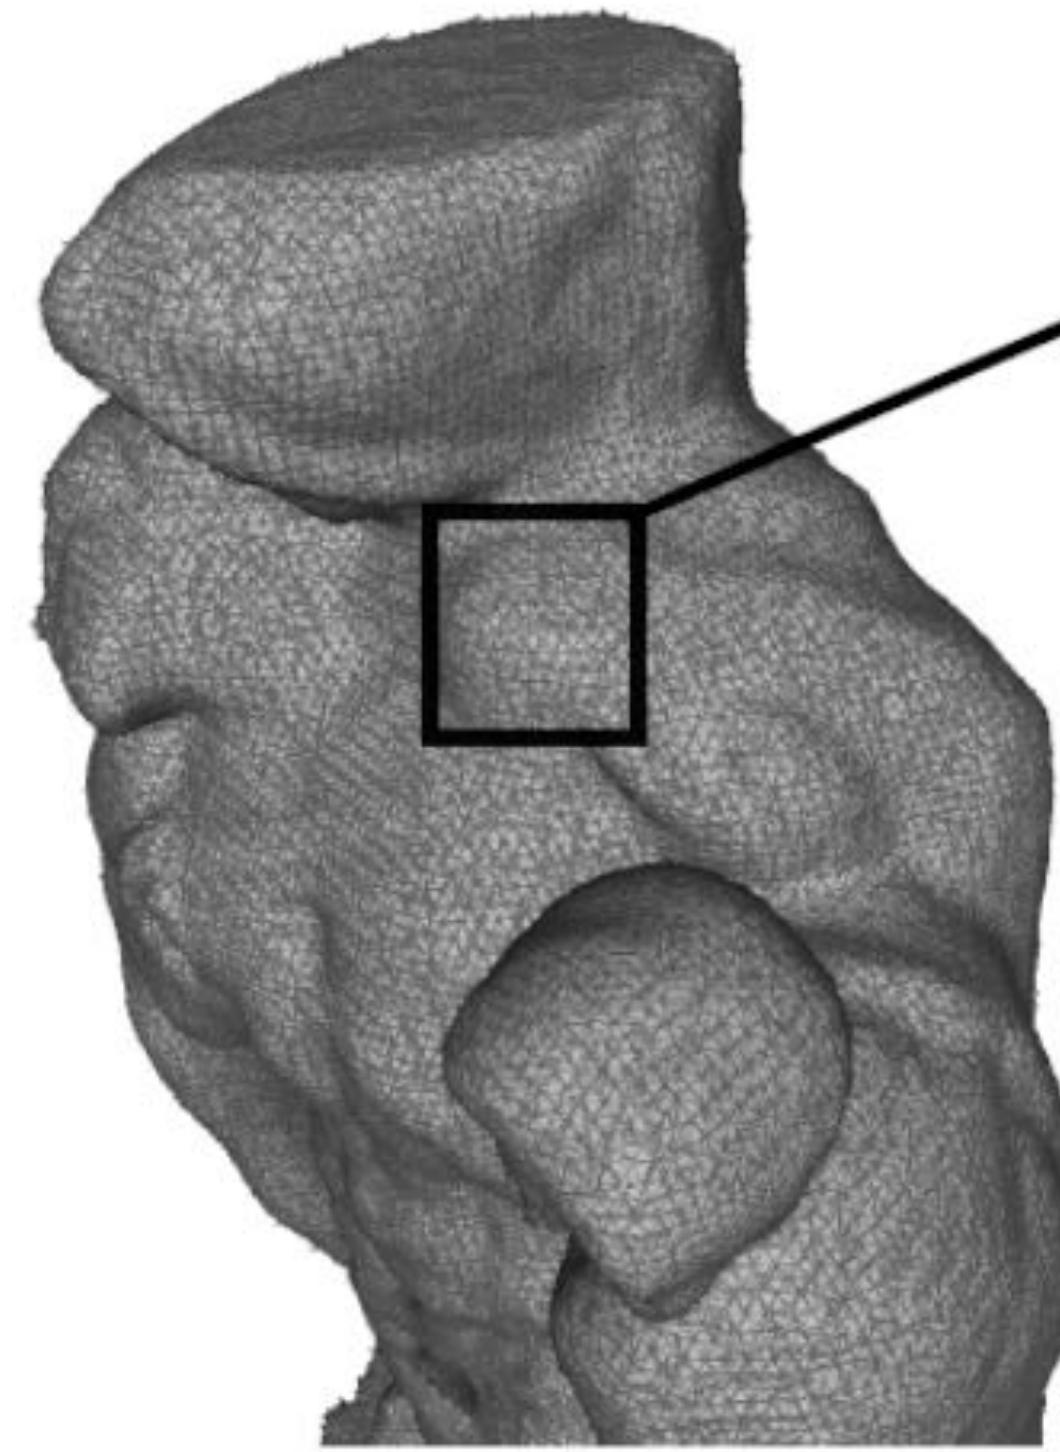
$$V = \{ v_1, v_2, \dots, v_n \} \subset \mathbb{R}^3$$

$$E = \{ e_1, e_2, \dots, e_k \} \subseteq V \times V$$

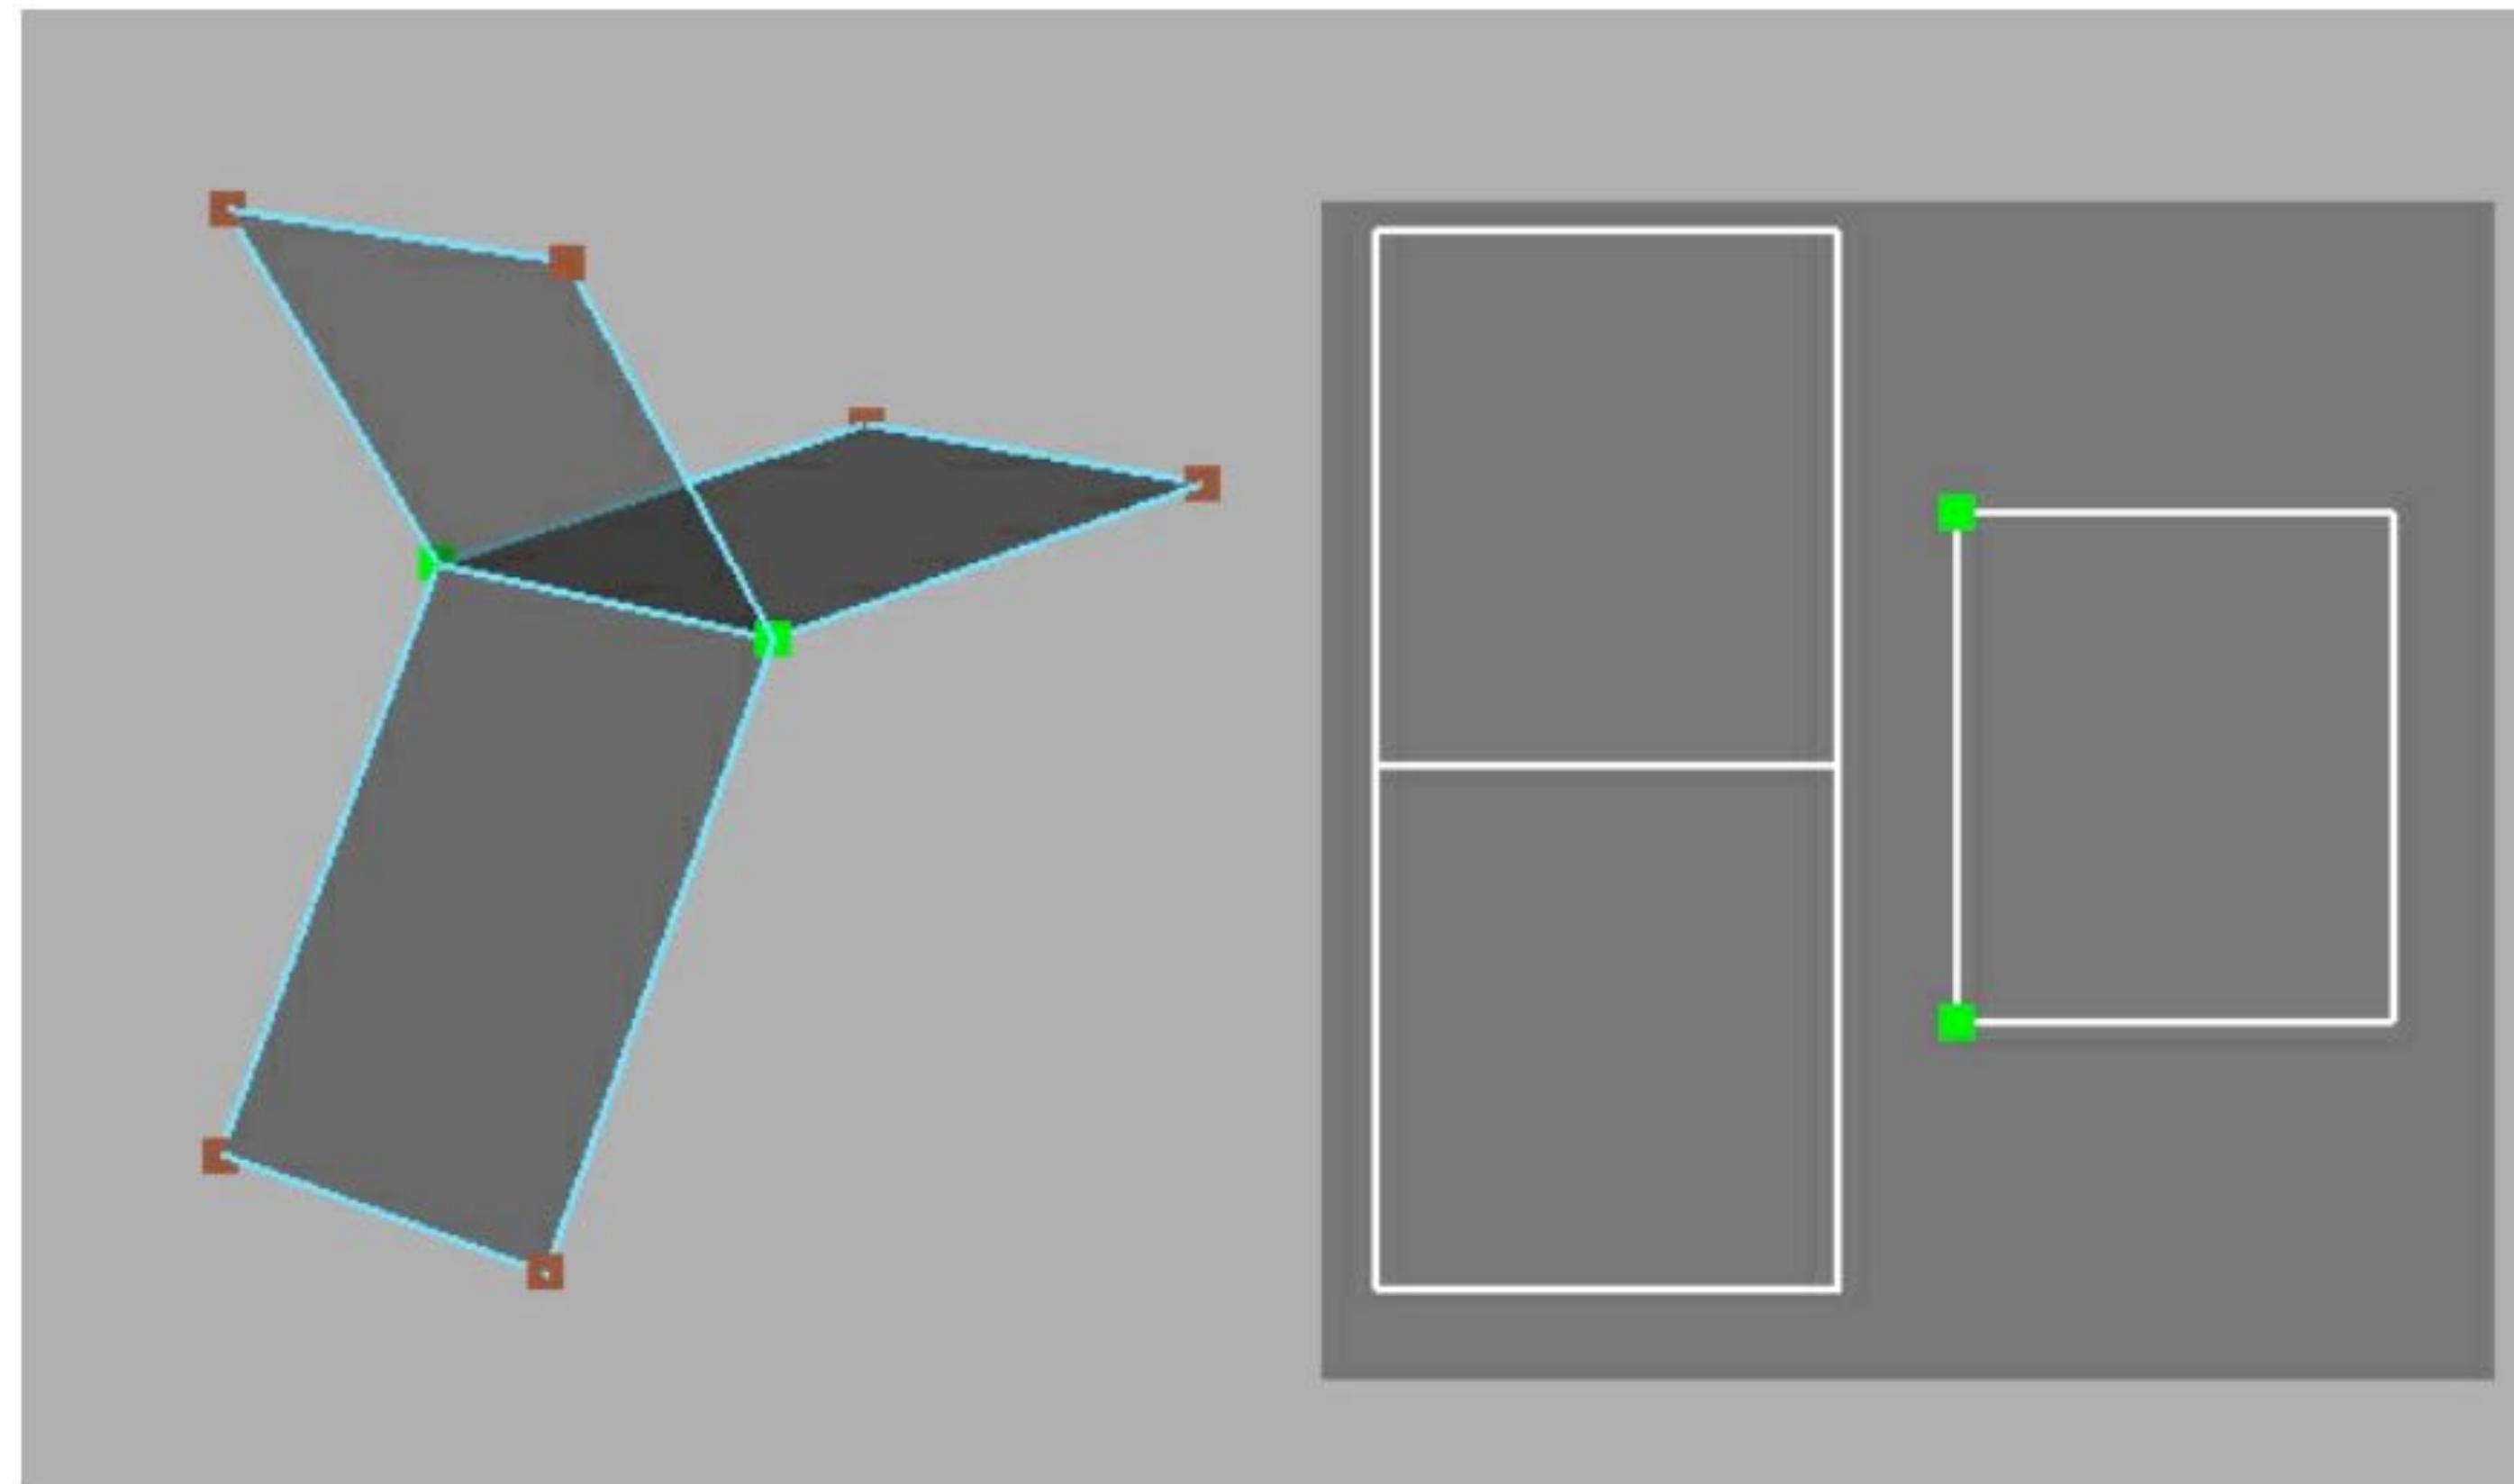
$$F = \{ f_1, f_2, \dots, f_m \} \subseteq V \times V \times V$$

Plus manifold conditions

Bad Surfaces

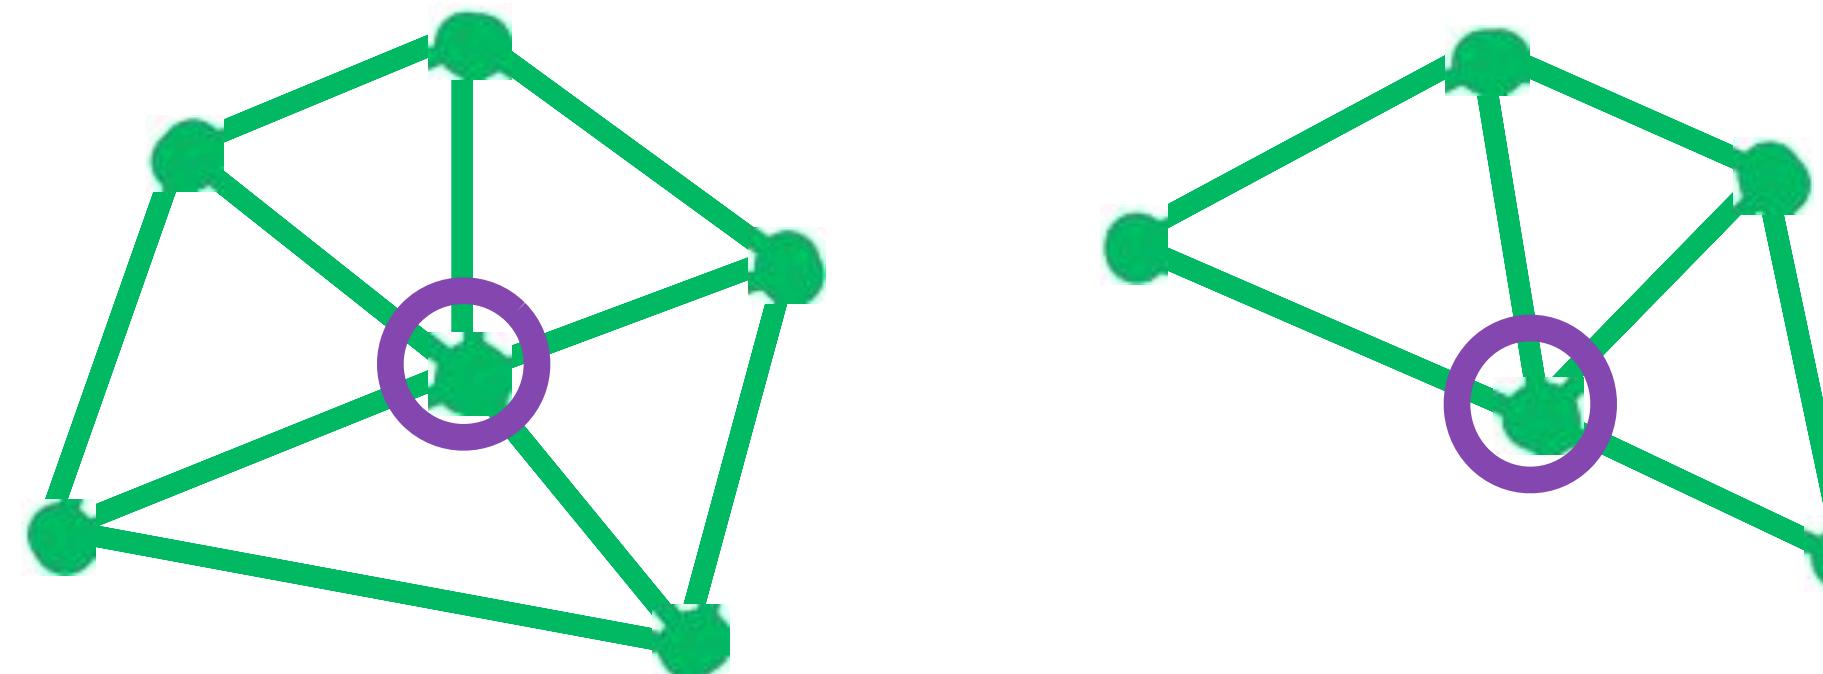


Nonmanifold Edge

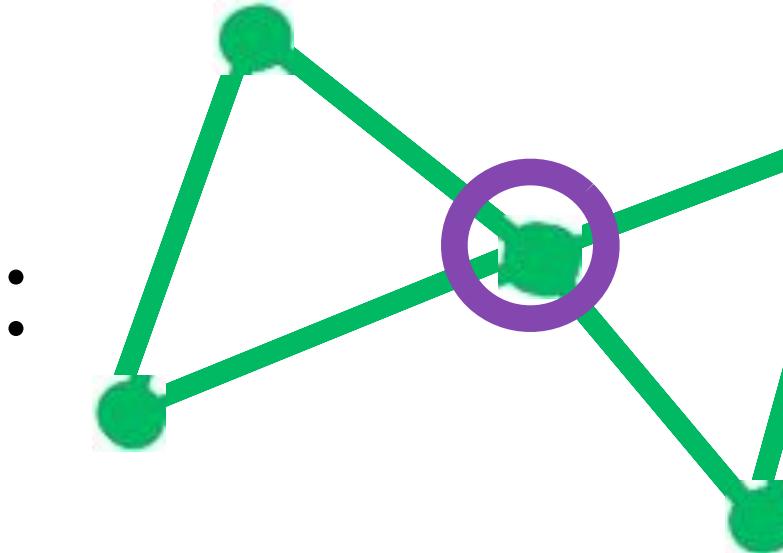


Manifold Mesh

1. Each **edge** is incident to one or two faces
2. **Faces** incident to a vertex form a closed or open fan

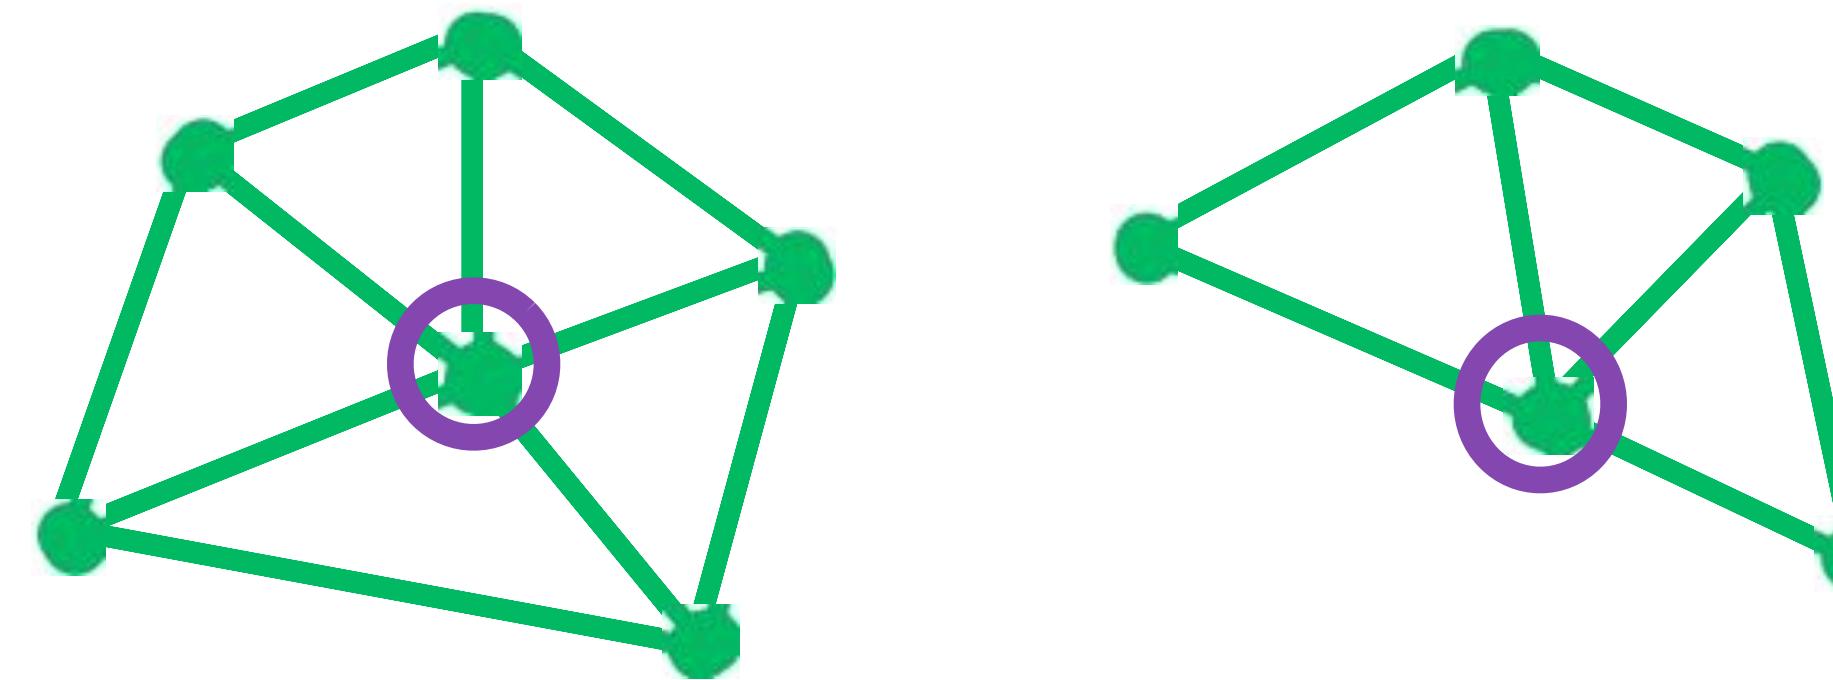


This is not a fan:

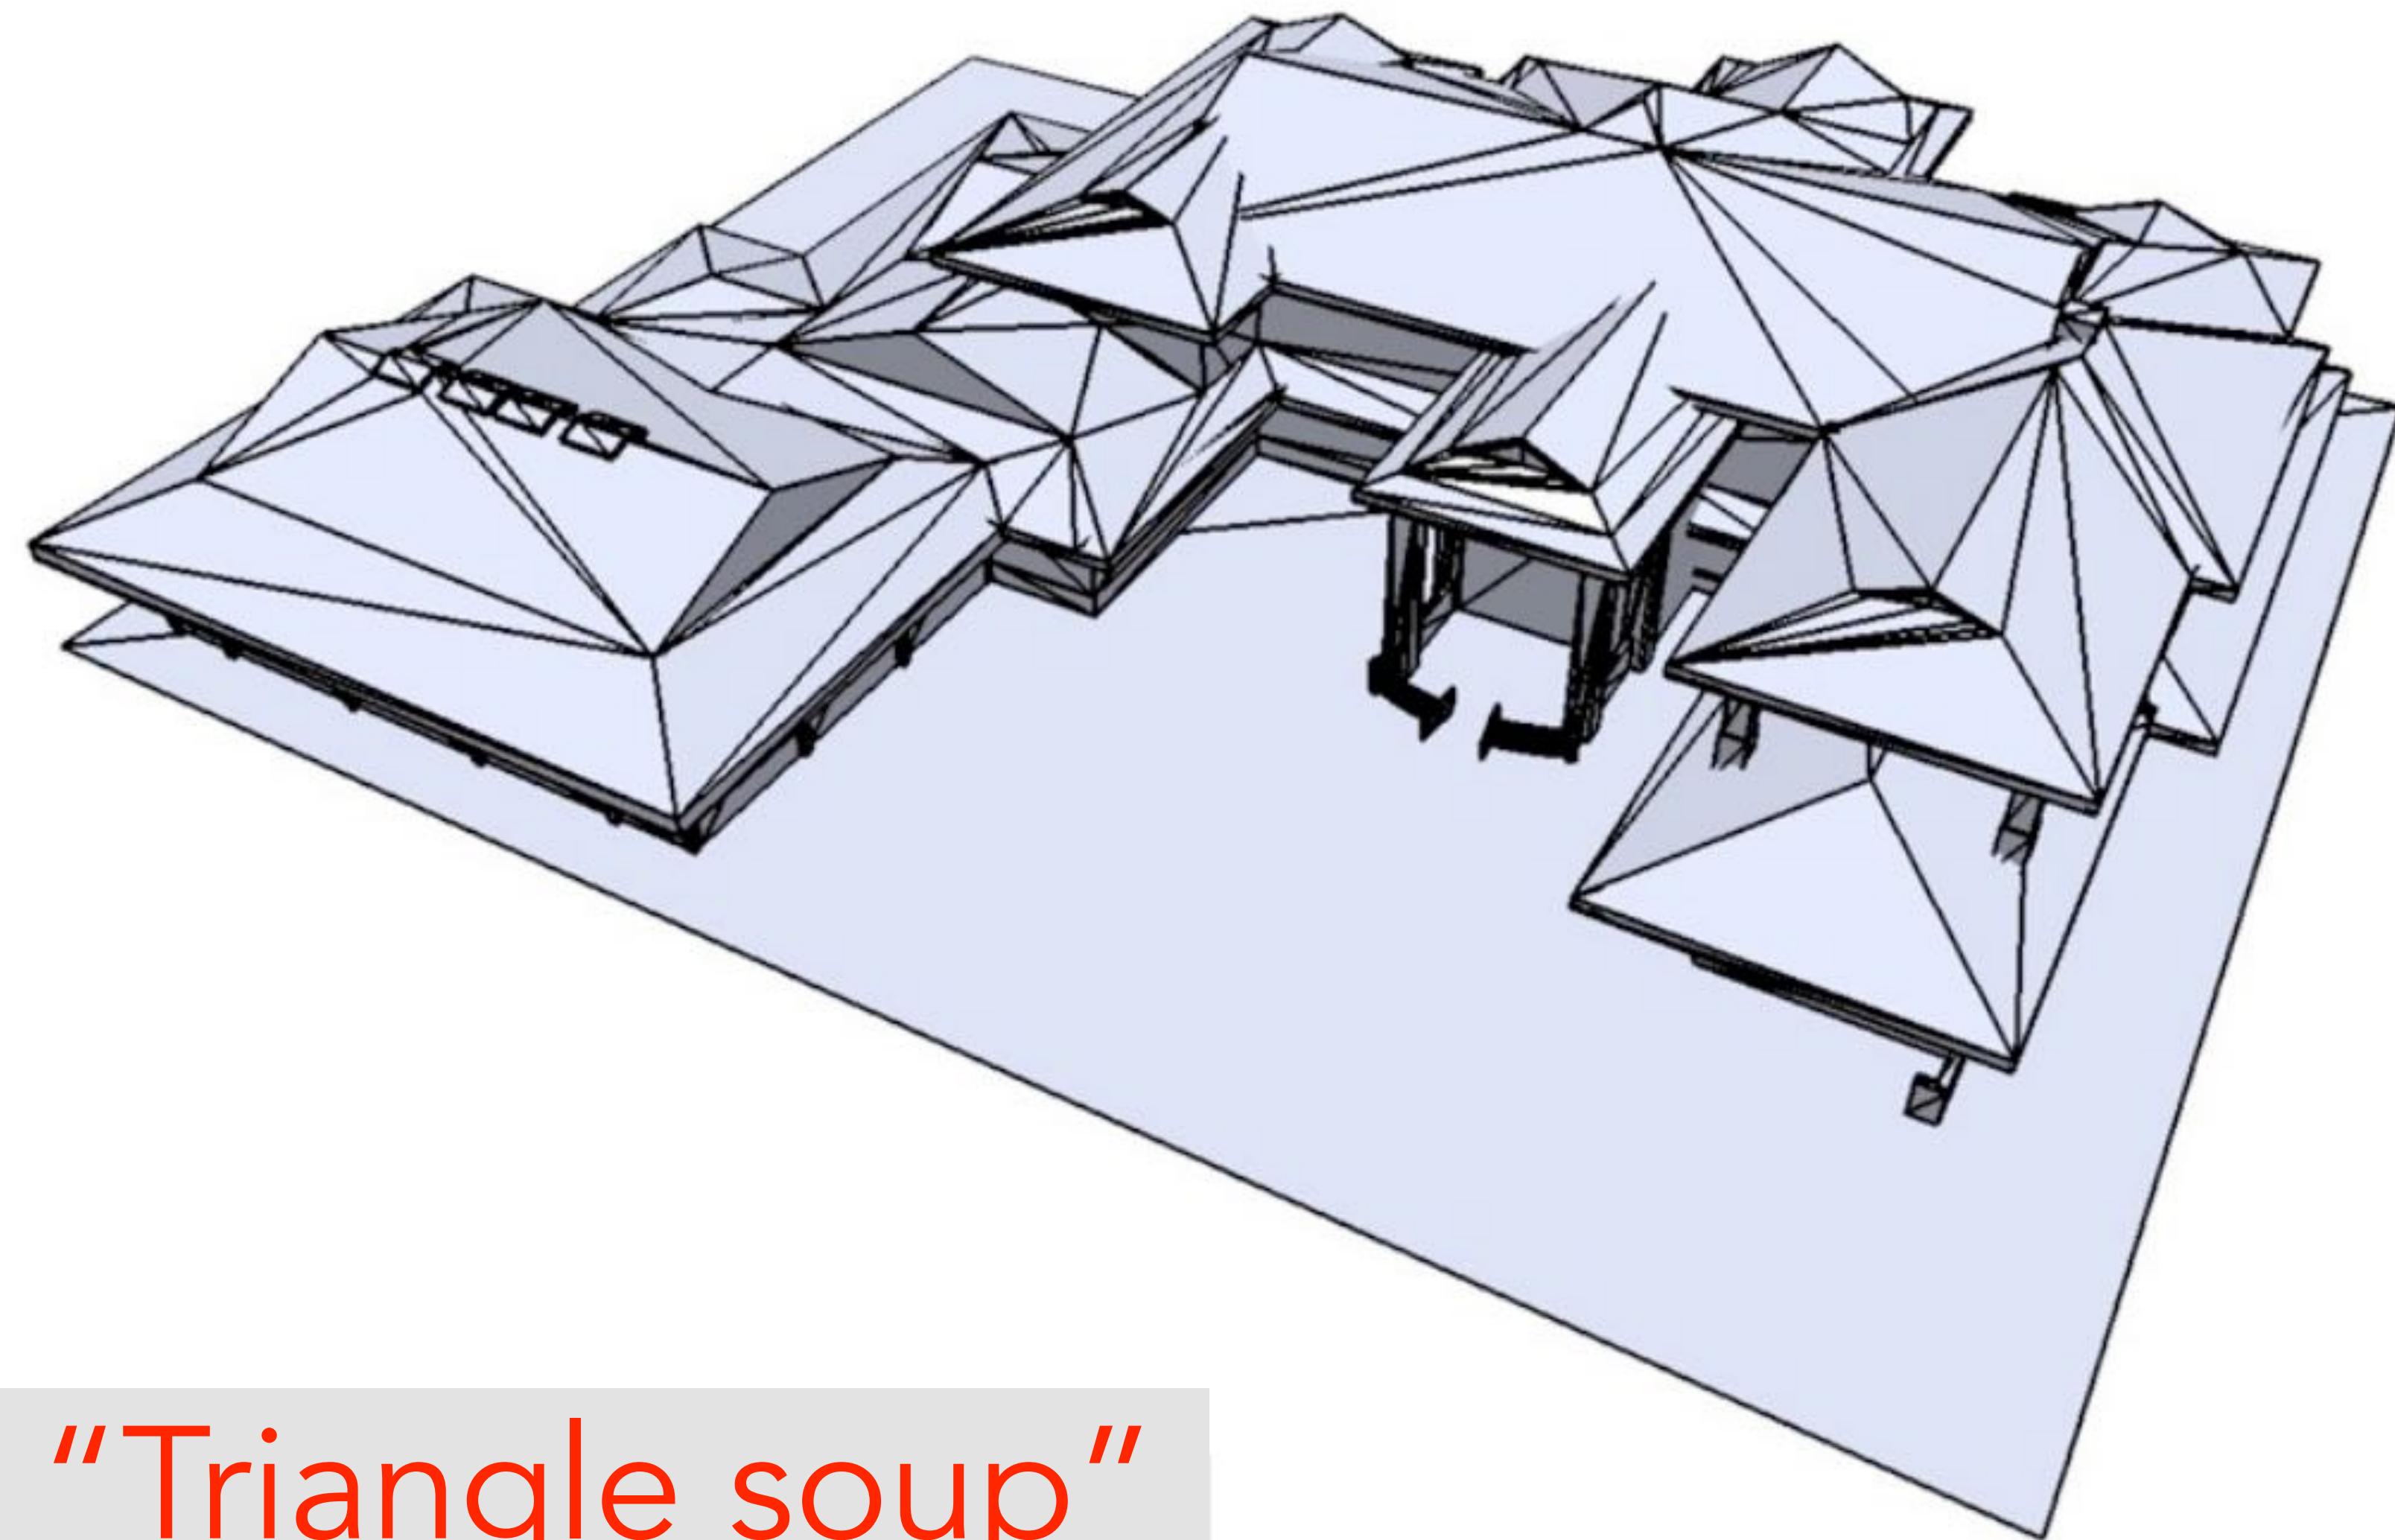


Manifold Mesh

1. Each **edge** is incident to one or two faces
2. **Faces** incident to a vertex form a closed or open fan

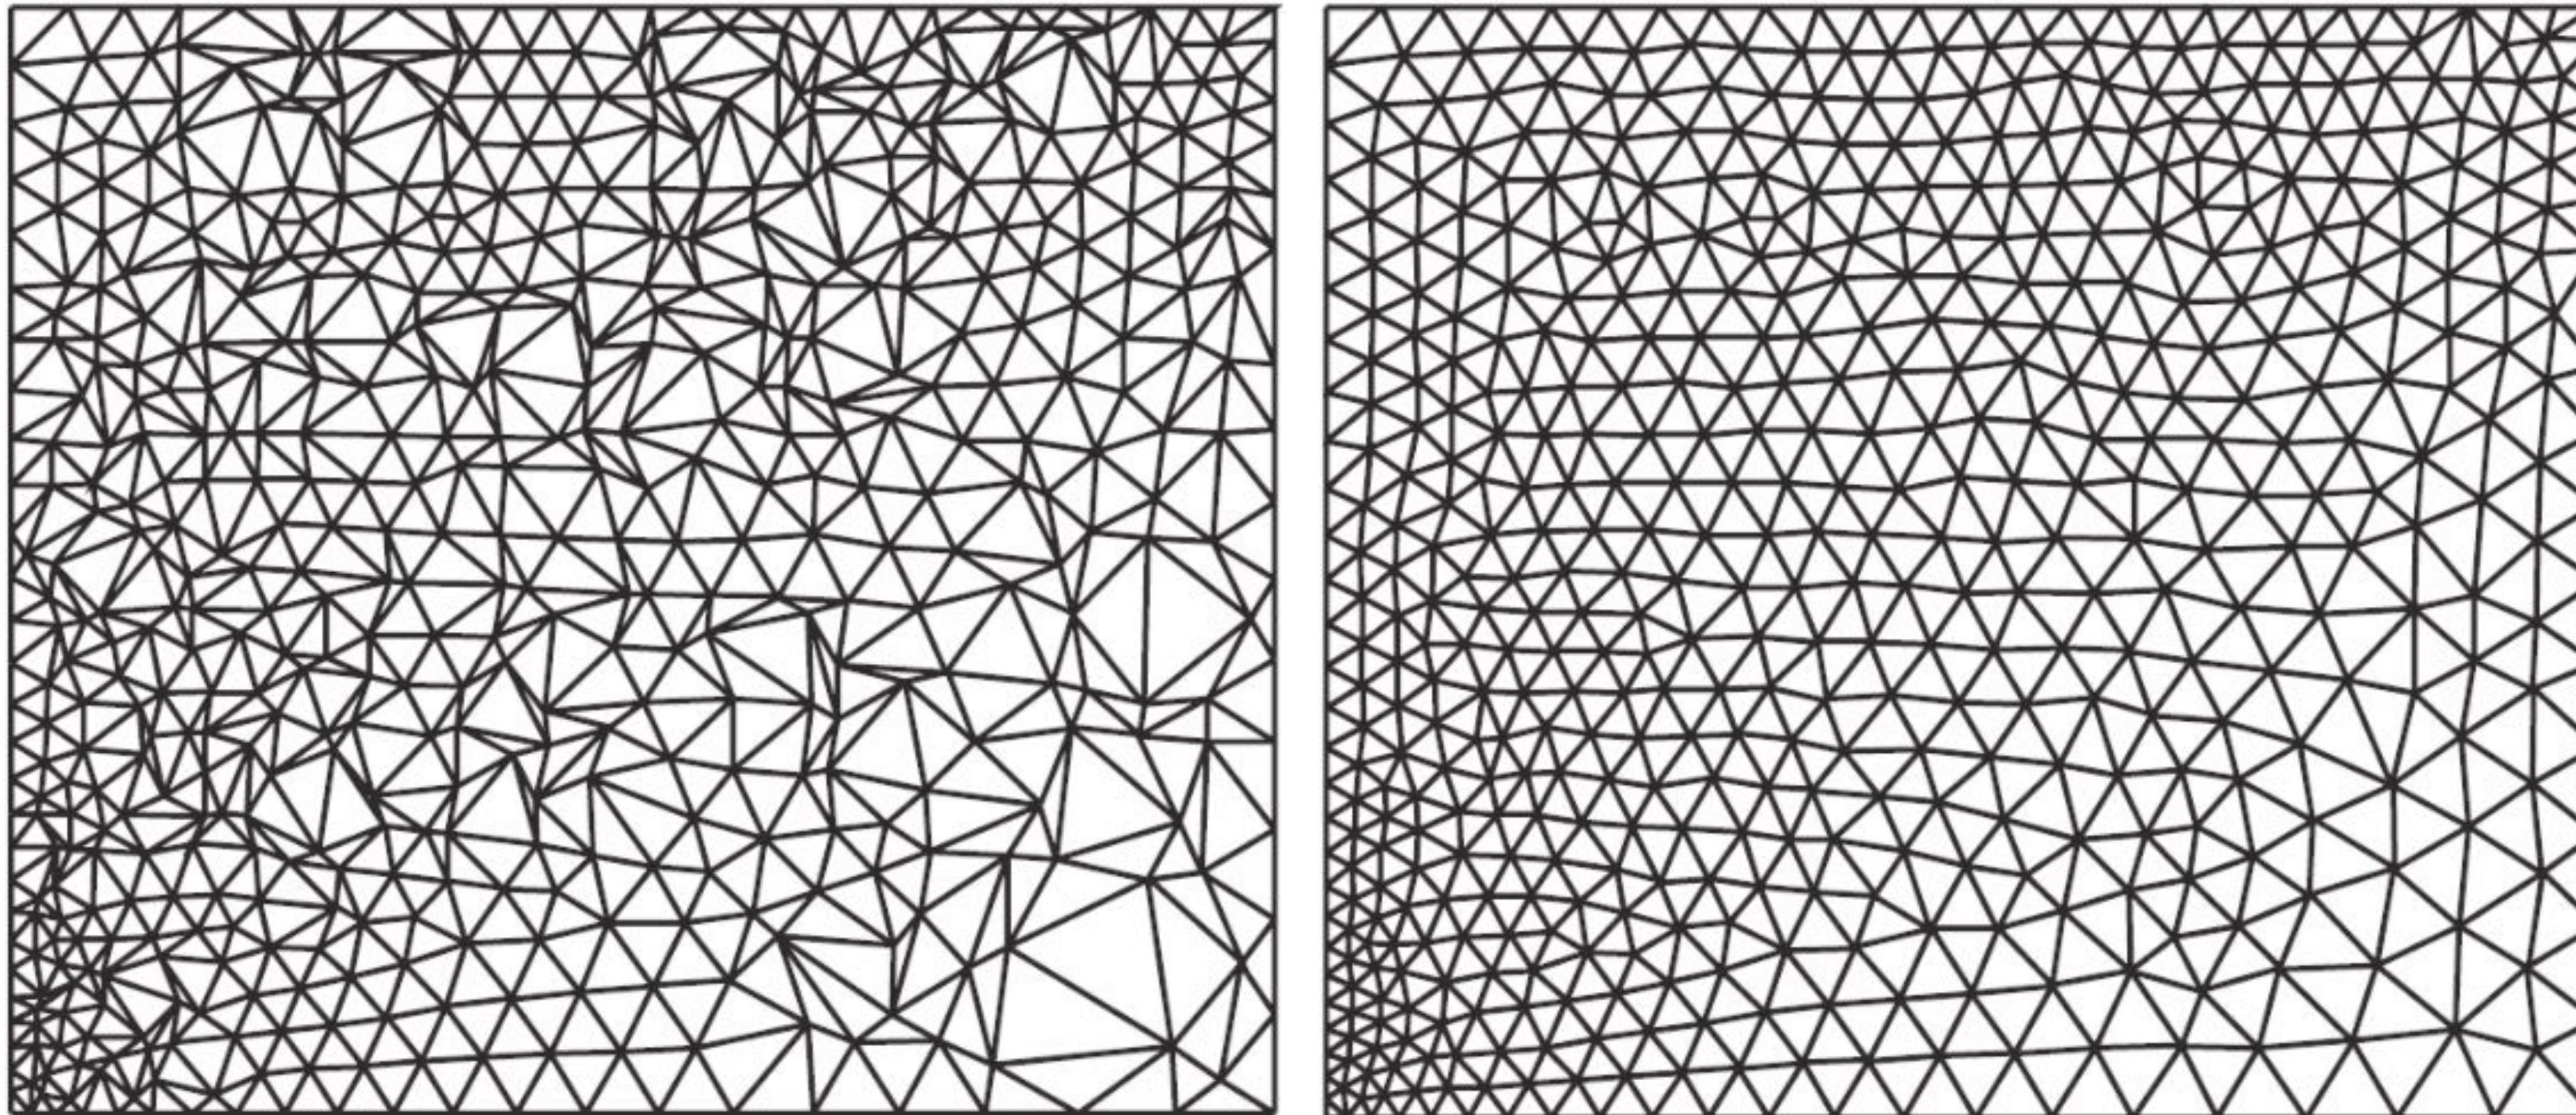


Assume meshes are manifold
(for now)



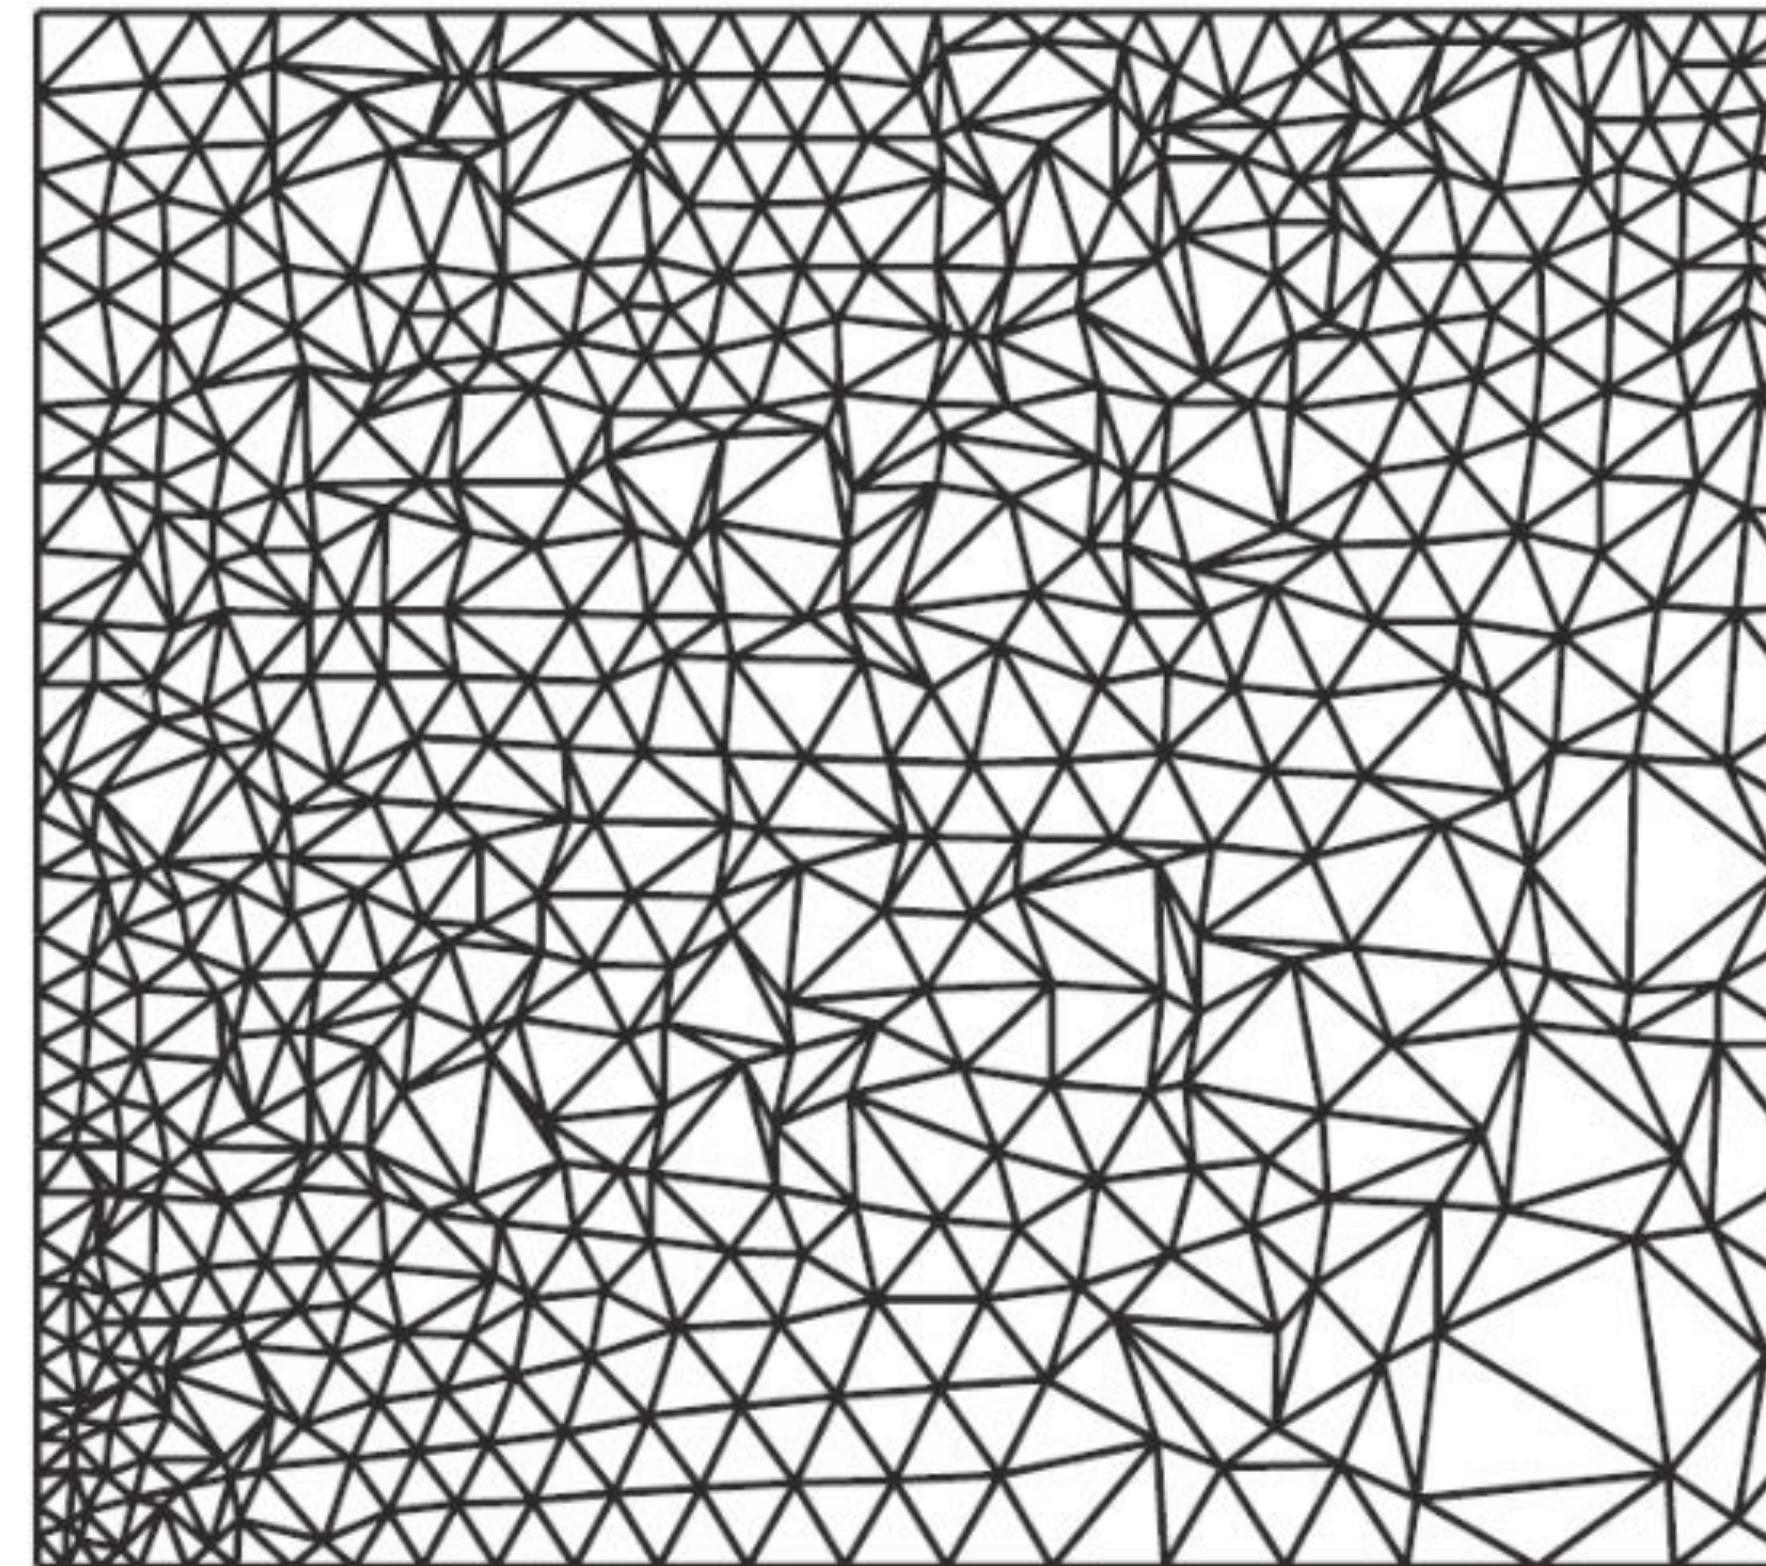
“Triangle soup”

Bad Meshes



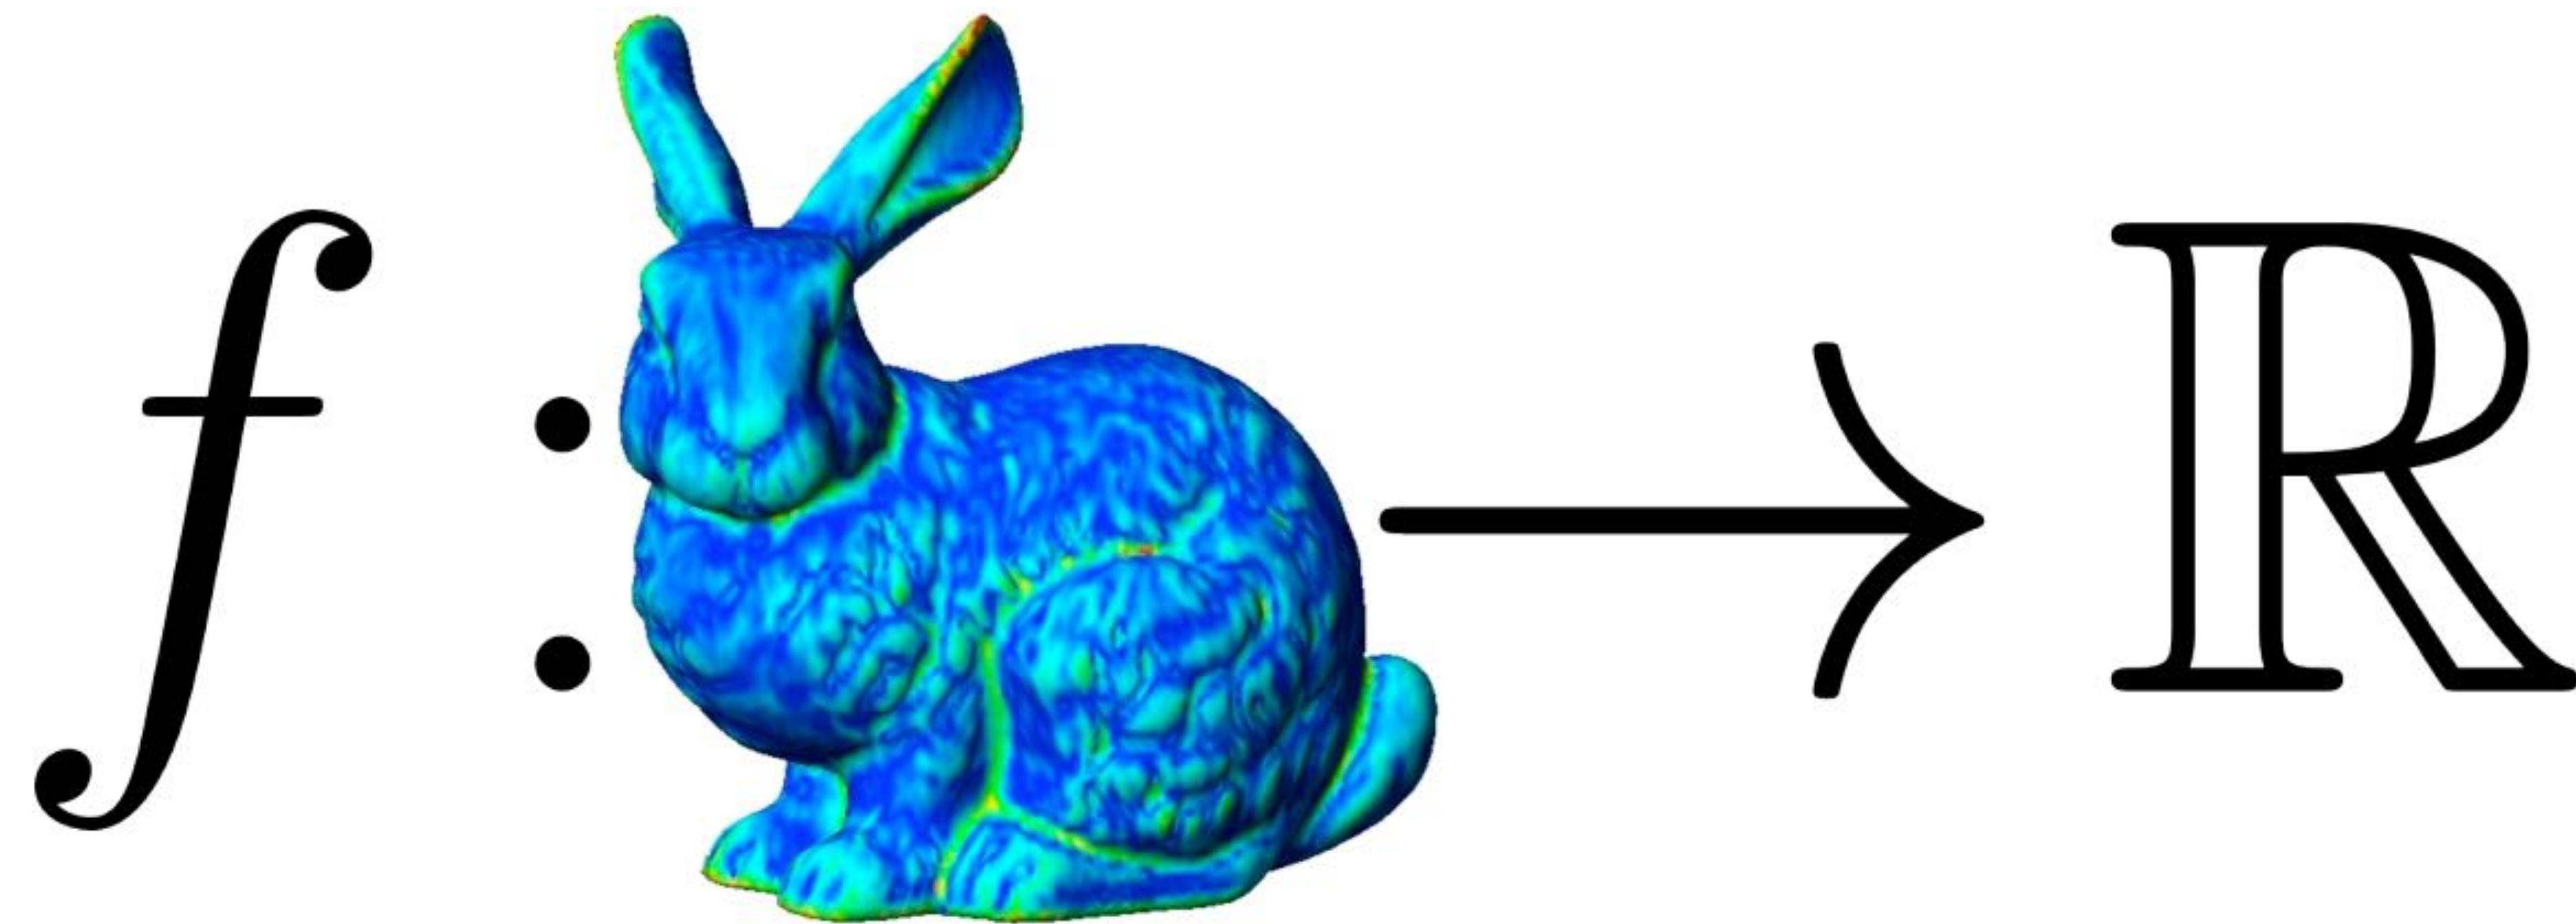
Nonuniform
areas and angles

Why is Meshing an Issue?



How do you interpret
one value per vertex?

Assume Storing Scalar Functions on Surface

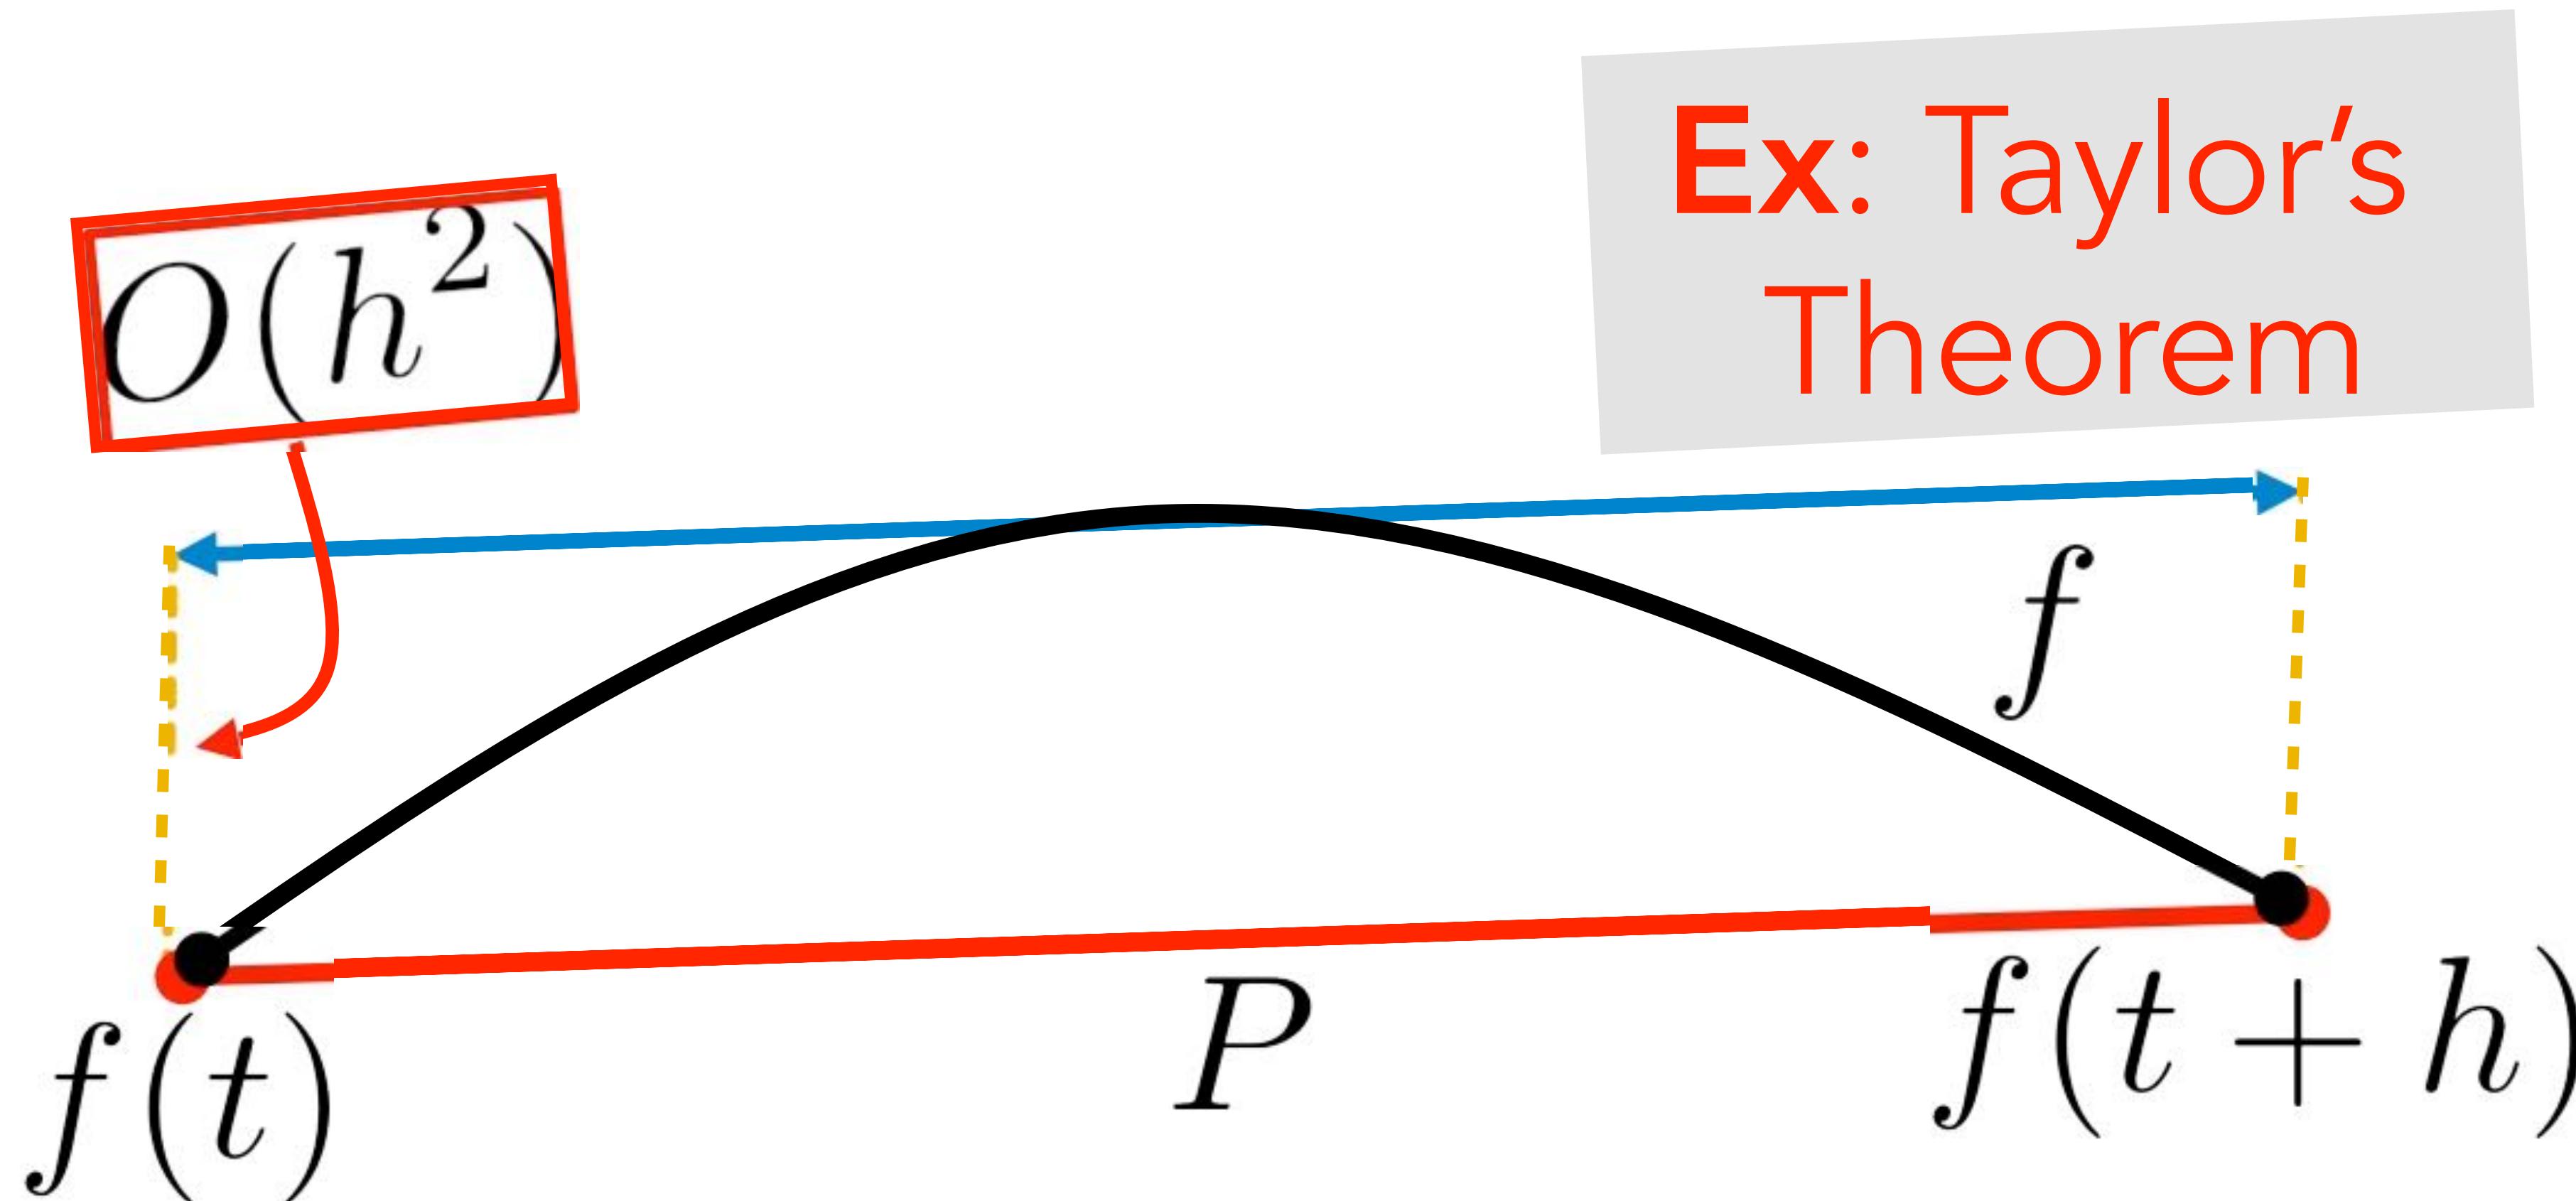


http://www.ieeta.pt/polymeco/Screenshots/PolyMeCo_OneView.jpg

Map points to real numbers

Slides credit: Hao Su

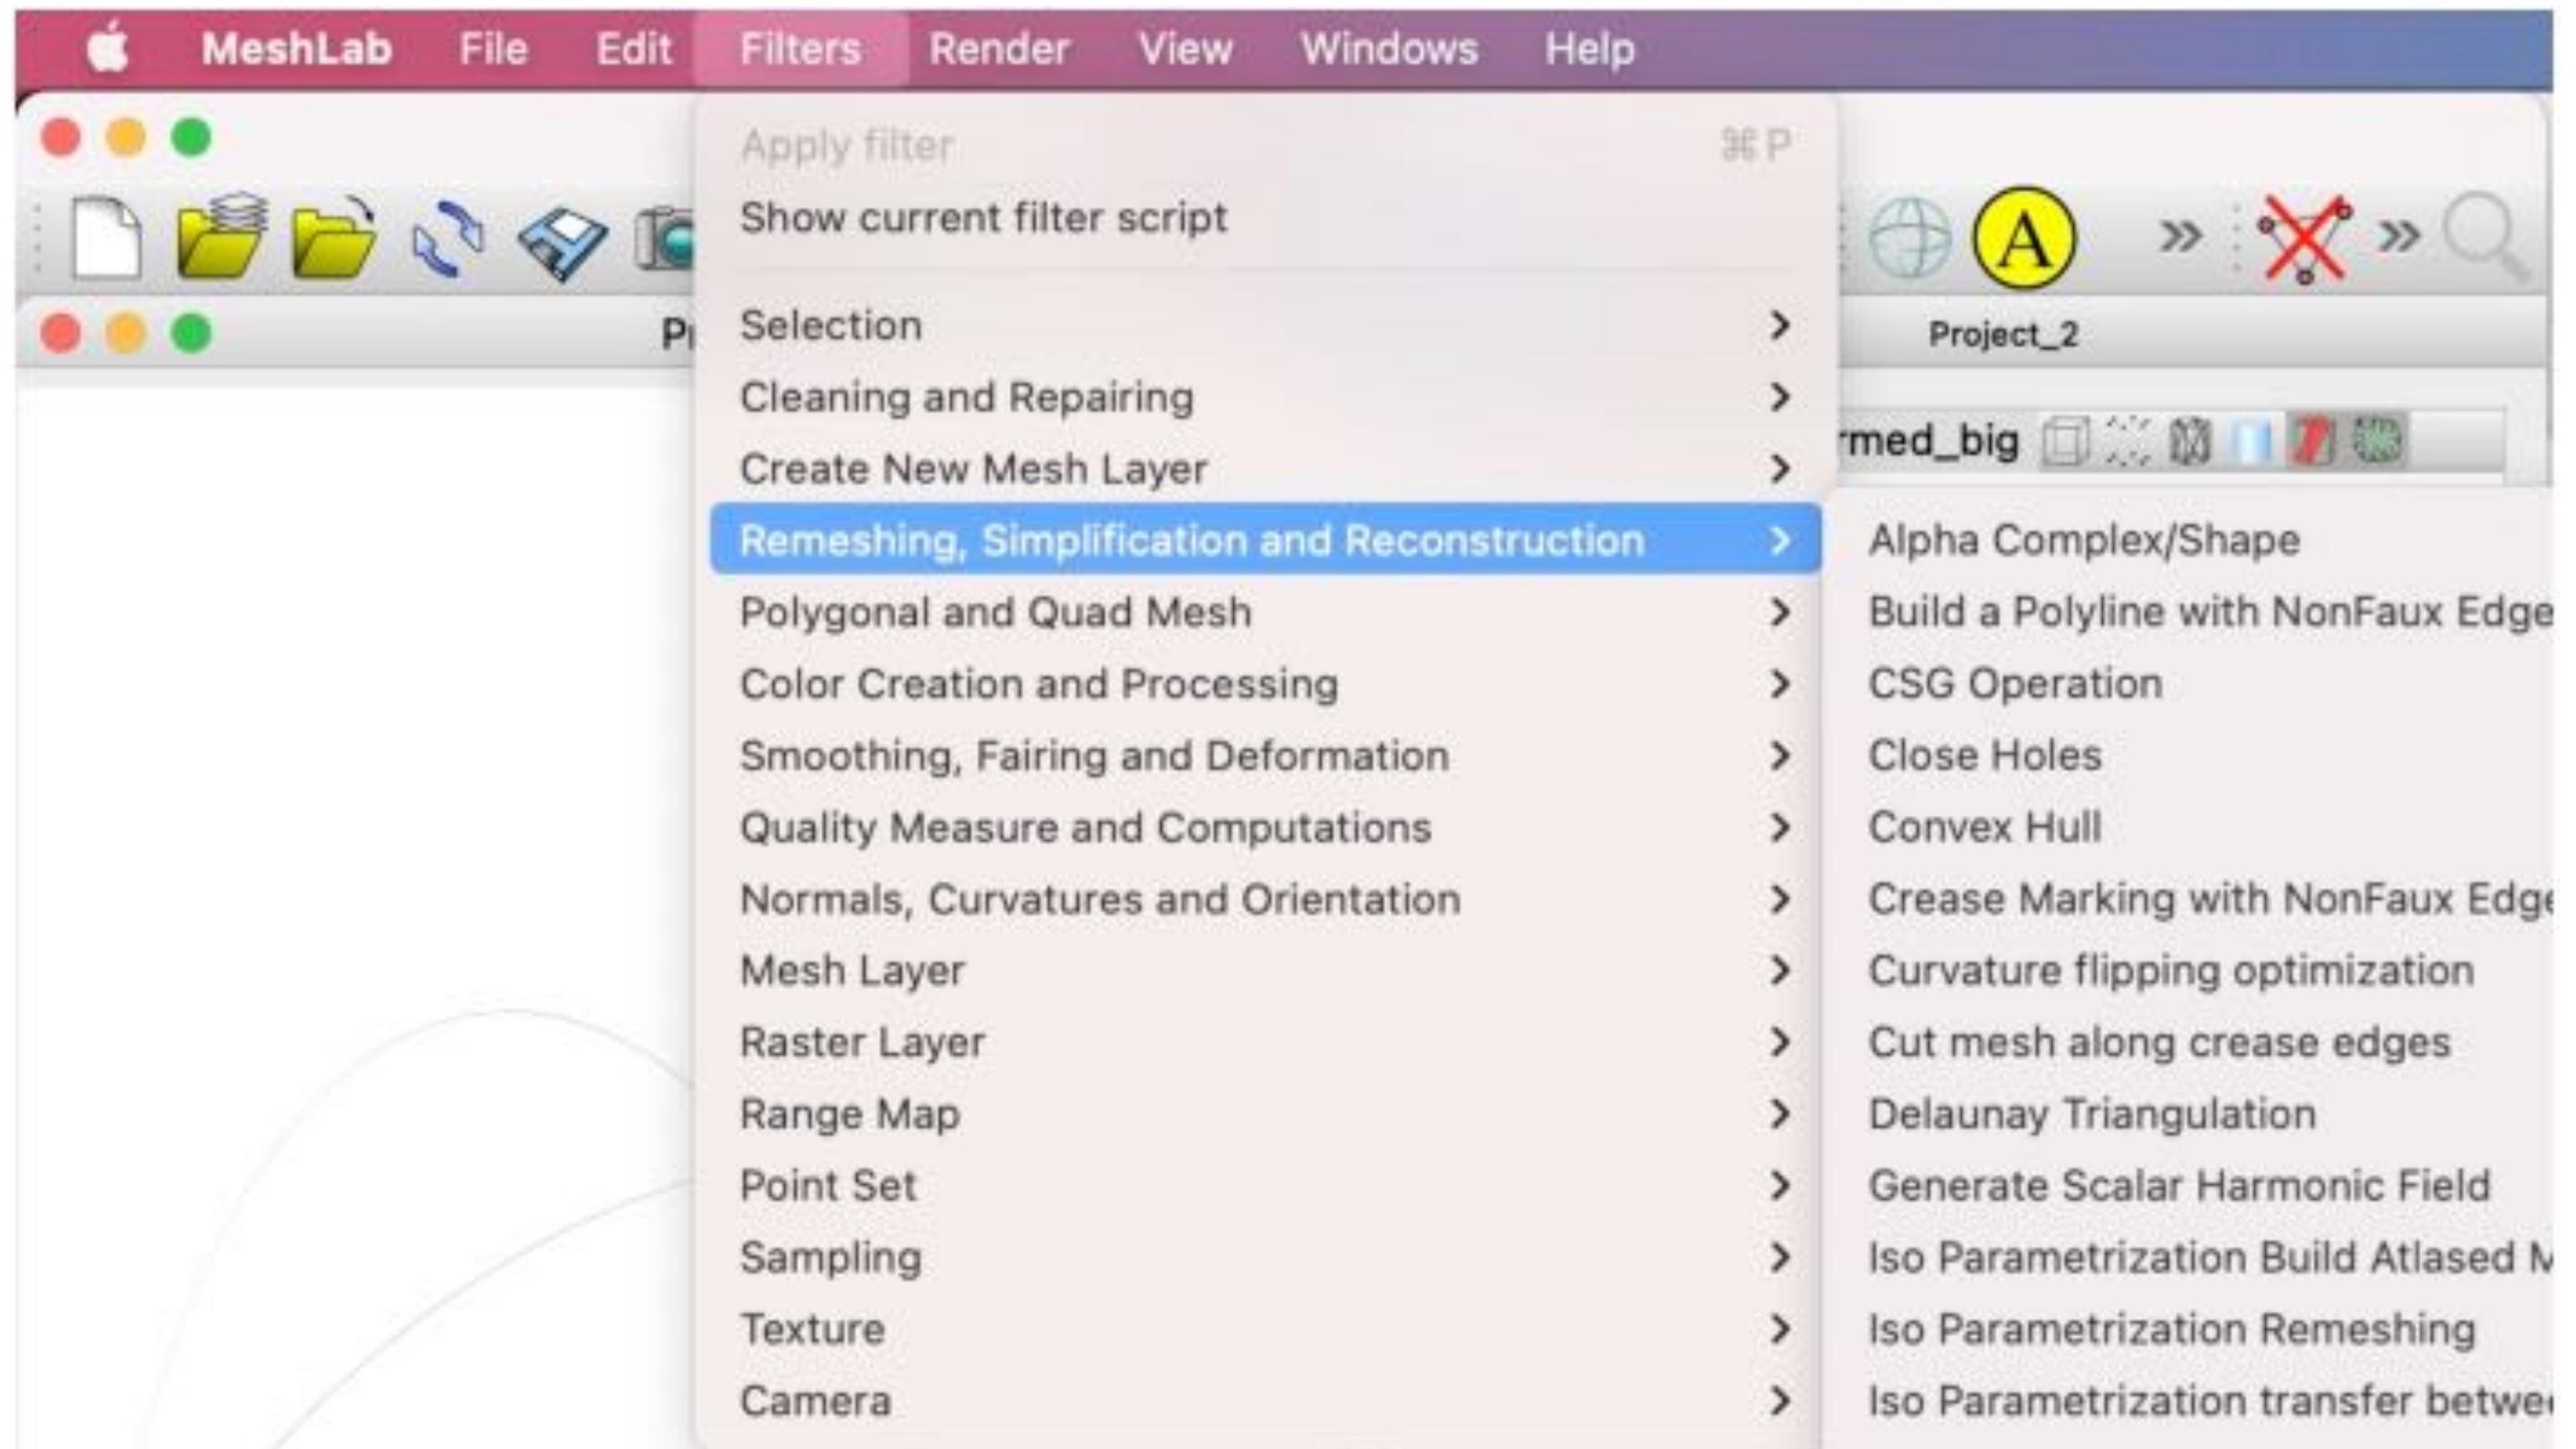
Approximation Properties



f : functions defined at vertices
(e.g., Gaussian curvature)

Techniques to Improve Mesh Quality

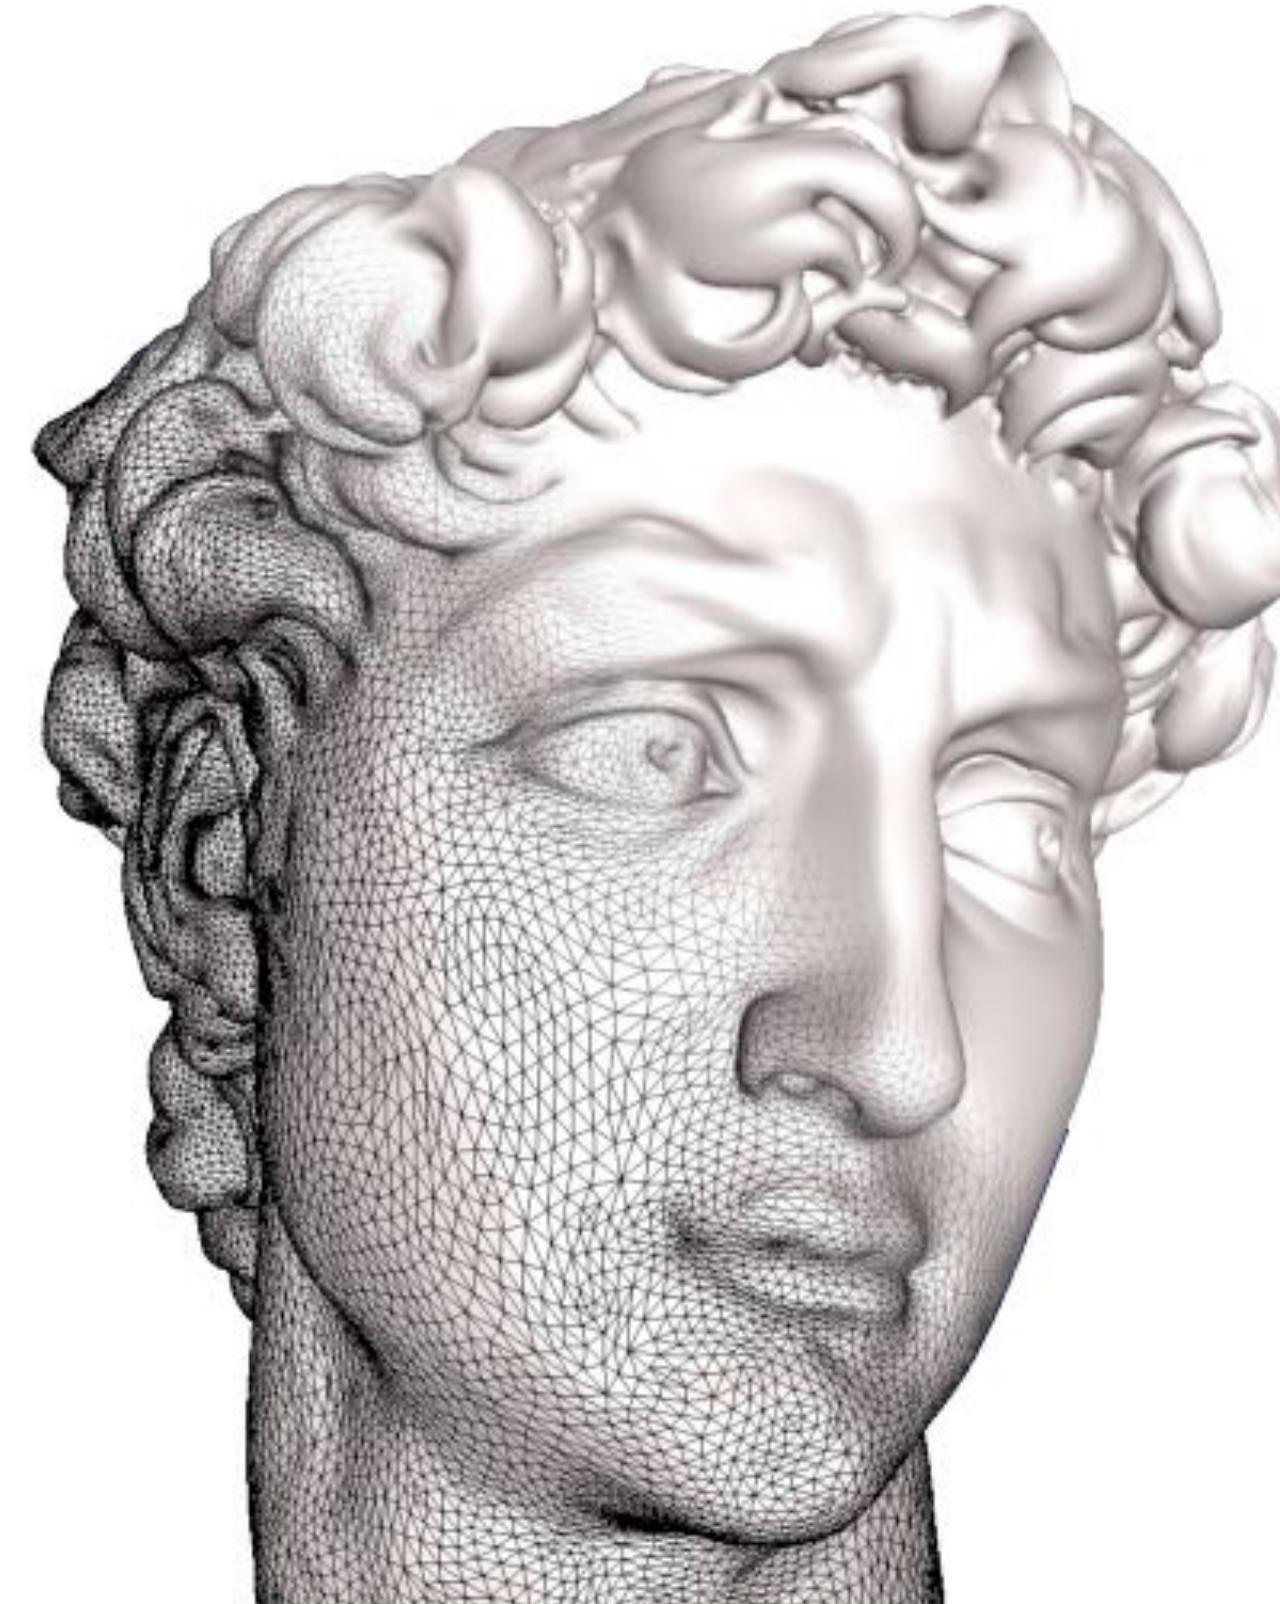
- Cleaning
- Repairing
- Remeshing
- ...



Outline

- Rasterized 3D representations
- *Mesh*
 - Representation
 - *Storage*
 - Curvature computation
- Point cloud
- Implicit representations

Data Structures for Surfaces



What should be stored?

- Geometry: 3D coordinates
- Topology
- Attributes
 - Normal, color, texture coordinates
 - Per vertex, face, edge

Simple Data Structures: Triangle List

- STL format (used in CAD)
- Storage
 - Face: 3 positions
- No connectivity information

Triangles			
0	x0	y0	z0
1	x1	x1	z1
2	x2	y2	z2
3	x3	y3	z3
4	x4	y4	z4
5	x5	y5	z5
6	x6	y6	z6
...

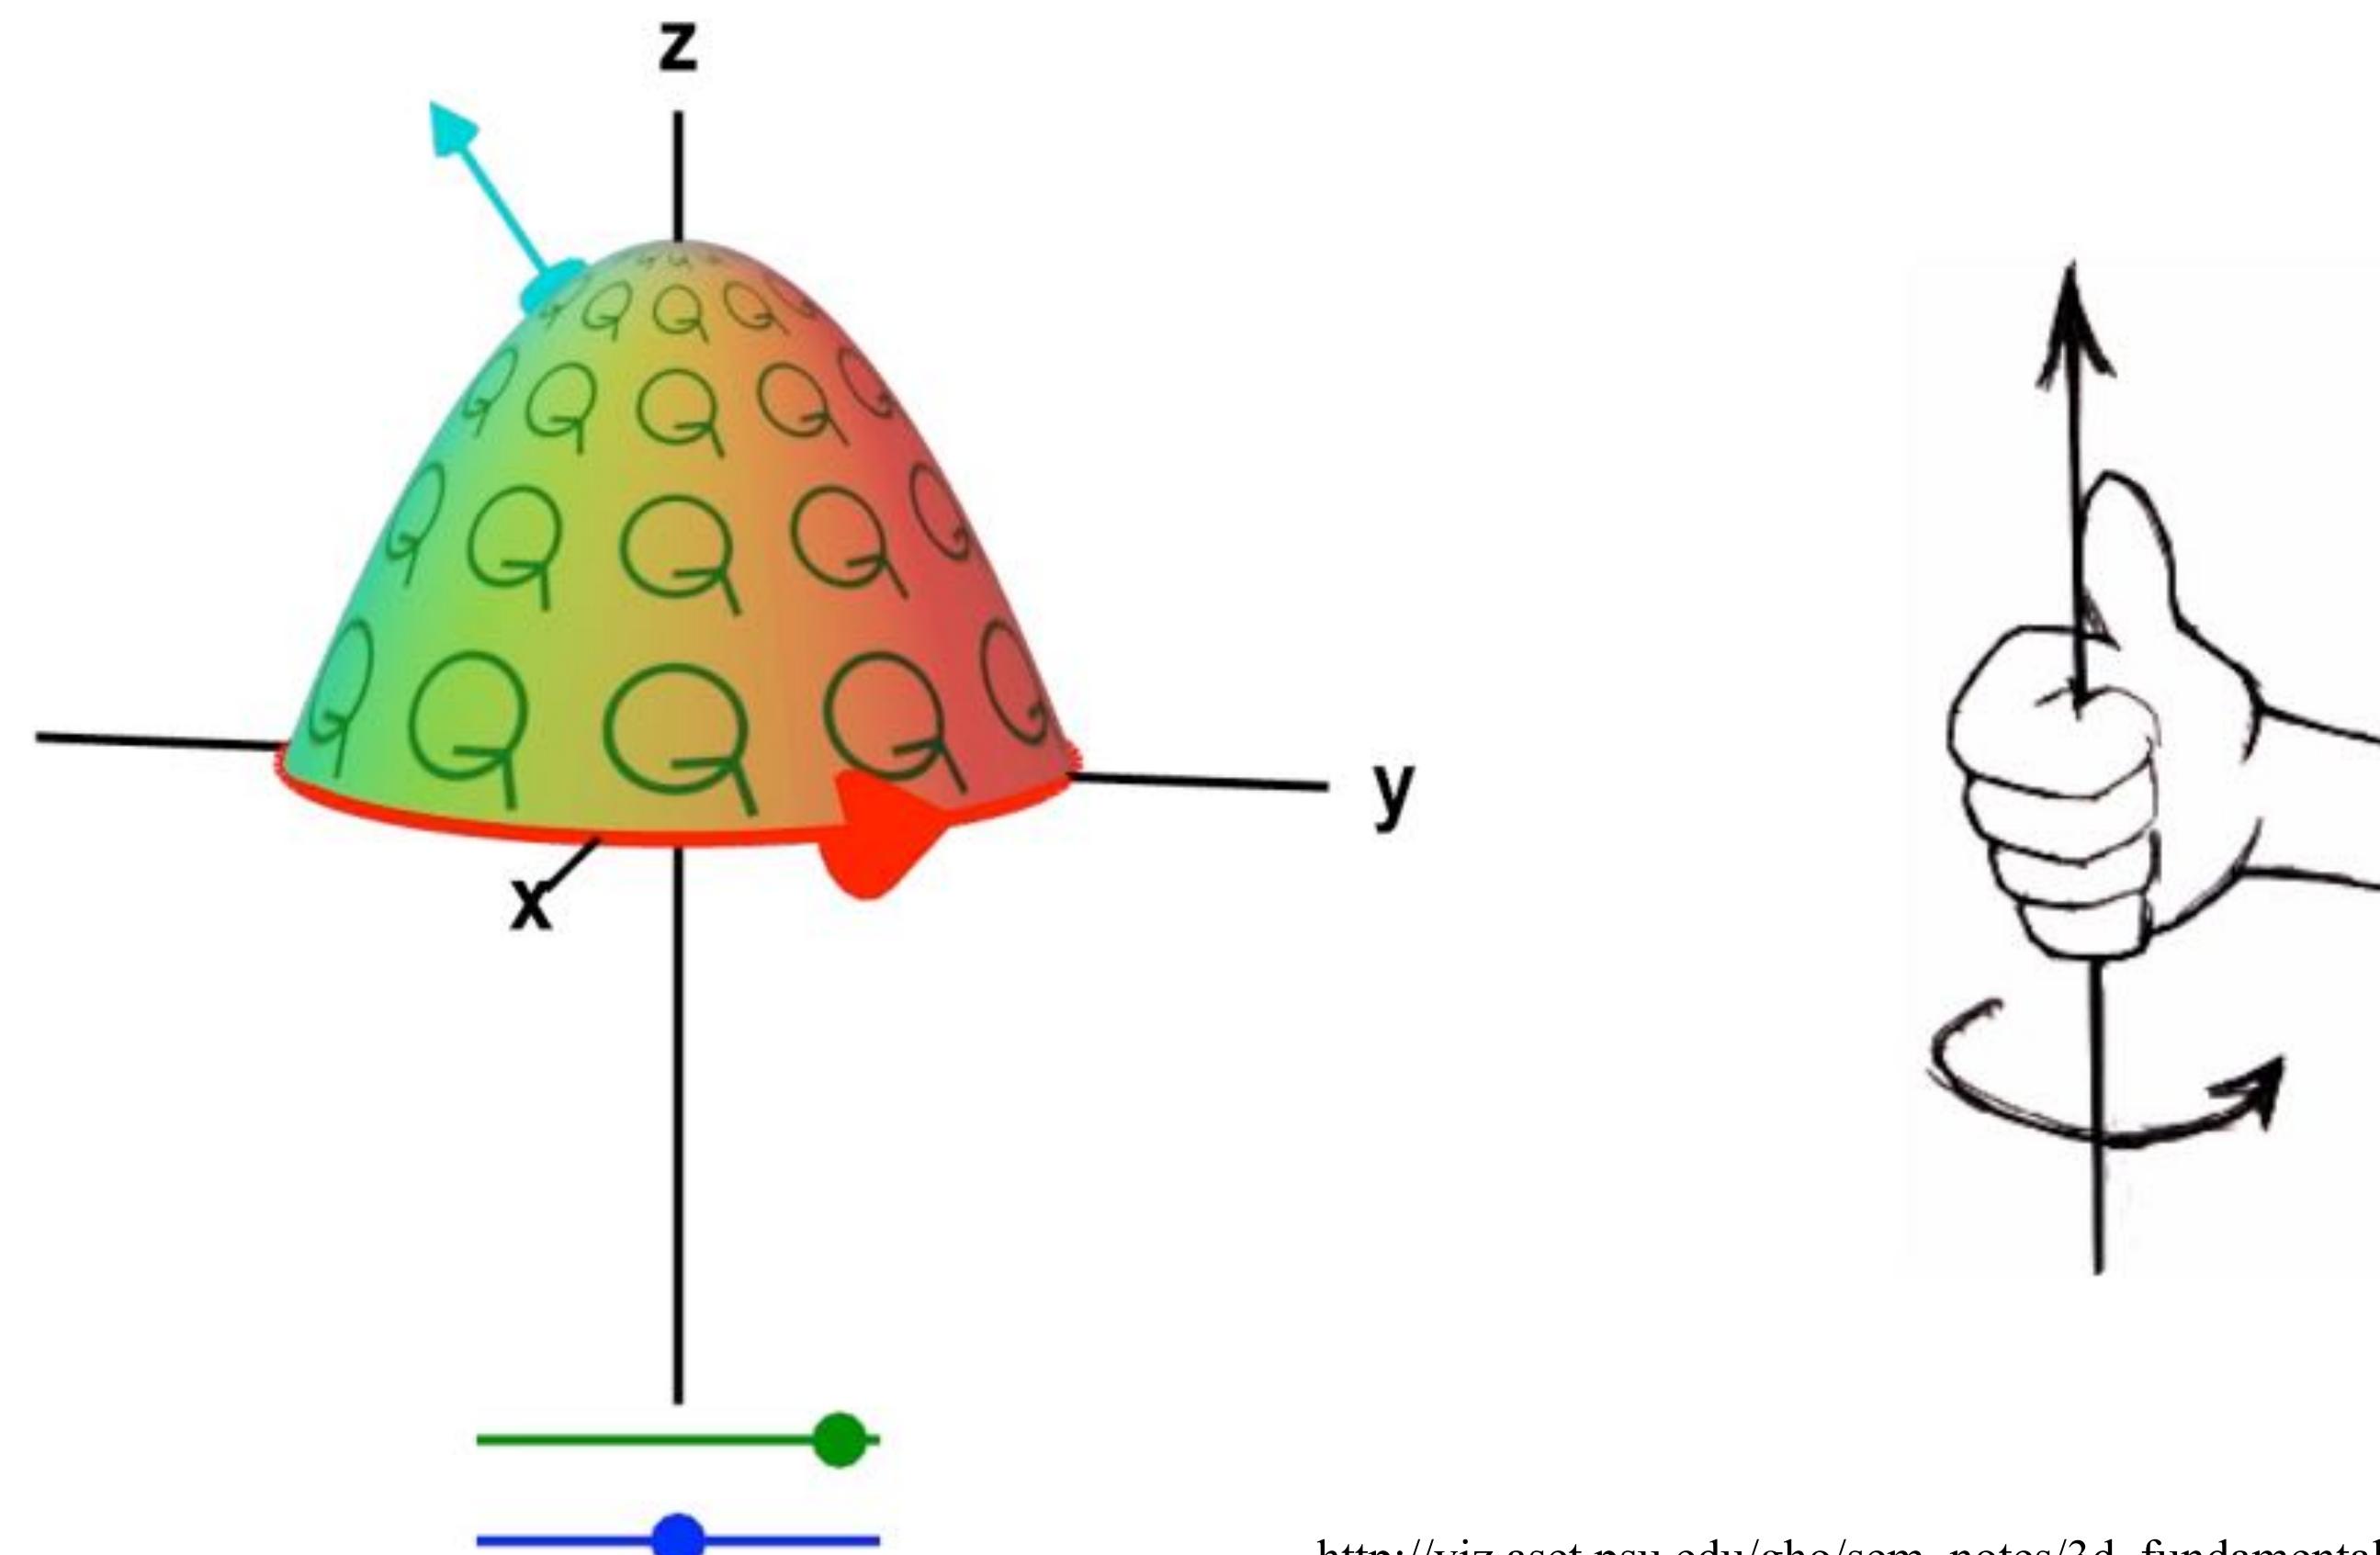
Simple Data Structures: Indexed Face Set

- Used in formats
 - OBJ, OFF, WR L
- Storage
 - Vertex: position
 - Face: vertex indices
 - Convention is to save vertices in counter-clockwise order for normal computation

Vertices				
v0	x0	y0	z0	
v1	x1	x1	z1	
v2	x2	y2	z2	
v3	x3	y3	z3	
v4	x4	y4	z4	
v5	x5	y5	z5	
v6	x6	y6	z6	
...

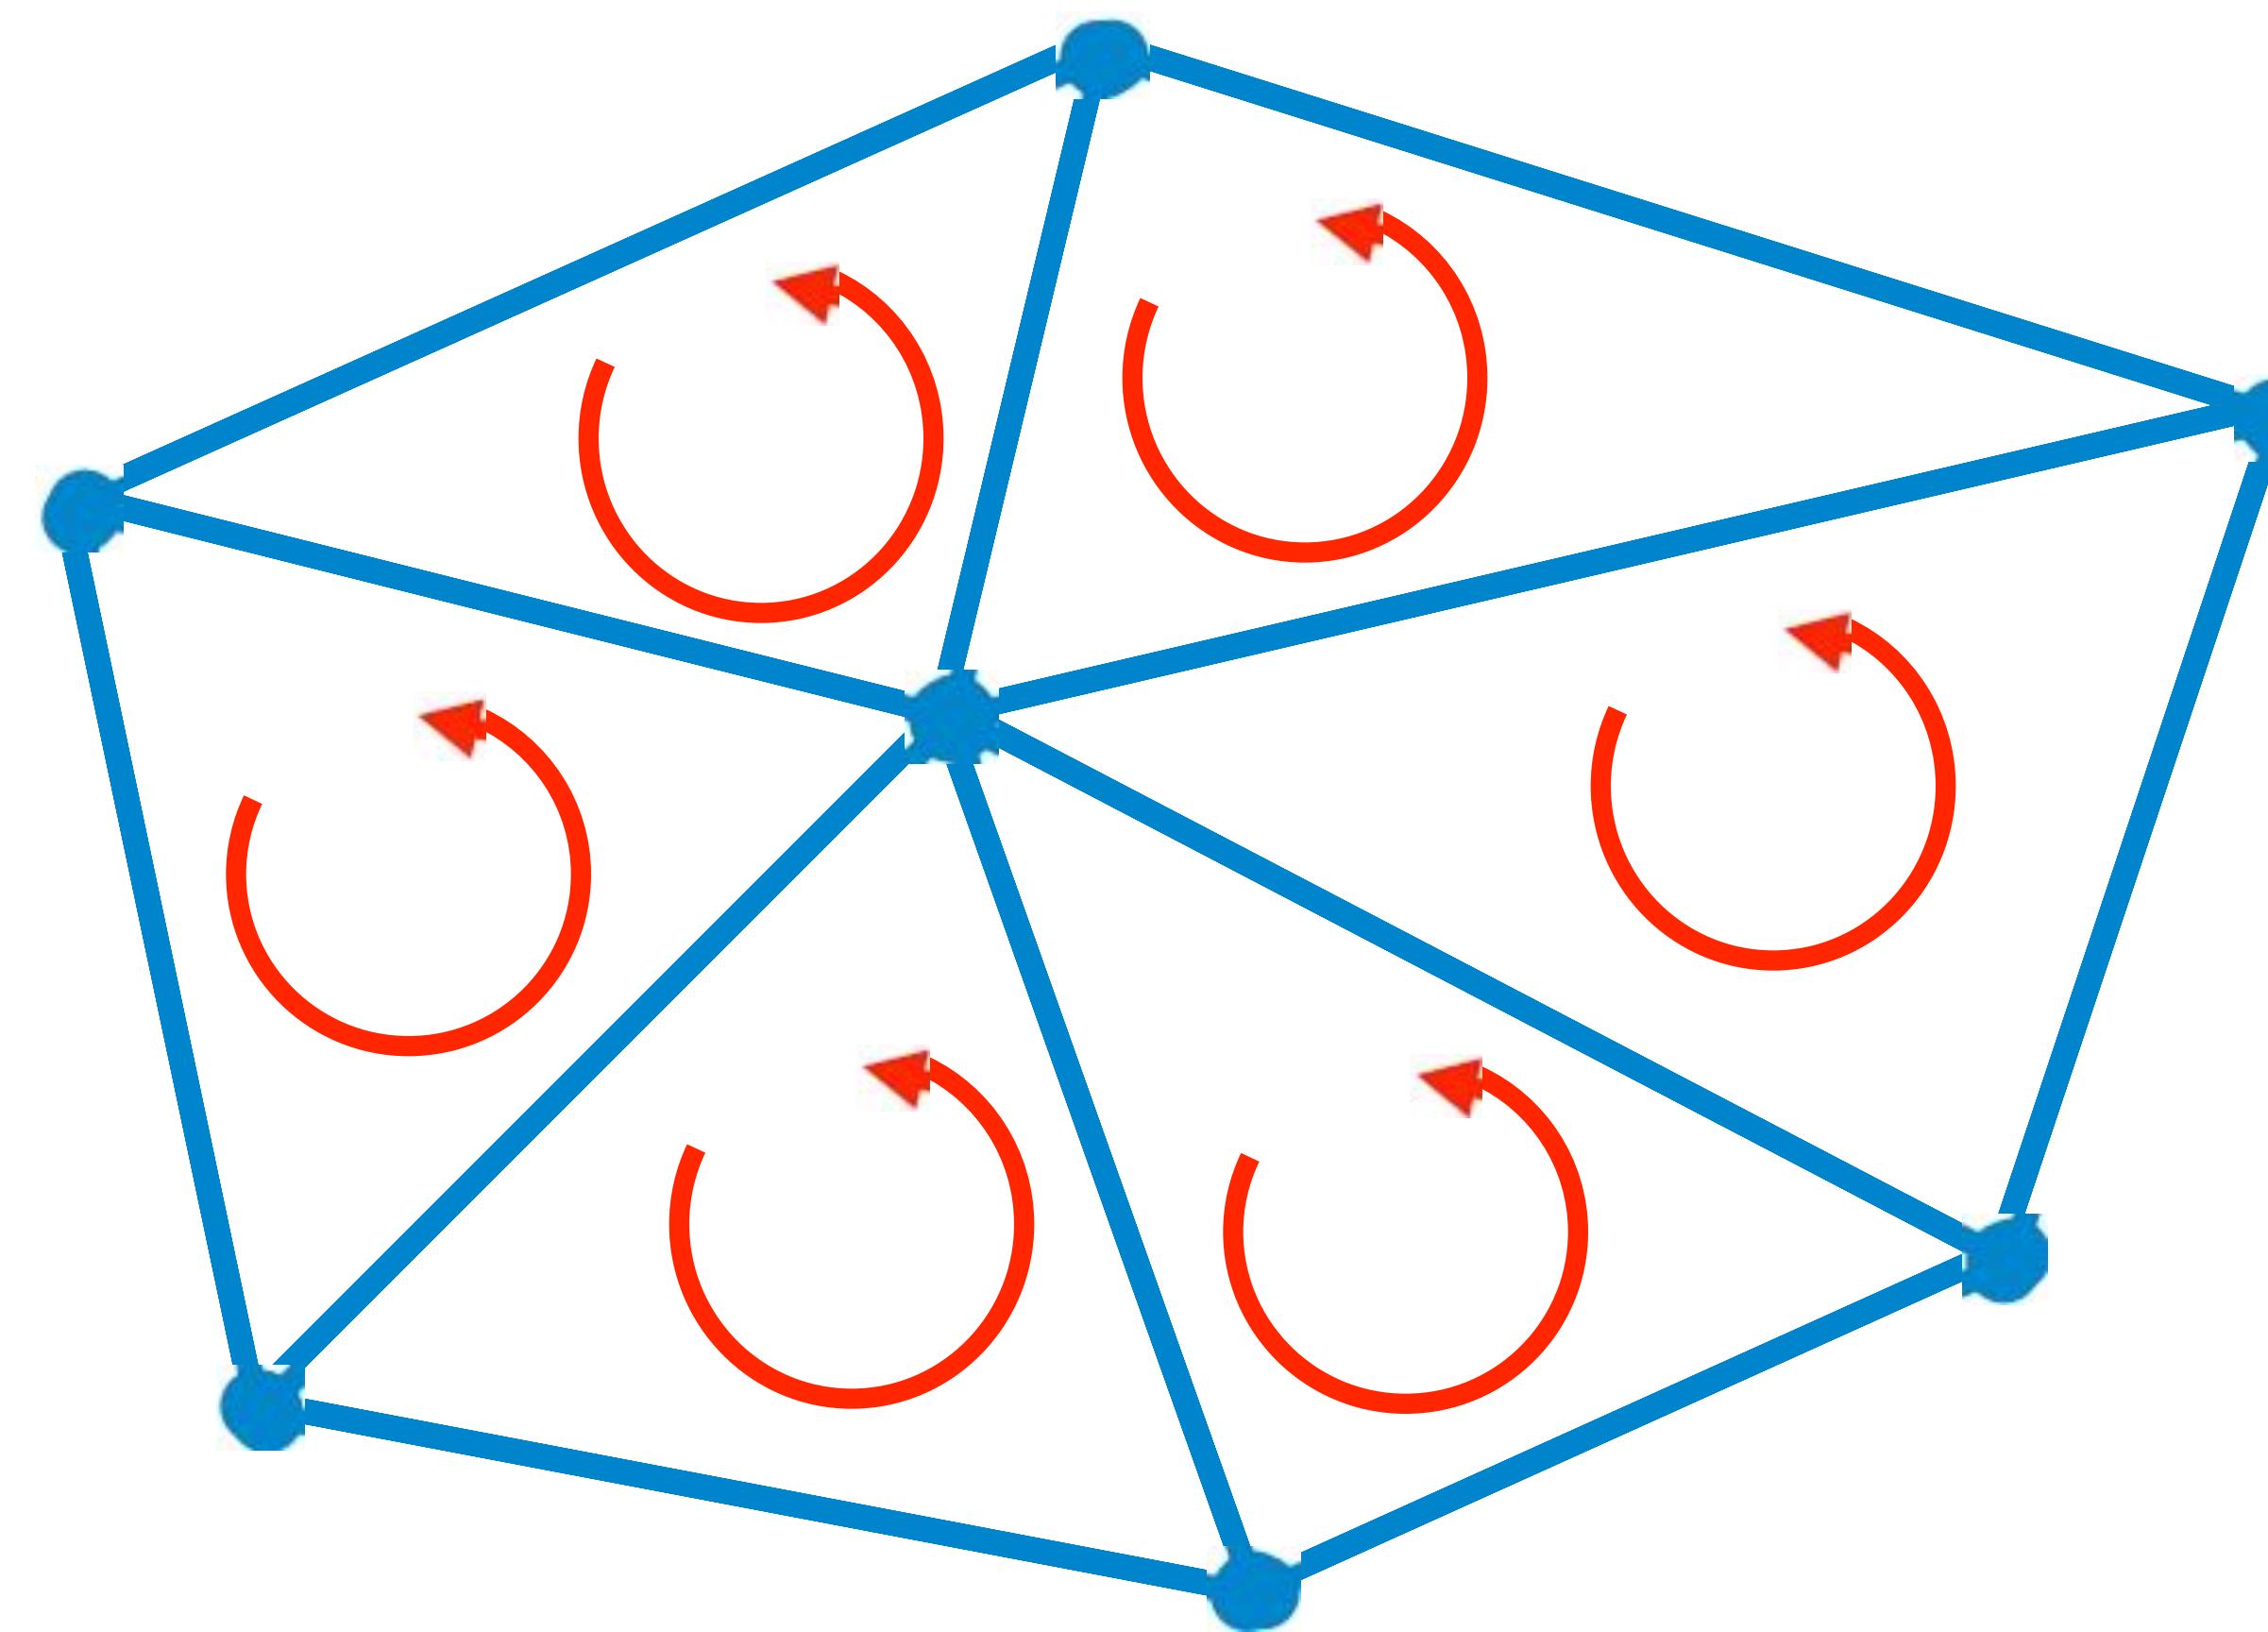
Triangles				
t0	v0	v1	v2	
t1	v0	v1	v3	
t2	v2	v4	v3	
t3	v5	v2	v6	
...

Right-Hand Rule

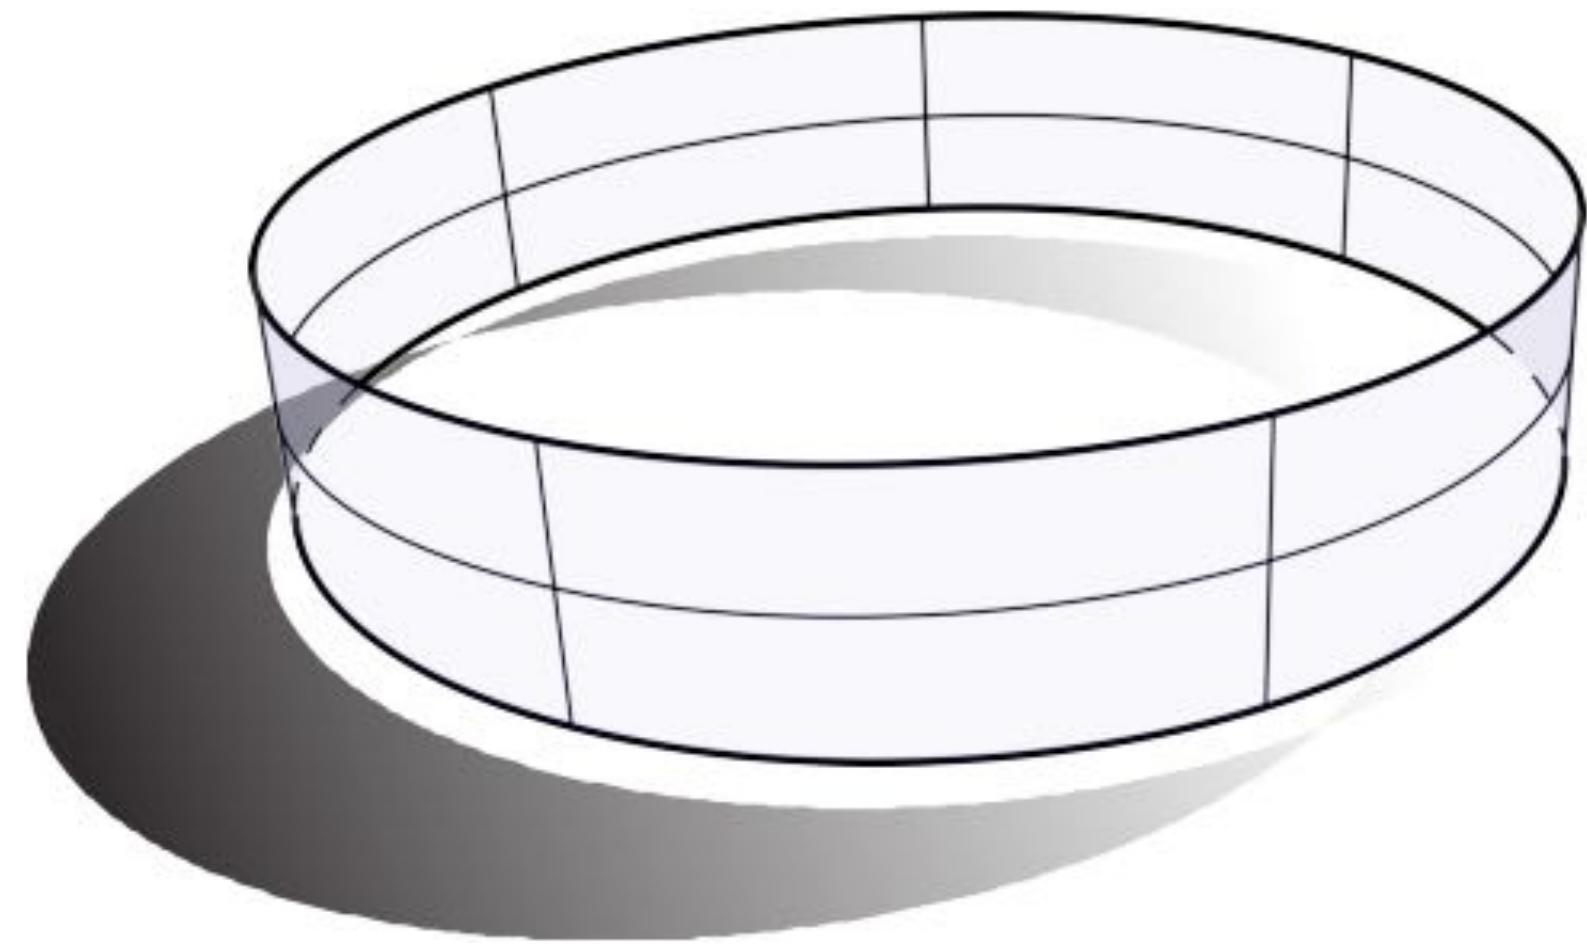


http://viz.aset.psu.edu/gho/sem_notes/3d_fundamentals/html/3d_coordinates.html
http://mathinsight.org/stokes_theorem_orientation

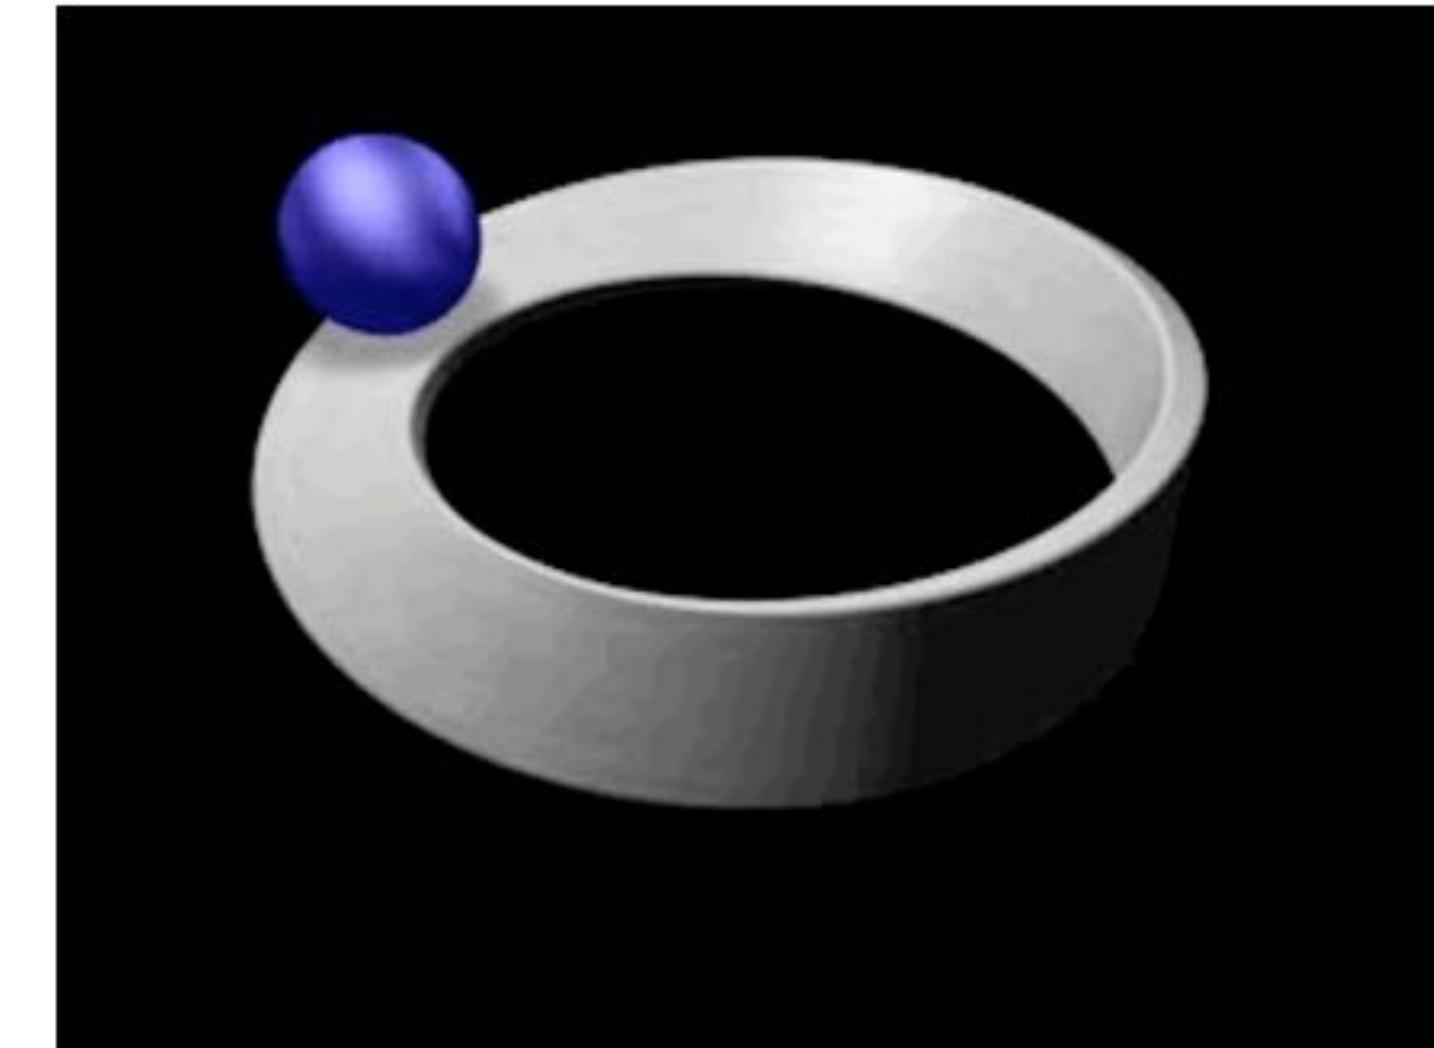
Normal Computation



Orientability



orientable



non-orientable

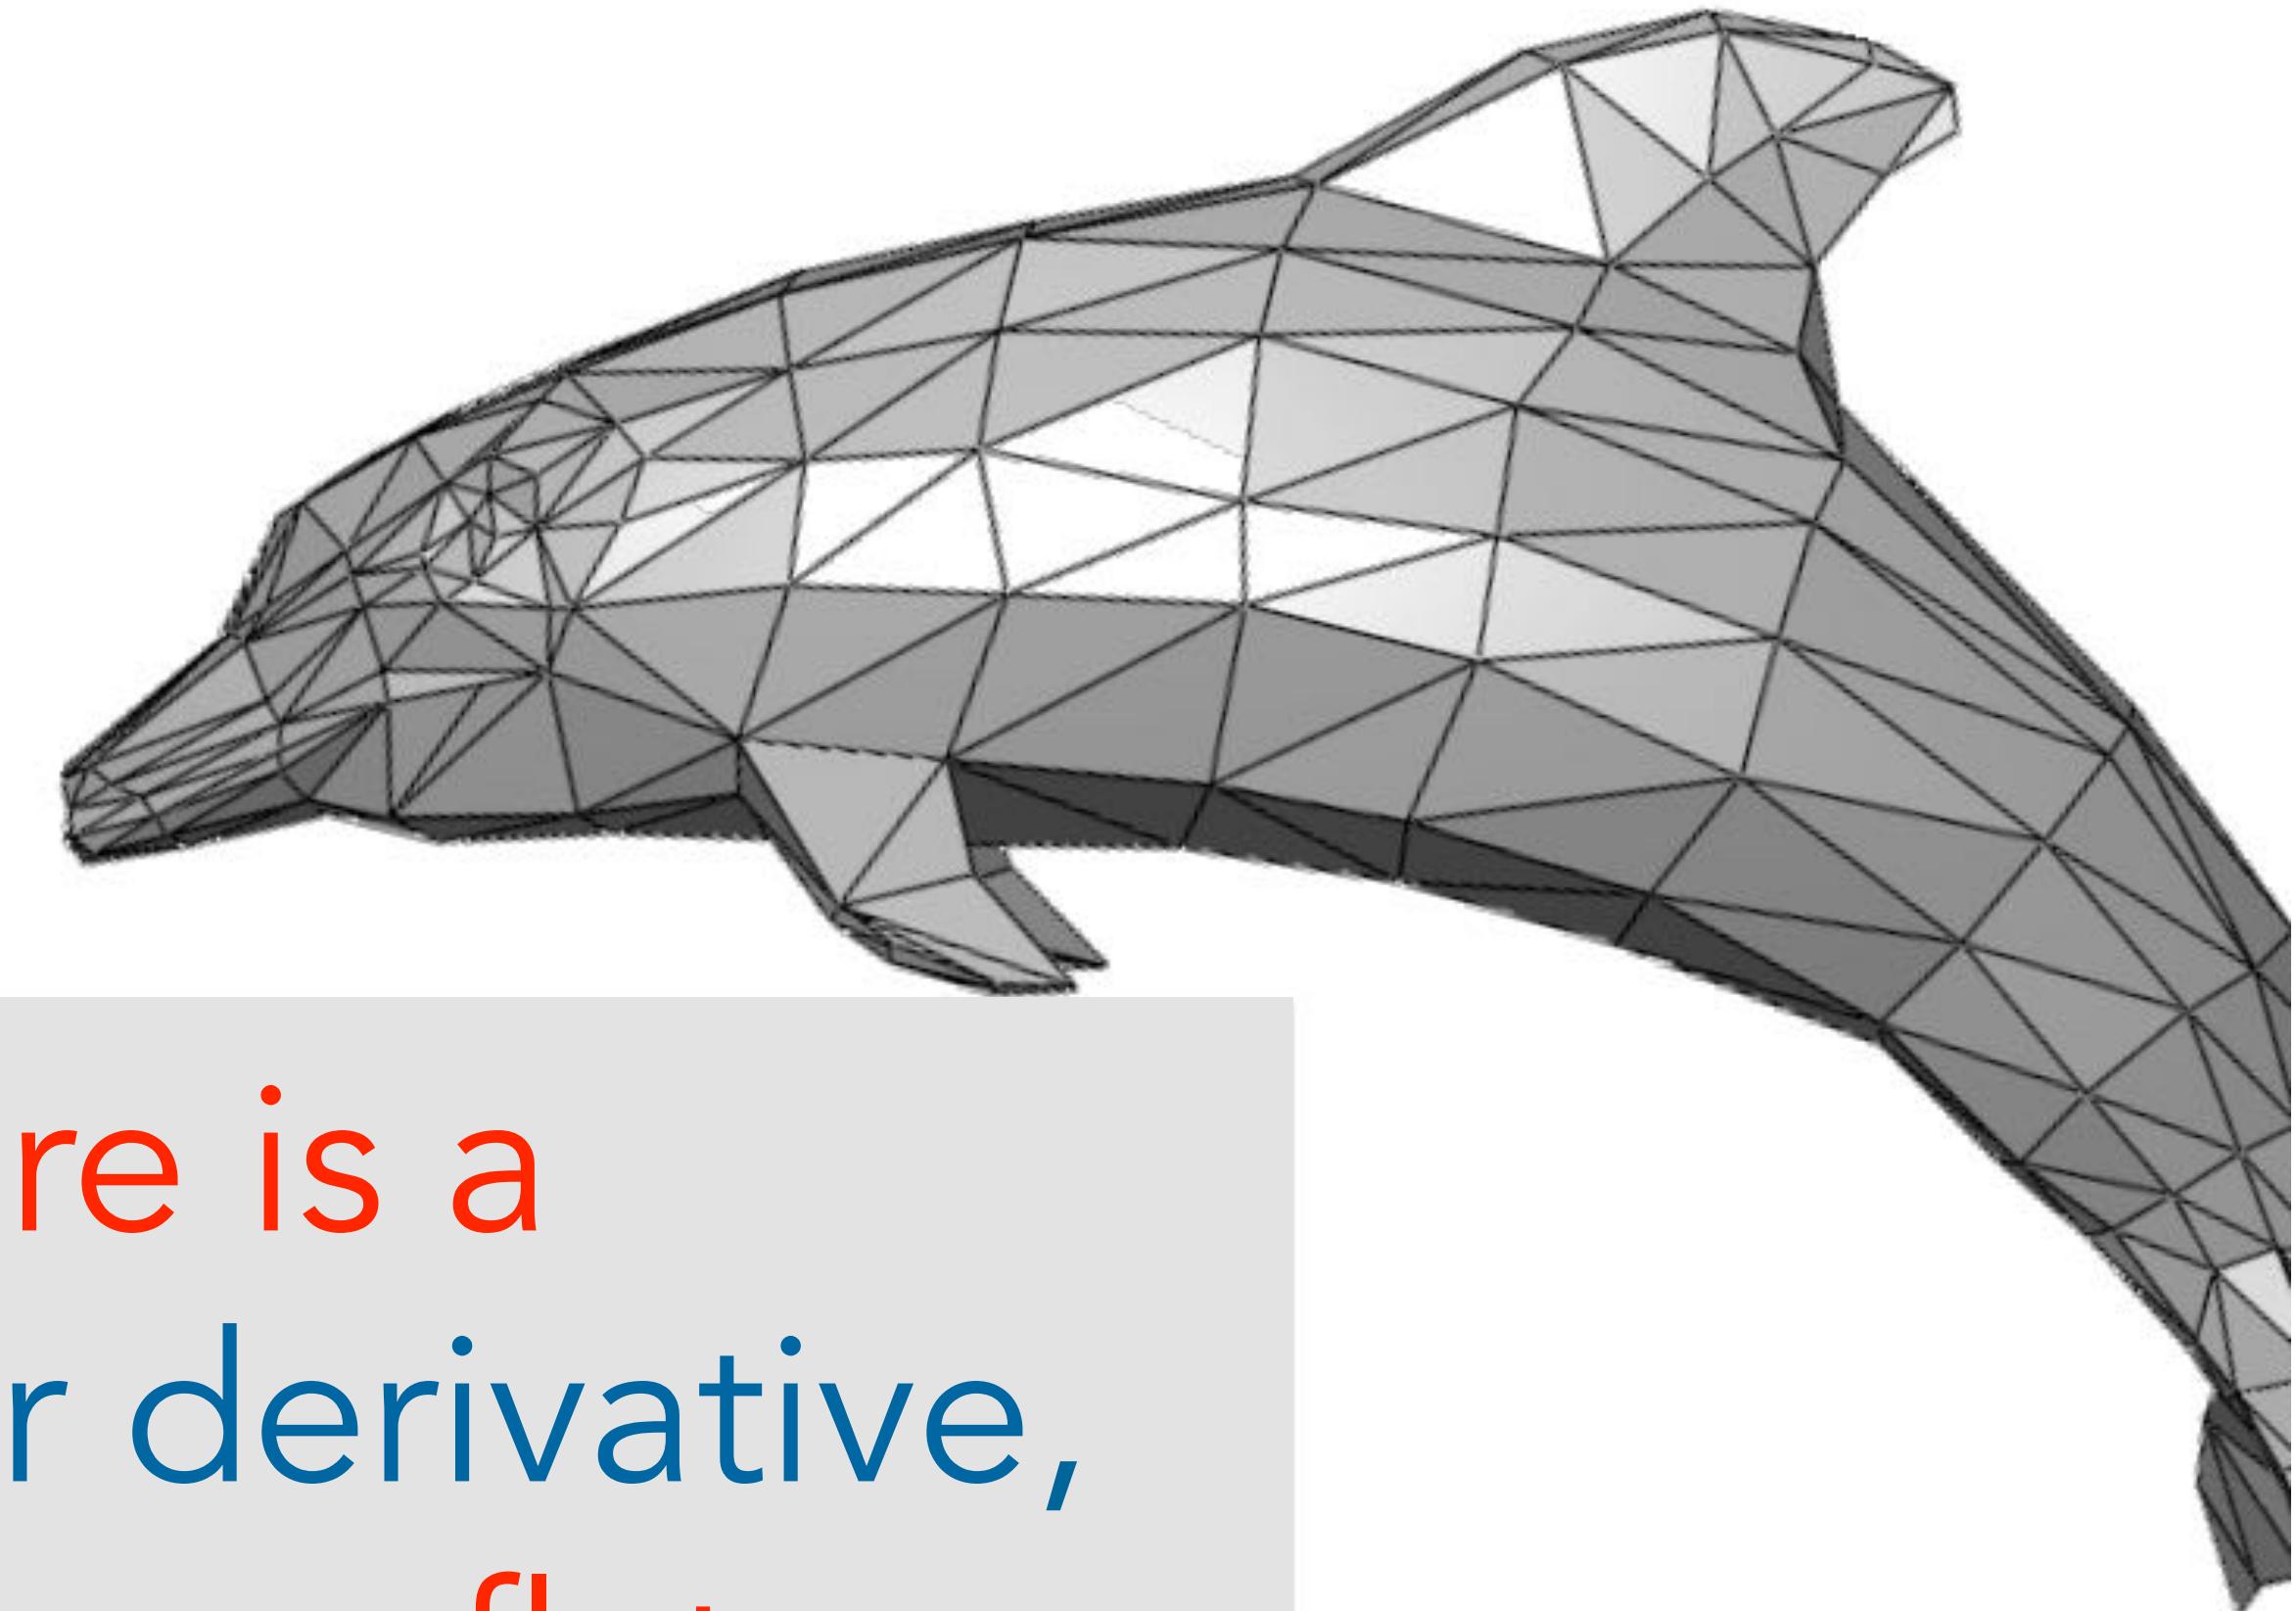
Summary of Polygonal Meshes

- Polygonal meshes are piece-wise linear approximation of smooth surfaces
- Good triangulation is important (manifold, equi-length)
- Vertices, edges, and faces are basic elements
- While real-data 3D are often point clouds, meshes are quite often used to visualize 3D and generate ground truth for machine learning algorithms

Outline

- Rasterized 3D representations
- *Mesh*
 - Representation
 - Storage
 - *Curvature computation*
- Point cloud
- Implicit representations

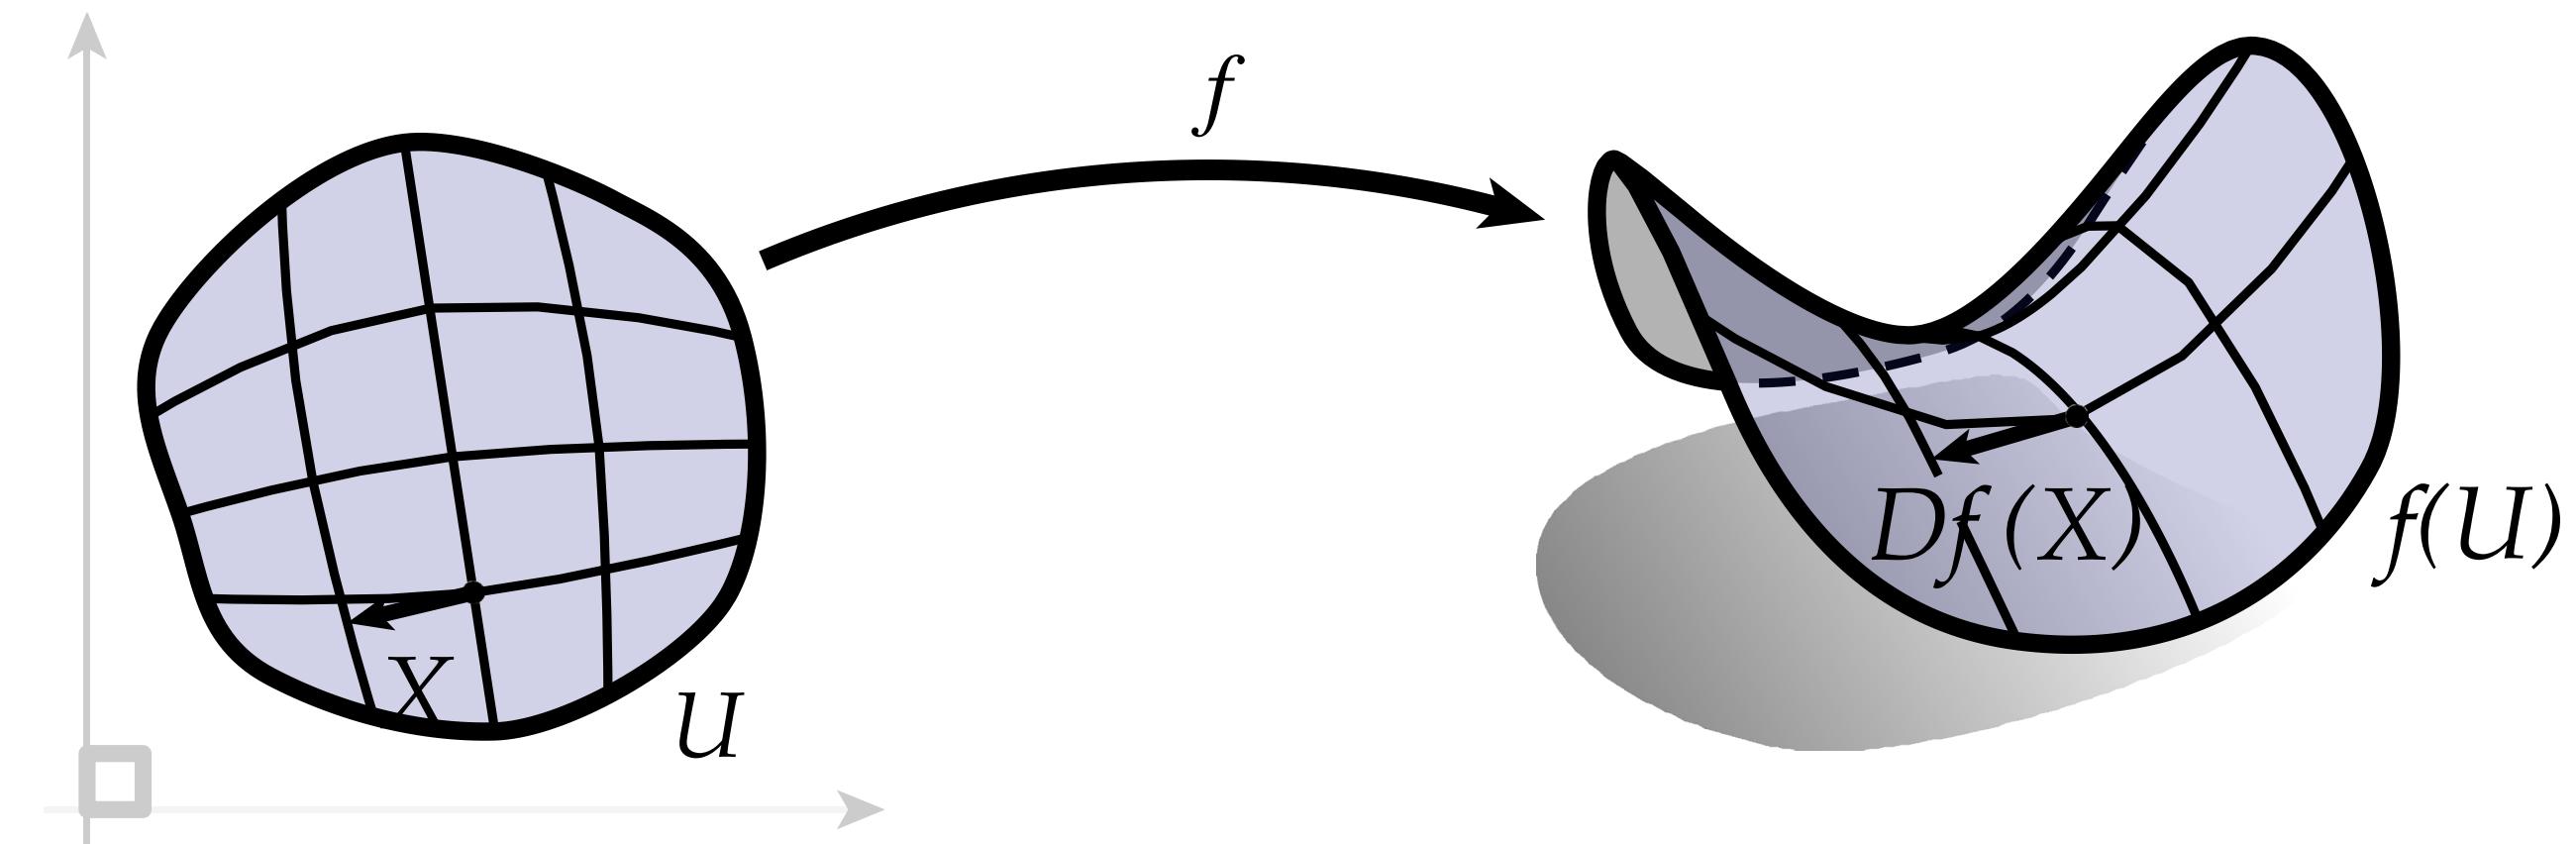
Challenge on Meshes



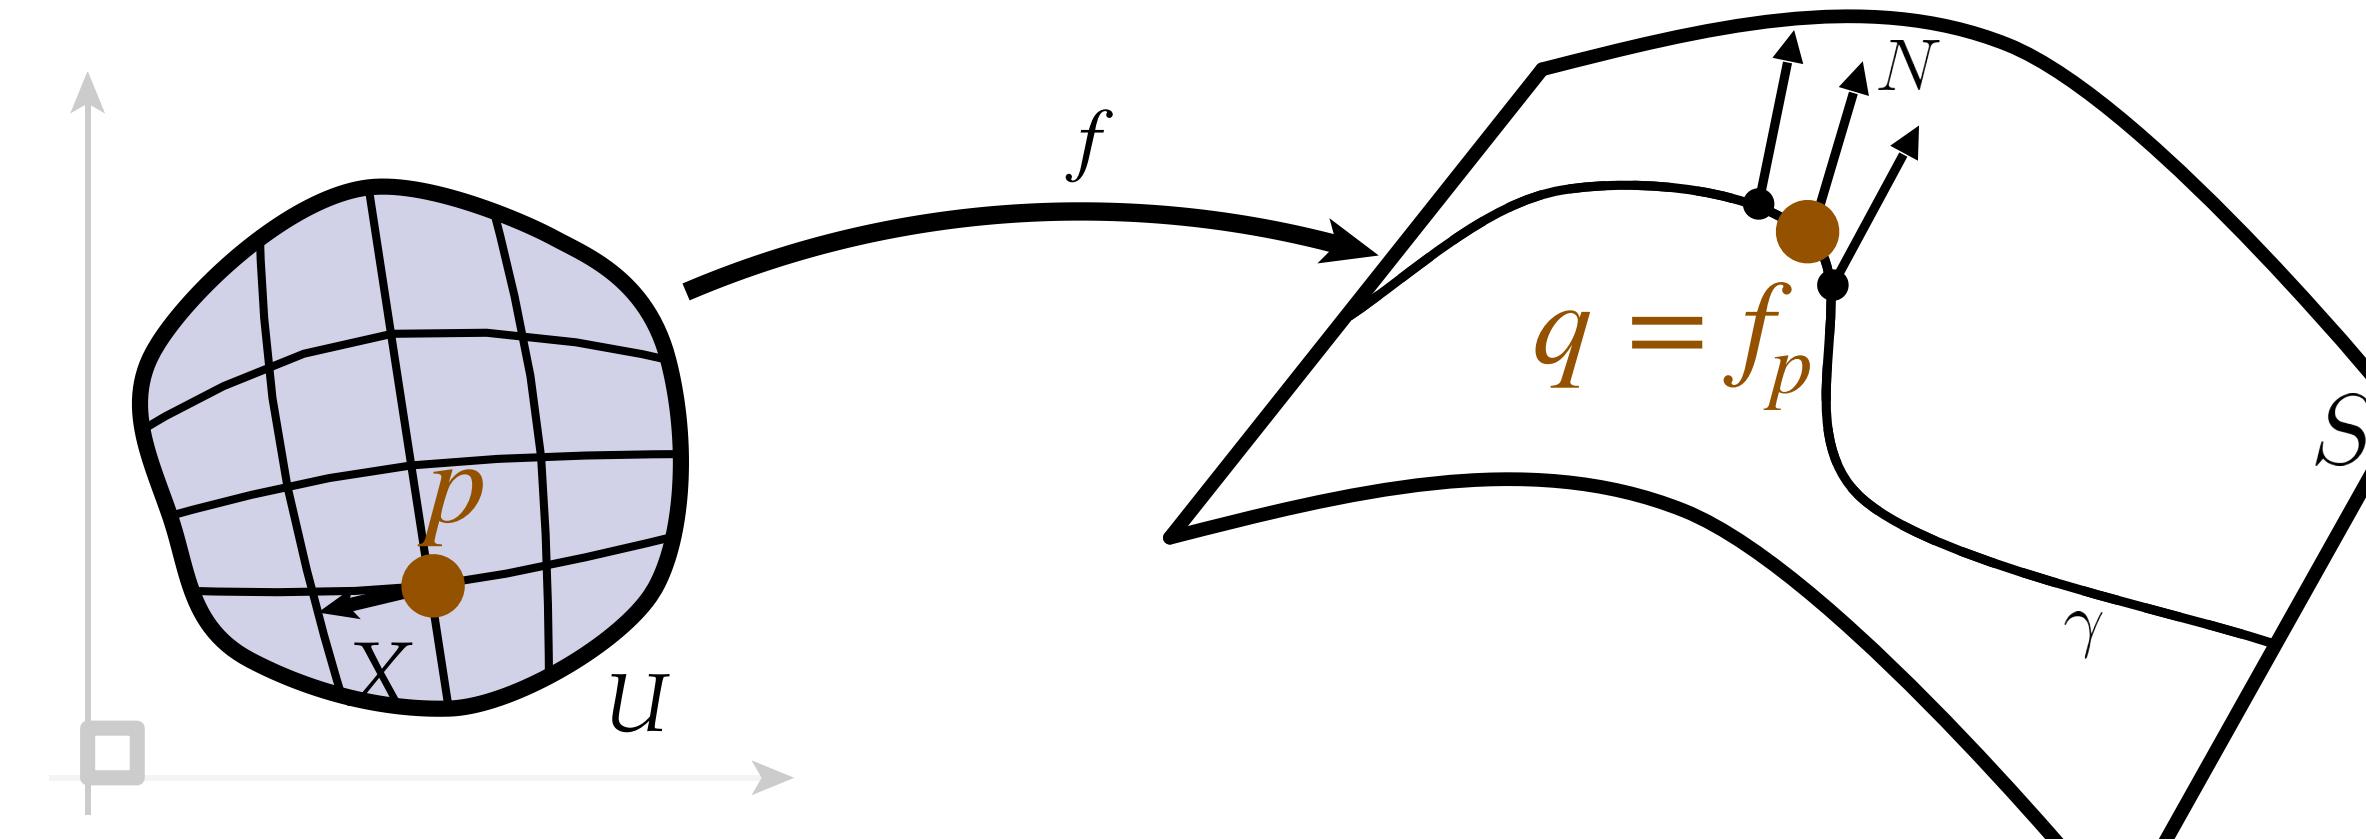
Curvature is a
second-order derivative,
but triangles are flat.

Differential Map in a Nutshell

Differential Map Df :



Differential of Normal DN :

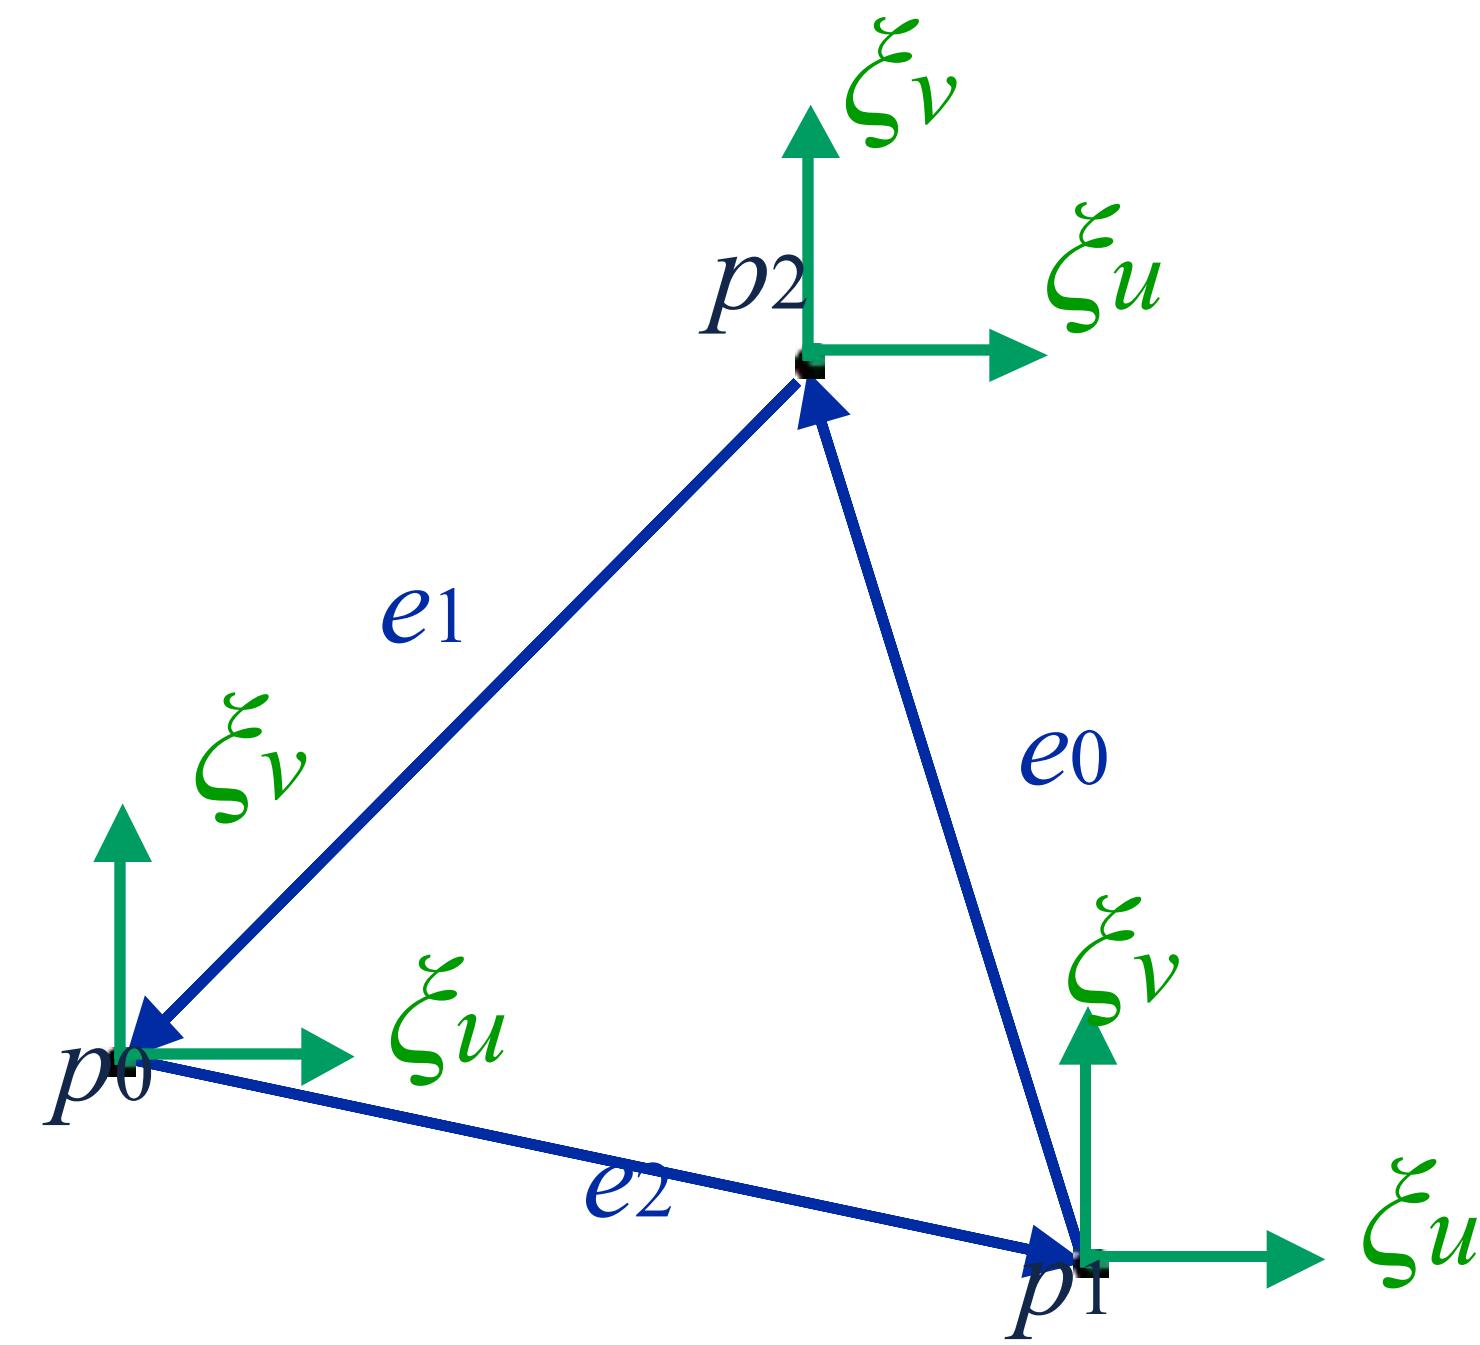


Shape Operator S :

$$\exists S \in \mathbb{R}^{2 \times 2} \text{ such that } DN_p = Df_p S$$

If we have S , we can compute principal curvatures!

Rusinkiewicz's Method

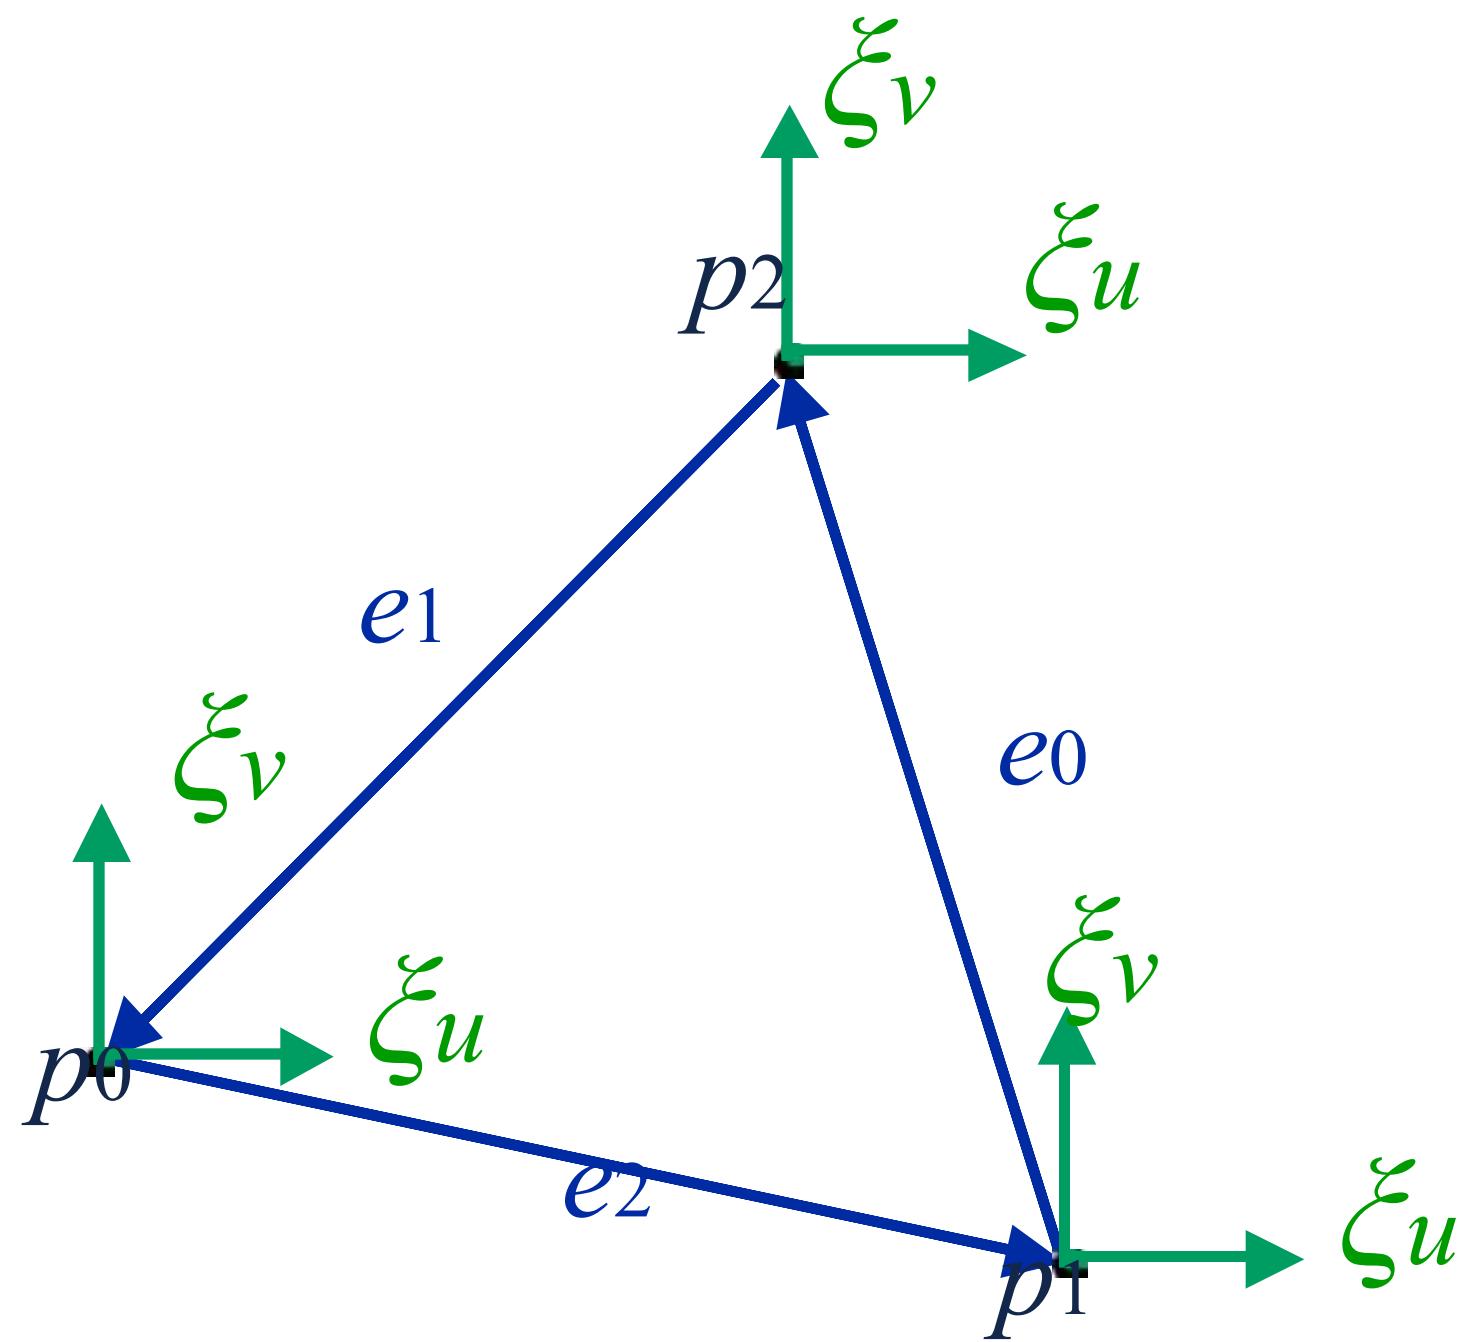


Assume a local $f: U \rightarrow \mathbb{R}^3$ at a small triangle

Assume that \mathbf{T}_{p_i} 's are roughly parallel

Assume that $Df[uv] = u \xi_u + v \xi_v$, i.e., $Df = [\xi_u, \xi_v]$
(U is also aligned with \mathbf{T}_{p_i} 's)

Rusinkiewicz's Method



Shape operator S satisfies: $DN = Df \cdot S$.

$$\therefore Df = [\xi_u, \xi_v], \therefore S = Df^T DN$$

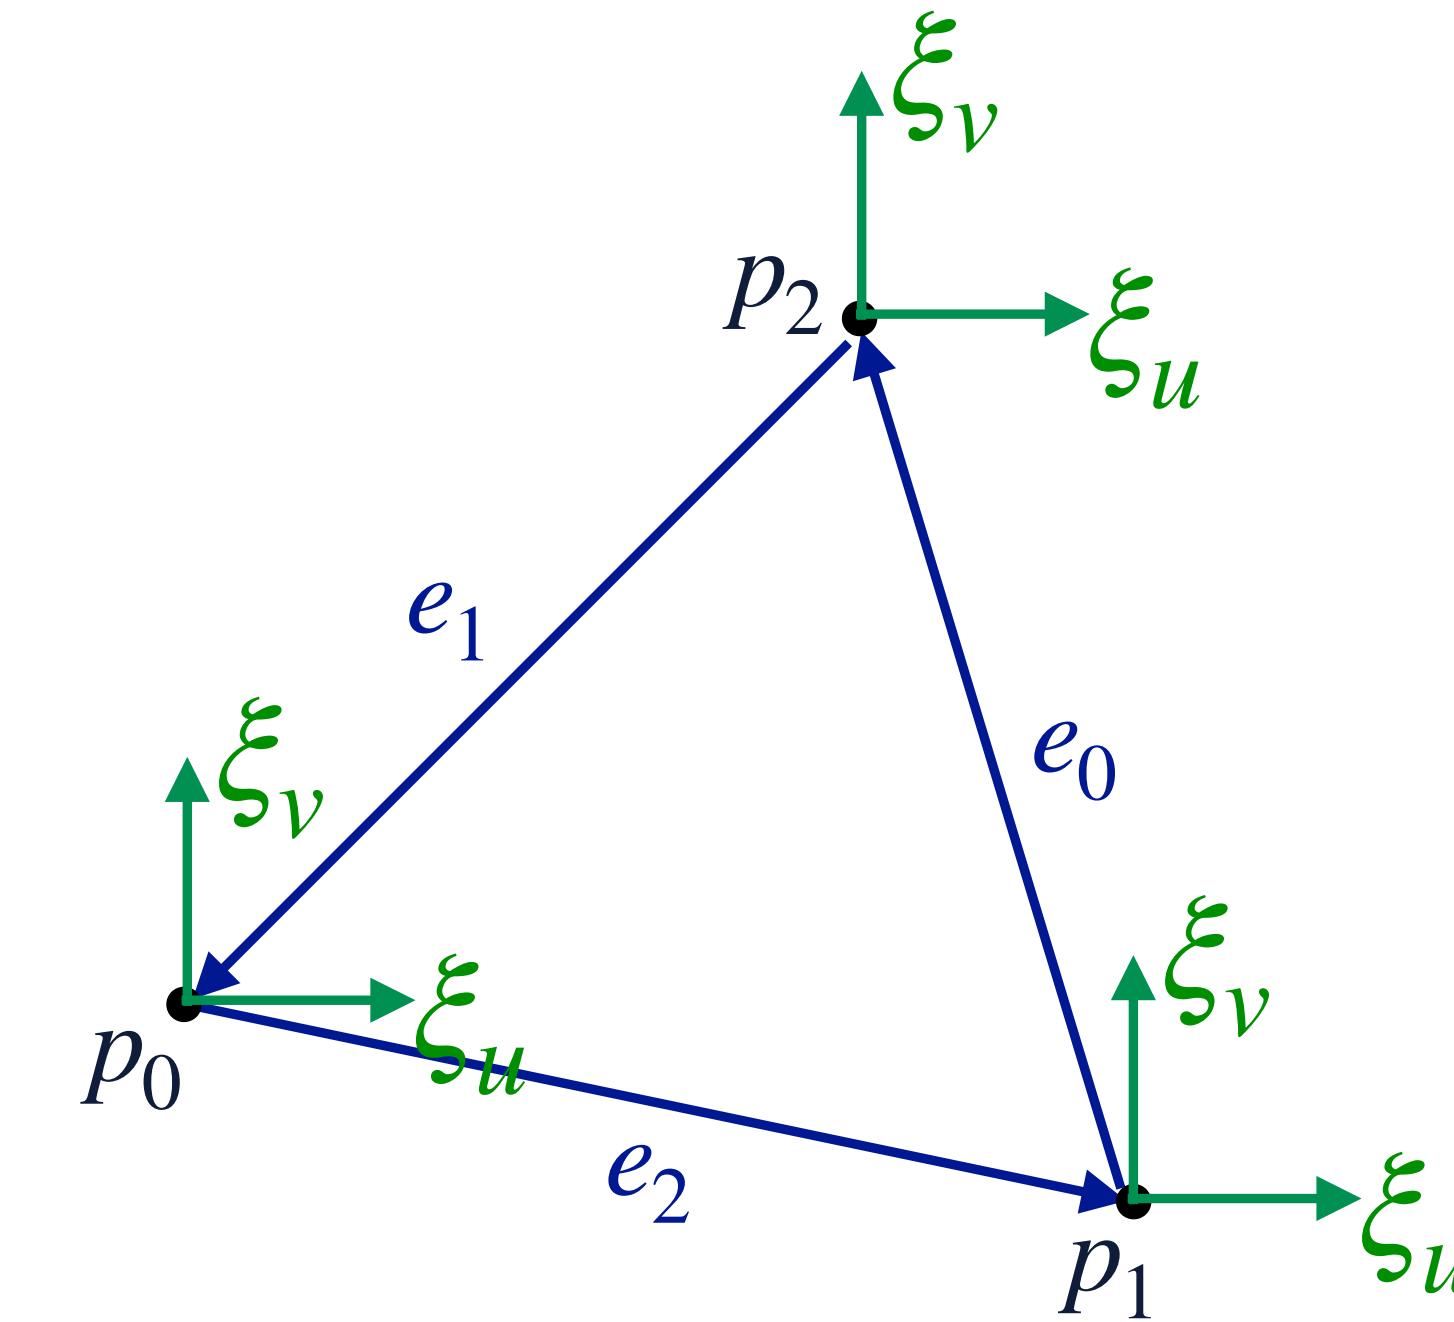
How to estimate S ?

Rusinkiewicz's Method

$$Df^T \left(DN \begin{bmatrix} u \\ v \end{bmatrix} \right) \approx Df^T \Delta \vec{n} \implies S \begin{bmatrix} u \\ v \end{bmatrix} \approx Df^T \Delta \vec{n}$$

$$\because Df \begin{bmatrix} u \\ v \end{bmatrix} = Y \in \mathbf{T}(\mathbb{R}^3) \text{ and } Df = [\vec{\xi}_u, \vec{\xi}_v] \quad \therefore \begin{bmatrix} u \\ v \end{bmatrix} = Df^T Y$$

$$\therefore S[Df]^T Y \approx [Df]^T \Delta \vec{n}$$



Rusinkiewicz's Method

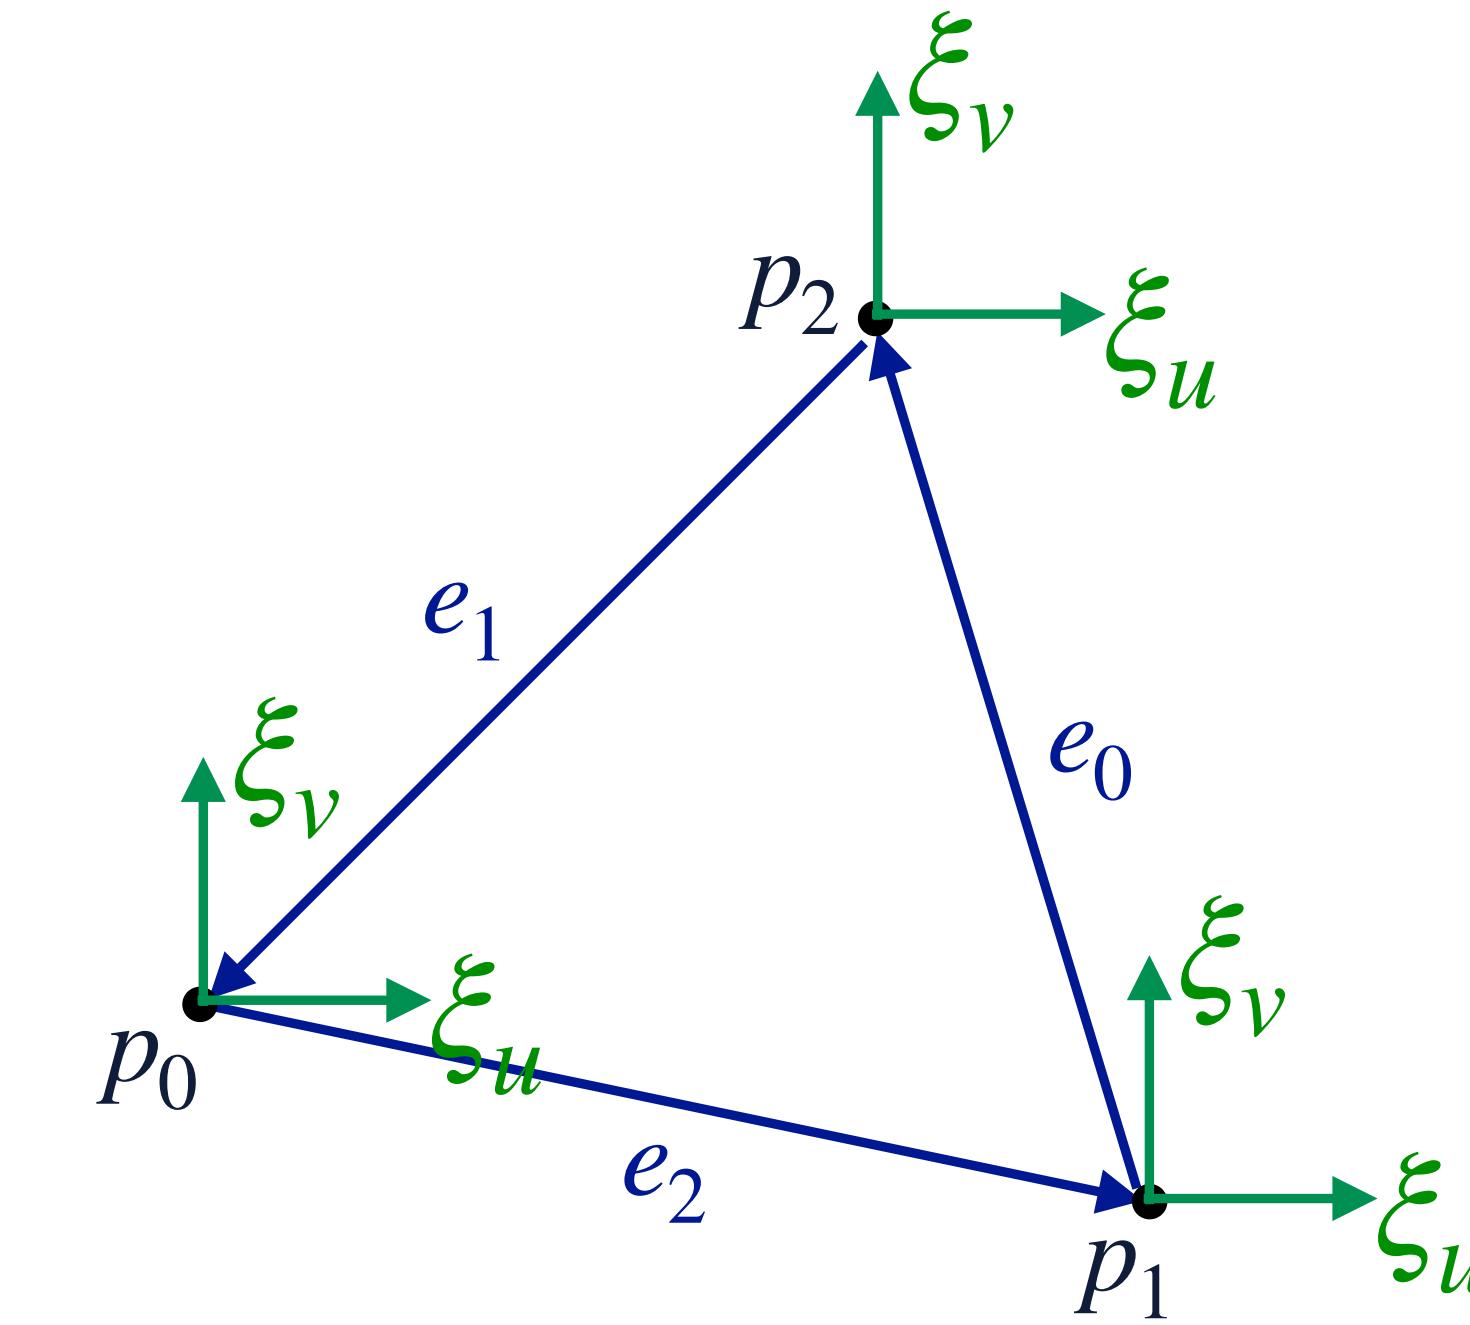
$$Df^T \left(DN \begin{bmatrix} u \\ v \end{bmatrix} \right) \approx Df^T \Delta \vec{n} \implies S \begin{bmatrix} u \\ v \end{bmatrix} \approx Df^T \Delta \vec{n}$$

$$\because Df \begin{bmatrix} u \\ v \end{bmatrix} = Y \in \mathbf{T}(\mathbb{R}^3) \text{ and } Df = [\vec{\xi}_u, \vec{\xi}_v] \quad \therefore \begin{bmatrix} u \\ v \end{bmatrix} = Df^T Y$$

$$\therefore S[Df]^T Y \approx [Df]^T \Delta \vec{n}$$

$$\begin{cases} S[Df]^T e_0 = Df^T (\vec{n}_2 - \vec{n}_1), \\ S[Df]^T e_1 = Df^T (\vec{n}_0 - \vec{n}_2), \\ S[Df]^T e_2 = Df^T (\vec{n}_1 - \vec{n}_0), \end{cases}$$

So we can solve $S \in \mathbb{R}^{2 \times 2}$ by least square(6 equations and 4 unknowns)



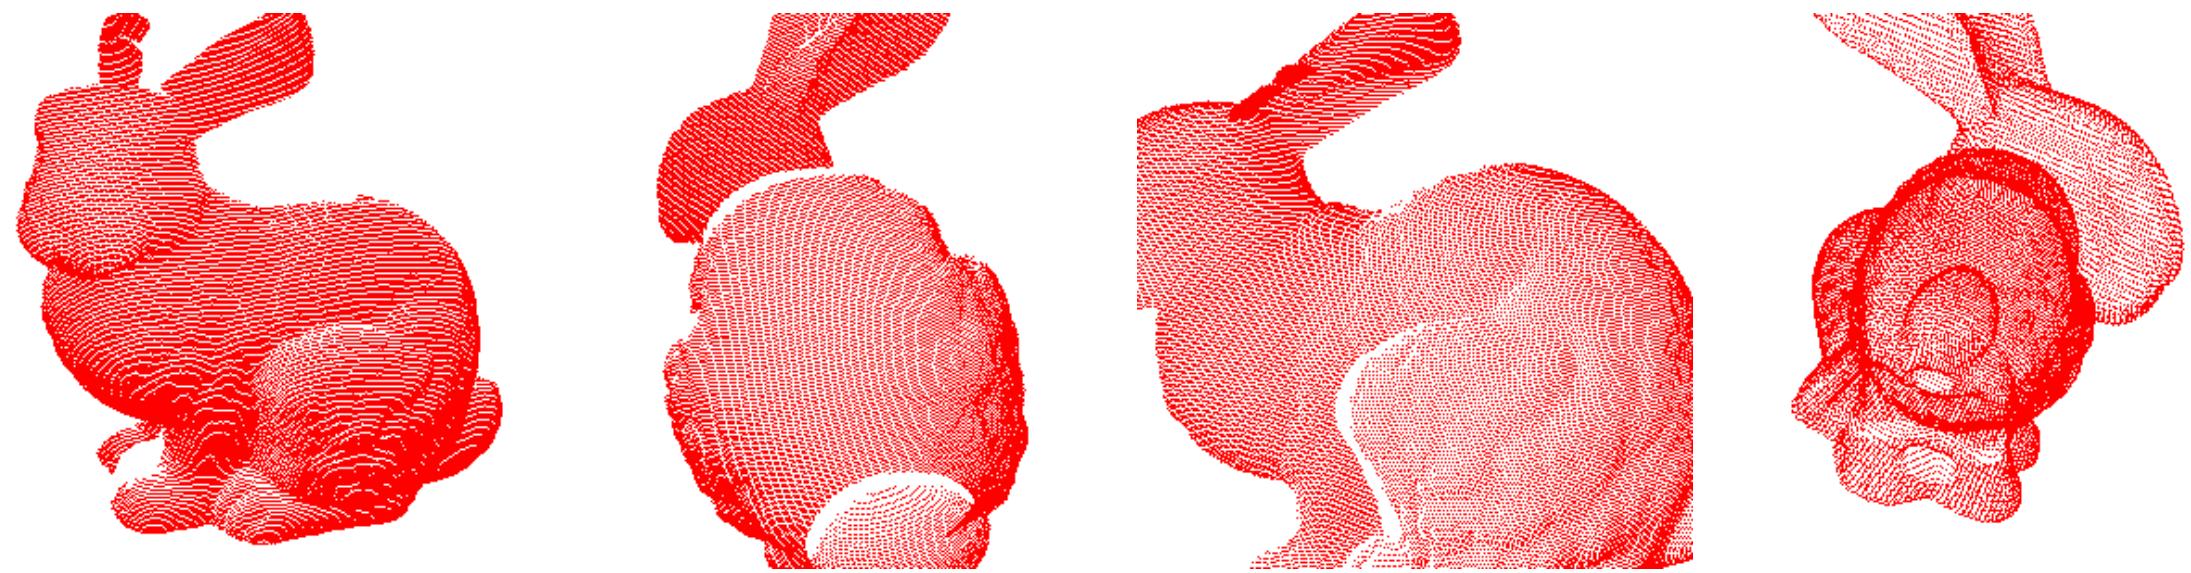
Summary of Mesh Curvature Estimation

- Rusinkiewicz's method is an effective approach for face curvature estimation
 - Szymon Rusinkiewicz, "Estimating Curvatures and Their Derivatives on Triangle Meshes", 3DPVT, 2004
- Good robustness to moderate amount of noise and free of degenerate configurations
- Can be used to compute curvatures for point cloud as well

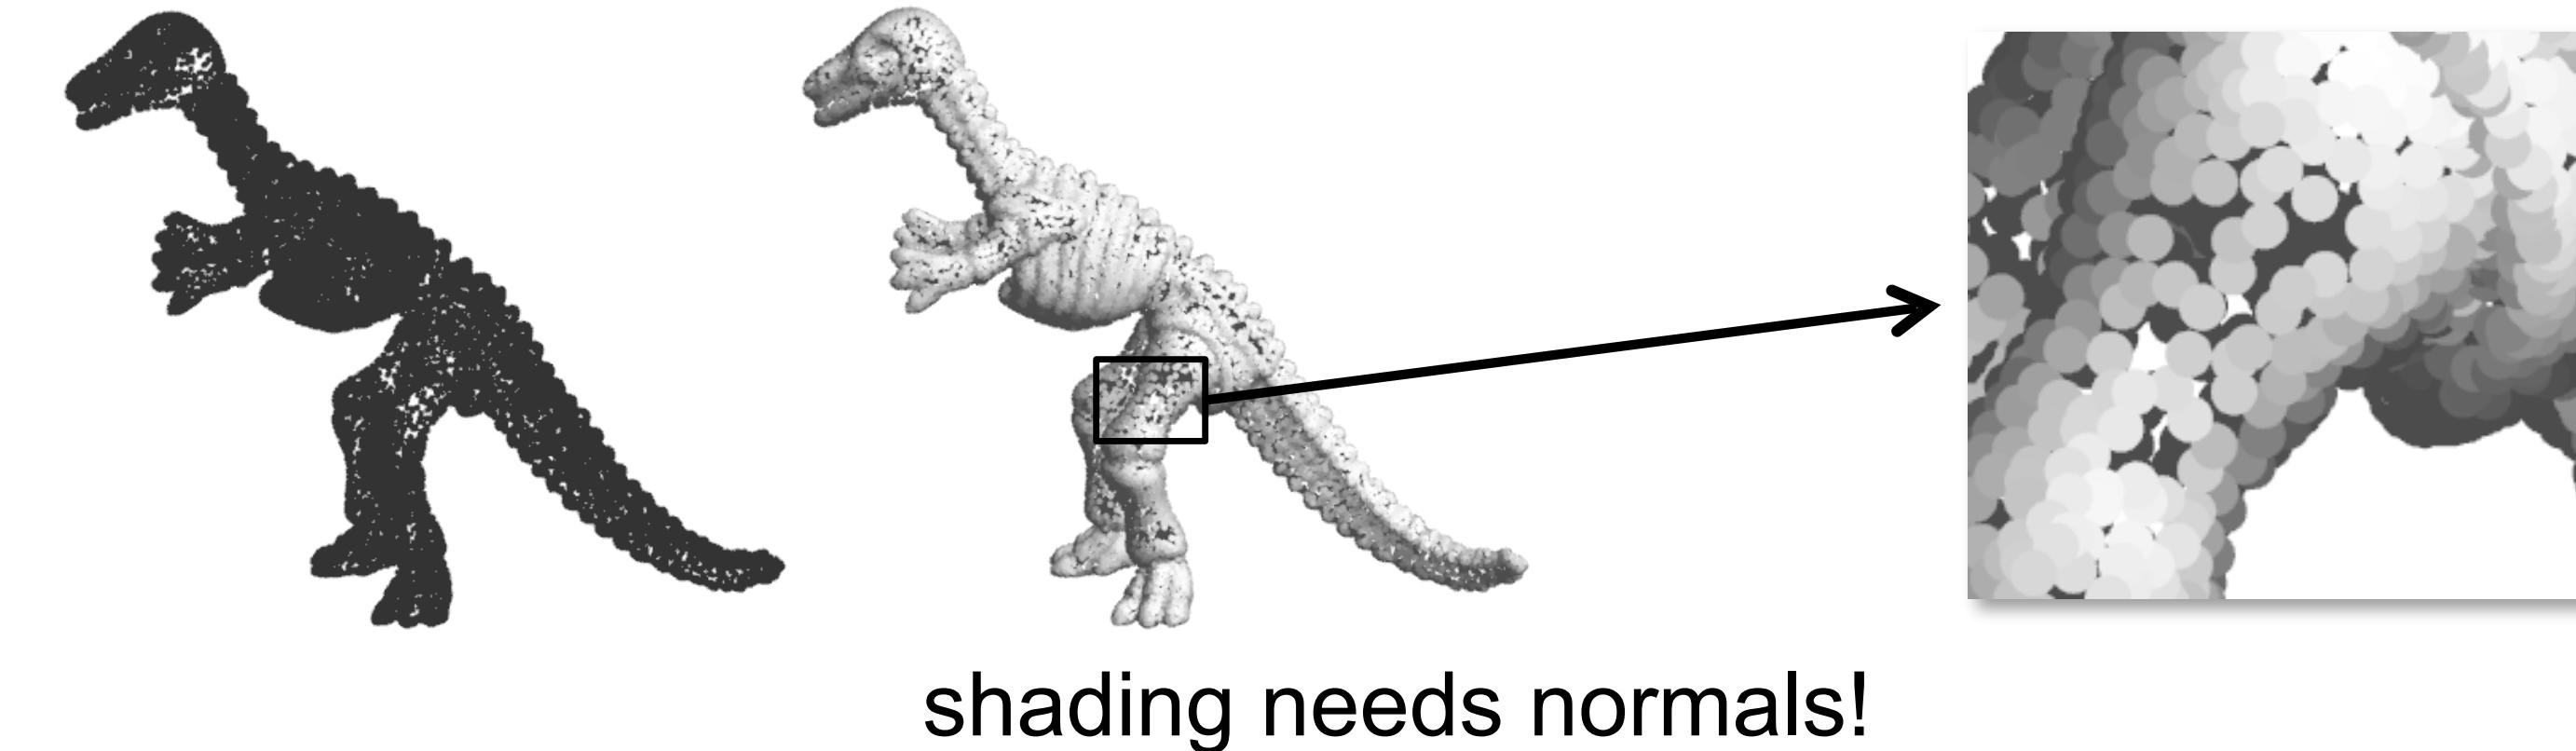
Outline

- Rasterized 3D representations
- Mesh
- *Point cloud*
 - *Representation*
 - Sampling points on surfaces
 - Normal computation
- Implicit representations

Point Cloud Representation



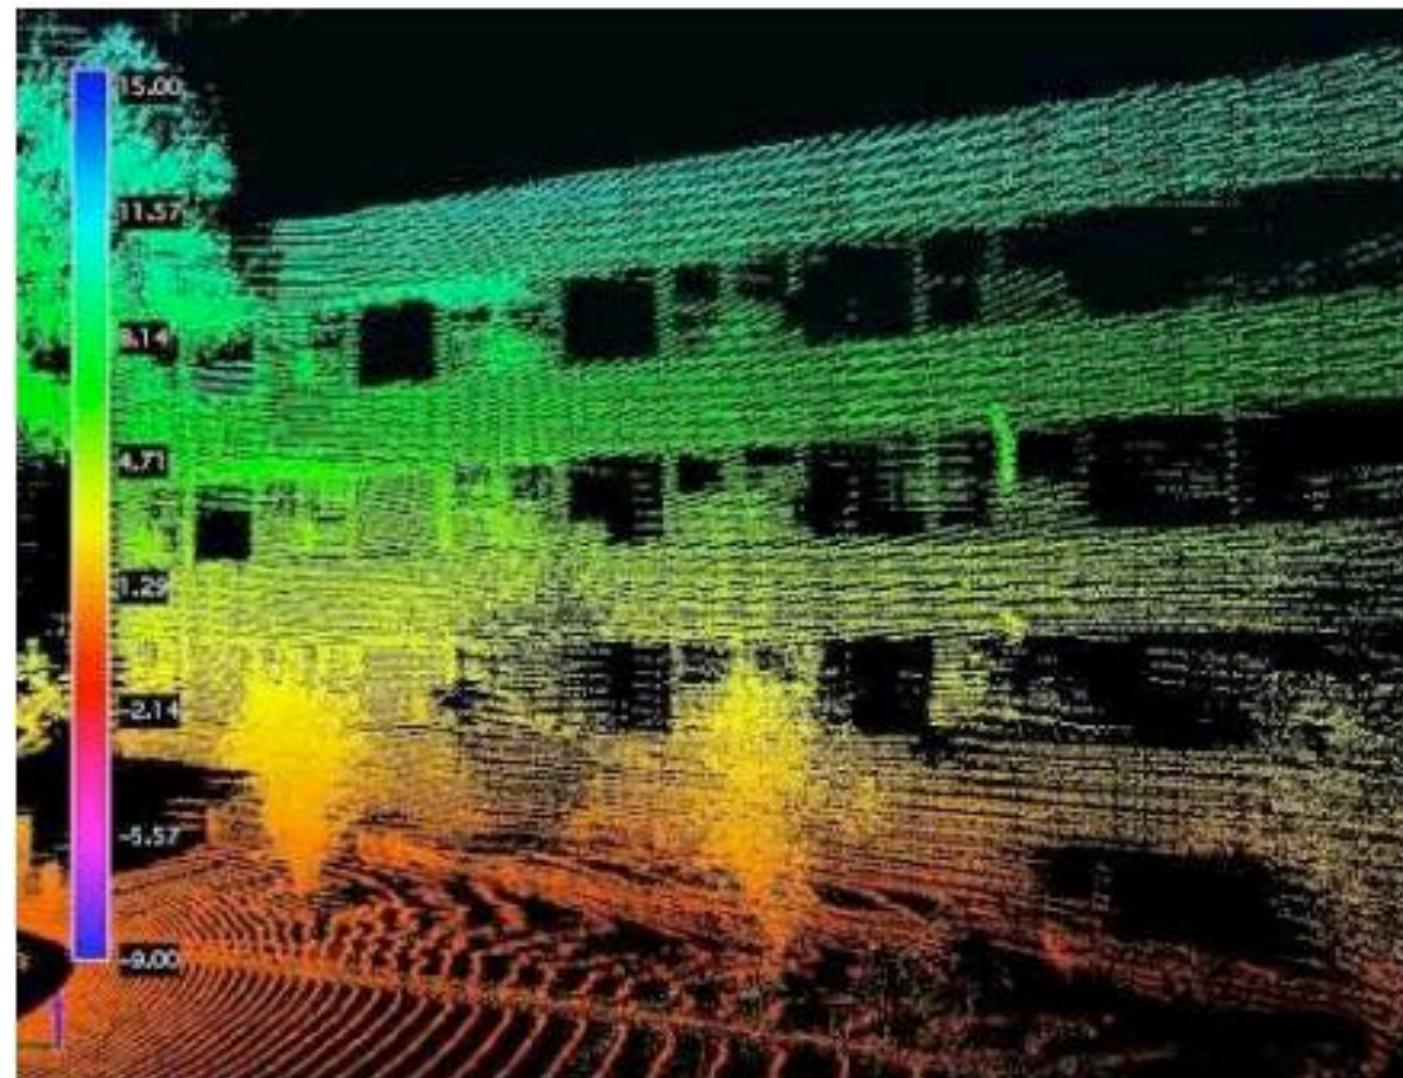
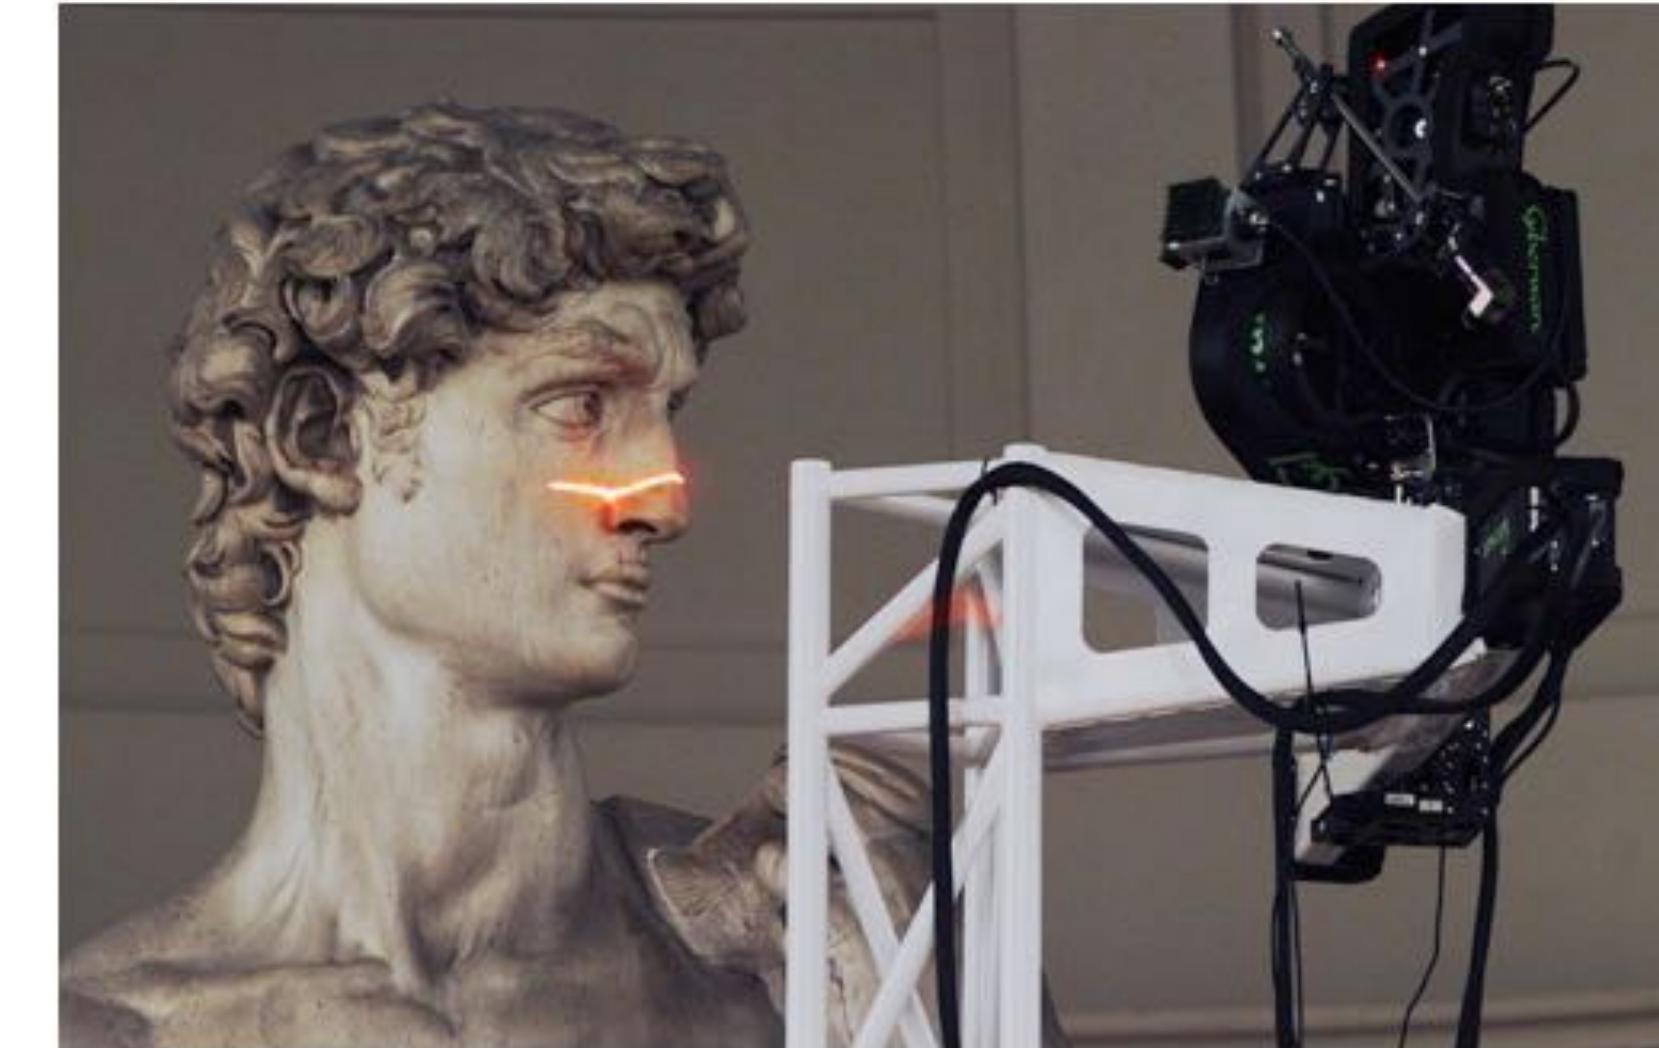
- Simplest representation: **only points**, no connectivity
- Collection of (x,y,z) coordinates, possibly with normal
- Points with orientation are called **surfels**



shading needs normals!

Acquiring Point Clouds

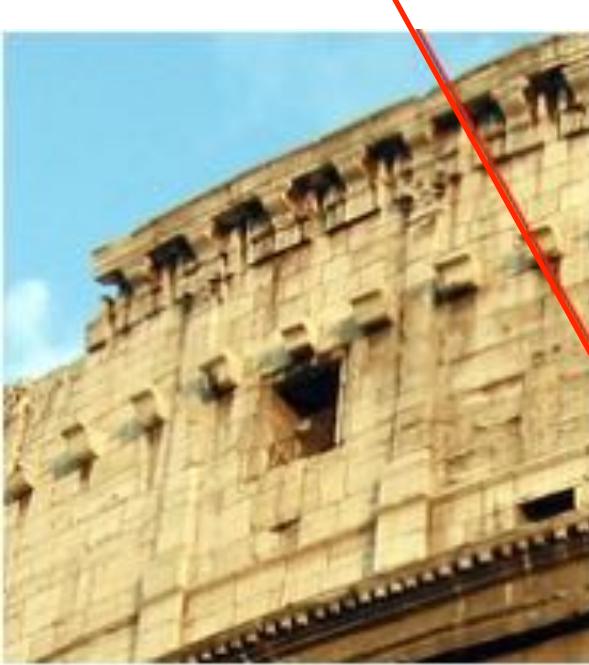
- From the real world
 - 3D scanning
 - Data is “striped”
 - Need multiple views to compensate occlusion



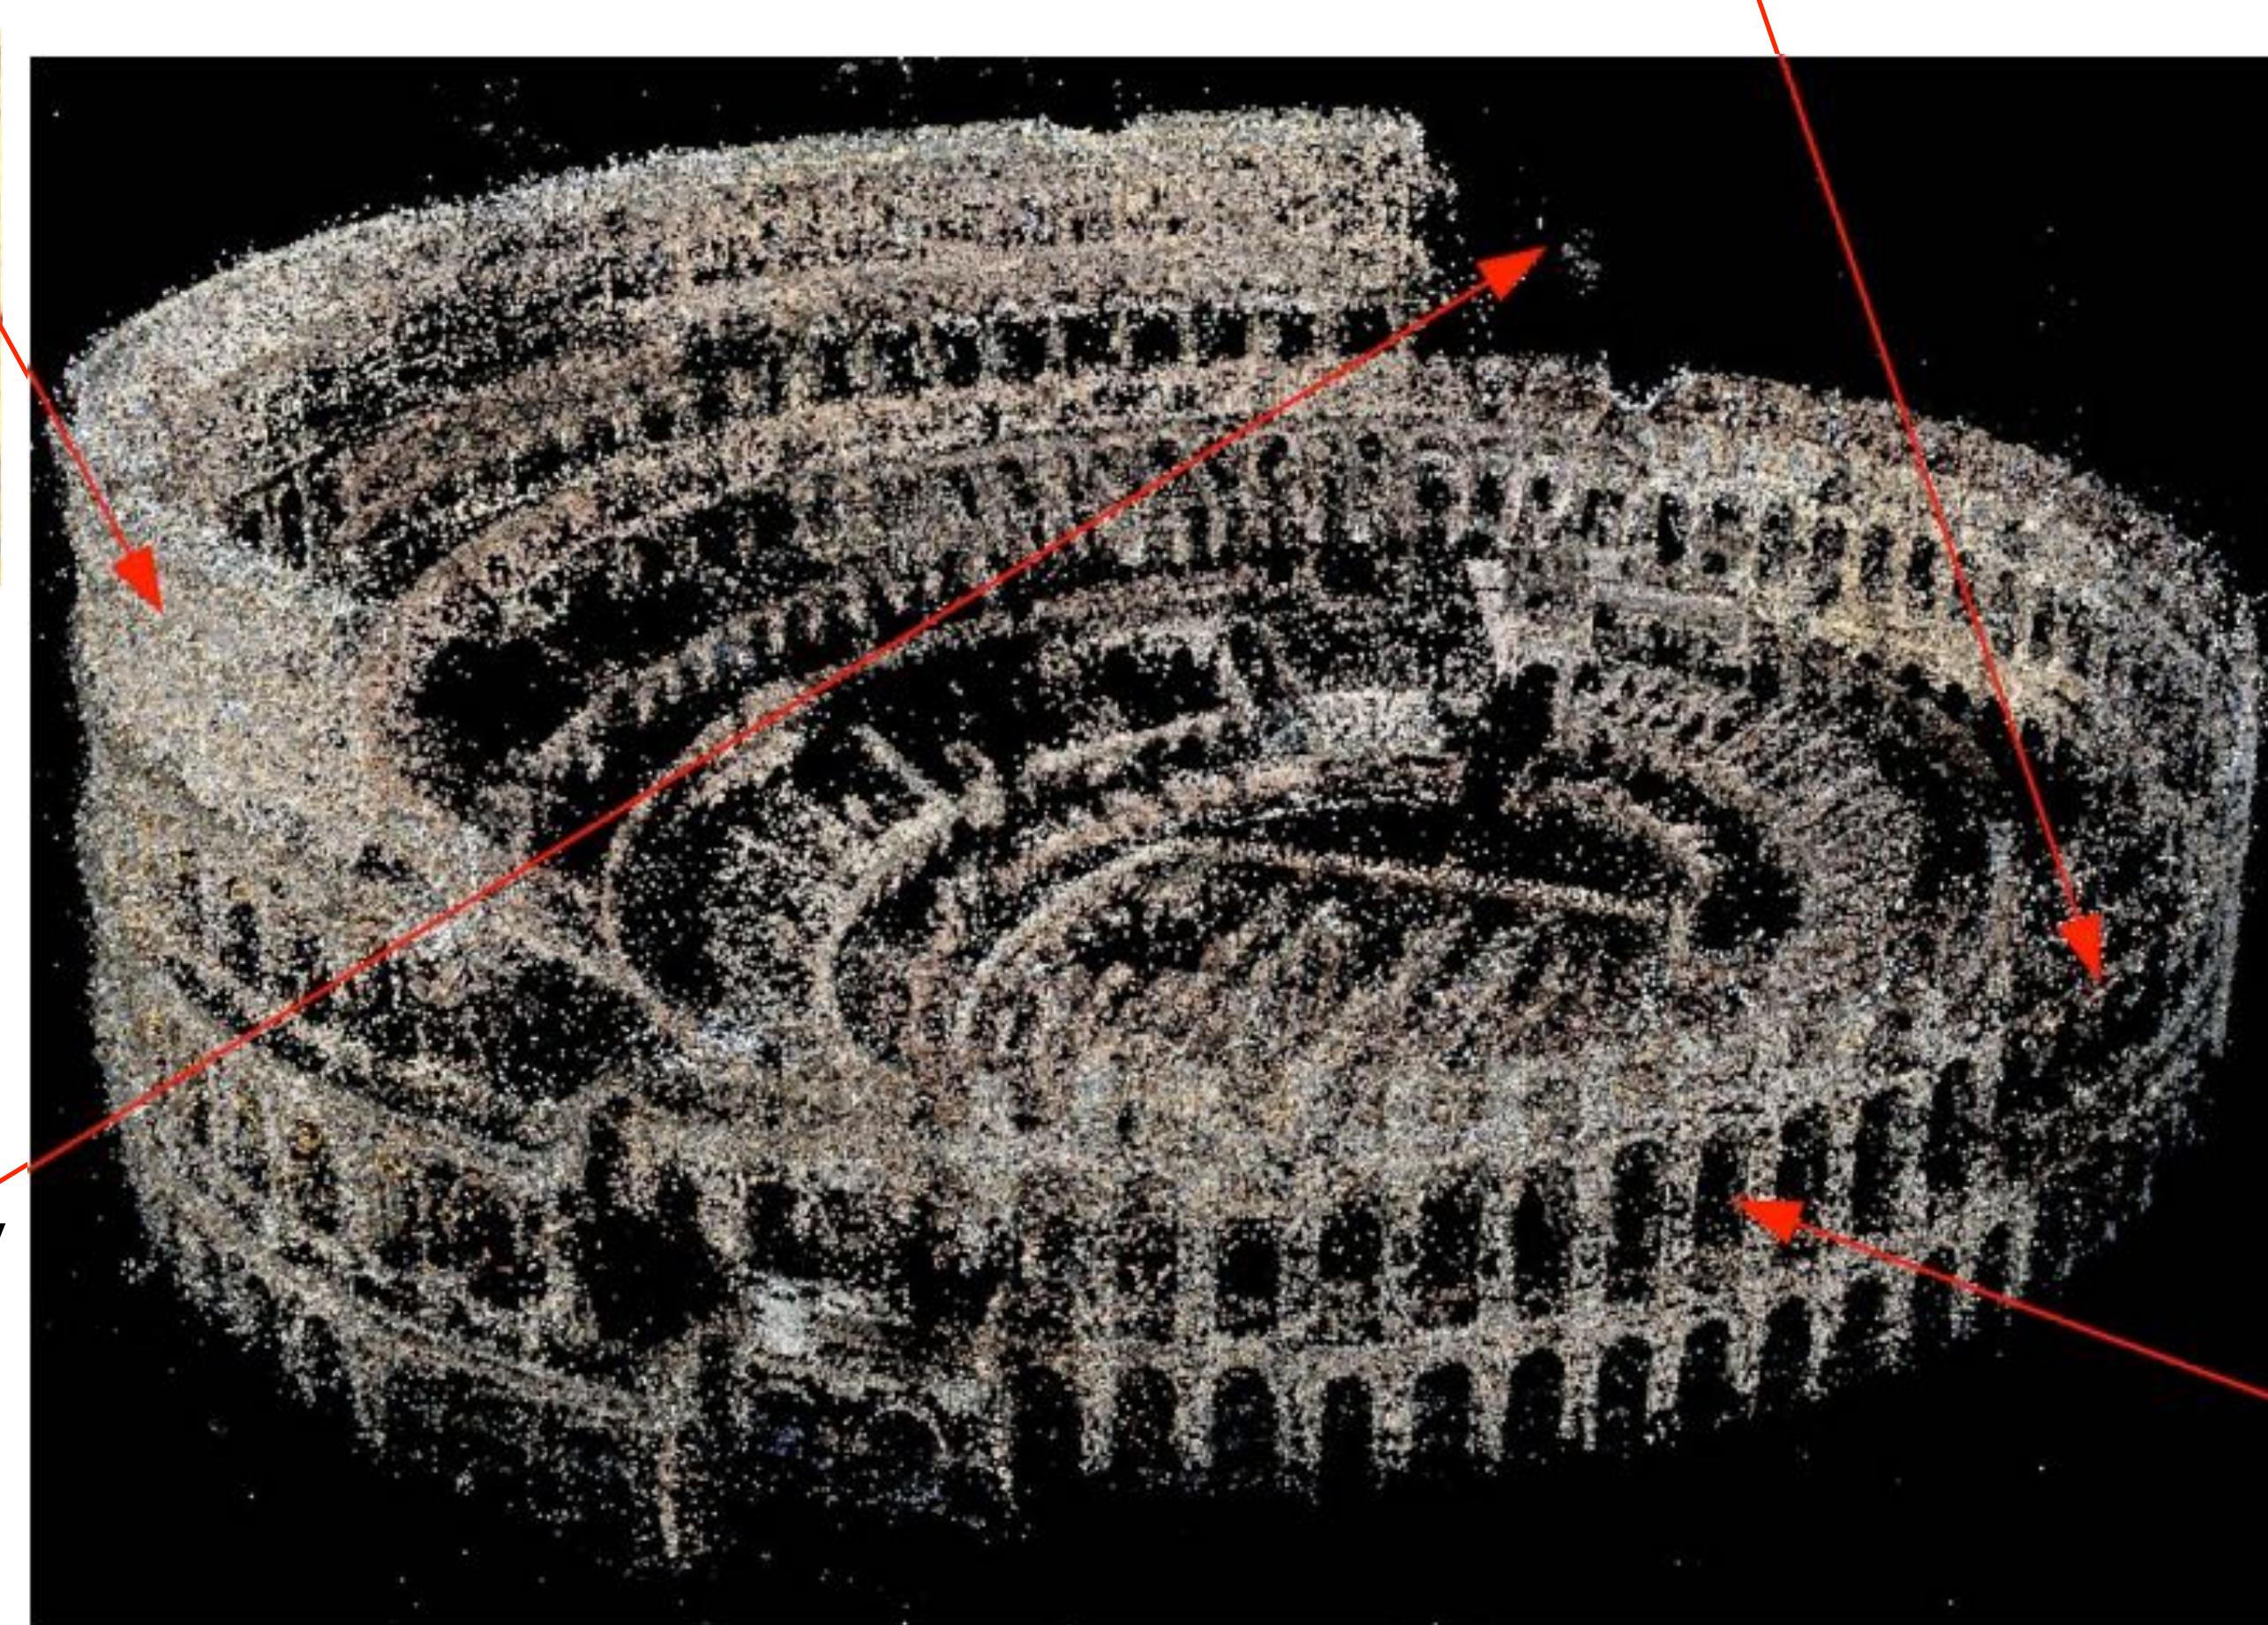
- Many techniques
 - Laser (LIDAR, e.g., StreetView)
 - Infrared (e.g., Kinect)
 - Stereo (e.g., Bundler)
- Many challenges: resolution, occlusion, noise, registration

Acquisition Challenges

Noise → Poor detail reproduction



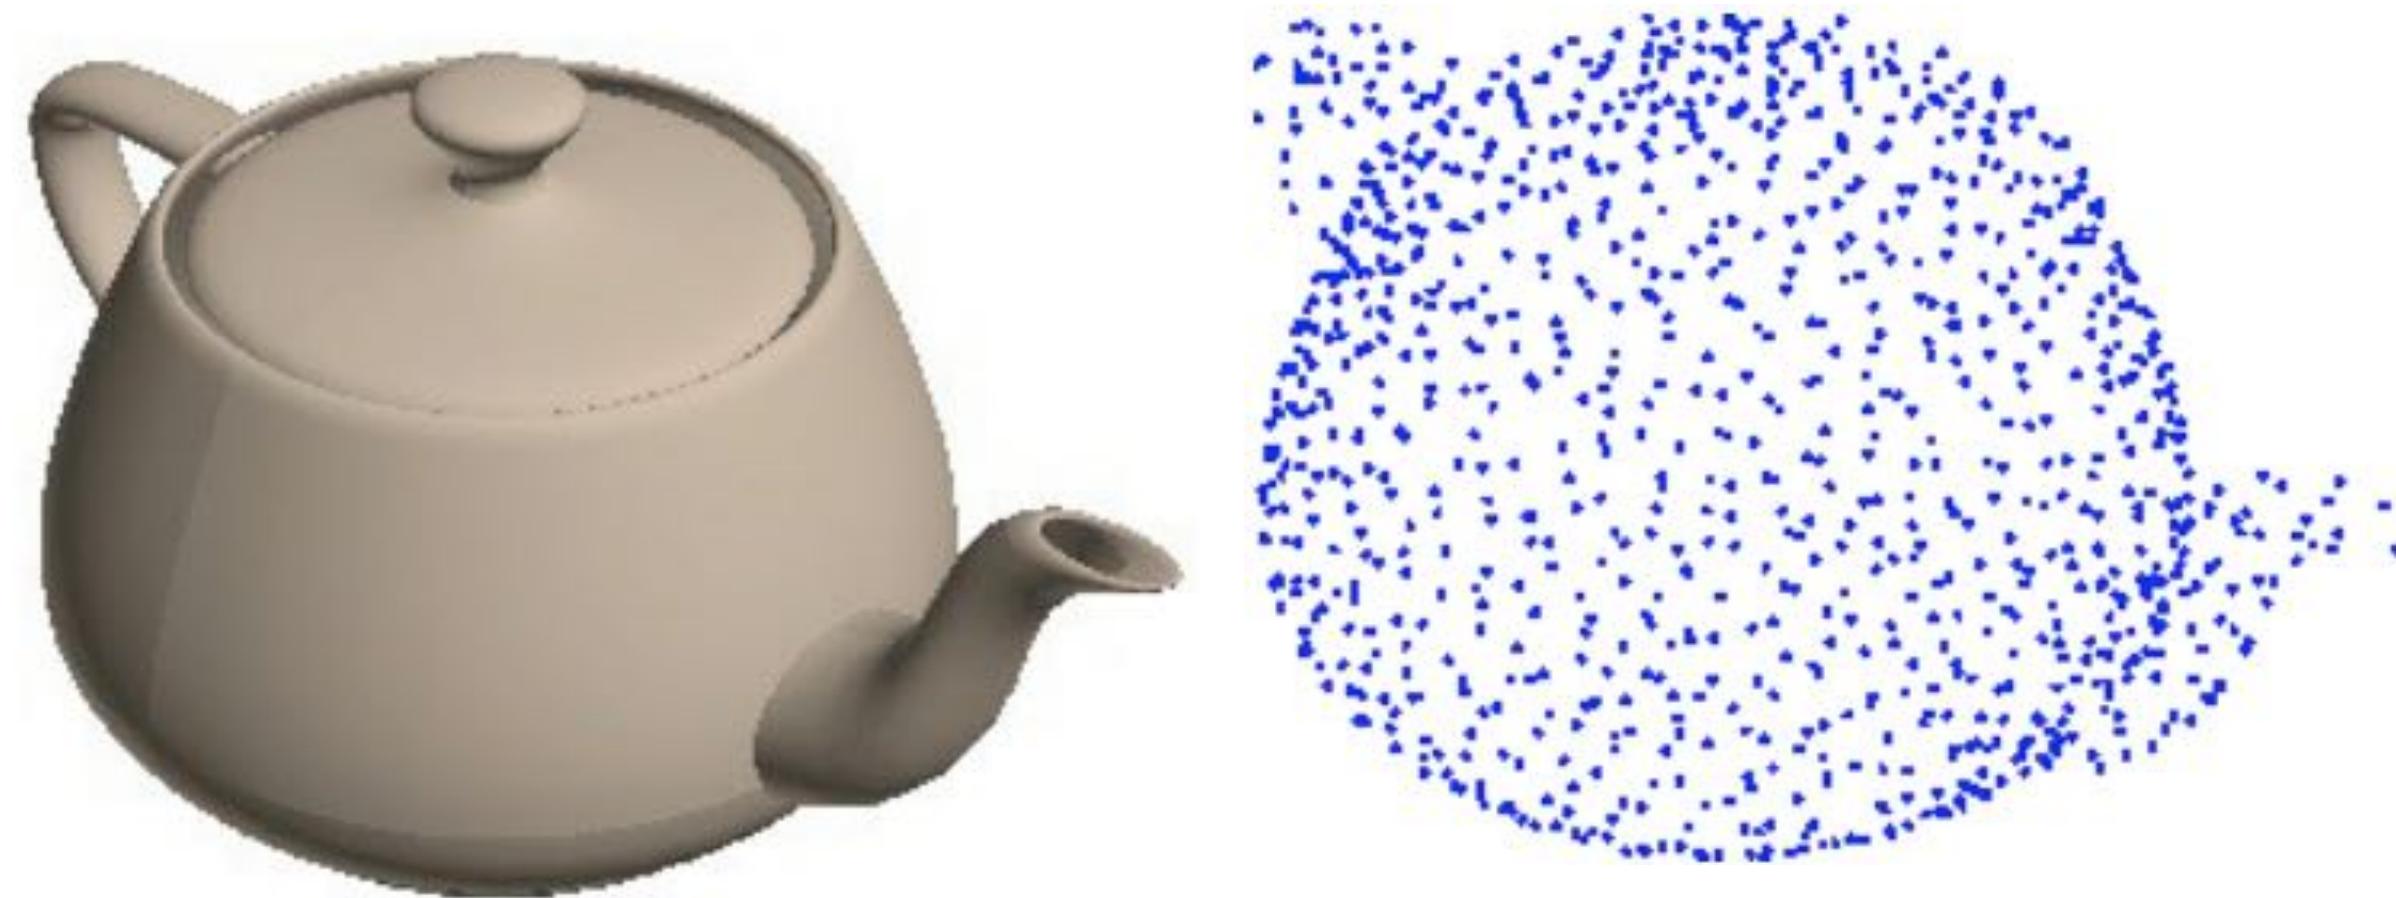
Low resolution further obscures detail



Occlusion → Interiors not captured

Acquiring Point Clouds

- From existing virtual shapes



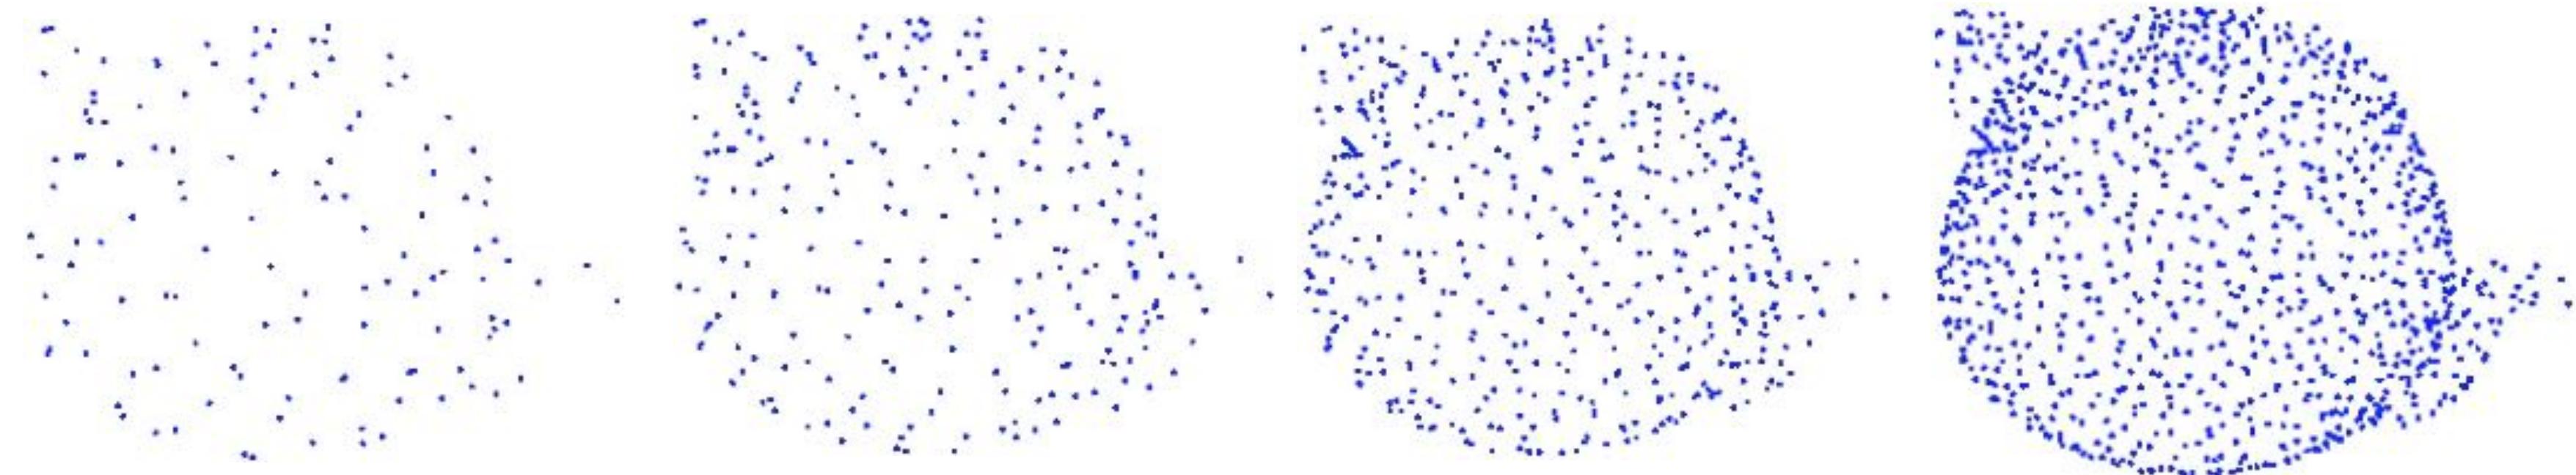
- Why would we want to do this?

Light-weight Shape Representation

Point cloud:

- Simple to understand
- Compact to store
- Generally easy to build algorithms

Yet already carries rich information!



$N = 125$

$N = 250$

$N = 500$

$N = 1000$

Outline

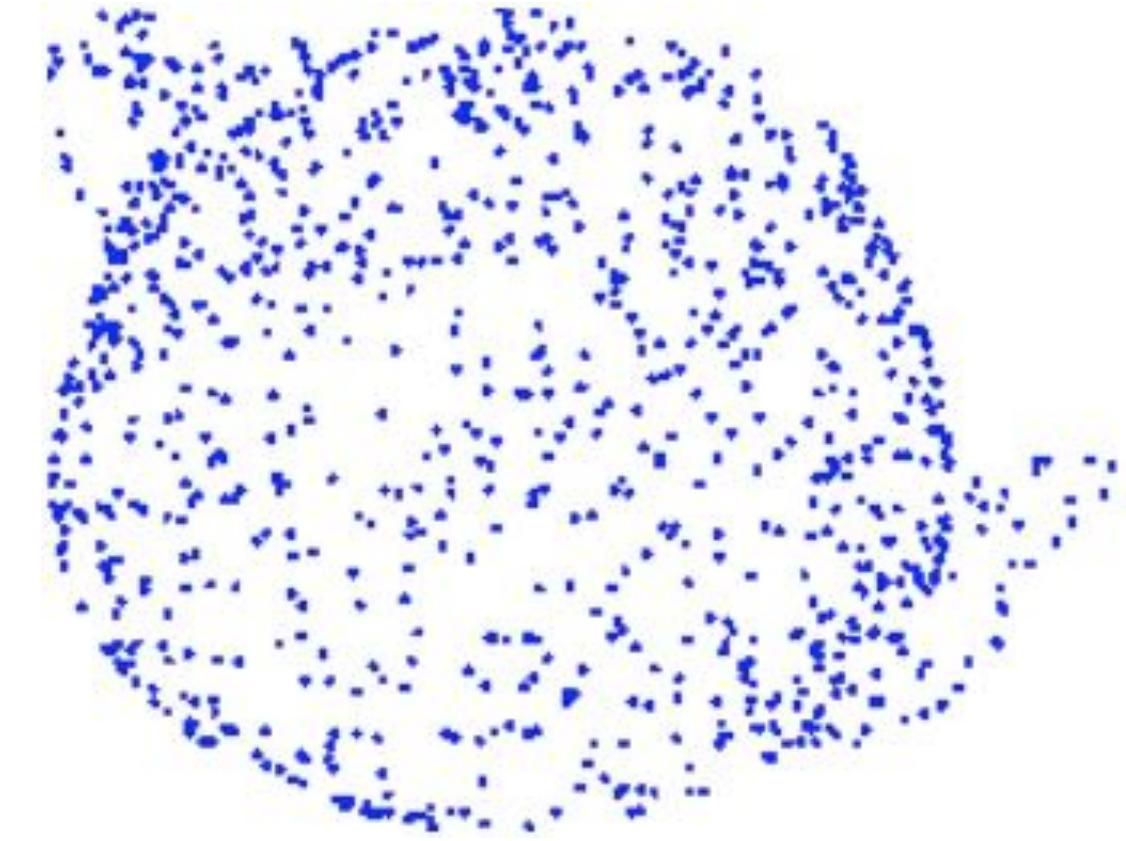
- Rasterized 3D representations
- Mesh
- *Point cloud*
 - Representation
 - *Sampling points on surfaces*
 - Normal computation
- Implicit representations

Application-based Sampling

- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap

Naive Strategy: Uniform Sampling

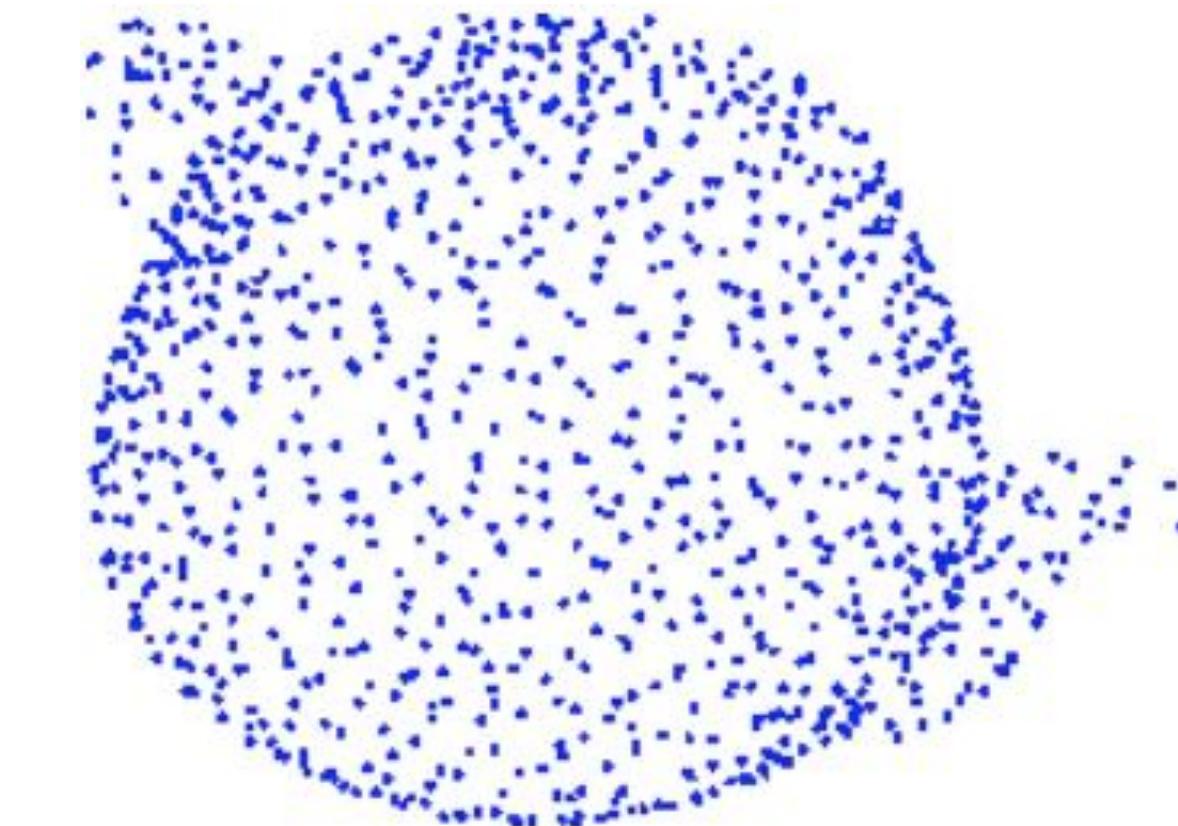
- Independent identically distributed (i.i.d.) samples by surface area:



- Usually the easiest to implement
- Issue: Irregularly spaced sampling

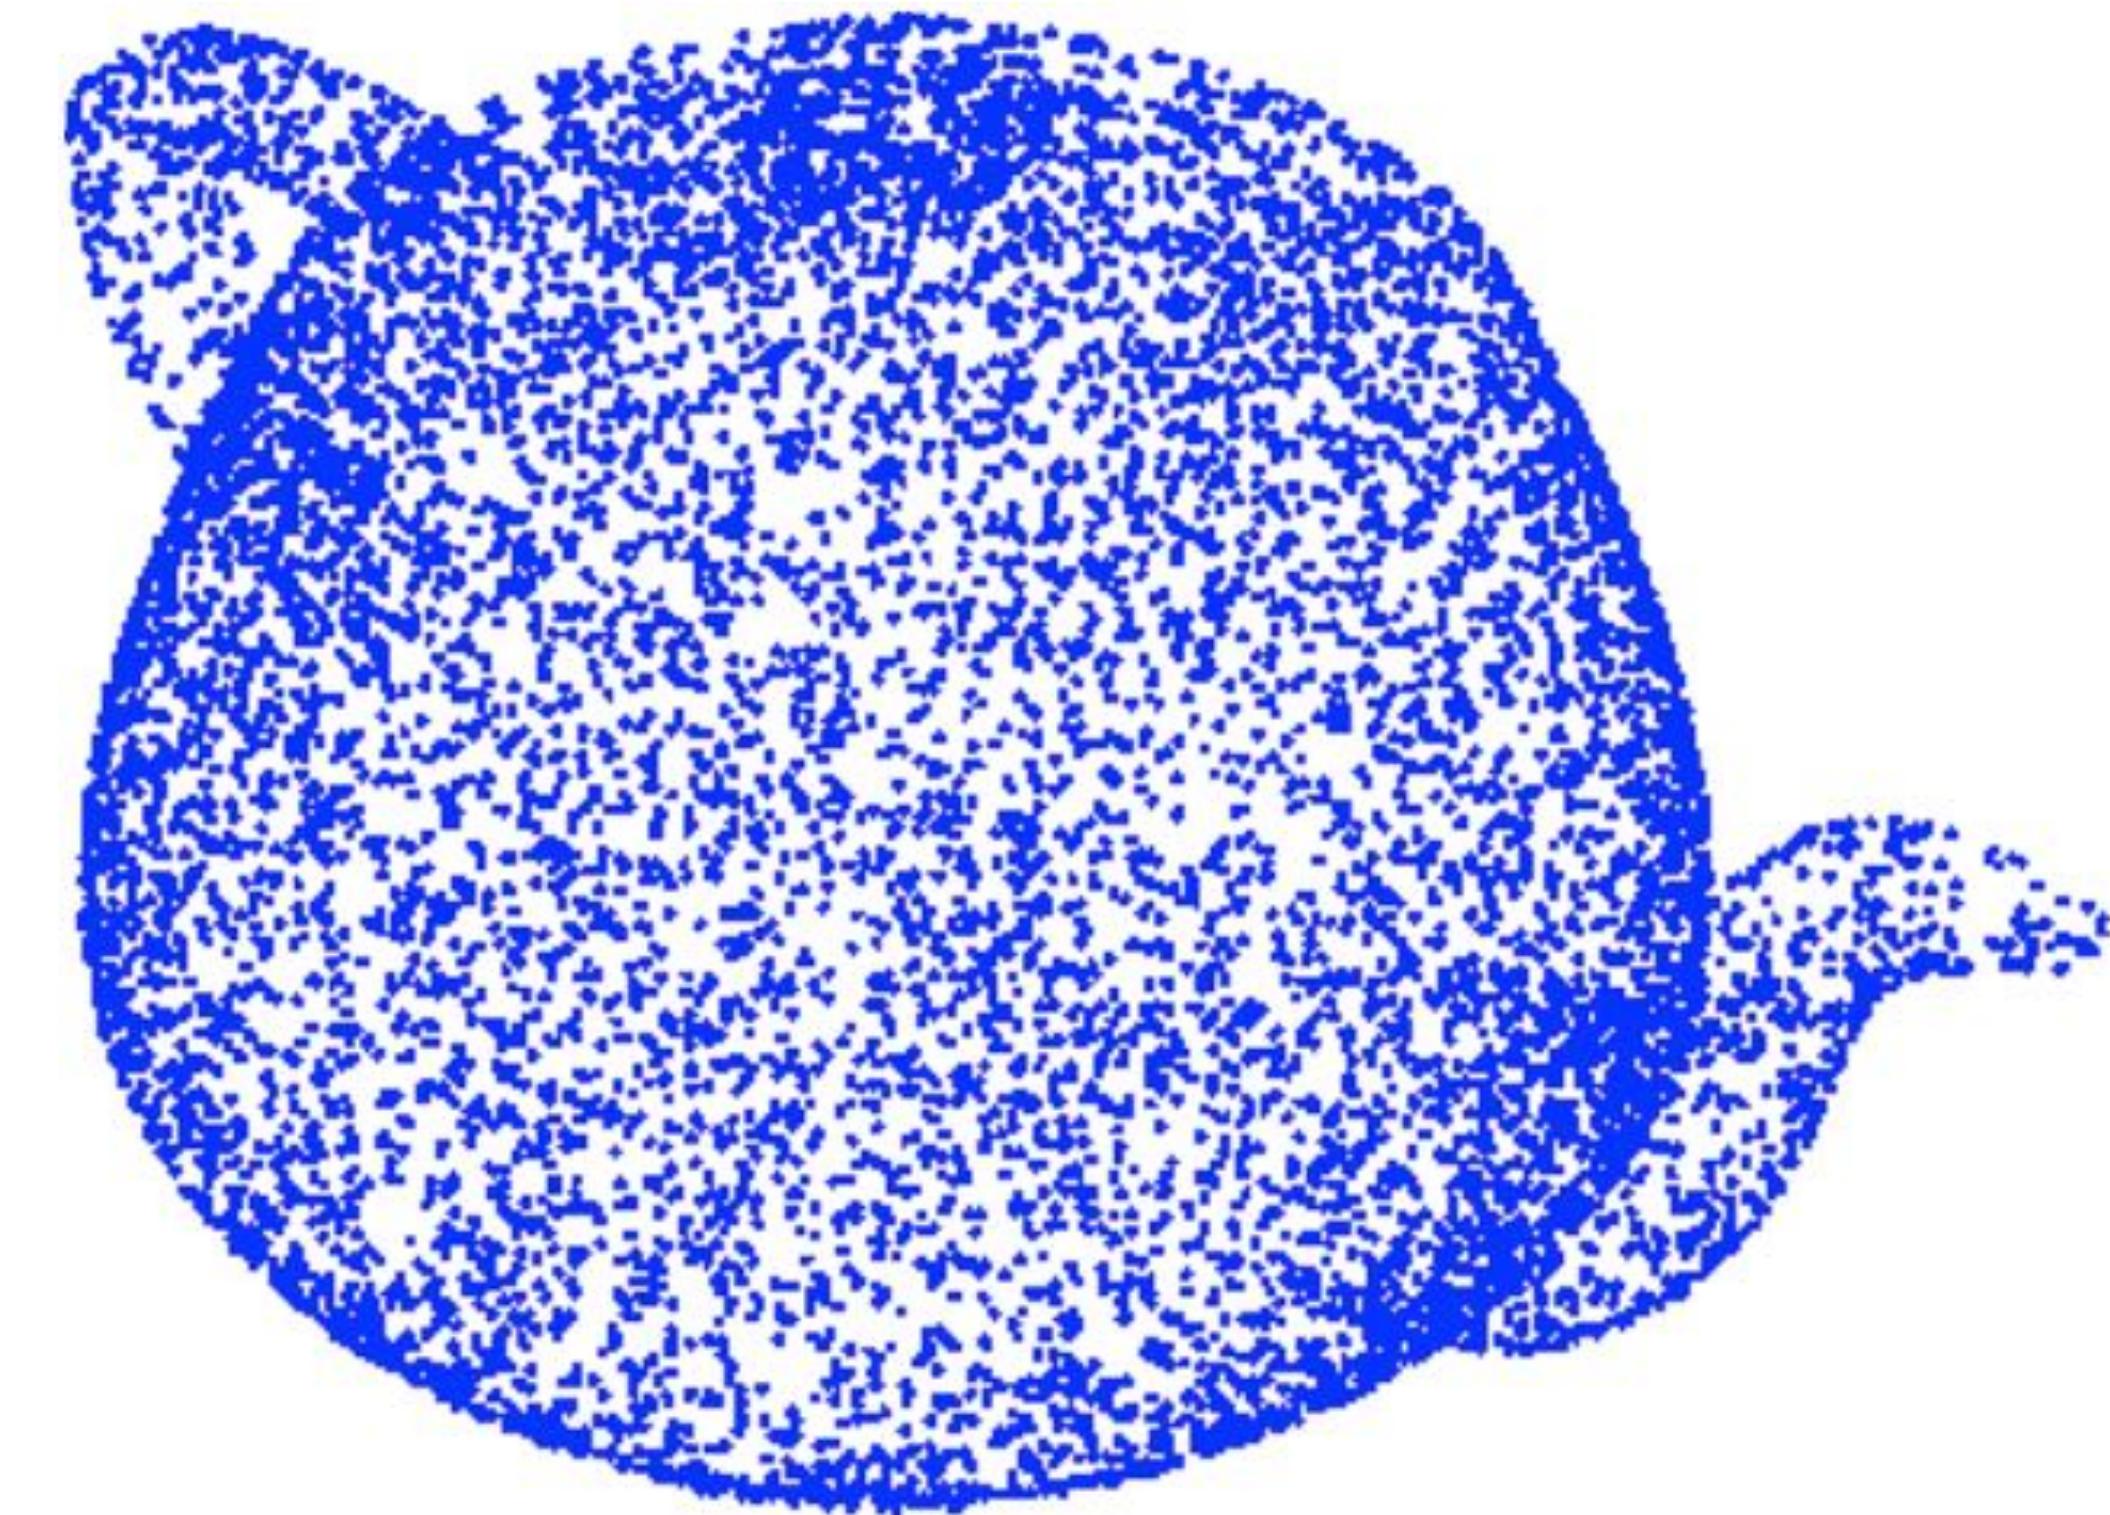
Farthest Point Sampling

- Goal: Sampled points are far away from each other
- NP-hard problem
- What is a greedy approximation method?



Iterative Furthest Point Sampling

- Step 1: Over sample the shape by any fast method
(e.g., uniformly sample $N=10,000$ i.i.d. samples)



Iterative Furthest Point Sampling

- Step 2: Iteratively select K points

U is the initial big set of points

$S = \{ \}$

add a random point from U to S

for $i=1$ to K

 find a point $u \in U$ with the largest distance to S

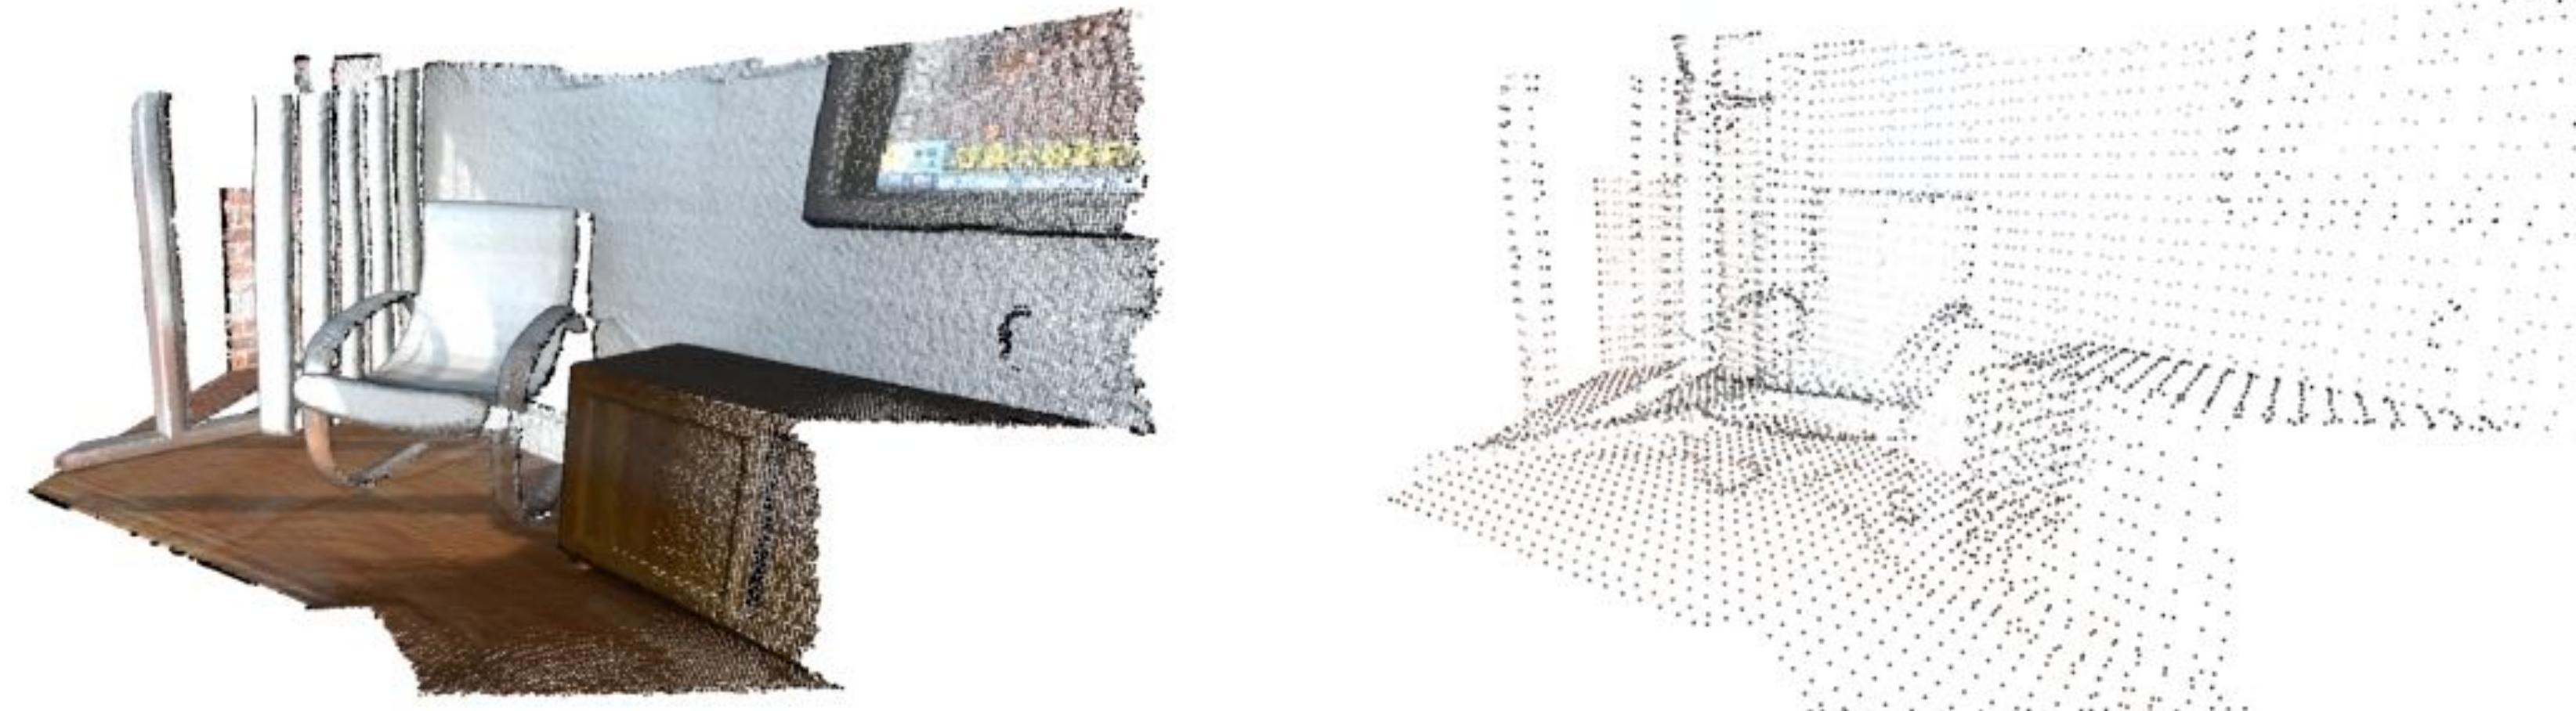
 add u to S

Issues Relevant to Speed

- Theoretically, naive implementation gives $\mathcal{O}(KN)$, but how to improve from $\mathcal{O}(KN)$ is an open question
- **Implementation can cause large speed difference**
 - As this is a serial algorithm in K , engineers optimize the efficiency in N (computing point-set distance)
 - CPU: Suggest using vectorization (e.g., numpy, `scipy.spatial.distance.cdist`)
 - GPU: By using shared memory, the complexity can be reduced to $\mathcal{O}(K(N/M + \log M))$, where M is the number of threads ($M=512$ in practice for modern GPU).

Voxel Downsampling

- Uses a *regular* voxel grid to downsample, taking one point per grid
- Allows higher parallelization level
- Generates regularly spaced sampling (with noticeable artifacts)



Credit to: Open3D

Issues Relevant to Speed

- Need to map each point to a bin. Often implemented as adding elements into a hash table
- $\mathcal{O}(N)$ (assuming that the inserting into hash table takes $O(1)$)
- On GPUs, parallelization reduces complexity of
 - Mapping each point to an integer value
 - Assign each value to an index so that the same value shares the same index
 - Aggregate indexes and form the output (called scattering in CUDA)

Application-based Sampling

- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap

Application-based Sampling

- For storage or analysis purposes (e.g., shape retrieval, signature extraction),
 - the objective is often to preserve surface information as much as possible
- For learning data generation purposes (e.g., sim2real),
 - the objective is often to minimize virtual-real domain gap

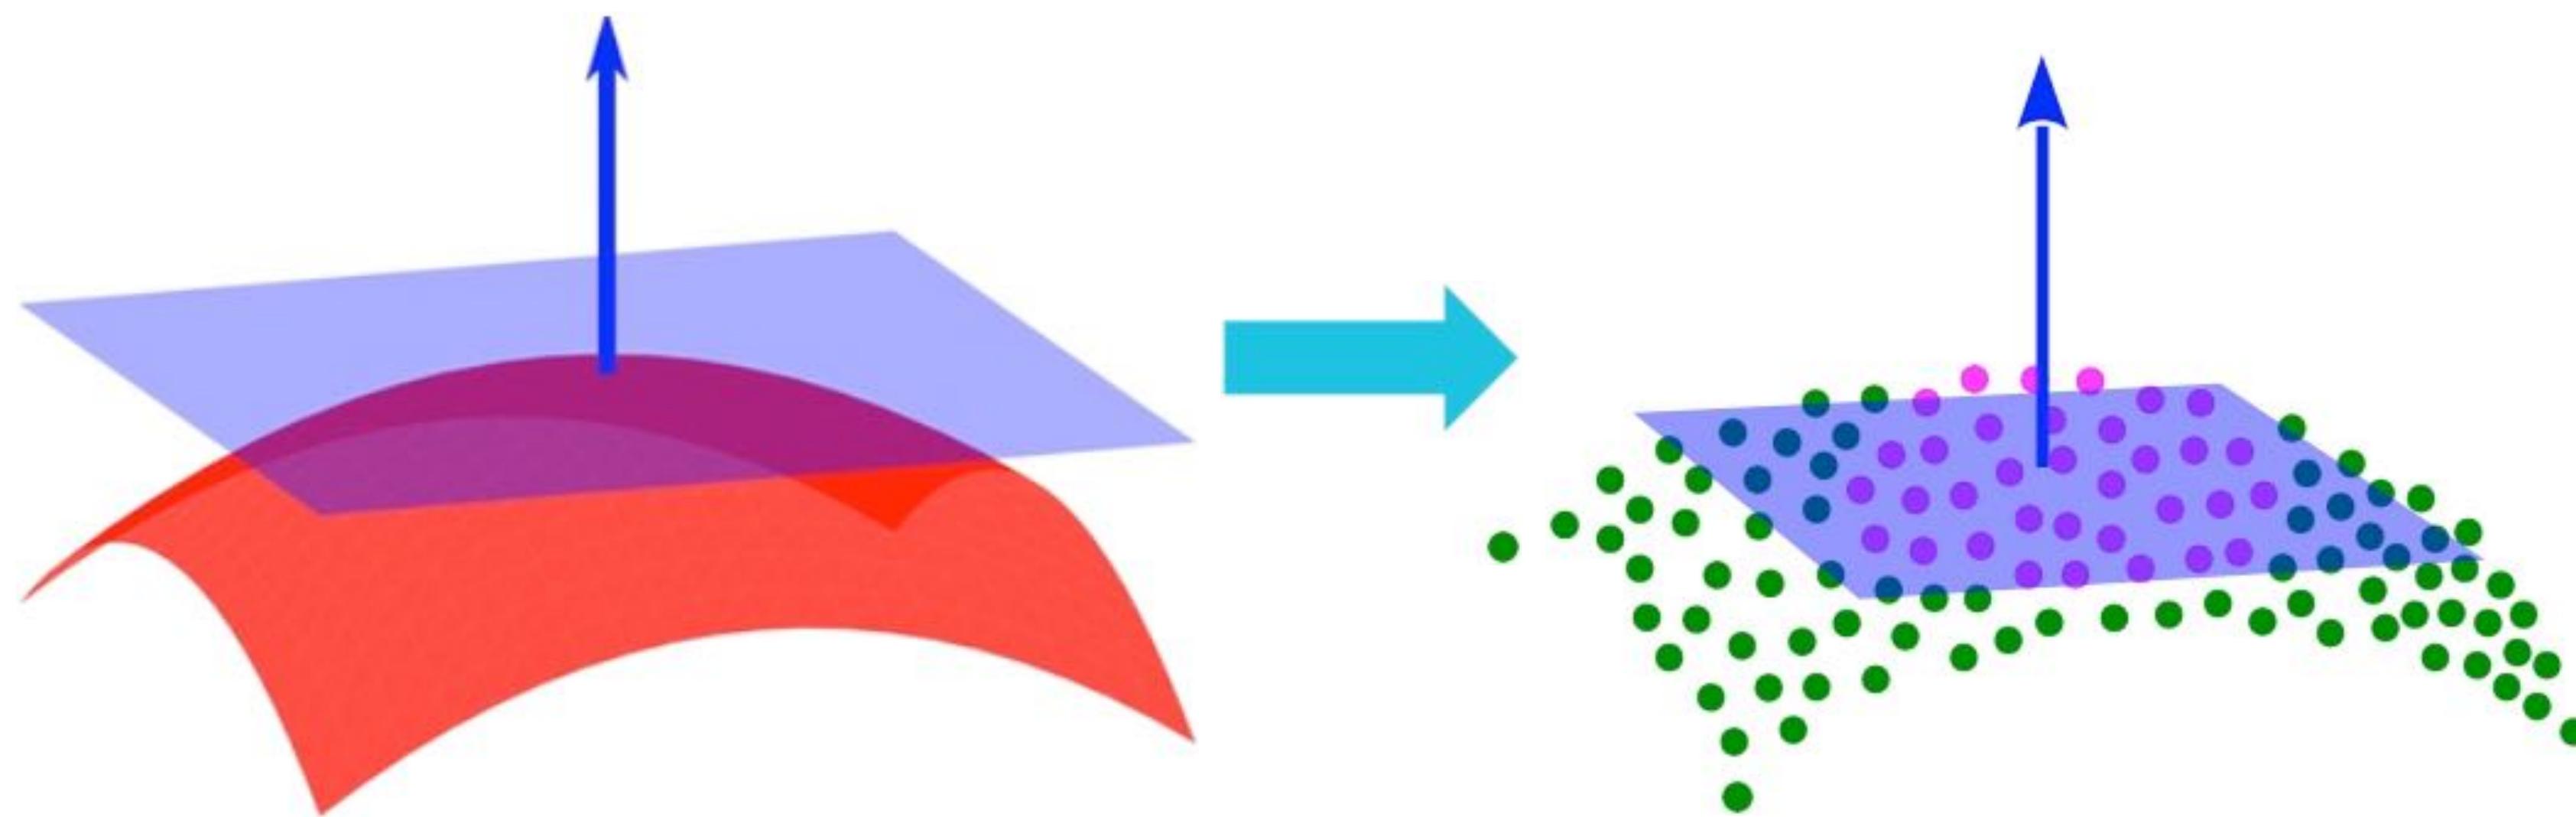
*A good research topic
(e.g., GAN? Adversarial training? Differentiable sampling?)*

Outline

- Rasterized 3D representations
- Mesh
- *Point cloud*
 - Representation
 - Sampling points on surfaces
 - *Normal computation*
- Implicit representations

Estimating Normals

- Plane-fitting: find the plane that best fits the neighborhood of a point of interest



Least-square Formulation

- Assume the plane equation is:

$$w^T(x - c) = 0 \quad \text{with} \quad \|w\| = 1$$

- Plane-fitting solves the least square problem:

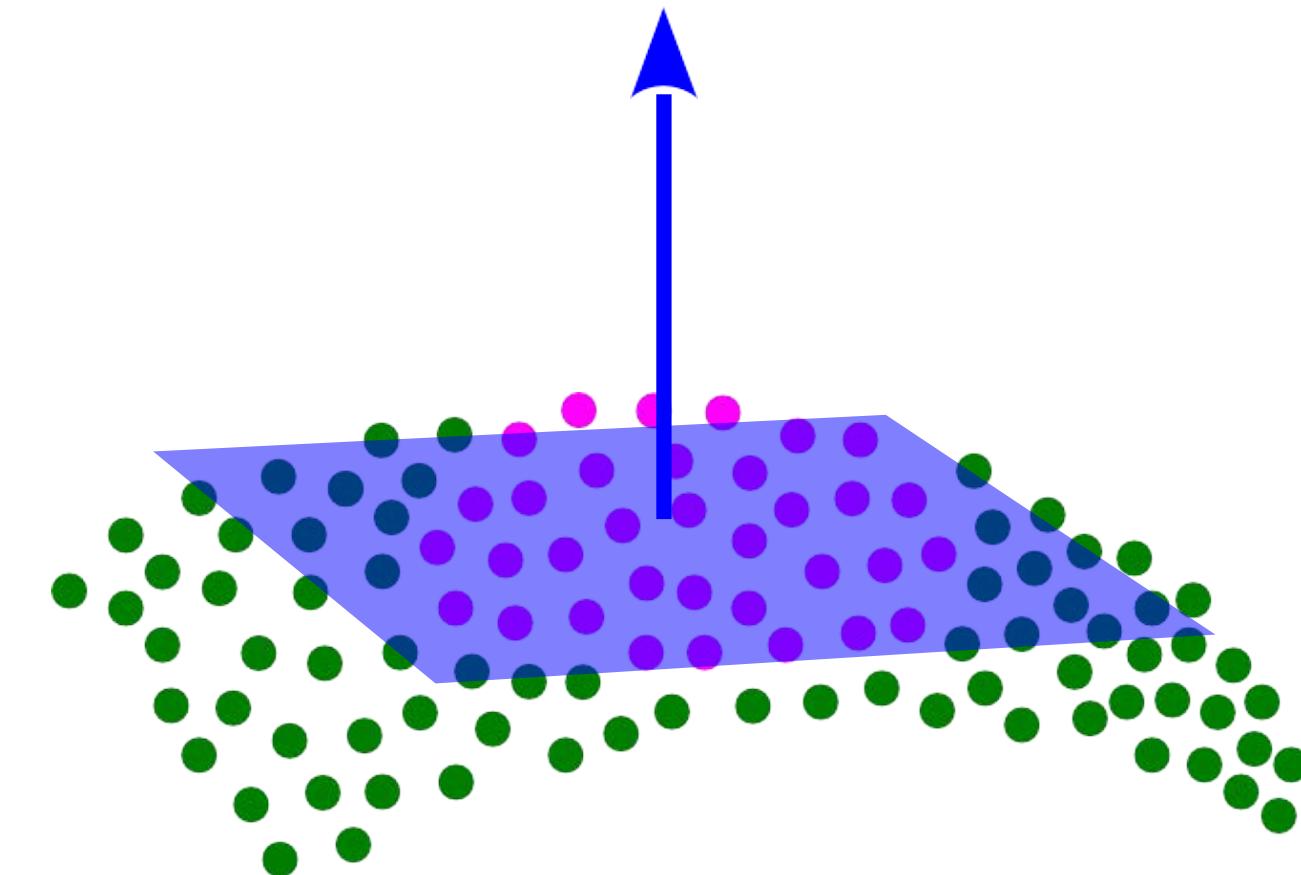
$$\begin{aligned} & \text{minimize} && \sum_i \|w^T(x_i - c)\|_2^2 \\ & \text{subject to} && \|w\|^2 = 1 \end{aligned}$$

where $\{x_i\}$ is the neighborhood of a point x
that you query the normal

Least-square Formulation

- Doing Lagrangian multiplier and the solution is:

- Let $M = \sum_i (x_i - \bar{x})(x_i - \bar{x})^T$ and $\bar{x} = \frac{1}{n} \sum_i x_i$,
- w : the smallest eigenvector of M
- $c = w^T \bar{x}$



- w also corresponds to the third principal component of M (yet another usage of PCA)
 - Where are the first and second principal components?

Summary of Normal Computation

- The normal of a point cloud can be computed through PCA over a local neighborhood
- Remark:
 - The choice of neighborhood size is important
 - When outlier points exist, RANSAC can improve quality

Outline

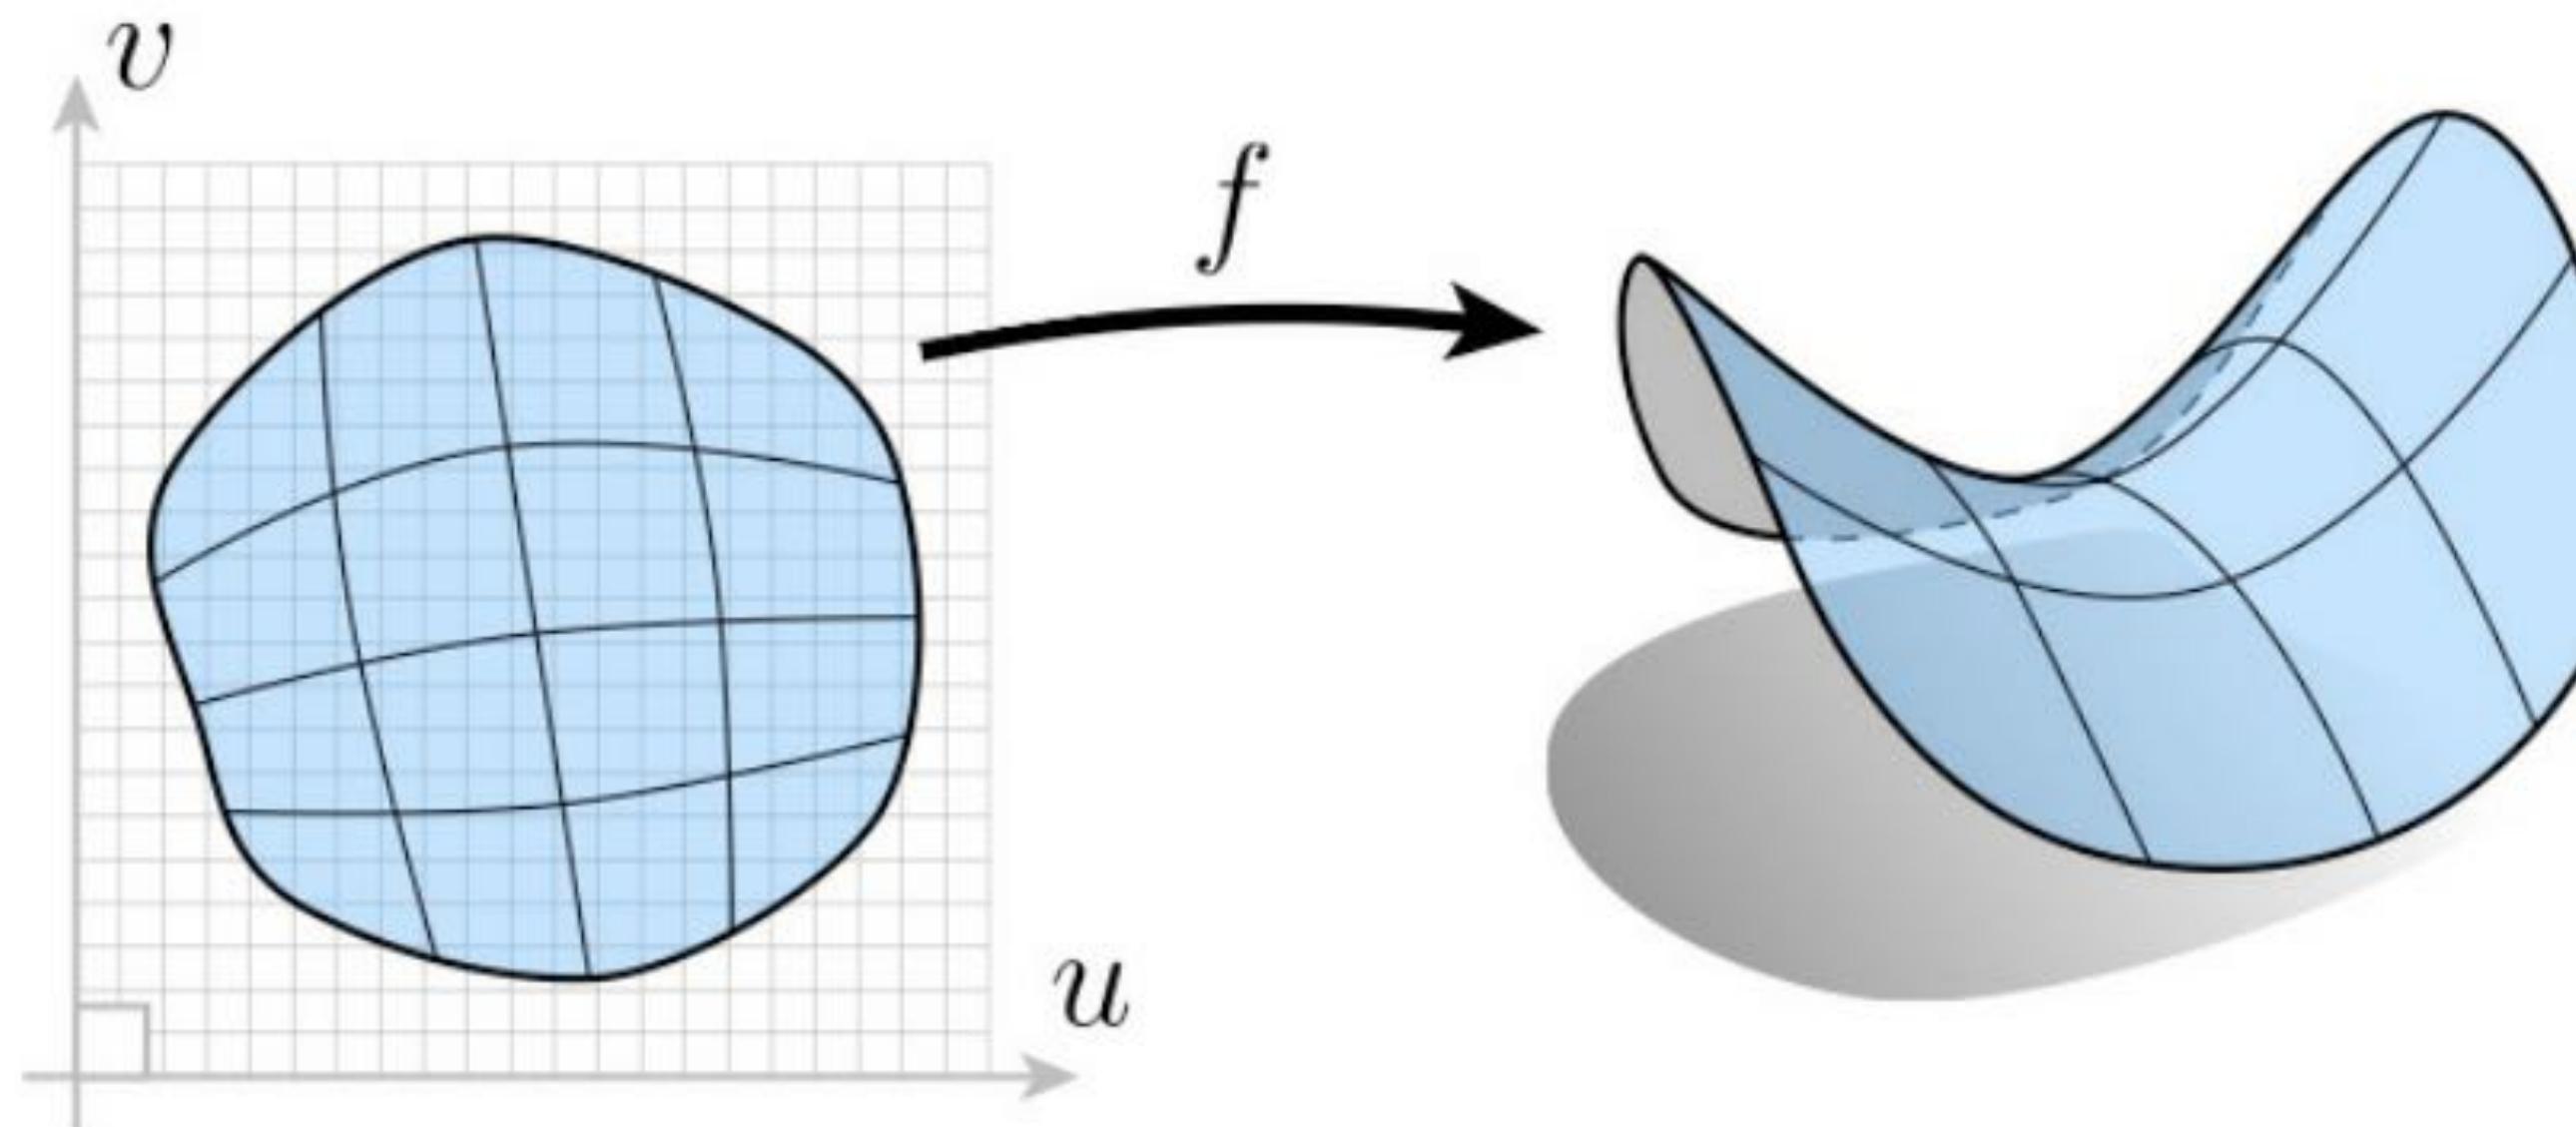
- Rasterized 3D representations
- Mesh
- Point cloud
- *Implicit representations*

Explicit Representation of Geometry

All points are given directly

Generally:

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^3; (u, v) \mapsto (x, y, z)$$

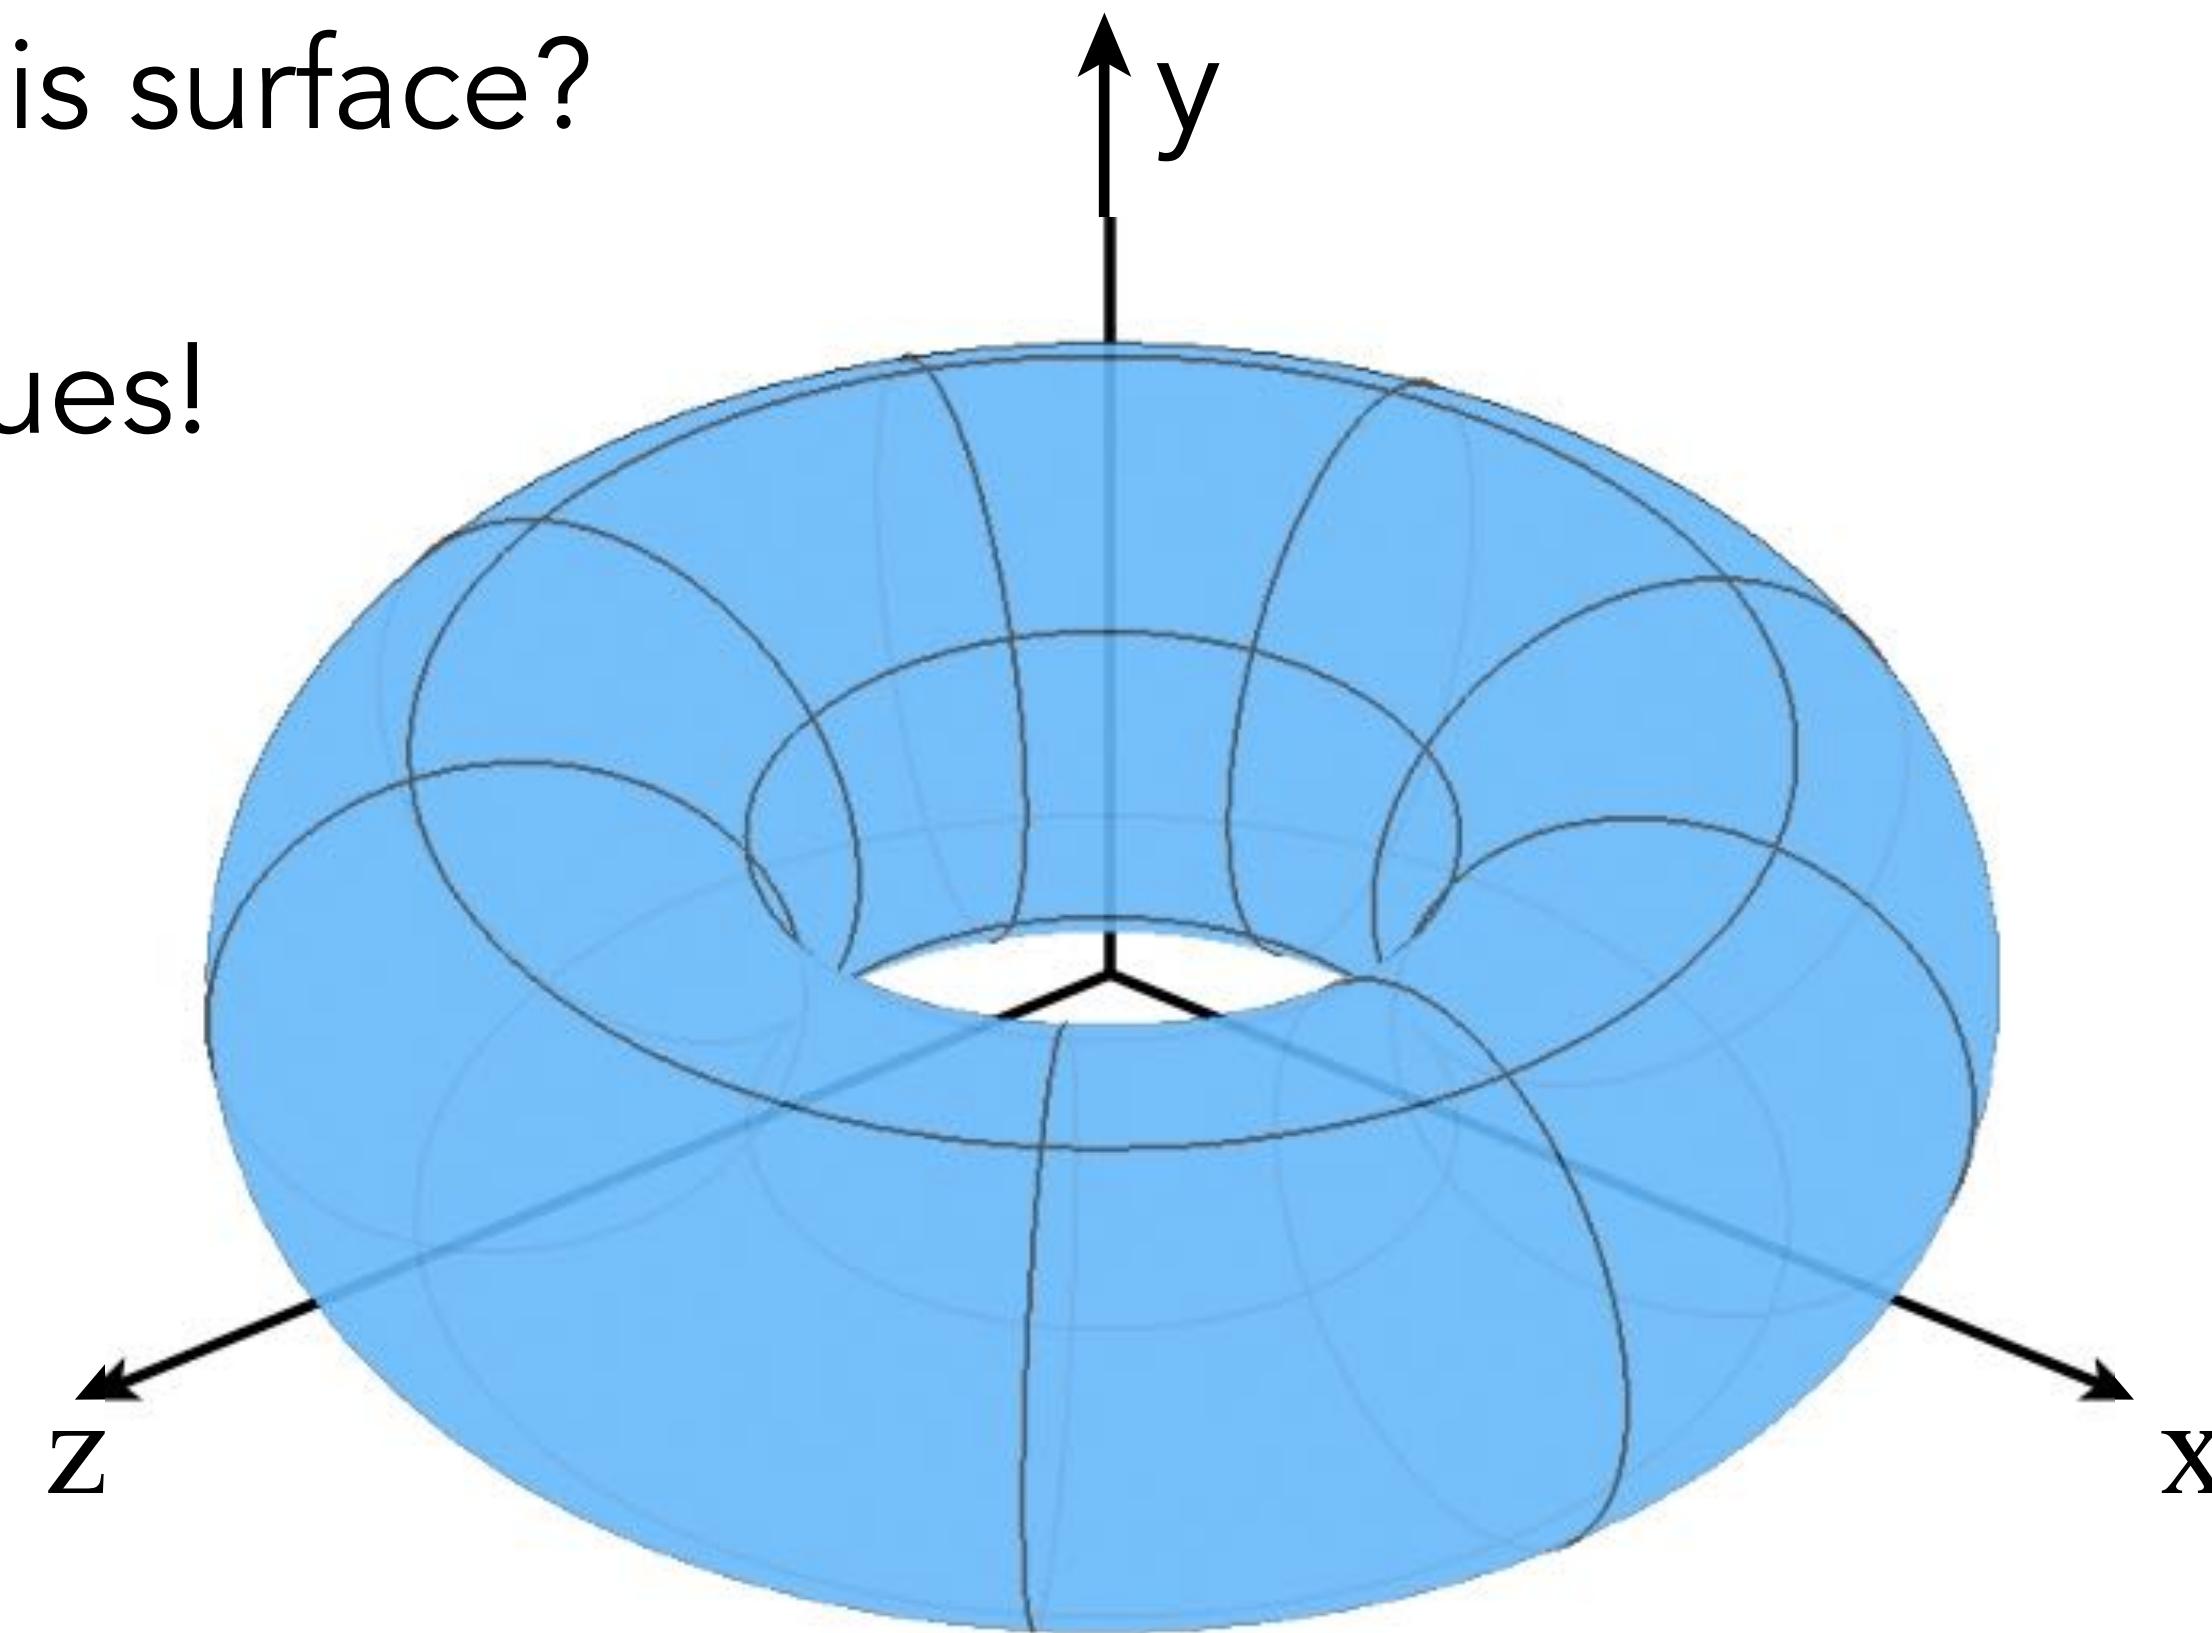


Explicit Surface – Sampling Is Easy

$$f(u, v) = ((2 + \cos u)\cos v, (2 + \cos u)\sin v, \sin u)$$

What points lie on this surface?

Just plug in (u, v) values!

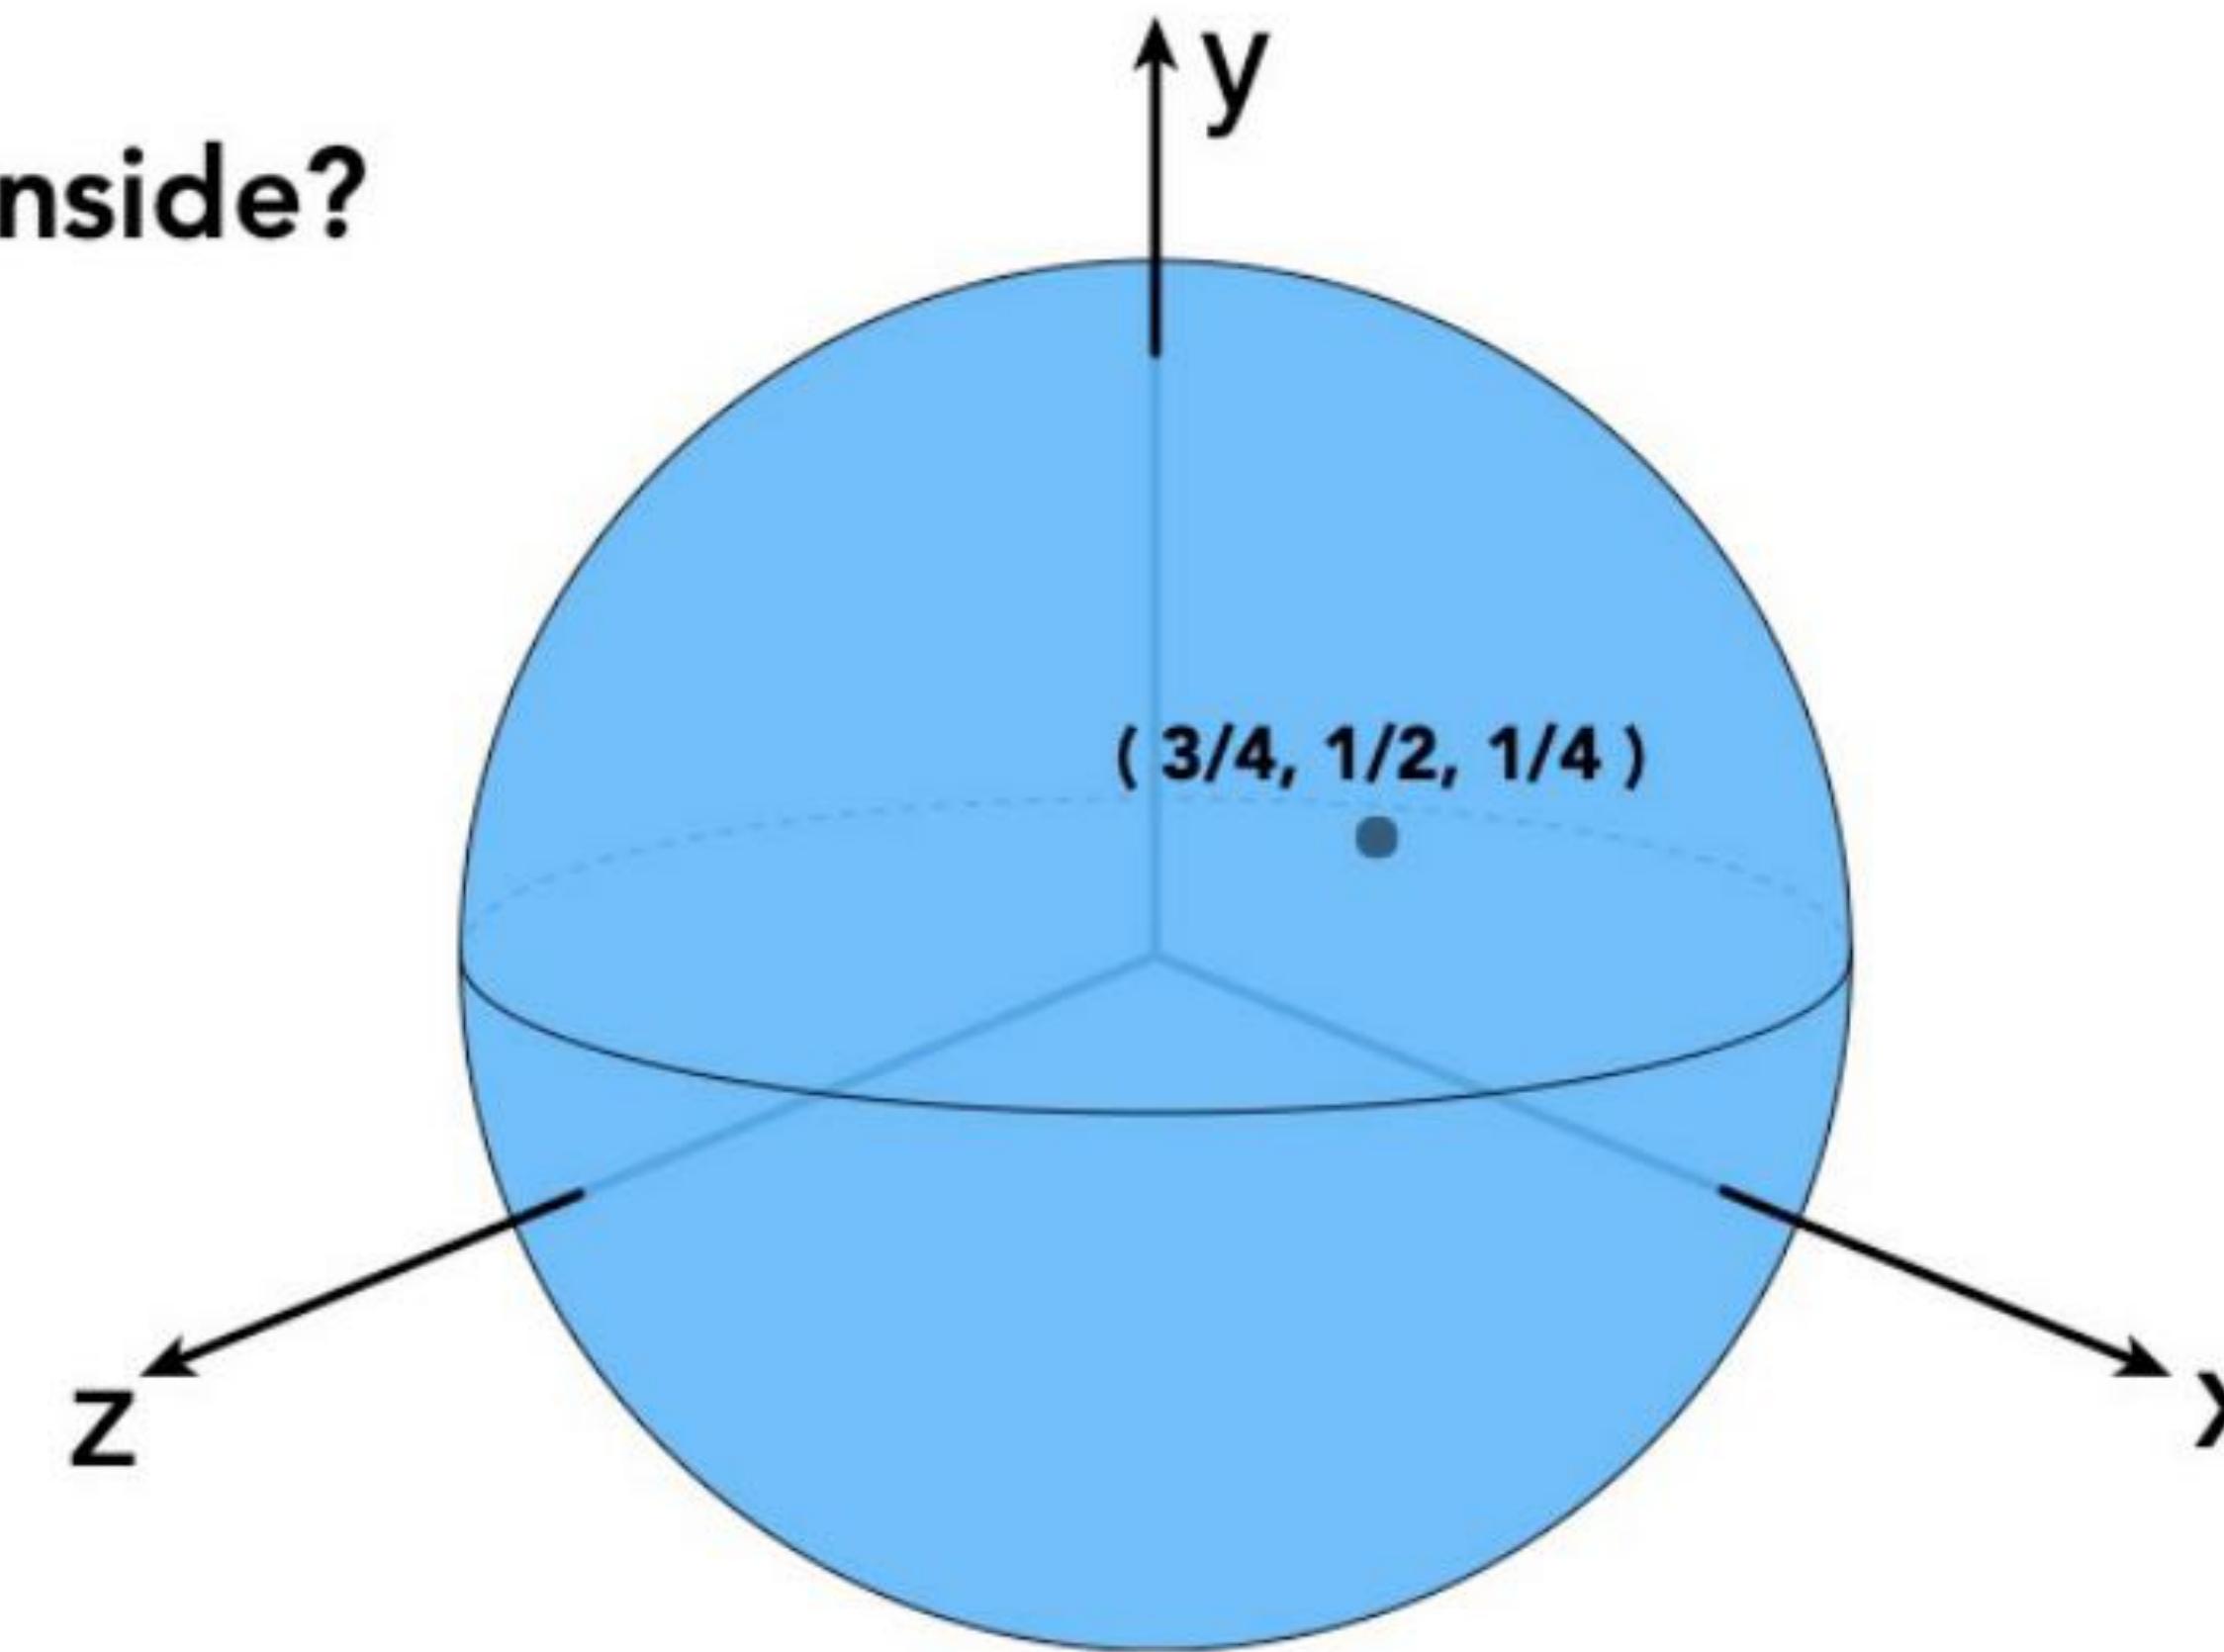


Explicit representations make some tasks easy

Explicit Surface - Inside/Outside Test Hard

$$f(u, v) = (\cos u \sin v, \sin u \sin v, \cos v)$$

Is $(3/4, 1/2, 1/4)$ inside?



Some tasks are hard with explicit representations

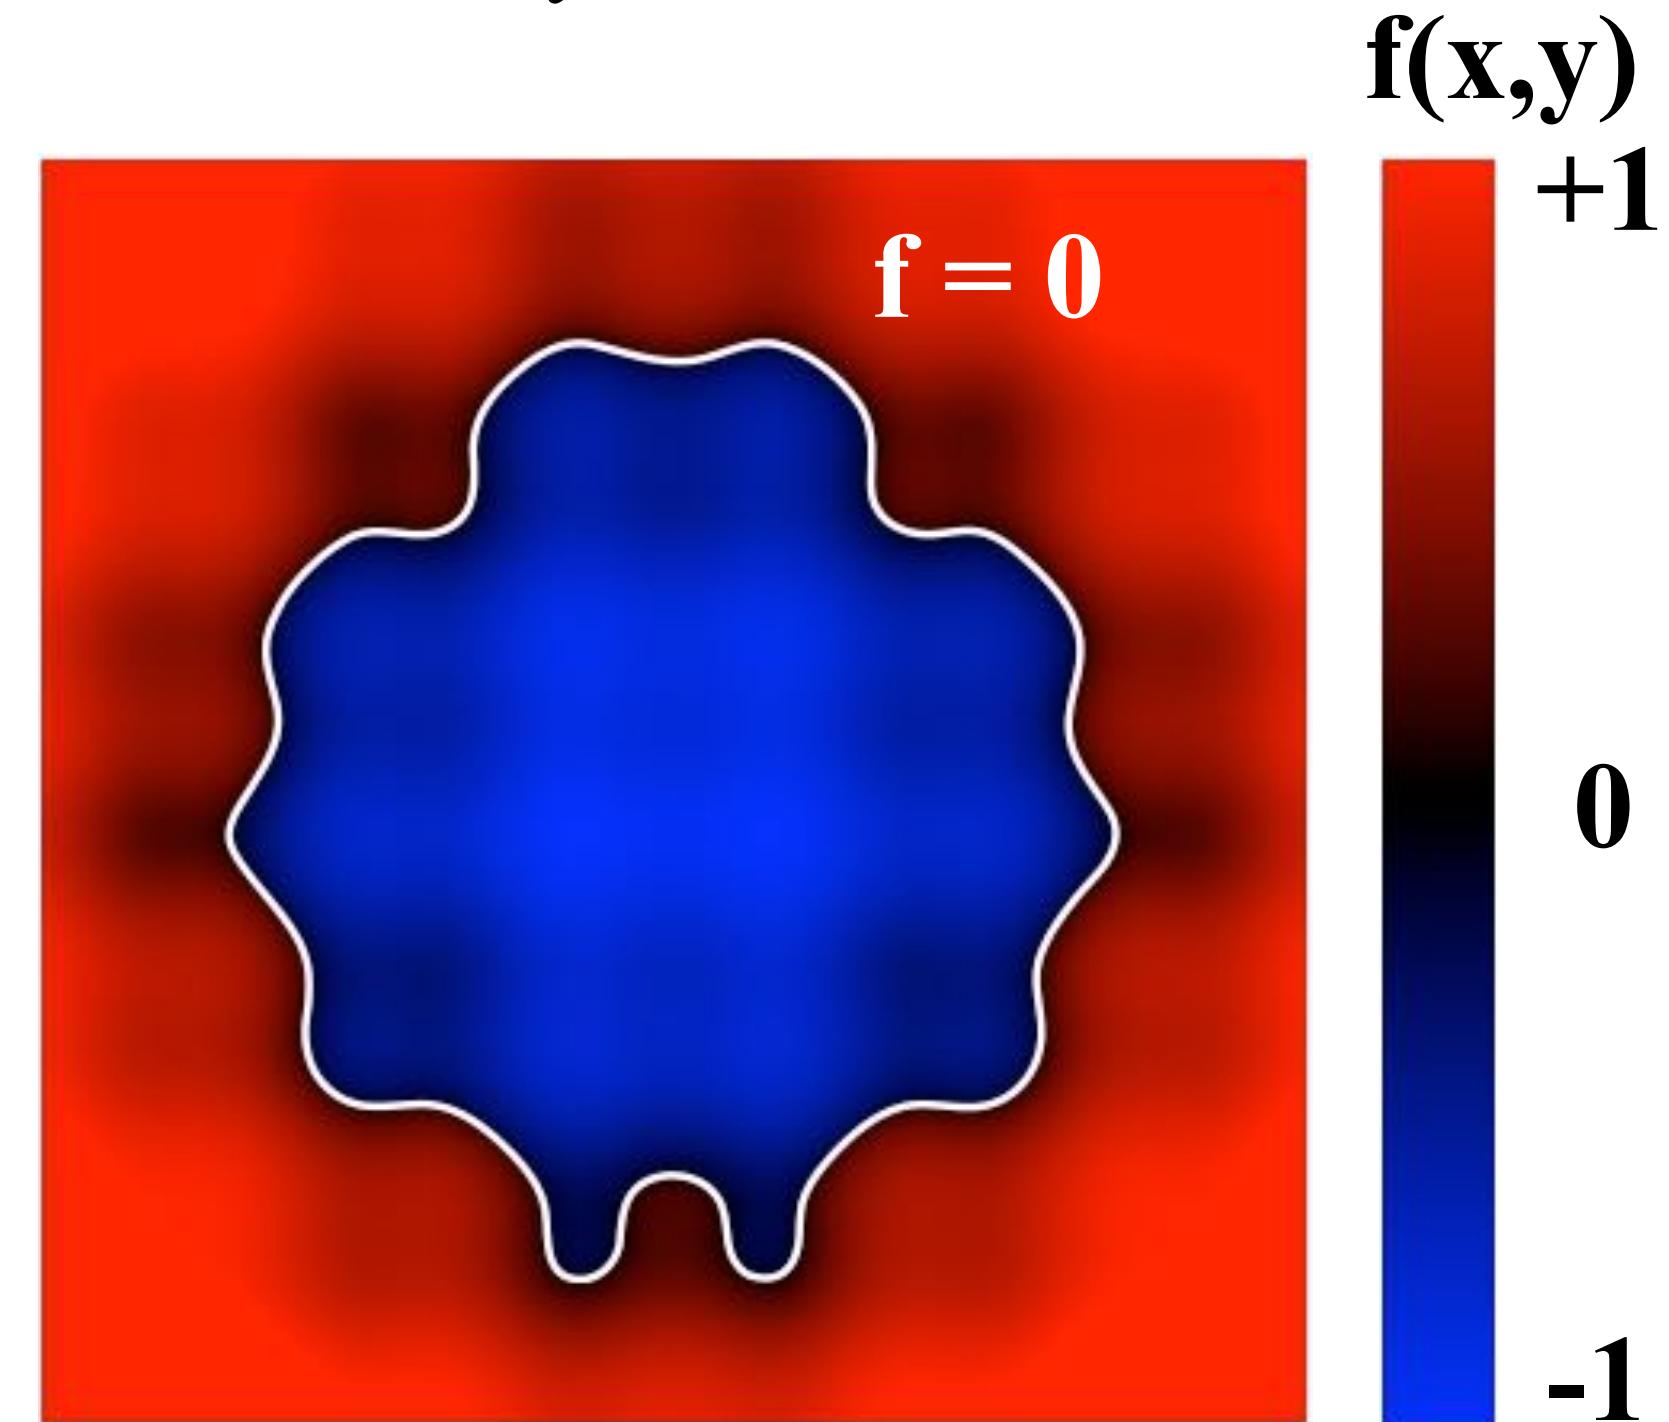
Implicit Representations of Geometry

Based on classifying points

- Points satisfy some specified relationship

E.g. sphere: all points in 3D, where $x^2 + y^2 + z^2 = 1$

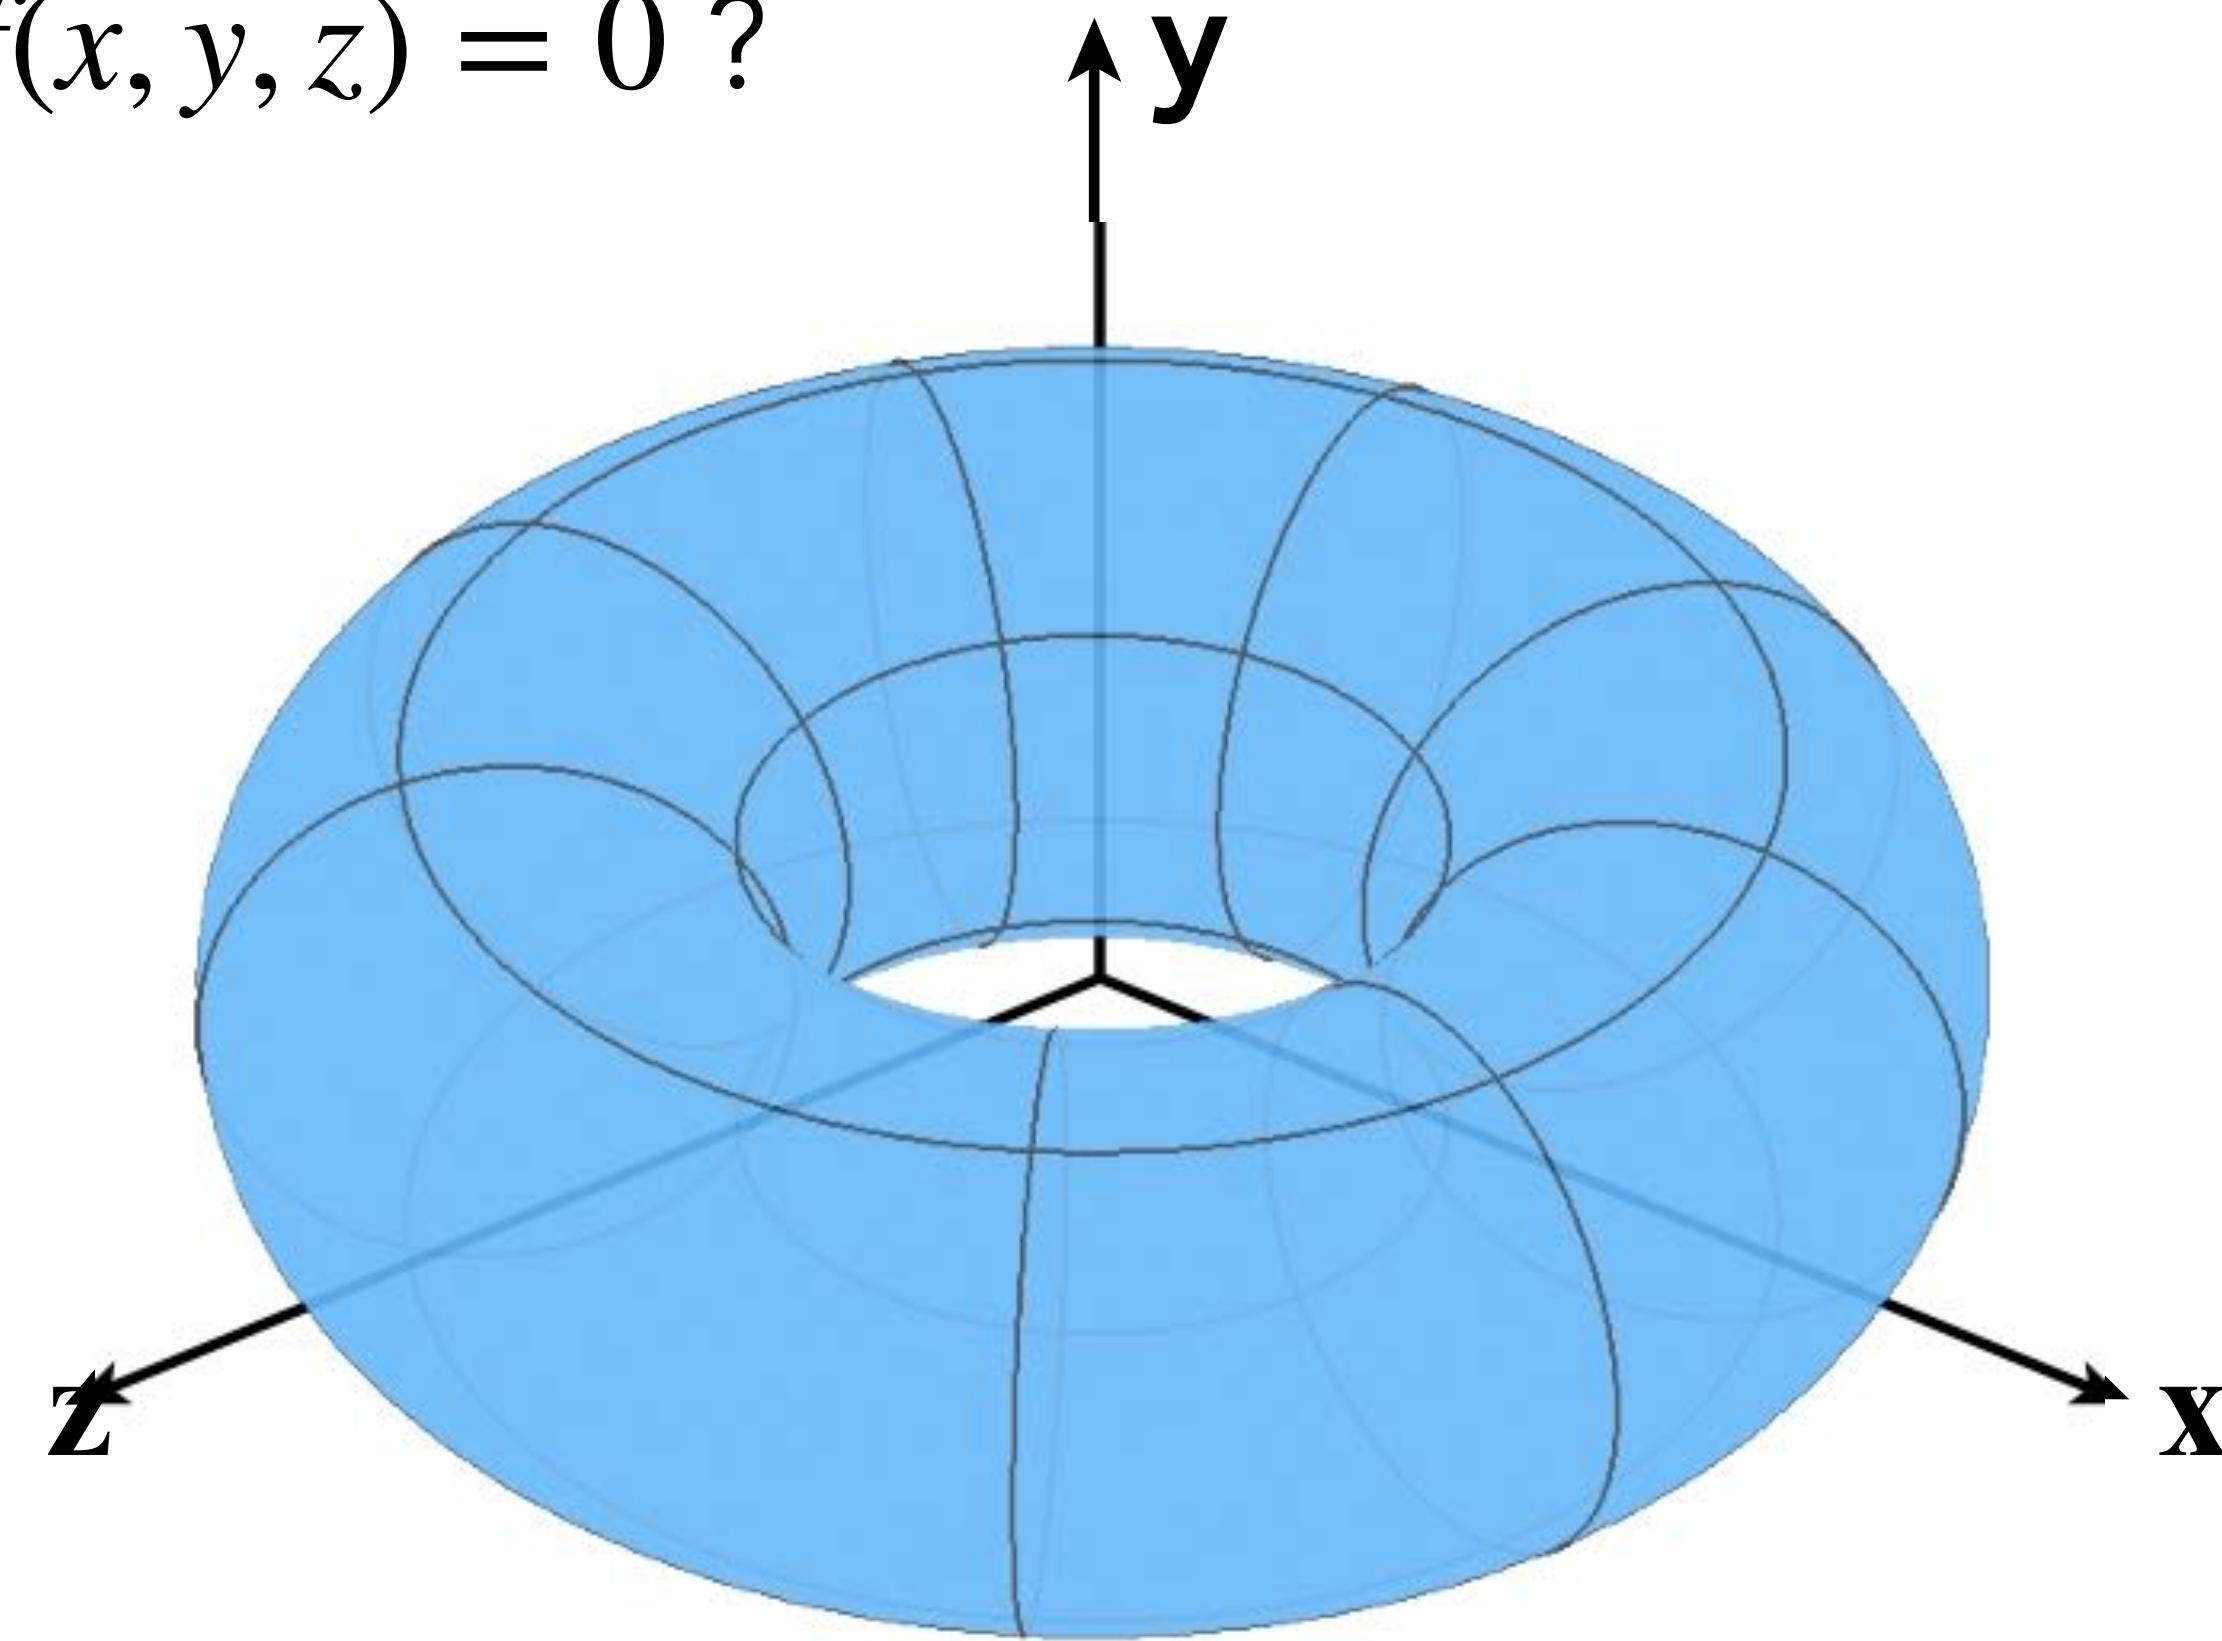
More generally, $f(x, y, z) = 0$



Implicit Surface – Sampling Can Be Hard

$$f(x, y, z) = (2 - \sqrt{x^2 + y^2})^2 + z^2 - 1$$

What points lie on $f(x, y, z) = 0$?



Some tasks are hard with implicit representations

Implicit Surface - Inside/Outside Tests Easy

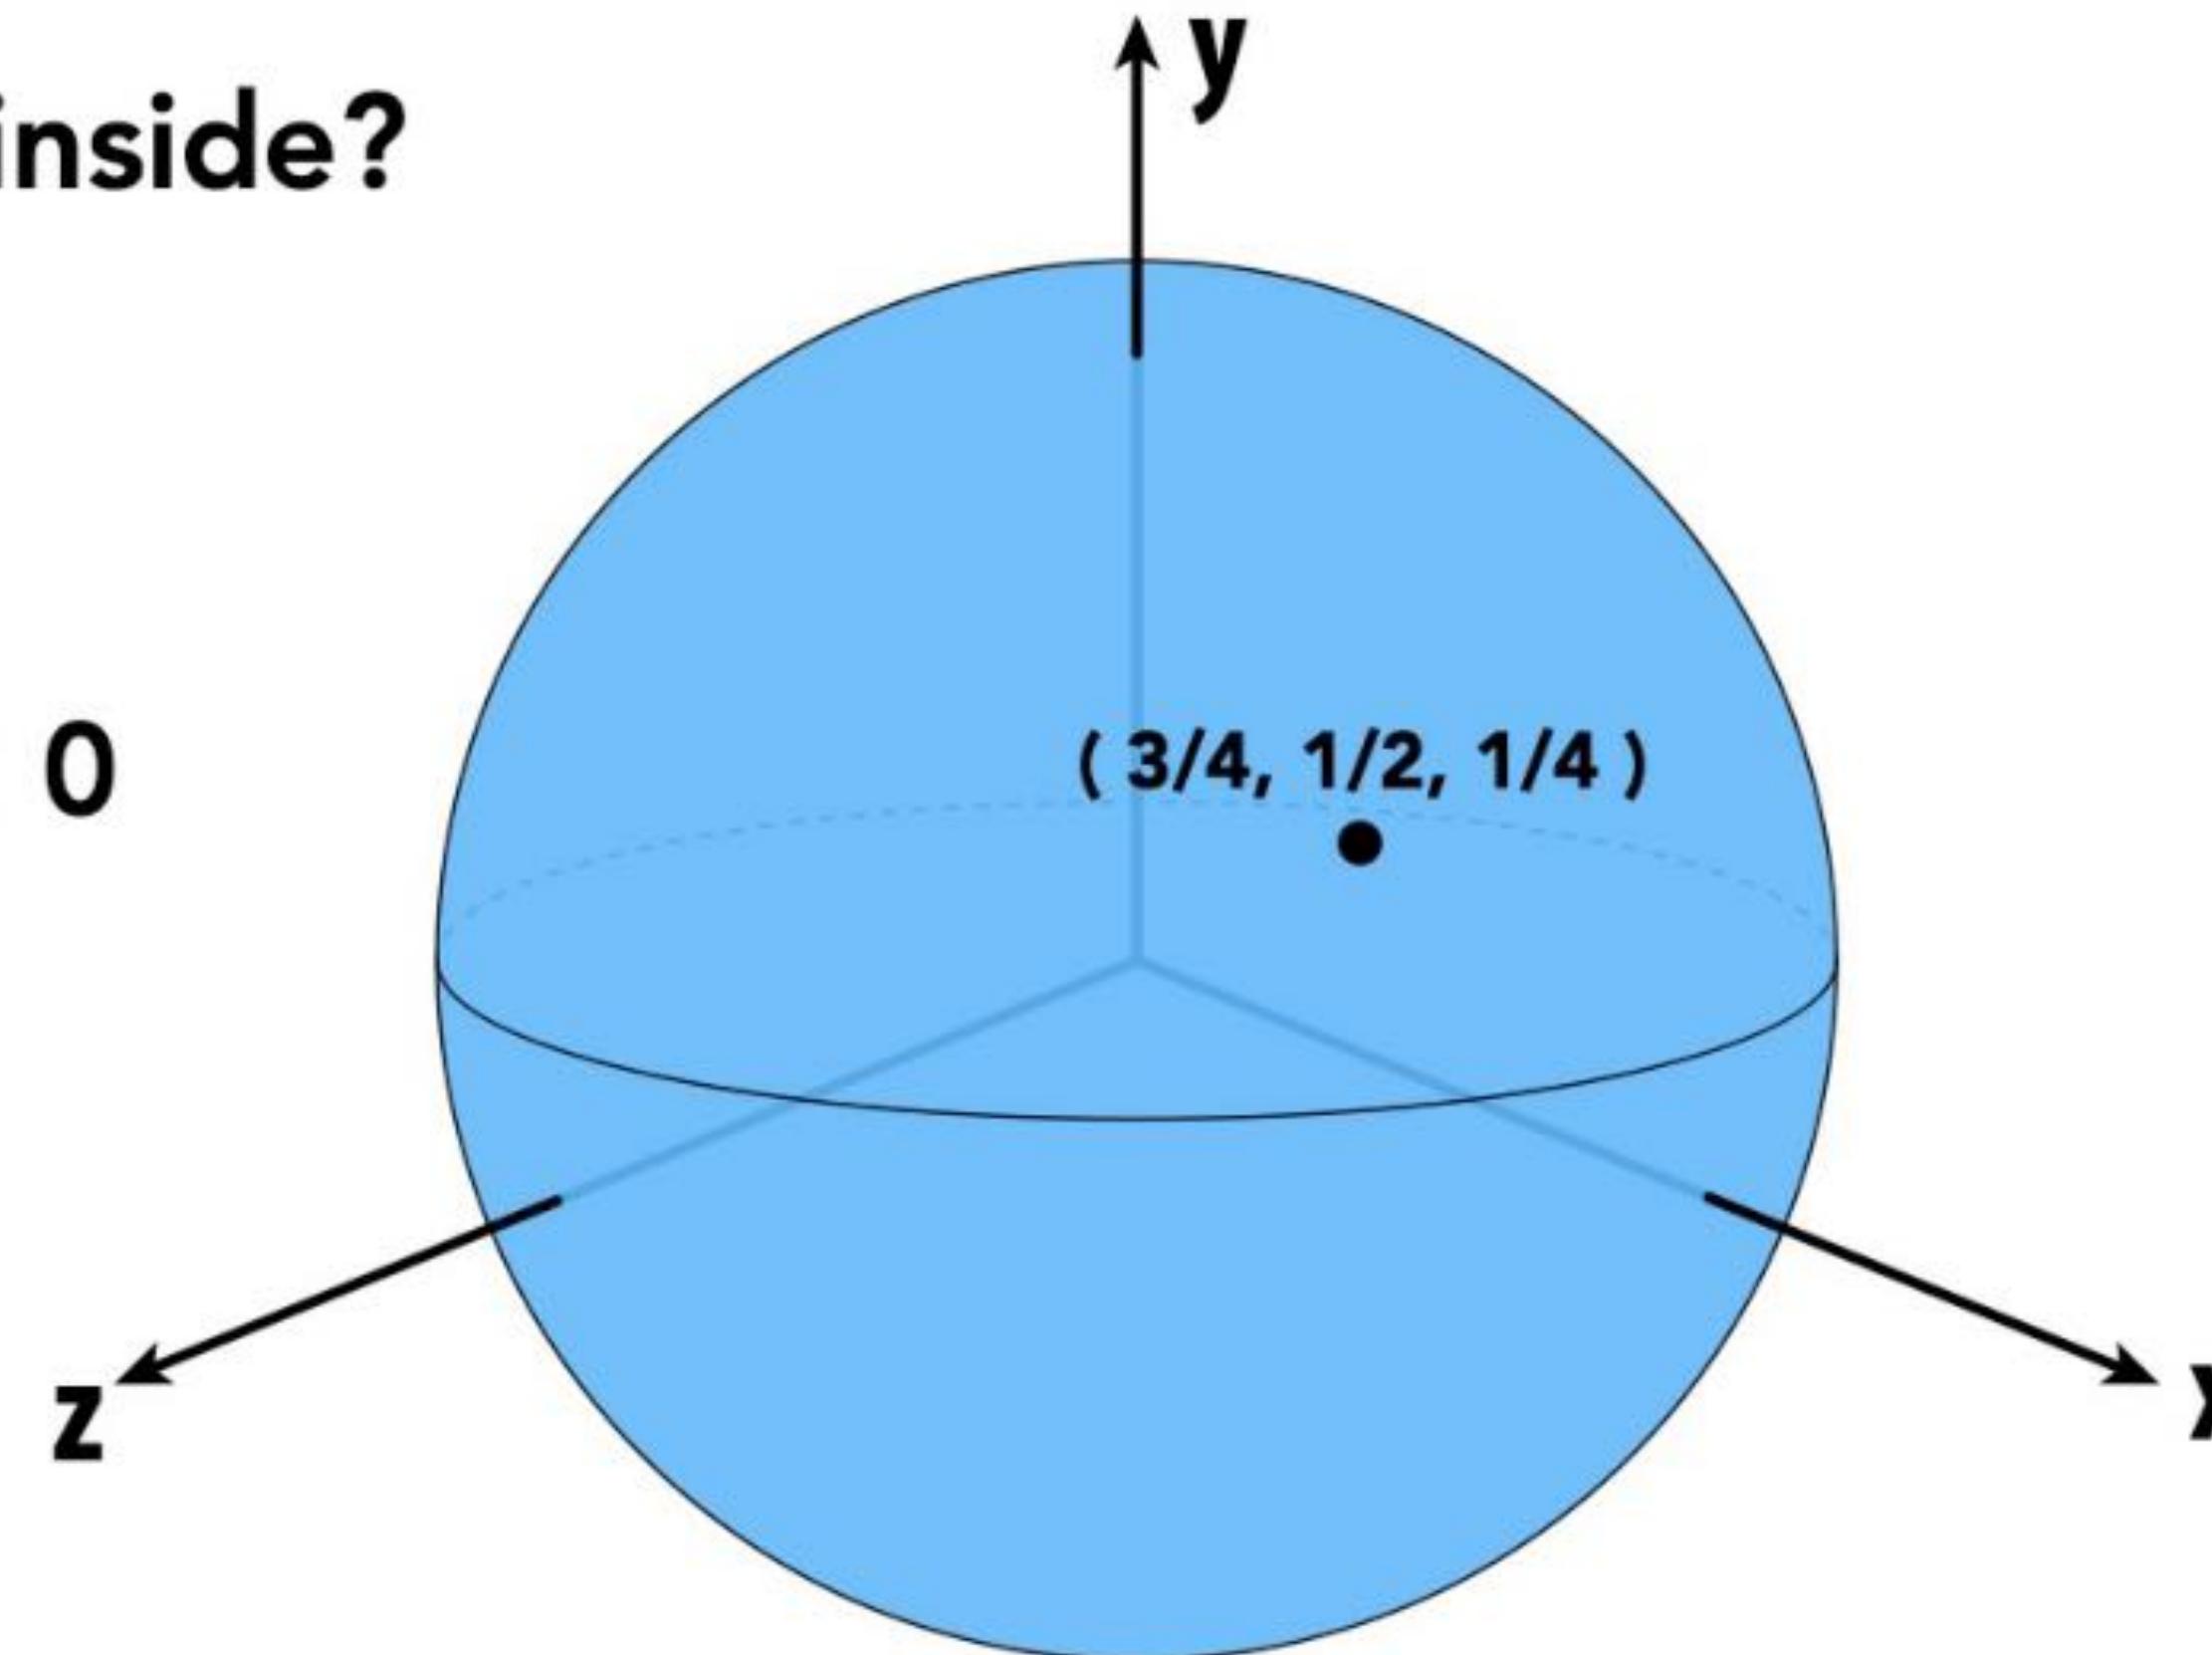
$$f(x, y, z) = x^2 + y^2 + z^2 - 1$$

Is $(3/4, 1/2, 1/4)$ inside?

Just plug it in:

$$f(x, y, z) = -1/8 < 0$$

Yes, inside.



Implicit representations make some tasks easy

Algebraic Surfaces (Implicit)

Surface is zero set of a polynomial in x, y, z



$$x^2 + y^2 + z^2 = 1$$



$$(R - \sqrt{x^2 + y^2})^2 + z^2 = r^2$$



$$(x^2 + \frac{9y^2}{4} + z^2 - 1)^3 =$$

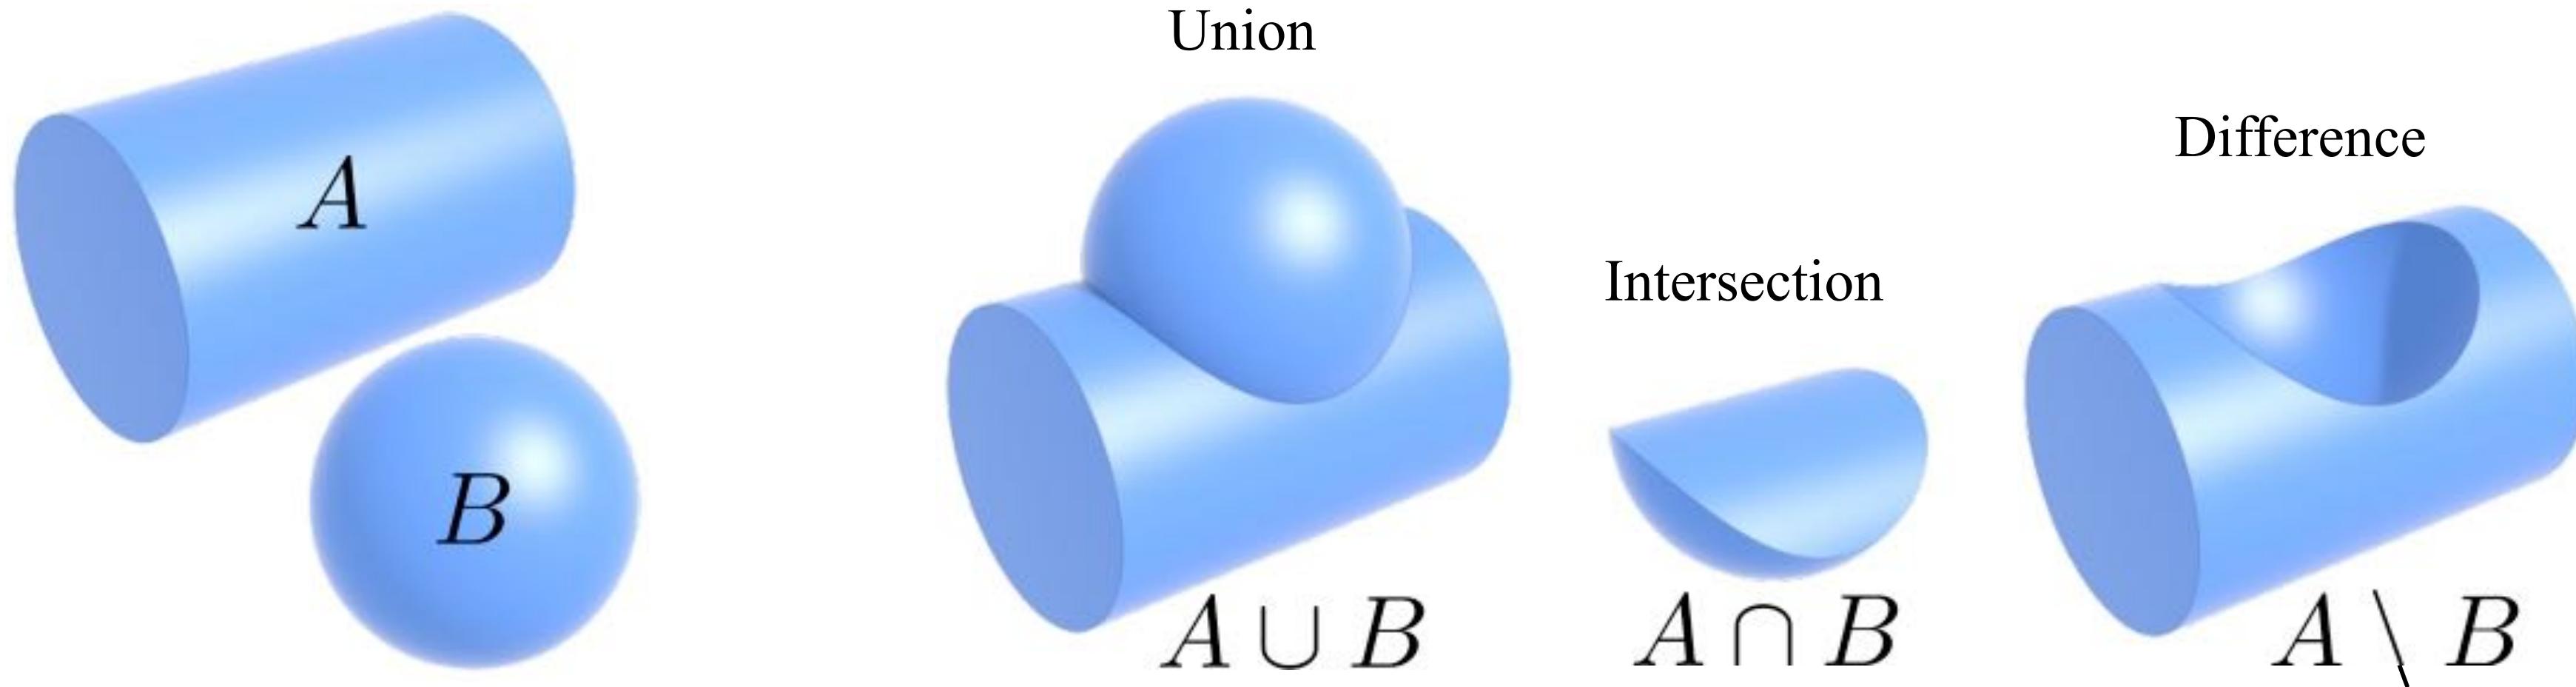
$$x^2 z^3 + \frac{9y^2 z^3}{80}$$



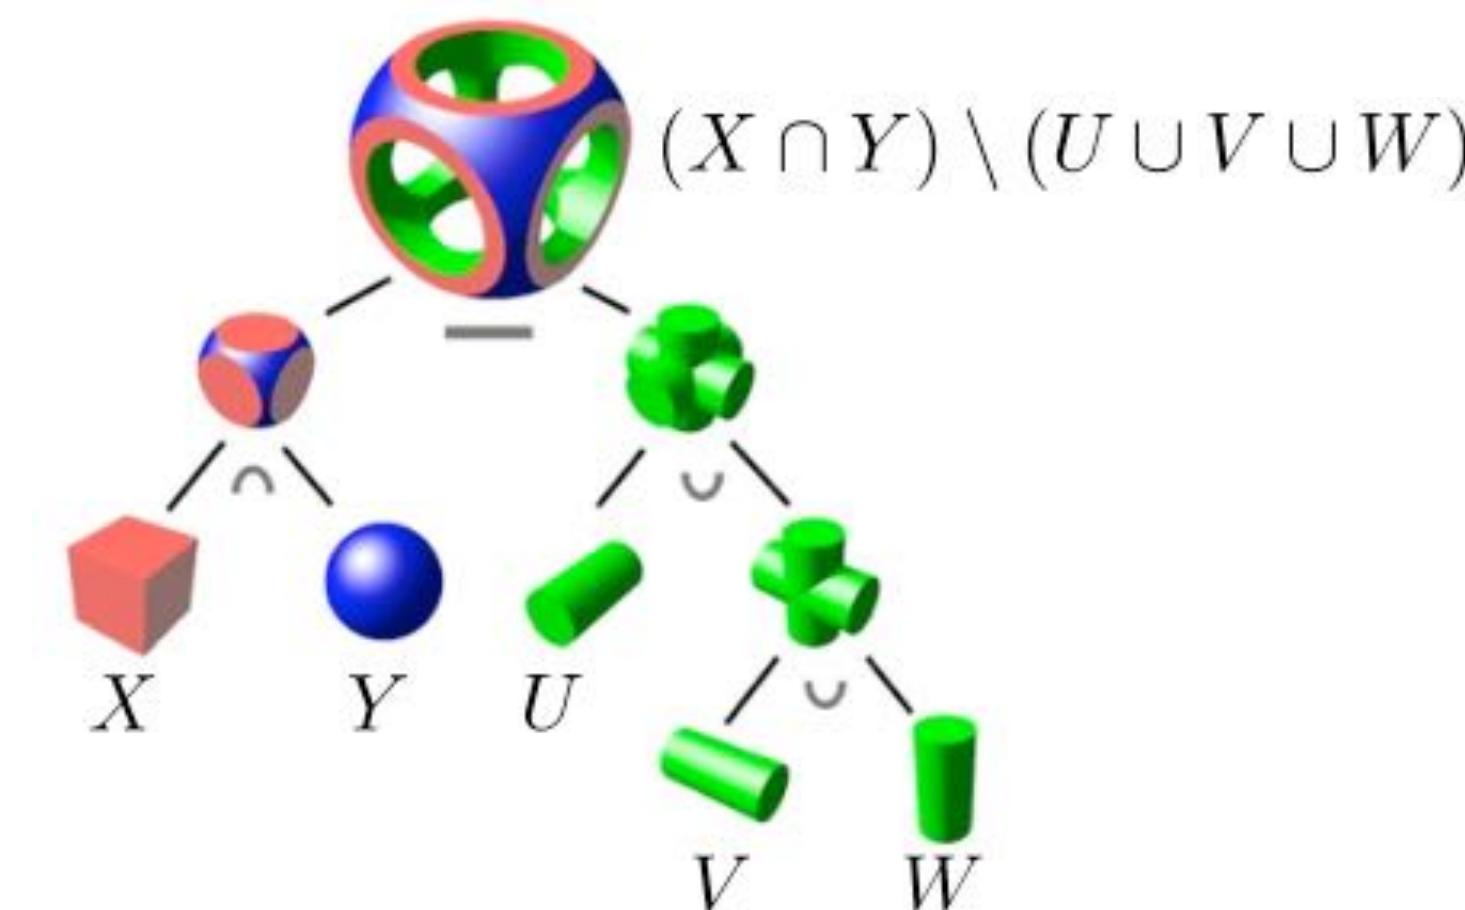
More complex shapes?

Constructive Solid Geometry (Implicit)

Combine implicit geometry via Boolean operations

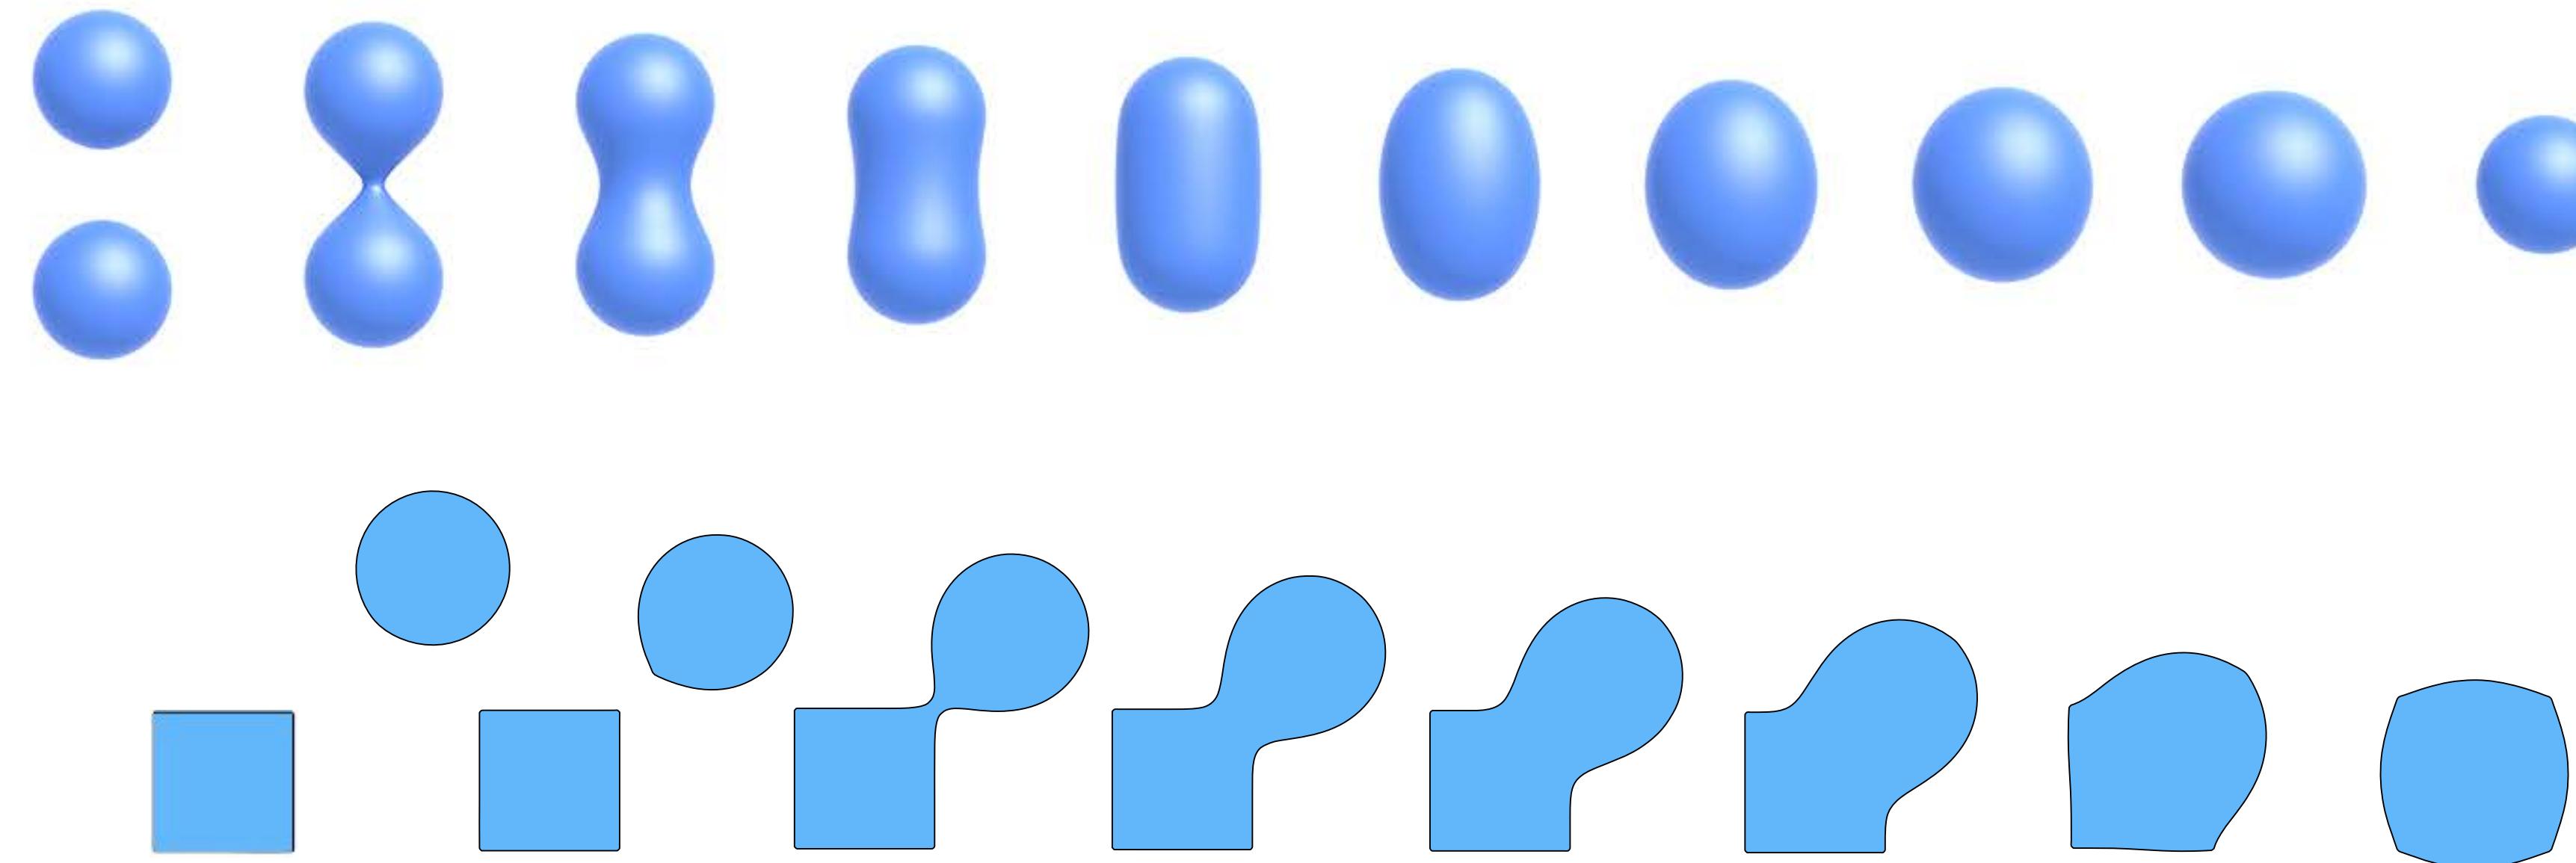


Boolean expressions :



Distance Functions

- Distance functions: giving minimum distance (could be signed distance) from anywhere to object
- Instead of booleans, gradually blend surfaces together using distance functions



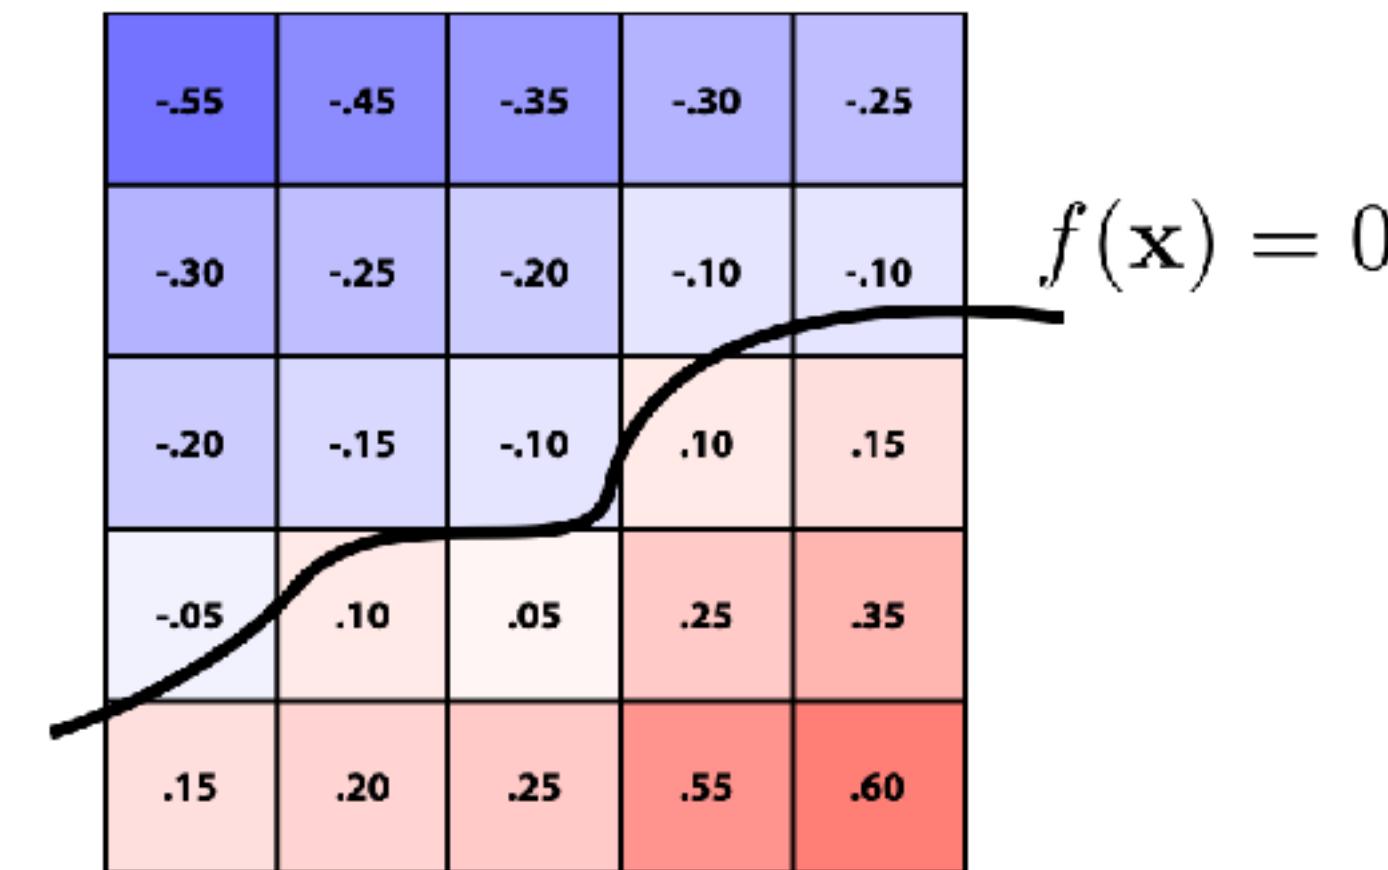
Scene of Pure Distance Functions (Not Easy!)



Level Set Methods

Closed-form equations are hard to describe complex shapes

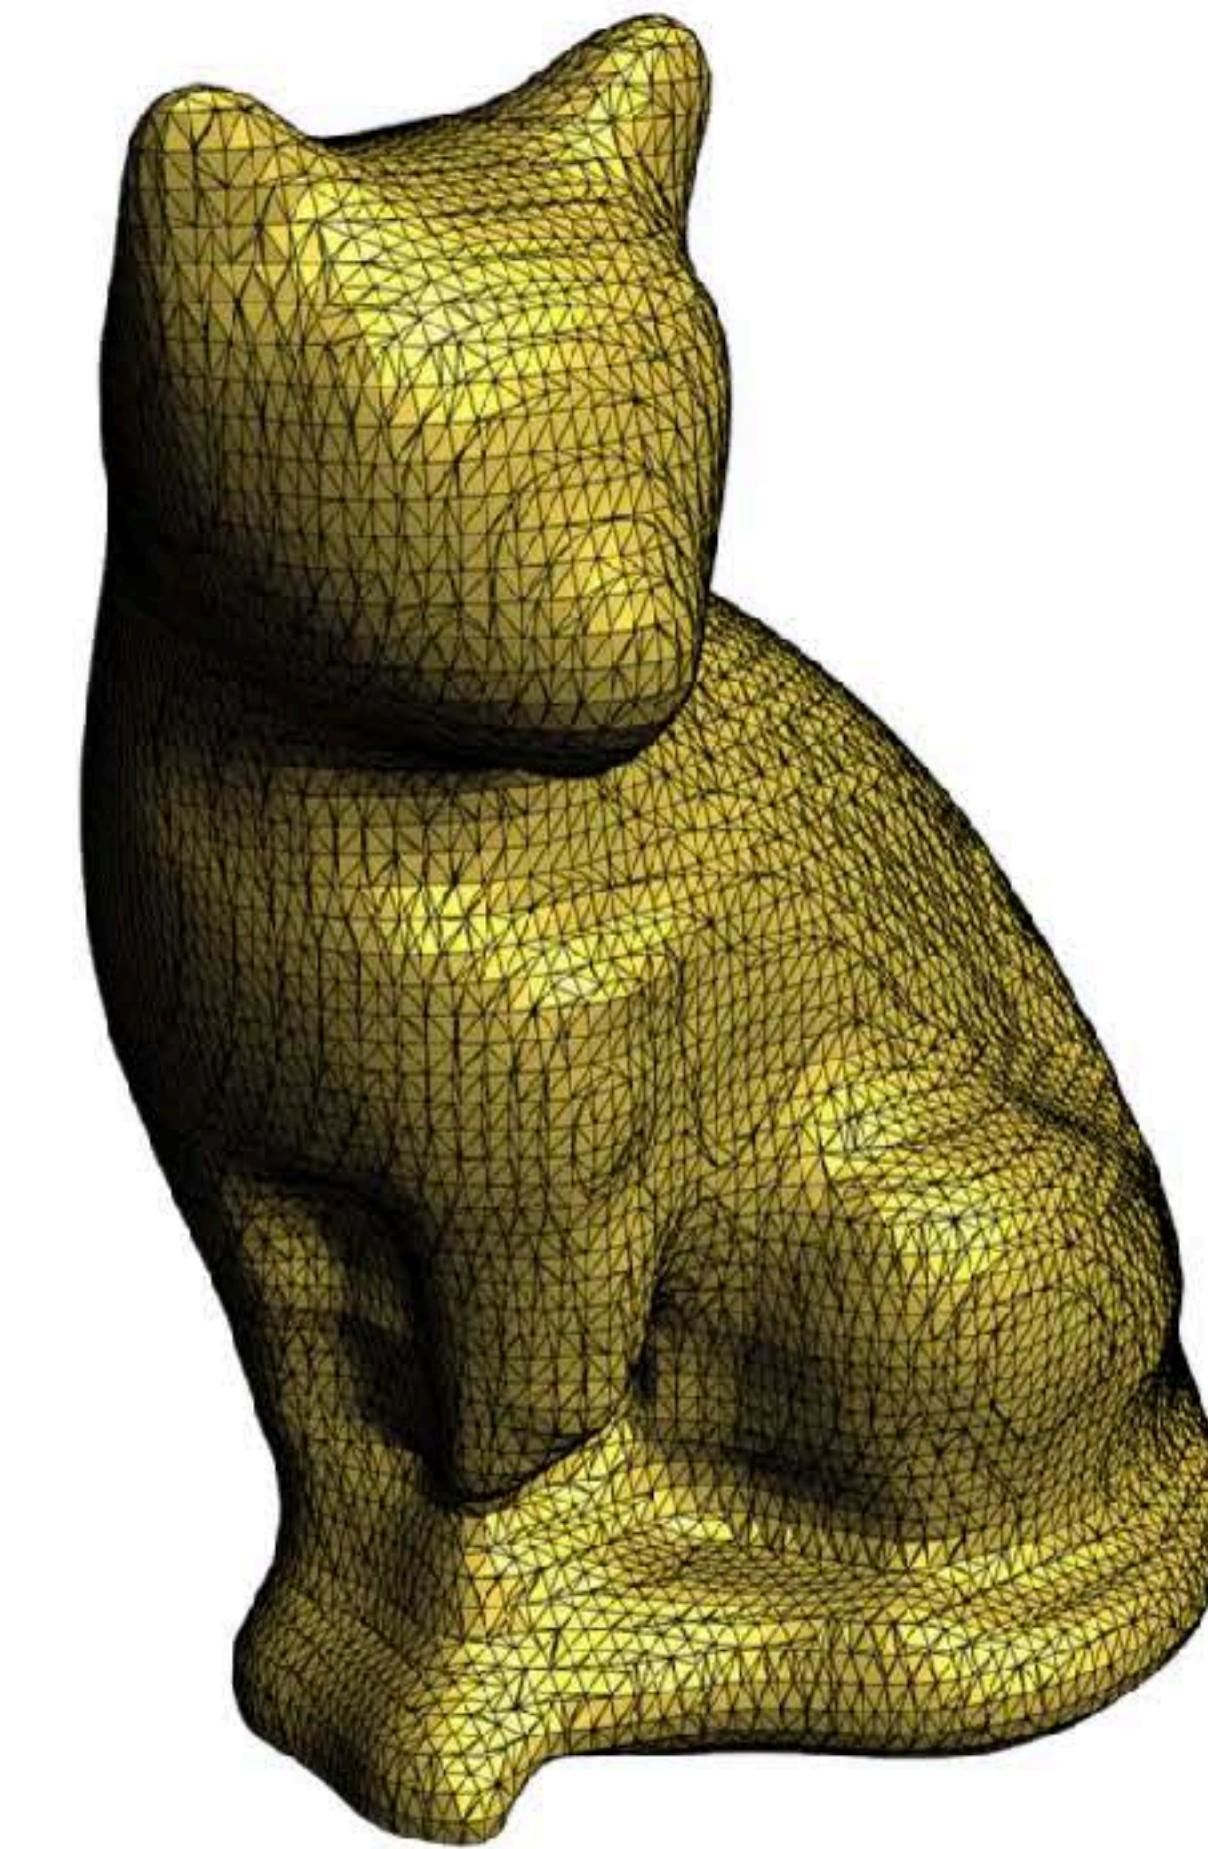
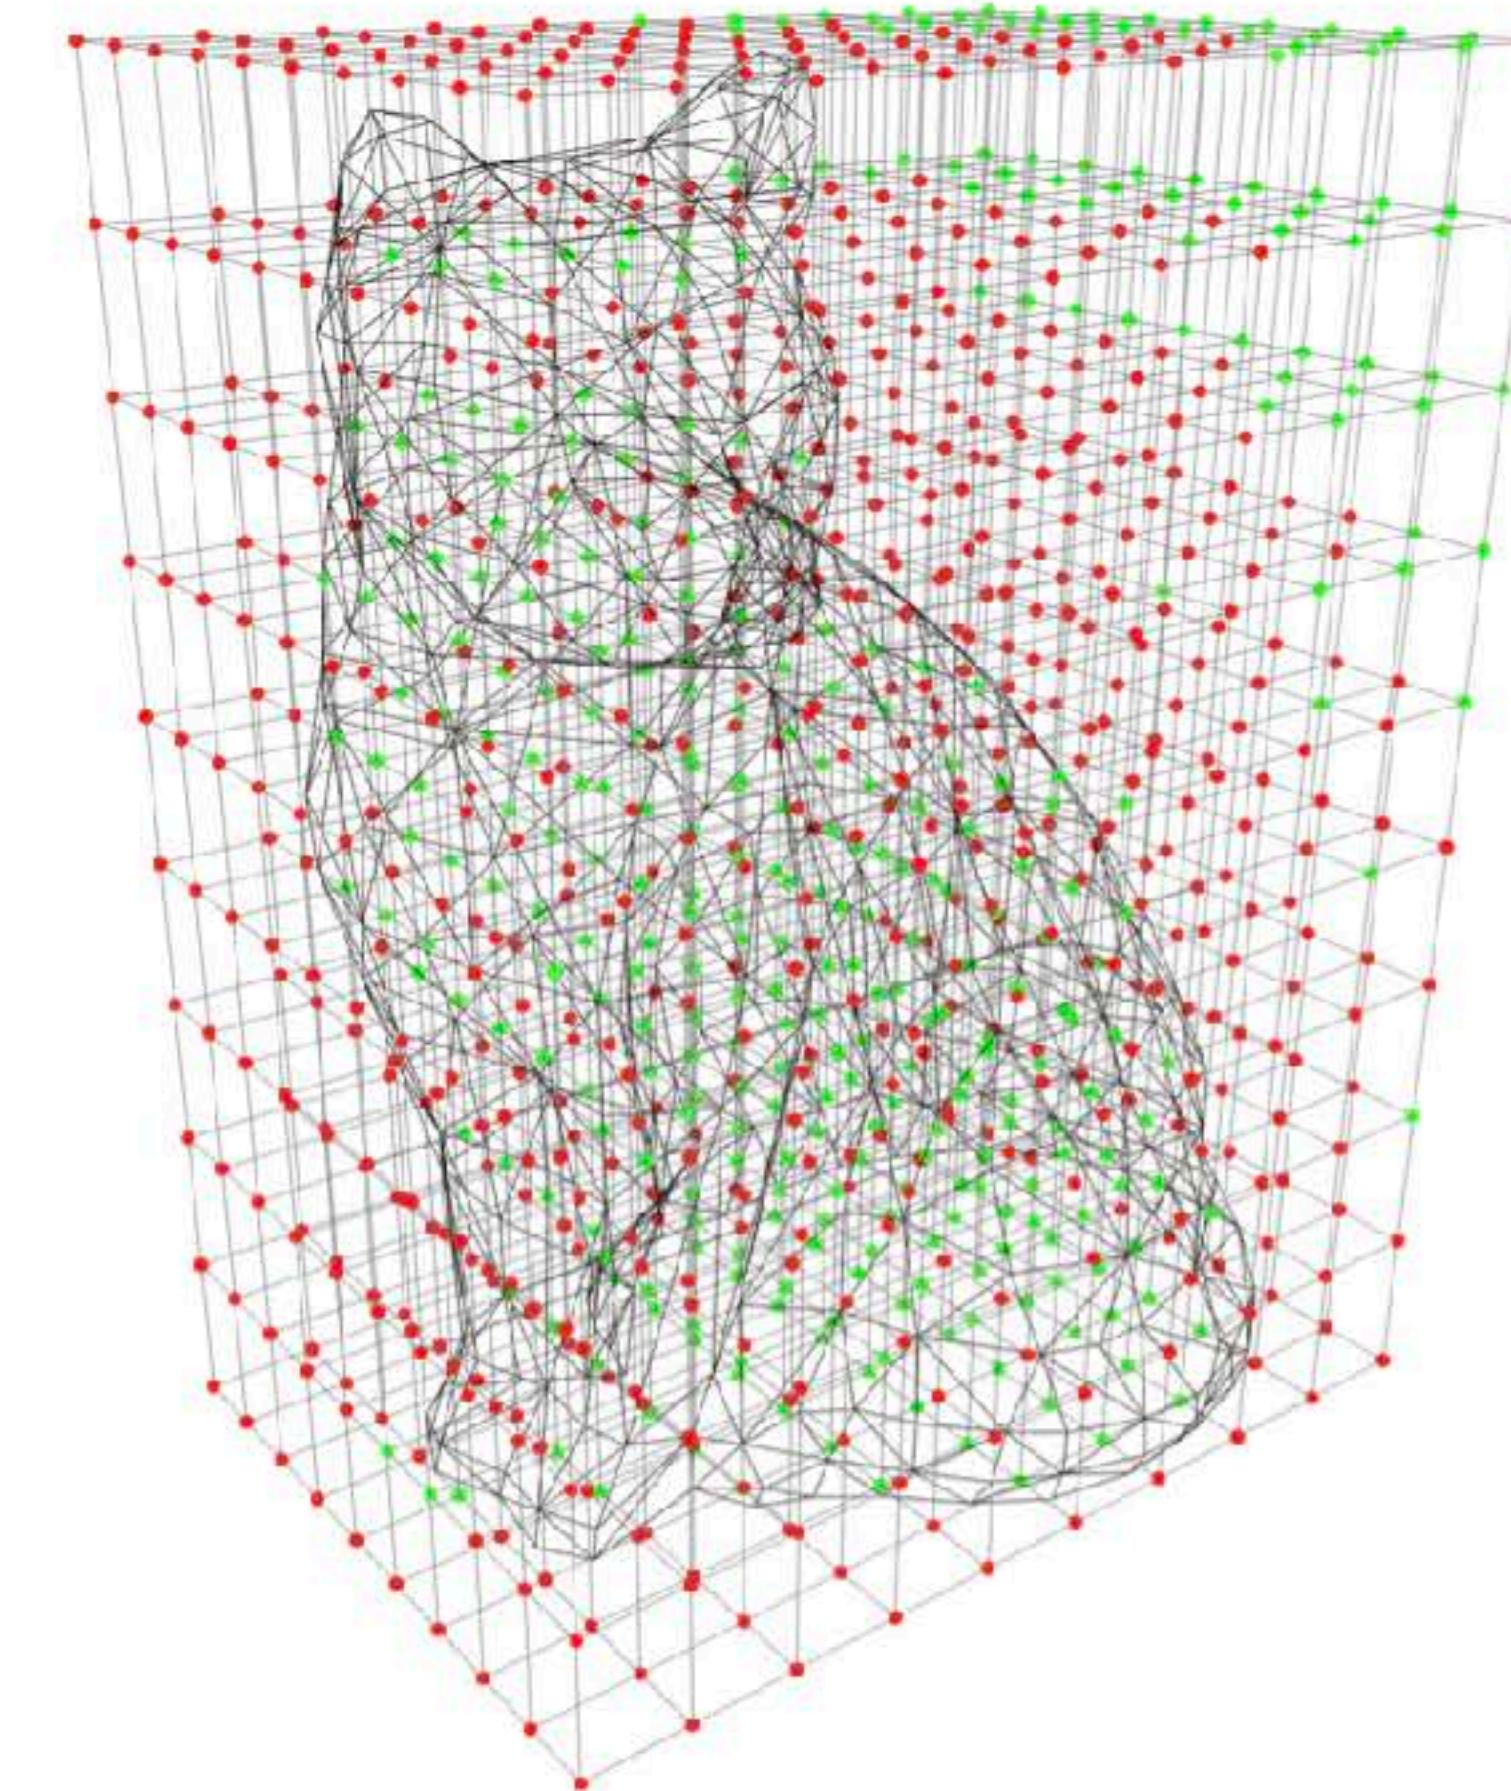
Alternative: store a grid of values approximating function



Surface is found where interpolated values equal zero

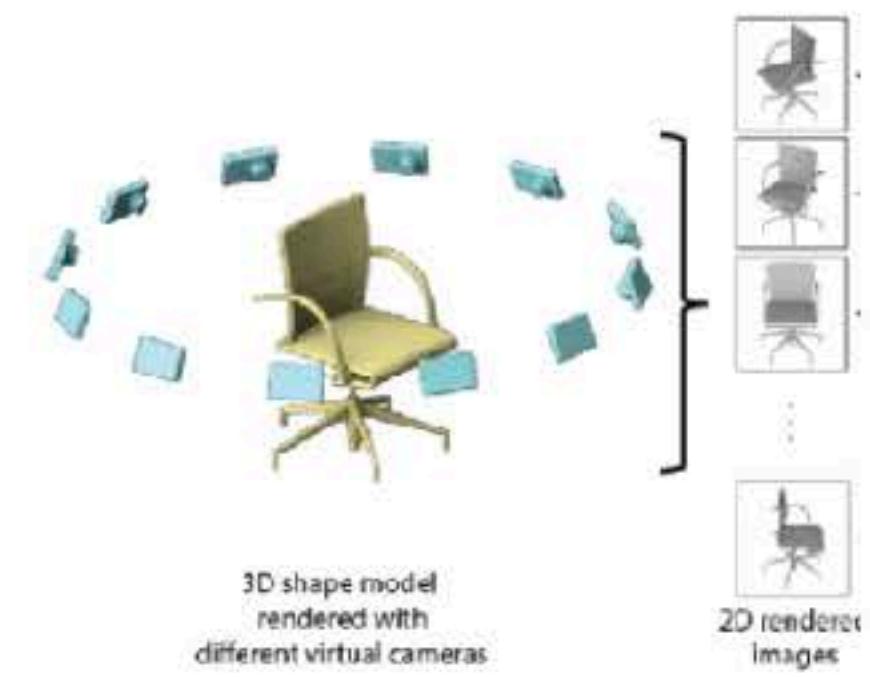
Provides much more explicit control over shape (like a texture)

Conversion between different representations

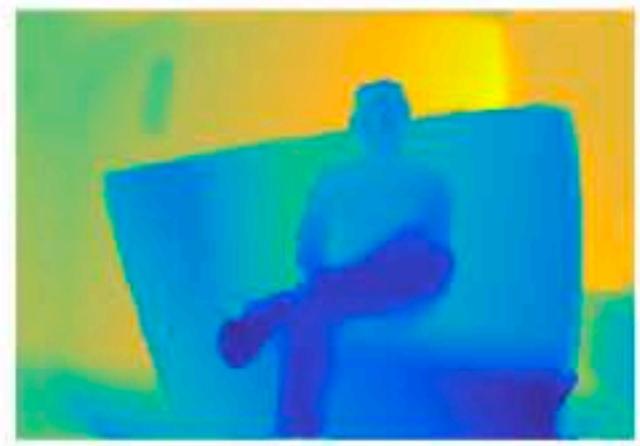


Summary

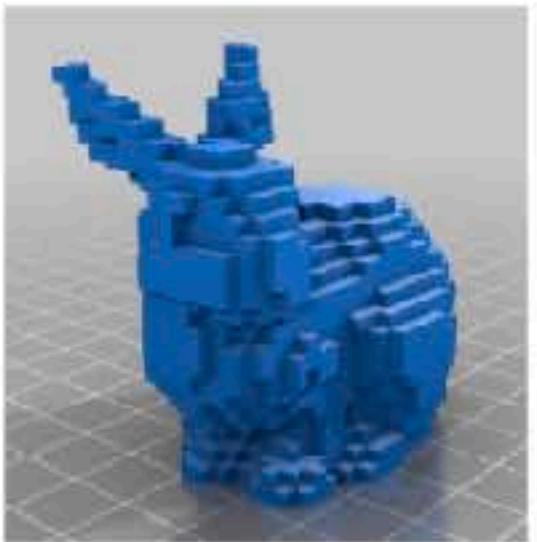
Rasterized form (regular grids)



Multi-view

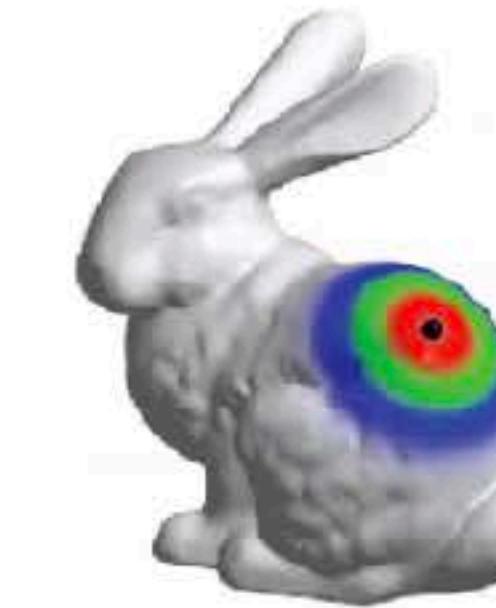


Depth Map

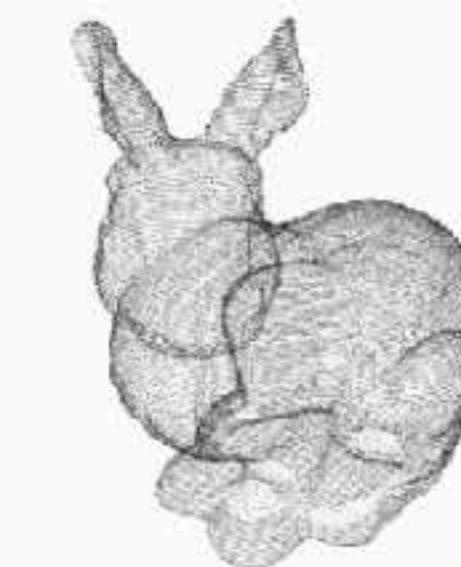


Volumetric

Geometric form (irregular)



Mesh



Point Cloud

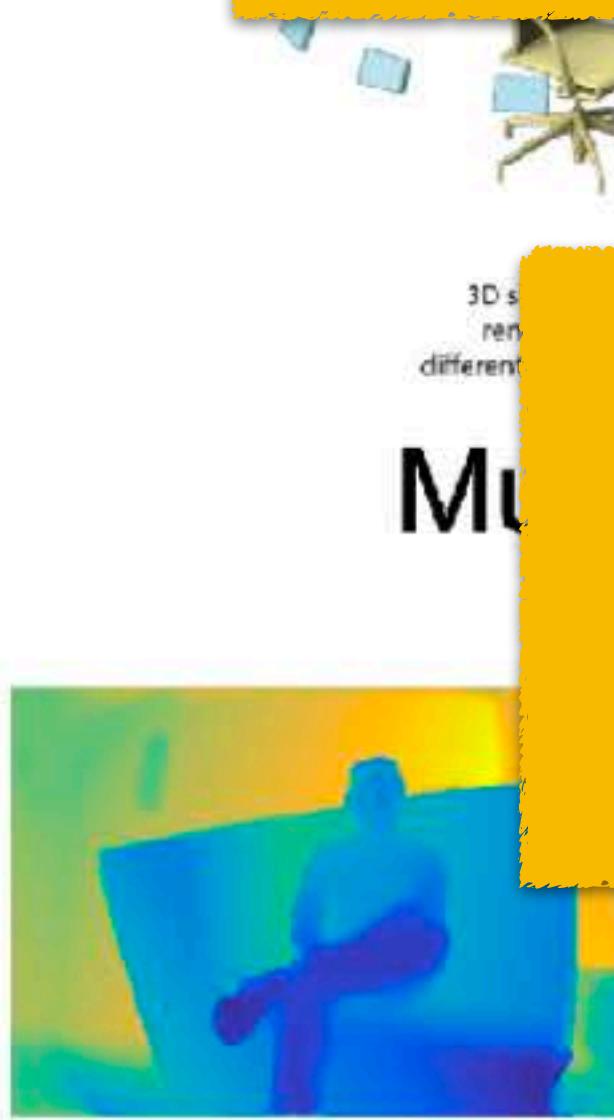
$$F(x) = 0$$

Implicit Shape

Summary

**Rasterized form
(regular grids)**

No “best” geometric representation



Depth Map



Volumetric

**Geometric form
(irregular)**



Point Cloud

$$f(x) = 0$$

Implicit Shape

Next time: 3D transformations

