



Lecture 3: 3D Transformations

Li Yi

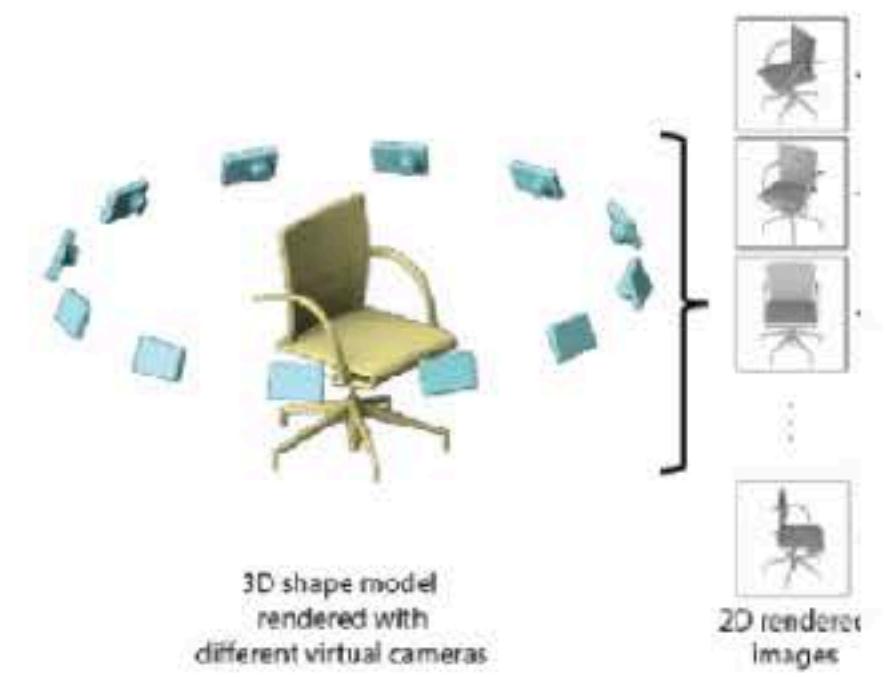
2025.03.06

Announcement

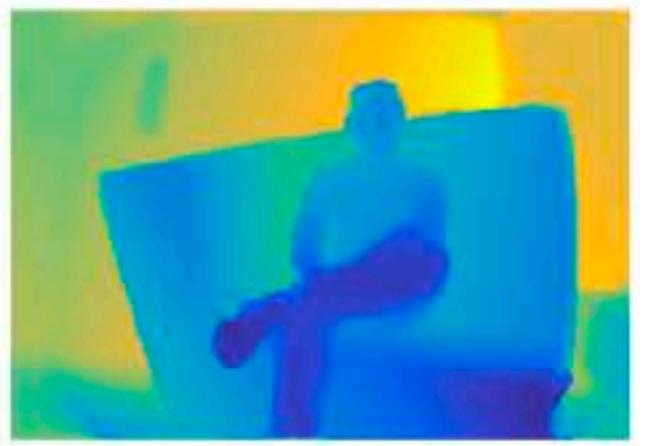
- Homework 1 released!
- Please provide your feedback through anonymous WeChat questionnaire
(<https://www.wjx.cn/vm/rVxWHKI.aspx>)

Recap

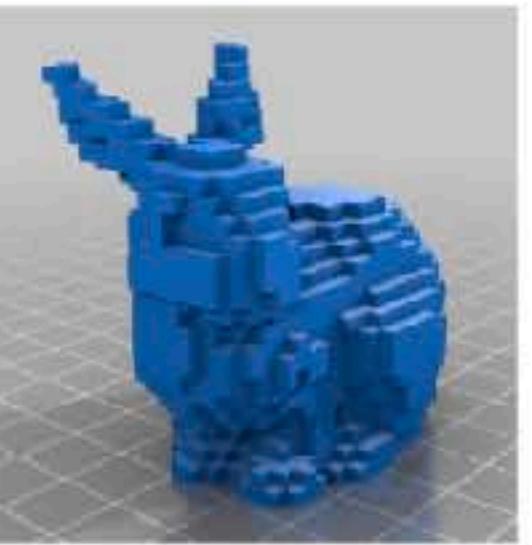
Rasterized form (regular grids)



Multi-view

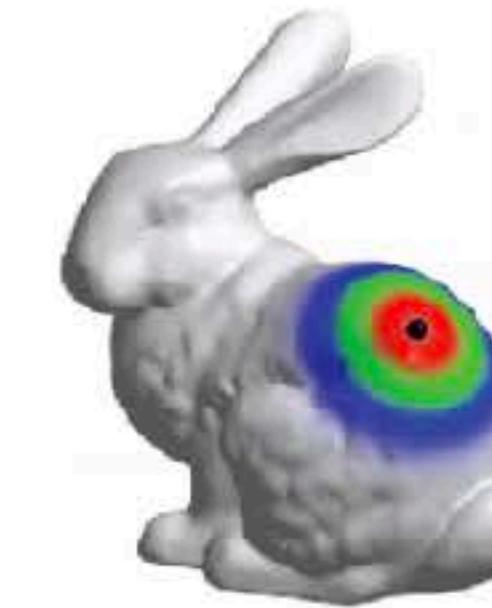


Depth Map

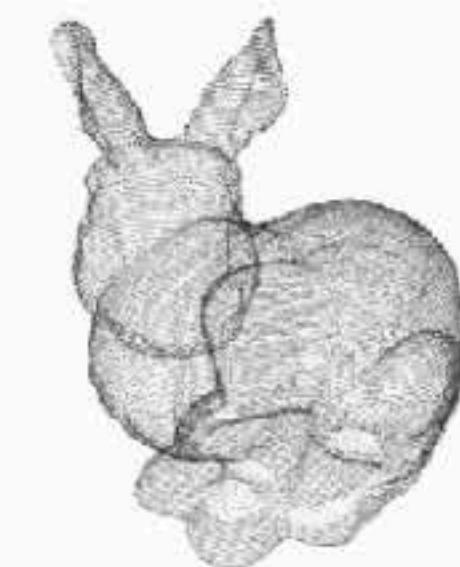


Volumetric

Geometric form (irregular)



Mesh



Point Cloud

$$F(x) = 0$$

Implicit Shape

Outline

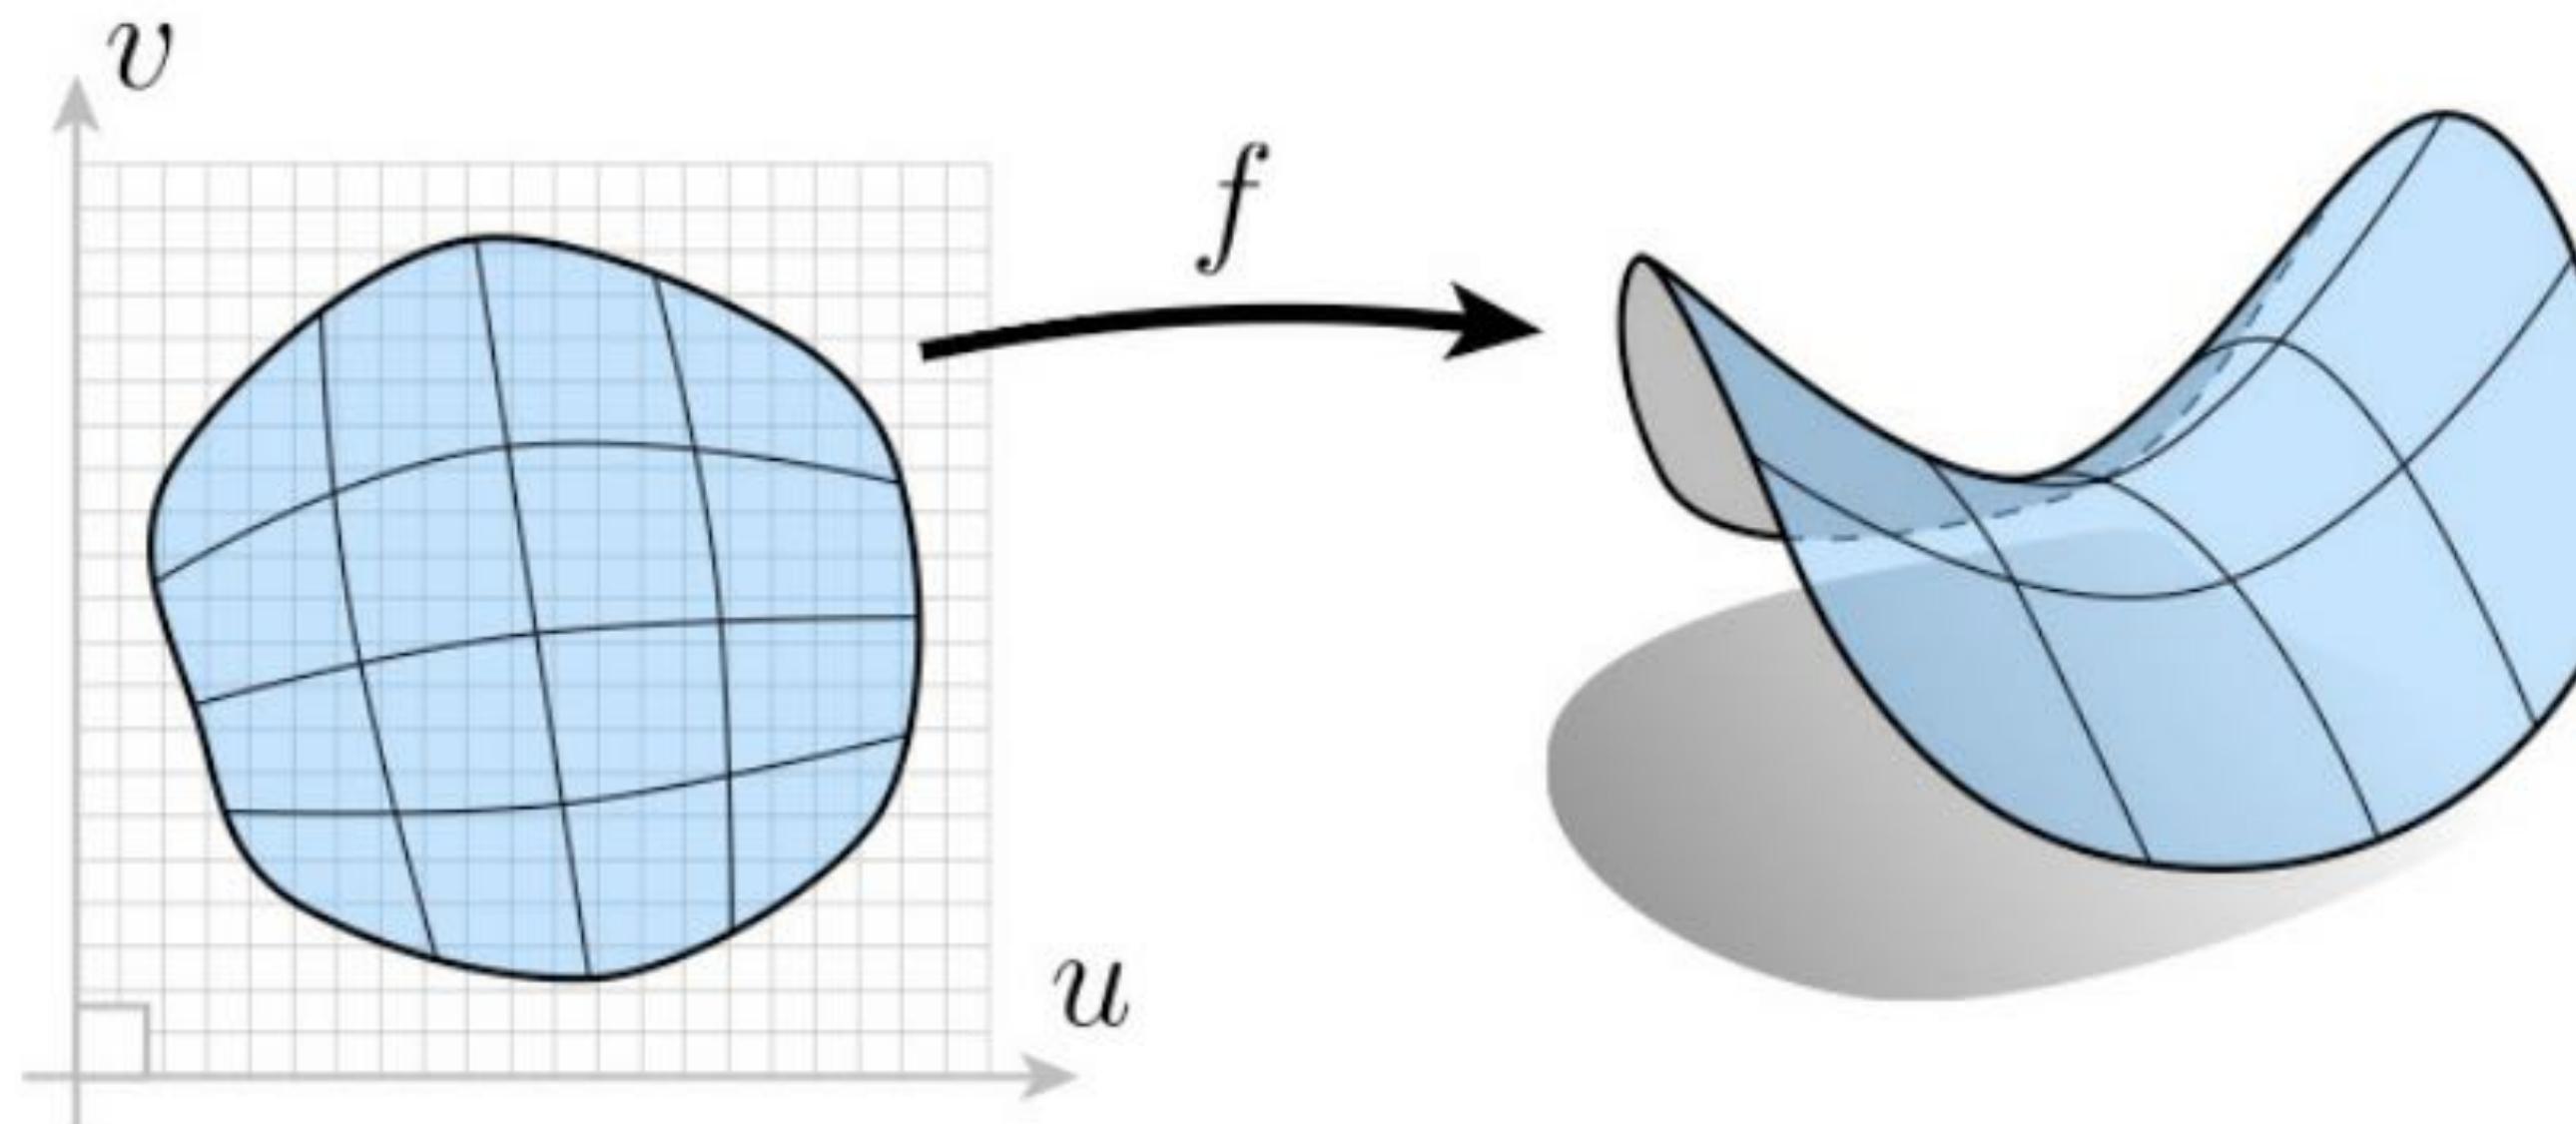
- Rasterized 3D representations
- Mesh
- Point cloud
- *Implicit representations*

Explicit Representation of Geometry

All points are given directly

Generally:

$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^3; (u, v) \mapsto (x, y, z)$$

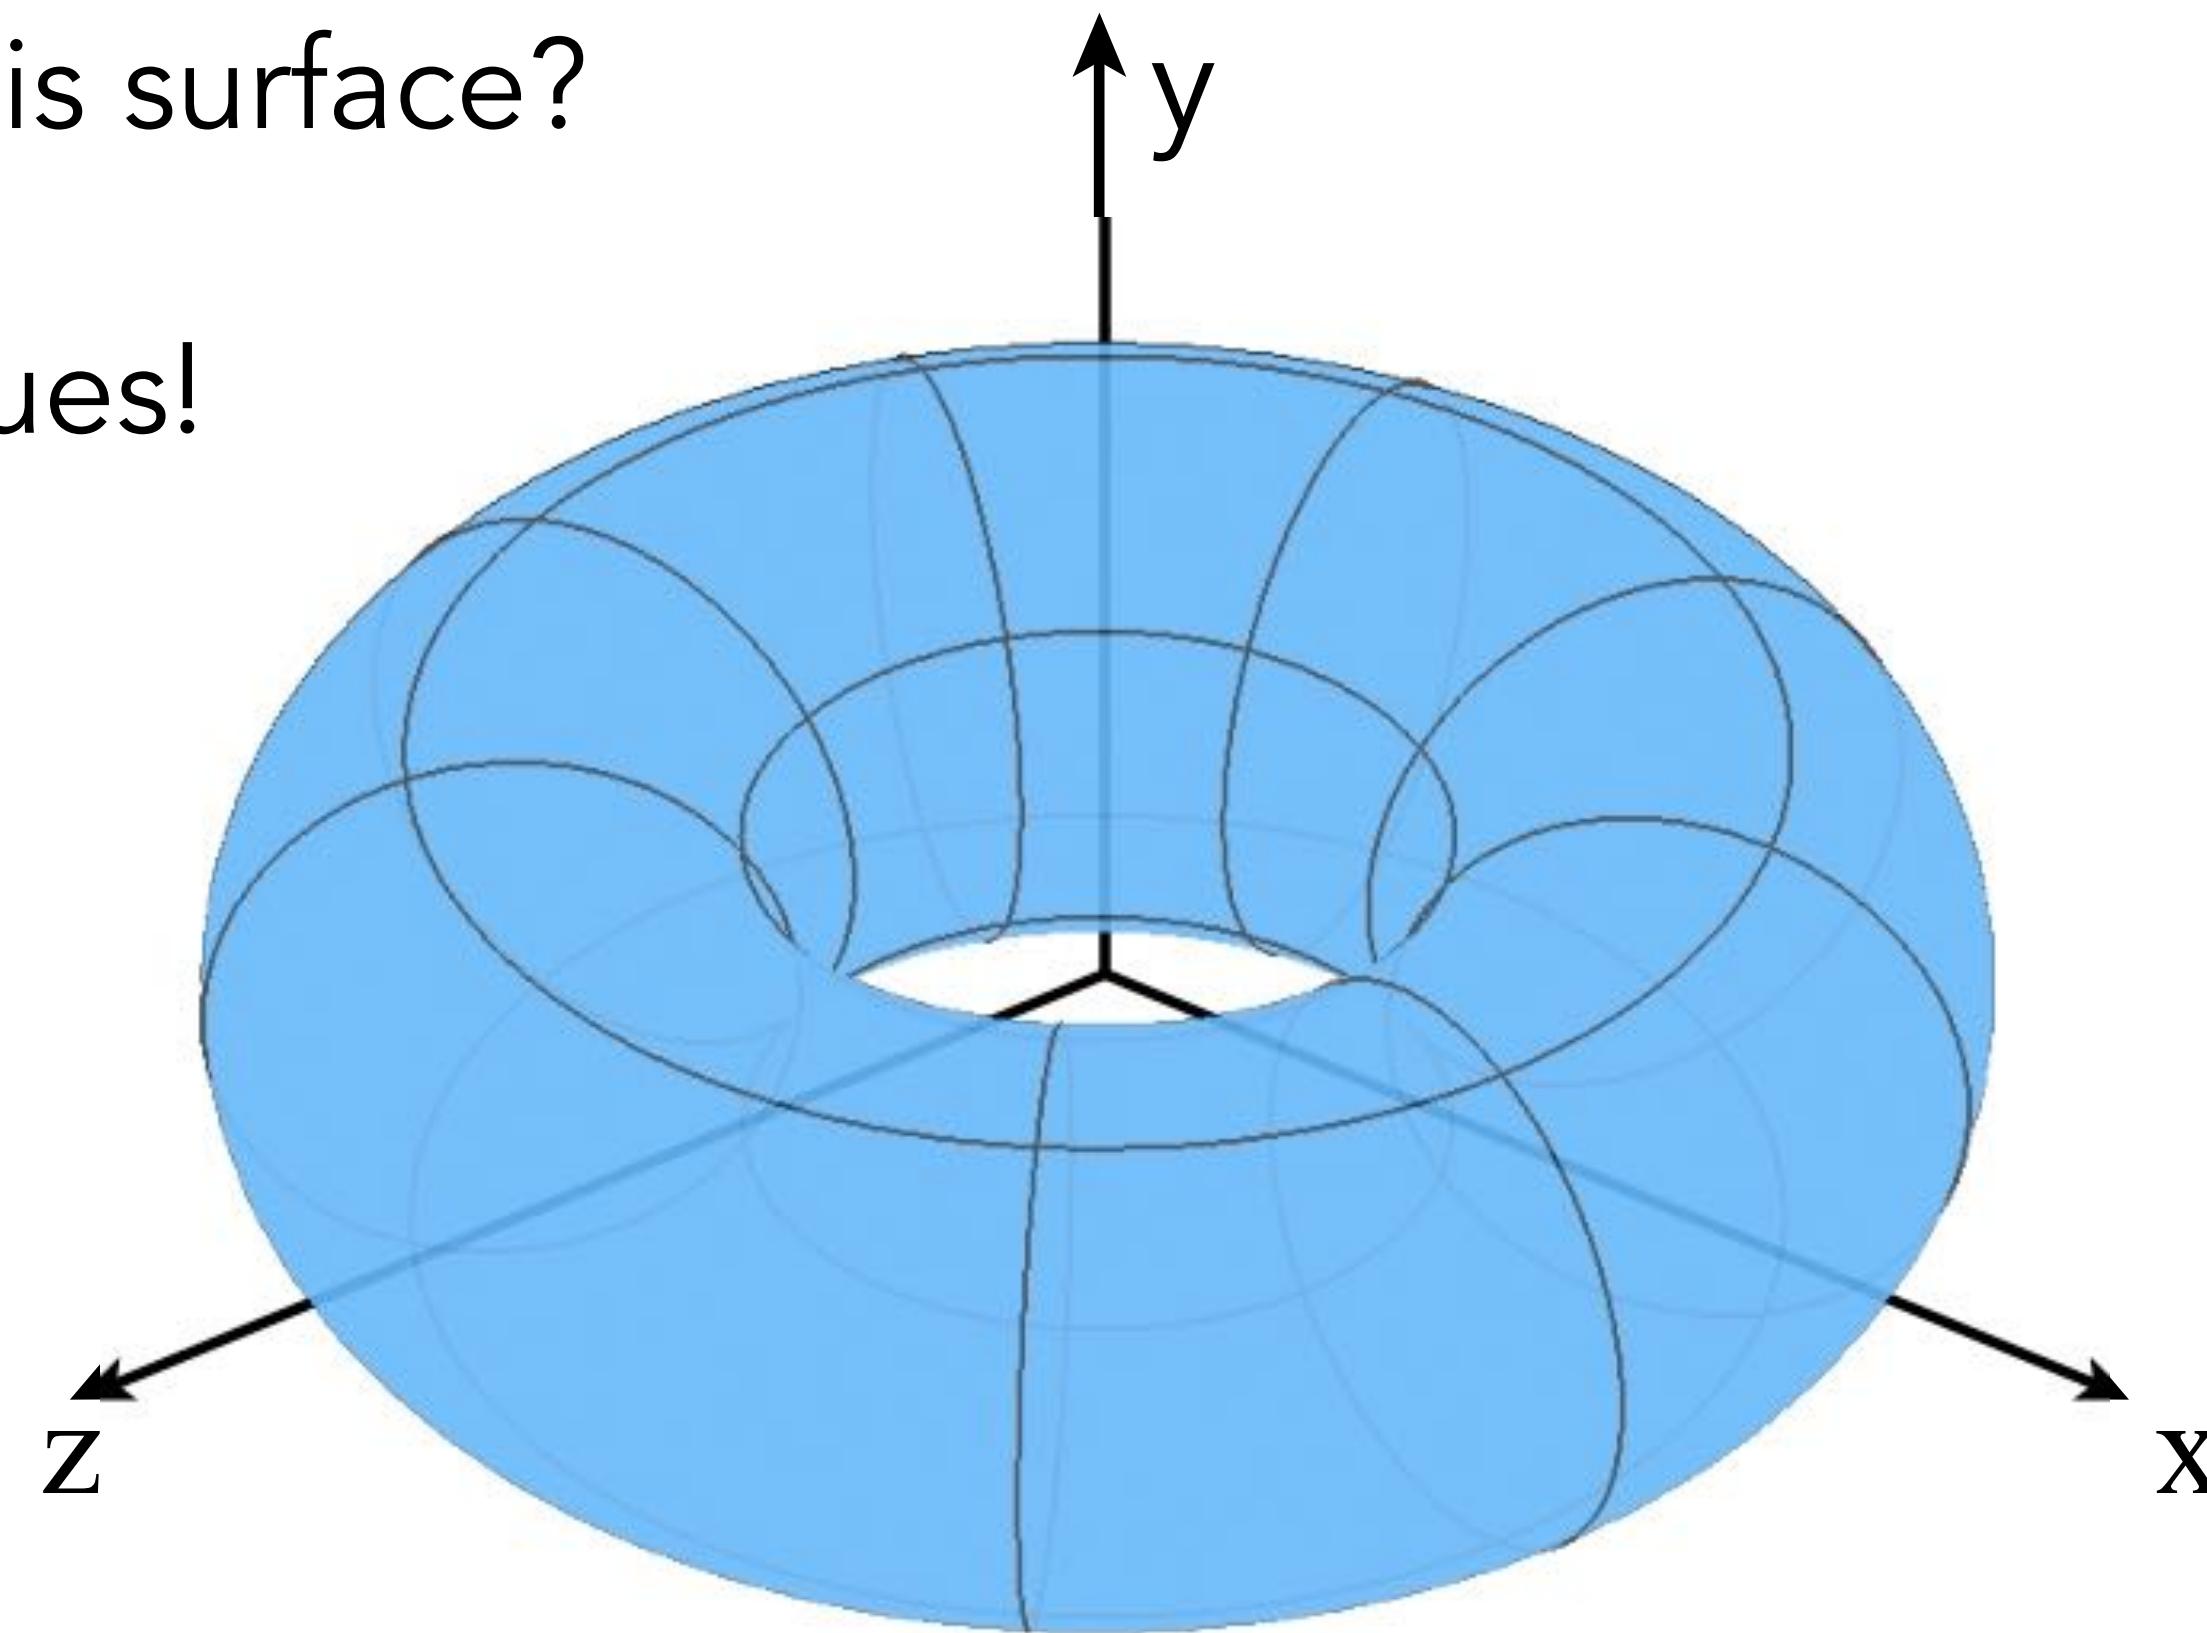


Explicit Surface – Sampling Is Easy

$$f(u, v) = ((2 + \cos u)\cos v, (2 + \cos u)\sin v, \sin u)$$

What points lie on this surface?

Just plug in (u, v) values!

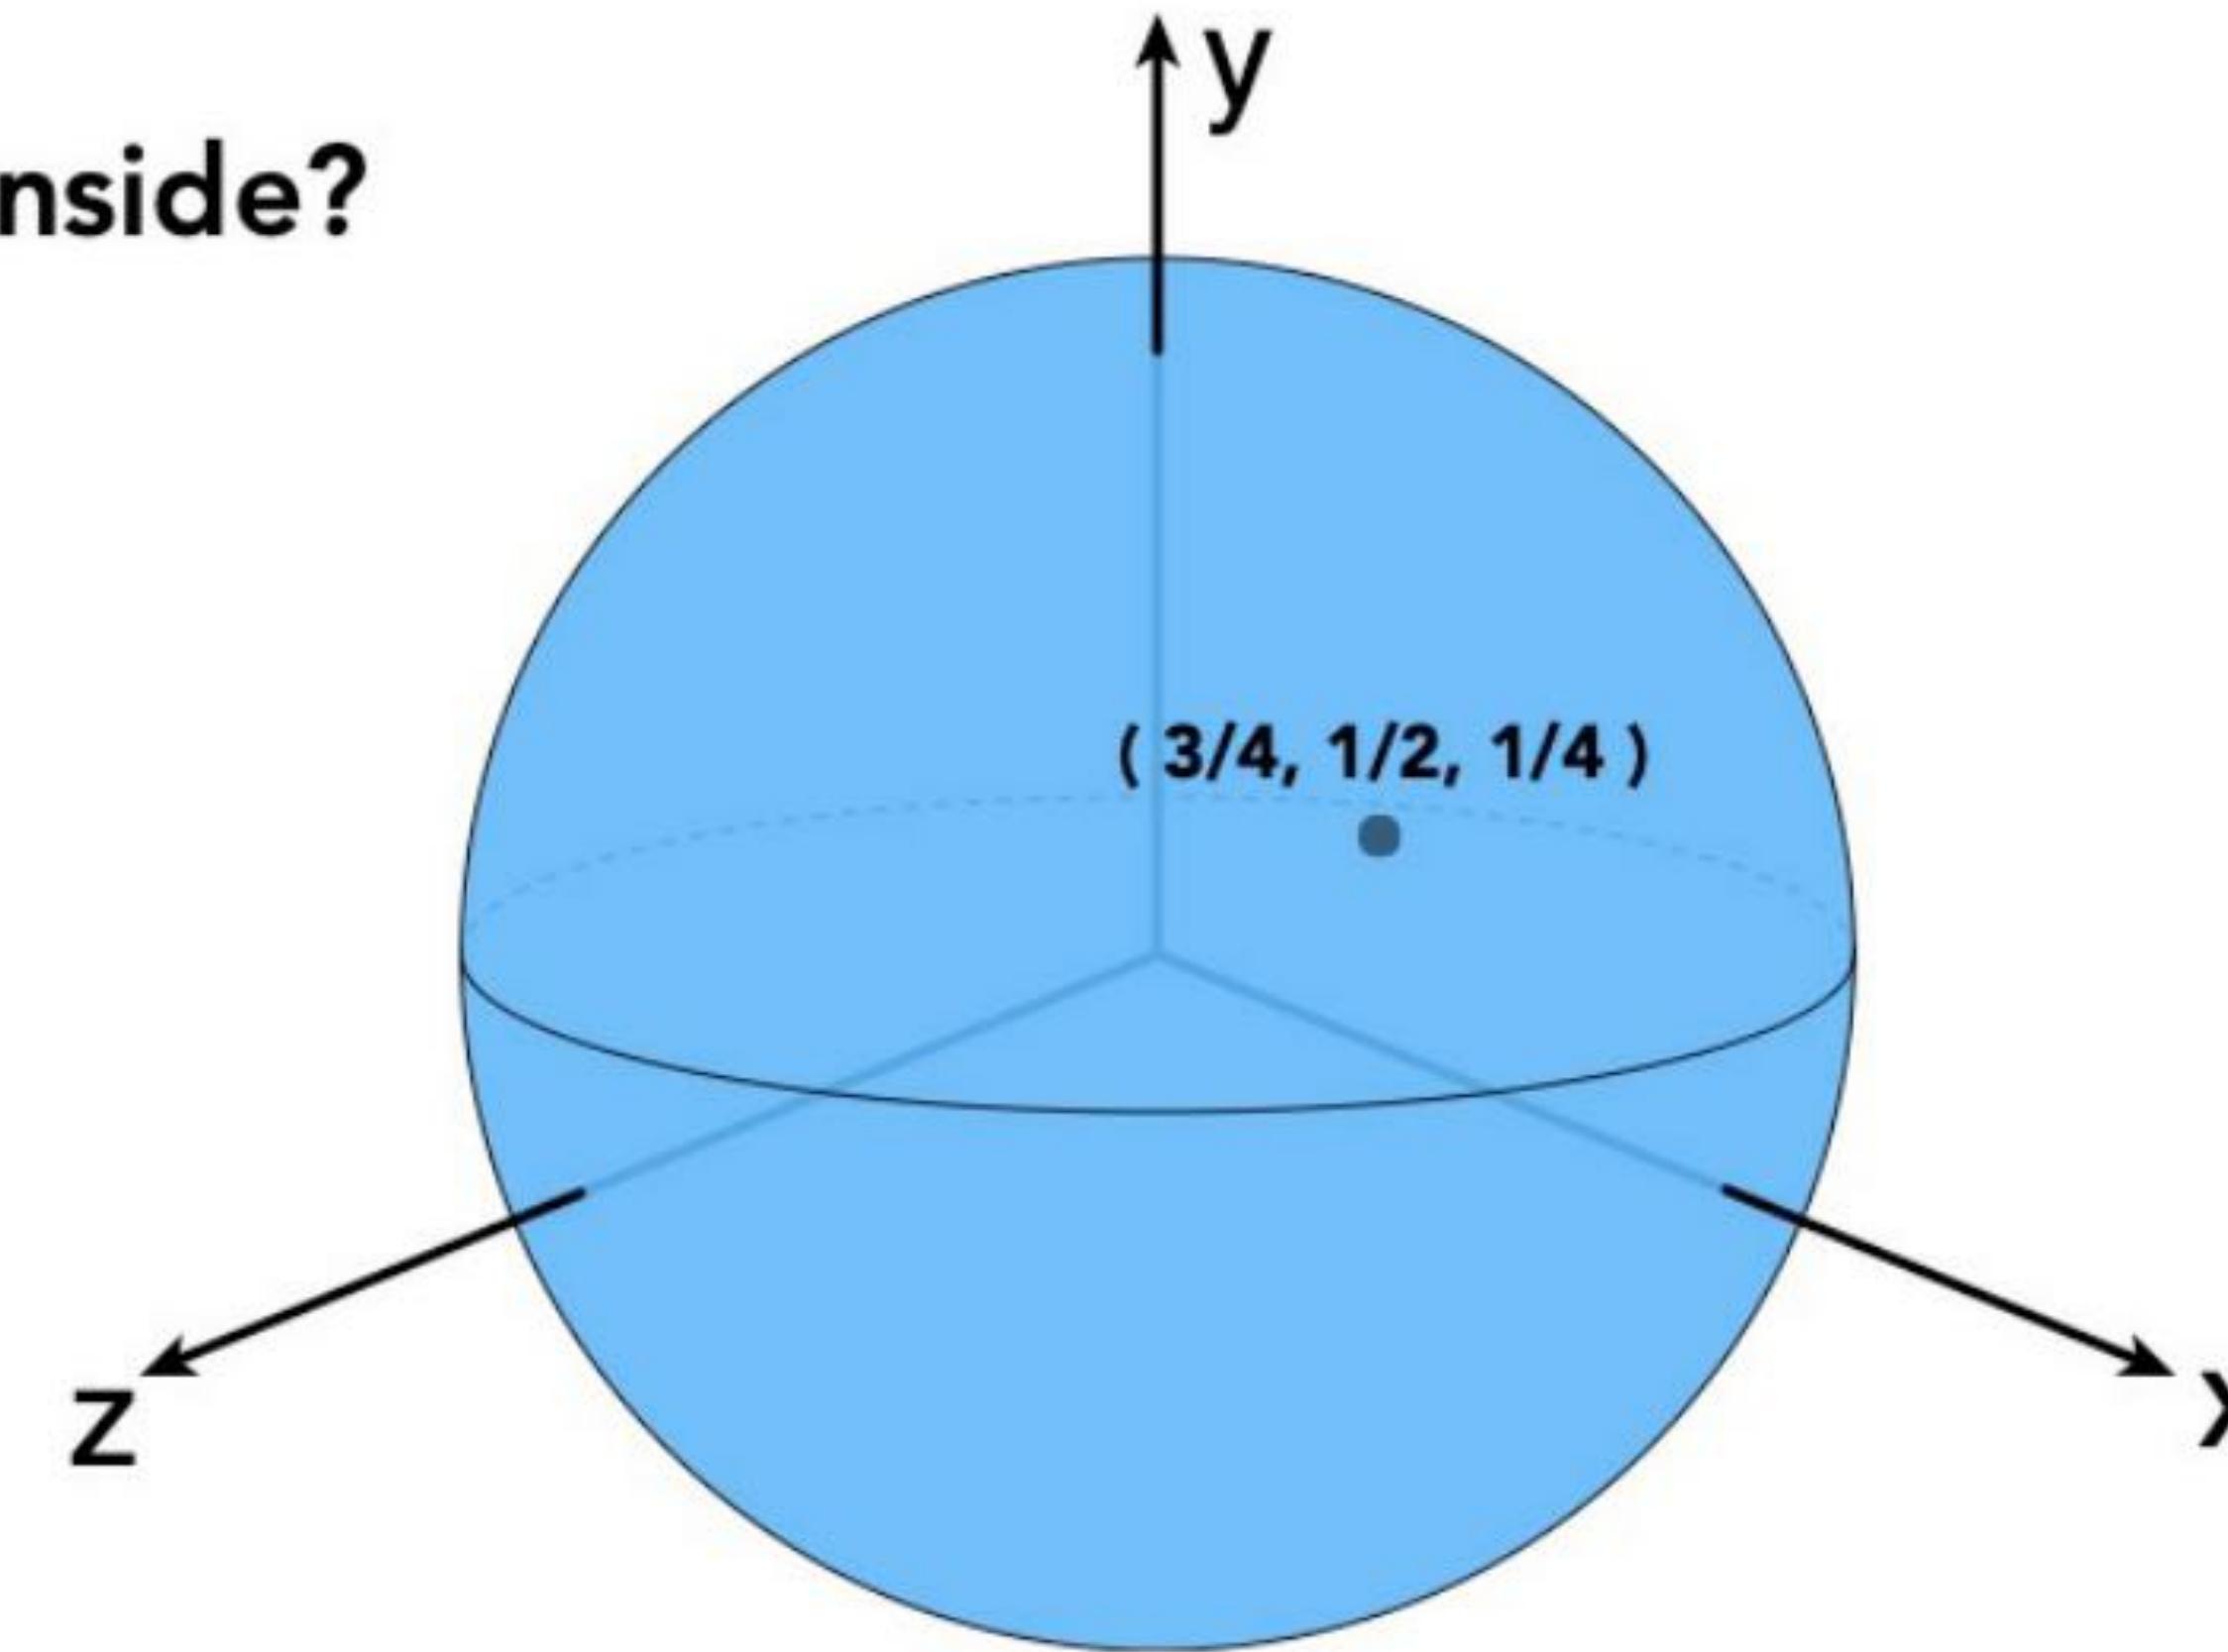


Explicit representations make some tasks easy

Explicit Surface - Inside/Outside Test Hard

$$f(u, v) = (\cos u \sin v, \sin u \sin v, \cos v)$$

Is $(3/4, 1/2, 1/4)$ inside?



Some tasks are hard with explicit representations

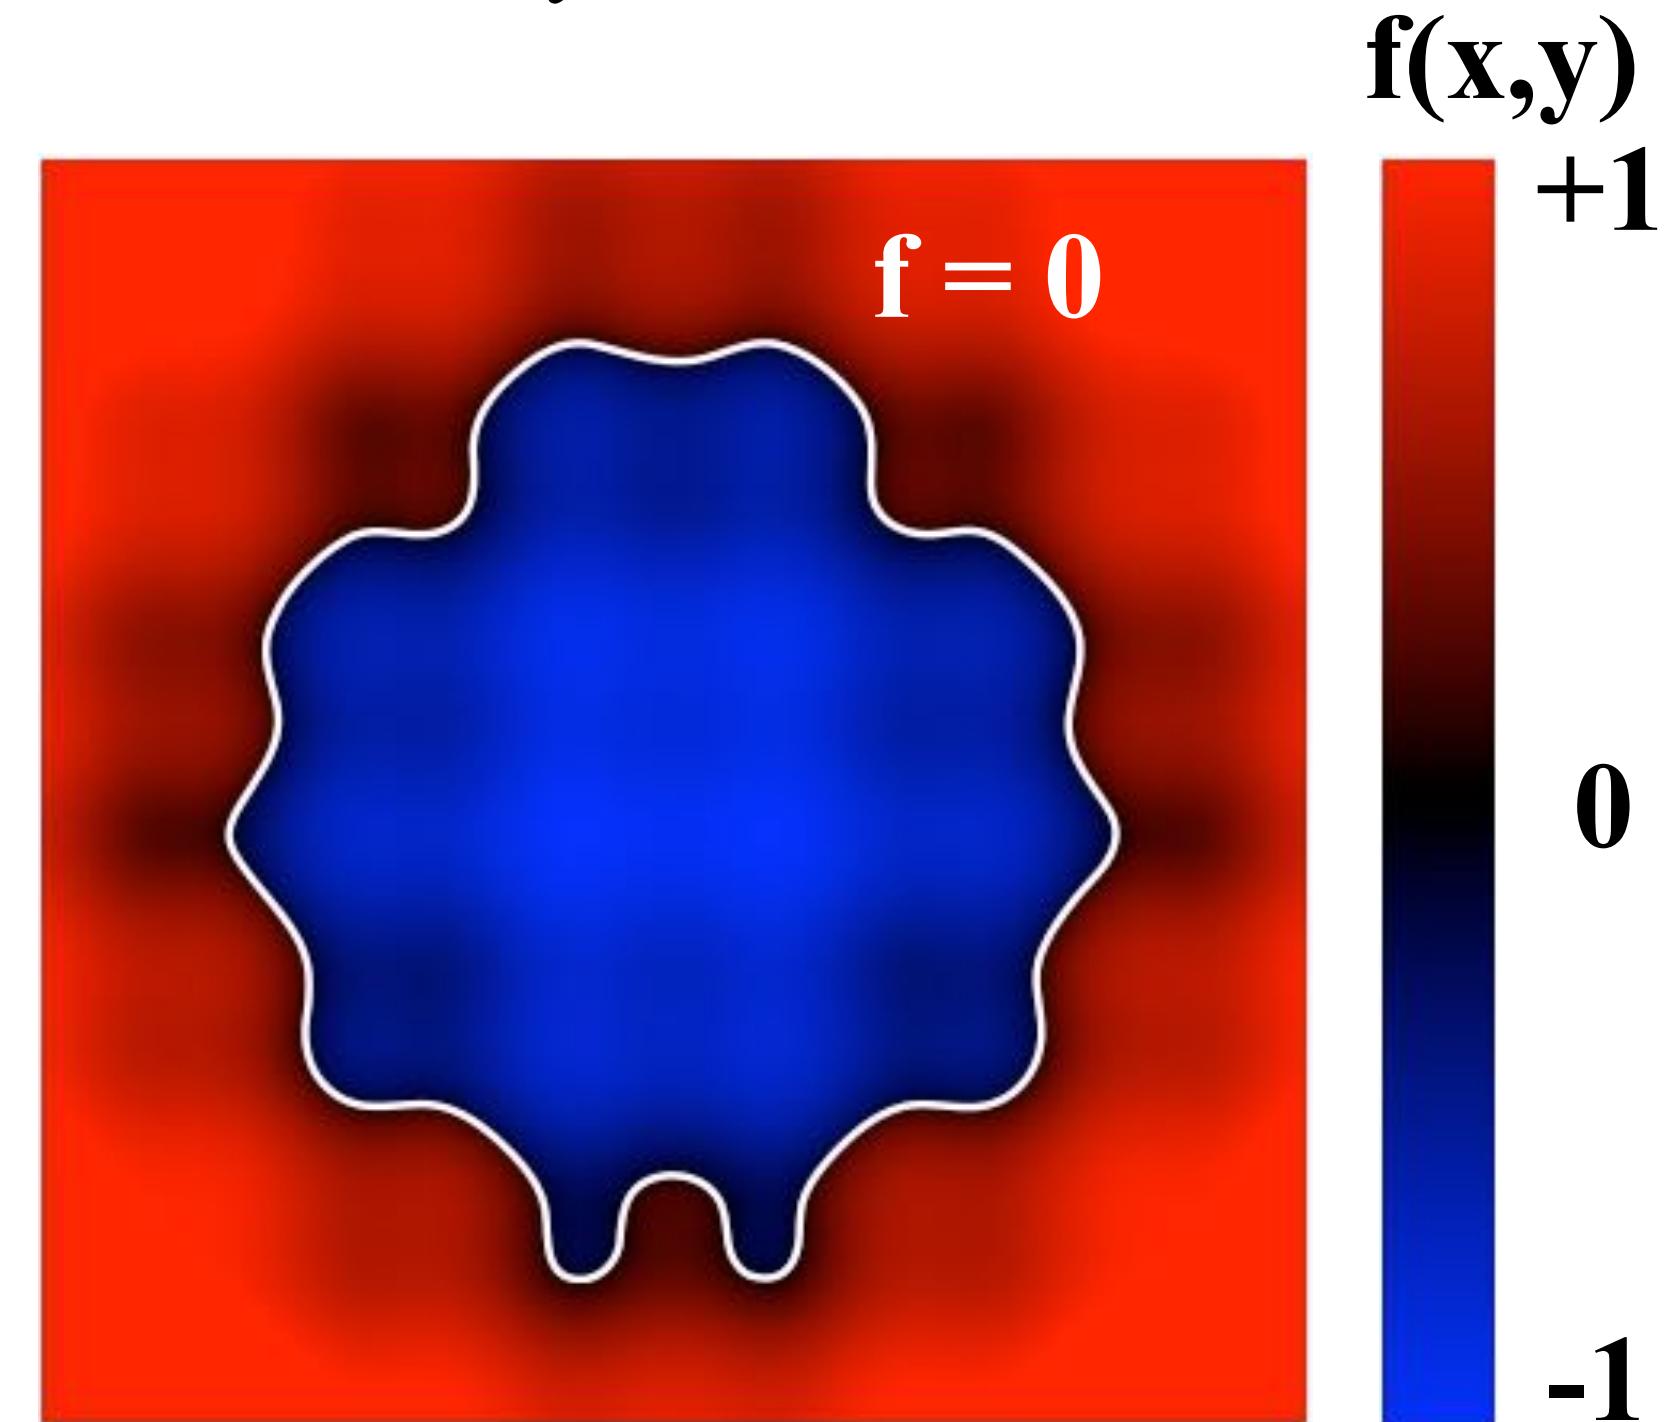
Implicit Representations of Geometry

Based on classifying points

- Points satisfy some specified relationship

E.g. sphere: all points in 3D, where $x^2 + y^2 + z^2 = 1$

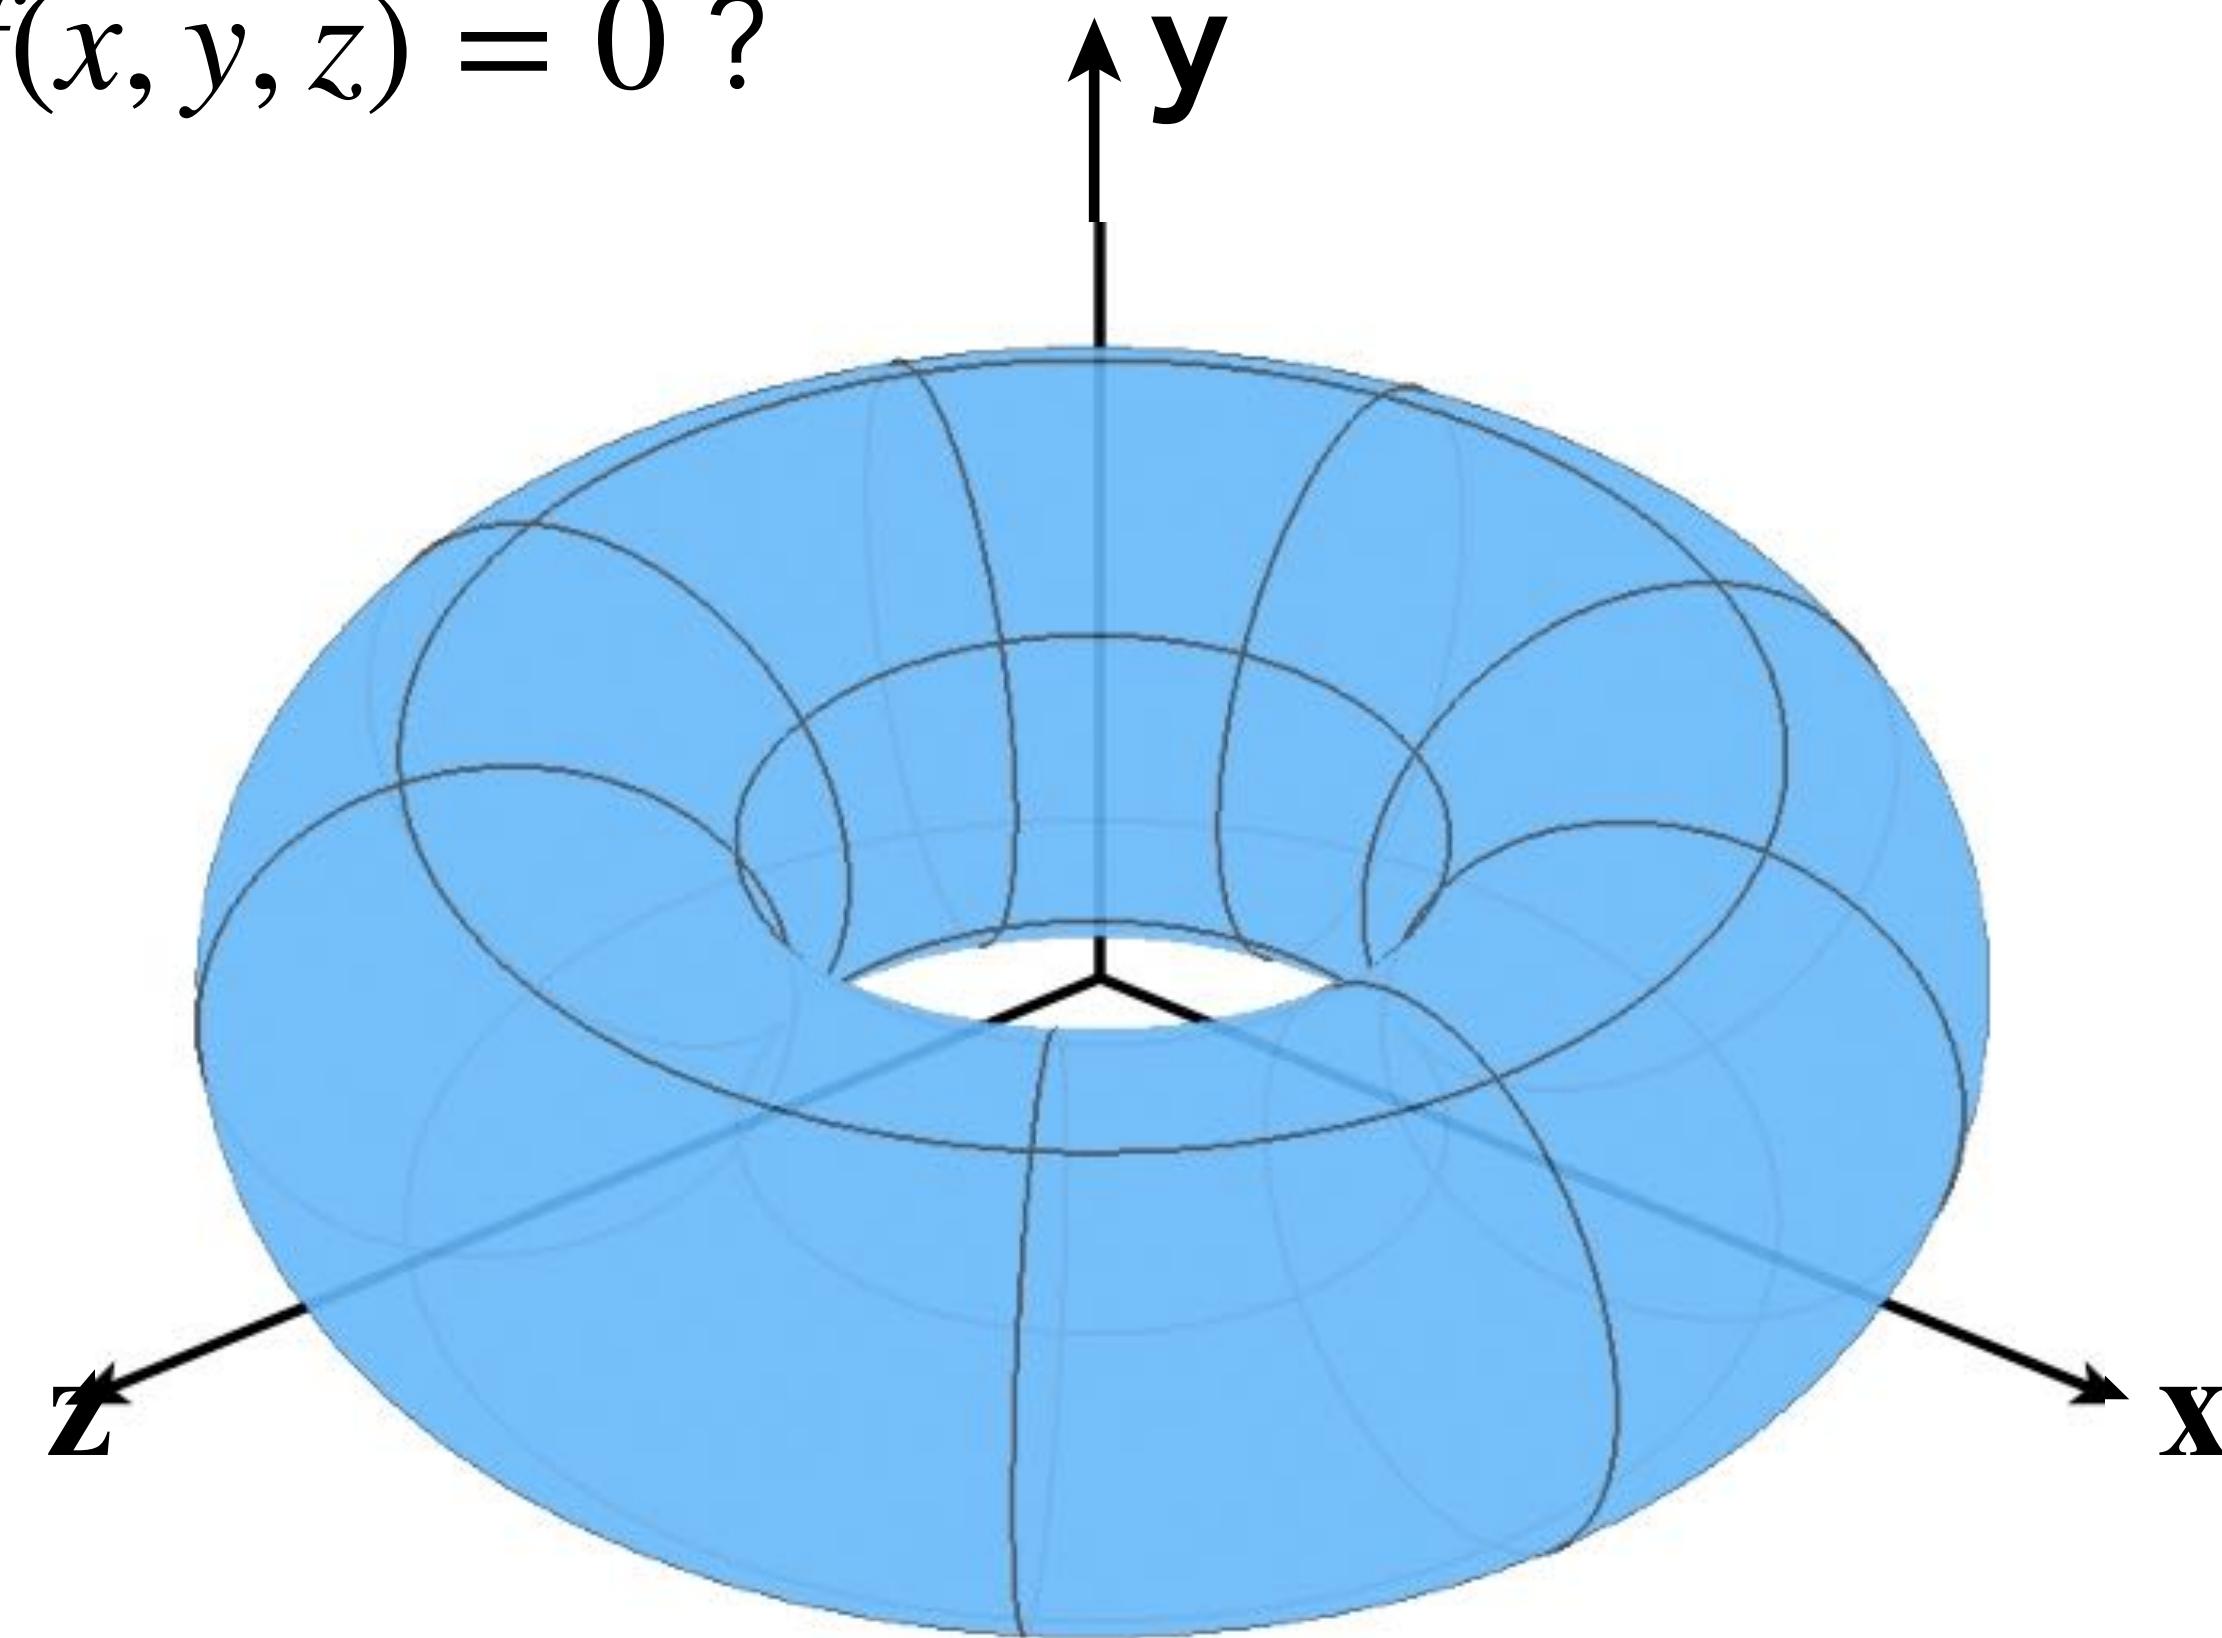
More generally, $f(x, y, z) = 0$



Implicit Surface – Sampling Can Be Hard

$$f(x, y, z) = (2 - \sqrt{x^2 + y^2})^2 + z^2 - 1$$

What points lie on $f(x, y, z) = 0$?



Some tasks are hard with implicit representations

Implicit Surface - Inside/Outside Tests Easy

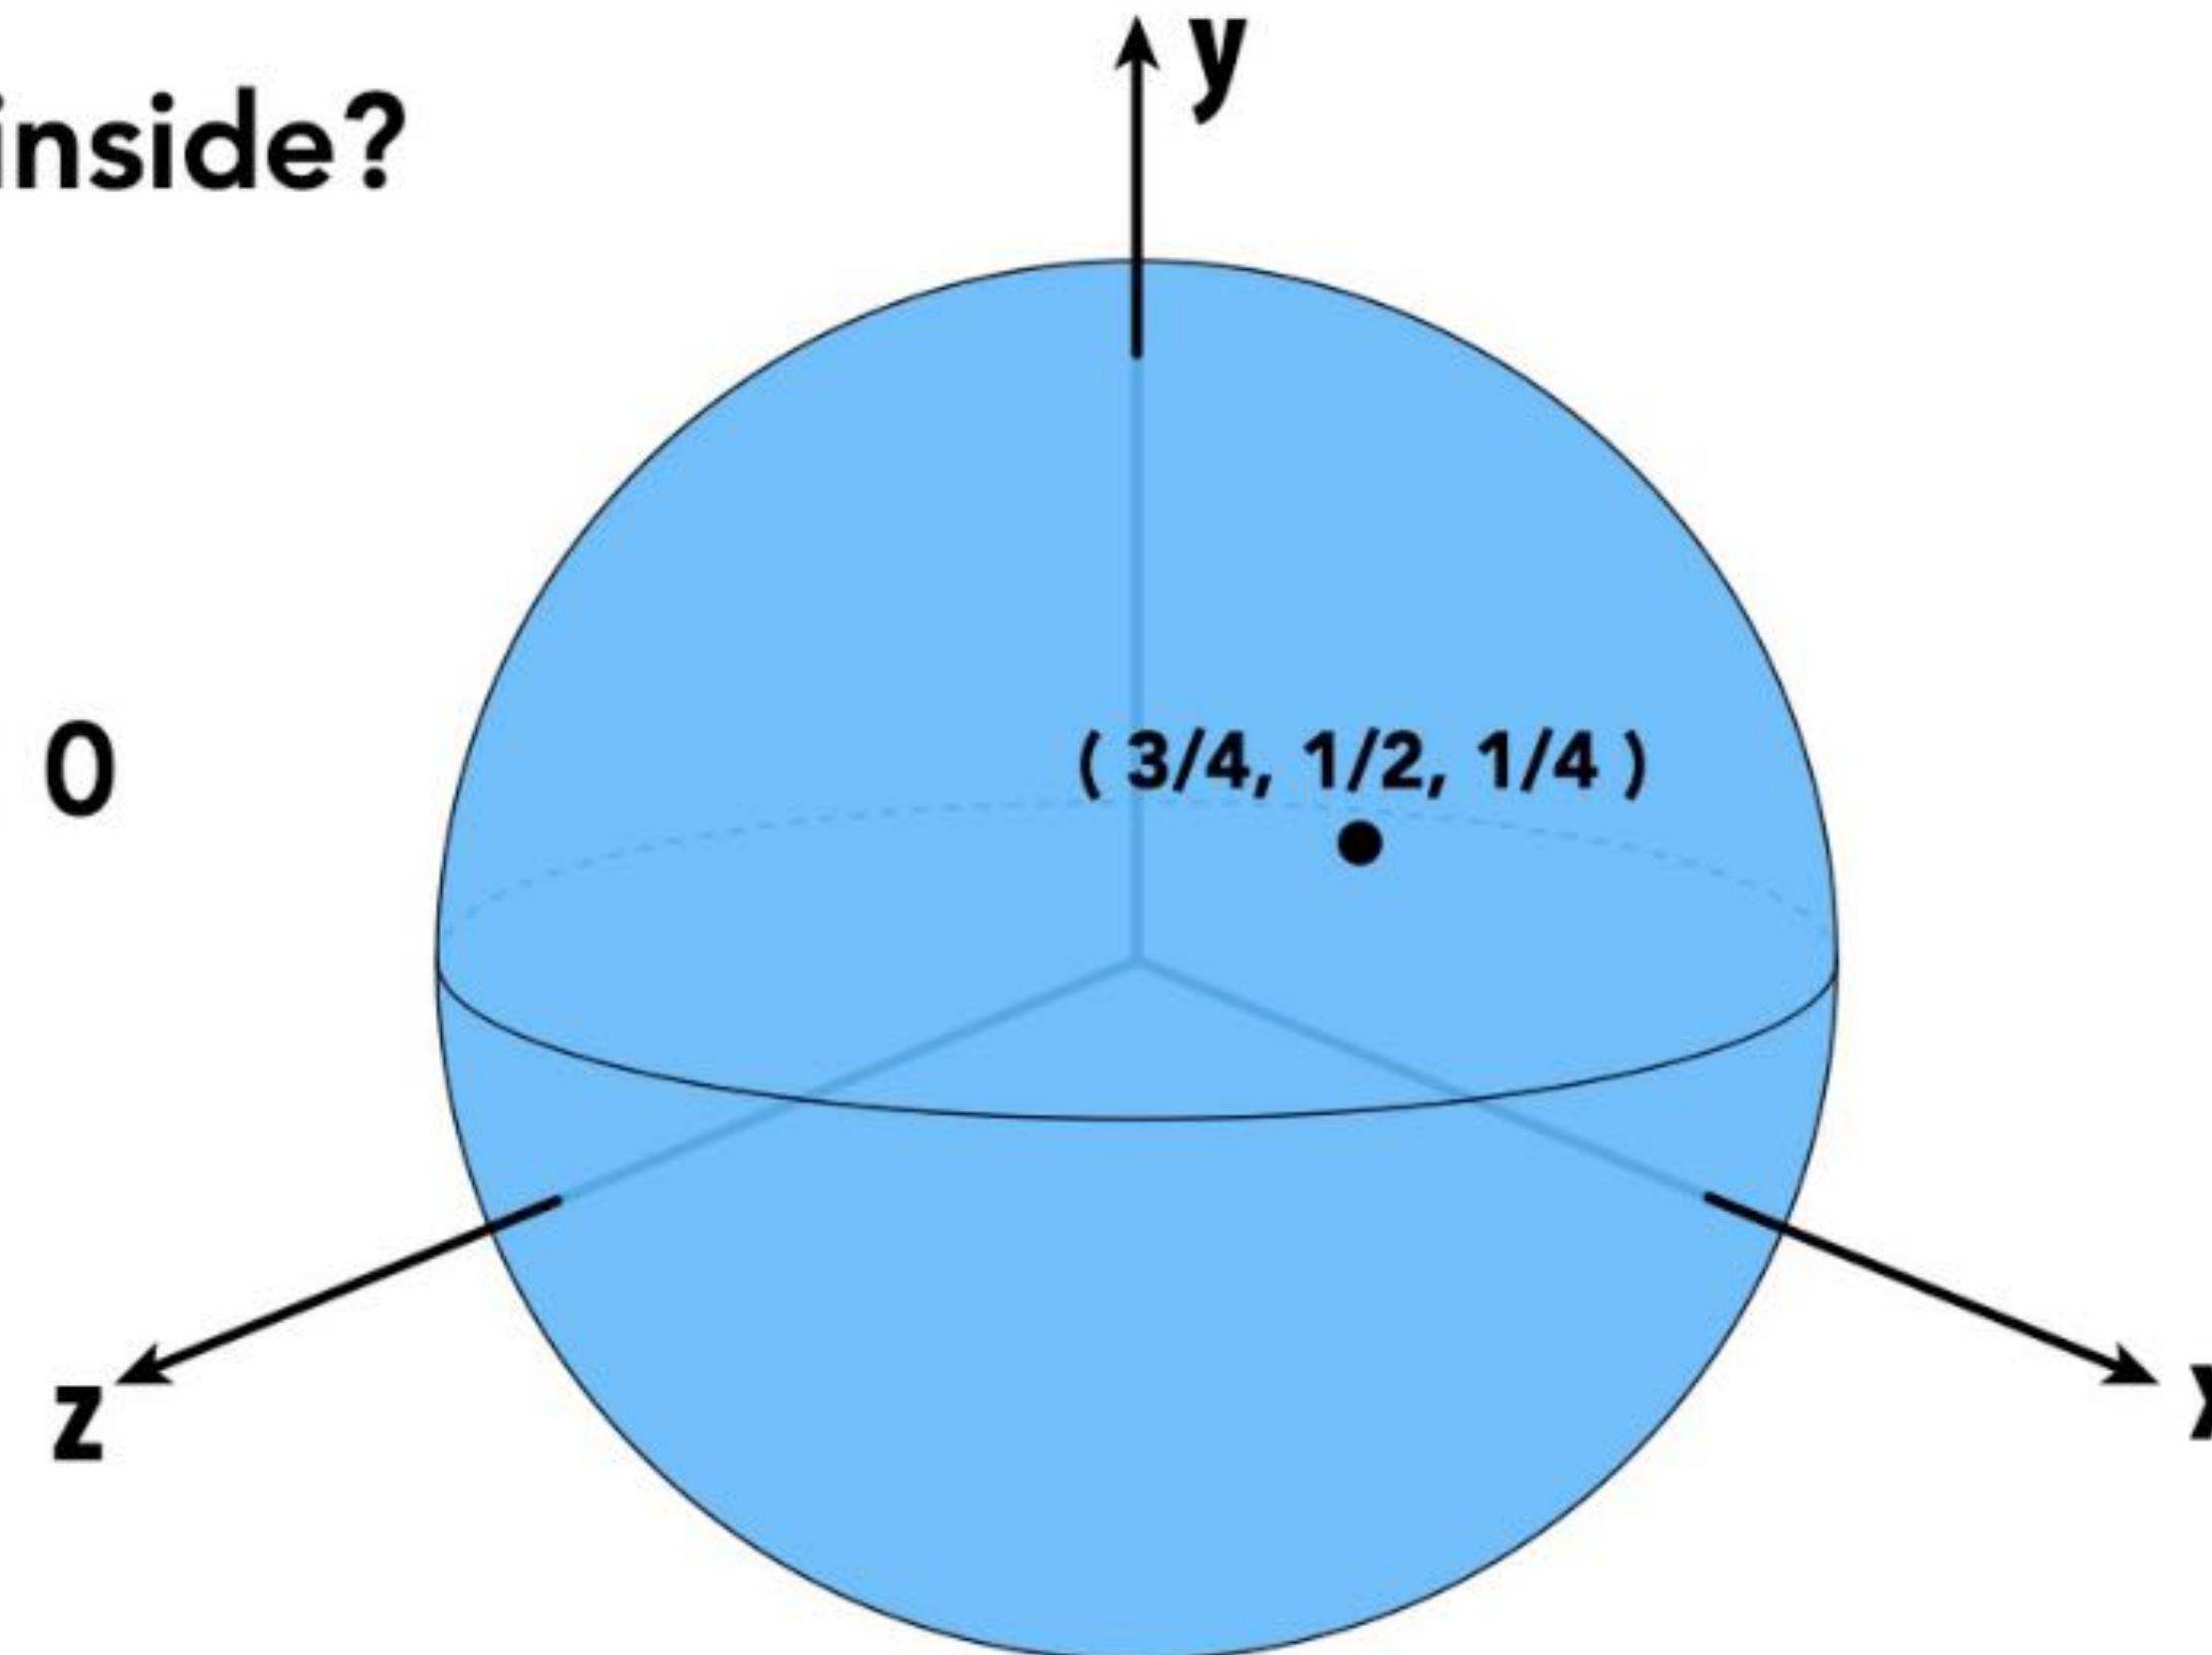
$$f(x, y, z) = x^2 + y^2 + z^2 - 1$$

Is $(3/4, 1/2, 1/4)$ inside?

Just plug it in:

$$f(x, y, z) = -1/8 < 0$$

Yes, inside.



Implicit representations make some tasks easy

Algebraic Surfaces (Implicit)

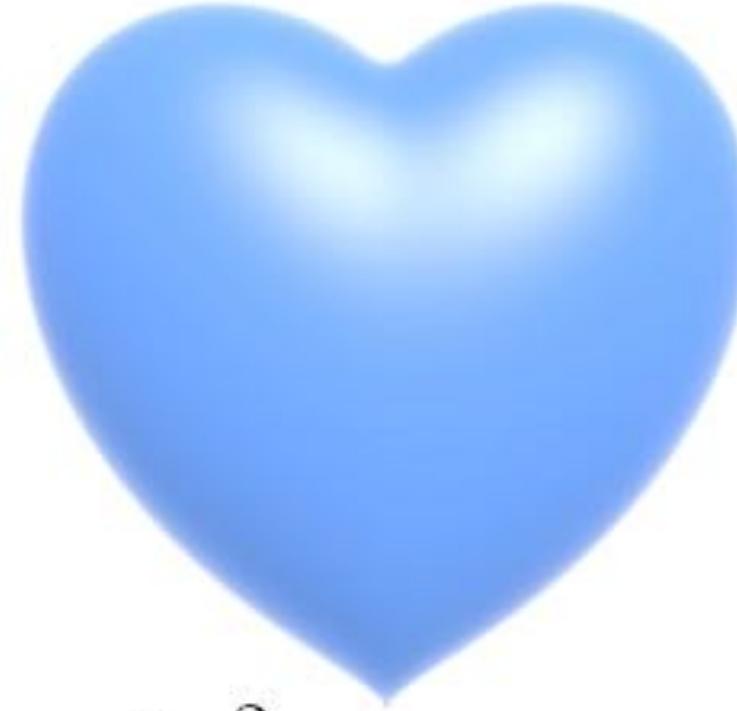
Surface is zero set of a polynomial in x, y, z



$$x^2 + y^2 + z^2 = 1$$



$$(R - \sqrt{x^2 + y^2})^2 + z^2 = r^2$$



$$(x^2 + \frac{9y^2}{4} + z^2 - 1)^3 =$$

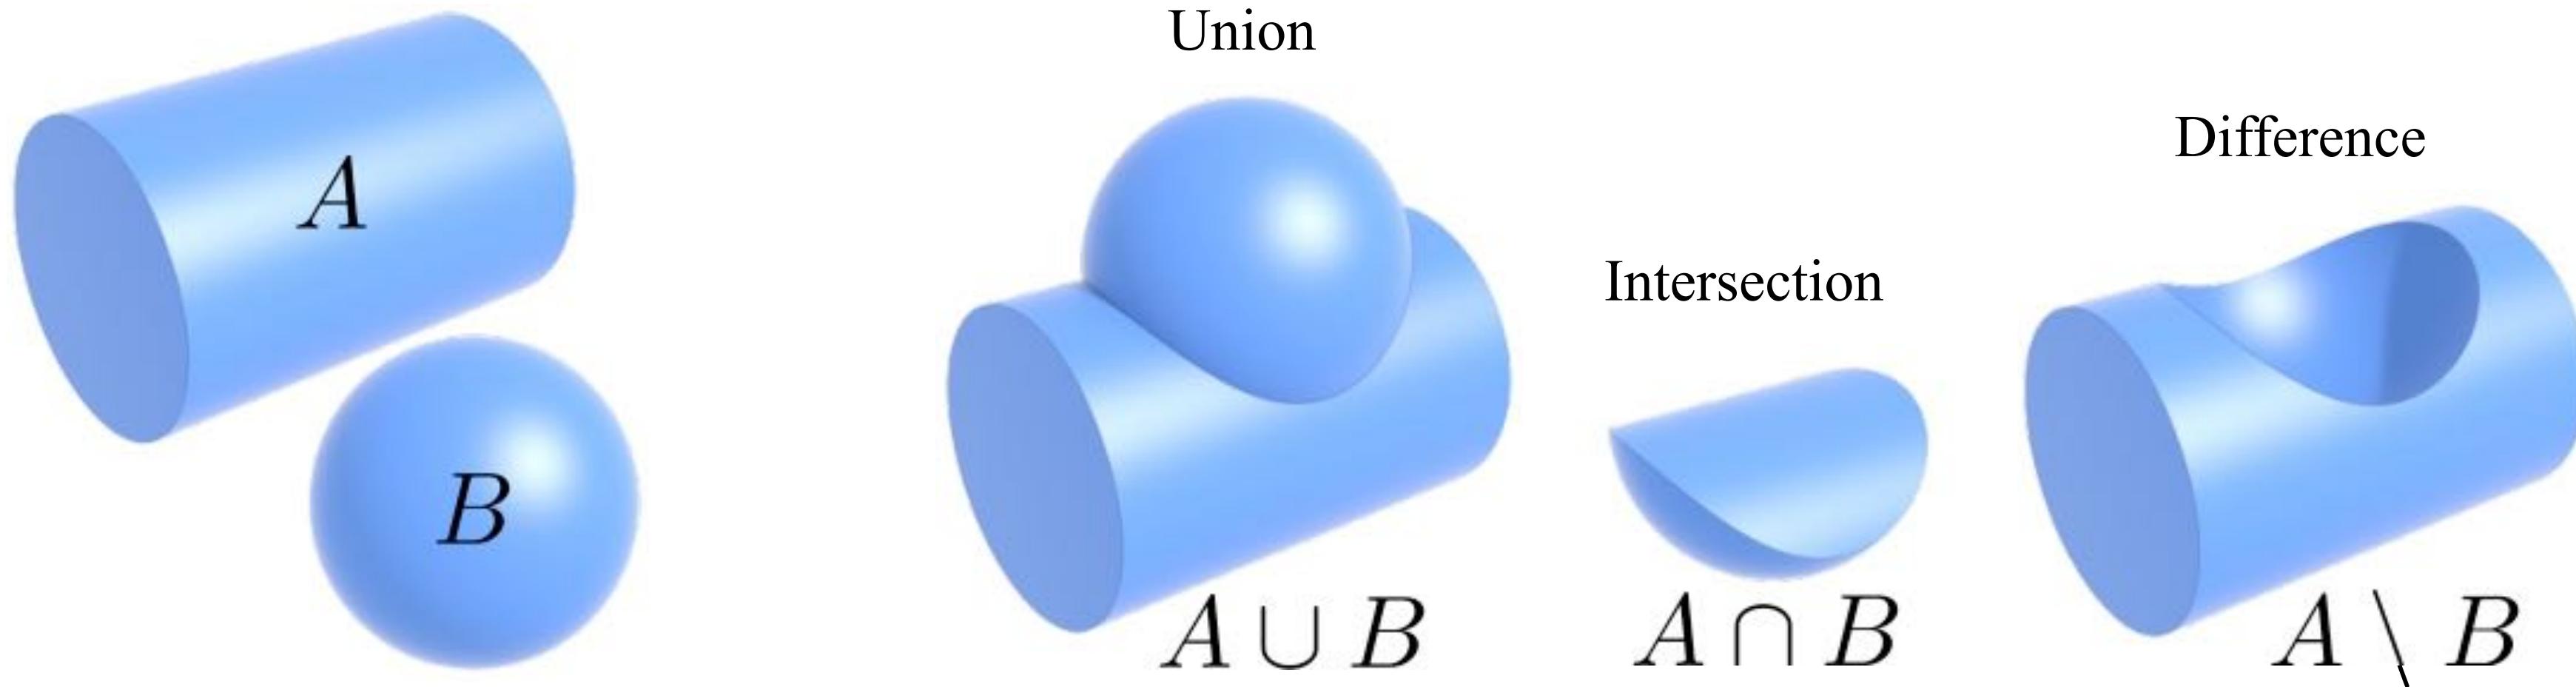
$$x^2 z^3 + \frac{9y^2 z^3}{80}$$



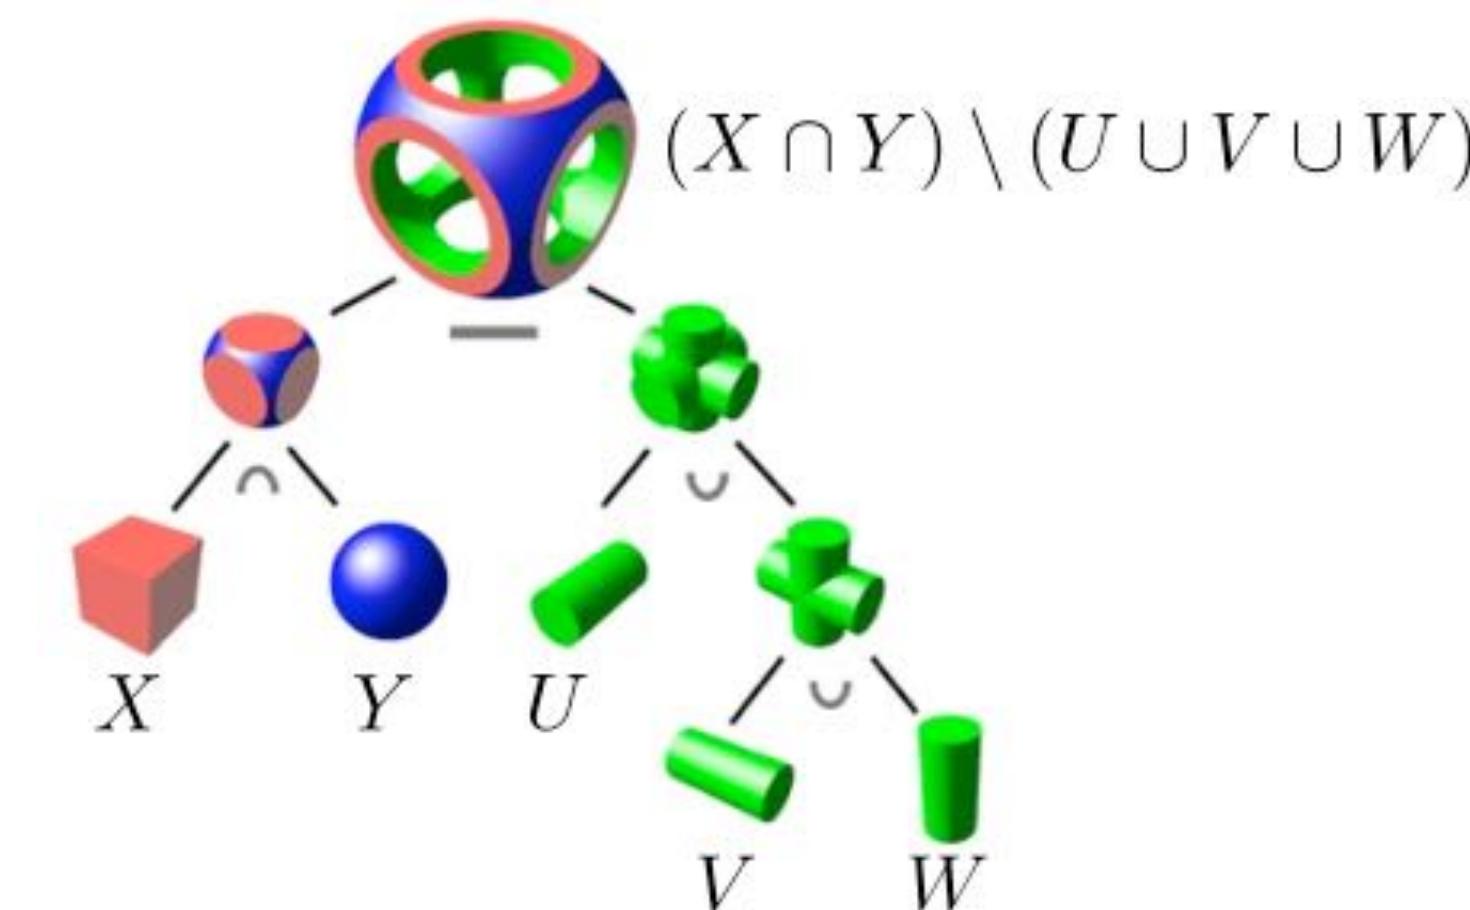
More complex shapes?

Constructive Solid Geometry (Implicit)

Combine implicit geometry via Boolean operations

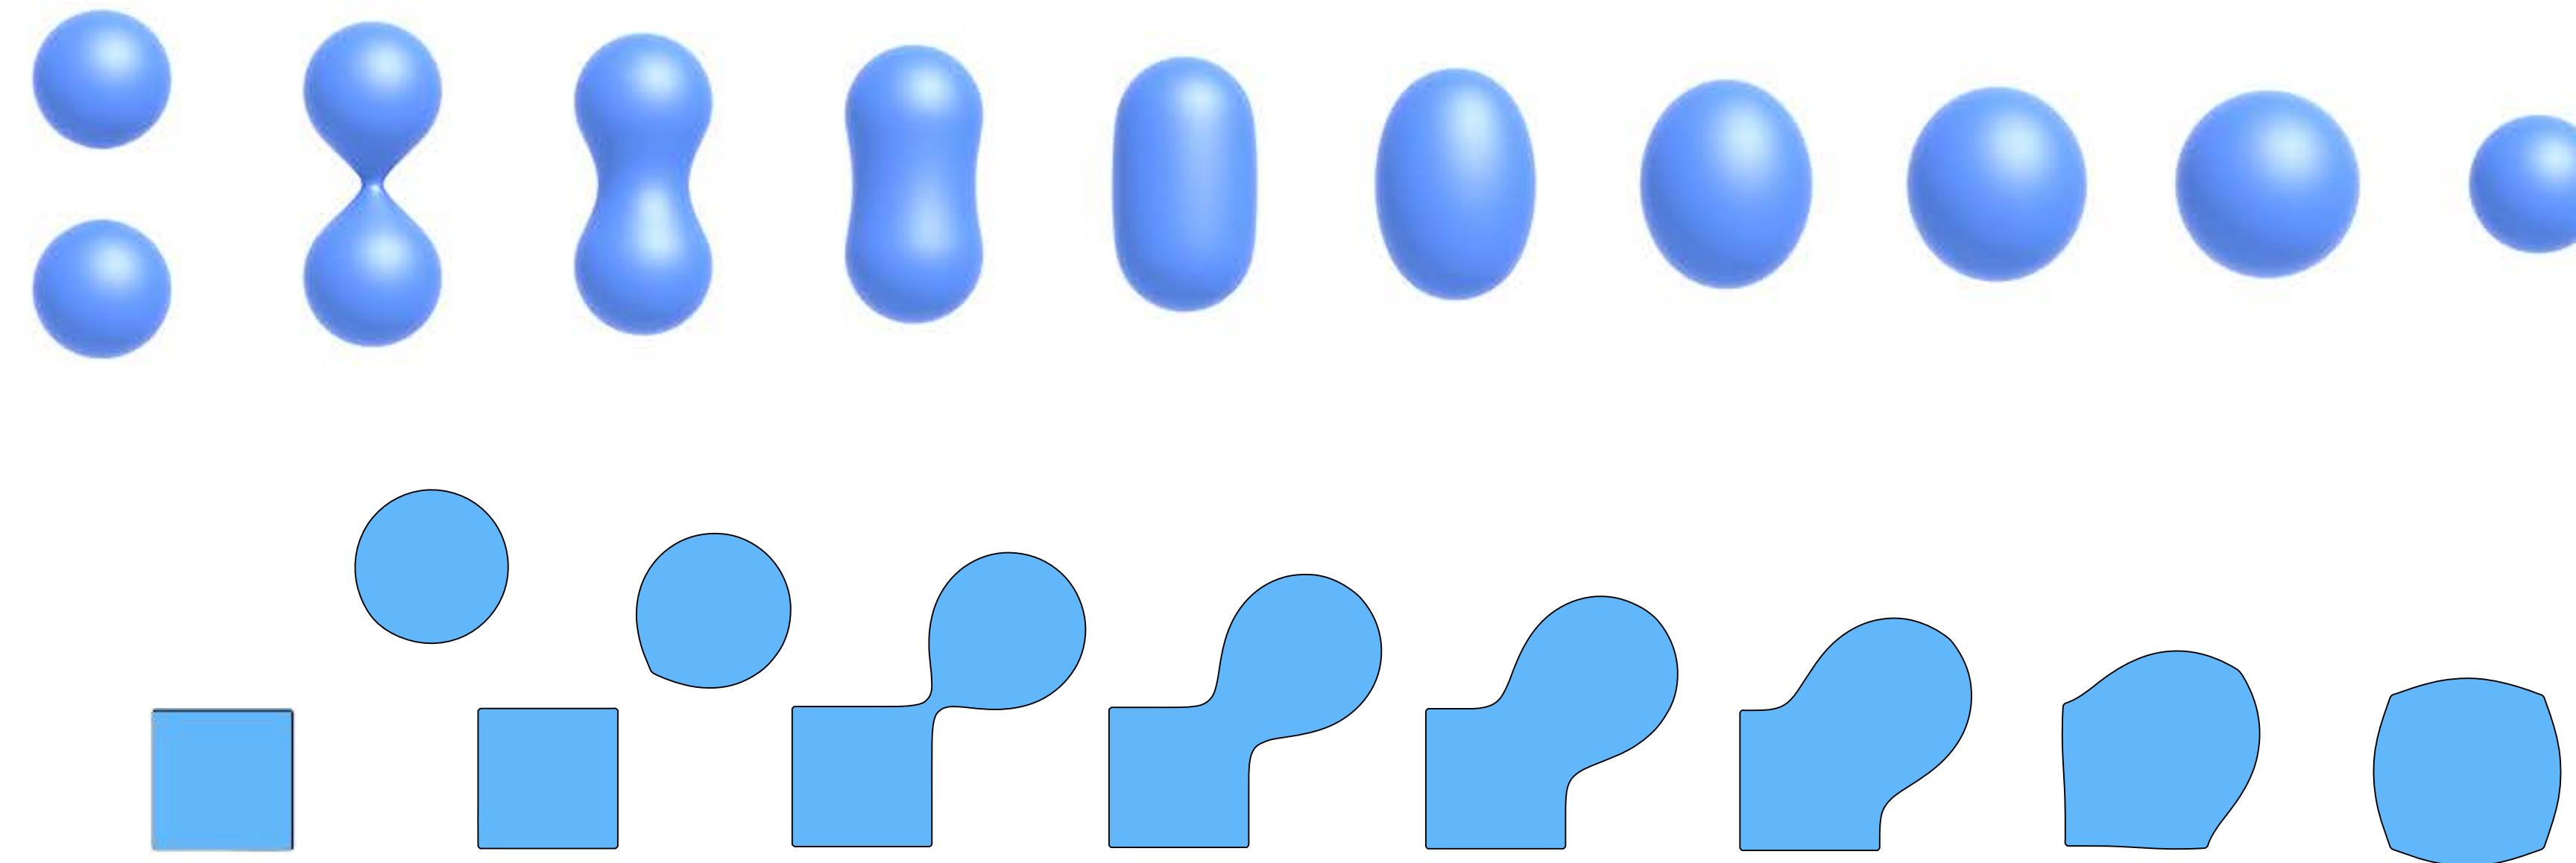


Boolean expressions :



Distance Functions

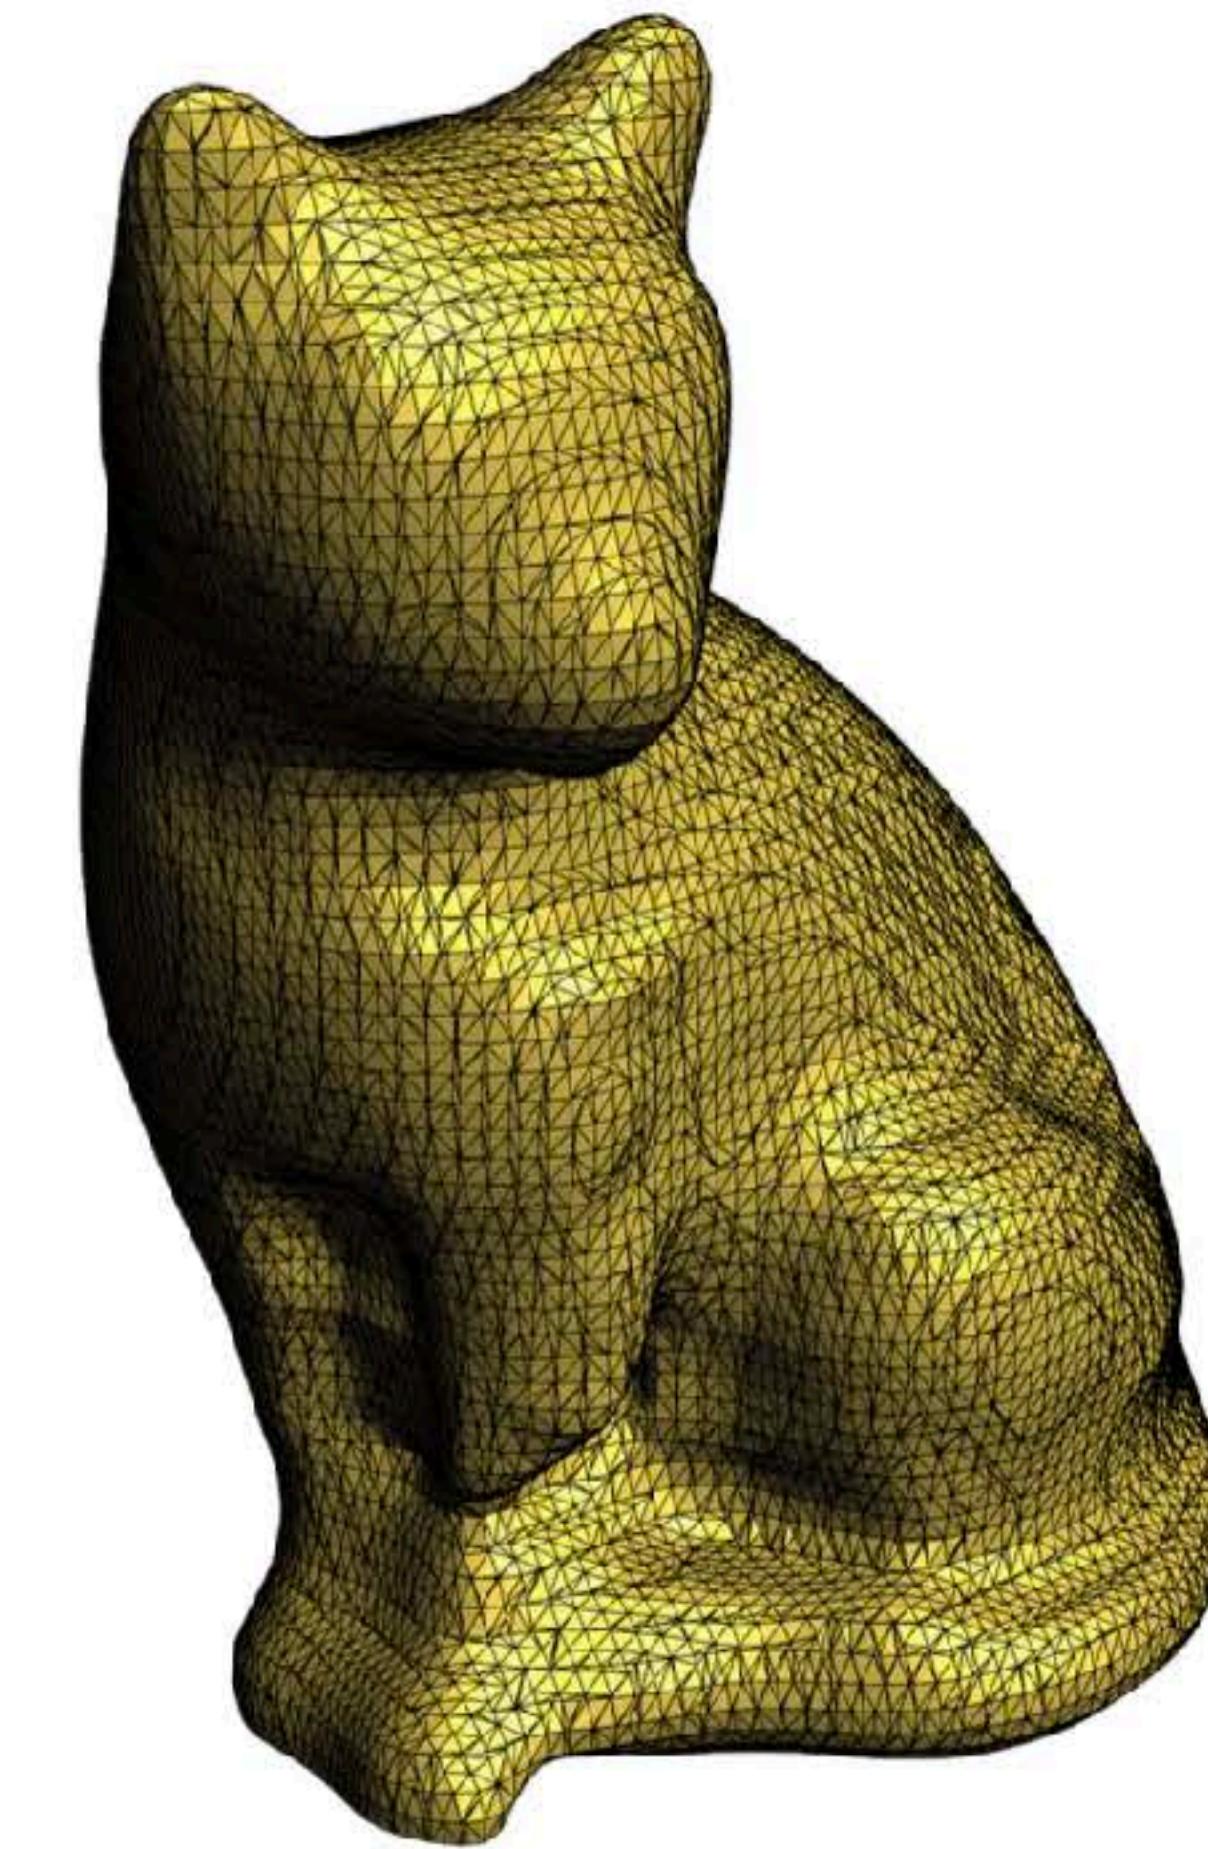
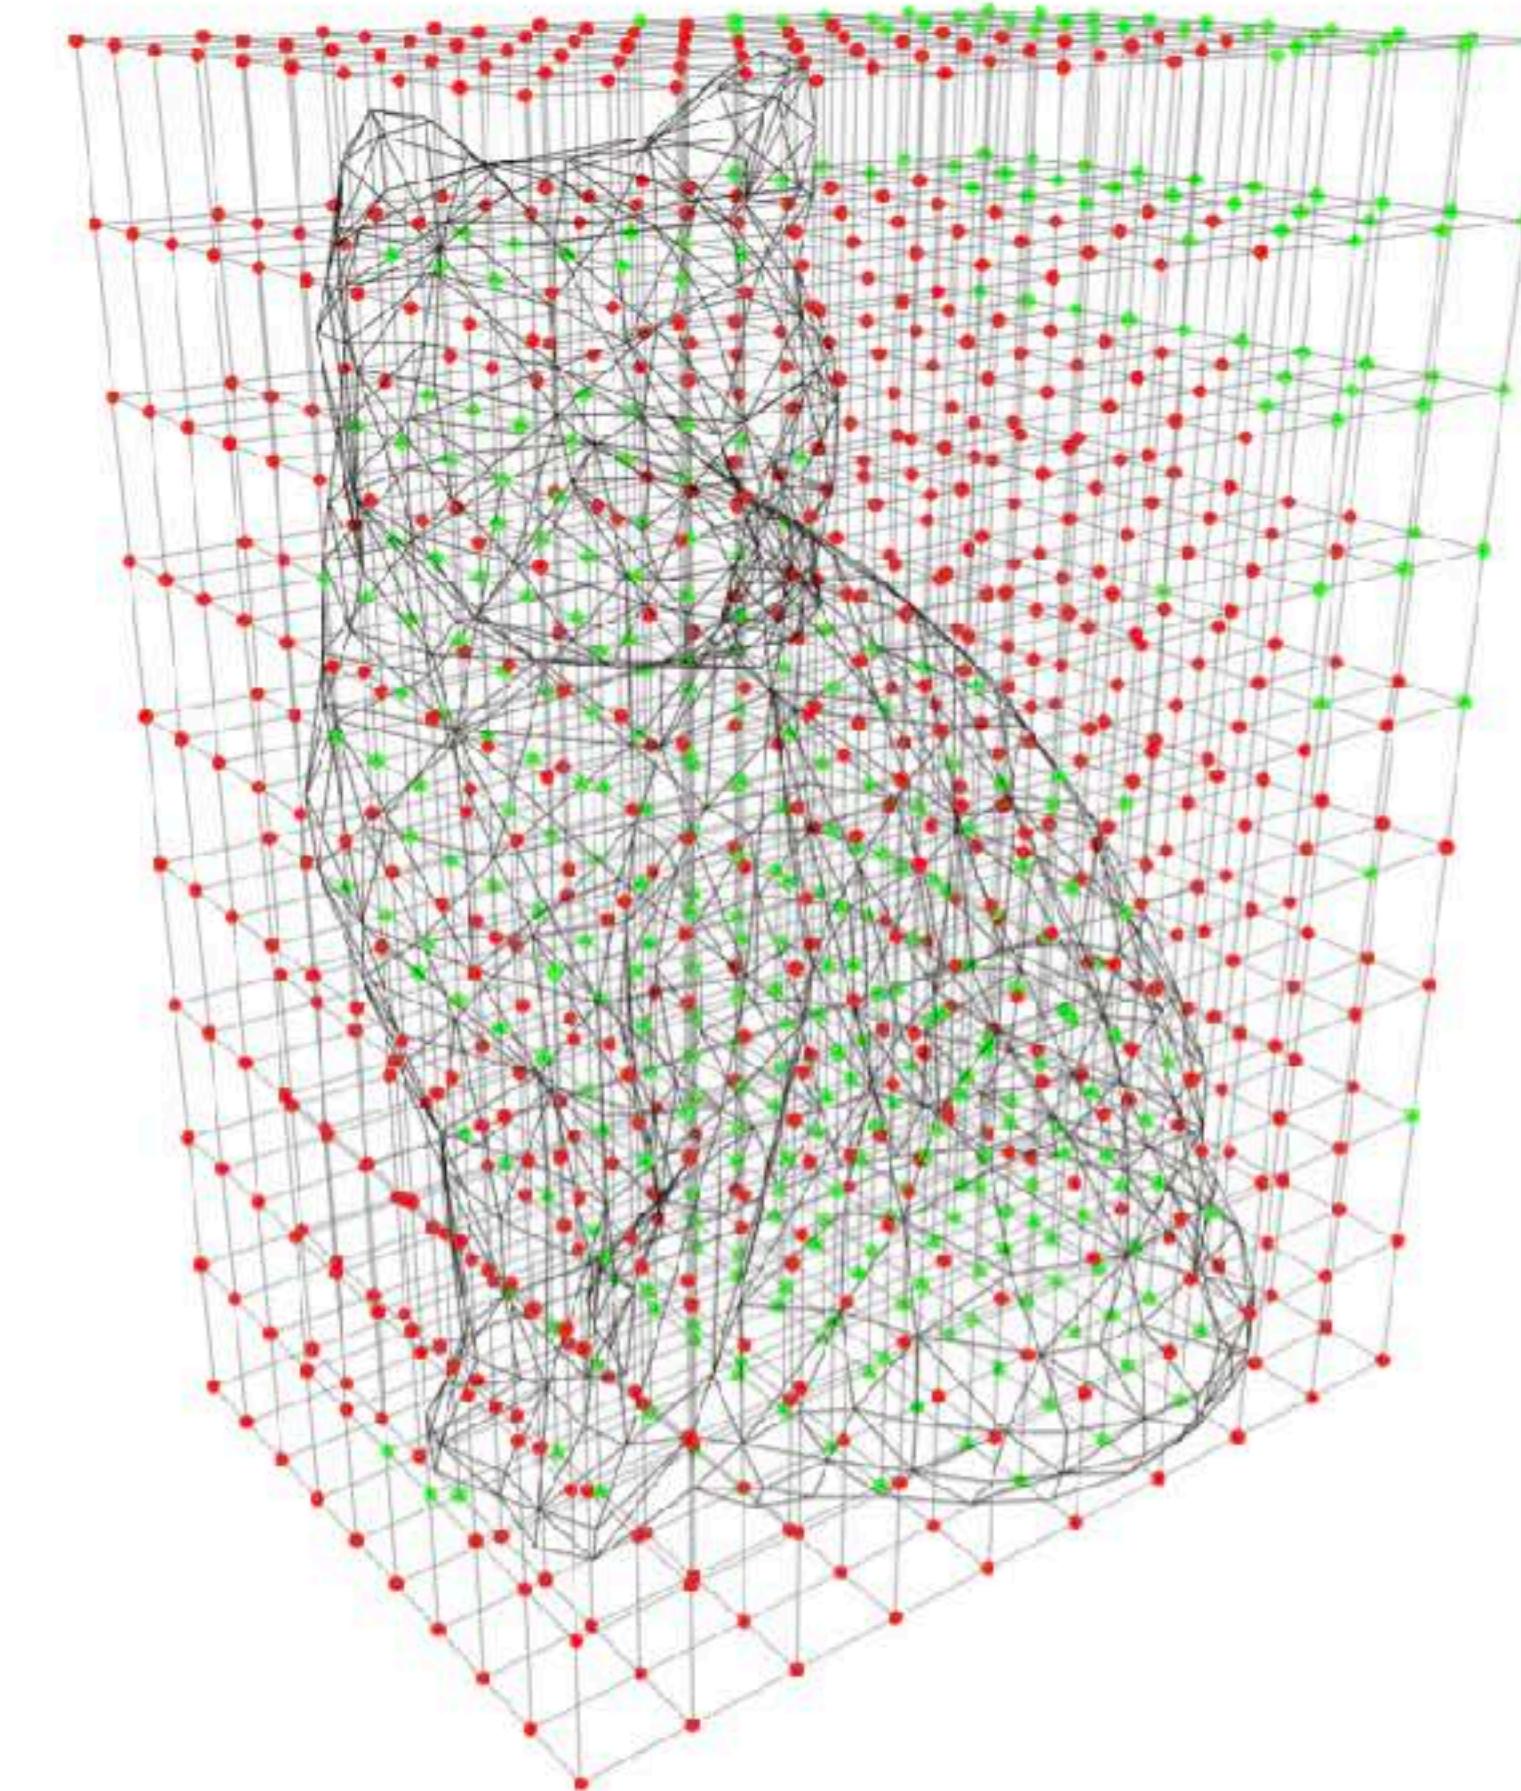
- Distance functions: giving minimum distance (could be signed distance) from anywhere to object
- Instead of booleans, gradually blend surfaces together using distance functions



Scene of Pure Distance Functions (Not Easy!)



Conversion between different representations



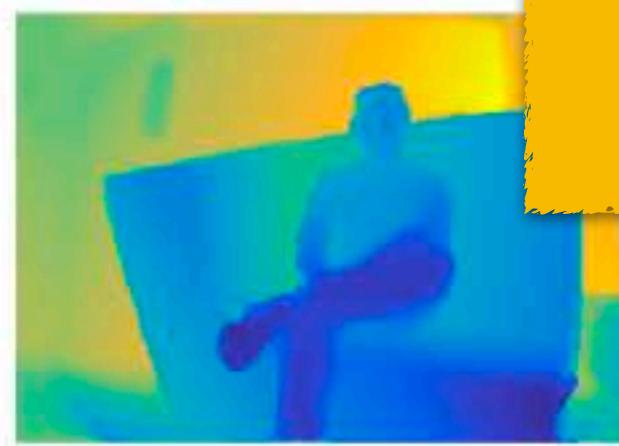
Take Home Message

Rasterized form
(regular grids)

Geometric form
(irregular)

No “best” geometric representation

Choose your representation
depending on the task!



Depth Map



Volumetric



Point Cloud

$$f(x) = 0$$

Implicit Shape

3D Spatial Relationships

- How to represent the relationships between objects?

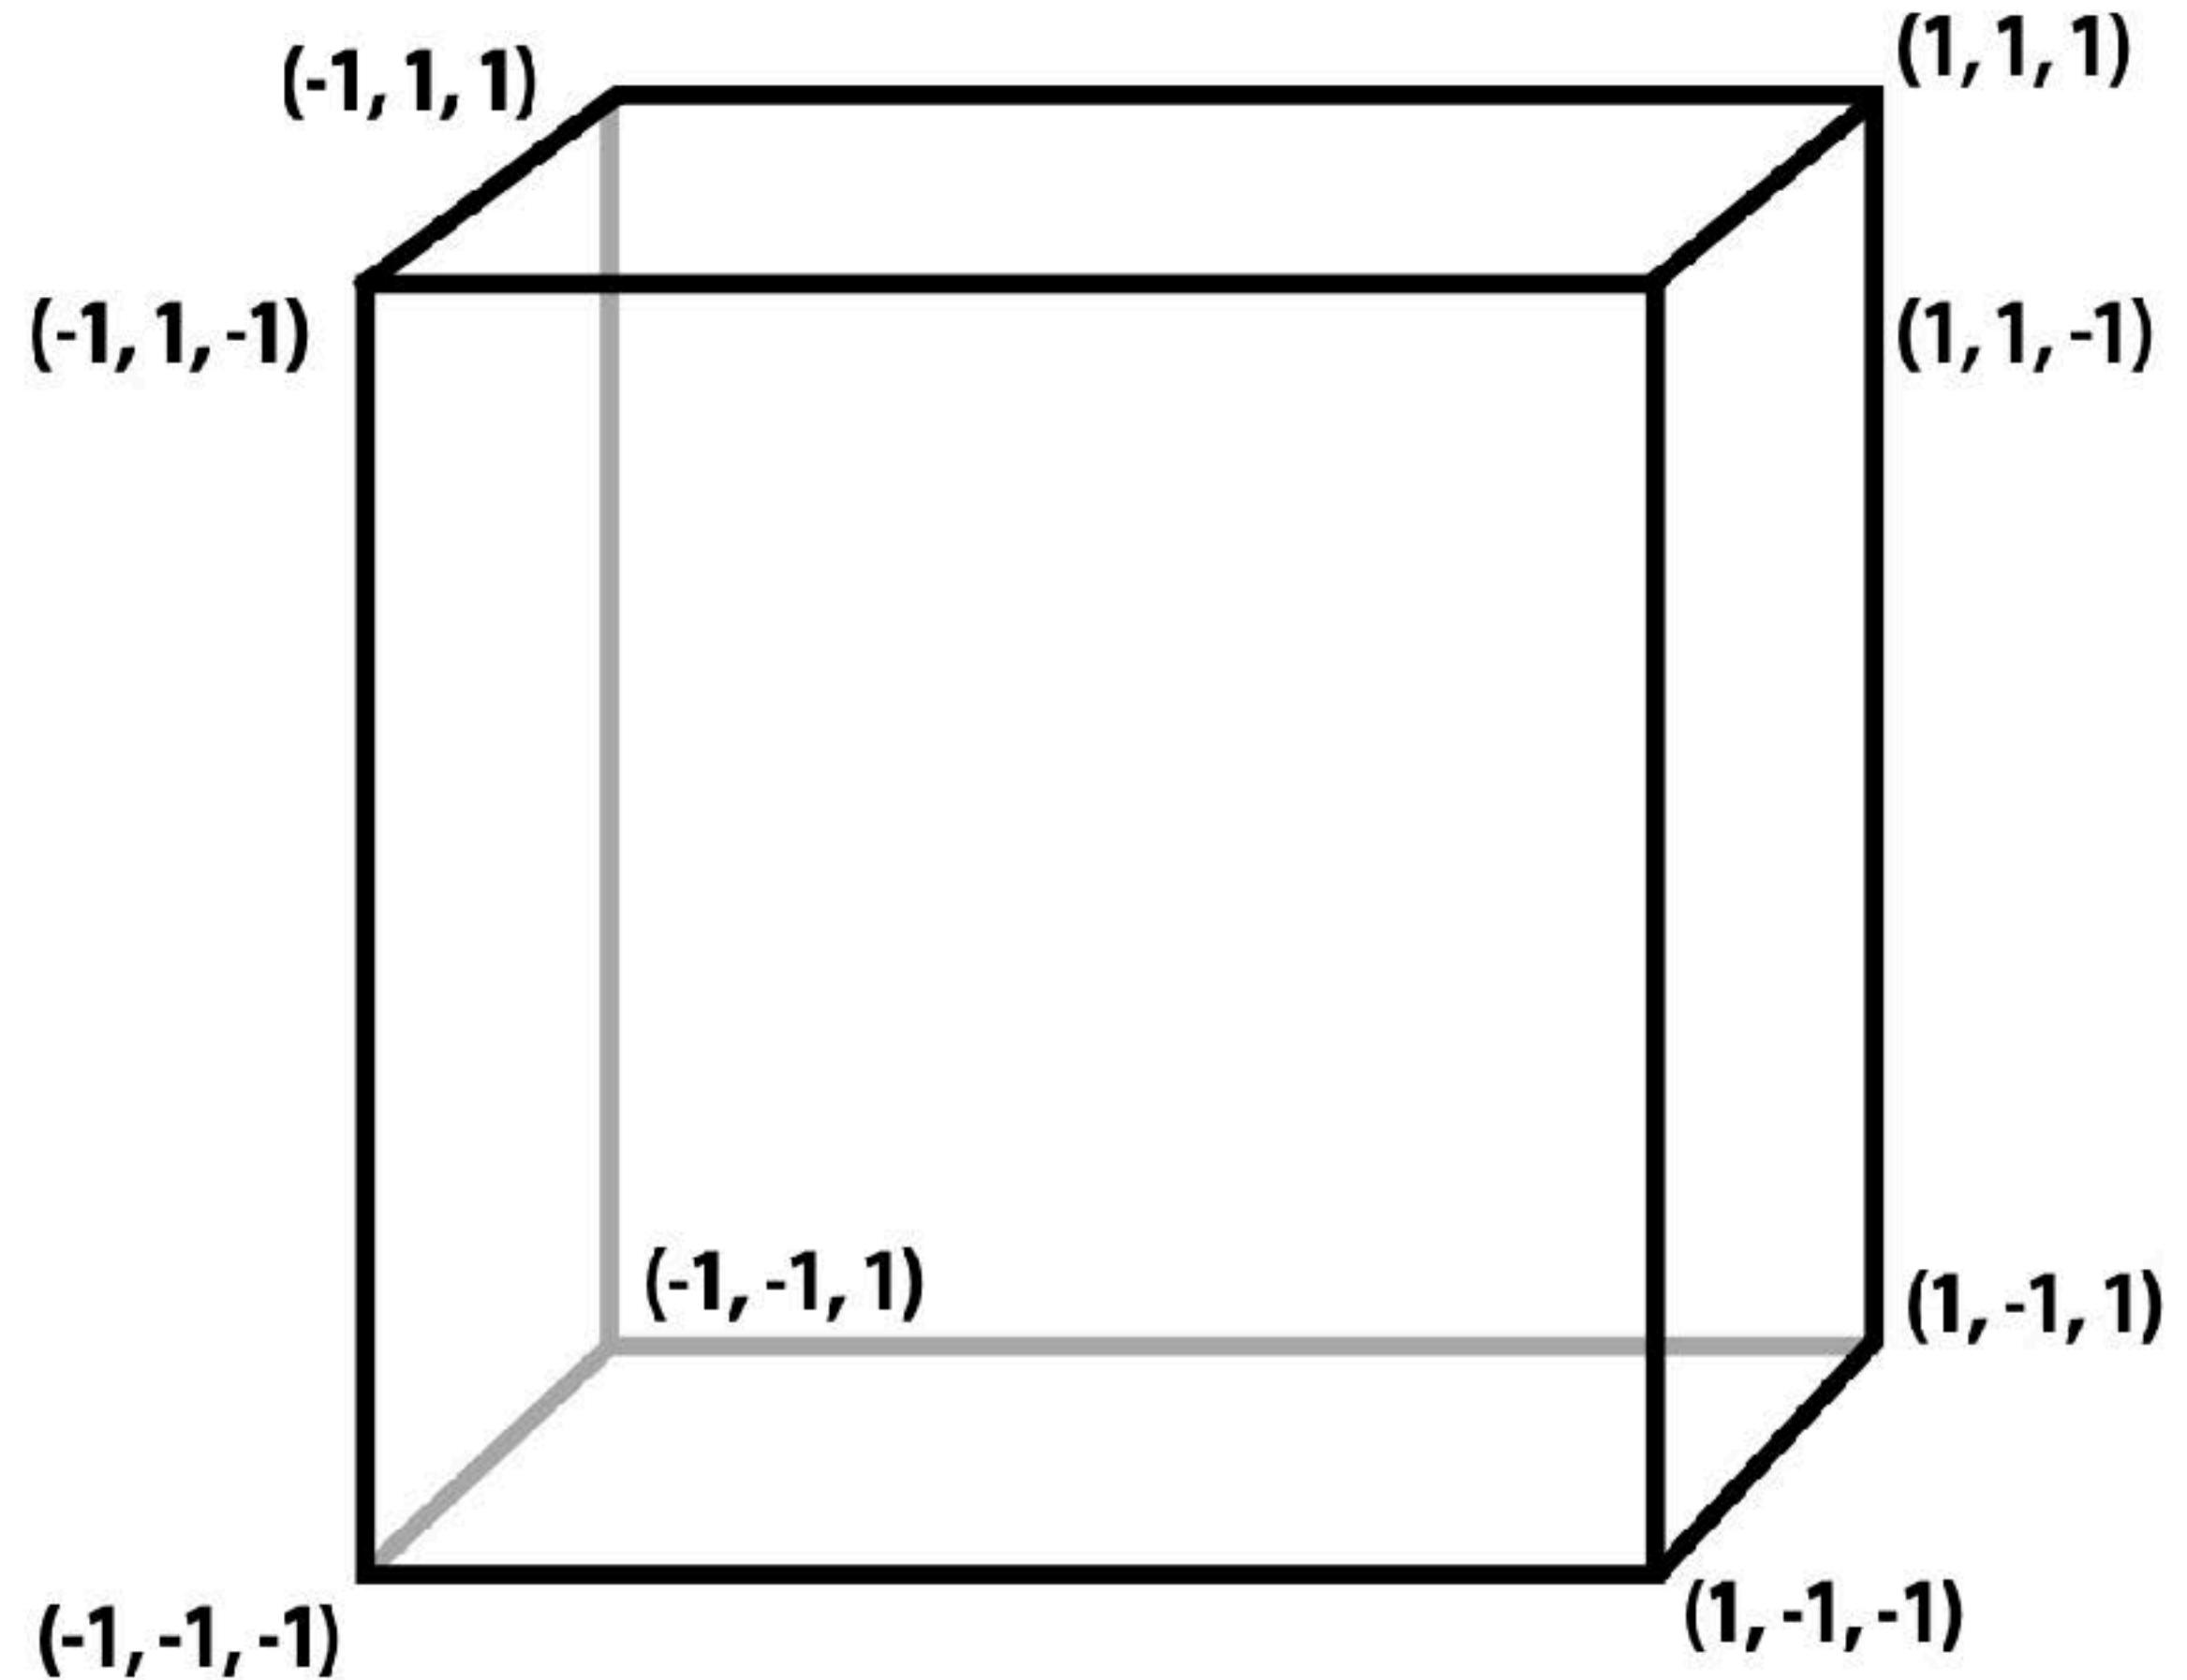


3D Spatial Relationships

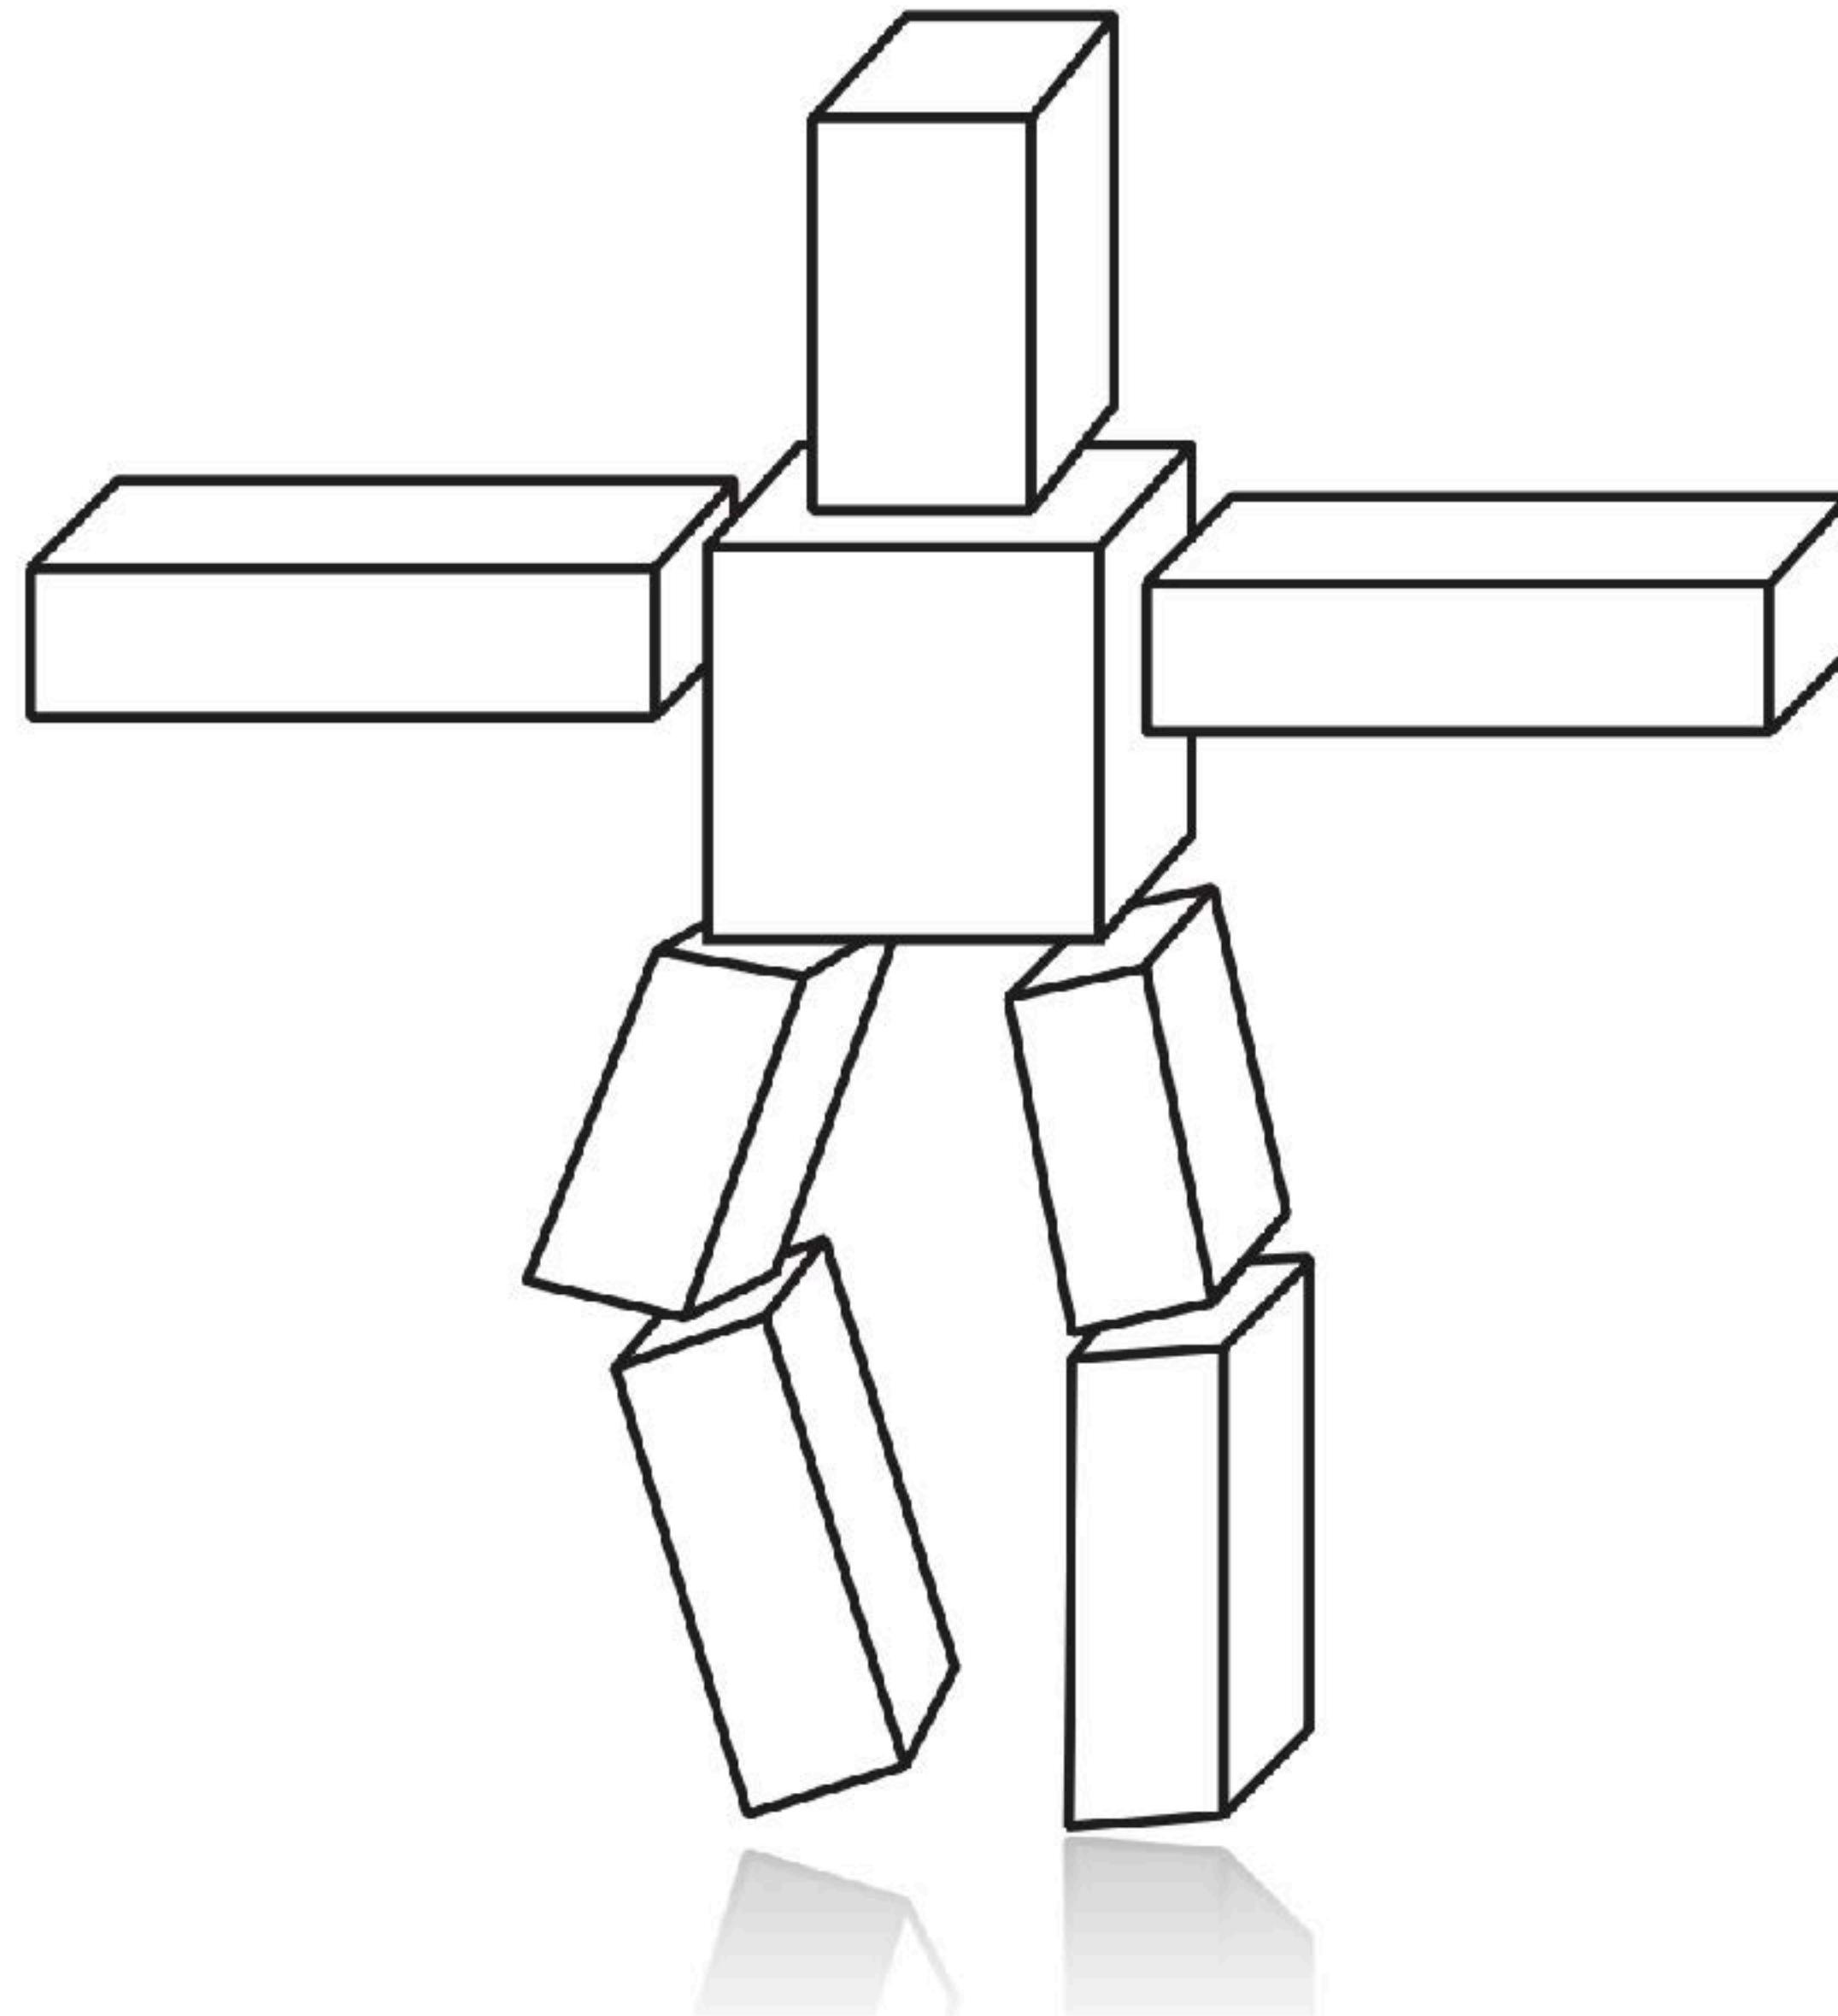
- How to represent the relationships between objects?



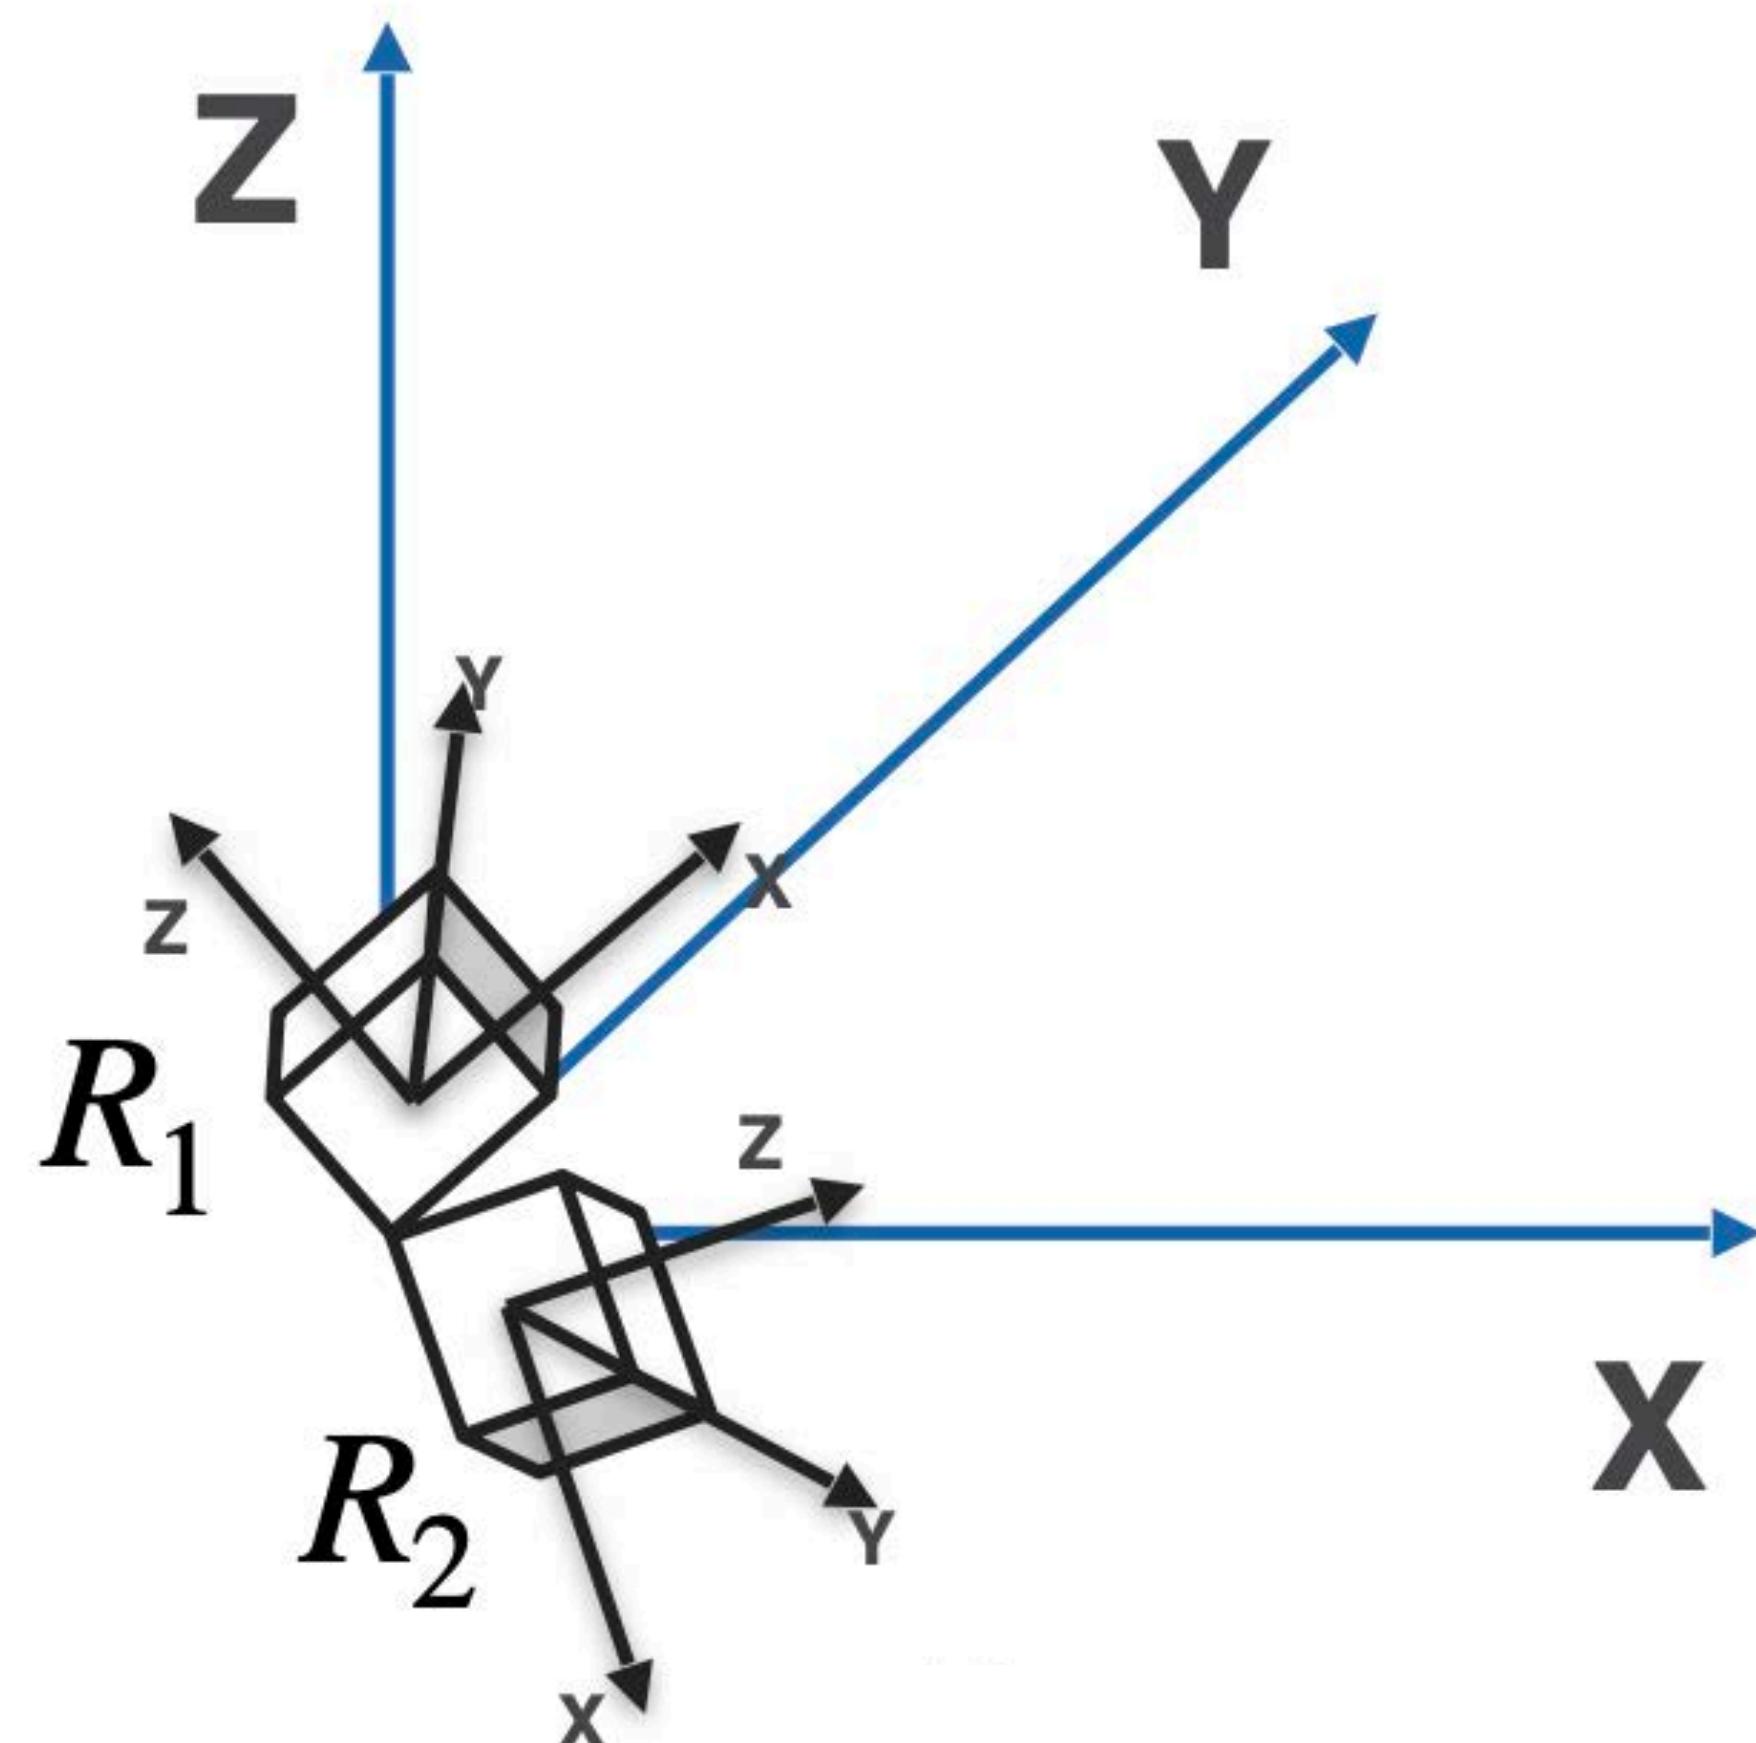
Cube



Drawing a Cube Person



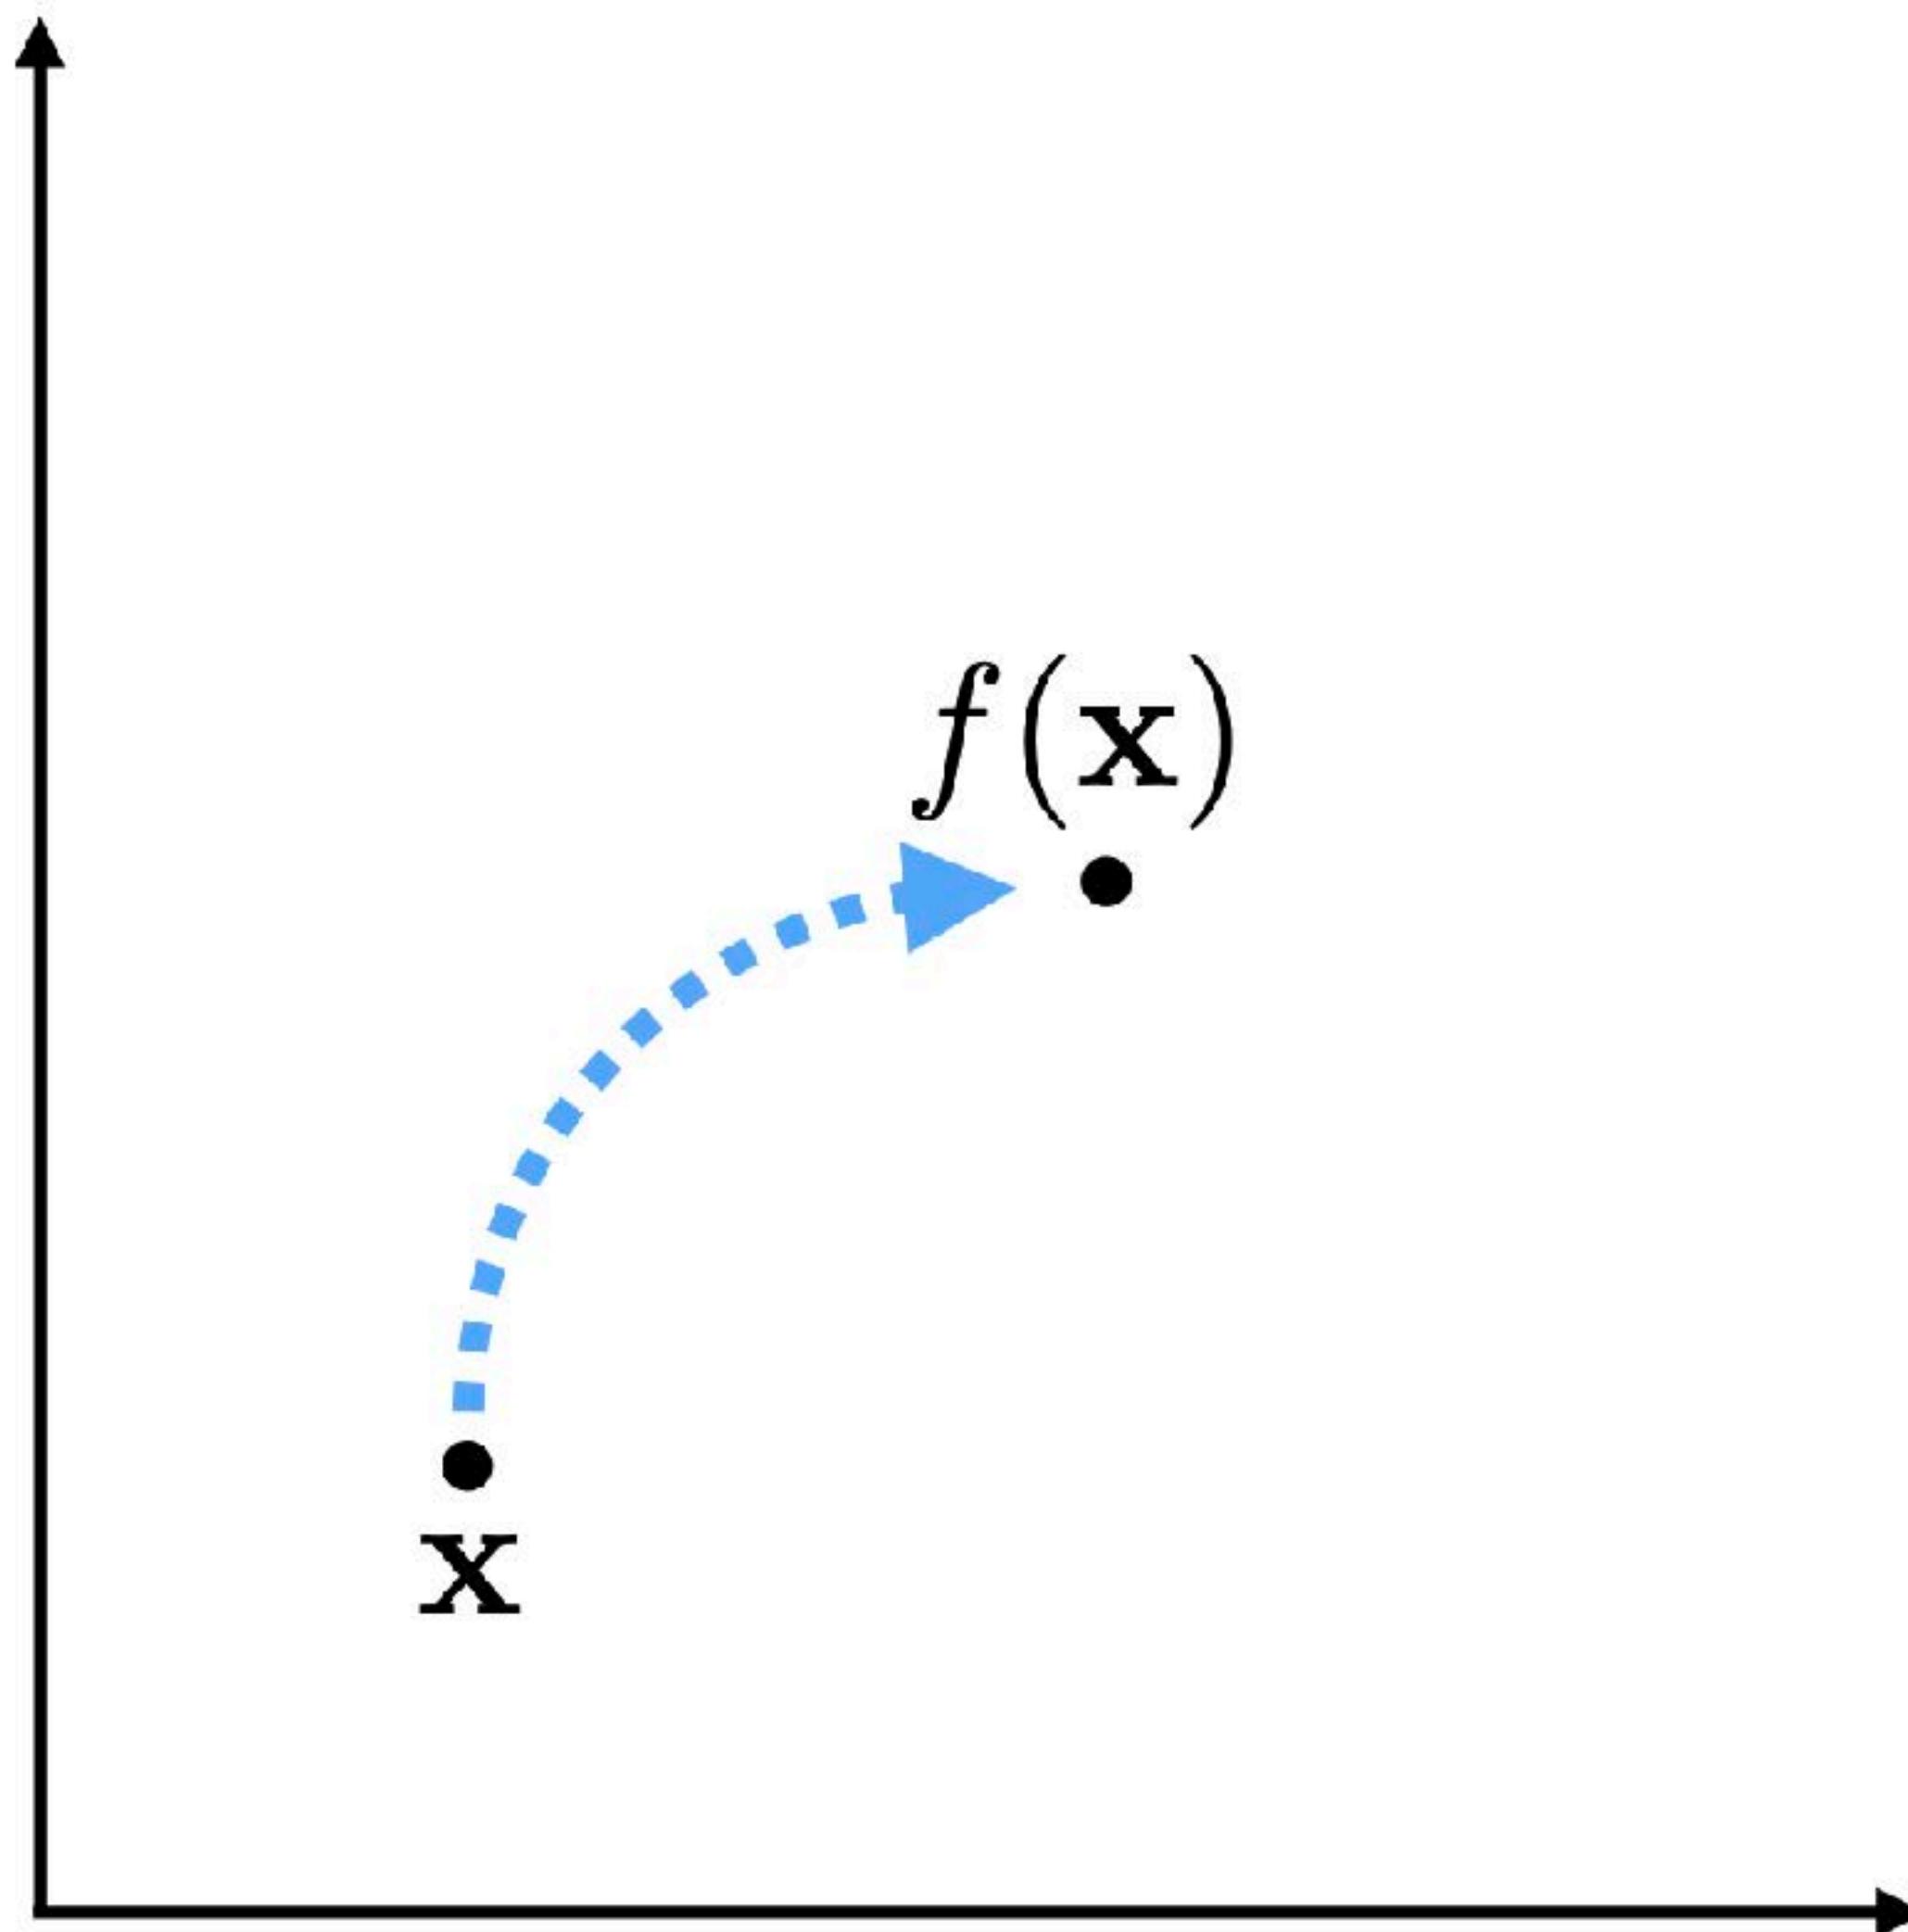
Today's focus: 3D transformations



Outline

- *Transformations and homogeneous coordinates*
- Rotation and $\text{SO}(n)$
- 3D Rotation representation
 - Euler Angles
 - Axis-Angle
 - Quaternion

Basic Idea: f Transforms x to $f(x)$



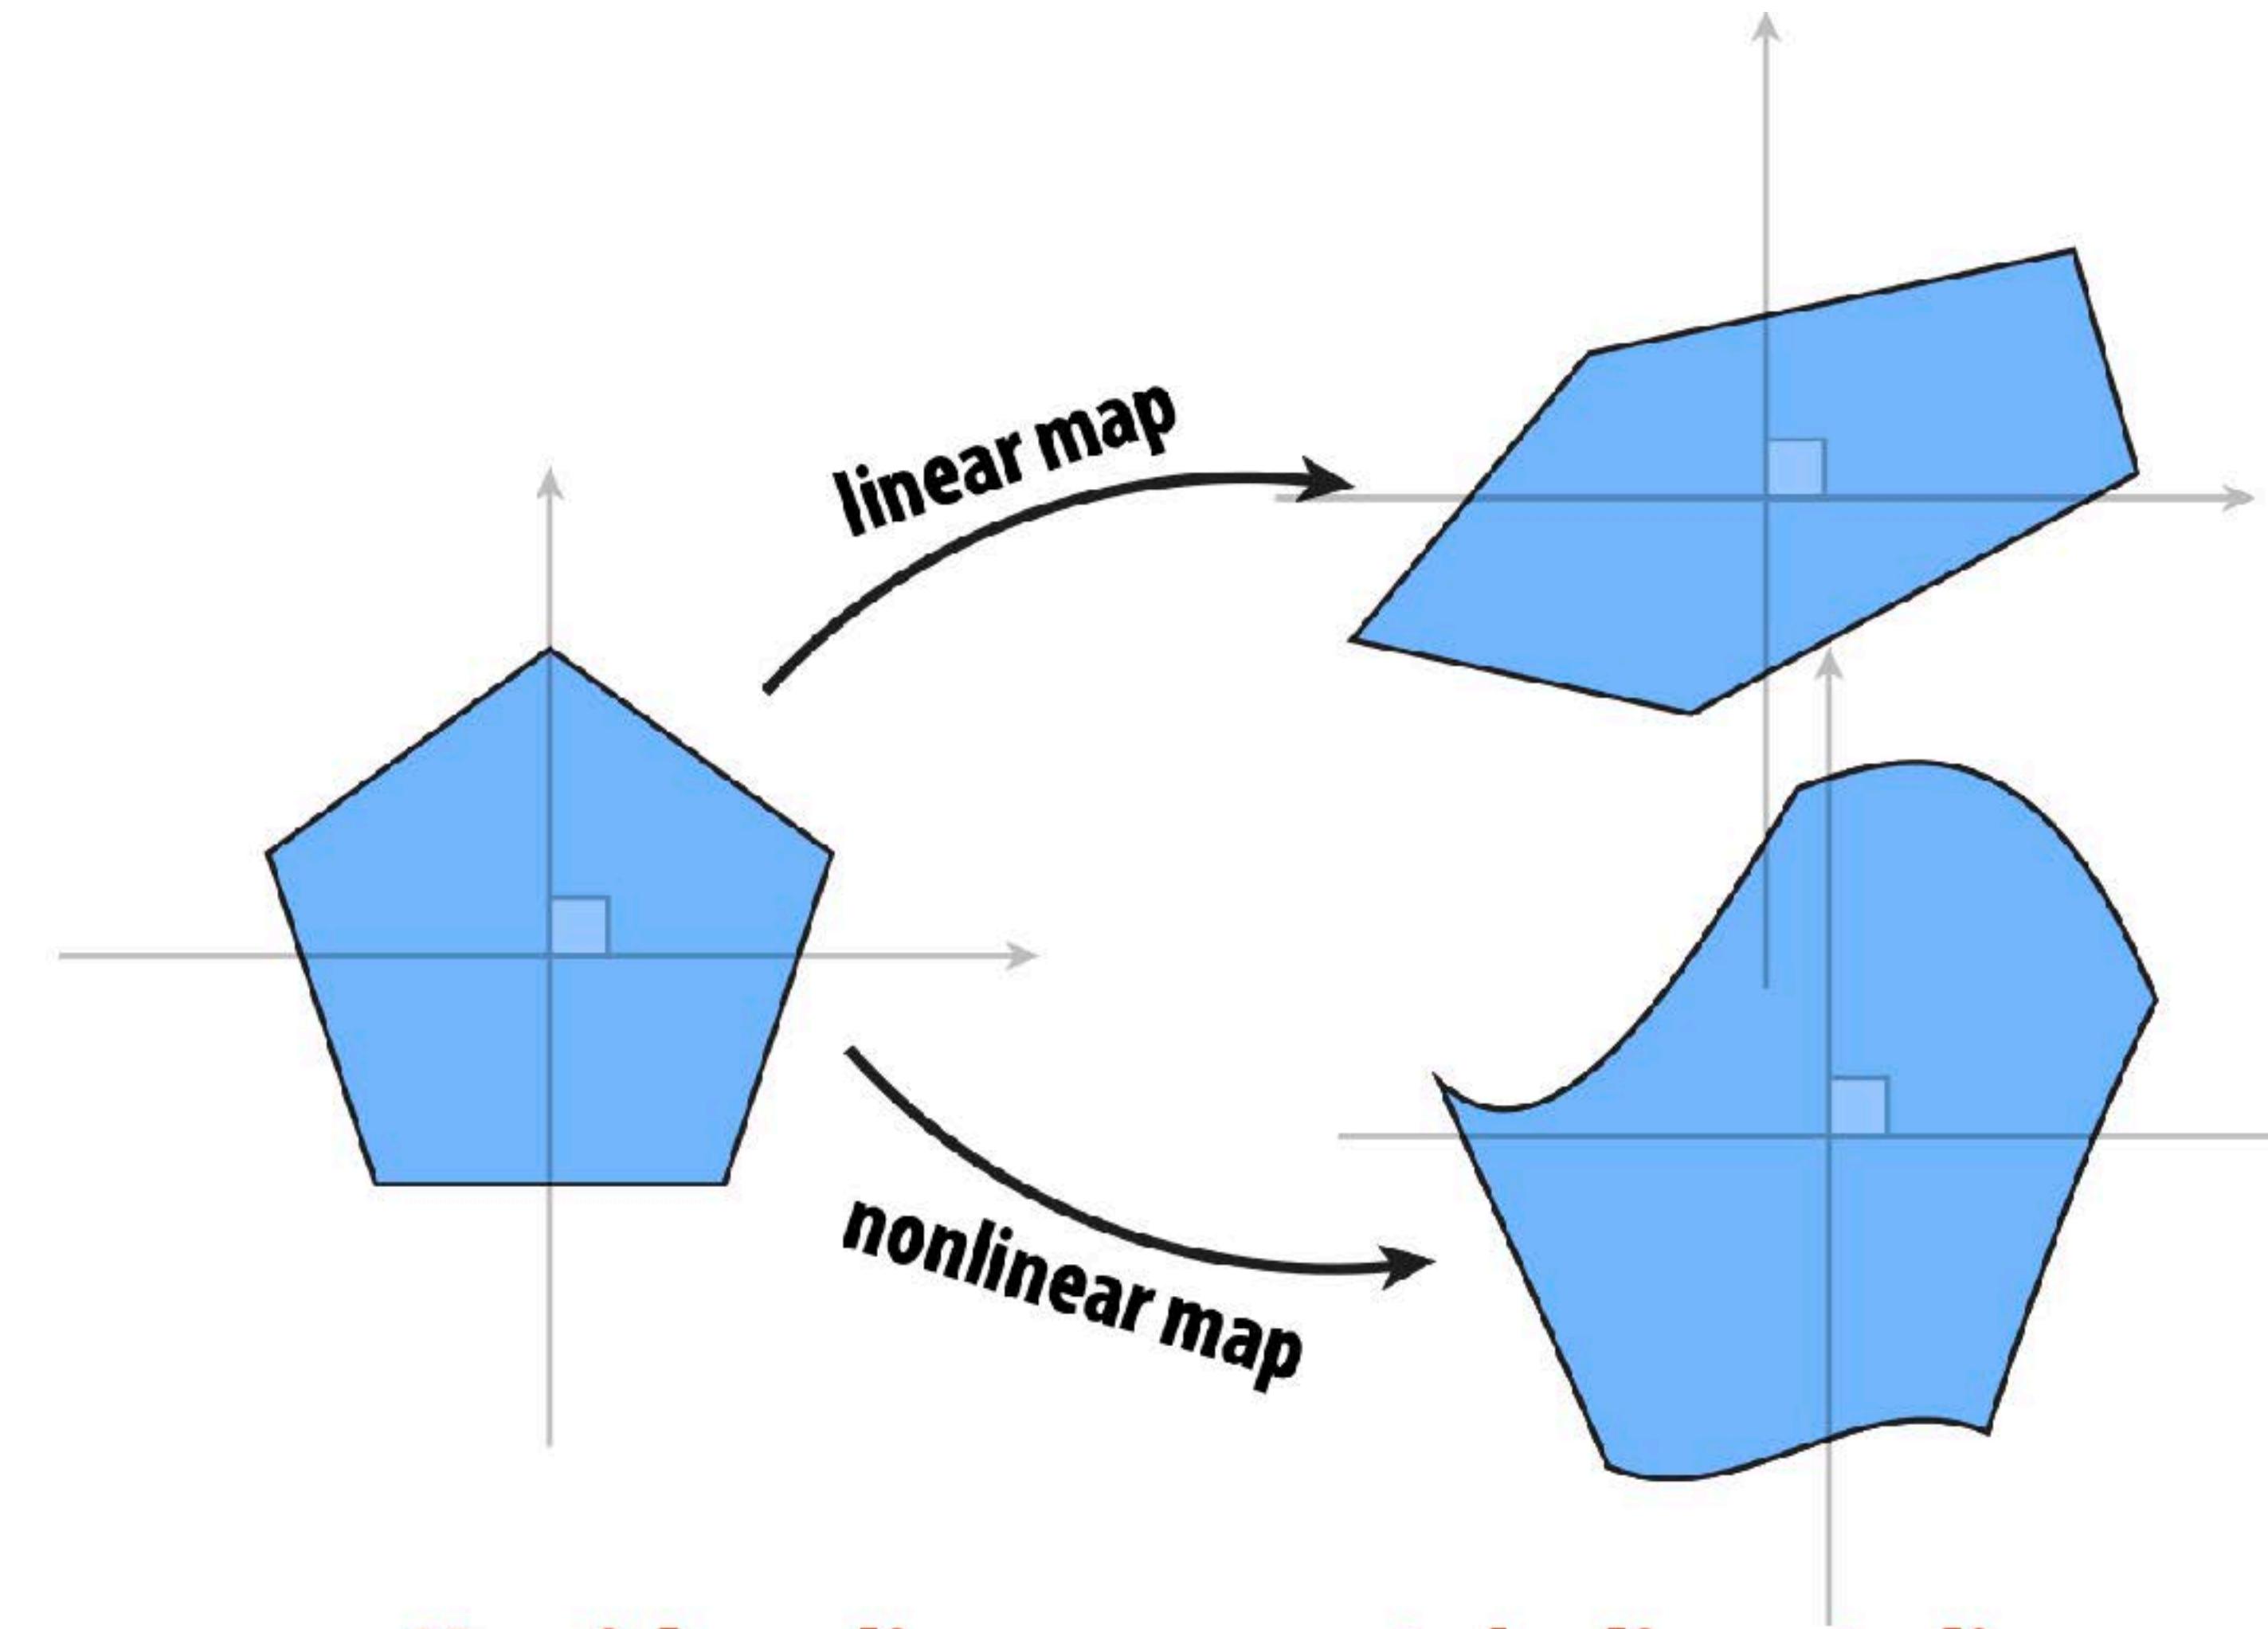
Linear Transformation

- Cheap to compute
- Composition of linear transformations is linear
 - Leads to uniform representation of transformations

$$f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$$

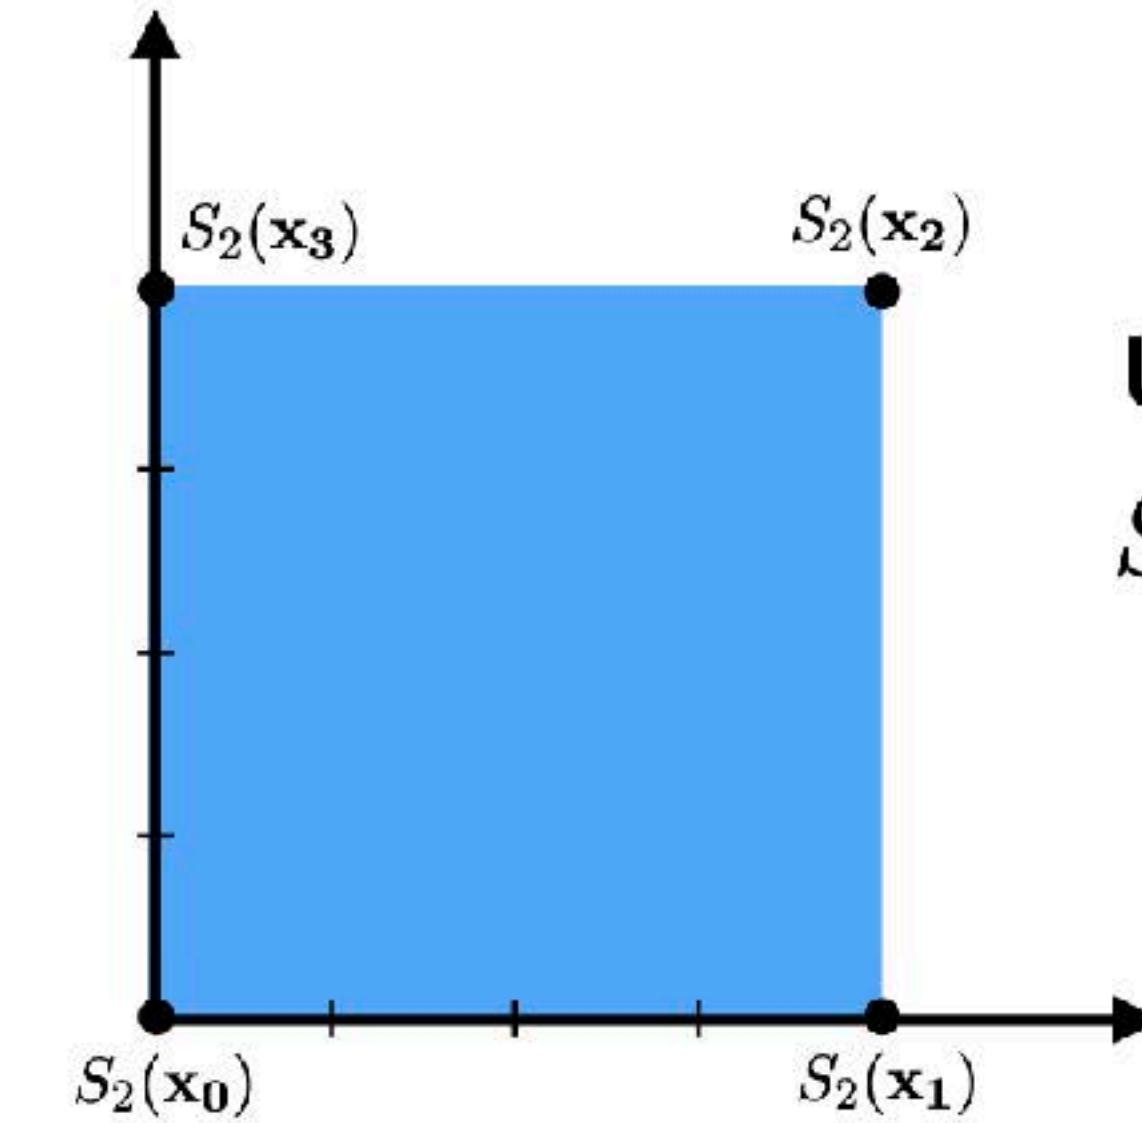
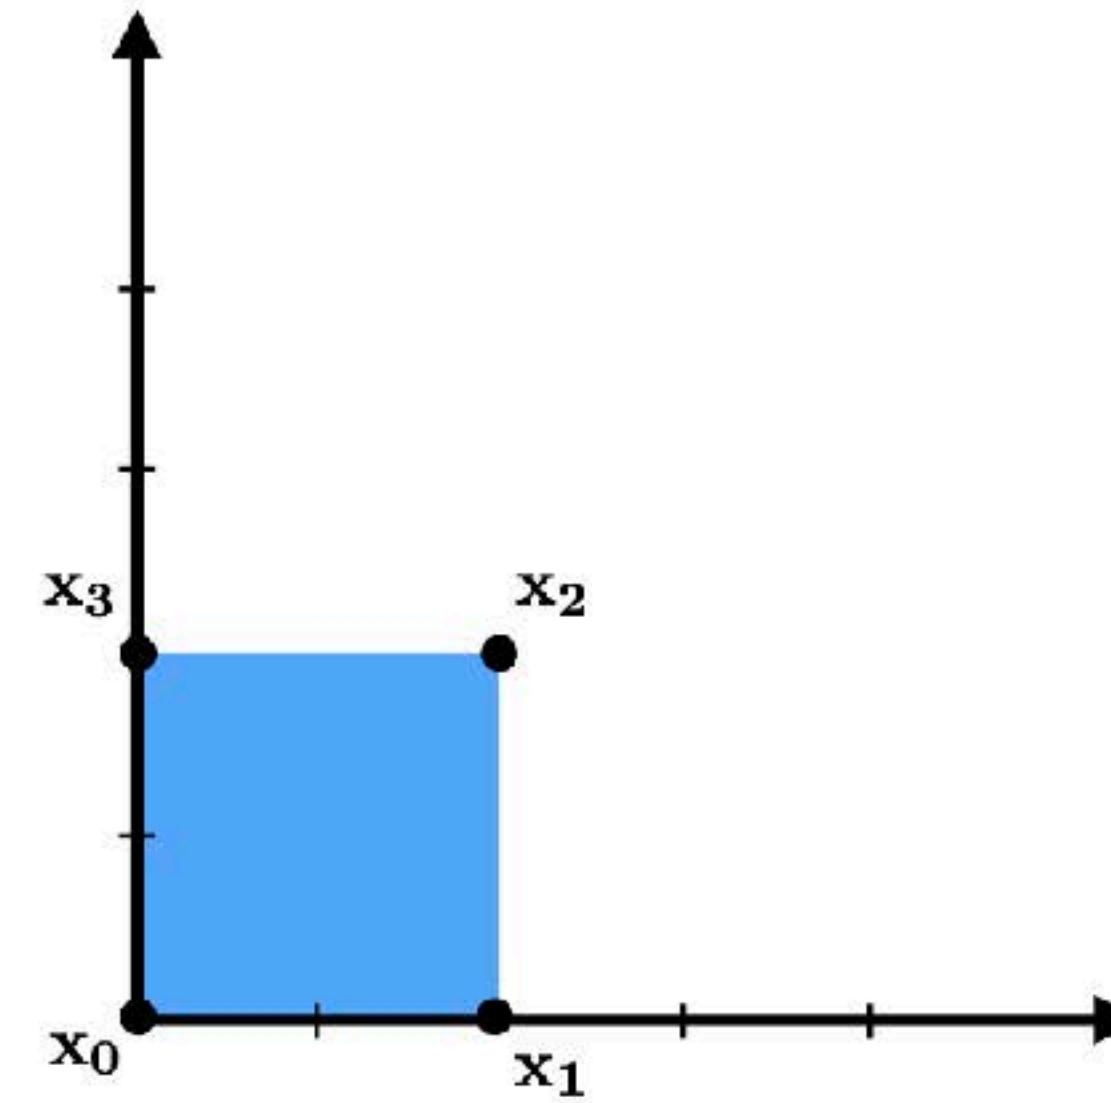
$$f(a\mathbf{x}) = af(\mathbf{x})$$

Linear Transformation Visualization



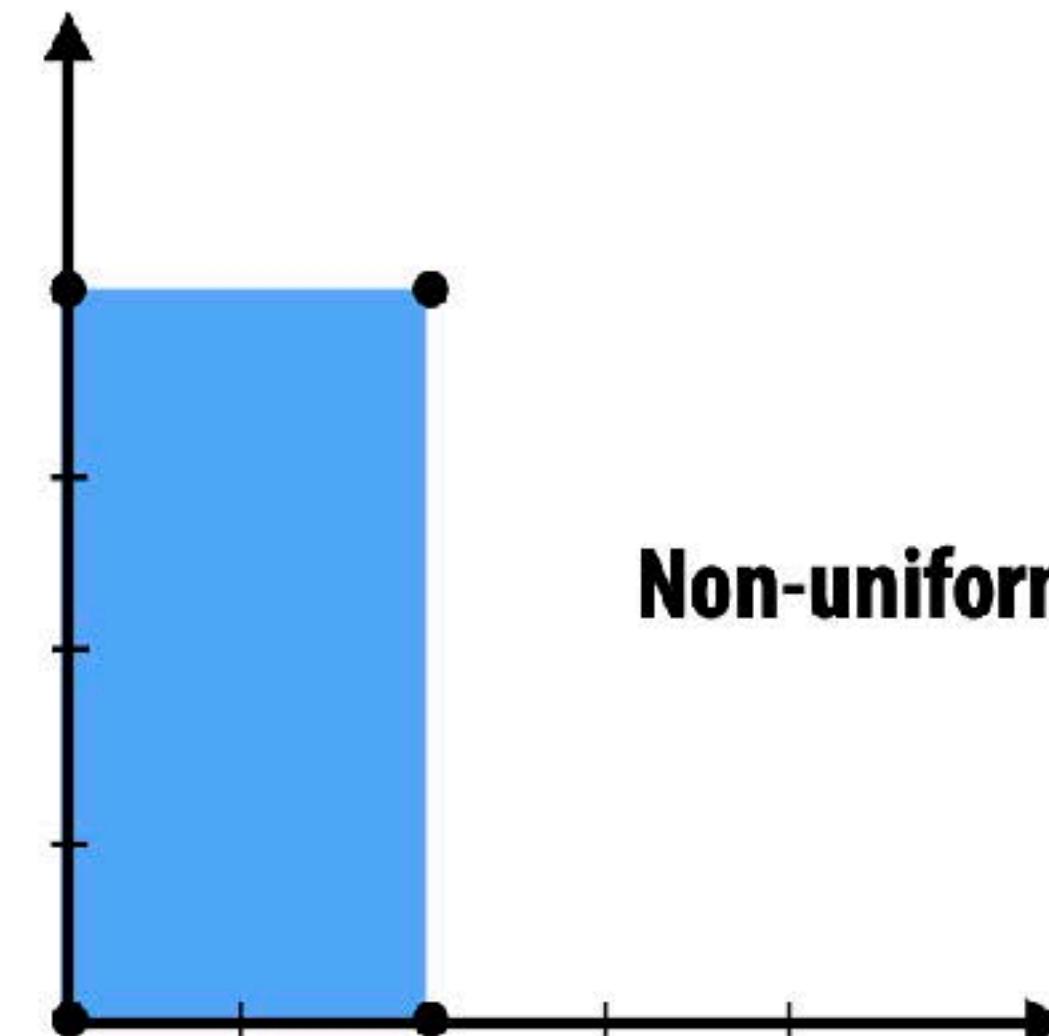
Key idea: *linear maps take lines to lines*

Scale



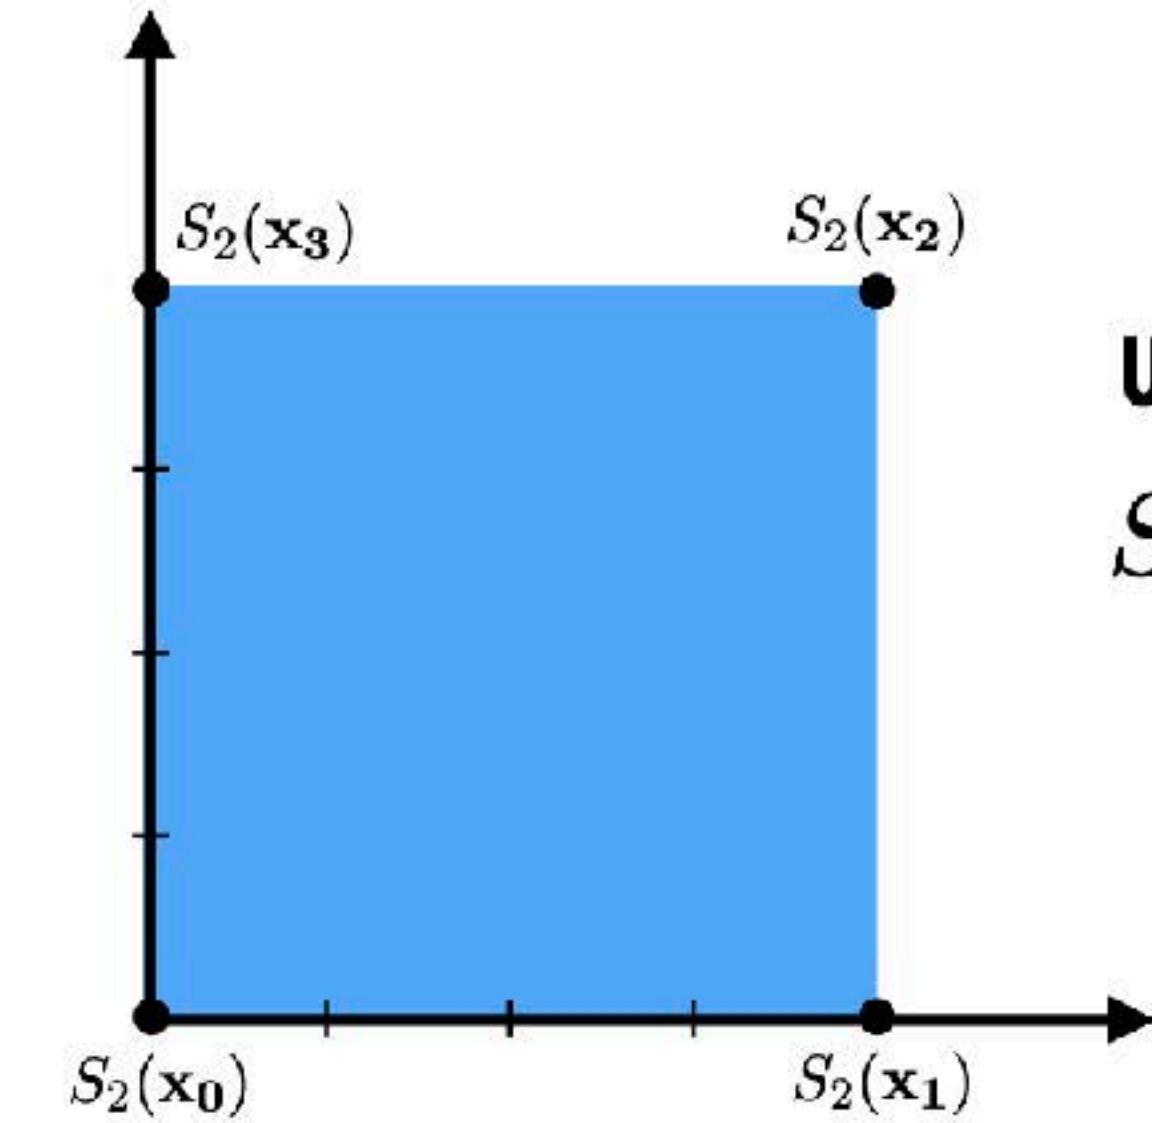
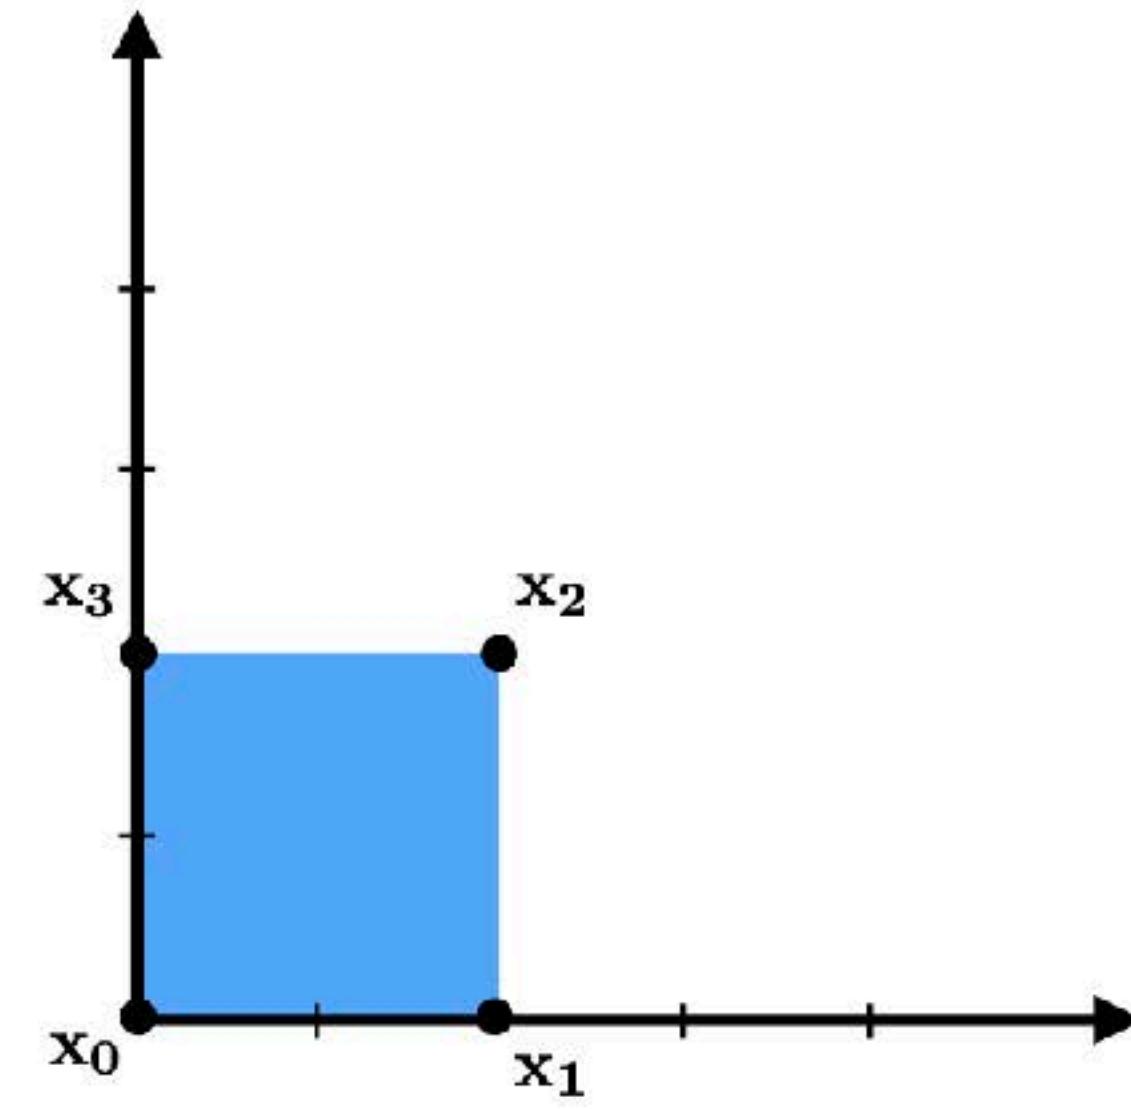
Uniform scale:

$$S_a(\mathbf{x}) = a\mathbf{x}$$



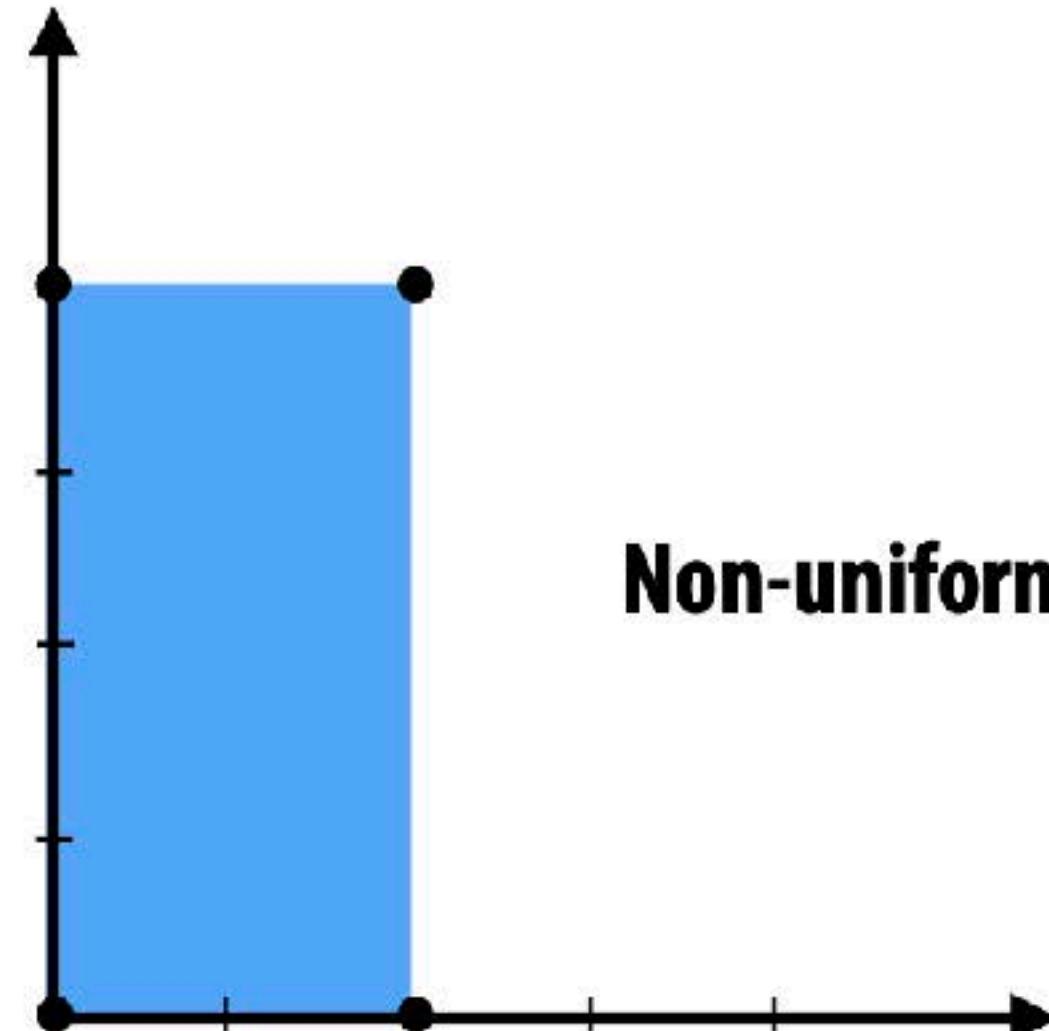
Non-uniform scale??

Scale



Uniform scale:

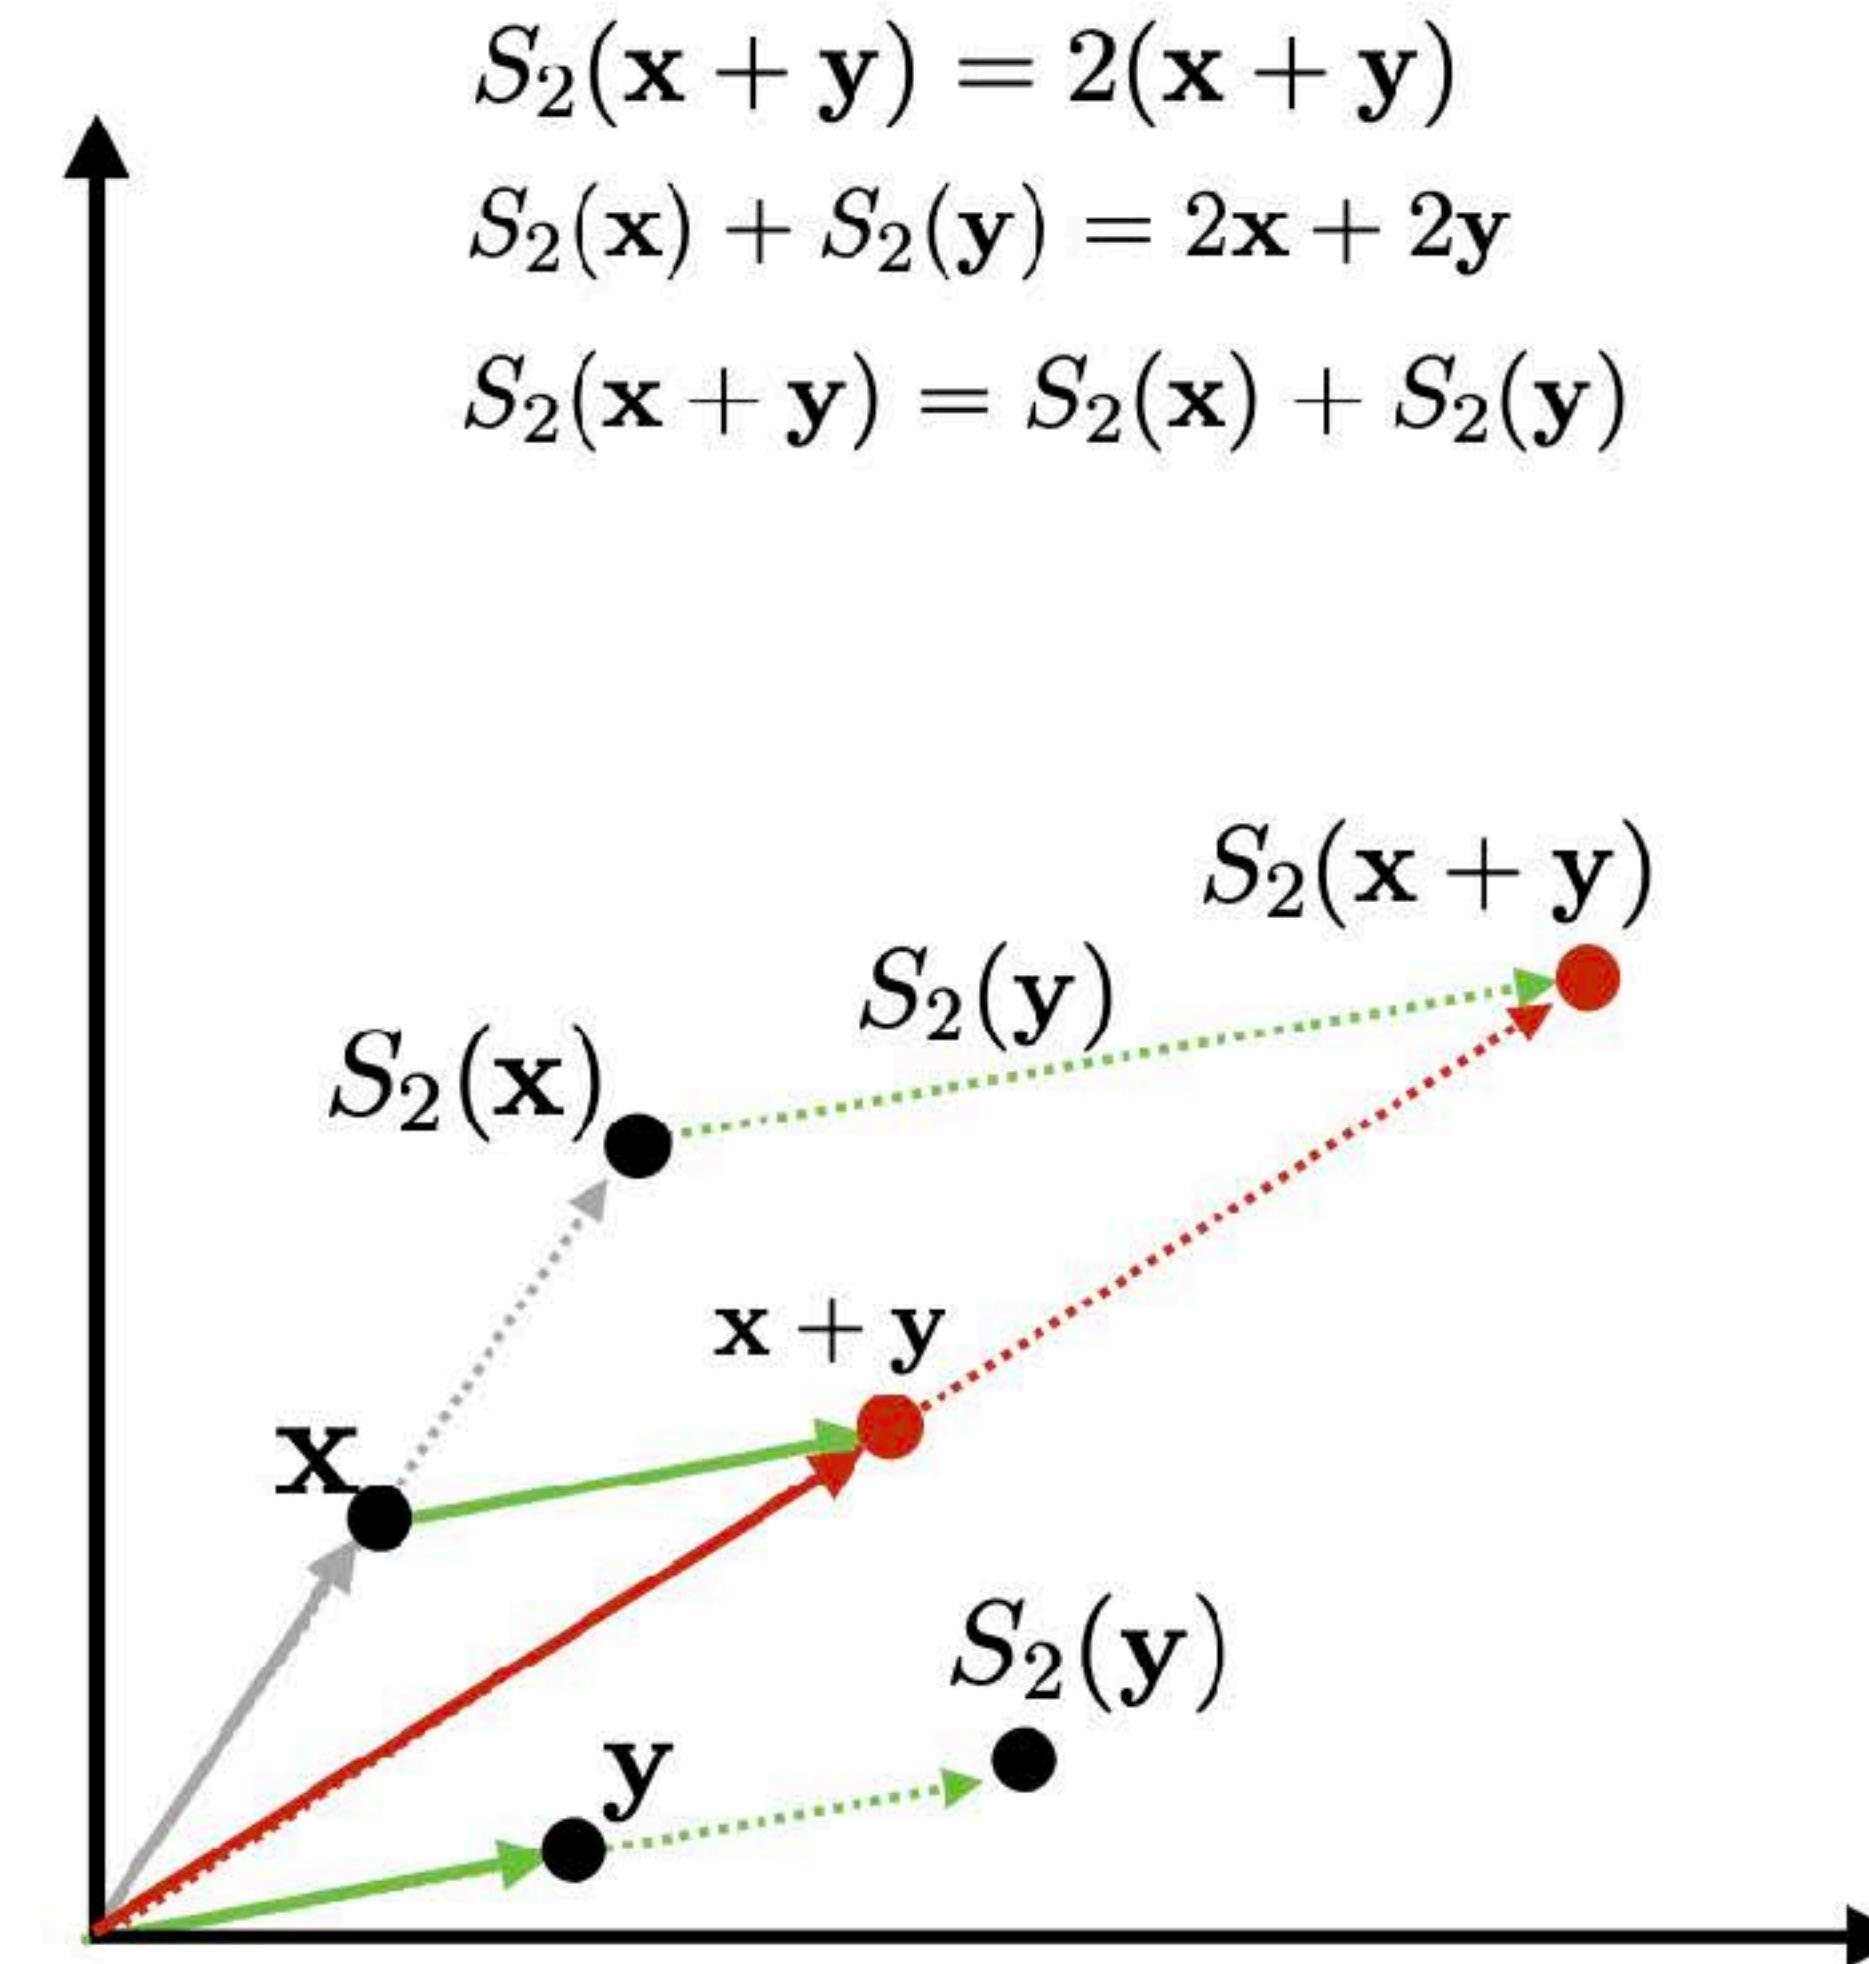
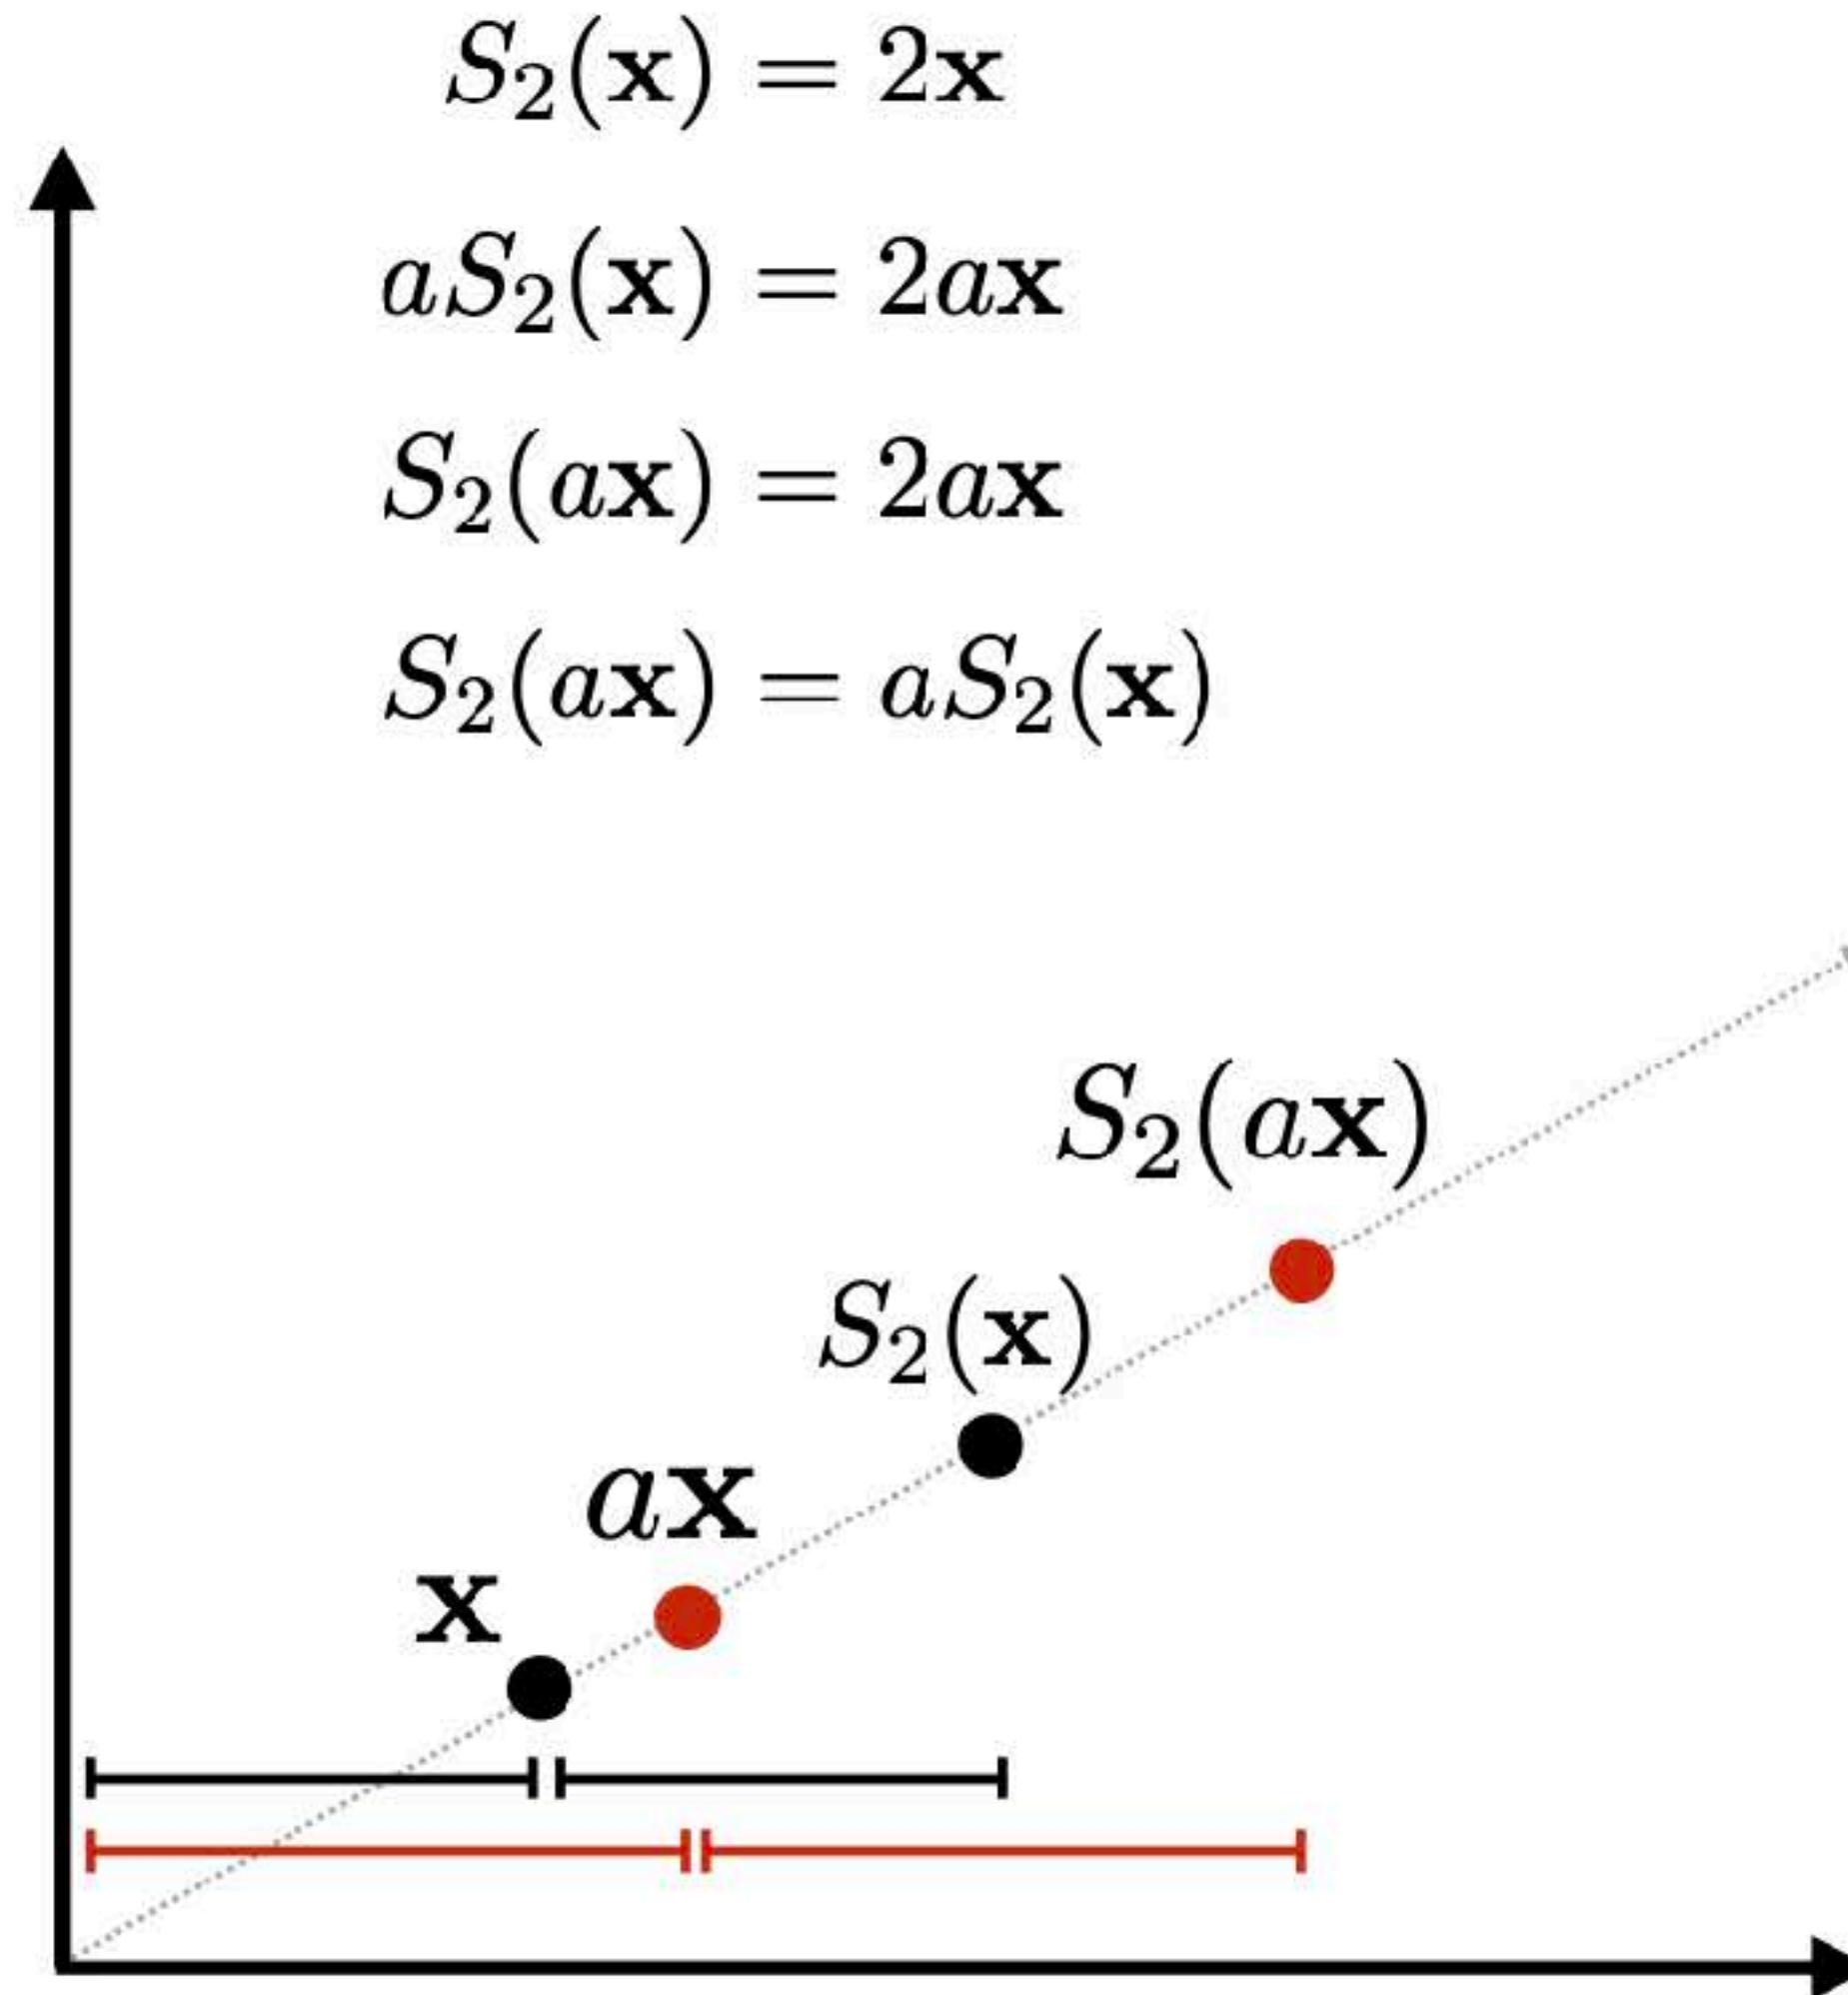
$$S_a(\mathbf{x}) = a\mathbf{x}$$



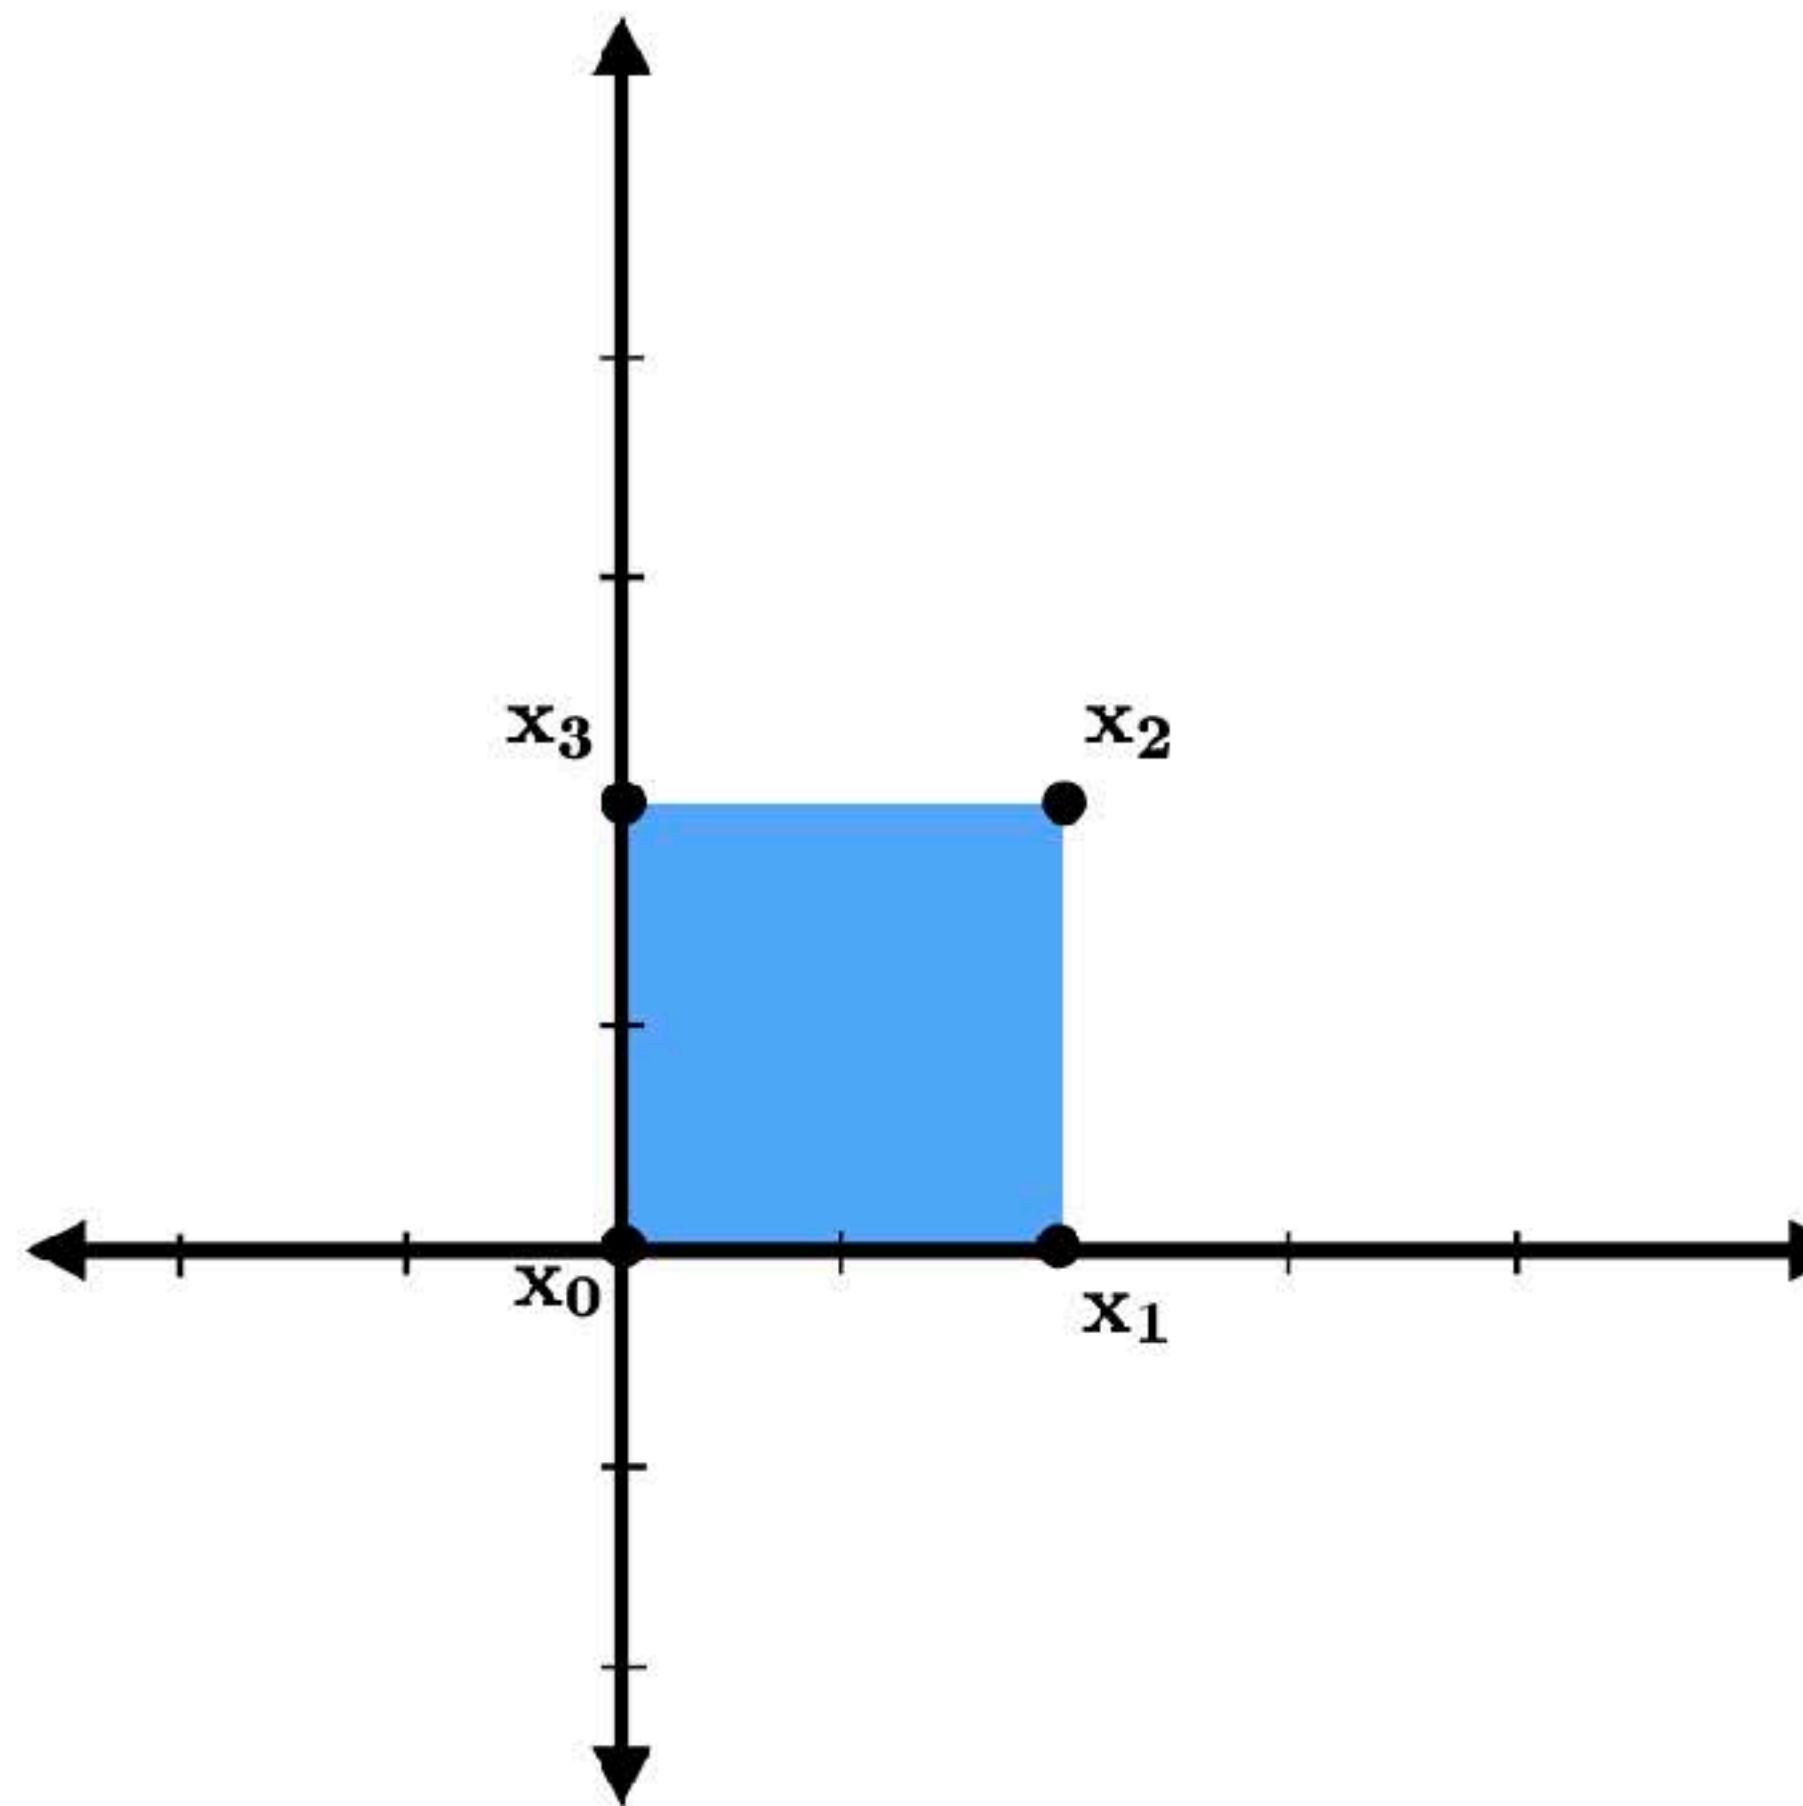
Non-uniform scale??

$$\mathbf{S}_s = \begin{bmatrix} \mathbf{s}_x & 0 \\ 0 & \mathbf{s}_y \end{bmatrix}$$

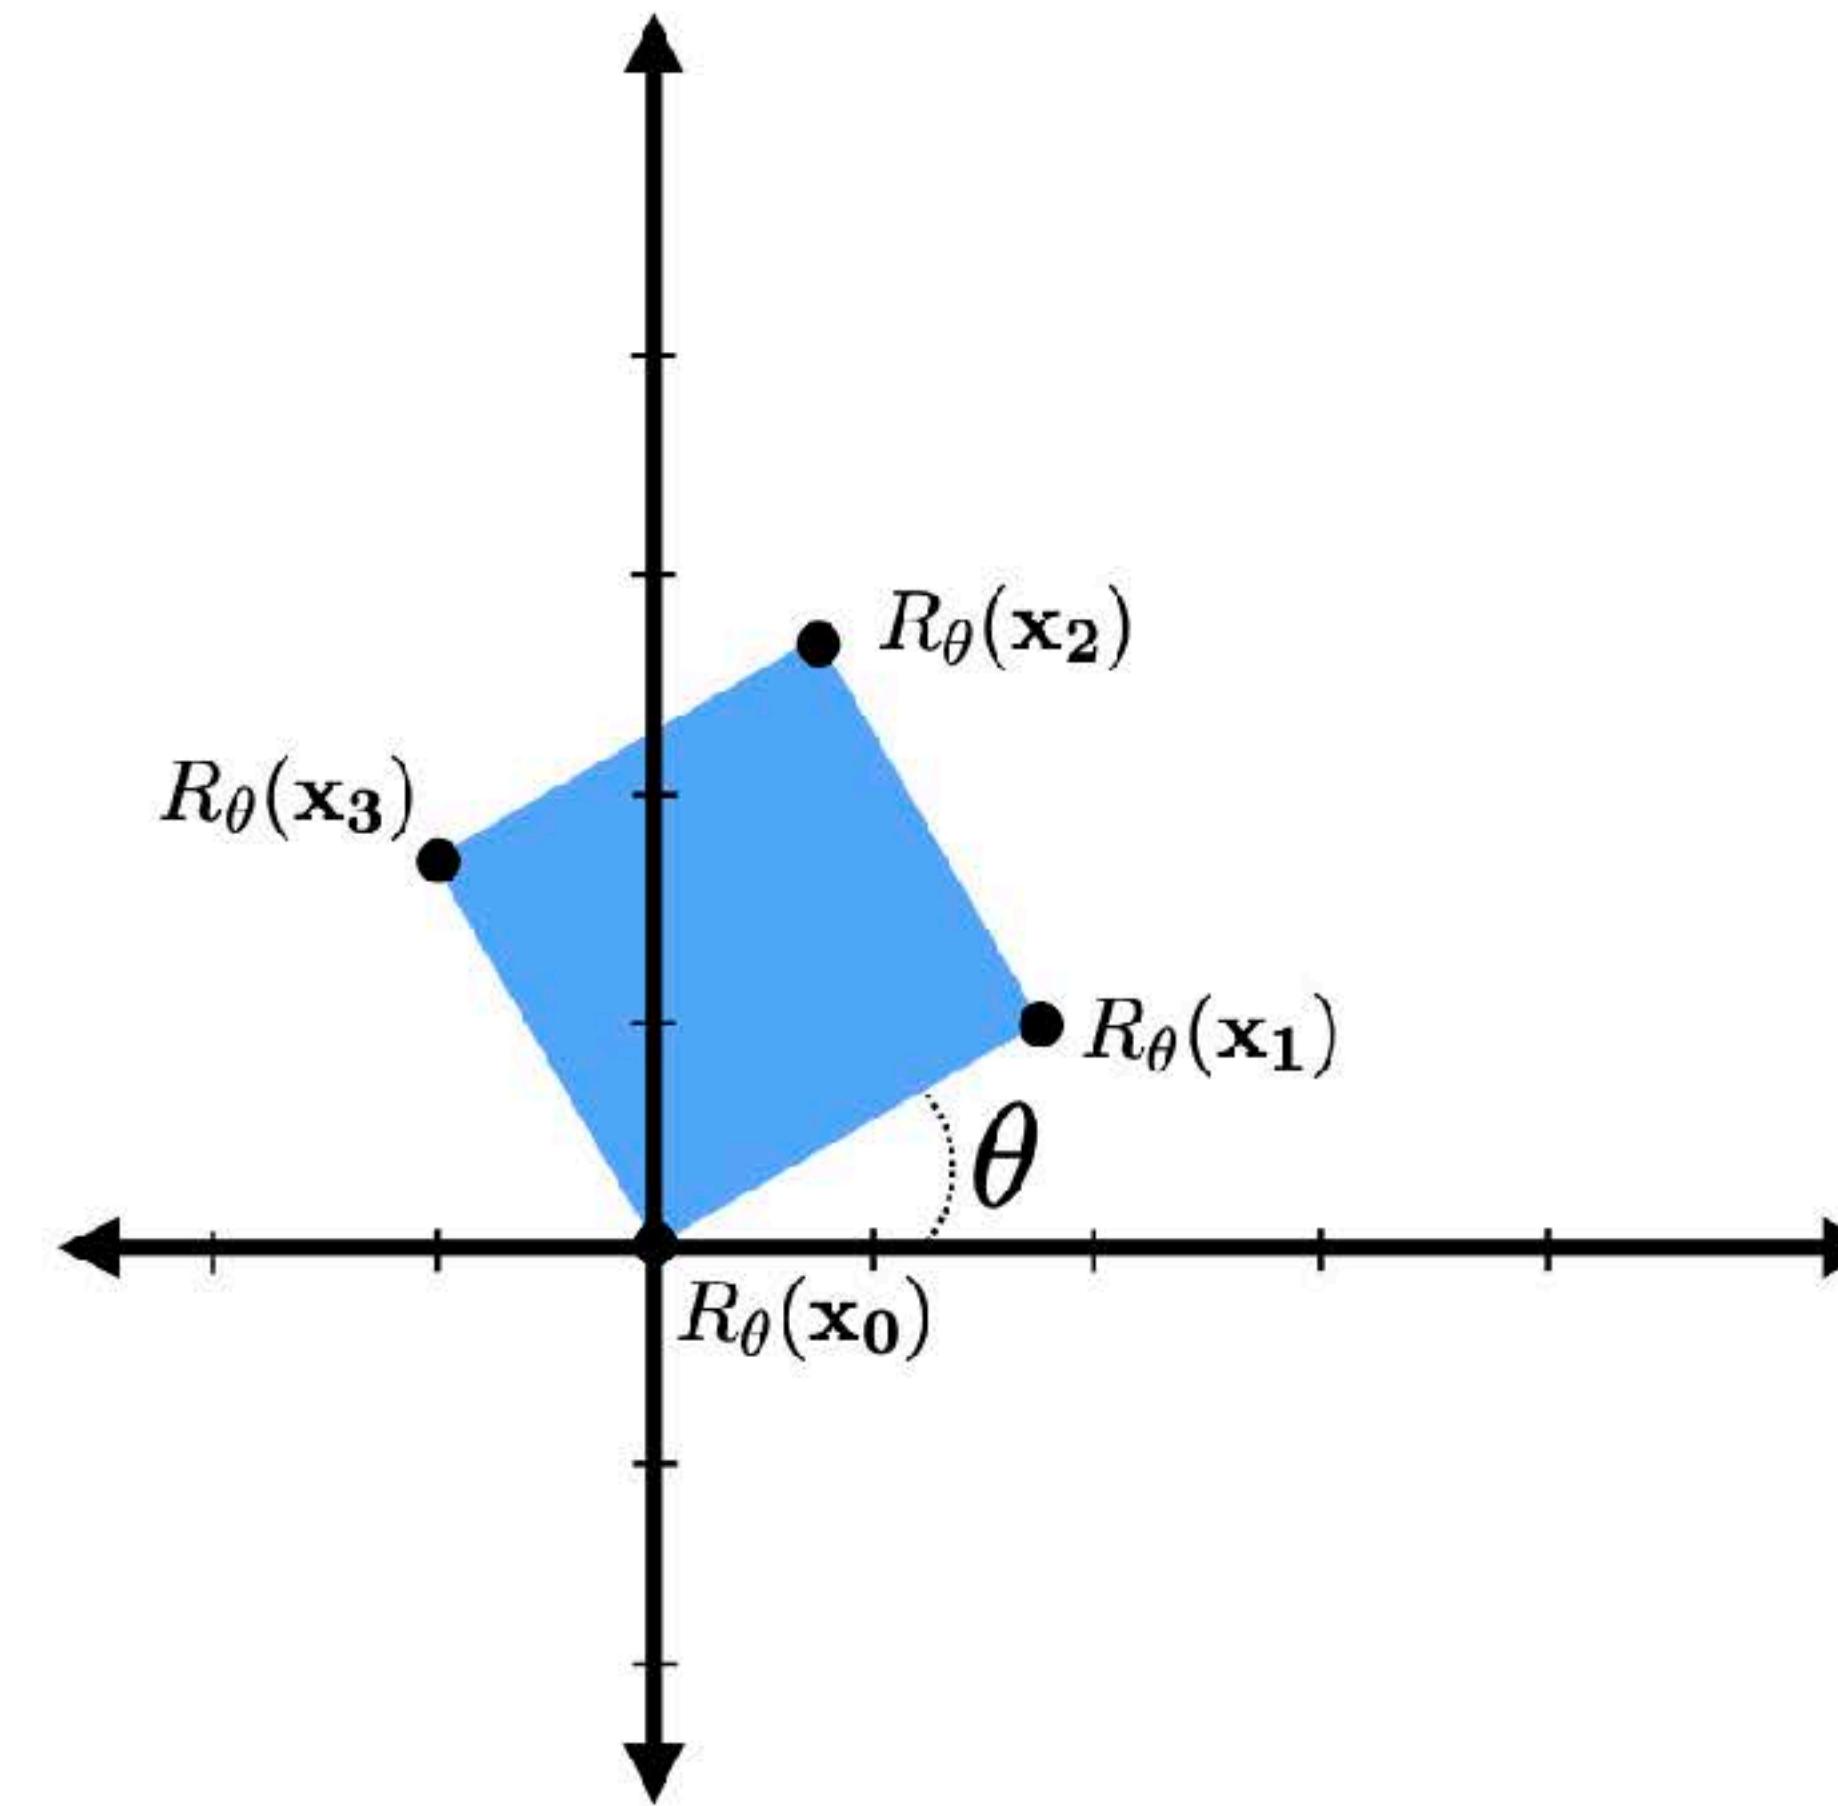
Linearity?



Rotation

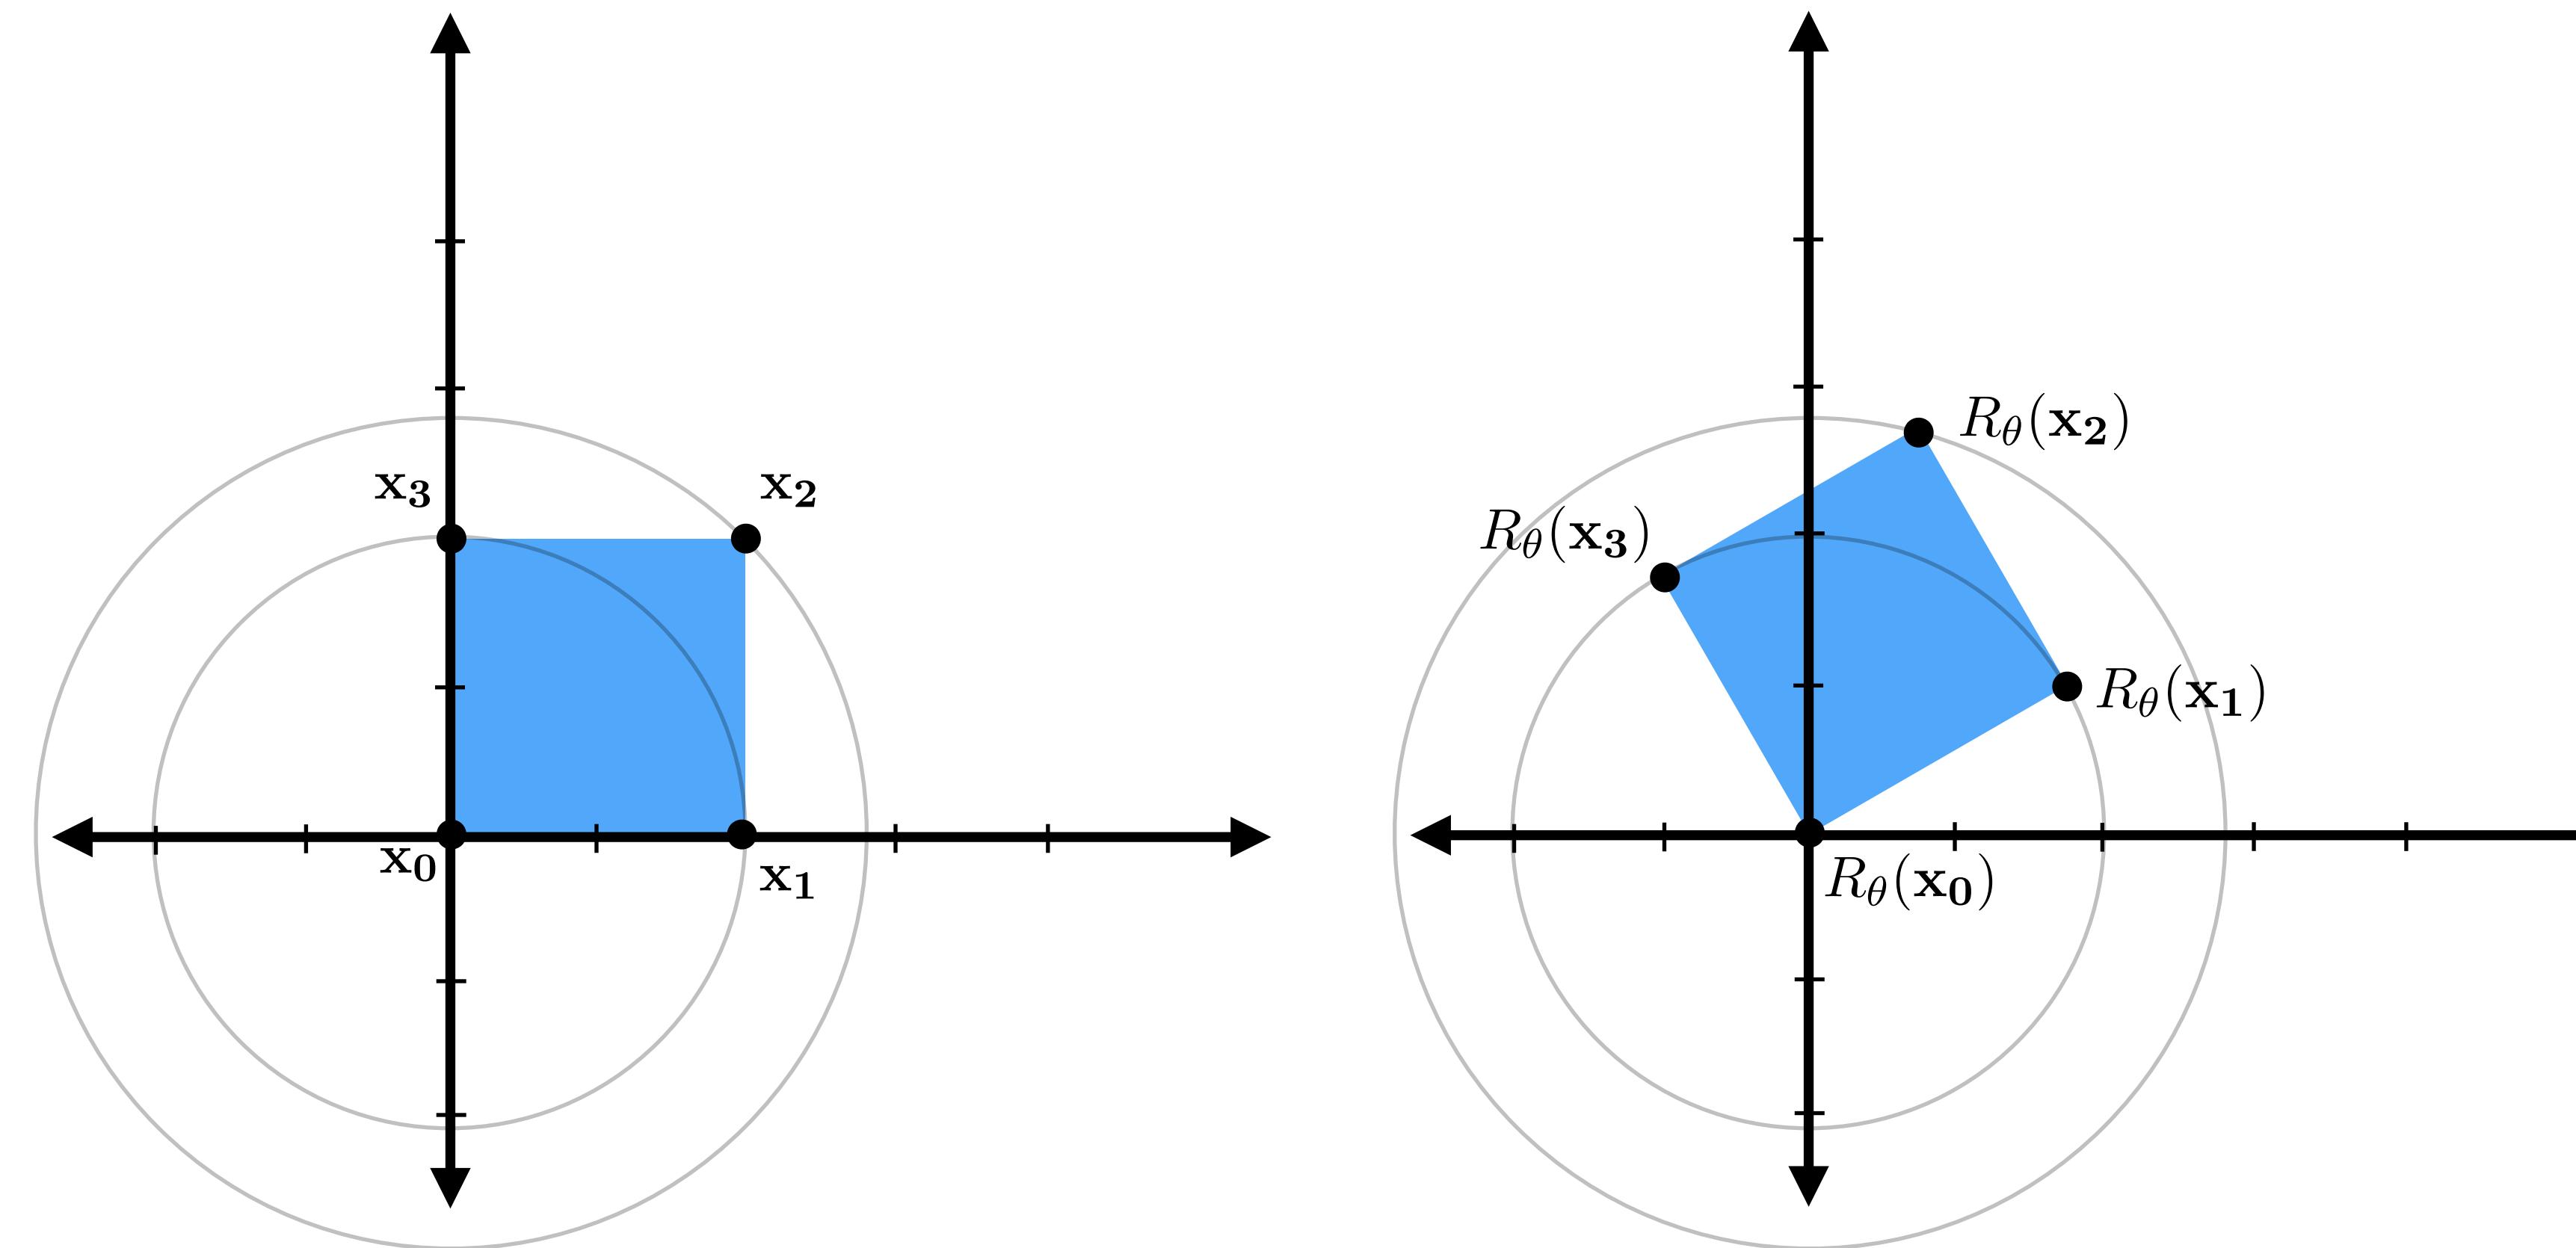


R_θ = rotate counter-clockwise by θ



$$\mathbf{R}_\theta = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

Rotation as circular motion

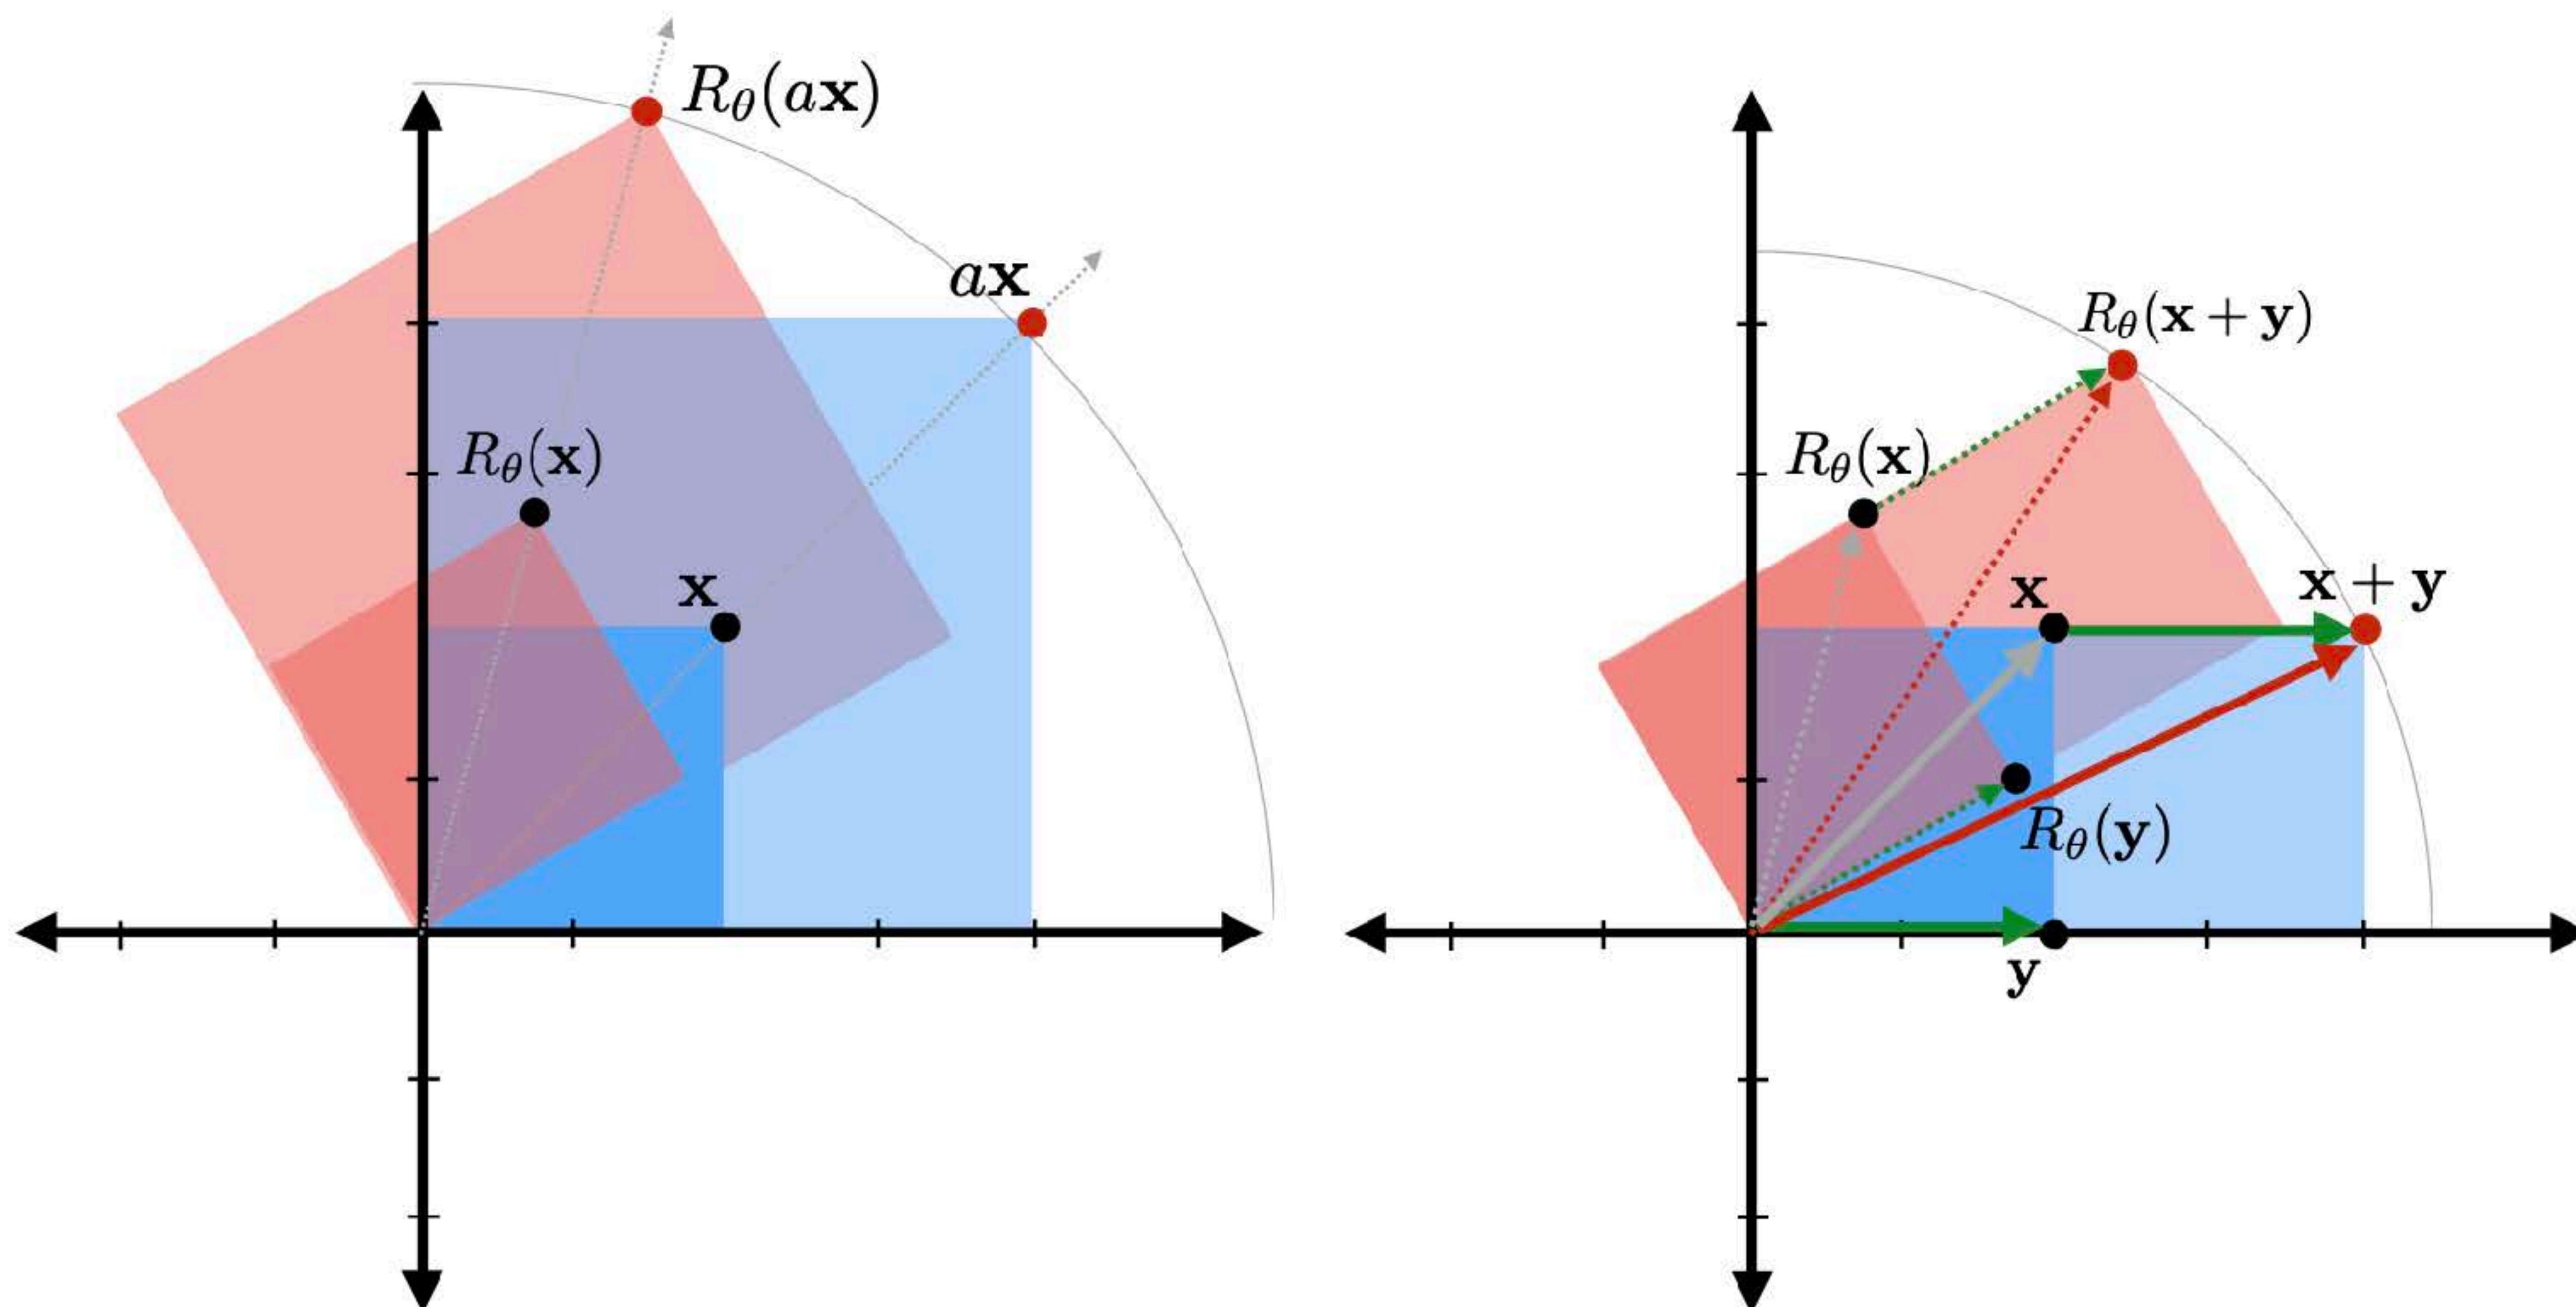


R_θ = rotate counter-clockwise by θ

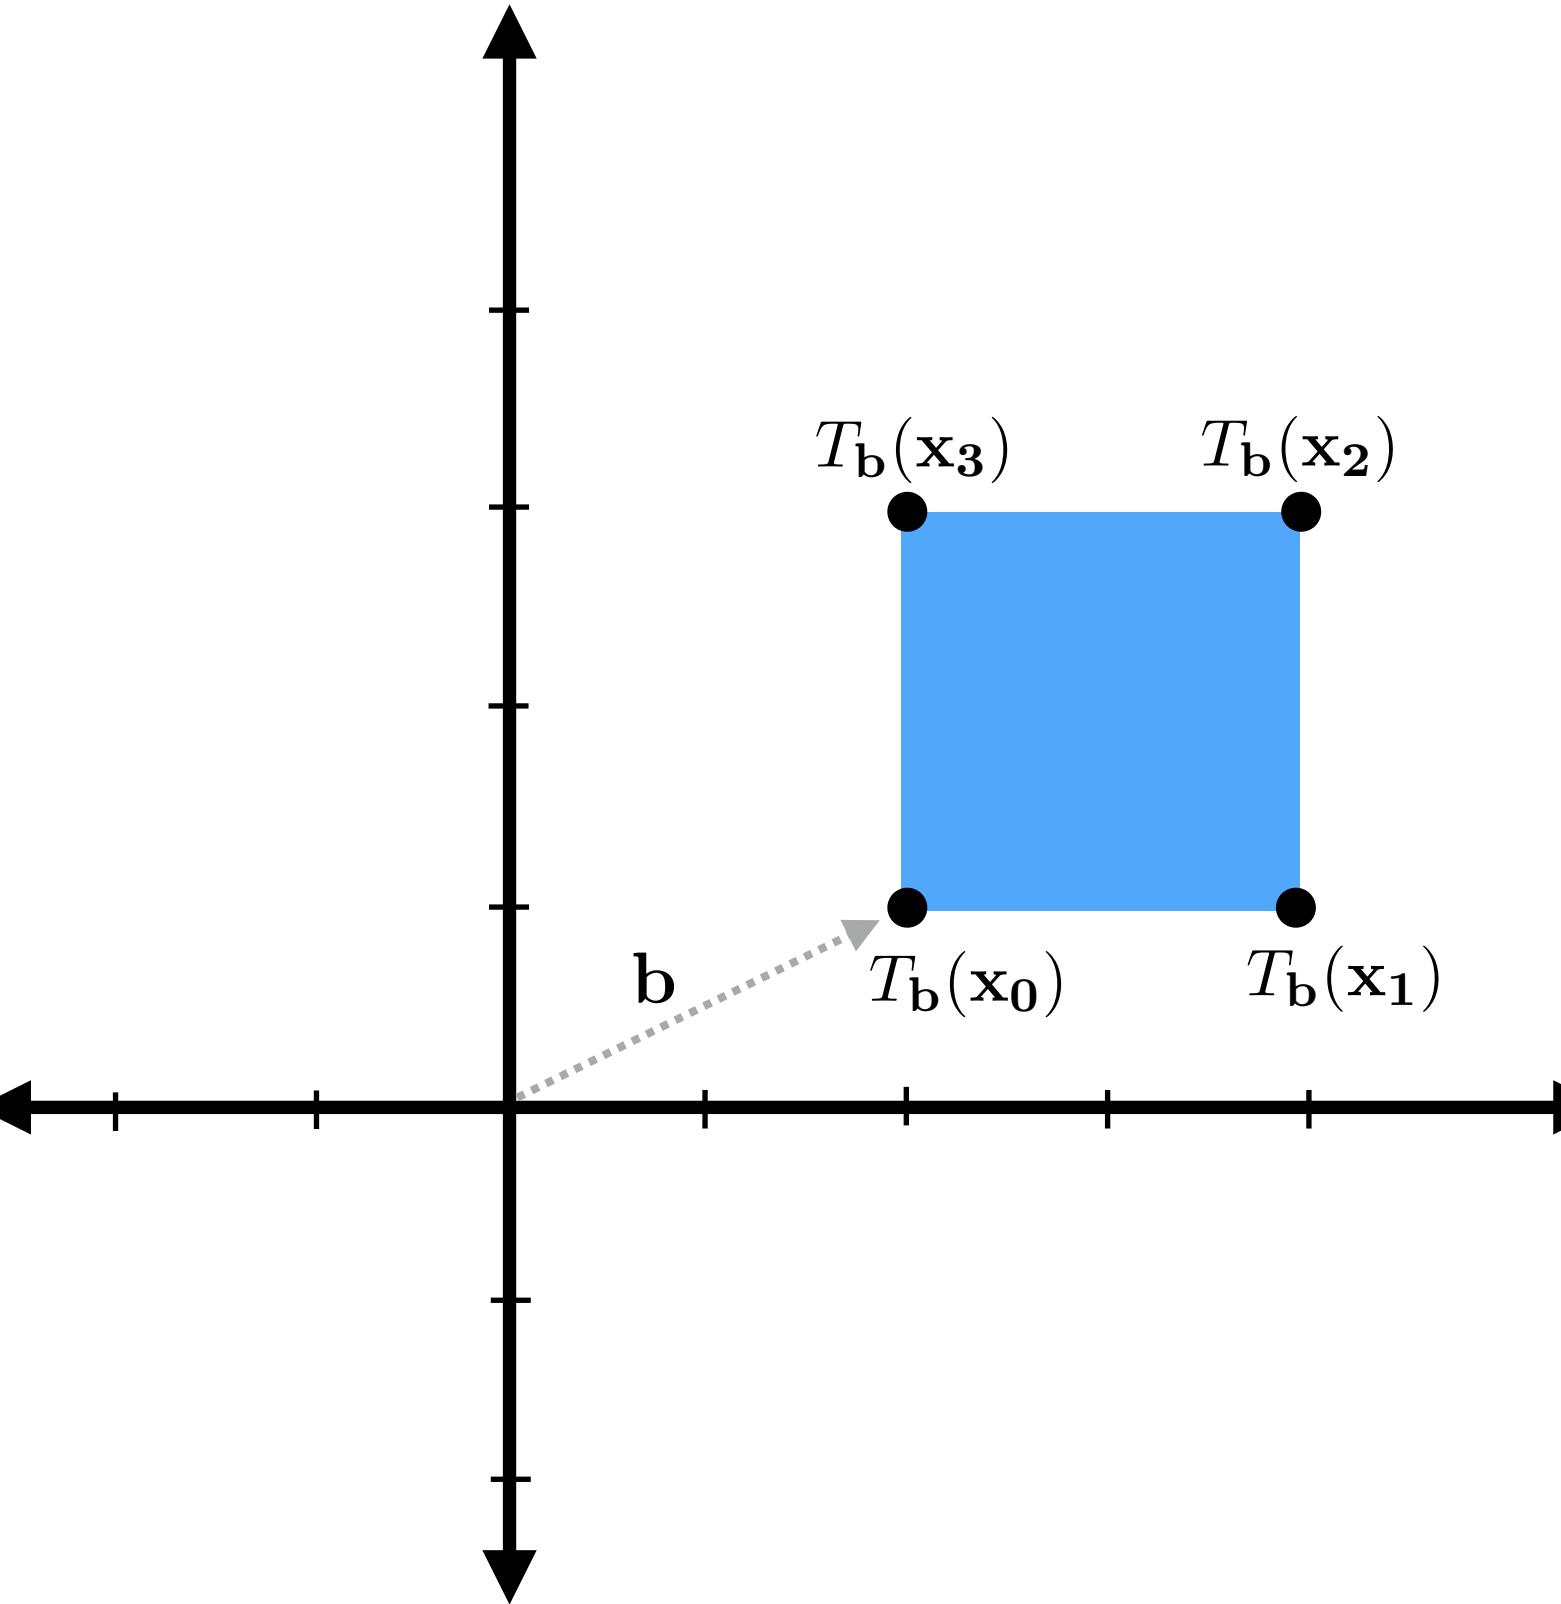
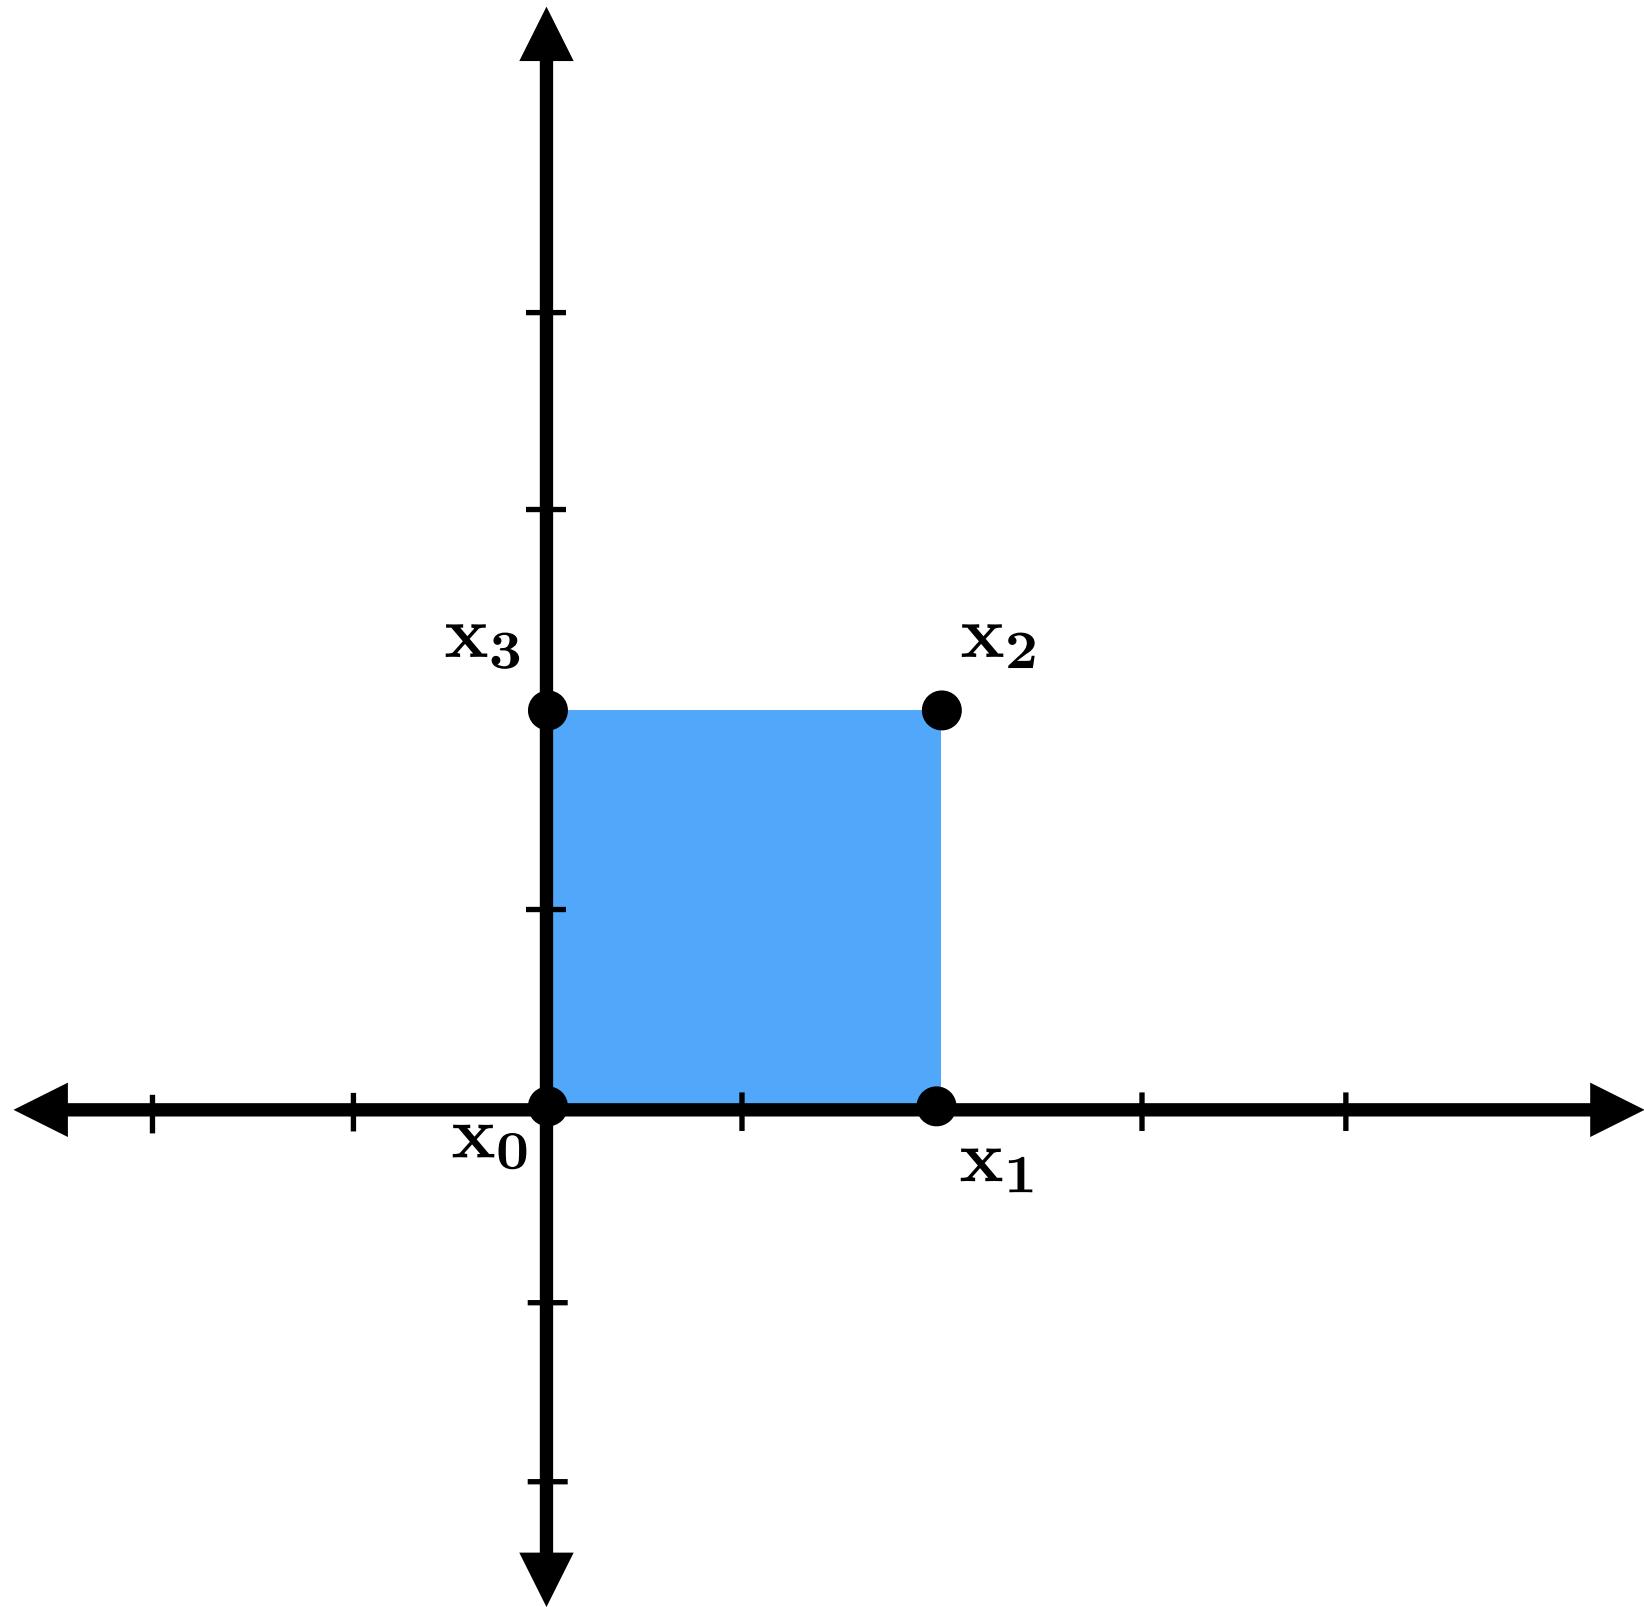
As angle changes, points move along *circular* trajectories.

Hence, rotations preserve length of vectors: $|R_\theta(\mathbf{x})| = |\mathbf{x}|$

Linearity?



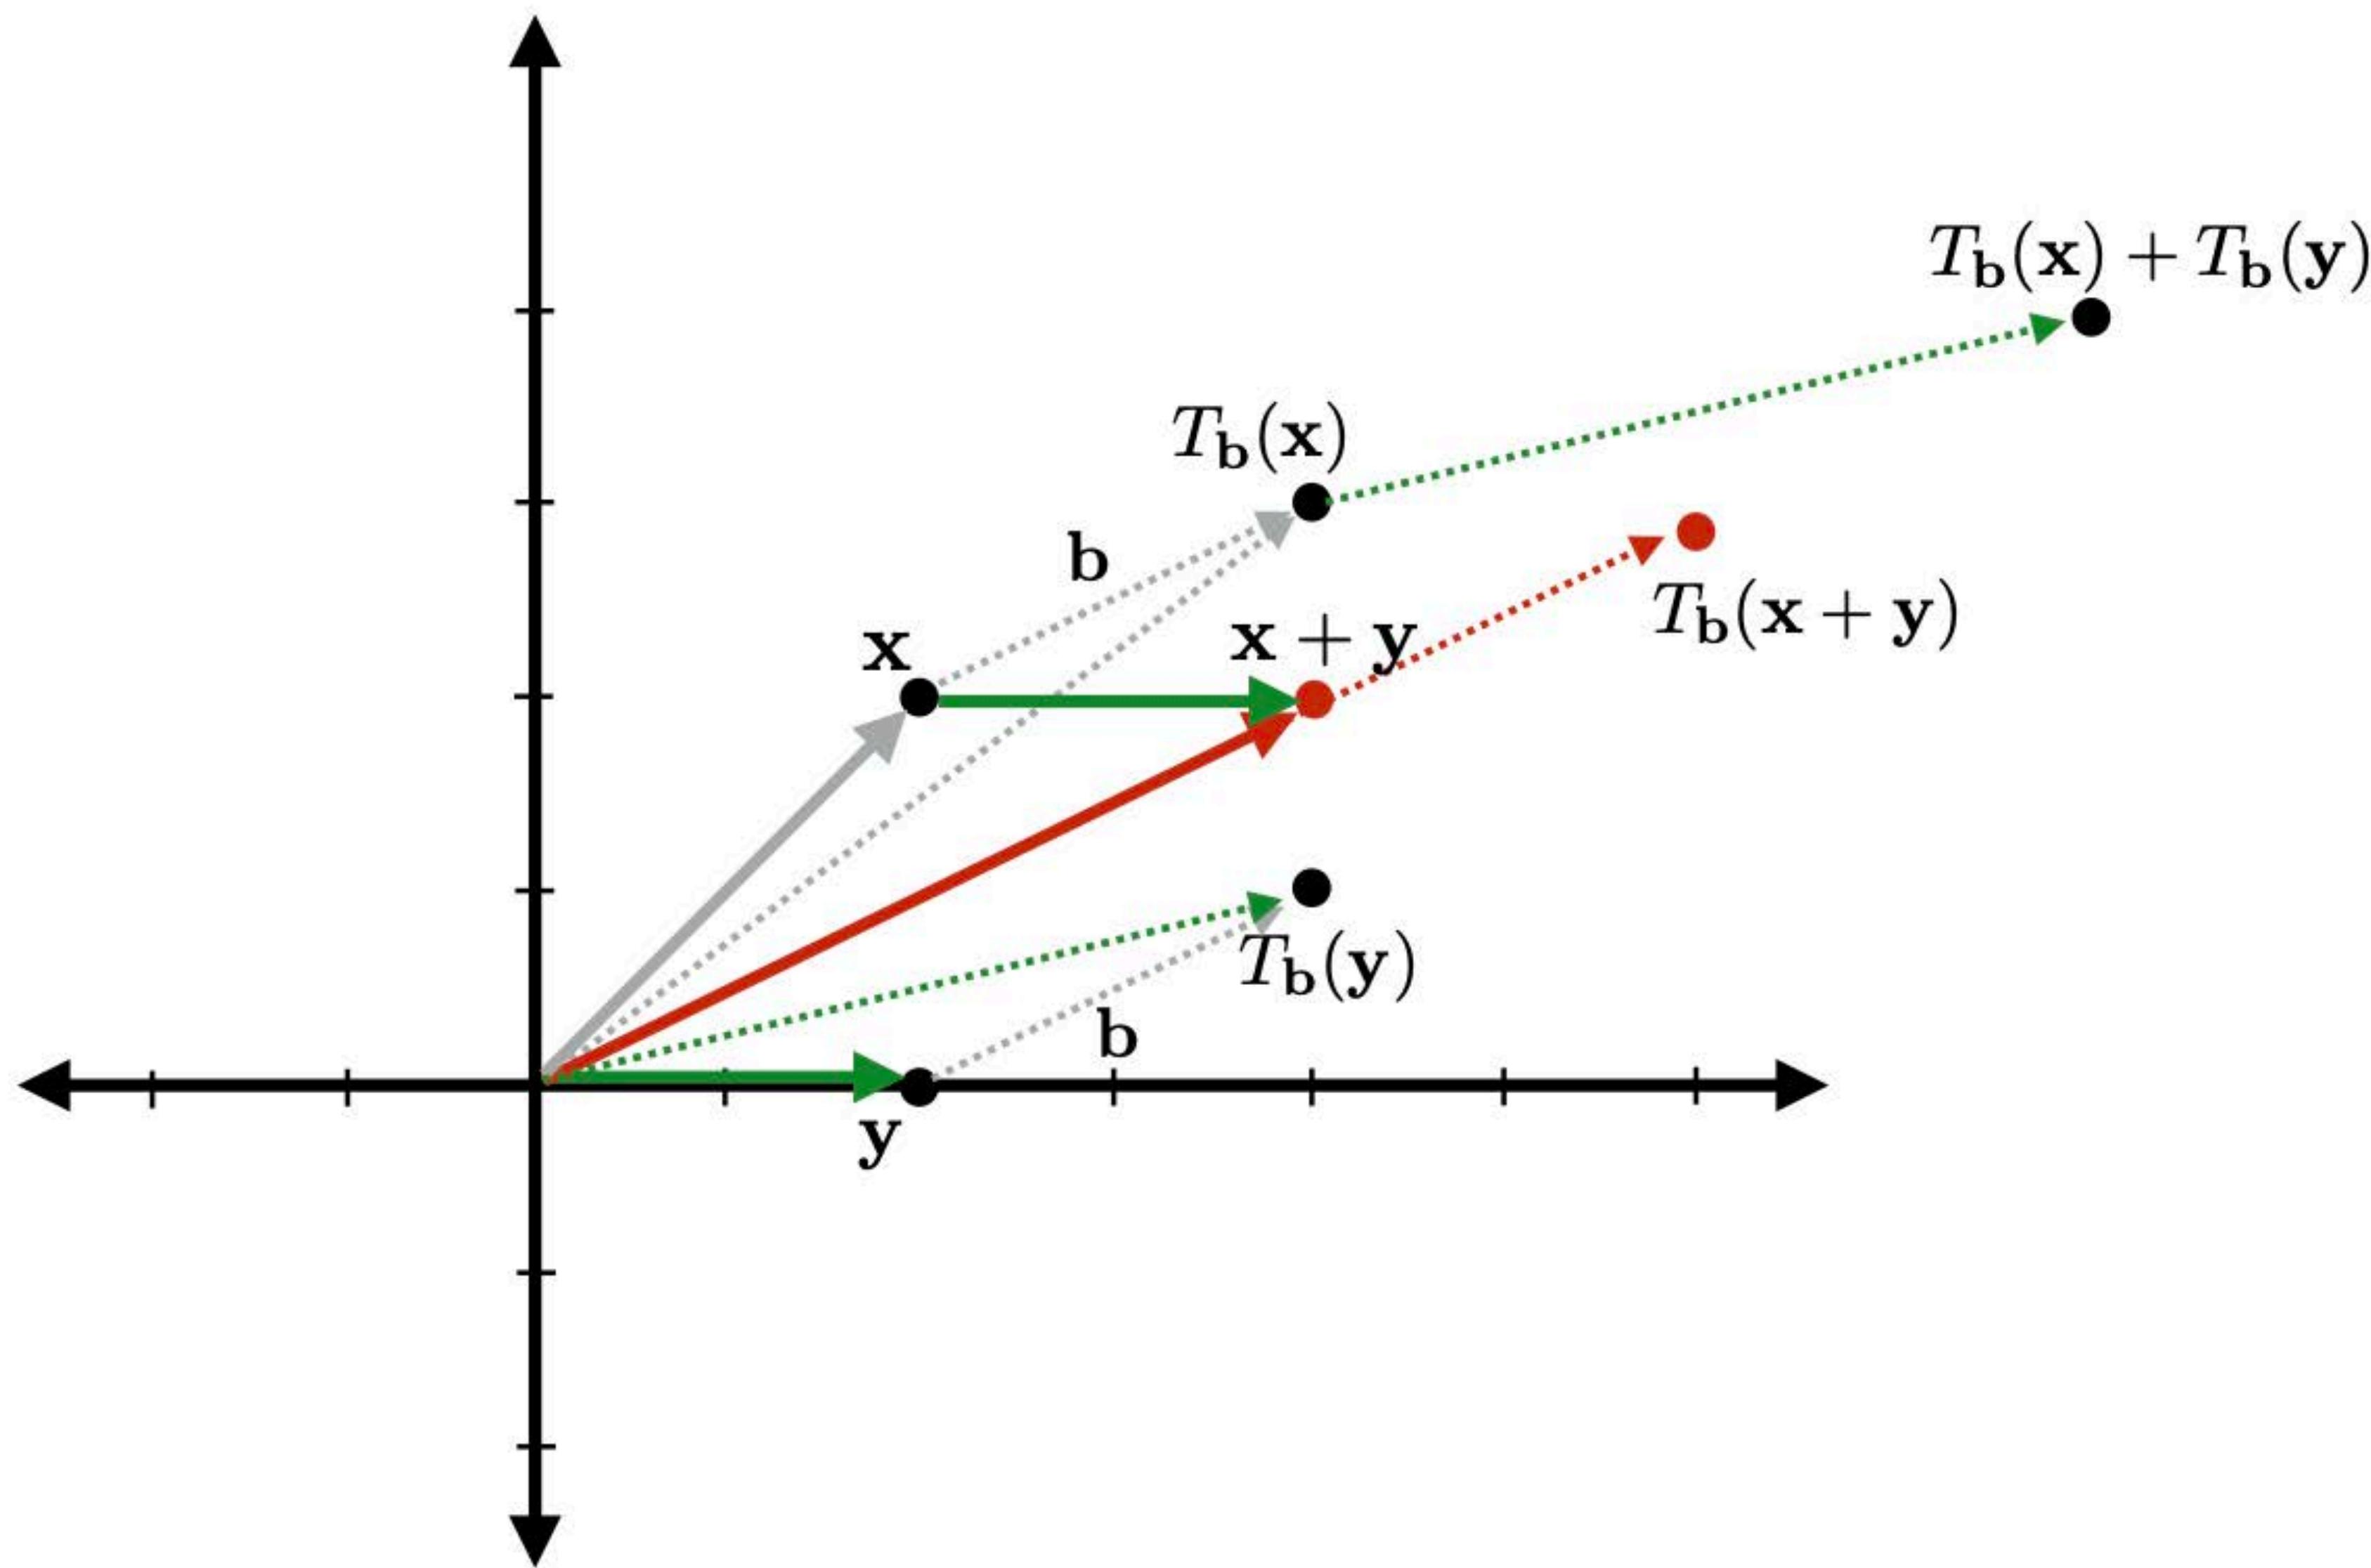
Translation



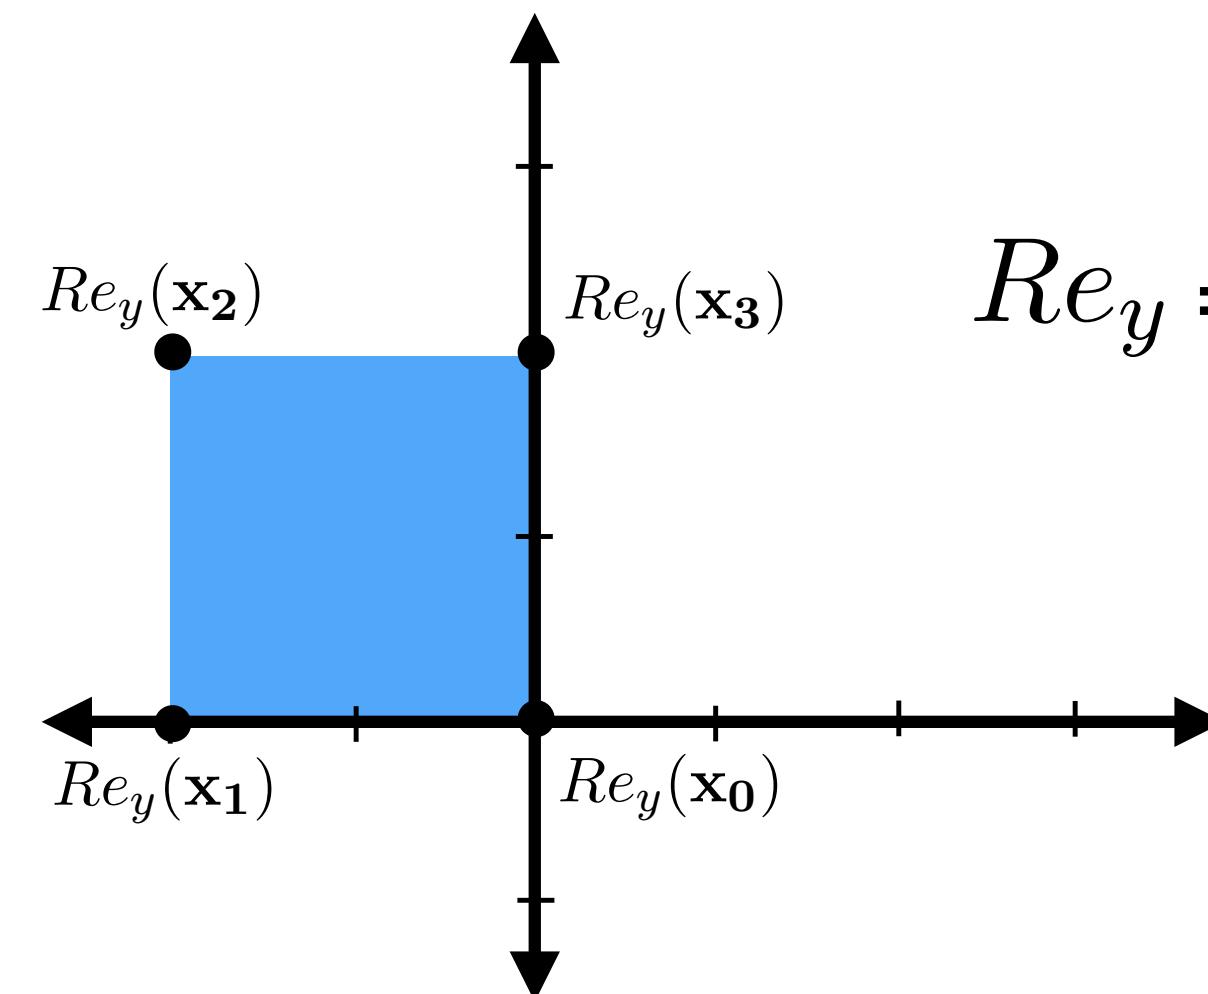
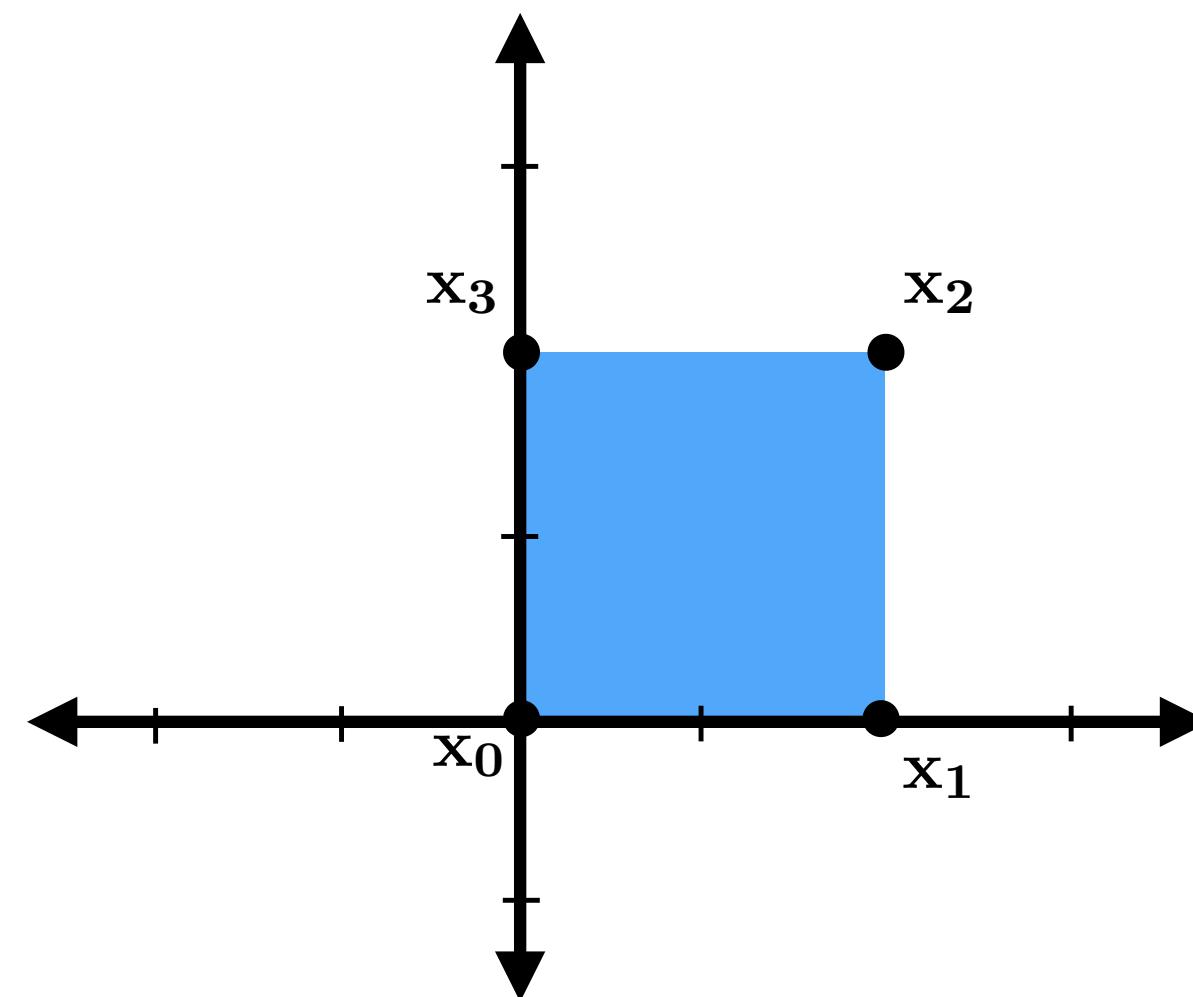
T_b — “translate by b ”

$$T_b(\mathbf{x}) = \mathbf{x} + \mathbf{b}$$

Linearity?

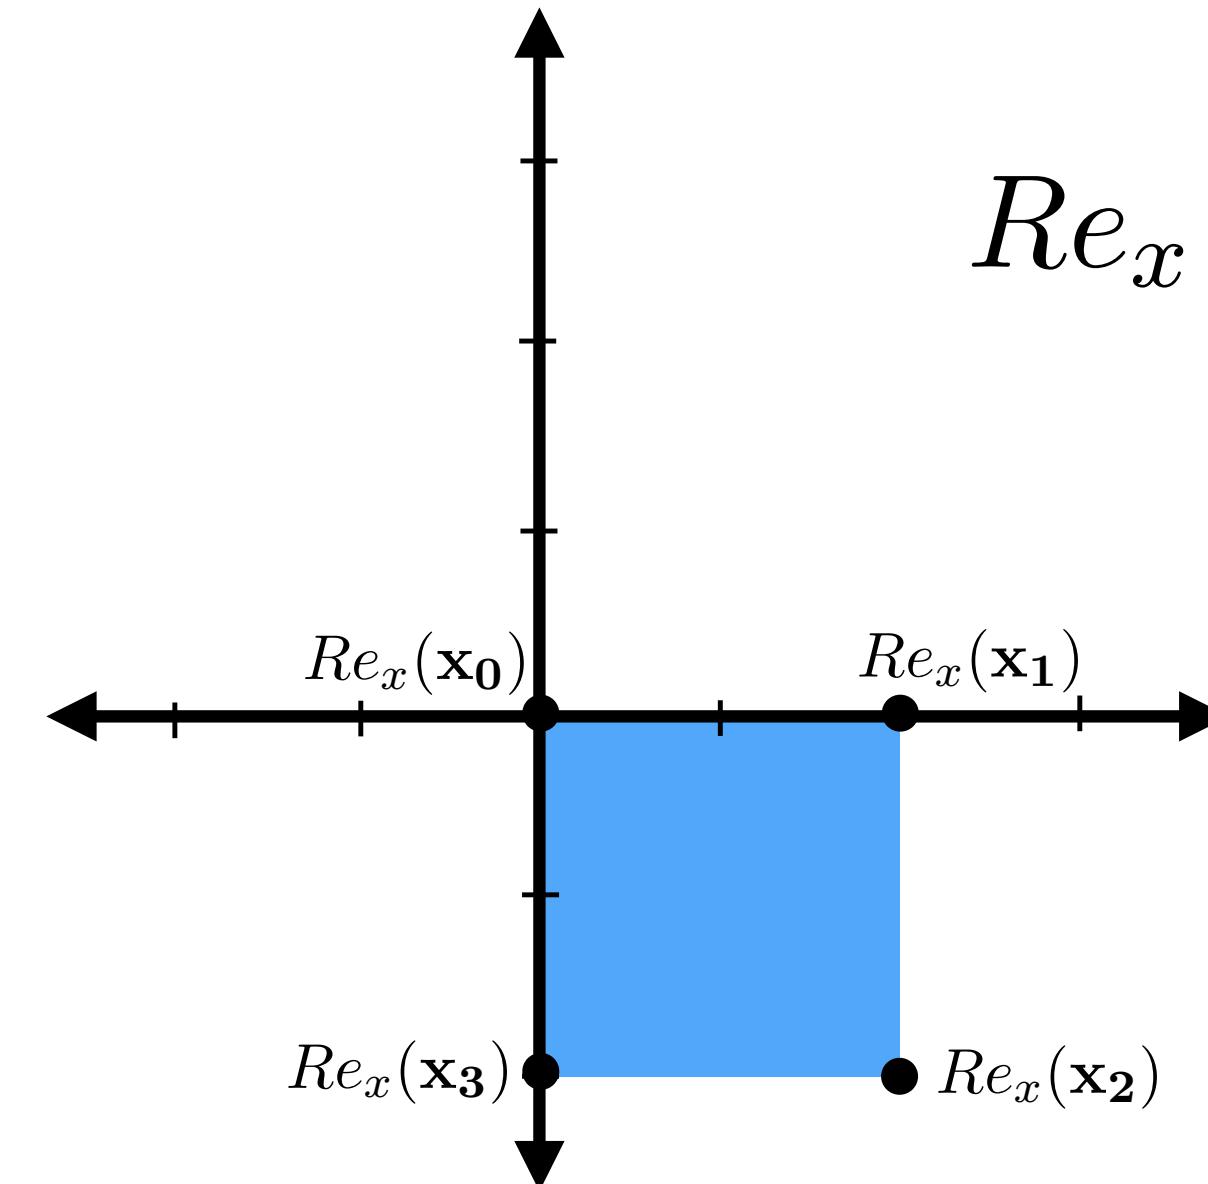


Reflection



Re_y = **reflection about y**

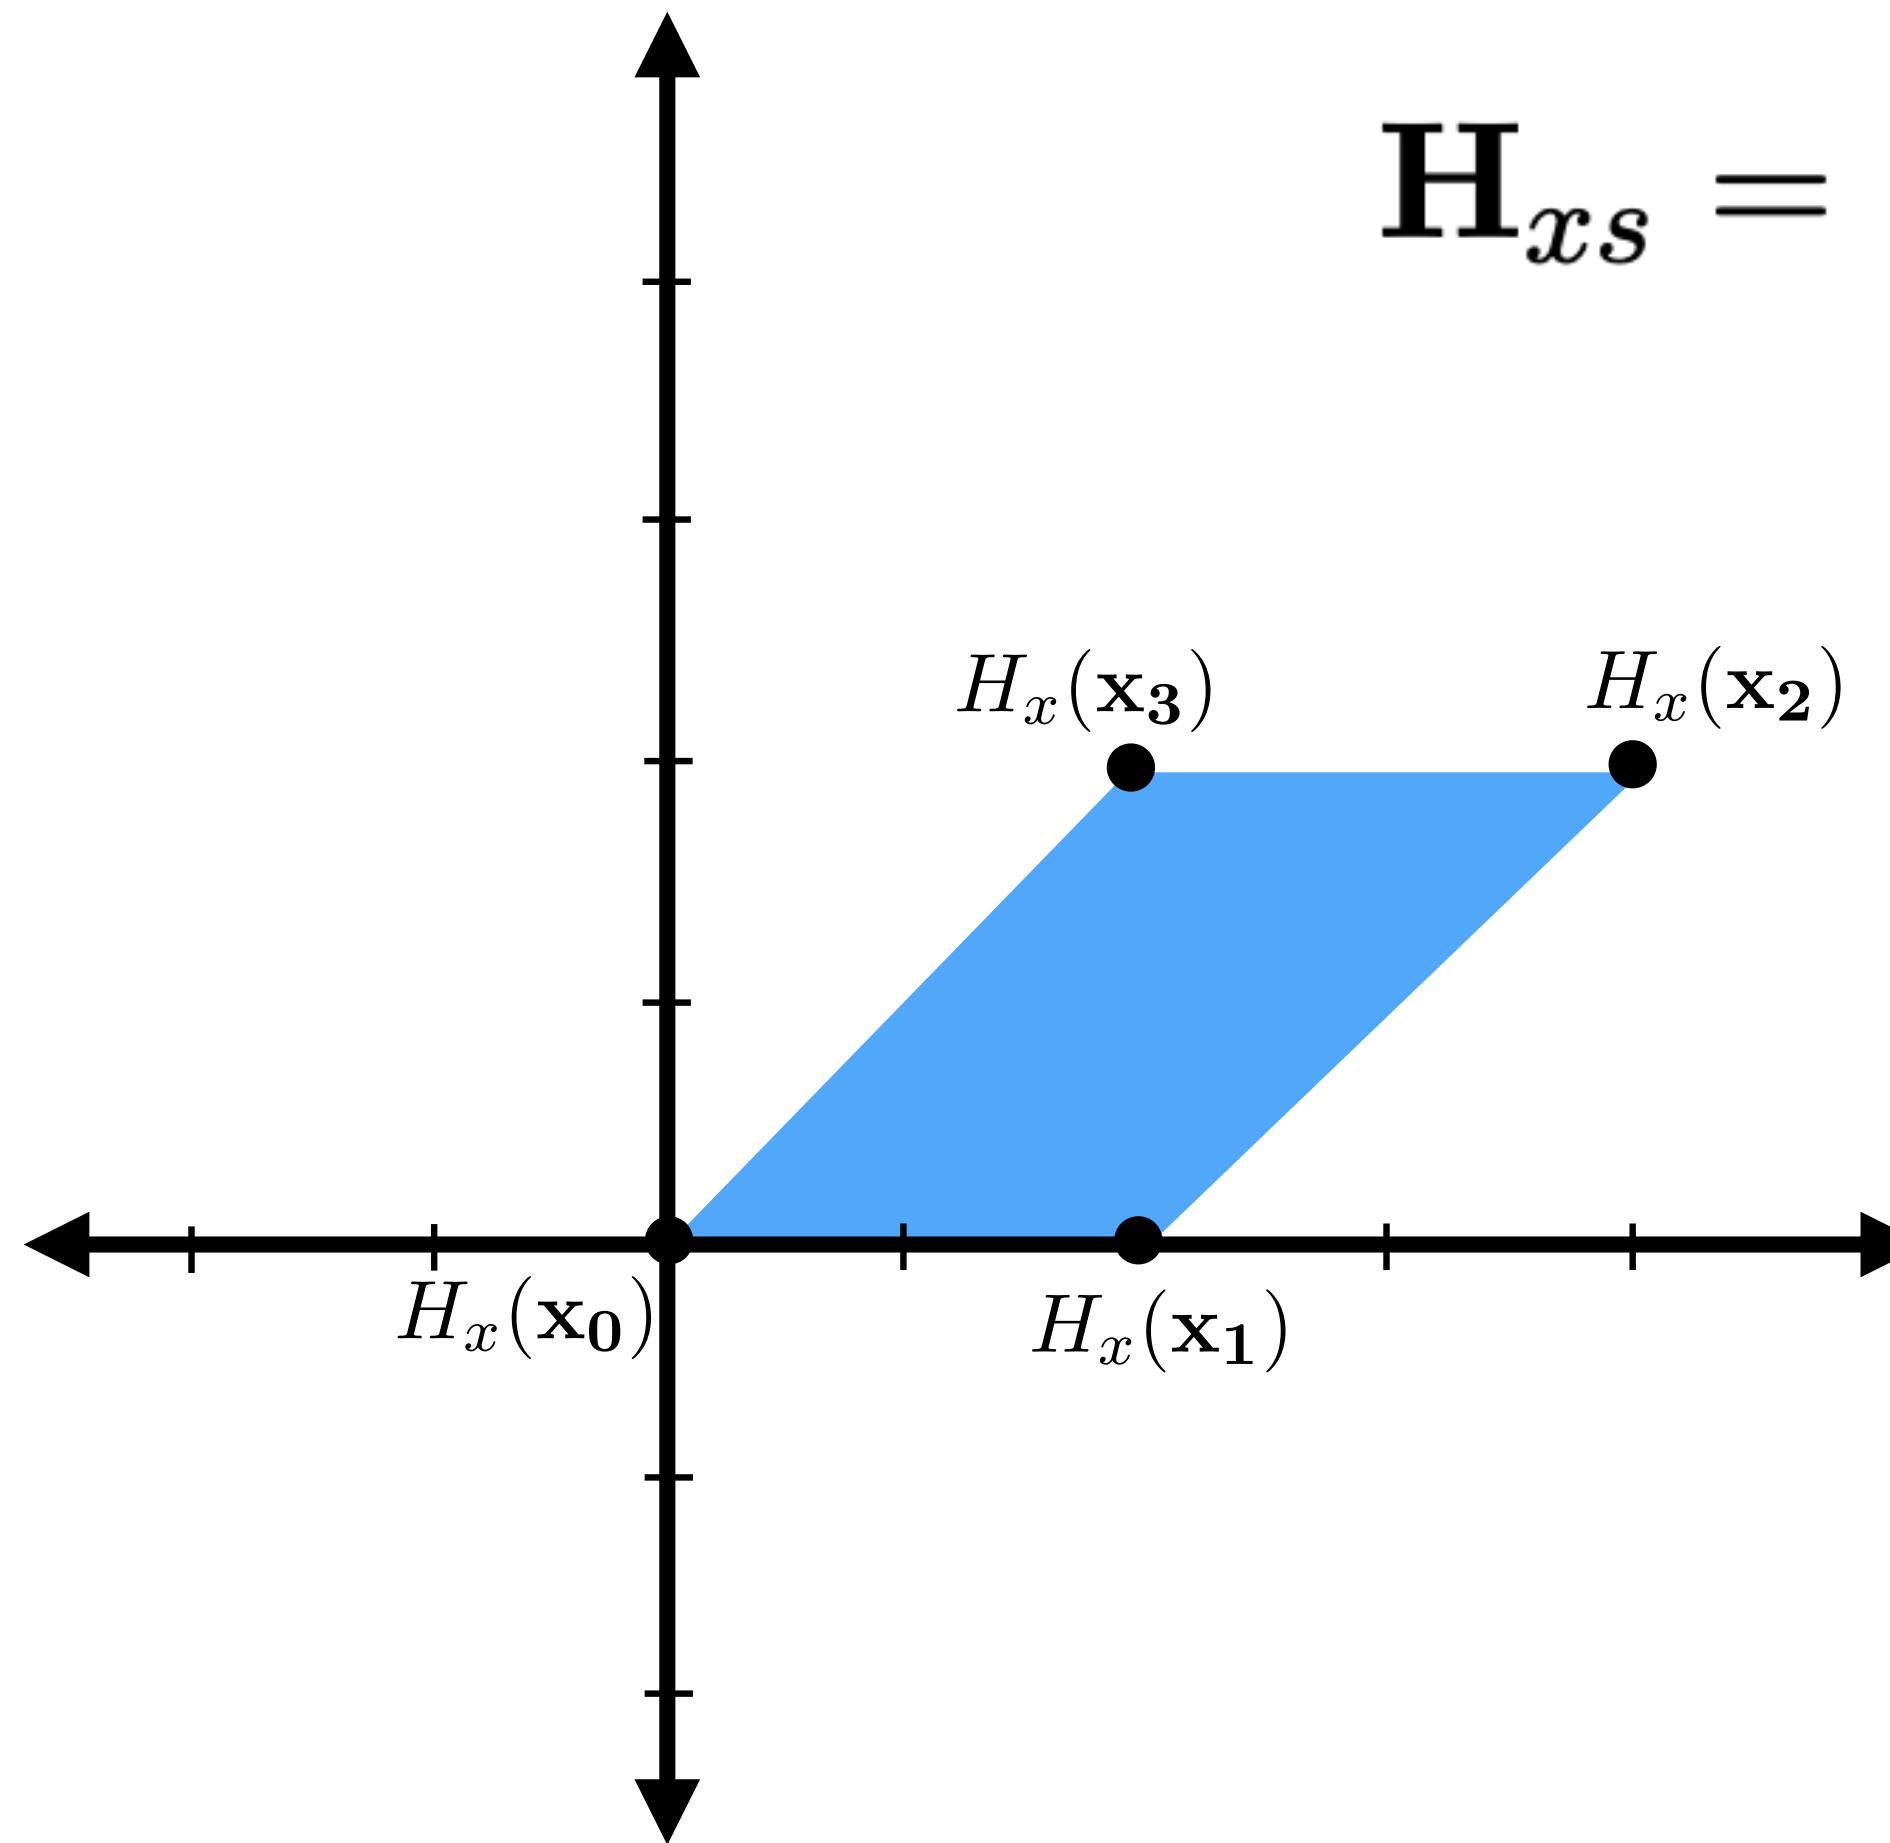
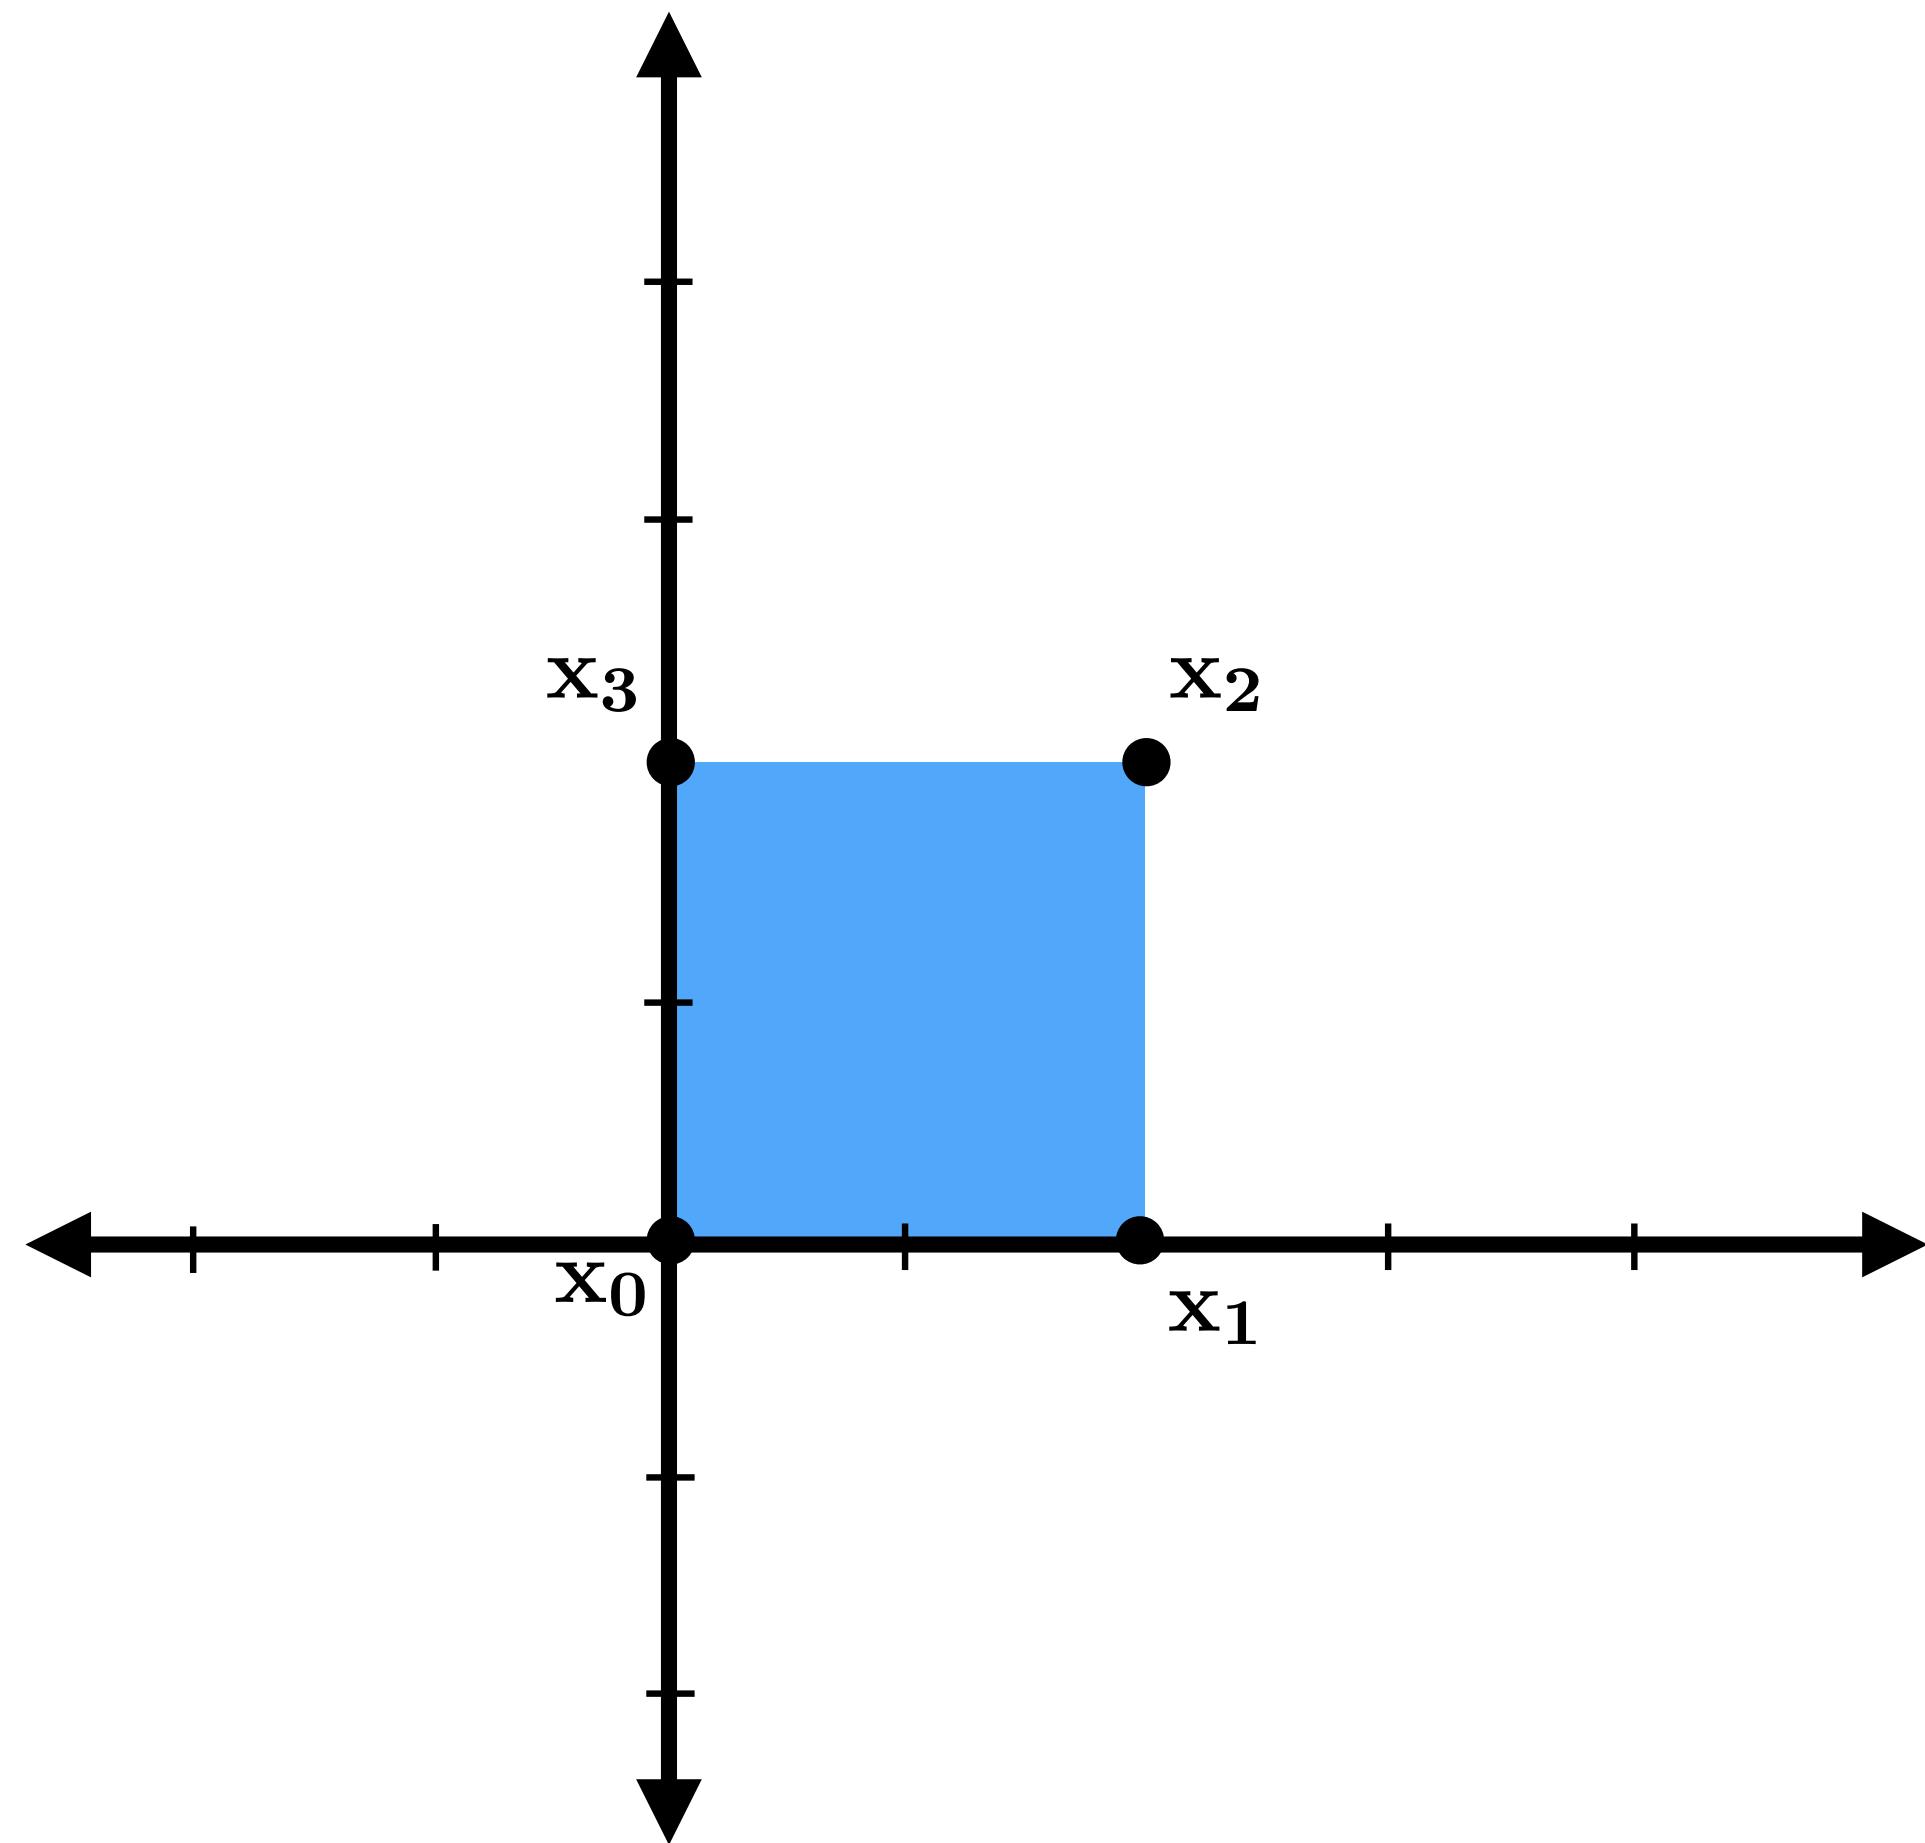
$$Re_y = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$$



Re_x = **reflection about x**

$$Re_x = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

Shear (in x direction)



$$\mathbf{H}_{xs} = \begin{bmatrix} 1 & s \\ 0 & 1 \end{bmatrix}$$

Summary of Transformations

Linear:

$$f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$$

$$f(a\mathbf{x}) = af(\mathbf{x})$$

Scale

Rotation

Reflection

Shear

Not linear:

Translation

Affine:

Composition of linear transform + translation

$$f(\mathbf{x}) = g(\mathbf{x}) + \mathbf{b}$$

Euclidean: (Isometries)

Preserve distance between points (preserves length)

$$|f(\mathbf{x}) - f(\mathbf{y})| = |\mathbf{x} - \mathbf{y}|$$

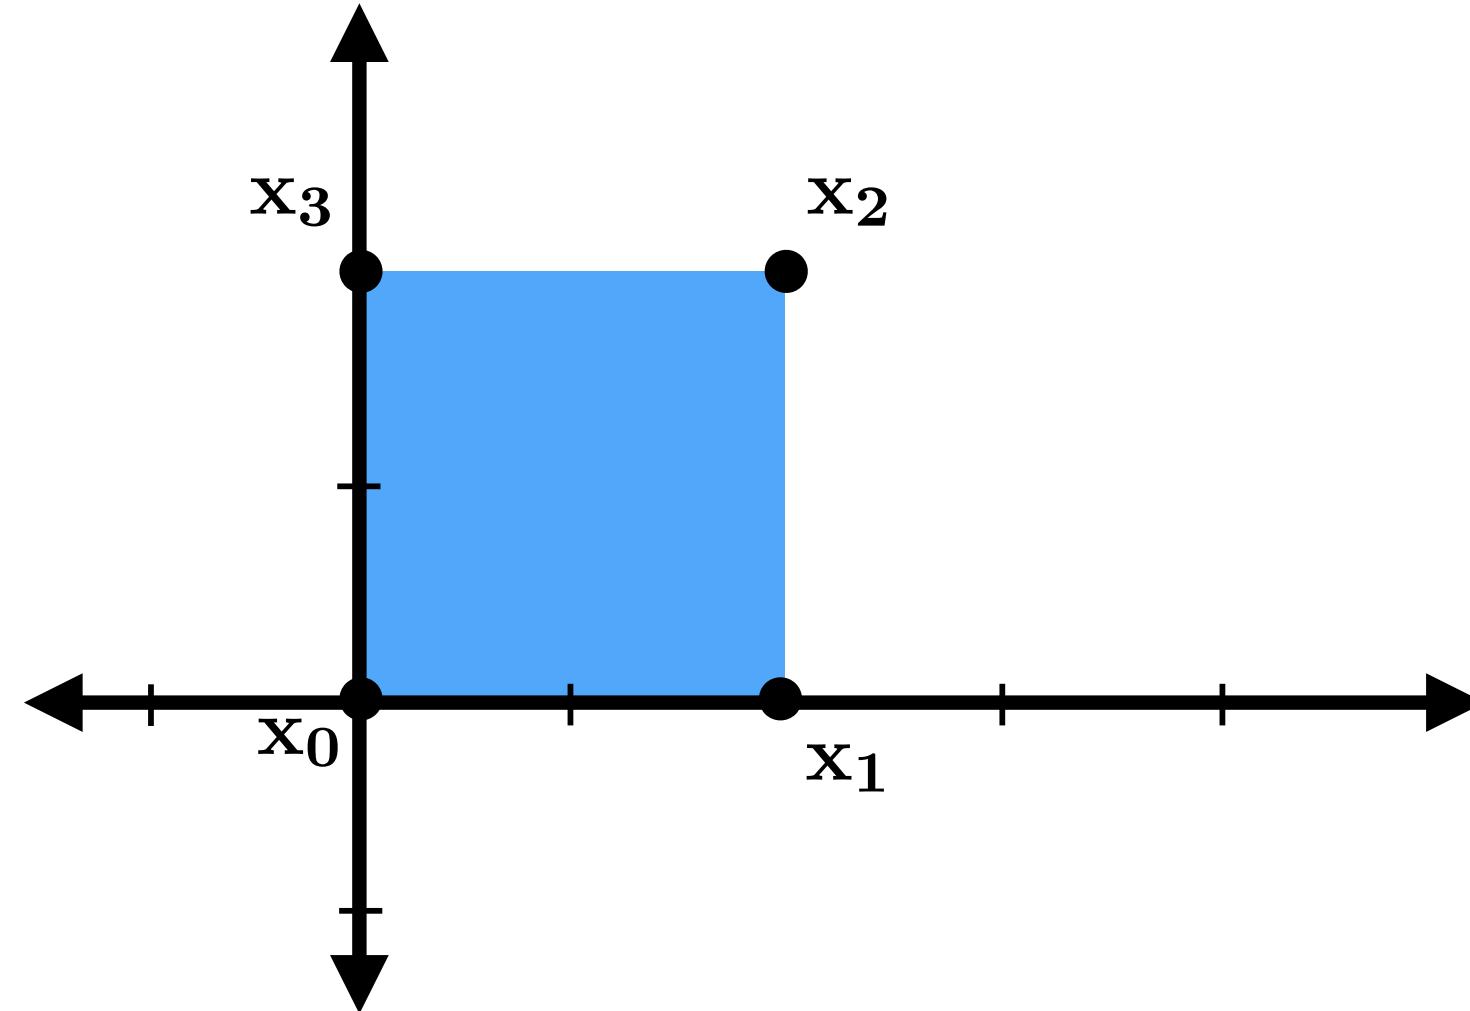
Translation

Rotation

Reflection

“Rigid body” transformations are distance-preserving motions that also preserve *orientation* (i.e., does not include reflection)

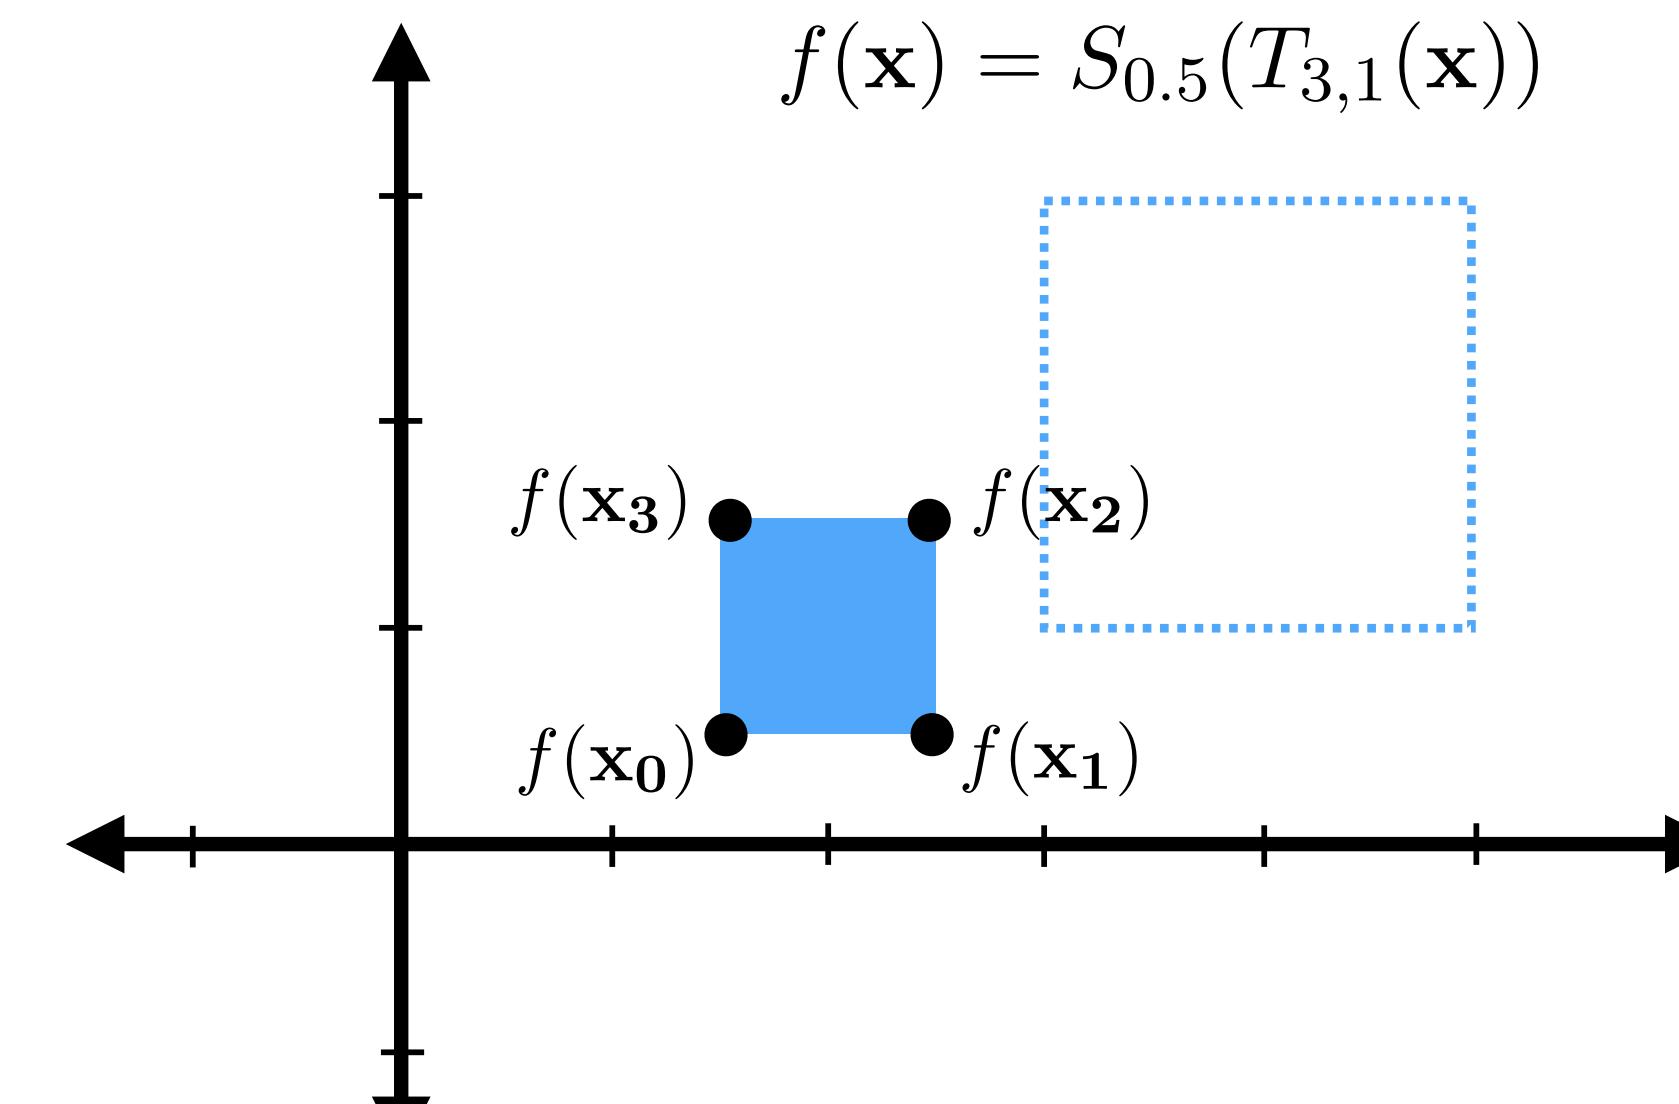
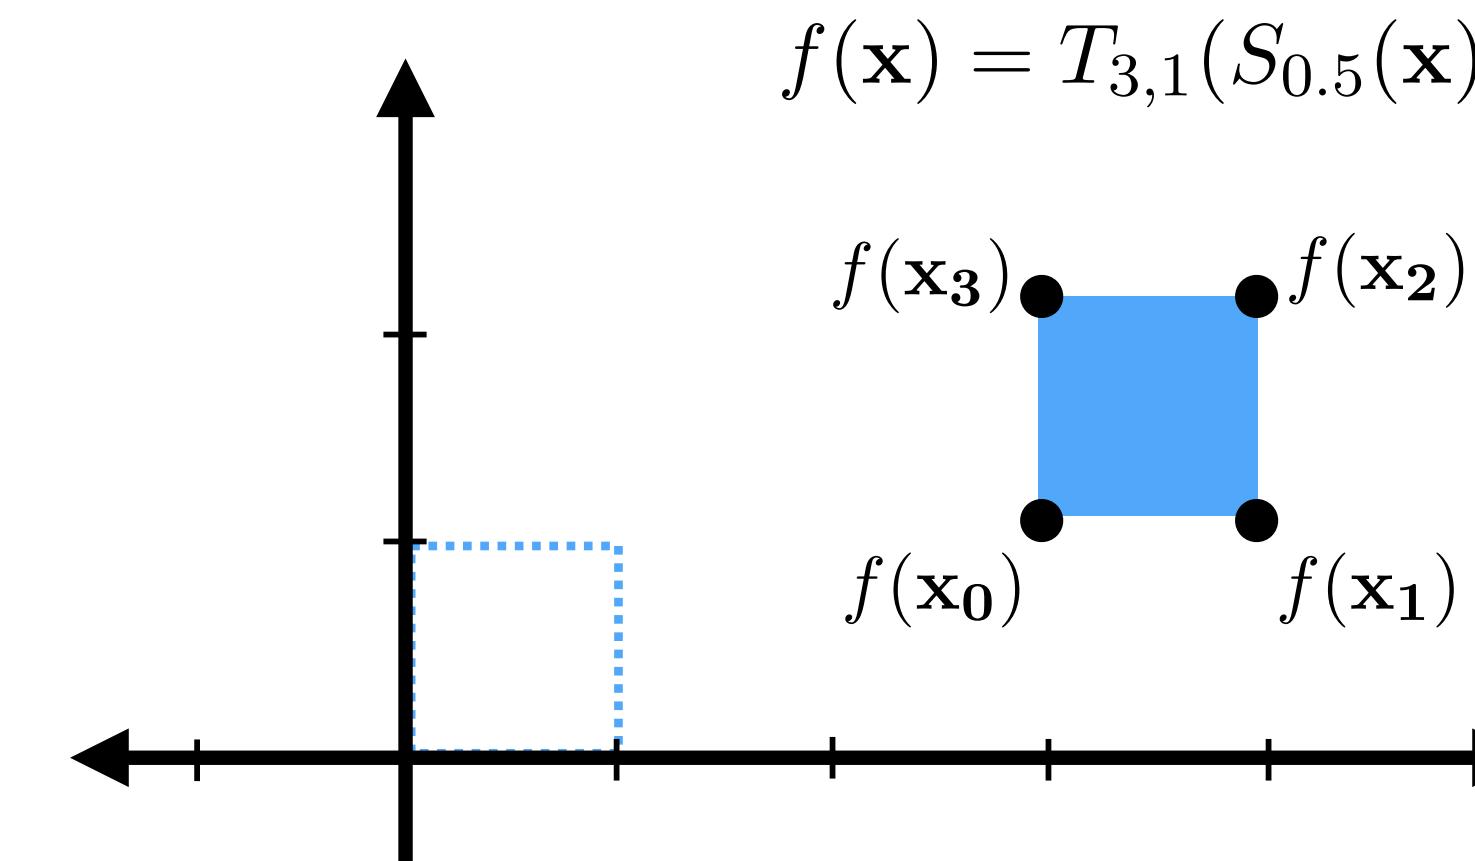
Composition of Transformations



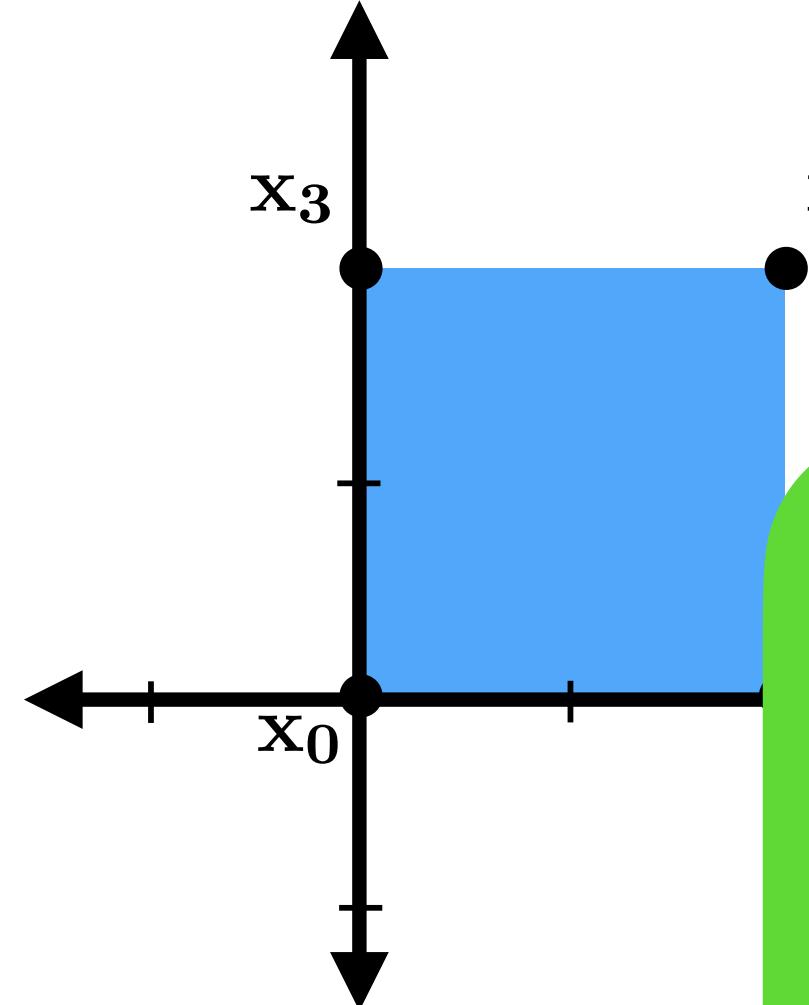
Note: order of composition matters

Top-right: scale, then translate

Bottom-right: translate, then scale



Composition of Transformations



We could leverage matrix multiplication without translation!

How to deal with translation then?

Note: order of composition matters

Top-right: scale, then translate

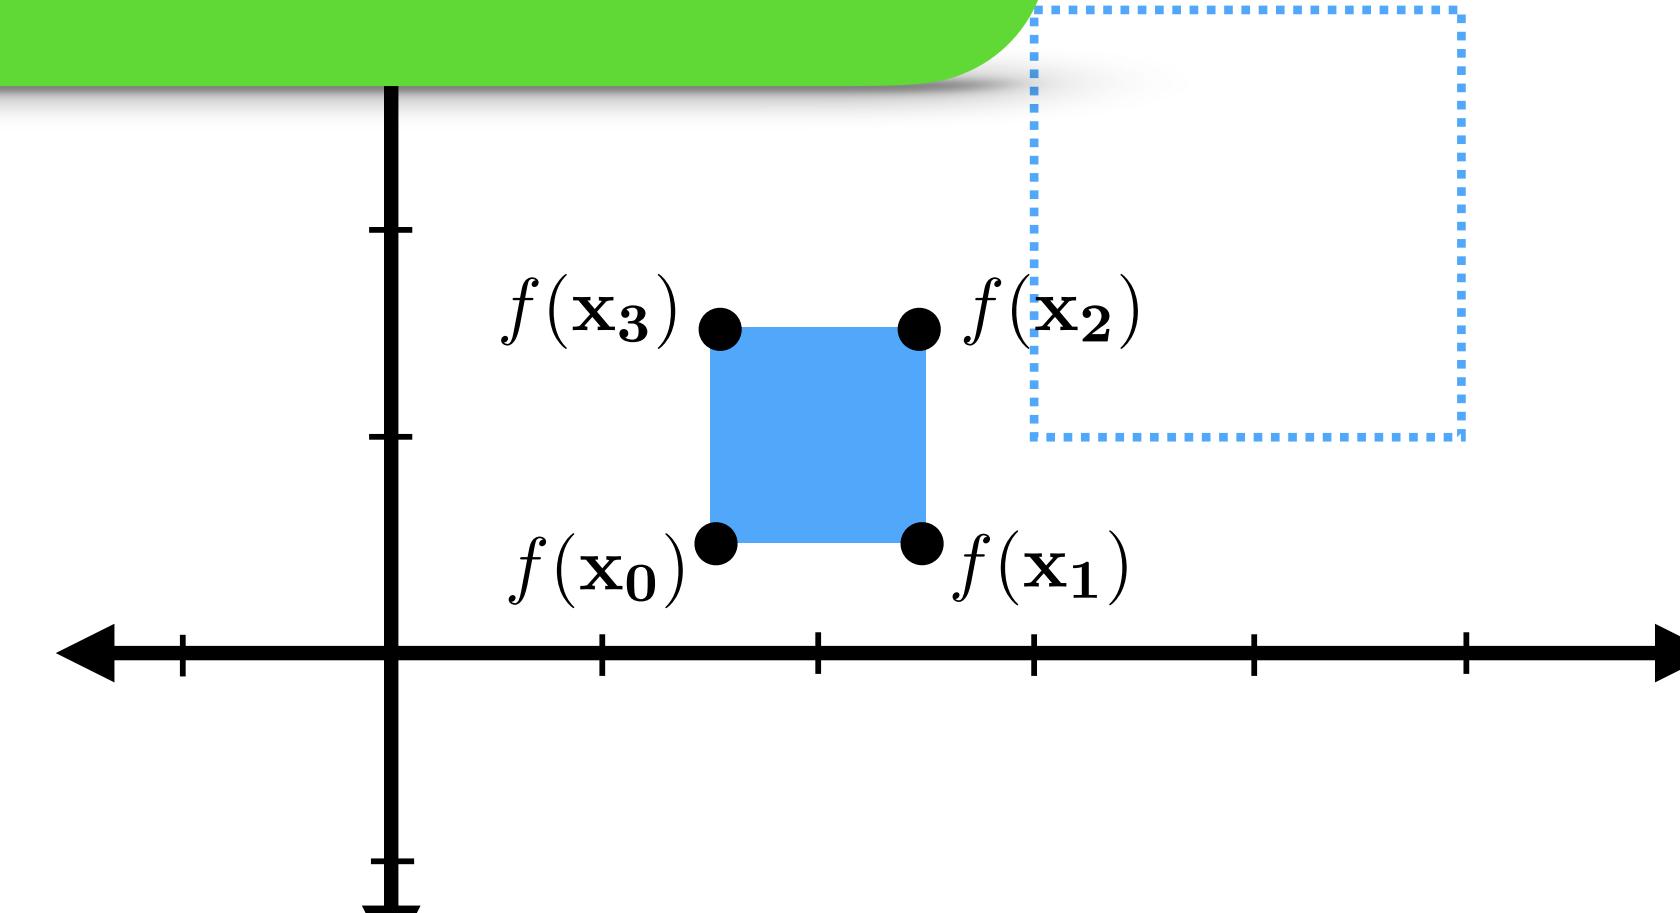
Bottom-right: translate, then scale

$$f(\mathbf{x}) = T_{3,1}(S_{0.5}(\mathbf{x}))$$

$$f(\mathbf{x}_3) \quad f(\mathbf{x}_2)$$

$$f(\mathbf{x}_1)$$

$$S_{0.5}(T_{3,1}(\mathbf{x}))$$



Homogeneous Coordinates (in 2D)

Idea: represent 2D points with THREE values (“homogeneous coordinates”)

So the point (x, y) is represented as the 3-vector: $[x \quad y \quad 1]^T$

And transformations are represented a 3×3 matrices that transform these vectors.

Recover final 2D coordinates by dividing by “extra” (third) coordinate

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow \begin{bmatrix} x/w \\ y/w \end{bmatrix}$$

Scale and Rotation with 2D-H

- For transformations that are already linear, not much changes:

$$\mathbf{S}_s = \begin{bmatrix} S_x & 0 & 0 \\ 0 & S_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad \mathbf{R}_\theta = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Notice that the last row/column doesn't do anything interesting. E.g., for scaling:

$$\mathbf{S}_s \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} S_x x \\ S_y y \\ 1 \end{bmatrix}$$

Now we divide by the 3rd coordinate to get our final 2D coordinates (not too exciting!)

$$\begin{bmatrix} S_x x \\ S_y y \\ 1 \end{bmatrix} \Rightarrow \begin{bmatrix} S_x x / 1 \\ S_y y / 1 \end{bmatrix} = \begin{bmatrix} S_x x \\ S_y y \end{bmatrix}$$

Translation with 2D-H

Translation expressed as 3x3 matrix multiplication:

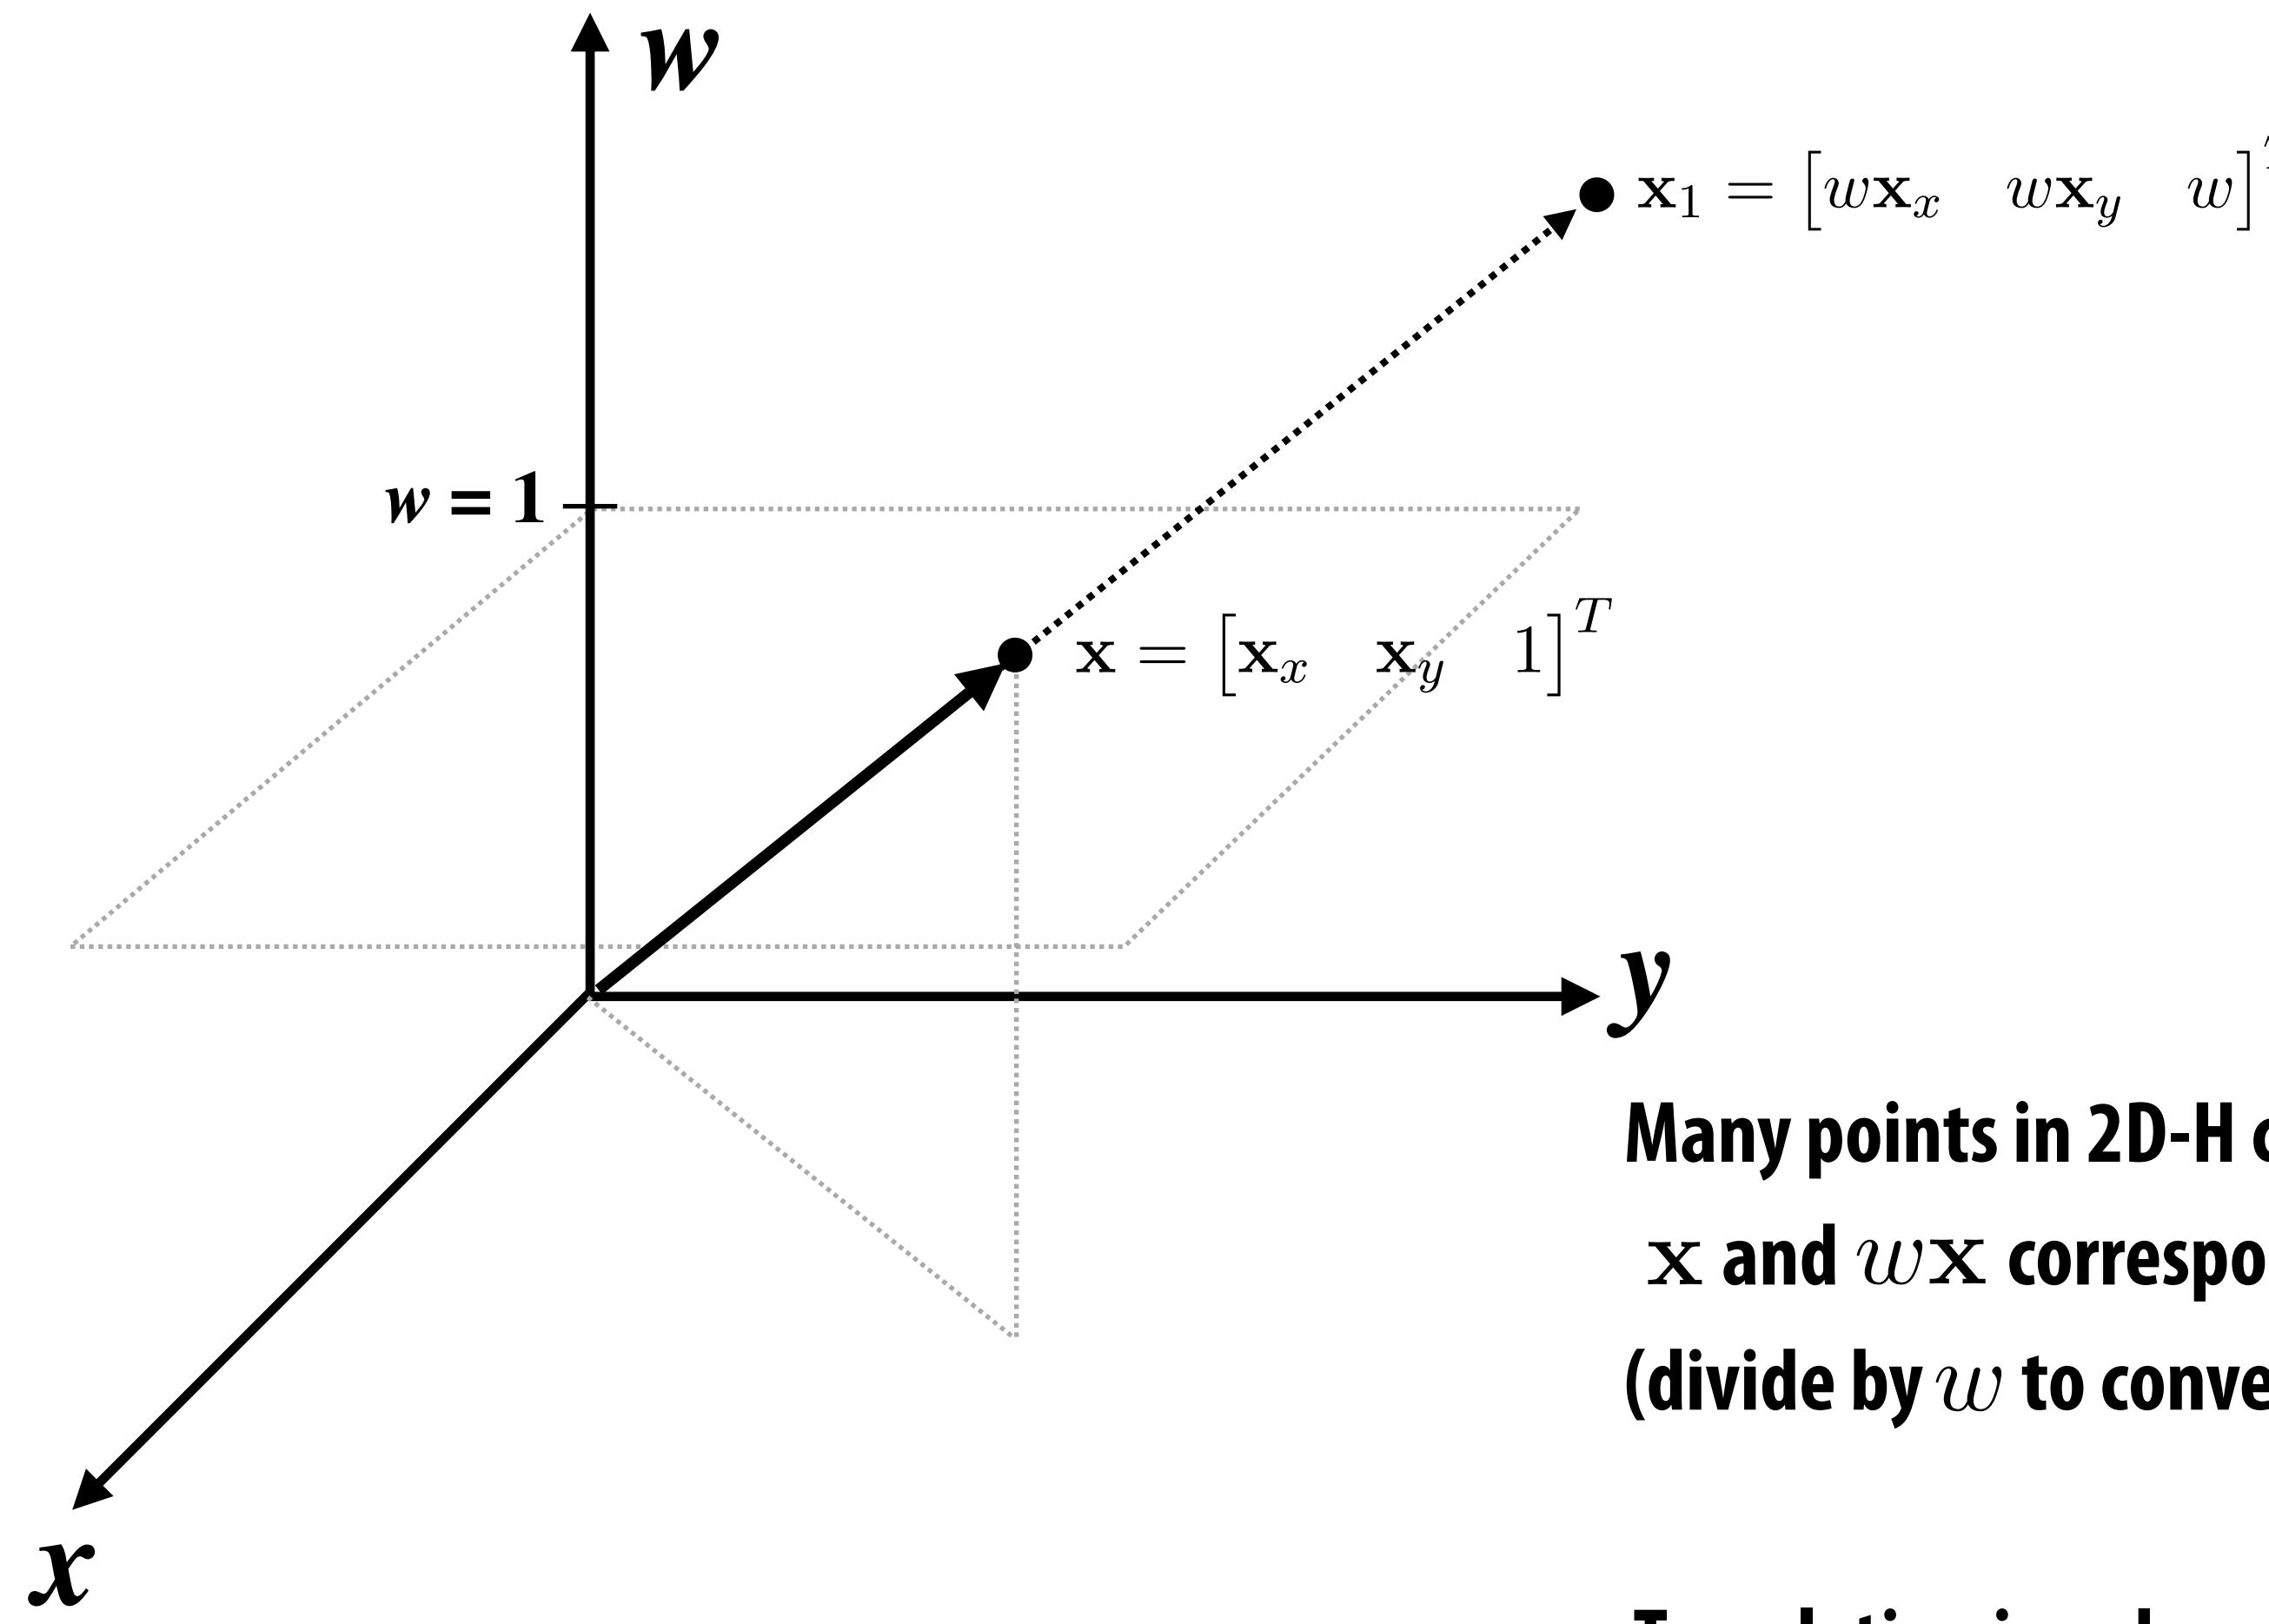
$$\mathbf{T}_b = \begin{bmatrix} 1 & 0 & b_x \\ 0 & 1 & b_y \\ 0 & 0 & 1 \end{bmatrix}$$

$$\mathbf{T}_b \mathbf{x} = \begin{bmatrix} 1 & 0 & b_x \\ 0 & 1 & b_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_x \\ x_y \\ 1 \end{bmatrix} = \begin{bmatrix} x_x + b_x \\ x_y + b_y \\ 1 \end{bmatrix}$$

**(remember: just a linear combination
of columns!)**

Cool: homogeneous coordinates let us encode translations as *linear* transformations!

Intuition behind Homogeneous Coordinates

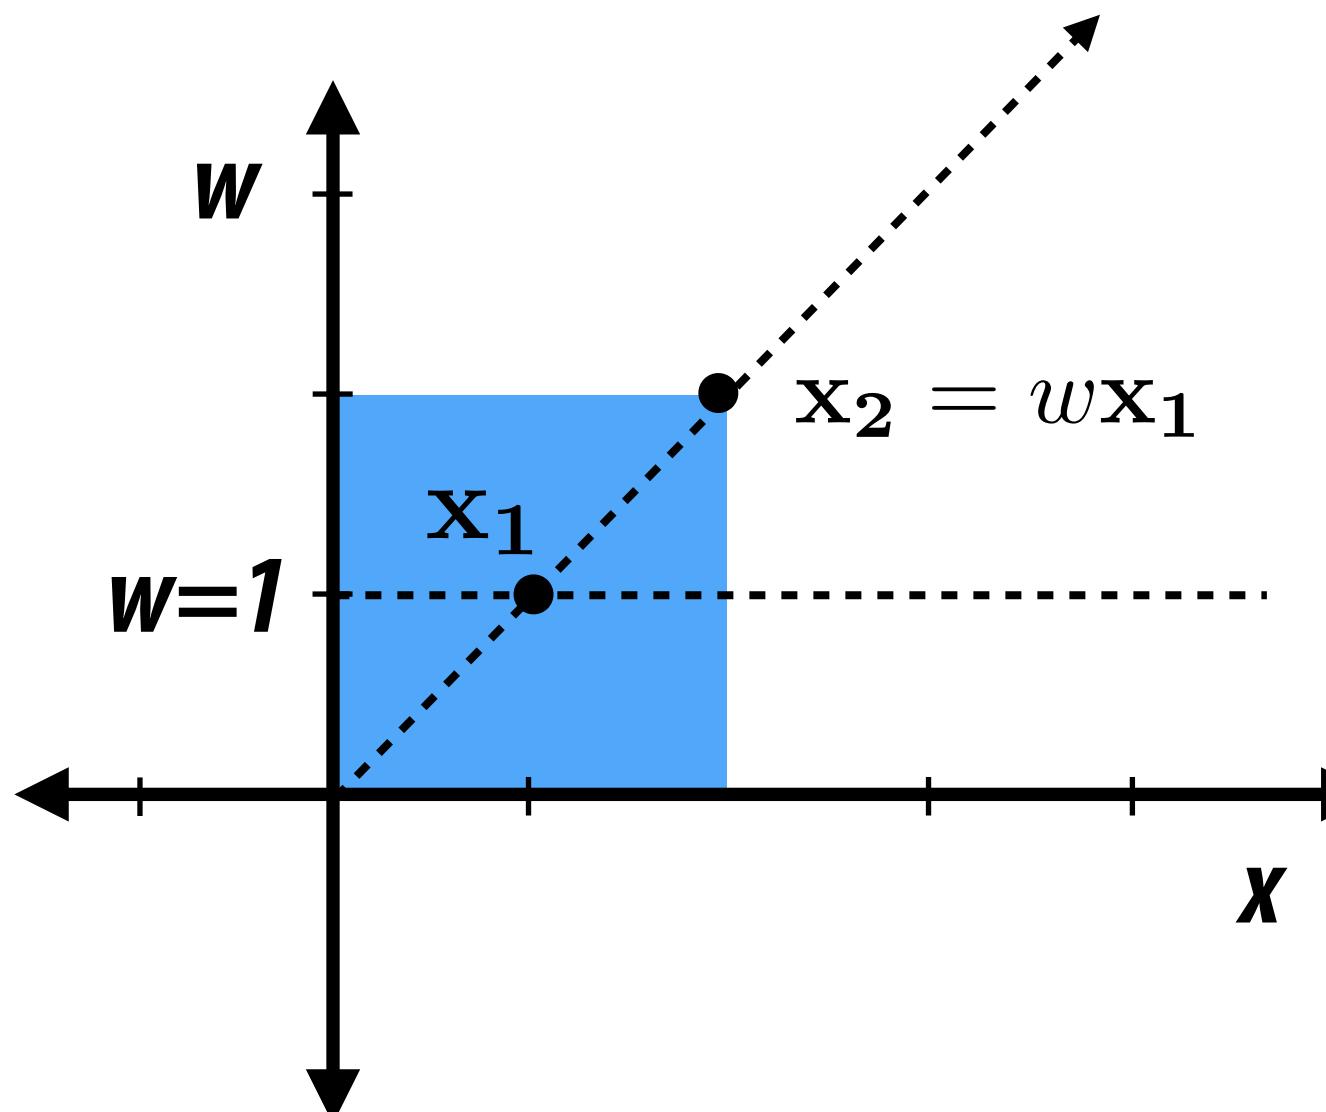


**Many points in 2D-H correspond to same point in 2D
 \mathbf{x} and $w\mathbf{x}$ correspond to the same 2D point
(divide by w to convert 2D-H back to 2D)**

Translation is a shear in x and y in 2D-H space

$$\mathbf{T}_b \mathbf{x} = \begin{bmatrix} 1 & 0 & b_x \\ 0 & 1 & b_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} w\mathbf{x}_x \\ w\mathbf{x}_y \\ w \end{bmatrix} = \begin{bmatrix} w\mathbf{x}_x + w\mathbf{b}_x \\ w\mathbf{x}_y + w\mathbf{b}_y \\ w \end{bmatrix}$$

Intuition behind Homogeneous Coordinates

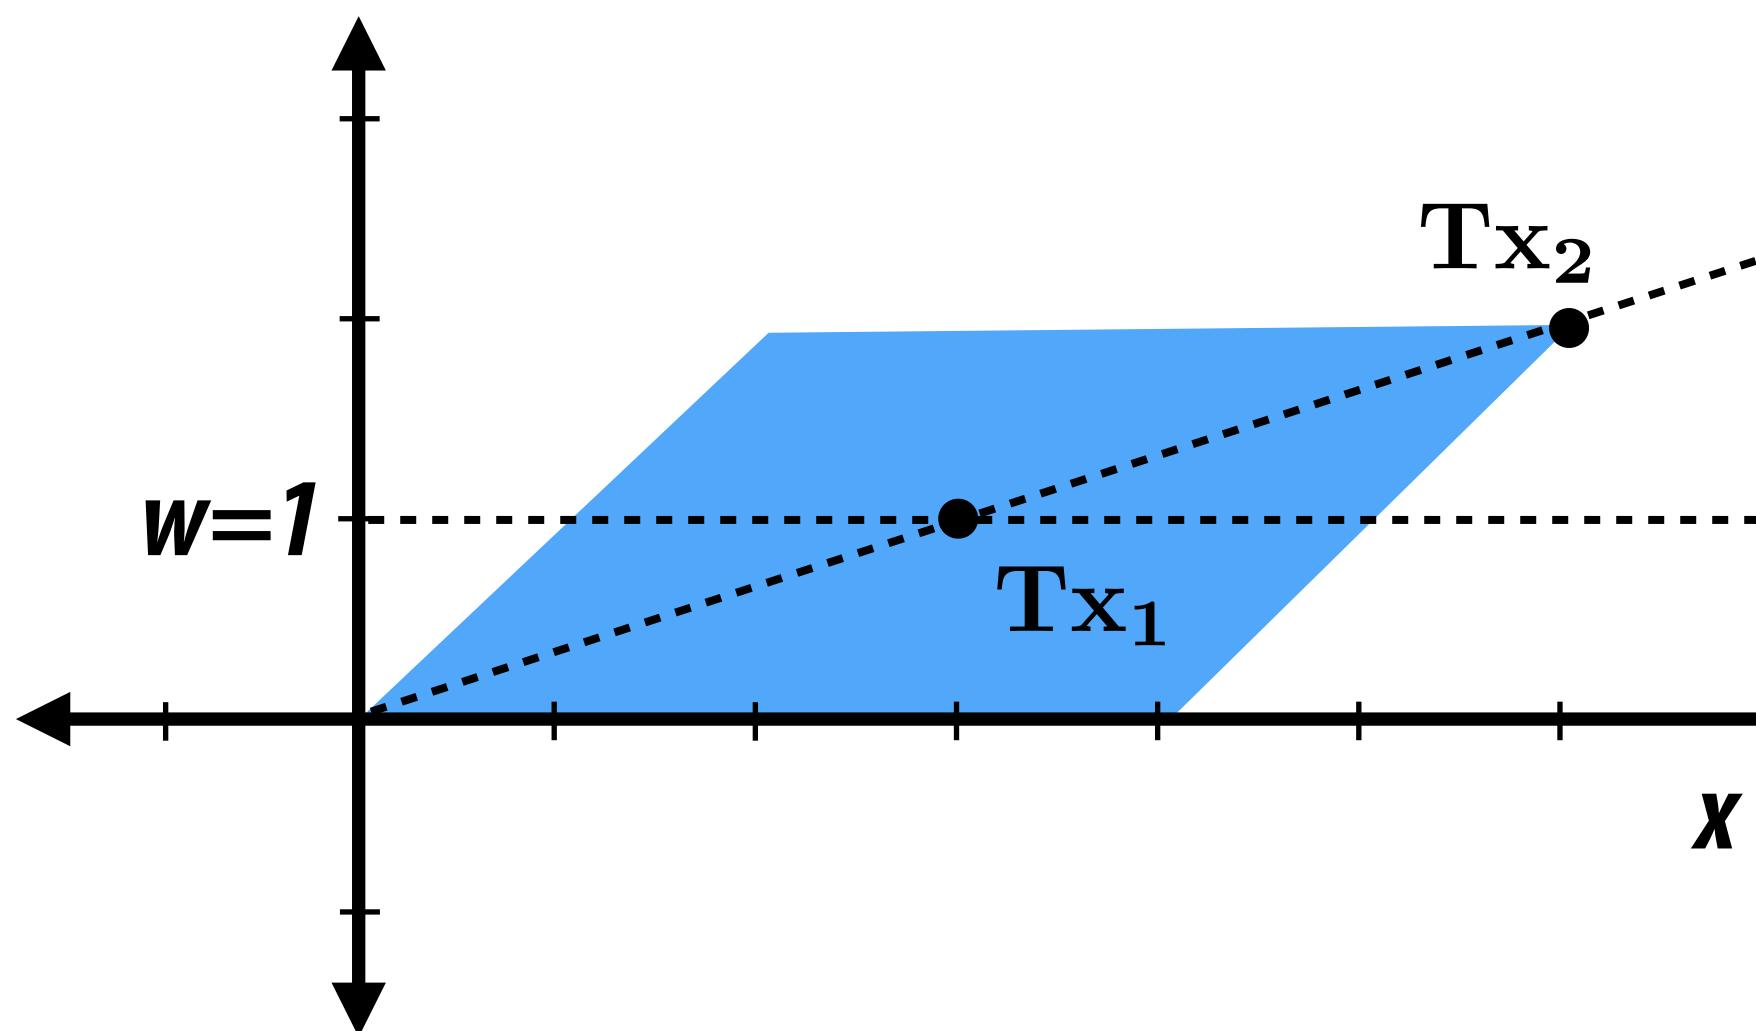


For simplicity, consider 1D-H:

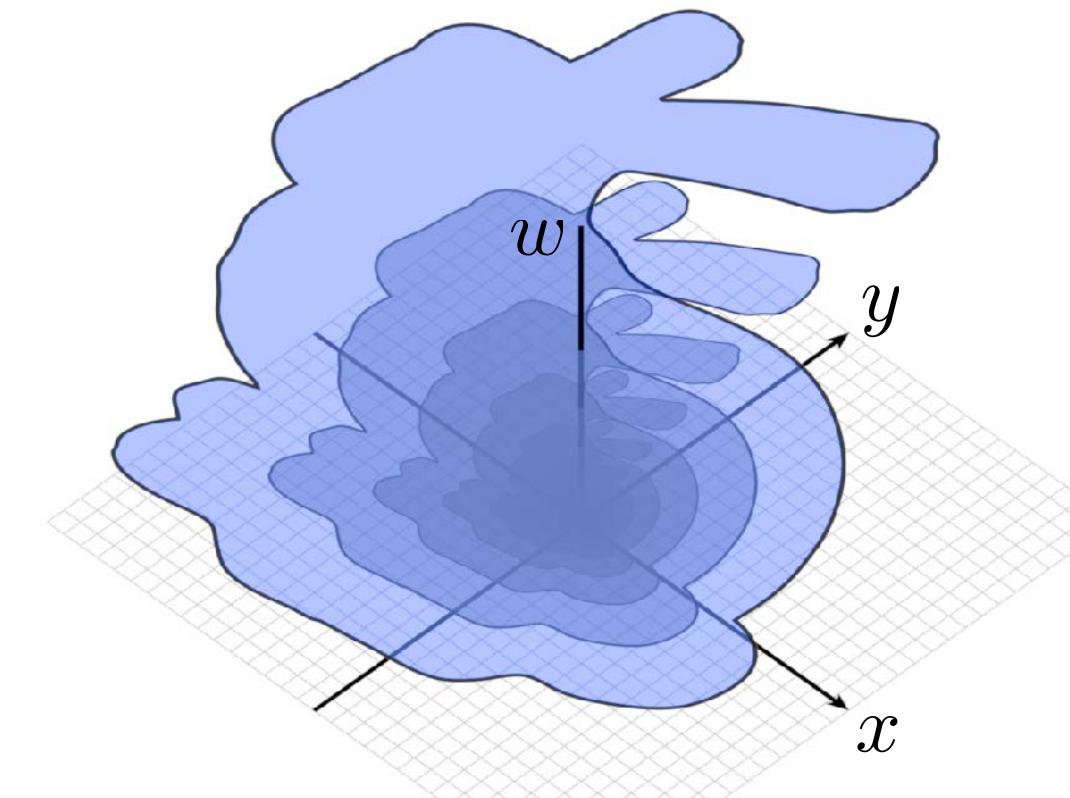
Translate by $t=2$: $T = \begin{bmatrix} 1 & t \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 \\ 0 & 1 \end{bmatrix}$

Recall: this is a shear in homogeneous x .

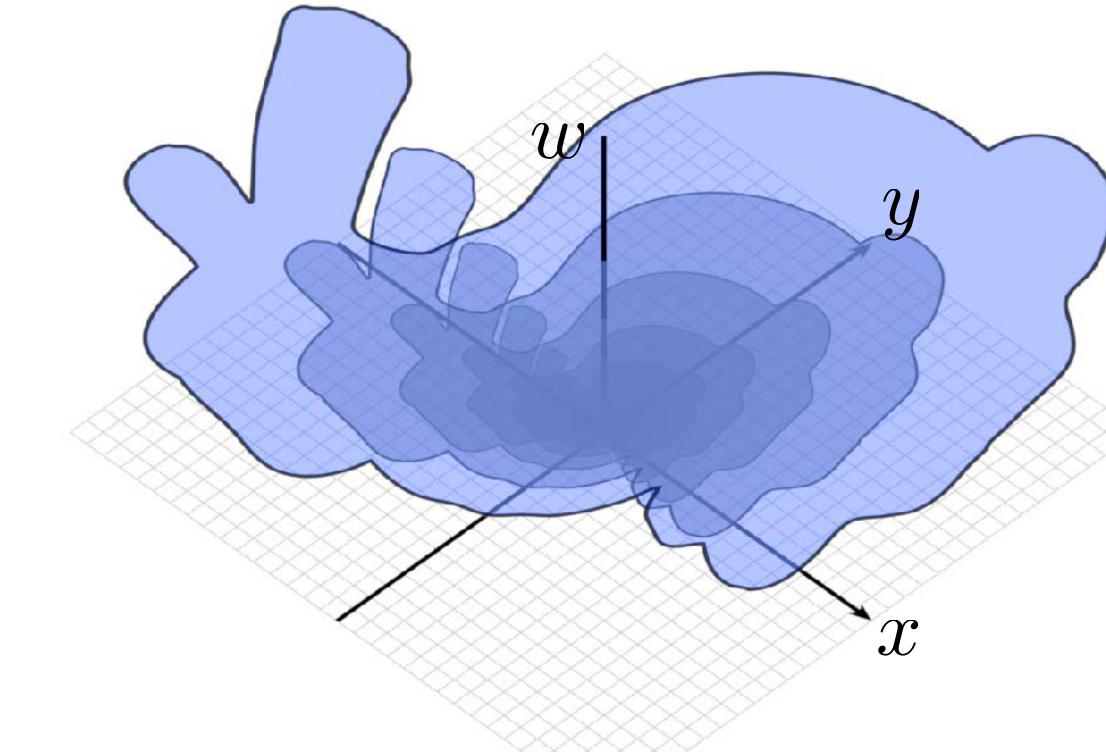
**1D translation is affine in 1D ($x + t$),
but it is linear in 1D-H**



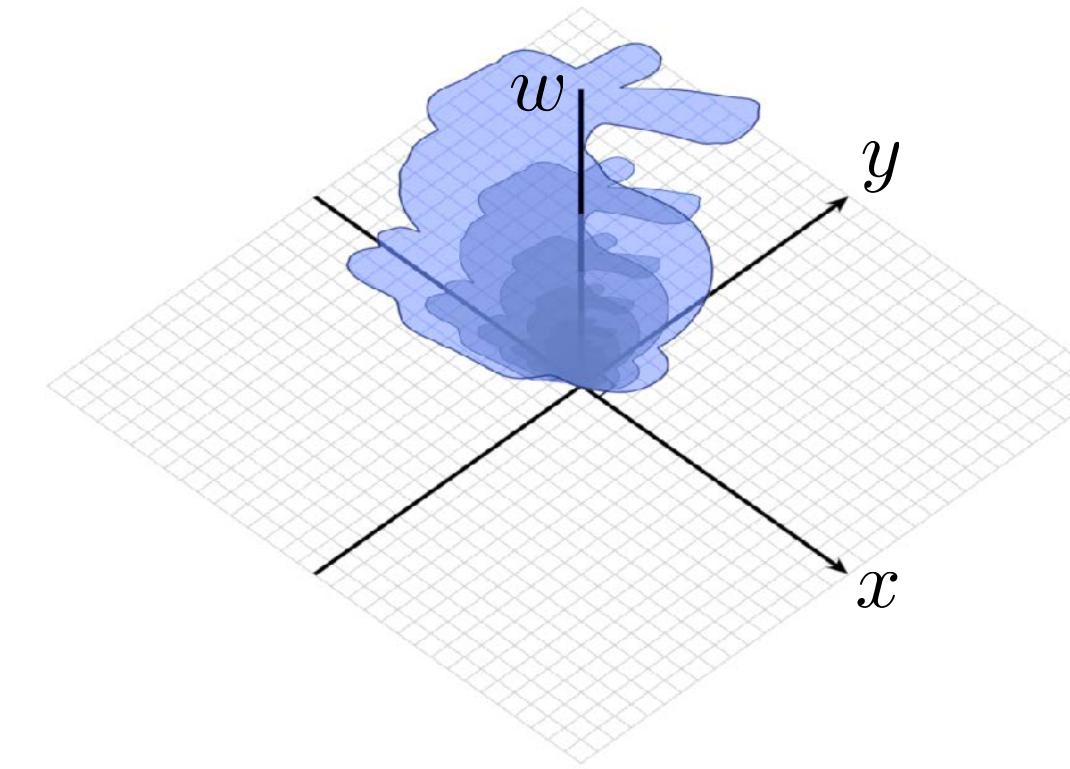
Visualizing 2D Transformations in 2D-H



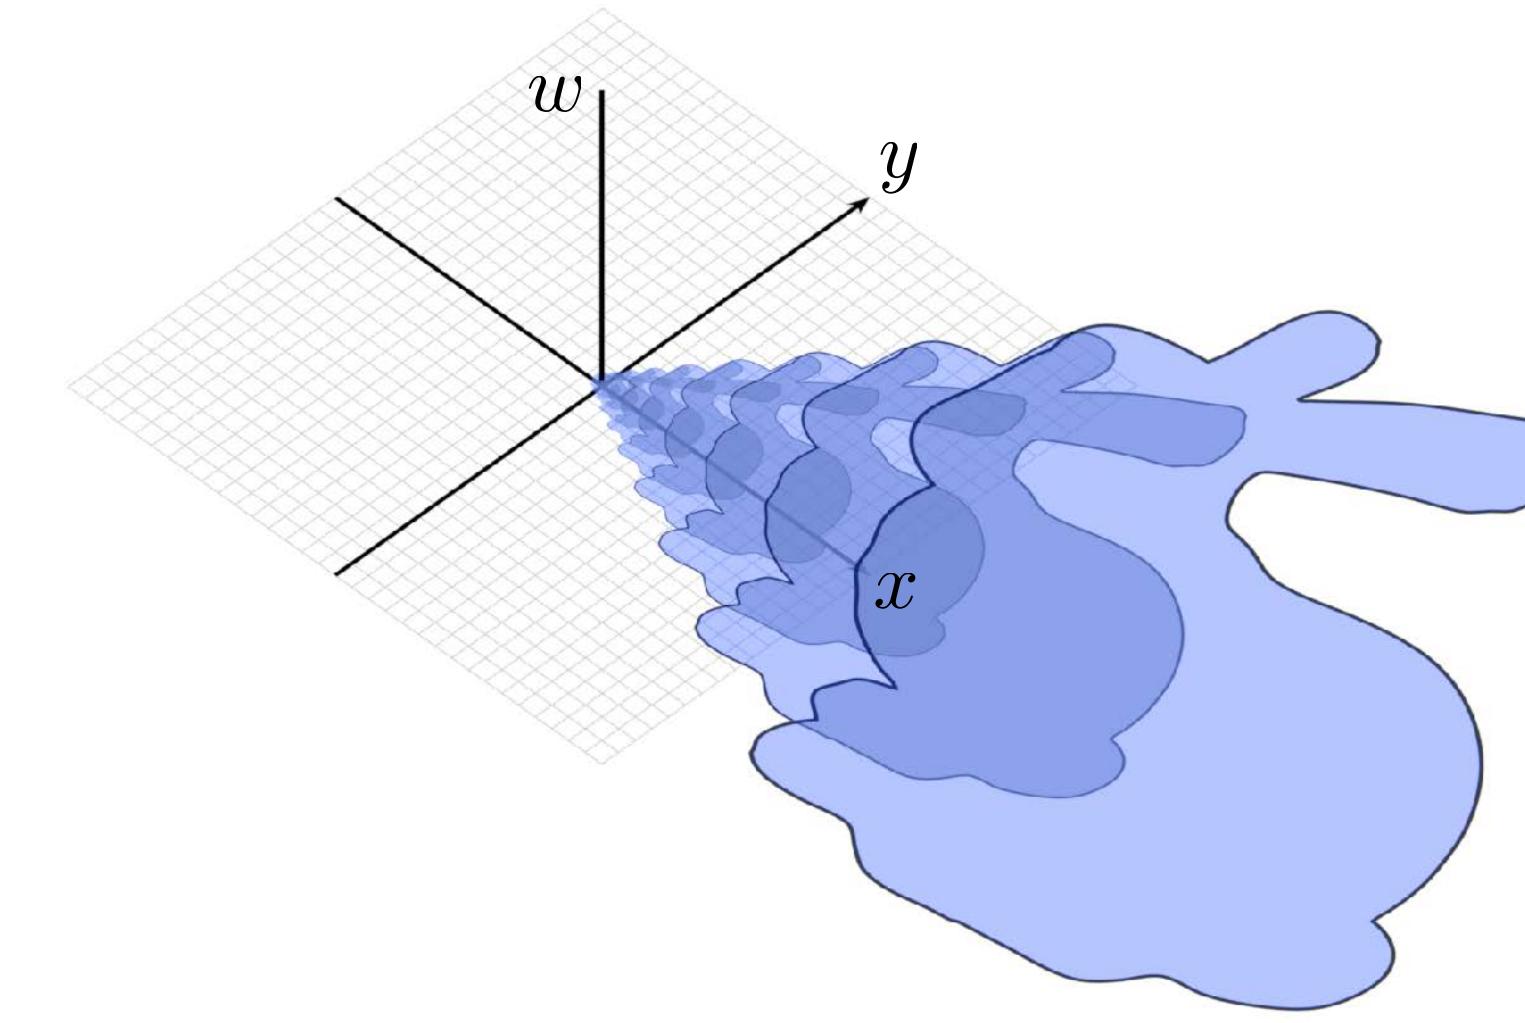
Original shape in 2D can be viewed as many copies, uniformly scaled by w.



2D rotation \leftrightarrow rotate around w



2D scale \leftrightarrow scale x and y; preserve w
(Question: what happens to 2D shape if you scale x, y, and w uniformly?)



**2D translate \leftrightarrow shear in 2D-H
(LINEAR!)**

Moving to 3D (and 3D-H)

Represent 3D transformations as 3x3 matrices and 3D-H transformations as 4x4 matrices

Scale:

$$\mathbf{S}_s = \begin{bmatrix} \mathbf{S}_x & 0 & 0 \\ 0 & \mathbf{S}_y & 0 \\ 0 & 0 & \mathbf{S}_z \end{bmatrix} \quad \mathbf{S}_s = \begin{bmatrix} \mathbf{S}_x & 0 & 0 & 0 \\ 0 & \mathbf{S}_y & 0 & 0 \\ 0 & 0 & \mathbf{S}_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Shear (in x, based on y,z position):

$$\mathbf{H}_{x,d} = \begin{bmatrix} 1 & d_y & d_z \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad \mathbf{H}_{x,d} = \begin{bmatrix} 1 & d_y & d_z & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Translate:

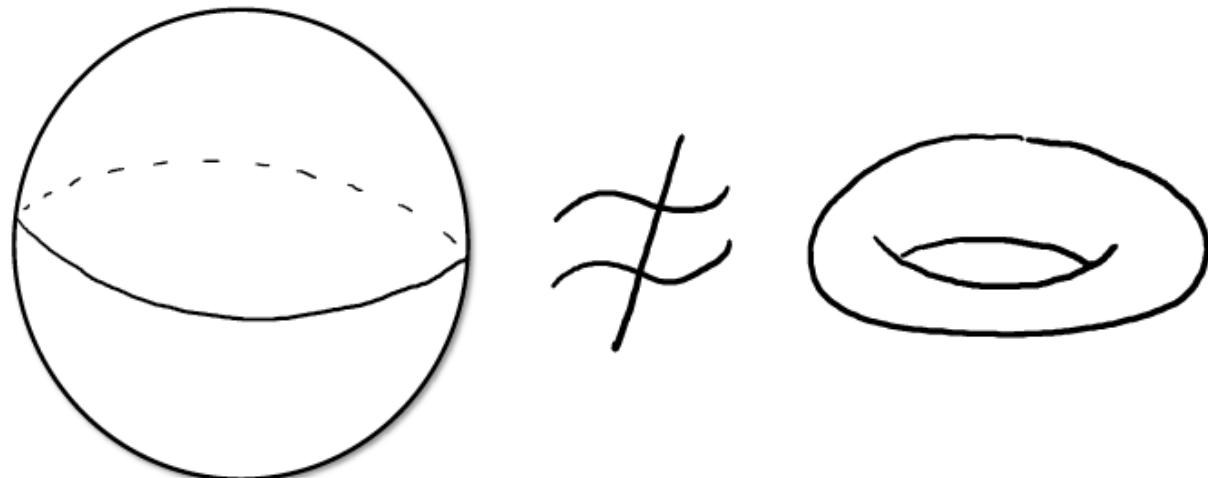
$$\mathbf{T}_b = \begin{bmatrix} 1 & 0 & 0 & b_x \\ 0 & 1 & 0 & b_y \\ 0 & 0 & 1 & b_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Outline

- Transformations and homogeneous coordinates
- *Rotation and $SO(n)$*
- 3D Rotation representation
 - Euler Angles
 - Axis-Angle
 - Quaternion

Prereq: Topology

- Topology: Structural Properties of a Manifold

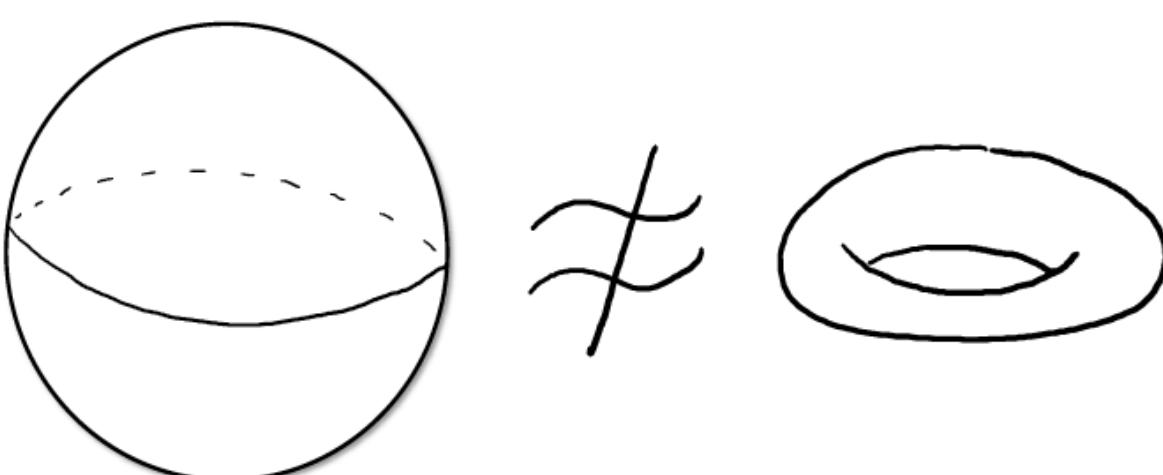
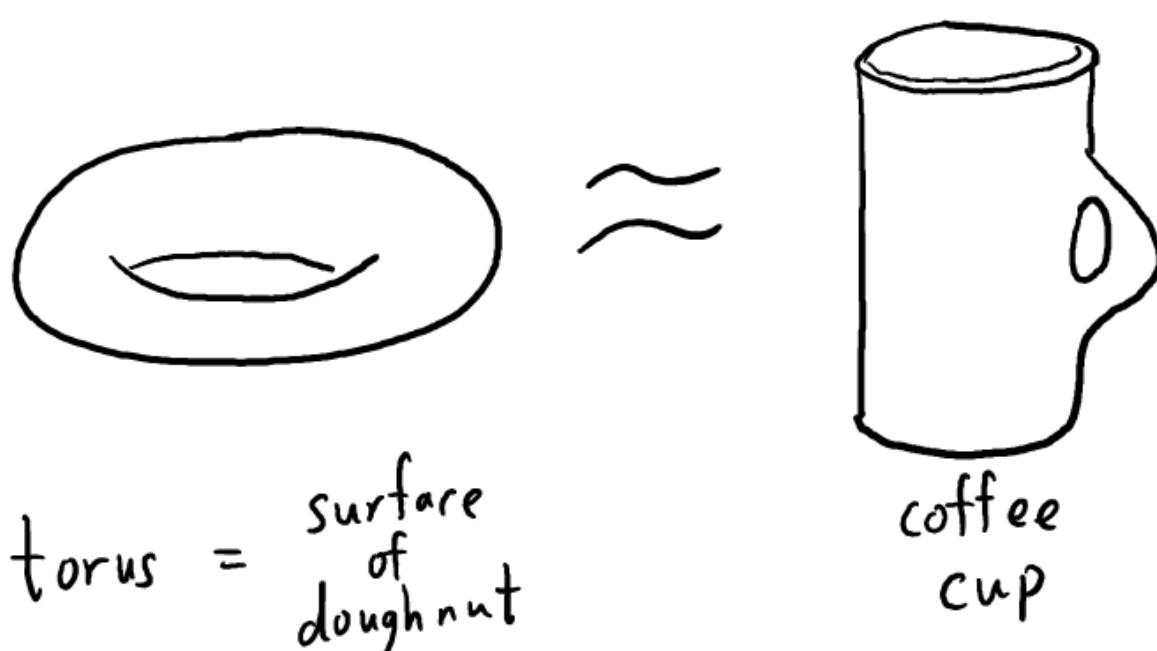
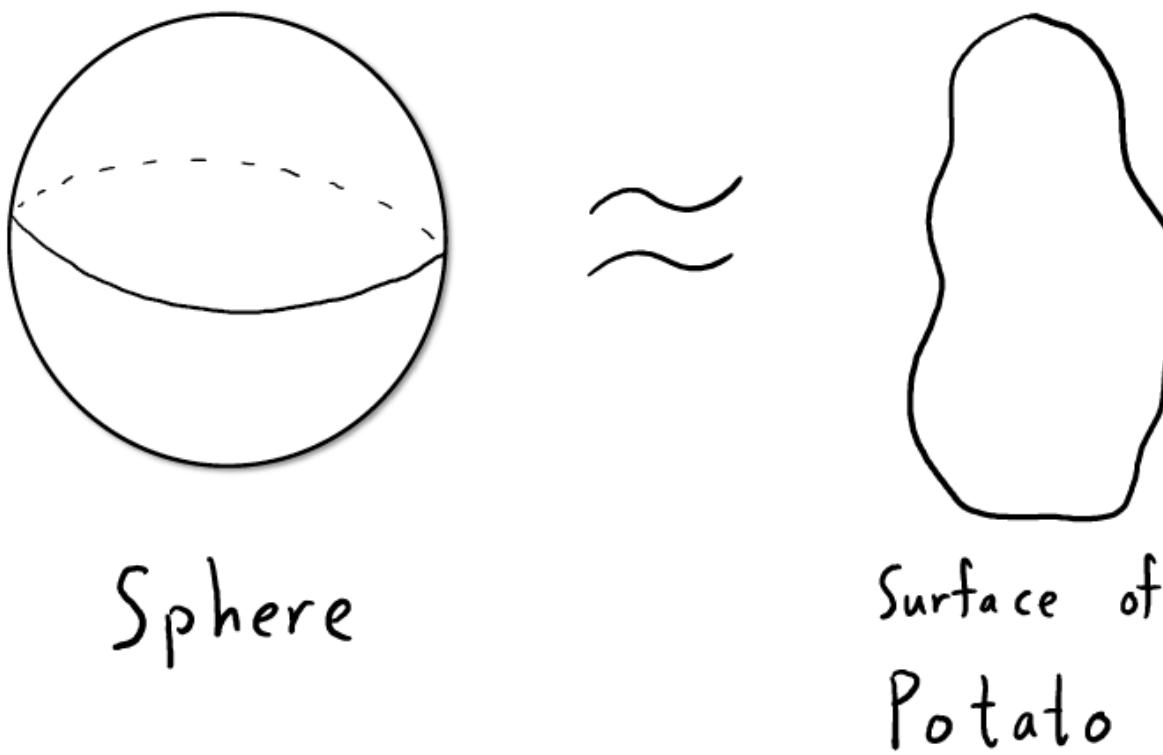


- Two surfaces M and N are *topologically equivalent* if there is a **differentiable bijection** between M and N



Prereq: Topology

- More examples:

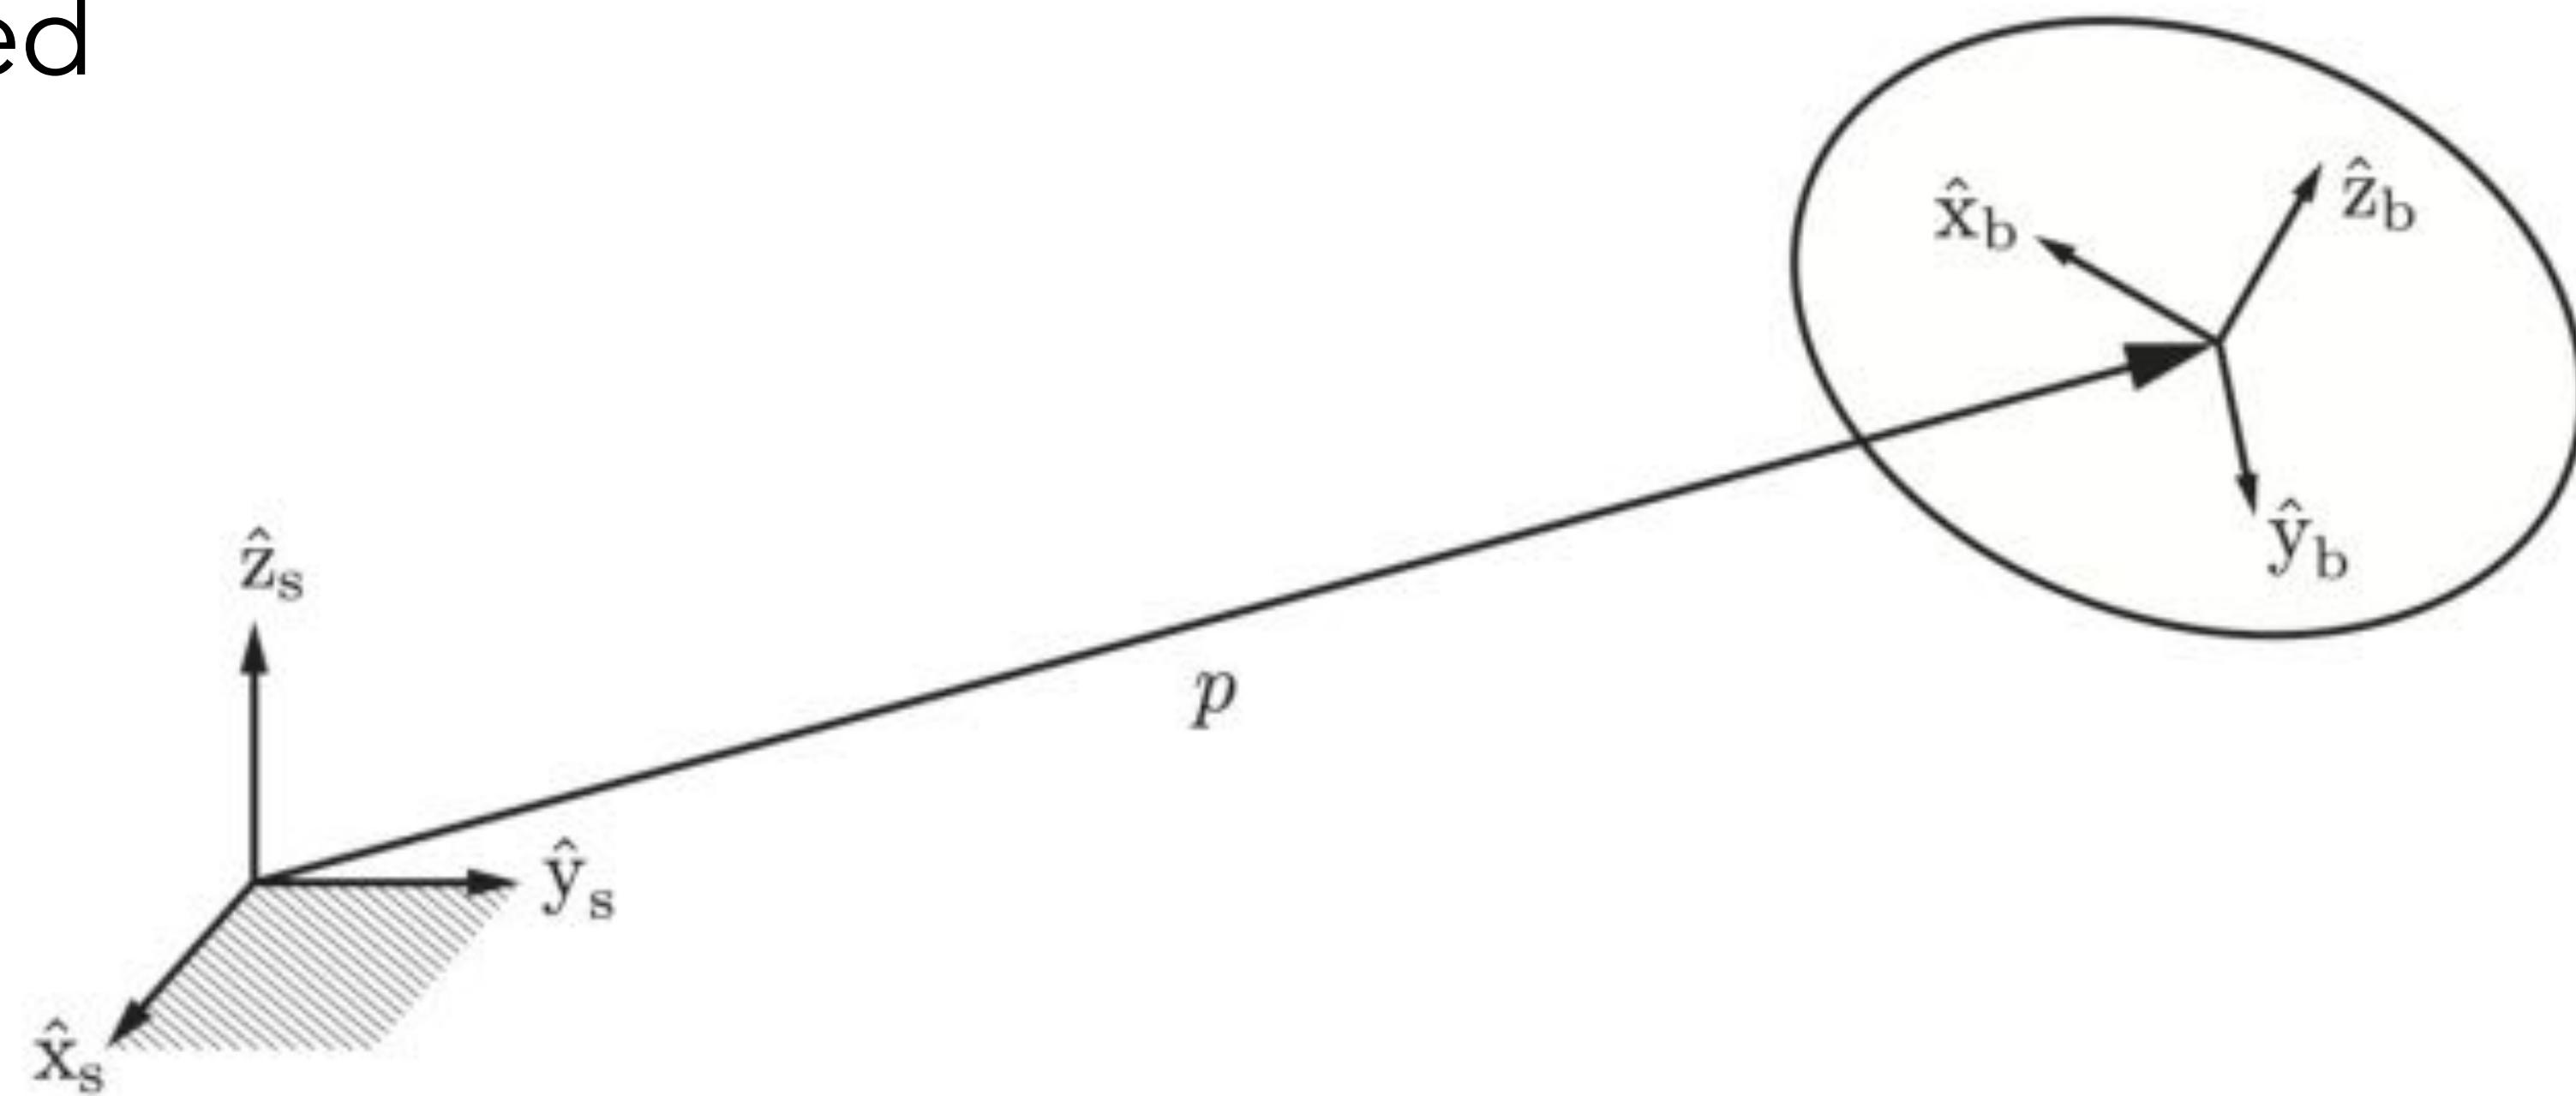


Orientation

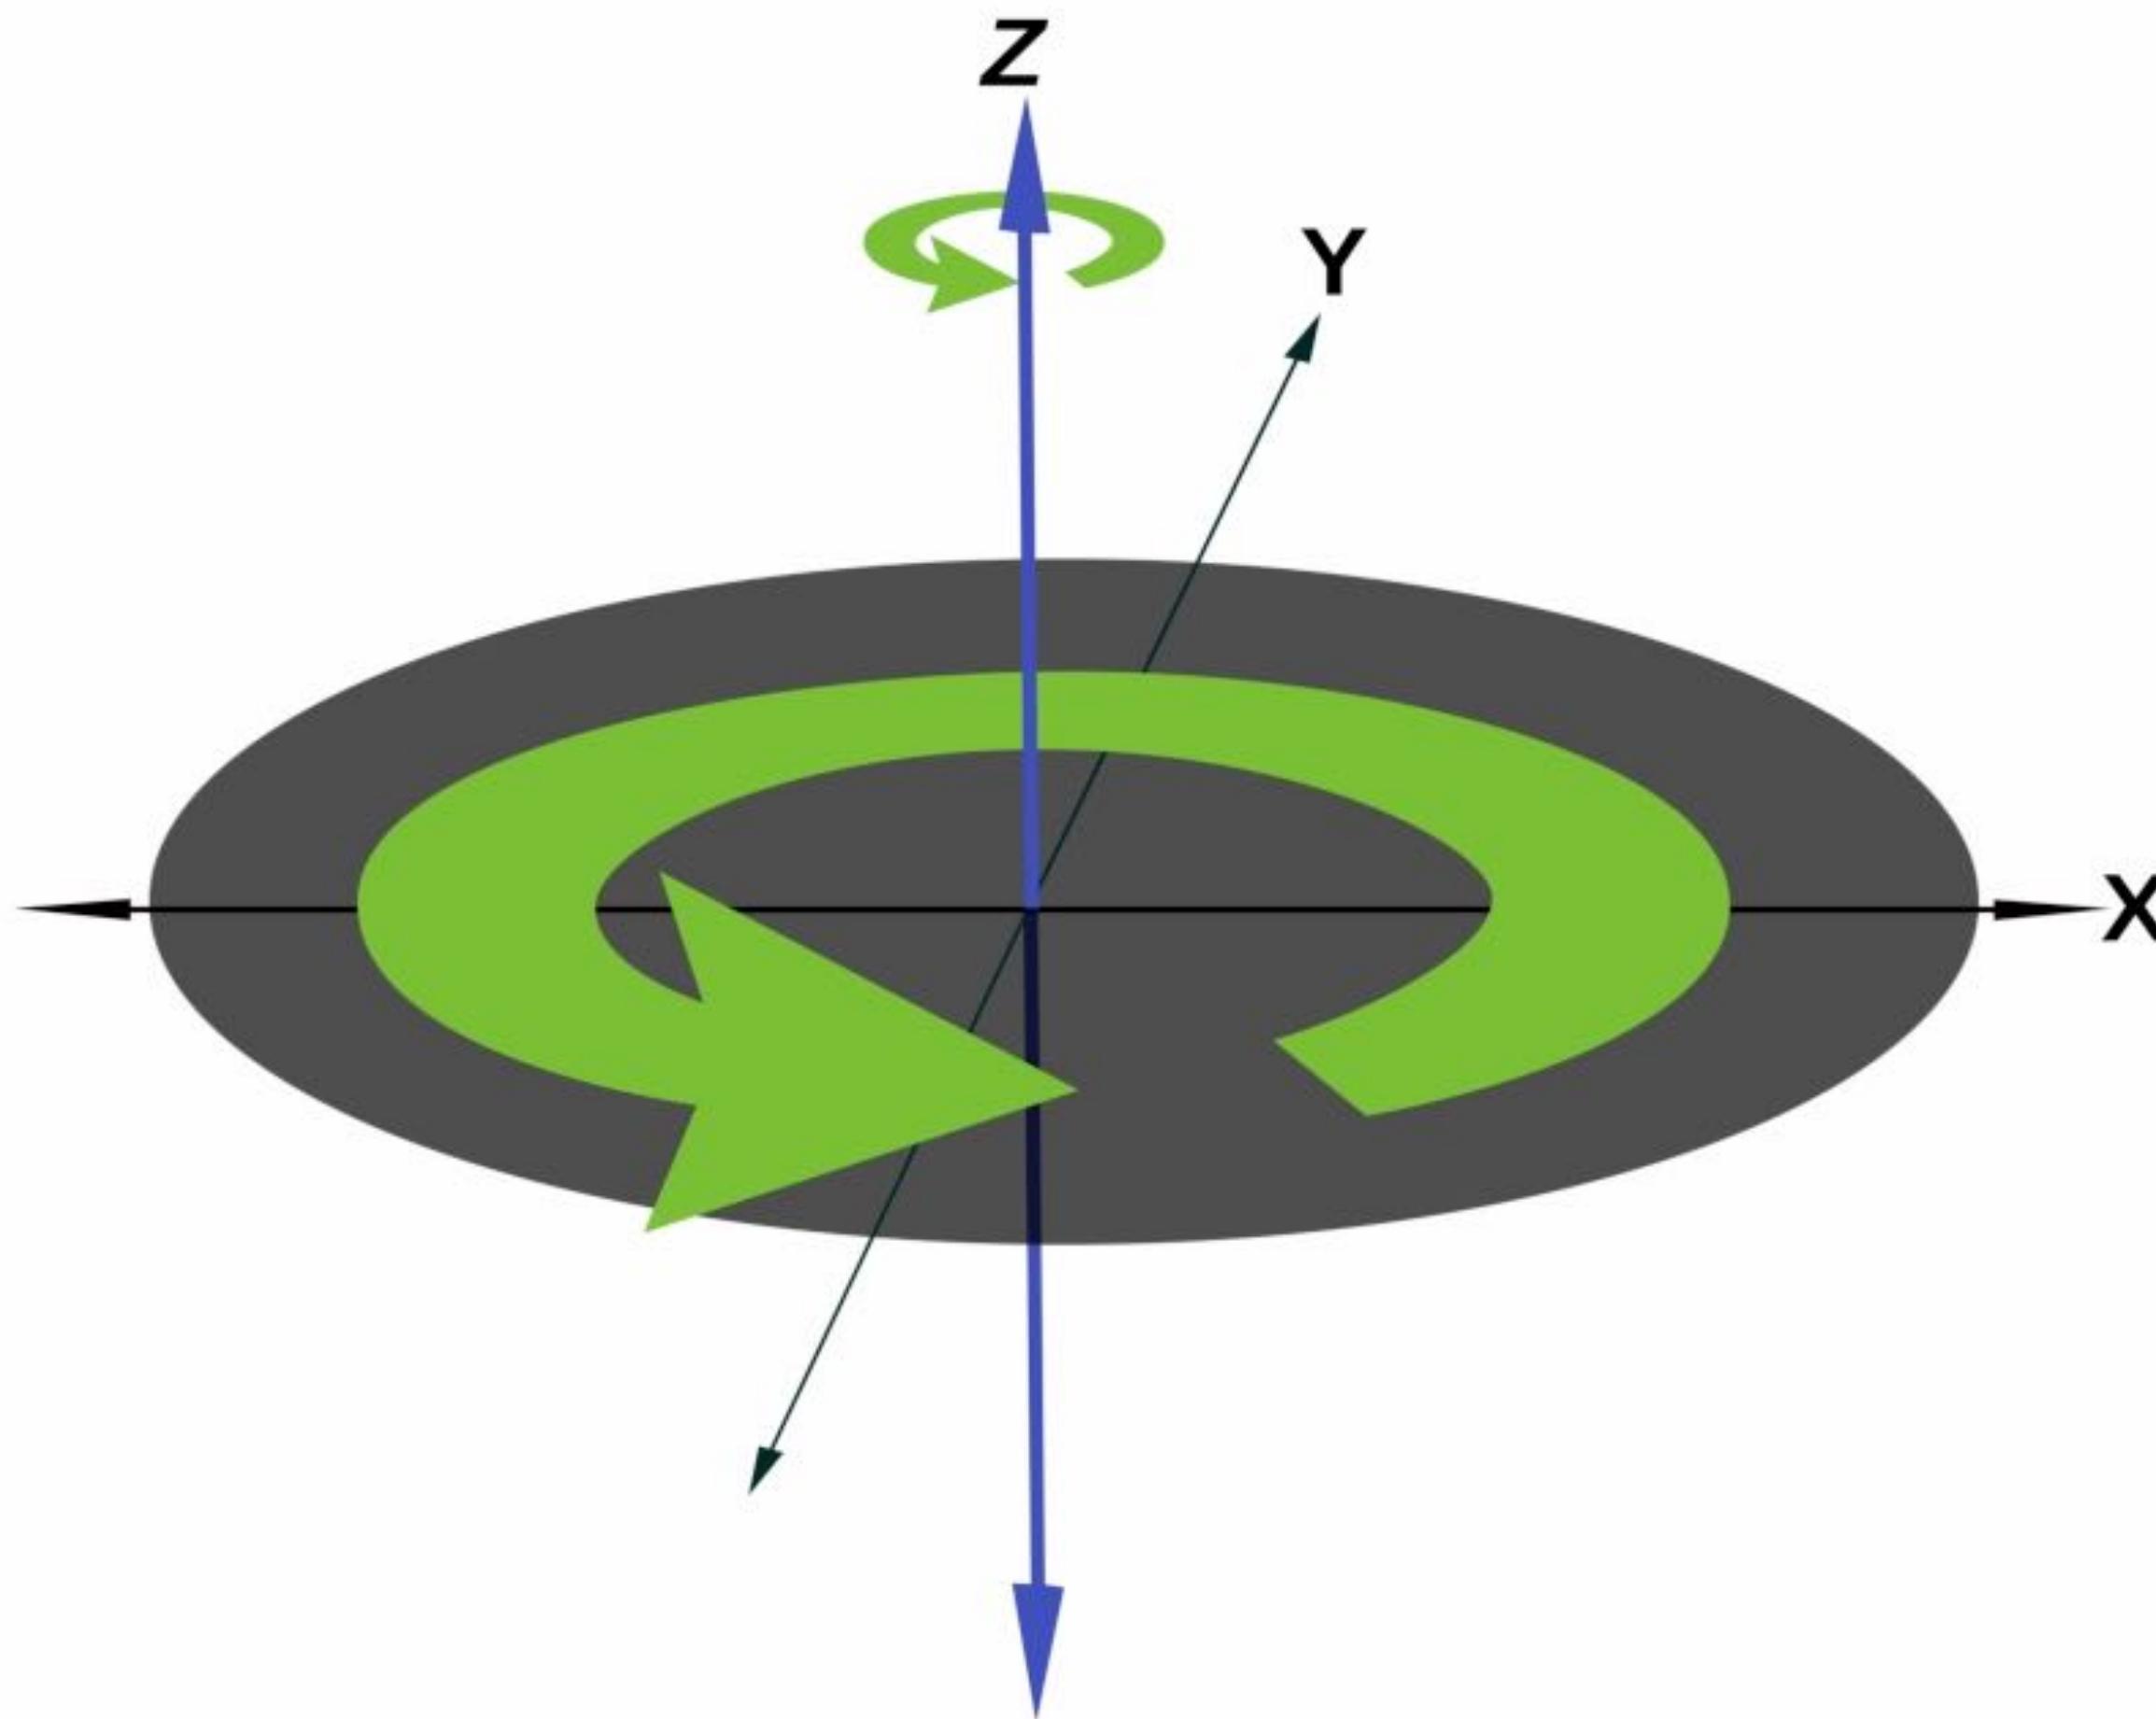
We use “rotation” to represent the relative orientation between two frames

For example,

- Space Frame: $\{ s \} = \{ x_s, y_s, z_s \}$
- Body Frame: $\{ b \} = \{ x_b, y_b, z_b \}$
- R_{sb} rotates the frame of the space to the frame of the body after the origins are aligned



Rotation in \mathbb{R}^2



1 Degree of Freedom

Rotation in \mathbb{R}^3



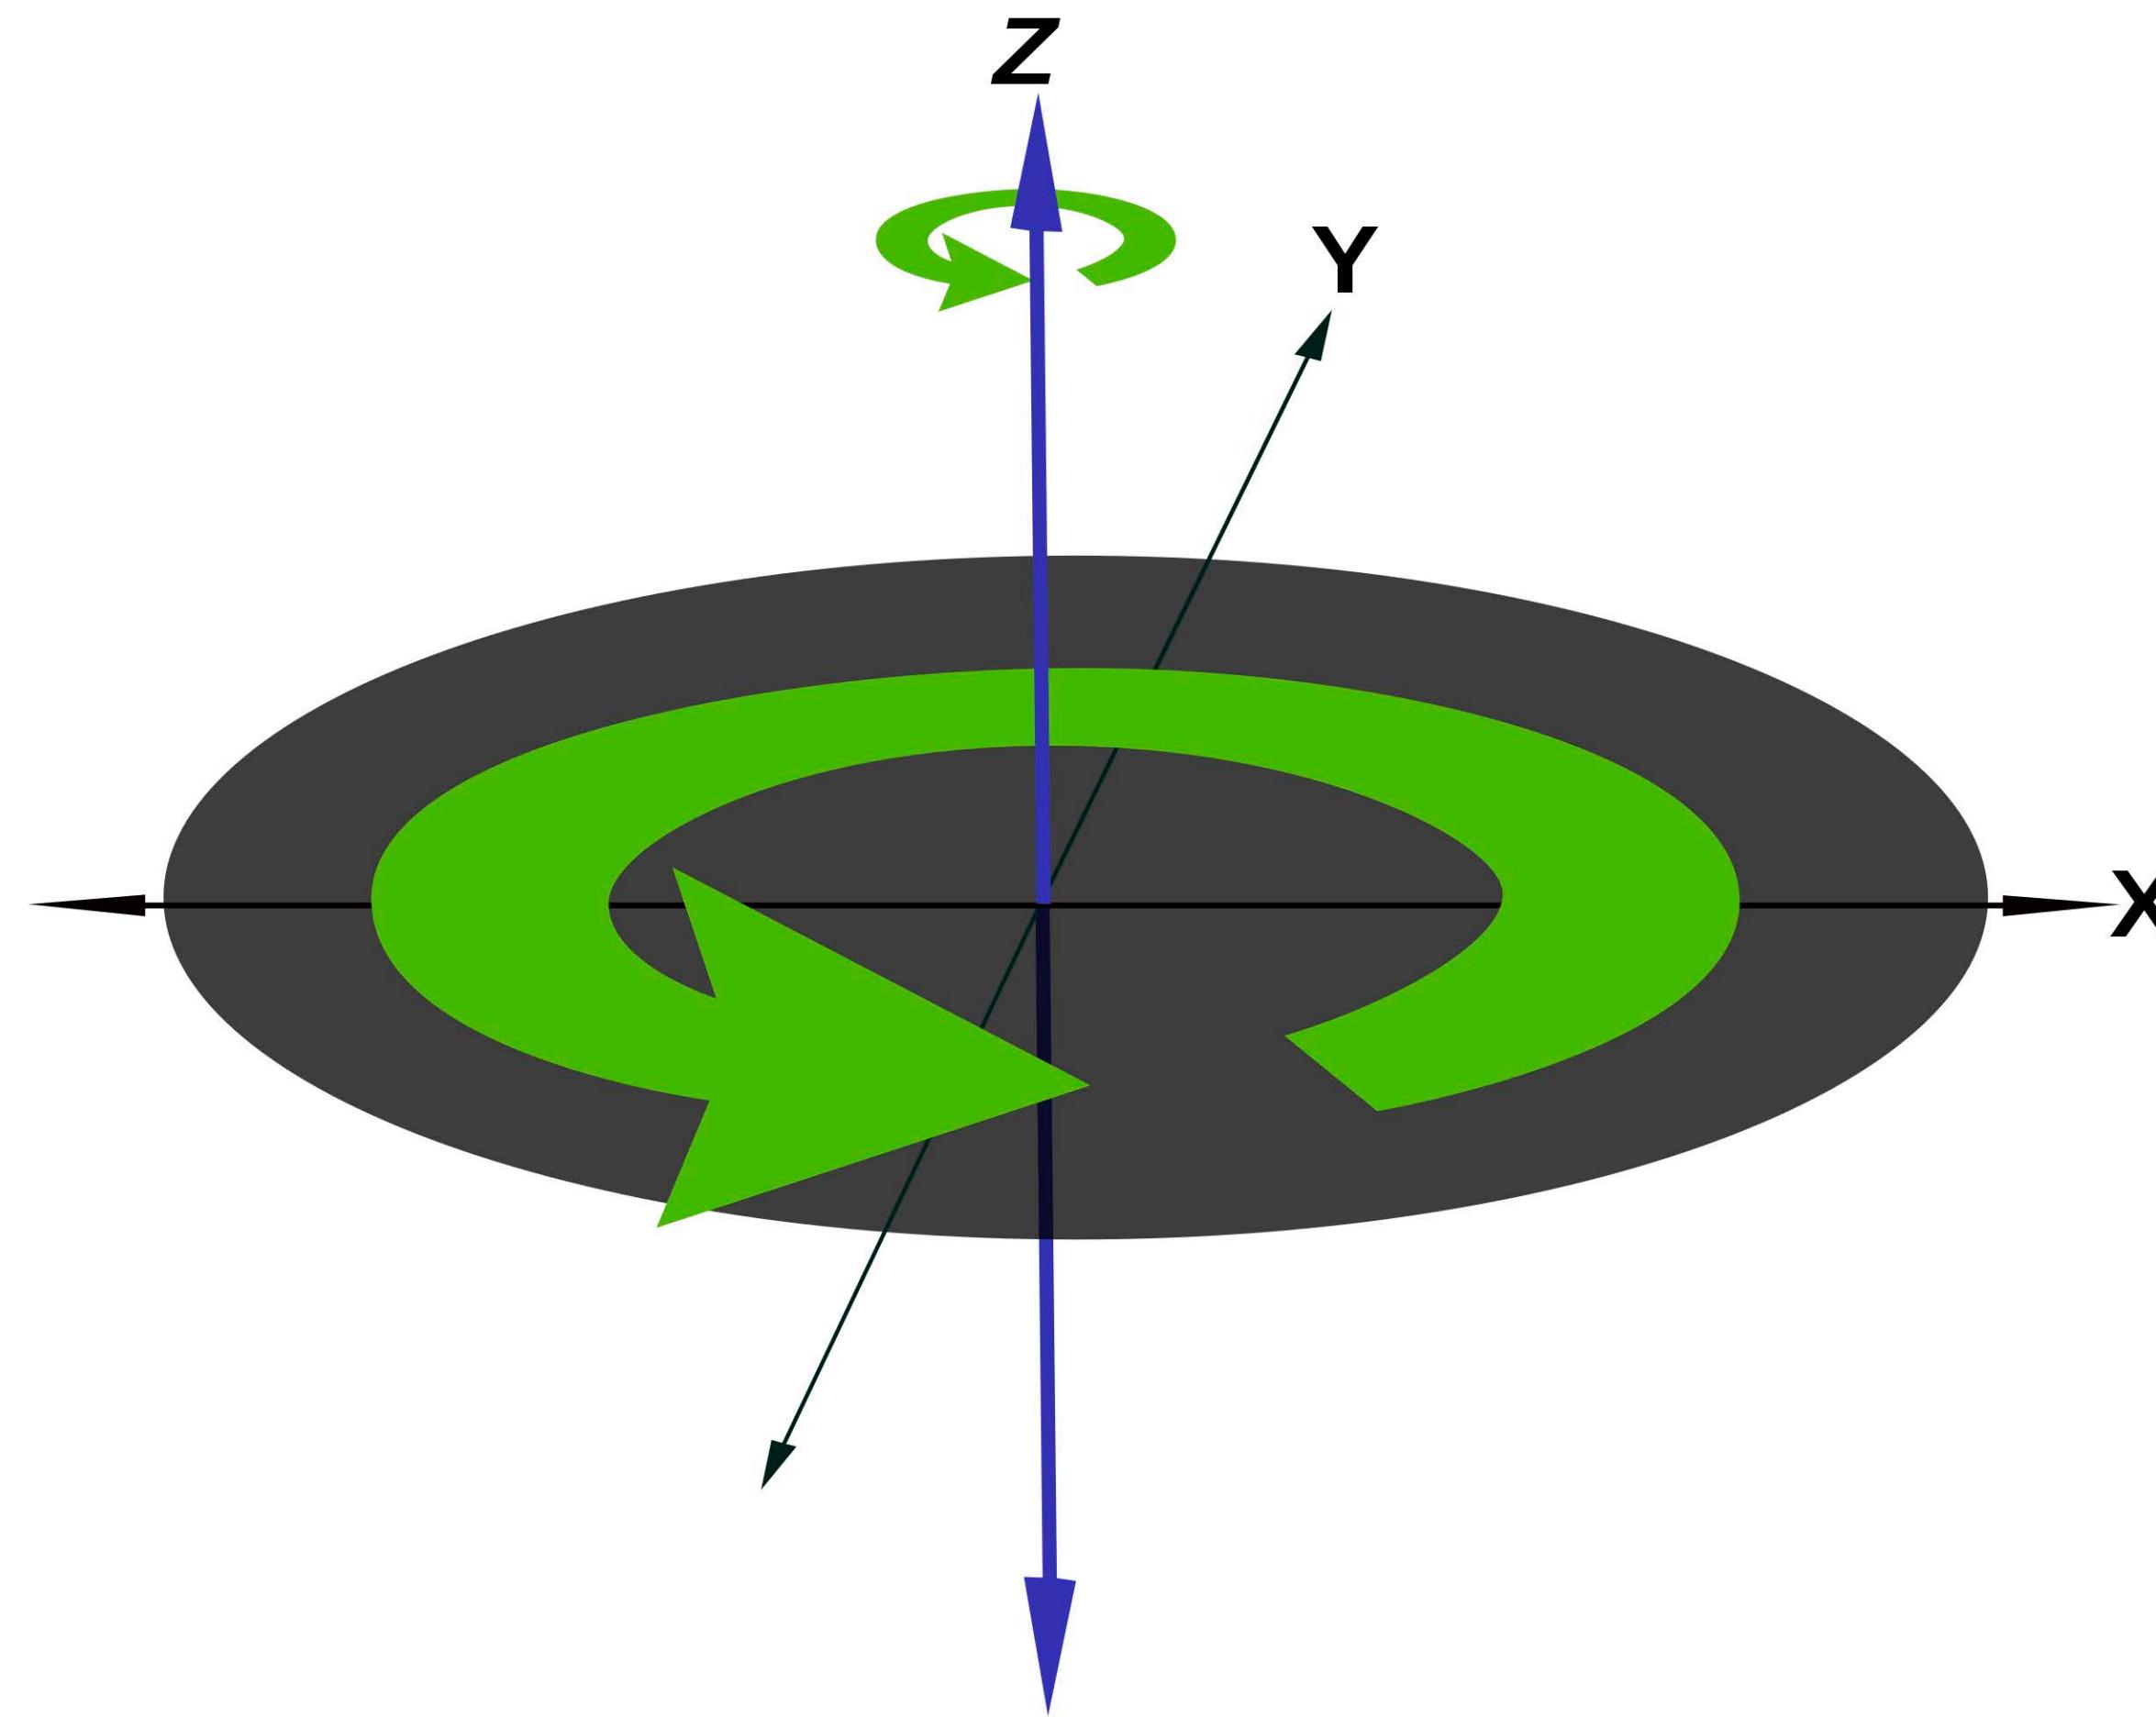
3 Degree of Freedoms

The Set of Rotations

- $SO(n) = \{R \in \mathbb{R}^{n \times n} : \det(R) = 1, RR^T = I\}$
- $SO(n)$: “Special Orthogonal Group”
 - “Group”: a group under the *matrix multiplication*
 - “Orthogonal”: $RR^T = I$
 - “Special”: $\det(R) = 1$
- $SO(2)$: 2D rotations, 1 DoF
- $SO(3)$: 3D rotations, 3 DoF

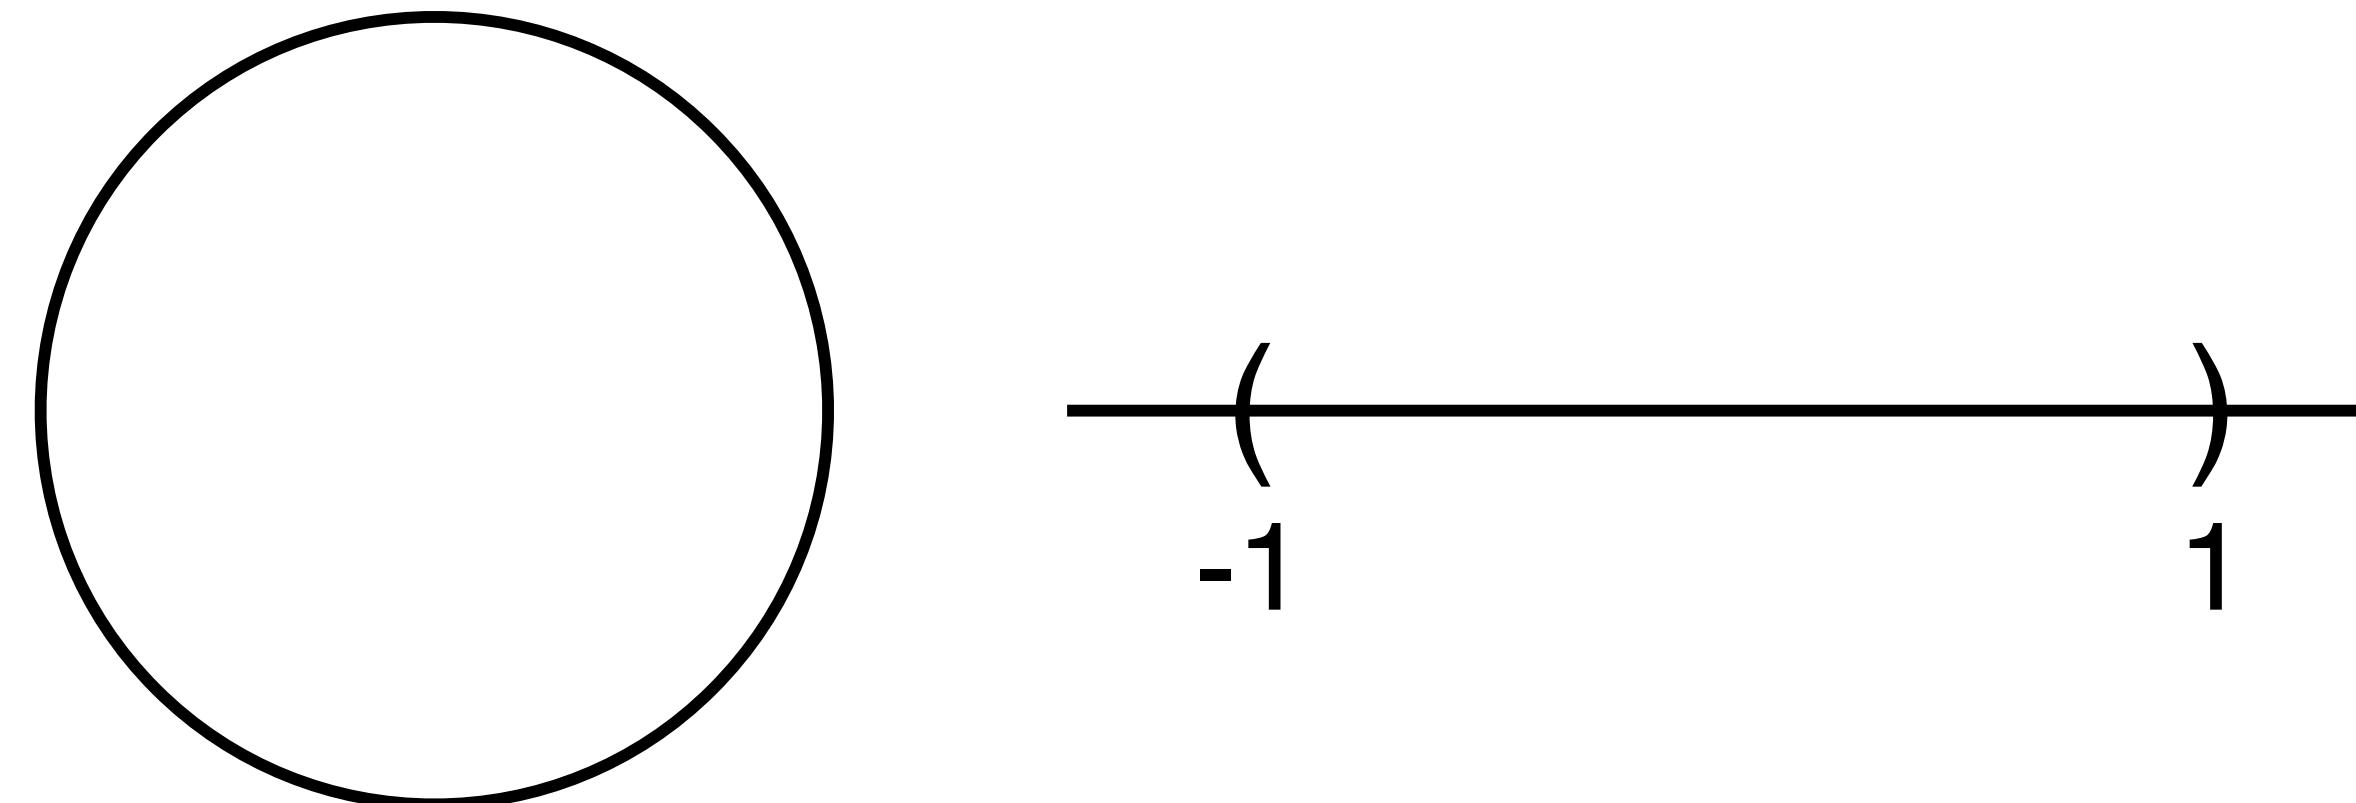
Topology of $SO(n)$

- The topology of $SO(2)$ is the same as a circle



Topology of $SO(n)$

- Circles do not have the same topology as $(-1,1)^n$
 \implies No differentiable bijections between $SO(2)$ and $(-1,1)^n$



- The topology of $SO(3)$ is also different from $(-1,1)^n$

Why do we care about the topology?

Parameterizing Rotation in Networks is Tricky

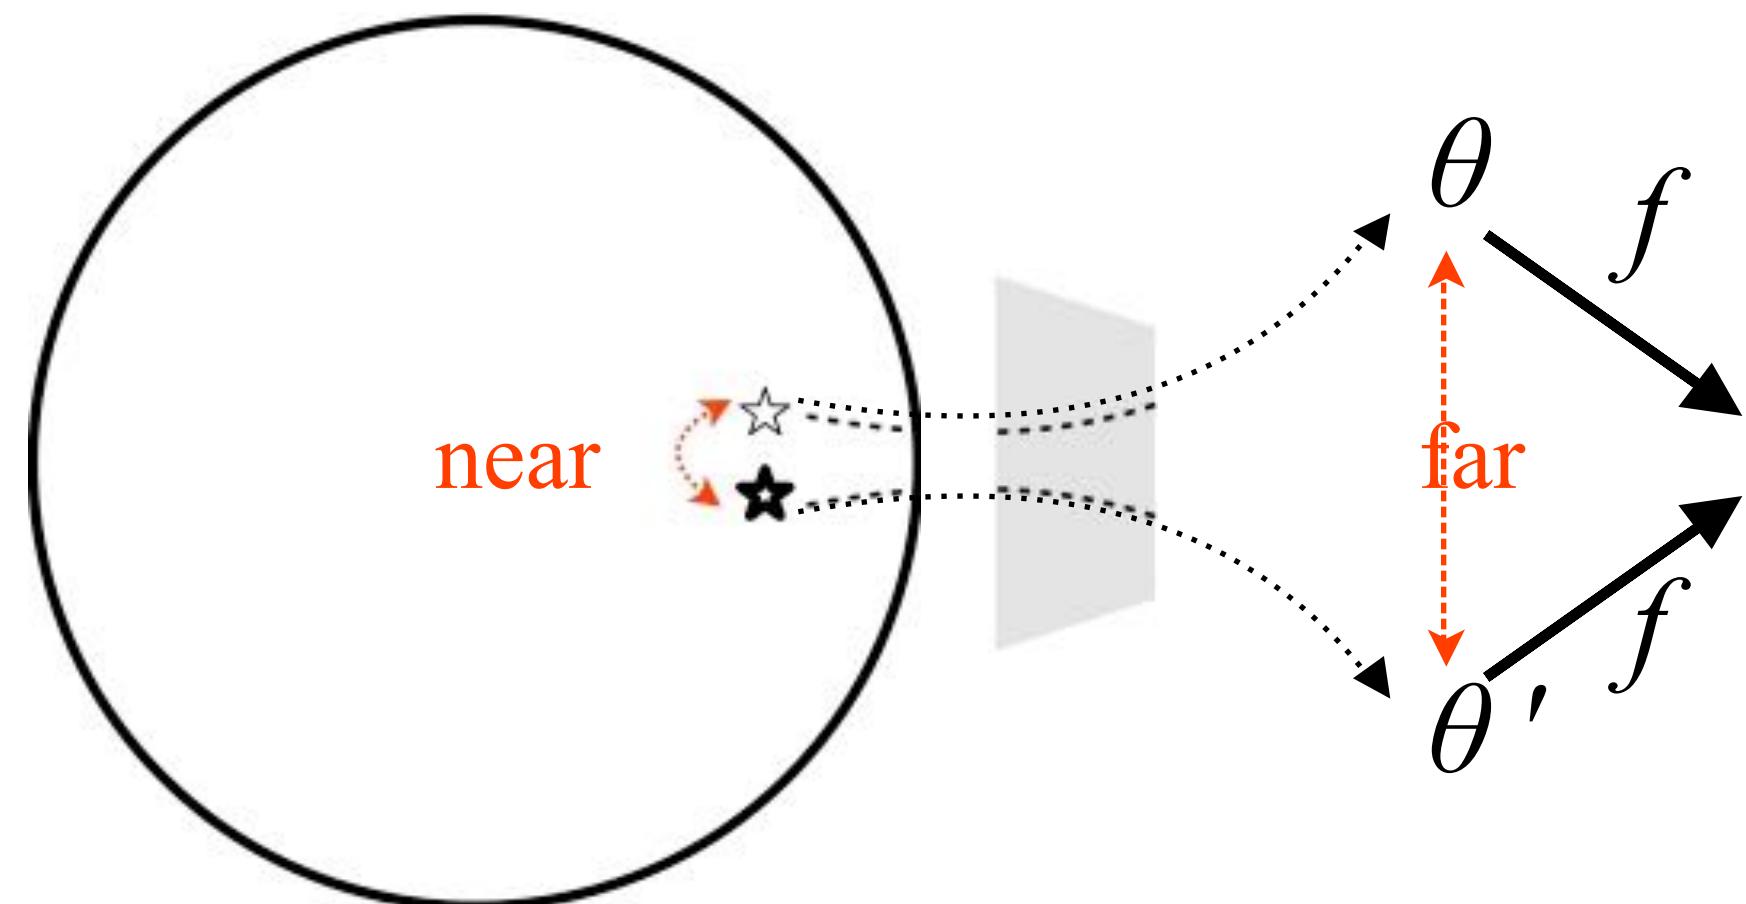
- An ideal parameterization $f(\theta) : U \mapsto SO(2)$ to use in networks:
 1. The domain is $(-l, l)^n$ (as network output)

Parameterizing Rotation in Networks is Tricky

- An ideal parameterization $f(\theta) : U \mapsto SO(2)$ to use in networks:

1. The domain is $(-l, l)^n$ (as network output)
2. f is a differentiable *bijection*

Otherwise:



- If input data points to network are close, but the θ predictions happen to be far after convergence, the network (a continuous function) will make awful predictions between the two data points !
- Need special network design to overcome the issue

Outline

- Transformations and homogeneous coordinates
- Rotation and $SO(n)$
- 3D Rotation representation
 - *Euler Angles*
 - Axis-Angle
 - Quaternion

Euler Angle is Very Intuitive



Euler Angle to Rotation Matrix

- Rotation about principal axis is represented as:

$$R_x(\alpha) := \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha \\ 0 & \sin \alpha & \cos \alpha \end{bmatrix}$$

$$R_y(\beta) := \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix}$$

$$R_z(\gamma) := \begin{bmatrix} \cos \gamma & -\sin \gamma & 0 \\ \sin \gamma & \cos \gamma & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- $R = R_z(\gamma)R_y(\beta)R_x(\alpha)$ for arbitrary rotation

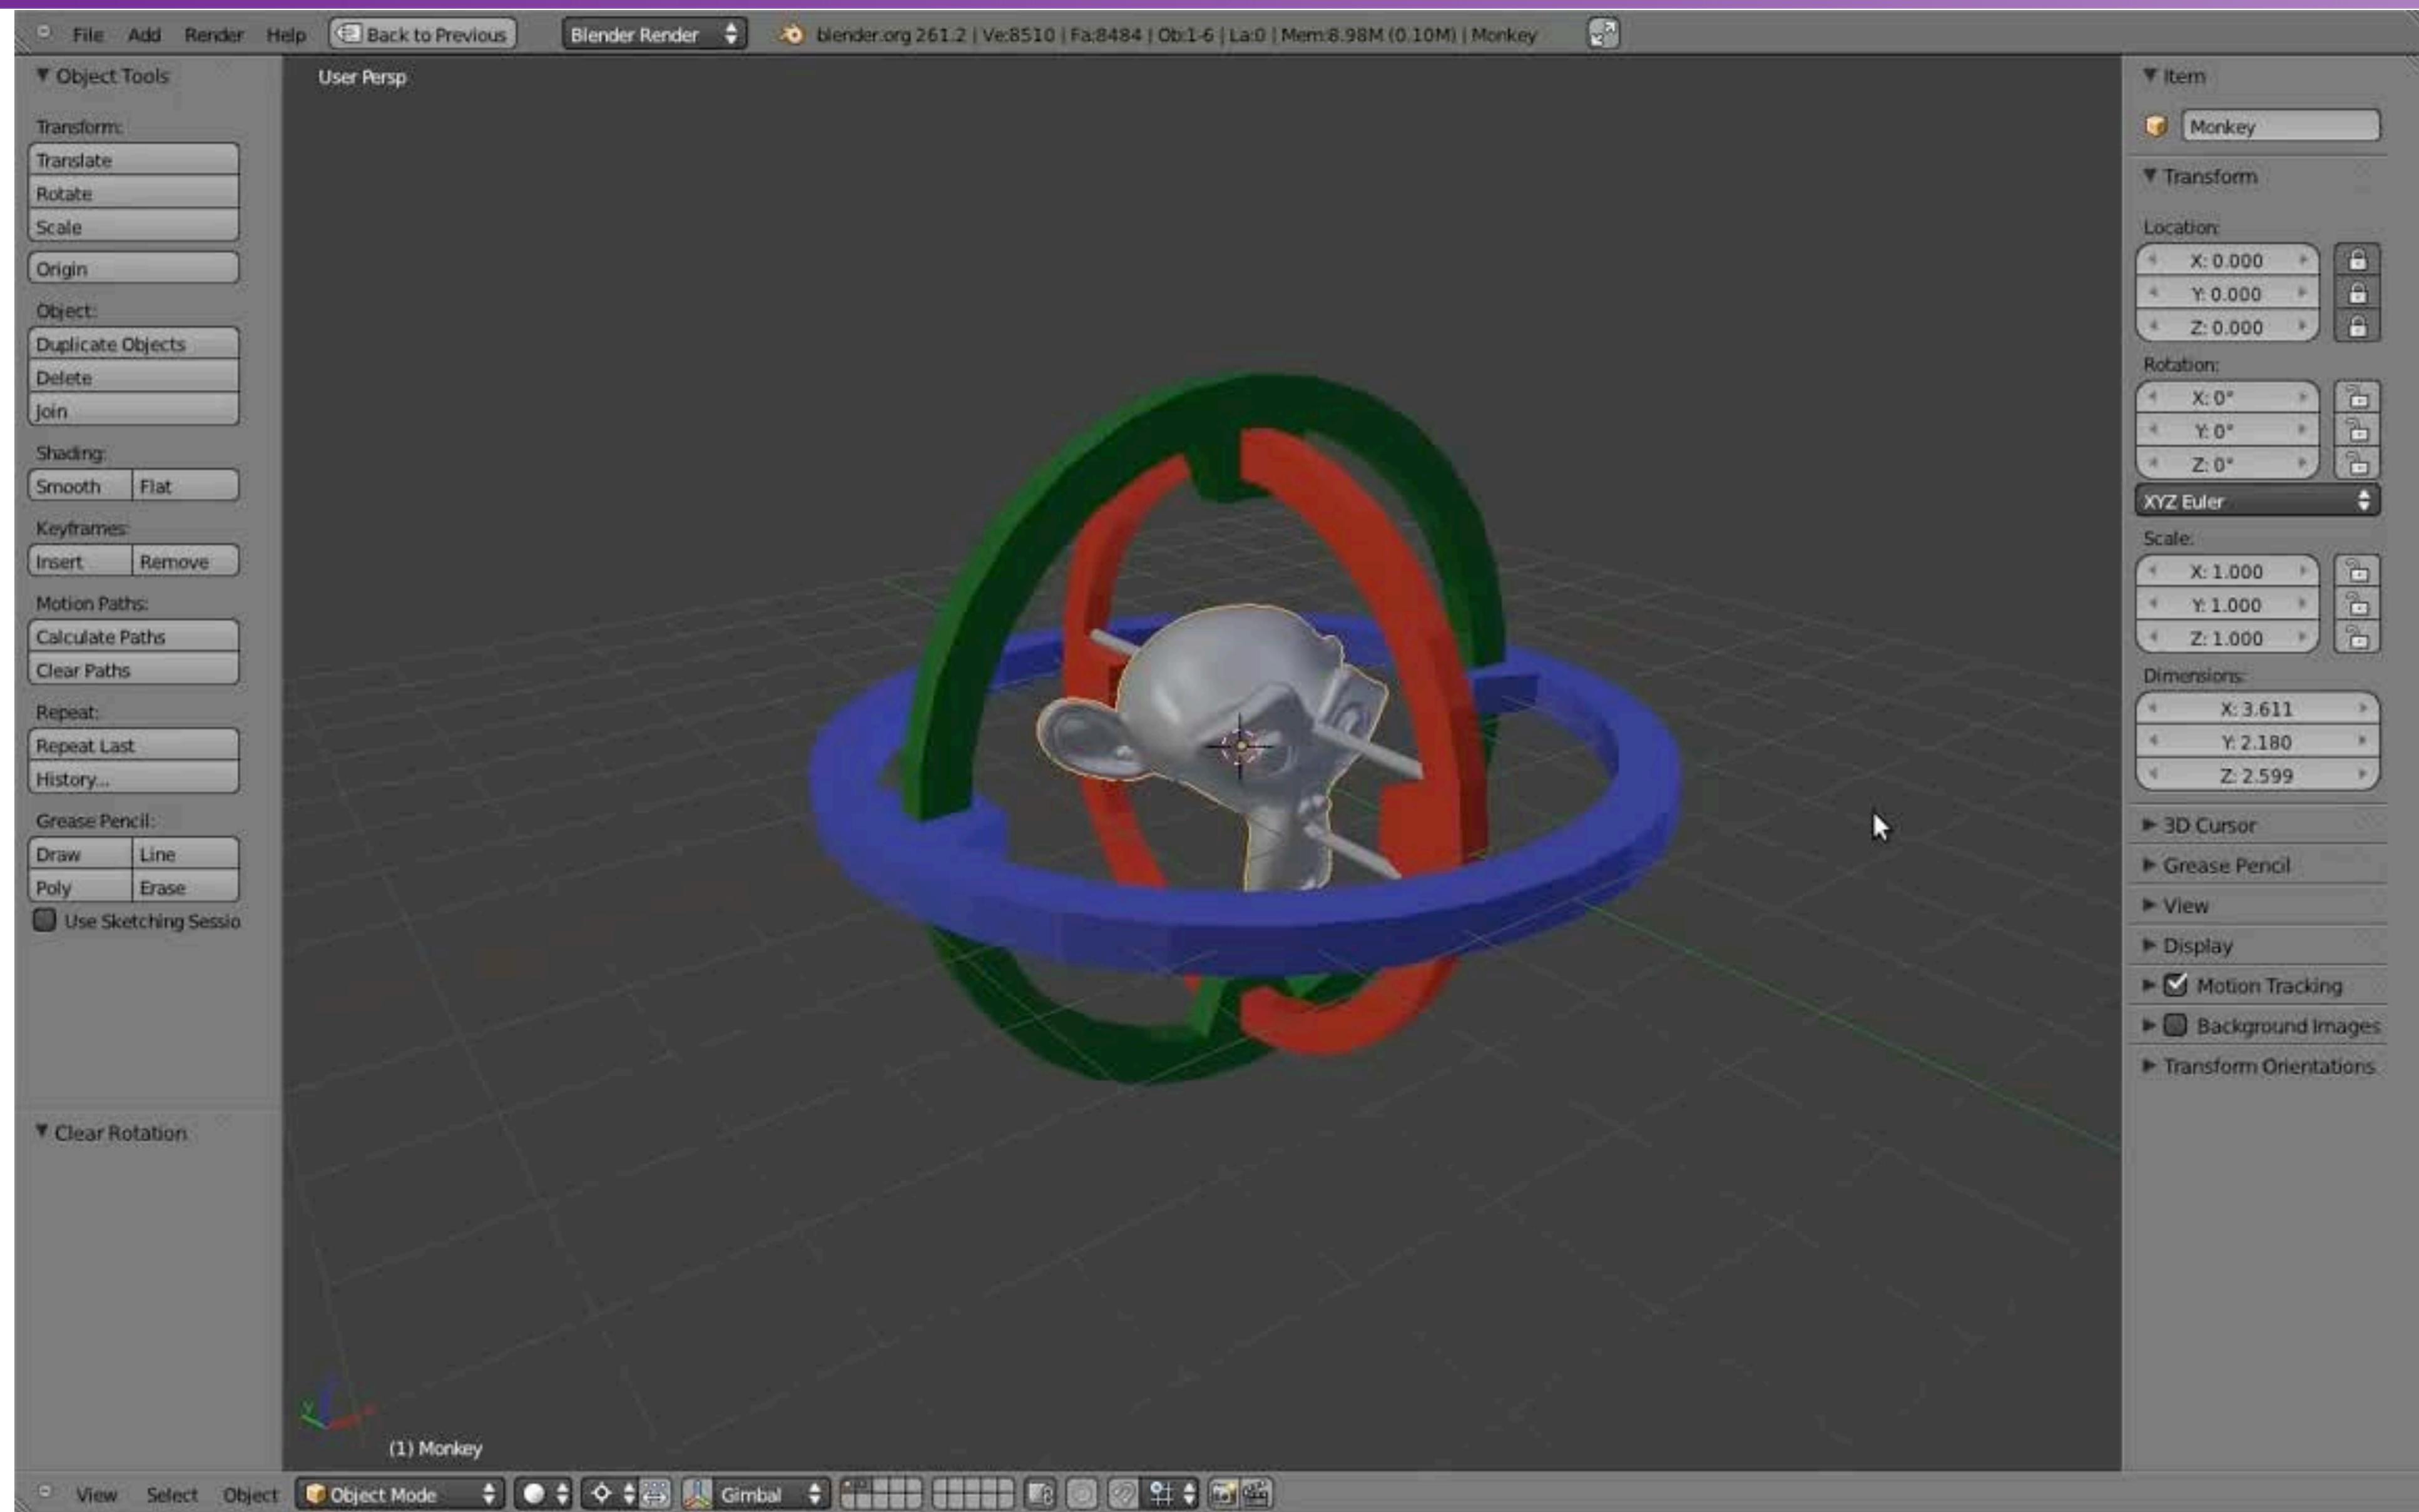
Gimbal Lock

- Euler Angle is **not unique** for some rotations. For example,

$$R_z(45^\circ)R_y(90^\circ)R_x(45^\circ) = R_z(90^\circ)R_y(90^\circ)R_x(90^\circ)$$

$$= \begin{bmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ -1 & 0 & 0 \end{bmatrix}$$

Gimbal Lock



Gimbal Lock

- For example: When $\beta = \pi/2$,

$$\begin{aligned} R &= R_z(\gamma)R_y(\beta)R_x(\alpha) \\ &= \begin{bmatrix} 0 & 0 & 1 \\ \sin(\alpha + \gamma) & \cos(\alpha + \gamma) & 0 \\ -\cos(\alpha + \gamma) & \sin(\alpha + \gamma) & 0 \end{bmatrix} \end{aligned}$$

since changing α and γ has the same effects, a degree of freedom disappears!

Summary

- Euler angle can parameterize every rotation and has good interpretability
- It is not a unique representation at some points
- There are some points where not every change in the target space (rotations) can be realized by a change in the source space (Euler angles)

Outline

- Transformations and homogeneous coordinates
- Rotation and $SO(n)$
- 3D Rotation representation
 - Euler Angles
 - *Axis-Angle*
 - Quaternion

Euler Theorem

- Any rotation in $SO(3)$ is equivalent to rotation about a fixed axis $\omega \in \mathbb{R}^3$ through a positive angle θ
- $\hat{\omega}$: unit vector of rotation axis ($\|\hat{\omega}\| = 1$)
- θ : angle of rotation
- $R \in SO(3) := Rot(\hat{\omega}, \theta)$

Euler Theorem

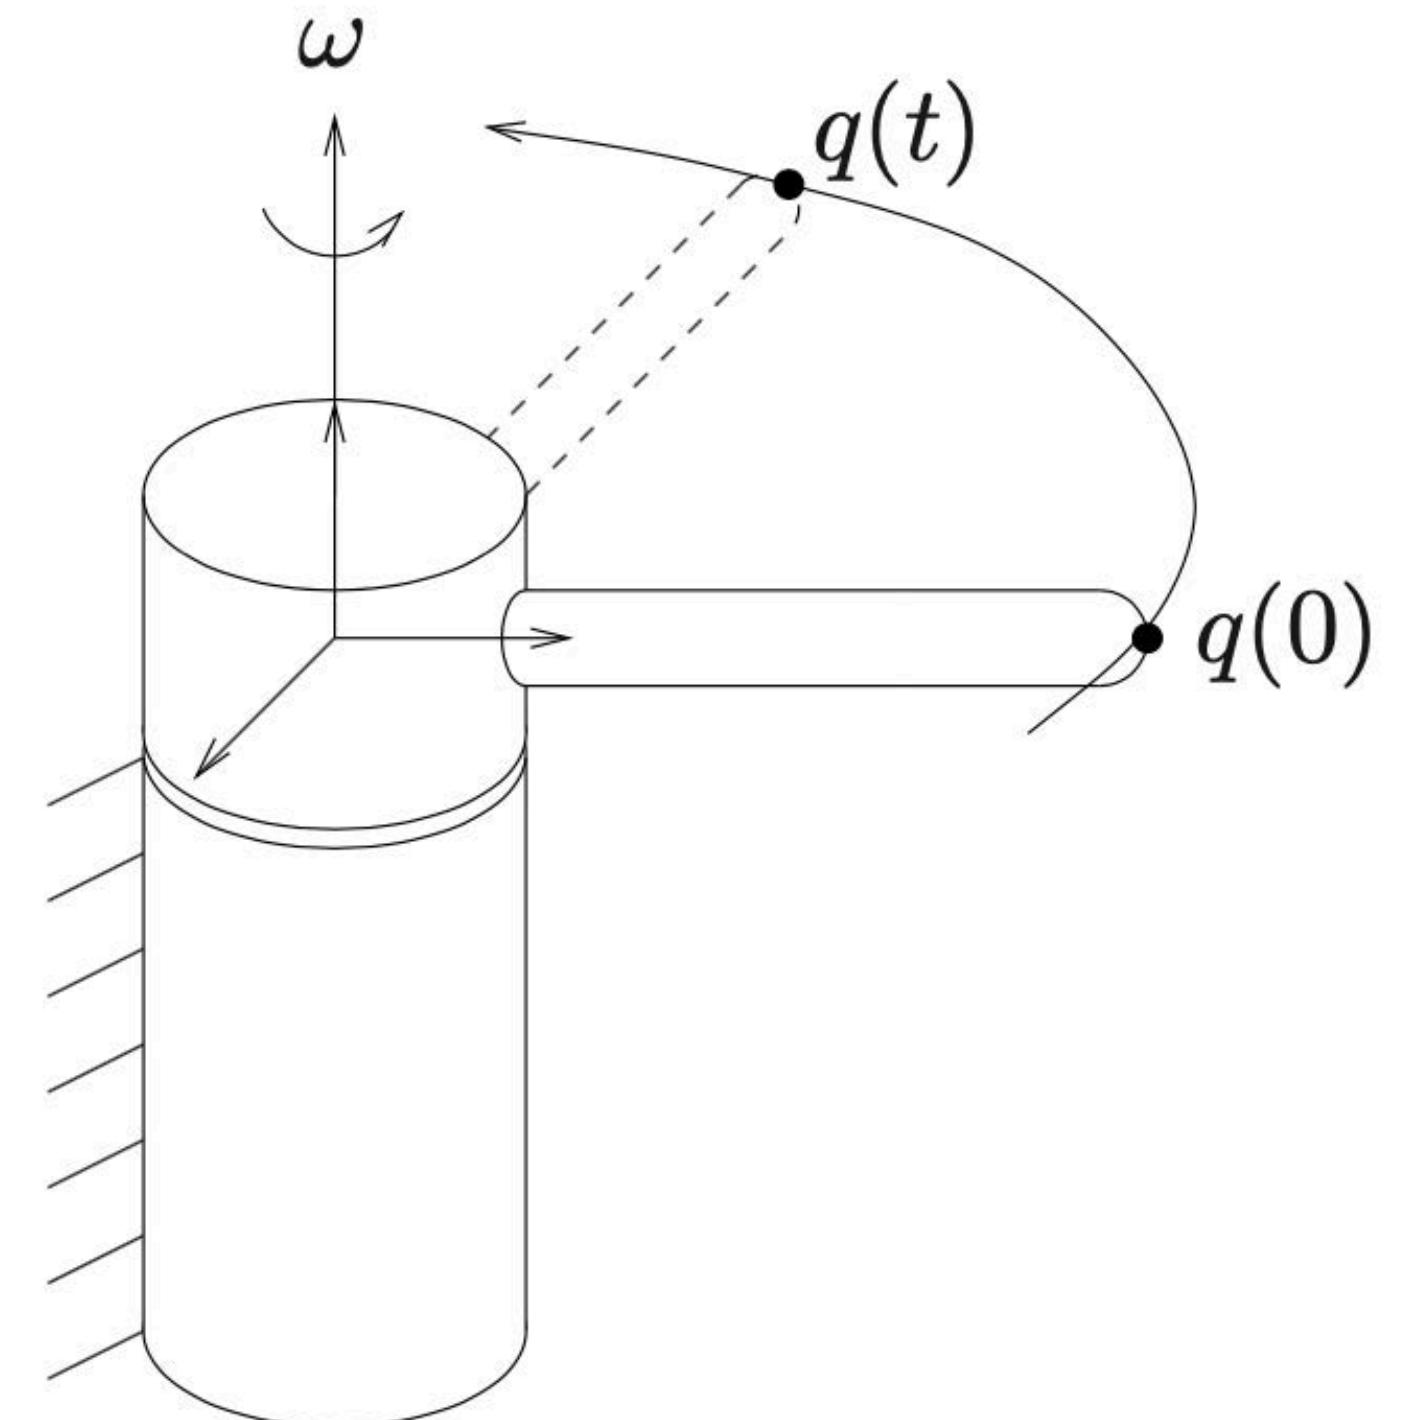
- Any rotation in $SO(3)$ is equivalent to rotation about a fixed axis $\omega \in \mathbb{R}^3$ through a positive angle θ
- $\hat{\omega}$: unit vector of rotation axis ($\|\hat{\omega}\| = 1$)
- θ : angle of rotation
- $R \in SO(3) := Rot(\hat{\omega}, \theta)$

Given $\hat{\omega}$ and θ , what is $R \in SO(3)$?

Given $\hat{\omega}$ and θ , what is $R \in SO(3)$?

- Consider a point q . At time $t = 0$, the position is q_0

- Rotate q with **unit** angular velocity around axis $\hat{\omega}$, i.e.,
 - $v = \hat{\omega} \times r$
 - $\dot{q}(t) = \hat{\omega} \times q(t) = [\hat{\omega}]q(t)$



Skew-Symmetric Matrix

- A is skew-symmetric $A = -A^T$

- Skew-symmetric matrix operator:

$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \\ \omega_3 \end{bmatrix}, \quad [\omega] := \begin{bmatrix} 0 & -\omega_3 & \omega_2 \\ \omega_3 & 0 & -\omega_1 \\ -\omega_2 & \omega_1 & 0 \end{bmatrix}$$

- Cross product can be a linear transformation
 - $a \times b = [a]b$

Given $\hat{\omega}$ and θ , what is $R \in SO(3)$?

$$\begin{aligned}\dot{q}(t) &= \hat{\omega} \times q(t) = [\hat{\omega}]q(t) \\ \Rightarrow q(t) &= e^{[\hat{\omega}]t}q_0 \text{ (solution of the ODE)}\end{aligned}$$

$$\begin{aligned}\|\hat{\omega}\| &= 1 \\ \Rightarrow \text{the swept angle } \theta &= \|\hat{\omega}t\| = t \\ \Rightarrow q(\theta) &= e^{[\hat{\omega}]\theta}q_0 \\ \Rightarrow Rot(\hat{\omega}, \theta) &= e^{[\hat{\omega}]\theta} = e^{[\hat{\omega}\theta]} \text{ (exponential map)}\end{aligned}$$

- $\vec{\omega} = \hat{\omega}\theta$ is also called **rotation vector** or **exponential coordinate**

Given $\hat{\omega}$ and θ , what is $R \in SO(3)$?

- Definition of Matrix Exponential:

$$e^{[\hat{\omega}]\theta} = I + \theta[\hat{\omega}] + \frac{\theta^2}{2!}[\hat{\omega}]^2 + \frac{\theta^3}{3!}[\hat{\omega}]^3 + \dots$$

- Sum of infinite series? **Rodrigues Formula**
 - Can prove that $[\hat{\omega}]^3 = -[\hat{\omega}]$
 - Then, use Taylor expansion of **sin** and **cos**
 - $e^{[\hat{\omega}]\theta} = I + [\hat{\omega}]\sin\theta + [\hat{\omega}]^2(1 - \cos\theta)$

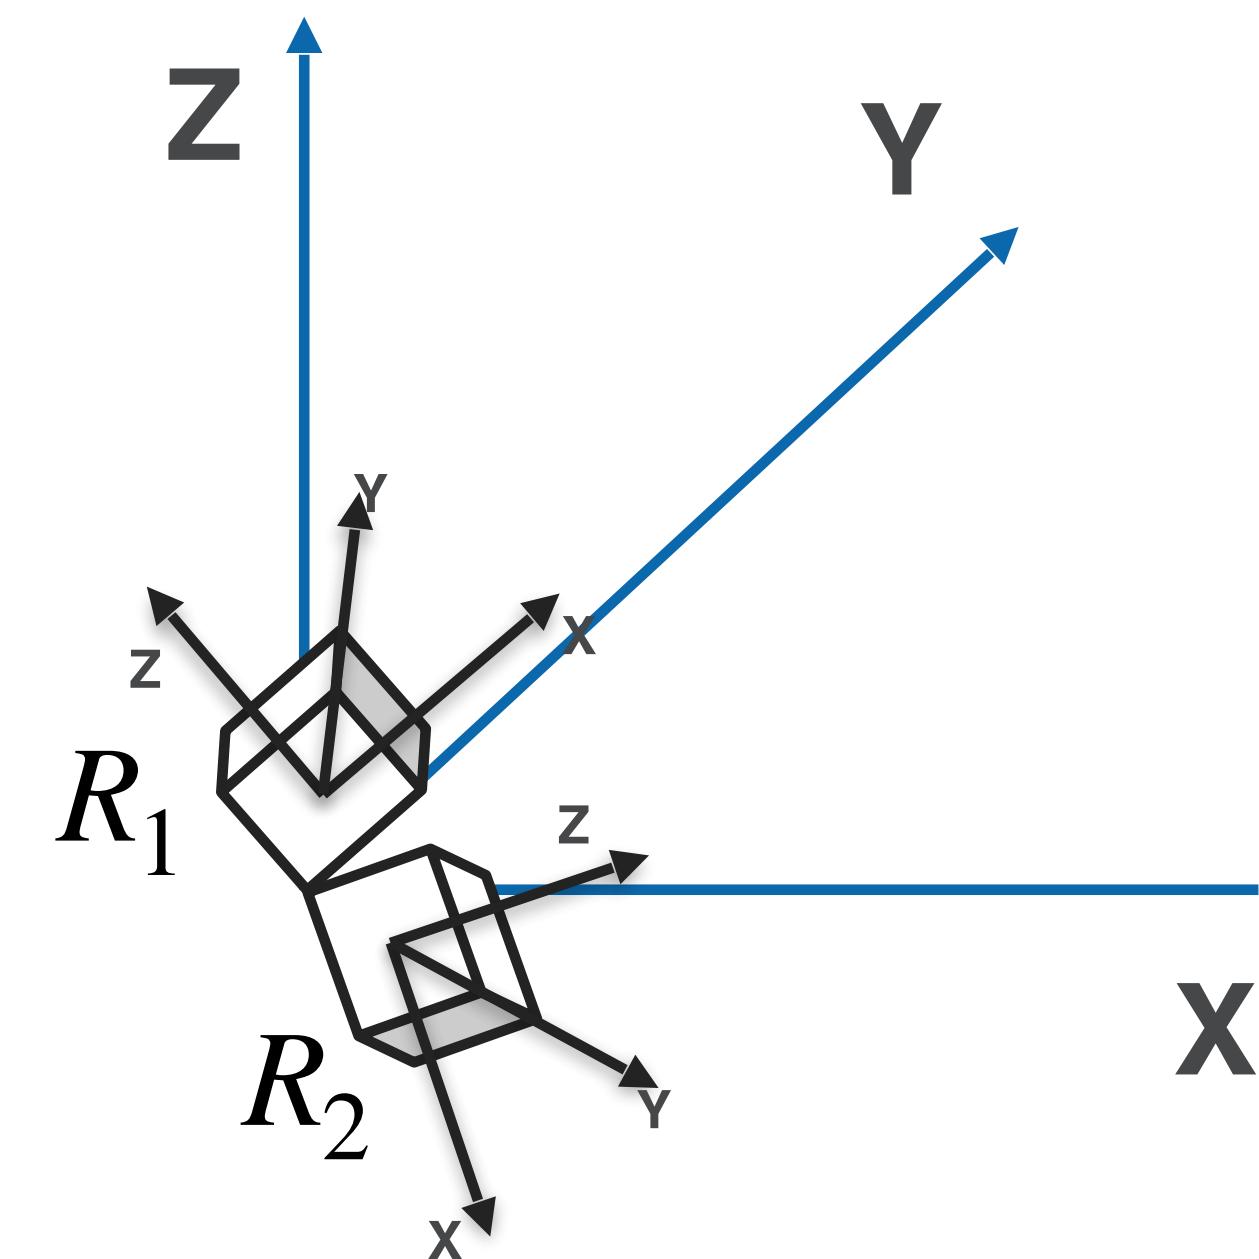
Given $\hat{\omega}$ and θ , what is $R \in SO(3)$?

- First question: Is there a **unique** parametrization?
 - No:
 1. $(\hat{\omega}, \theta)$ and $(-\hat{\omega}, -\theta)$ give the same rotation
 2. when $R = I$, $\theta = 0$ and $\hat{\omega}$ can be arbitrary
- When 2 does not happen, and if we also restrict $\theta \in (0, \pi]$, a unique parameterization exists:
 - when $\text{tr}(R) \neq -1$, can be computed by
$$\theta = \arccos \frac{1}{2}[\text{tr}(R) - 1], \quad [\hat{\omega}] = \frac{1}{2 \sin \theta}(R - R^T)$$
 - when $\text{tr}(R) = -1$, they are the cases that $\theta = \pi$ for rotations around x/y/z axis

Distance between Rotations

- How to measure the distance between rotations (R_1, R_2) ?
- A natural view is to measure the (minimal) effort to rotate the body at R_1 pose to R_2 pose:

$$\therefore (R_2 R_1^T) R_1 = R_2 \quad \therefore \text{dist}(R_1, R_2) = \theta(R_2 R_1^T) = \arccos \frac{1}{2} [\text{tr}(R_2 R_1^T) - 1]$$



Inspection from Learning Perspective

- When used in networks, one prominent issue is:
 - Suppose that you are estimating $\theta\hat{\omega}$ as a 3D vector
 - To keep a unique parameterization, you assume that $\theta \in (0, \pi]$
 - Your current solution is $\pi\hat{\omega}$
 - $(\pi - \epsilon)(-\hat{\omega})$ is mapped to a neighborhood point in $SO(3)$, but it is not in the neighborhood of the domain, hence gradient descent could not achieve it

Summary of Axis-Angle

- Axis-Angle is an intuitive rotation representation
- By adding a constraint to the domain of θ , the parameterization can be unique at most points
- Can be converted to and from rotation matrices by exponential map and its inverse (when possible)
- Induced a distance between rotations which is a metric in $\text{SO}(3)$ (independent of parameterization)

Outline

- Transformations and homogeneous coordinates
- Rotation and $SO(n)$
- 3D Rotation representation
 - Euler Angles
 - Axis-Angle
 - *Quaternion*

Mathematical Definition

- Recall the complex number $a + bi$
- Quaternion is a more generalized complex number:

$$q = w + xi + yj + zk$$

- w is the real part and $\vec{v} = (x, y, z)$ is the imaginary part

- Imaginary: $i^2 = j^2 = k^2 = ijk = -1$

- anti-commutative :

$$ij = k = -ji, \quad jk = i = -kj, \quad ki = j = -ik$$

Properties of General Quaternions

- In vector-form, the product of two quaternions:
For $q_1 = (w_1, \vec{v}_1)$ and $q_2 = (w_2, \vec{v}_2)$

$$q_1 q_2 = (w_1 w_2 - \vec{v}_1^T \vec{v}_2, w_1 \vec{v}_2 + w_2 \vec{v}_1 + \vec{v}_1 \times \vec{v}_2)$$

- Conjugate: $q^* = (w, -\vec{v})$
- Norm: $\|q\|^2 = w^2 + \vec{v}^T \vec{v} = q q^* = q^* q$
- Inverse: $q^{-1} := \frac{q^*}{\|q\|^2}$

Unit Quaternion as Rotation

- A **unit quaternion** $\|q\| = 1$ can represent a rotation
 - Four numbers plus one constraint $\rightarrow 3$ DoF
- Geometrically, the shell of a 4D sphere

Unit Quaternion as Rotation

- Rotate a vector \vec{x} by quaternion q :
 1. Augment \vec{x} to $x = (0, \vec{x})$
 2. $x' = qxq^{-1}$
- Compose rotations by quaternion:
 - $(q_2(q_1xq_1^*)q_2^*)$: first rotate by q_1 and then by q_2
 - Since $(q_2(q_1xq_1^*)q_2^*) = (q_2q_1)x(q_1^*q_2^*)$, we conclude that **composing rotations is as simple as multiplying quaternions!**

Conversation between Quaternions and Angle-Axis

- Exponential coordinate → Quaternion:

$$q = [\cos(\theta/2), \sin(\theta/2)\hat{\omega}]$$

Quaternion is very close to angle-axis representation!

- Exponential coordinate ← Quaternion:

$$\theta = 2 \arccos(w), \quad \hat{\omega} = \begin{cases} \frac{1}{\sin(\theta/2)} \vec{\nu} & \theta \neq 0 \\ 0 & \theta = 0 \end{cases}$$

Conversation between Quaternion and Rotation Matrix

- Rotation \leftarrow Quaternion

$$R(q) = E(q)G(q)^T$$

where $E(q) = [-\vec{v}, wI + [\vec{v}]]$ and
 $G(q) = [-\vec{v}, wI - [\vec{v}]]$

- Rotation \rightarrow Quaternion
 - Rotation \rightarrow Angle-Axis \rightarrow Quaternion

Inspection from Learning Perspective

- Each rotation corresponds to two quaternions (“double-covering”)
- Need to normalize to unit length in networks. This normalization may cause big/small gradients in practice

More about Quaternion

- Quaternion is computationally cheap:
 - Internal representation of Physical Engine and Robot
 - Pay attention to convention (w, x, y, z) or (x, y, z, w)?
 - (w, x, y, z): SAPIEN, transforms3d, Eigen, blender, MuJoCo, V-Rep
 - (x, y, z, w): ROS, PhysX, PyBullet

Summary of Quaternion

- Very useful and popular in practice
- 4D parameterization, compact and efficient to compute
- Good numerical properties in general

Summary of Rotation Representations

	Inverse?	Composing?	Any local movement in $\text{SO}(3)$ can be achieved by local movement in the domain?
Rotation Matrix	✓	✓	N/A
Euler Angle	Complicated	Complicated	No
Angle-axis	✓	Complicated	?
Skew-symmetrical Matrix	✓	Complicated	?
Quaternion	✓	✓	✓

? means no singularity with single exceptions

Useful Resources

- On the Continuity of Rotation Representations in Neural Networks
- A useful torch library that you can play with is “kornia”
- Use with cautious to its numerical properties
- “ceres” is a C++ library that is quite useful

Next Time: Learning-based SFM and SLAM

